


THE CURRENT BRIEFING OF

MERIDIAN

Twelve desks. One world.

TUESDAY 29 SEPTEMBER 2026



HE STAYED IN NEW YORK. THREE DAYS AFTER THE PLAN WAS REFUSED, THE MEDIATORS WENT BACK TO WORK ON IT (P3)

THE LONG READ

Two hundred and ninety-four accounts, built in May, and not one has said anything political yet

THE ATELIER

Five Italian mill names at one address near Como, and the Swiss one is only a name

THE MACRO DESK

The long end did not lead the selloff. The two-year note did, by eleven basis points

THIS ISSUE · NO. 117

Contents

- p3–4** The World — the refused plan goes back to the mediators; Bangkok’s toll turns out to exclude Bangkok

p5 Singapore — week five of the haze, and a trigger nobody has reached

p6 The Atelier — the dress shirt: count the stitches, pinch the collar

p7 Technology — the second chip company in twenty-five days to buy the layer above it

p8 The Home Lab — forty-two times faster, and one approving review short

p9 The Kit — on sale today, with its headline feature in early access

p10 The Good Life — a ceramic pebble from 1988, on a borrowed movement

p11 Screen & Sound — the number-one film is seven years old and selling a
- trailer

p12–13 The Family Desk — which app it happened on, and why eleven is the worst age

p14 The Macro Desk — the curve flattened, which is a statement about the Fed

p15–17 The Rabbit Hole — orienteering: the map is the opponent, and the flagship shoe is a consumable

p18 Fitness — ten thousand is the wrong variable

p19–20 The Diary — ten days to the race, and ten seats at the noodle bar

p21–23 The Long Read — the brittle place

p24 The Last Word — loaded, and not yet fired

FROM THE DESK

“If in the first act you have hung a pistol on the wall, then in the following one it should be fired.”

Anton Chekhov, as recalled by Ilia Gurliand, 1904 — for an issue about things that are finished, and waiting

Two hundred and ninety-four accounts have been built and have said nothing. A forty-two-fold speed-up sits in a personal fork. A flagship shipped this morning with its best feature unfinished. And a plan that was publicly refused is being redrafted by the people who carried it.

THE STRIP · AT MONDAY’S CLOSE

US 2-year	4.92%	▲ 11bp
US 10-year	5.24%	▲ 7bp
US 30-year	5.56%	▲ 7bp
VWRA	\$192.30	▼ 0.63%
AGGU	\$5.71	▼ 0.35%
DXY	101.19	▲ 0.22%
USD/SGD	1.2772	▼ 0.03%
Brent	\$105.28	▲ 0.9%
Bitcoin	\$83,588	▼ 1.26%

Close of Monday 28 September. Yields are the Treasury’s own par curve, read cell by cell against its header row. Gold fell through \$4,200 to a seven-week low; the exact close is held on **p14**. Not advice.¹

THE ARCHIVE · WHAT NO. 117 IS STANDING ON

The number we would not settle. No. 116 set Bangkok’s governor against the press on the flood toll and declined to resolve it. The gap turns out to be a counting boundary (**p4**).

The tenor, again. No. 115 printed the seven-year yield as the ten-year and No. 116 corrected it. Reading the table a third time finds the consensus story about Monday the wrong way round (**p14**).

SOURCES 1 Levels at Monday 28 September’s close, each sourced on **p14**. Yields are the US Treasury’s own daily par curve, read against its verbatim header row. VWRA’s level was cross-checked against the fund’s sterling line, its quote page having mislabelled the currency.

THE MERIDIAN DISPATCH • NEW YORK

Three days after he refused it, the same seven days are being redrafted

Qatar's mediators were to sit separately with each side in New York on Monday or Tuesday, on an amended version of the plan Washington turned down on Saturday. Neither capital has withdrawn anything.

The rejection was public, immediate, and by Monday beside the point. On Saturday, on the White House driveway, the American president refused Iran's offer to reopen the Strait of Hormuz within seven days: "I rejected their deal. They want to make a deal where they open the strait immediately because they're losing so badly." By Monday, officials on both sides confirmed that Qatar's mediators were still working, and would sit separately with Iran's foreign minister in New York and with the American side within the day — on an amended version of the same seven-day plan.¹

The terms have not moved: the blockade of Iranian ports lifted, sanctions on oil sales waived, roughly **twelve billion dollars** of frozen assets released, a ceasefire wide enough to take in Lebanon and Yemen. What has moved is one word. The running order this desk printed on Sunday is now a draft with edits on it, and **Abbas Araghchi** has stayed in New York to receive them.²

\$4.49

THE AMERICAN PUMP PRICE ON 26 SEPTEMBER, AGAINST \$4.10 A MONTH EARLIER. THE PRESIDENT CALLED IT "A LITTLE GLITCH." BRITISH DIESEL SET AN ALL-TIME RECORD OF 198.32P A LITRE, WHICH THE RAC BLAMED ON THE WAR, AND AN AUGUST SURVEY FOUND **81 PER CENT** OF AMERICANS SAYING THE WAR HAD RAISED THEIR COST OF LIVING³

A procedural oddity sits underneath. On Saturday Araghchi said the refusal had reached him through a television set rather than a mediator: "nothing has been conveyed to us." It had been broadcast; it had not been delivered. On Sunday, on American television, he took both registers in one breath — "even if it comes to a doomsday war," then, without a pause, "we stand ready for diplomacy. It is up to President Trump to choose."⁴

DOHA — AL JAZEERA

Reads the offer as dead and the refusal as a choice. Its Washington correspondent **Mike Hanna** called the proposal "dead in the water." **Mohamad Elmasry** of the Doha Institute gave the transactional reading — it was "kind of throwing a bone to Trump" — and warned Tehran may make him pay for refusing.⁵

31 / 69

THE SHARE OF AMERICANS APPROVING AND DISAPPROVING OF THE PRESIDENT'S HANDLING OF THE WAR, IN NBC'S POLLING. ON THE COST OF LIVING IT IS 36 AGAINST 62⁷

BEIJING — MINISTRY OF FOREIGN AFFAIRS

Declines the frame. Asked on **24 September** about Iran refusing free navigation while the blockade stands, spokesman **Guo Jiakun** named no party at fault: "Protraction of the U.S.-Iran conflict does not serve the common interests of the parties concerned or the international community." No Chinese state commentary on the rejection itself appeared between 26 and 29 September.⁶

THE SCEPTIC • PAIN IS NOT THE SAME AS CONSENT

Wait for better terms, runs the consensus case, because pressure compounds. Two Iran specialists say that reads the wrong variable. **Andreas Krieg** of King's College London grants the premise and refuses the conclusion: the danger "is that he is confusing Iranian pain with Iranian willingness to capitulate. Those are not the same thing." **Negar Mortazavi** of the Center for International Policy: the economy is under enormous pressure, but Tehran "will not necessarily accept demands it considers fundamental red lines."⁸

One number will not agree with itself, and it is worth printing as a warning. Fox News has transits about **75 per cent below** pre-conflict levels; NBC has nearly **40 million barrels** through in a single 48 hours. Both ran the same weekend, neither states its unit, and no authority publishes a reconciled figure.⁹

SOURCES 1 NBC News, 26 September 2026, for the rejection and the quoted remarks; the Monday state of play from the Associated Press wire of 28 September as carried by three outlets. 2 NBC News and Al Jazeera, 26 September, for the four conditions and the twelve-billion-dollar figure. 3 Fox News, 26 September, for the pump price; ABC News for the president's remark. 4 Araghchi on Al Jazeera, 26 September, and NBC's *Meet the Press*, 27 September. 5 Al Jazeera, 27 September. 6 Chinese Ministry of Foreign Affairs press conference, 24 September, transcript read at source; Global Times carried no Iran item on 28 or 29 September. 7 RAC via the BBC, 28 September; NBC News; Institute for Global Affairs, fielded 11–14 August, n=1,000. 8 Al Jazeera, 27 September. 9 Fox News and NBC News, both 26 September; neither states its unit, and this desk prints both rather than choosing.

AROUND THE DESKS · BANGKOK

The national toll is eight, and none of the eight died in Bangkok

No. 116 printed the governor's count against the press count and refused to settle it. The gap has a shape: it is a boundary, not a contradiction, and the capital's three dead sit outside it.

Waters began easing across the Thai capital on Monday, four days after every district there was declared a disaster zone. The government's national count stands at **eight dead** since mid-September, all in three eastern provinces and none in Bangkok. The capital's own casualties are itemised separately, and there are three: a woman of 58 electrocuted in a Khlong Chan car park; a man of 73 who drowned in Bueng Kum after slipping while moving his belongings; and a bedridden man of 78, in the same district, who could not be evacuated.¹

So the discrepancy this desk declined to settle last issue was a counting boundary. Neither figure was wrong; neither was complete.¹

1.7m

PEOPLE AFFECTED ACROSS 27 THAI PROVINCES, IN MORE THAN 500,000 HOUSEHOLDS. IN BANGKOK ALONE SOME 329,000 HOUSEHOLDS — ABOUT 700,000 PEOPLE — WERE HIT OVER THE WEEKEND, AND MORE THAN 80,000 WERE STILL UNABLE TO LEAVE THEIR HOMES ON MONDAY²

Governor **Chadchart Sittipunt** said on Sunday that full drainage could take “up to a week.” The government spokesman **Ekkapob Pianpises** confirmed the recession while noting that “some areas in Bangkok are still facing high water levels.” The cabinet declared 28 and 29 September special holidays in four provinces. Across the border in Myanmar's southern Tanintharyi region at least eight more died and as many as thirty are missing; three of the recovered bodies were teenagers. Taken together the two countries have lost at least sixteen people and seen roughly **2.5 million** affected.³

THE CHATTER · WHAT RESIDENTS SAID THEY WERE TOLD

The complaint is not about the water. It is about the notice. **Yanika Luangcharoen**, 44: “We weren't prepared at all. We only knew it was going to rain, but there was no warning whatsoever.” A flat manager, **Suranee Panoh**, on Monday: “We can't go outside. The water level is so high — it would come up to neck level for tall people.”³

THE SCEPTIC · A TOLL THAT EXCLUDES ITS LARGEST CITY

A national figure of eight, while the capital separately records three, is not a lie; it is a number that will be quoted without its footnote. Deaths are attributed to declared disaster provinces, and a metropolitan administration reports its own. The honest form of the sentence is at least eleven dead in Thailand and at least eight in Myanmar — and the Myanmar figure is the one likely to move, because thirty people are still missing there.¹

ELSEWHERE · IN BRIEF

Ebola past 8,000. Cases in the Democratic Republic of Congo topped eight thousand and spread into two further regions. **Tedros Adhanom Ghebreyesus** of the WHO noted that vaccines and treatments exist but that, without incentives, people avoid medical centres altogether. A separate vessel capsizing in the DRC killed at least twelve.⁴

Taiz. Seven killed and forty wounded, children among them, in airstrikes on the Yemeni province, per the Houthi-run health ministry; government forces said they had struck Houthi positions. Nearly seven hundred have died in Yemen since fighting escalated last month.⁴

Fairford. Five British men in their twenties were arrested early Sunday on suspicion of preparing acts of terrorism after approaching the Gloucestershire airbase from which American bombers have flown Iran strikes since 1 March. Bomb-disposal robots searched three vans; dozens of residents were evacuated. *Single-sourced through news aggregators — no police or ministry statement was obtainable, and this desk could not reach a British outlet directly.*⁵

Washington. The US ambassador to China said the president had asked Xi Jinping about buying American weapons. The State Department then said flatly that “U.S. law prohibits arms sales to China and there is no offer or plan to sell arms to China.”⁴

SOURCES 1 France24/AFP, 28 September 2026, for the national count of eight and the three itemised Bangkok deaths; Watchers.news for the combined total. The separation of national and metropolitan counts is this desk's reading: no Thai authority publishes a reconciled figure. 2 The same two sources for the household and provincial figures. 3 France24/AFP, carrying Chadchart, Ekkapob and both residents, and for Tanintharyi and the four-province holiday. 4 Havana Times international briefs, 28 September, an aggregator digest; each item single-sourced. 5 Havana Times and World News Group, plus a Euronews bulletin attributing the BBC. No primary police or ministry statement was reachable; “explosives” appears in one of three and is not printed.

THE LITTLE RED DOT • THE HAZE, WEEK FIVE

The plan announced on 4 September has a trigger at two hundred. Singapore has not been near it

The only thing cancelled so far was cancelled by a private organiser, at three in the morning, two hours before the gun.

At two o'clock on Sunday morning the twenty-four-hour PSI read 126 in the Central region. At three, Decathlon and Metasport shortened the Kiprun half-marathon from 21.1 kilometres to ten. The gun was at five. The shorter races ran unchanged; the awards ceremony was scrapped and every category made non-competitive, so there are no results. Entries had cost **S\$78 to S\$98**, and two entrants told a local outlet they were offered neither refund nor compensation.¹

By six on Monday evening the agency had the twenty-four-hour PSI between 68 and 94, with the one-hour fine-particle band “elevated”, from 64 to **138 micrograms per cubic metre**. The forecast for today is “high moderate to mid unhealthy” despite expected showers — and cloud cover is preventing a full satellite assessment of how far the haze extends.²

101

THE PSI AT WHICH AIR IS CLASSED UNHEALTHY, AND THE LEVEL SINGAPORE HAS SPENT MUCH OF FIVE WEEKS ABOVE. THE COMMUNITY CENTRES OPEN AS HAZE REFUGES ONLY ABOVE **200** — A FIGURE THIS EPISODE HAS NOT APPROACHED³

That gap is where the desk's attention belongs. The multi-agency plan published on 4 September is real and specific: purifiers in every primary and secondary classroom, enclosed exam rooms when the air is Unhealthy, purifiers and coolers in hospitals and nursing homes, preschools minimising outdoor activity, employers advised to reduce strenuous outdoor work. But school operations and examination schedules are explicitly unchanged, and the one public refuge opens only above twice the current reading.³

THE CHATTER • THE ARGUMENT IS ABOUT THE REFUND

Runners were broadly with the organisers on the call itself — “the responsible thing to do,” one said. The anger went to the money: comment threads filled with questions about refunds, and the two entrants who spoke on the record had been offered nothing. The wider texture is a city checking an app again for the first time since 2019.¹

THE SCEPTIC • IS 101 A TRIGGER, OR A NOTIFICATION?

Five weeks in, the apparatus announced on 4 September has produced almost no visible public intervention; everything cancelled has been cancelled by private organisers. Two questions follow. Is Unhealthy a threshold that does something, or a number that is published? And an air-purifier strategy is an indoor strategy: it protects children in classrooms and patients in wards, and does nothing for the outdoor workforce, where the language is “advise” rather than require. The Institute of International Affairs warned in June of a high risk of a severe event this year. It was forecast; the response is still reactive.⁴

Elsewhere. Taxi and private-hire drivers were given a formal route into salaried work on Monday: a pathway to bus captain, announced by Senior Minister of State for Transport **Sun Xueling** and built with the Land Transport Authority and the unions. It runs two to six months and arrives with enhanced starting salaries from **1 January 2027**. It is the first state-brokered off-ramp out of platform driving into a unionised job, with a dated pay rise attached.⁵

IN BRIEF

Six charged over unpermitted assemblies, including two well-known actors, a journalist and three activists, over gatherings outside the State Courts in February 2025 and at Changi Prison in November 2024. No trial date reported.⁶

Still awaited. SP Group's fourth-quarter tariff was unpublished when this issue closed; the last four years put it on the 29th or 30th (p14).⁷

SOURCES 1 Mothership, 27–28 September 2026, for the truncation, the timings, the readings, the entry prices, the refund position and the reaction. The Singapore broadsheets all refuse this desk, so reporting here rests on Mothership, AsiaOne and the agencies' own pages. 2 National Environment Agency haze update, 28 September, and haze.gov.sg at 5.10pm. 3 NEA's multi-agency release of 4 September, for the thresholds and every measure. 4 This desk's reading of those measures against five weeks of readings; the Institute of International Affairs' Haze Outlook 2026, June. 5 Mothership, 28 September, on the bus-captain pathway. 6 Mothership, 25–28 September, each in-brief item read at source. 7 SP Group's media index, checked twice; MAS's advance release calendar.

BUY IT ONCE · THE DRESS SHIRT

Five Italian mill names, one address in Guanzate

The cloth label is the least reliable thing on a shirt. The two tests that work cost nothing: count the stitches on the cuff, and pinch the collar.

Start with the number everyone misreads. On a shirt, **120s** is not threads per square inch — it is the fineness of the yarn, and above about 100 it implies two yarns twisted together. So a “200s” single-ply is not better than a 120/2; it is thinner and weaker. **Carl Goldberg**, who has cut shirts in New York for decades, buys 100/2 and 120/2, calls 140/2 “nice but not as durable,” and is blunt about the fashionable single-ply compact yarns: “thin and slimey.” That is the ceiling, from a man with a reason to sell you dearer.¹

Now the label. **Thomas Mason** was founded in England in 1796 and has been Italian-owned since 1992 — an English name that has not been an English mill for thirty-four years. Its owner’s production-sites page lists eight facilities: two weave and dye in **Egypt**, a third weaves in the **Czech Republic**, and the spinning is in **Hungary**.²

THE WHITE LABEL · THIRTEEN HOUSES, ONE COMPANY

The sharper case is quieter. **C. Tessile S.p.A.** of Guanzate sells shirting under Canclini 1925, **Grandi & Rubinelli** (acquired July 2019), **Hausammann + Moos 1811** — “a historical Swiss company” by its own description, which joined the group in **2015** — and seven more names. None says where the cloth is woven. Goldberg is unsentimental: “Hausman & moos. Is just a name. Nothing woven in Switzerland.”³

Which is why the hand beats the label. **Count the stitches**: eighteen an inch is the floor, Ascot Chang states twenty-two. **Look at the side seam** — one line of stitching outside means single-needle, sewn twice and pressed flat; two means the fast way. **Pinch the collar leaf**: a fused collar moves as one board, an unfused one shifts.⁴

Note where the floor has moved: the cheapest shirt Kamakura lists is **US\$135**, the seventy-nine-dollar shirt it was famous for is gone, and its extra-long-staple tier is **US\$190**.⁵ The quality entry runs

US\$125–200 — Proper Cloth from \$125; Edit Suits from **S\$139**; CYC from **S\$200** — and buys two-ply cloth, single-needle seams, a split yoke and shell buttons. Real custom starts at Ascot Chang’s **US\$375**, and what you buy is the paper pattern, kept for **ten years**.⁶

The tropical question the literature never answers — it is written for English summers — is what works at thirty degrees and eighty per cent humidity. Evaporation is already throttled, so wicking has nowhere to send the water; the job is airflow and not sticking. **Voile** is the mechanism: high-twist yarns bulked by creping, so the cloth is woven with fewer yarns per inch and stays soft. The catch is opacity, which is why the shirtmaker **Thomas Auer** steers clients to twill: most “resent the idea of slight transparency.”⁷

THE SCEPTIC · WHERE THE MONEY BUYS A WORD

“Sea Island” is not a protected term, and forum members have watched it applied to any fine cotton for years. Charles Tyrwhitt lists non-iron shirts at \$129 and sells them at **\$74.75**, which makes the list price scenery. On Thomas Pink, **rathebaindc** puts it plainly — at “\$150–\$200, you’re better off with a quality MTM.” The same split runs at the top, where long-time buyers rate Charvet’s bespoke and dismiss its ready-to-wear as trading on the name. Buy the commission, not the label.⁸

NEXT ON THE BENCH — YOU PICK TOMORROW’S CATEGORY

Hand tools & screwdrivers — hardened tips, Wera against PB Swiss, and who really grinds them.

Mattresses & slat bases — latex against pocket springs, and what humidity does to both.

Enamelware & camp mugs — vitreous enamel over steel, and the three makers still firing it.

SOURCES 1 Proper Cloth’s shirting reference for the yarn-count and ply definitions; Carl Goldberg, as Shirtmaven, on Styleforum. 2 Albini Group’s own Thomas Mason and production-sites pages, read 29 September 2026. 3 C. Tessile S.p.A.’s own brand and company pages — the 2015 and July 2019 dates are the company’s. Goldberg’s claim that the group owns no weaving is his; the corporate pages state no weaving location at all. 4 Proper Cloth on stitch density and seams; Ascot Chang’s shirt page for twenty-two stitches an inch and its use of both fused and unfused interlinings — fused is a style choice, not a corner cut. 5 Kamakura’s live list, read 29 September 2026. 6 Live price lists for Proper Cloth, Edit Suits, CYC and Kevin Seah; Ascot Chang’s FAQ for the ten-year pattern retention, though it publishes no custom prices: the US\$375–650 range is from dated first-hand accounts. 7 Kabbaz and Auer, Styleforum; Singapore Customs and IRAS. 8 Styleforum on Sea Island cotton, Thomas Pink and Charvet; Charles Tyrwhitt’s own US page. Permanent Style refuses this desk.

SILICON & SOVEREIGNTY · SANTA CLARA

The second chip company in twenty-five days to buy the layer above it

AMD is paying about eight point two billion dollars in its own stock for World Labs. It is signed, it is not closed, and the open-model commons now sits inside two GPU vendors' balance sheets.

Monday's agreement is definitive, all-stock, and expected to close by the end of the year subject to regulatory approval — which is to say it is a promise with a condition attached. AMD will acquire **World Labs**, the spatial-intelligence lab **Fei-Fei Li** founded two years ago; she becomes AMD's executive vice-president and chief scientist, reporting to **Lisa Su**. **Justin Johnson** and **Ben Mildenhall** — a co-author of the original neural-radiance-fields paper — continue to lead the team.¹

What AMD says it is buying, in its own words, is “spatial-intelligence models that generate, reconstruct and simulate interactive 3D environments from text, image and video inputs, as well as technology for robotic learning and simulation.” Su's stated reason is narrower and more telling: “Building the compute platforms for the next generation of AI requires a deep understanding of how models are evolving.” The chip designer is buying a telescope pointed at its own roadmap.¹

25 days

BETWEEN NVIDIA'S US\$12.93BN AGREEMENT TO BUY HUGGING FACE, DISCLOSED TO THE SEC ON 3 SEPTEMBER, AND THIS ONE. BOTH ARE SIGNED; NEITHER HAS CLOSED²

World Labs' own announcement commits both parties to “an end-to-end open AI ecosystem spanning hardware, software, platforms, and widely accessible open models,” and says the decision grew out of an existing GPU-optimisation partnership — the stated motive being to get closer to the hardware. The two shipping products, **Marble** and **Atlas**, are named in that post and given no roadmap in either. AMD's press release does not mention them at all.³

THE CHATTER · WHAT A NAME IS WORTH

On Hacker News the first question was the obvious one. **AbuAssar**: “Is a 2-year-old company worth 8 billion USD?” **avaer** answered it without flattery — “a name like Fei-Fei Li is probably worth a chunk of that” — and argued these are “public bets, not valuations based on product economics.” The same commenter gave the sharpest reading of the upside and the risk in one line: the deal could “jumpstart competition with NV on a model/simulation ecosystem,” but it “could also mean the death of the interesting things WL was doing with Marble + Atlas.” And then, honestly: “it's hard to read PR.”⁴

THE SCEPTIC · WHAT THE DISCLOSURE DOES NOT CARRY

AMD is paying in its own shares for a company with no disclosed revenue, no disclosed headcount and no published product commitment, and the “open ecosystem” language lives in a blog post rather than in anything filed. Two tests will settle it. If Marble and Atlas are sunset or fenced behind ROCm within two quarters of closing, the openness was marketing. If regulators — who have Nvidia and Hugging Face in front of them at the same moment — issue a second request, the year-end close slips. What would confirm the bull case is duller and more convincing: an AMD roadmap disclosure naming spatial and world-model workloads as an MI-series design target.⁵

ALSO THIS WEEK

Meta takes a chief executive off a public company.

Chirantan “CJ” Desai stepped down as MongoDB's chief executive after less than a year to become Meta's chief enterprise platform officer, reporting to Mark Zuckerberg and running a new enterprise arm built around the Muse agent. MongoDB fell about a quarter on the news; Meta about four per cent.⁶

Stolen model access is now a commodity. A dark-web vendor resells frontier-model access at five to fifteen per cent of list by farming free cloud credits, operating as a proxy that can read and alter every prompt passing through it. Google's threat-intelligence group reports a major increase this year. The argument for running your own weights has rarely been cleaner (p8).⁷

SOURCES 1 AMD's investor-relations release, 28 September 2026, for the consideration, the all-stock structure, the year-end close subject to approval, the titles and both quotations. 2 Nvidia's Form 8-K for the period 2 September 2026, for the Hugging Face consideration. 3 World Labs' own announcement. 4 Hacker News item 49883760; usernames as posted. 5 This desk's reading of what each document does and does not commit to. TechCrunch, CNBC and Bloomberg independently confirm price, structure and title; World Labs' valuation, funding and headcount were not obtainable, and a claim that this is AMD's second-largest acquisition is uncorroborated and not printed. 6 CNBC and Seeking Alpha, 28 September. 7 PYMNTS, September 2026, reporting Google Threat Intelligence.

THE HOME LAB · THE MERGE QUEUE

A forty-two-fold speed-up that exists only in a personal fork

The code is written, benchmarked and public. What it does not have is a second approving review.

On 26 September **Hayder Tirmazi** published a set of optimisations to llama.cpp's prompt-lookup speculative drafter, measured on an M4 Pro against the WikiText-103 test set. Removing three map copies in favour of const-reference reads gave **4.5 to 25.6 times** faster drafting with acceptance unchanged to within 0.15 of a percentage point. Swapping the inner map for a sorted vector added twice more. Then **Daniel Lemire** joined in, and the combined figure reached about **140×**.¹

The headline number is drafting latency falling from **165.48 microseconds per token to 1.18**. As of today none of it is upstream. A search of the main repository's pull requests by that author returns nothing at all.²

THE SCEPTIC · MICROSECONDS AGAINST MILLISECONDS

Be careful what this is. It is a speed-up on the *drafter*, not on inference. Prompt-lookup drafting is already a microsecond-scale operation next to a forward pass measured in milliseconds, so the saving is about 164 microseconds per drafted token. The question was asked plainly on Hacker News by **Soy** — “How much does this speedup inference for the end user in terms of tk/s?” — and in a thread of 83 points it was never answered. Until somebody publishes an end-to-end benchmark this is beautiful data-structure work of unquantified benefit. The genuinely useful result is the quiet one: static-cache load for the 541MB corpus fell from **3.76 seconds to 0.23**.³

Why it is not merged is a smaller, sadder story. A maintainer, tagged by accident on Tirmazi's own fork rather than upstream, replied “Blocked for making me waste my time.” Tirmazi explained and apologised; on 28 September the thread was locked to collaborators. A personal block and a locked thread are established. A repository-level ban is not, and this desk will not print one.²

1 of 2

APPROVING REVIEWS ON PULL REQUEST 29500, THE SYCL OPTIMISATION THIS DESK LOGGED AS PENDING LAST ISSUE. IT RECEIVED ITS FIRST APPROVAL ON 27 SEPTEMBER; THE REPOSITORY REQUIRES TWO. THAT IS THE WEEK'S DELTA.⁴

That request is worth reading for the shape of its numbers rather than its headline. It ports a multi-column matrix-vector optimisation to a three-bit quantisation in the Intel backend. On an Arc A770 a matmul at batch four goes from **978 to 361 microseconds**. At batch one it goes from 250 to 249, which is nowhere. End to end that is 15 to 22 tokens per second on a short prompt, and **9.4 to 10.1** at 22,000 tokens of context. The sellable number is the first. The one an Arc owner running long contexts will feel is the second.⁴

THE CHATTER · THE OTHER HELD THREAD, STILL HELD

This desk logged **Xing4.0-29B-A4B** — China Telecom's Apache-licensed mixture-of-experts model, trained entirely on Ascend silicon — as awaiting independent reproduction. It still is: one write-up puts it exactly, that the comparisons “come directly from the model card” and “the results haven't been reproduced by a neutral third party.” Note what the card concedes: on its own numbers the model loses to Qwen3.6-35B-A3B on both SWE-bench Verified and AIME, winning on the agentic benchmarks. A much narrower claim than the one circulating. On Hugging Face, **grow-nexus** notes all five inference-framework pull requests were still pending — engine support in flight while the benchmarks were quoted.⁵

One tip, flagged as a tip. A single developer reports that moving mixture-of-experts layers off the CPU onto an Intel Arc B390 integrated GPU took generation on a 35-billion-parameter model from **14.7 to 32.5 tokens per second**, against documentation implying the offload penalty was about ten per cent. His caveat is the honest part: with roughly five gigabytes free he must revert past 49,000 tokens of context. One machine, no replication.⁶

SOURCES 1 Hayder Tirmazi's write-up, 26 September 2026, read in full, for the rig, corpus, multipliers and acceptance figures; the first optimisation is independently visible as a diff with its own benchmark table. 2 The fork's pull request and the main repository's search by that author, read as HTML on 29 September. The Hacker News framing of a ban overstates the evidence, and none is printed. 3 Hacker News item 49859982. 4 Pull request 29500, read as HTML on 29 September: opened 26 September, first approving review 27 September, still open; figures are the author's own. 5 The model card's benchmark table as reported by a write-up stating the non-reproduction plainly; Hugging Face discussions on the official repository. 6 A single developer's post of 28 September. Unreplicated; printed as a tip.

GEAR WE'RE WATCHING · ON SALE TODAY

The feature you are paying the premium for shipped in early access

The Sonos Ace Ultra reaches general sale in Singapore today at S\$699. Headphone Linking, the thing that justifies the price, is compressed stereo, needs Wi-Fi, and dissolves your speaker groups.

The hardware is not in doubt. A new 40mm driver in each cup, ten microphones against the original's eight, a 1,060mAh cell rated at up to **35 hours** with cancelling on and 50 with it off, Bluetooth 6.0 with multipoint and aptX Adaptive Lossless, spatial audio with head tracking, and an eight-band equaliser where the original had two. It weighs 322 grams and costs **US\$449**, or **S\$699** here.¹

Two things here are genuinely unusual: the ear cushions are user-replaceable at **US\$39**, and so is the battery, at **US\$25**. In a category that treats a dead cell as the end of the product, that is the buy-once argument made in parts.¹

S\$699

AGAINST A SONY FLAGSHIP AT AROUND US\$460 AND A BOSE AT US\$449, BOTH OF WHICH THREE SEPARATE REVIEWERS PLACE AHEAD OF IT ON NOISE CANCELLING. THERE IS NO INGRESS RATING AT THIS PRICE²

The reviews split, and the split is the useful part. **SlashGear** recommends it — the new driver, the deeper bass, the re-profiled cushions for spectacle wearers — while noting the hundred-dollar premium “feels incremental for casual users.” **Gizmodo** gave it three and a half out of five under a headline calling them headphones “only hardcore Sonos fans could love”: sound “solid but not excellent,” cancellation “middle of the pack” despite a claimed doubling, and earcup pressure that hurt on a long flight. **SlashGear** praised the comfort on a similar flight, so the honest reading is head-shape sensitivity: try before you buy.²

THE CHATTER · WHAT SONOS ITSELF LISTS AS LIMITATIONS

Headphone Linking — the feature that separates this from every rival — is in **early access**, announced by a Sonos staffer on the company's own forum a fortnight before general sale. The constraints are Sonos's own, not a reviewer's: “lossless or Spatial Audio quality is not supported” over the link, so it is compressed stereo; it needs Wi-Fi, so it fails on the move; only one pair can receive per session; **speaker grouping dissolves** when audio links to the headphones; and voice assistants cannot drive it. **Gizmodo's** reviewer could not get handover working at all on first setup.³

THE SCEPTIC · AN ECOSYSTEM ACCESSORY PRICED AS A FLAGSHIP

Strip it down and the premium buys one feature — unfinished, compressed, Wi-Fi-bound, app-initiated, exclusive to people who already own Sonos speakers, and destructive of the grouping those speakers exist for. Reviewers agree the cancellation trails Sony, Bose and Sennheiser, and there is no water resistance. For a reader migrating a house onto open local control, the headline capability is the most closed part of the product. What would change the verdict is precise: Headphone Linking leaving early access with lossless and spatial intact. Until then, buy it for the drivers and the replaceable parts, not the trick.⁴

ALSO ON THE BENCH

Sennheiser Momentum True Wireless 5, S\$479. Seven-millimetre drivers, Bluetooth 6.0, aptX Lossless to 24-bit/96kHz, six hours with cancelling on. The reason it is here: **user-replaceable batteries in both the buds and the case** — close to unique in true wireless, and the strongest buy-once story in this issue. The local reviewer's reservation is the price, above the previous generation's launch at a moment when Sony's flagship sits well under S\$400, and he calls the head-tracked spatial audio “gimmicky.”⁵

Still on the shelf. The Sonos Beam Ultra soundbar, covered here last issue, also reaches general availability today at S\$1,299.¹

SOURCES 1 TC Acoustic, the official Singapore distributor, for the S\$699 price, the 29 September date and the specification; Gear Patrol for the US and European prices and the replaceable cushion and battery prices; HardwareZone Singapore for the local timeline and the Beam Ultra date. Sonos's own site refuses this desk. 2 Gizmodo and SlashGear, both 22 September 2026; TechRadar's page truncated on fetch and only its headline verdict is used. 3 Sonos's own community forum, the early-access thread opened by staff on 15 September; every limitation quoted is from that post. Setup difficulties from Gizmodo. 4 This desk's verdict, from the three reviews and Sonos's own stated limits. 5 Techgoondu, 22 September. 6 TechRadar, Engadget and DPRreview; two pages truncated, so no Singapore price is printed.

OBJECTS & ESCAPES · PARIS

A ceramic pebble from 1988, and the movement inside it was built for something else

Louis Vuitton has revived the watch that predates its own watchmaking by fourteen years. The case is the achievement; the calibre is a hand-me-down.

Most people date Louis Vuitton's watchmaking to the Tambour of 2002. It is wrong by fourteen years. In 1988 the house made the **Montre LV I** and **LV II**, designed by the Italian architect **Gae Aulenti**, on quartz alarm movements supplied by IWC — which also gave Vuitton access to the ceramic case technology IWC had introduced two years earlier. The name everyone uses, “Monterey”, is a corruption: non-French speakers mispronouncing *montre*.¹

Monday's revival takes the 37mm LV II and rebuilds it properly. The case is a lug-less pebble in high-tech ceramic — zirconium-oxide powder and pigment compressed into a solid block, machined on three axes, fired at **1,000°C**, then ground and polished by hand rather than moulded. Eleven millimetres thick. The crown sits at twelve, knurled in a *clous de Paris* pattern, and the strap threads through a loop on the caseback — Aulenti's original answer to having no lugs. Black ceramic takes a white dial and red syringe hands; green takes yellow.²

174

COMPONENTS IN THE LFT MA01.02, BUILT BY LA FABRIQUE DU TEMPS IN GENEVA: 23MM ACROSS, 5.4MM TALL, 26 JEWELS, FOUR HERTZ, 45 HOURS OF RESERVE, AN EIGHTEEN-CARAT ROSE-GOLD ROTOR²

Not a limited edition: permanent collection, boutiques only. The price is where the sources part company. Two independent titles give **€23,000**; a third gives **€19,000 “before taxes.”** Those reconcile as with-and-without value-added tax, but no source says so, Vuitton's own pages refused this desk, and no Singapore price has been published. Both figures are printed here rather than one chosen.³

THE CHATTER · THE REVIEWERS, ON THE RECORD

JX Su of SJX — who is not given to enthusiasm — writes that “I like the styling of the 1988 originals, and the modern-day reissues both manage to capture the look and feel of the originals,” and rates the LFT family as “excellent examples of industrial-*haute horlogerie*, and comparable in construction and finish to peers in Geneva.” **Michael Sonsino** at Oracle of Time makes the material case: the ceramic construction “makes it incredibly light, resilient, scratch-proof and the colour will never fade,” and the crown-at-twelve layout gives it “the appearance of a pocket or pendent watch.” Enthusiast forums are unreachable from this desk, so this is reviewer opinion, not community sentiment, and is labelled as such.⁴

THE SCEPTIC · THE MOVEMENT CAME FROM SOMEWHERE ELSE

The sharpest criticism comes from the reviewer who likes it. **JX Su** calls it “decently priced, albeit not quite a value proposition,” then puts the knife in a technical sentence: “The LFT MA01.02 is, however, a relatively small movement, as it was originally developed for the Spin Time. This explains its relative height and short power reserve.” Read plainly: nineteen to twenty-three thousand euros for a time-only watch with forty-five hours of reserve, in a case two millimetres thicker than the quartz original it venerates, on a calibre inherited from a jumping-hour complication. The case is the object here. The movement is what fits inside it.⁴

ALSO THIS WEEK

Bose brings back the wire, with cancellation. Announced Monday at **£99 / US\$99**, on sale 15 October: noise-cancelling *wired* earbuds, bus-powered with no battery at all, two microphones per bud, IPX4. The catch: output is fixed at 24-bit/48kHz whatever you feed it, so “lossless” means bit-perfect to a 48kHz ceiling. No app, so no equaliser.⁵

SOURCES 1 Monochrome and Oracle of Time, 28 September 2026, for the 1988 originals, Aulenti's authorship, the IWC connection and the etymology. 2 Monochrome, SJX and Oracle of Time agree on the case, the firing temperature, the hand finishing and the full calibre specification. 3 Two titles give €23,000; a third gives €19,000 before taxes. Vuitton's own site erred for this desk, so both are printed. No Singapore price exists in any source read. 4 SJX and Oracle of Time, quoted verbatim. Watch forums and two enthusiast titles refuse this desk, so no community sentiment is claimed. 5 What Hi-Fi?, 28 September, including the fixed output specification. 6 Monochrome, 17 September; the price contradiction is within that single article. 7 A trade preview of the November auction calendar, not opened at the houses' own pages; indicative.

CULTURE DESK · THE WEEKEND OF 25–27 SEPTEMBER

The number-one film in America is seven years old and is selling a trailer

Three new wide releases opened well and reviewed better. A re-cut library title beat all of them.

The weekend's biggest film was *Avengers: Endgame*, released in April 2019. Re-cut with added scenes and exclusive footage from *Avengers: Doomsday*, it took **\$26,106,643** from 3,060 screens — a per-screen average of \$8,532, against \$3.8m in previews and \$11.2m on Friday. Overseas added \$60m for an \$86m global weekend, lifting the film's lifetime worldwide total to about **\$2.88bn** and putting it roughly **\$39m** from *Avatar*'s all-time record.¹

That last sentence is a records-table artefact worth naming. The \$26m is being credited to the 2019 film's cumulative total rather than tracked as a separate release, which is what generates the headline — and *Avatar* has itself been re-released repeatedly to defend the same crown. **Paul Dergarabedian** of Comscore gives the commercial logic without embarrassment: bringing the film back “with an exclusive look at *Avengers: Doomsday* is a smart way to give audiences a reason to revisit.” *Doomsday* opens on 18 December. Disney is not selling a film; it is selling a trailer with a film attached.¹

DOMESTIC TOP FIVE · ACTUALS, 25–27 SEPTEMBER

Avengers: Endgame (Encore)	\$26.1m	3,060
Resident Evil	\$23.0m	▼ 62%
Heart of the Beast	\$20.0m	new
Primetime	\$19.2m	new
Forgotten Island	\$13.0m	new

Total weekend market \$122.9m. Actuals, not estimates — the two trackers differ slightly and the reconciled figures are used.¹

The films it beat were not weak. **Primetime** (A24) opened to \$19.2m from 2,874 screens with an 87 per cent critical score from 194 reviews — Robert Pattinson playing Chris Hansen, under Lance Oppenheim's direction, with the aggregator's own consensus praising a film that “interrogates the slippery boundary between justice and lurid spectacle.” **Forgotten Island**, DreamWorks' animation drawn from Filipino folklore, with H.E.R. and **Lea Salonga** voicing, scored **95 per cent** from critics and

97 from audiences. And *Resident Evil*, in its second weekend, holds a 96 per cent score across 271 reviews — a video-game adaptation reviewing like that is its own story.²

THE CHATTER · THE AWARDS, AND A FIRST-TIME WINNER AT NUMBER FOUR

Sunday's MTV Video Music Awards, hosted by Snoop Dogg at the Peacock Theater: Video of the Year to **Taylor Swift** for “The Fate of Ophelia”, Artist of the Year to **Madonna**, Song of the Year to **BTS** for “Swim”, Best New Artist to **Sienna Spiro**. Madonna took seven awards from thirteen nominations — the nomination count equalling Lady Gaga's single-year record — and performed at the ceremony for the first time in 23 years. Nirvana took the Video Vanguard. The juxtaposition worth keeping: Spiro won Best New Artist in the same week her album sits at number four on the British chart, behind beabadoobee's *Pylon* at one.³

THE SCEPTIC · FOUR MORE SONGS, AND THE SAME SHOW

Taylor Swift released *The Life of a Showgirl: The Encore* on 25 September with four new tracks written and recorded in Sweden after the parent album, and premiered the “Patient Zero” video at the awards two nights later. **Cassidy Sollazzo**, reviewing for *Consequence*, declines the occasion: the record “sounds a lot like the show we just sat through,” and its author has become “a predictable, bordering-on-bland caricature.” On the new material: “It relays the giddiness of young love, but shouldn't we expect a little more from her than that?” The parent album had the biggest opening week in history and thirty-four variants; the encore is four songs and a video. Both facts belong in the same sentence.⁴

IN BRIEF

San Sebastián closed its 74th edition on 26 September with the Golden Shell to **Mike Leigh** for *Tender Loving Care*; the Silver Shell for direction to Amanda Kernell; leading-performance shells shared by Kate O'Flynn and Asta Kamma August. Donostia lifetime awards to Werner Herzog and Naomi Watts. The last edition under departing director José Luis Rebordinos, who leaves behind a new award for original music.⁵

SOURCES 1 The Numbers' weekend chart for 25–27 September 2026 for the actuals and screen counts; Variety's Australian edition, 28 September, for previews, the global weekend, the lifetime total and Dergarabedian. Variety's main domain is paywalled; Box Office Mojo's estimates differ and the actuals are used. A claim about the biggest weekend—39 total since 2015 traces only to a blocked outlet and is not printed. 2 Rotten Tomatoes pages for each film, for scores, review counts, cast and the quoted consensus. Budgets and exit-poll grades were unobtainable. 3 The 2026 MTV Video Music Awards record, corroborated on two points by a second outlet. 4 Stereogum, 25 September, for the release; *Consequence*'s review, same date, quoted verbatim. 5 The 74th San Sebastián festival's award record. 6 What Hi-Fi?, 25 September — headline and date confirmed; the body did not render and is not quoted.

RAISING THEM WELL • THE PLATFORM QUESTION

Somebody finally asked teenagers which app it happened on

The answers differ enormously by platform — which is the first genuinely actionable finding in this field all year, and a pilot study with no odds ratios in it.

Most research on adolescents and social media treats “social media” as one substance. A paper published on 18 September does something simpler and more useful: it names twenty platforms and asks which ones the harm happened on. The sample is a nationally representative American panel of **472 adolescents aged 13 to 17**, mean age 14.9, surveyed about the previous four weeks.¹

The aggregates first. **44.7 per cent** had witnessed someone harassed; **29.8** had been harassed themselves. Online exposure exceeded offline for violent imagery (10.7 against 1.6 per cent), unwanted sexual images (13.4 against 2.0) and self-harm content (9.7 against 5.2) — but direct harassment was near-identical online and off, 16.4 against 18.5. And **51.5 per cent** said technology disrupted their sleep.¹

BY PLATFORM • PER CENT REPORTING, PAST FOUR WEEKS

	FELT WORSE ABOUT SELF	SLEEP DISRUPTED
Instagram	16.6	27.8
TikTok	13.4	42.2
Text messaging	5.6	10.2

YouTube scored 4.1 on the first column. Witnessing harassment: TikTok 22.7, Instagram 20.5, Discord 19.7. Unwanted approaches were most prevalent on Snapchat.¹

A four-to-one spread between Instagram and YouTube on feeling worse about yourself, and another between TikTok and messaging on sleep, is a difference a parent can act on. It is also the kind a self-report survey is worst at measuring, and the authors say so.¹

472

ADOLESCENTS IN THE SAMPLE — A PILOT, NOT A DEFINITIVE SURVEY. THE AUTHORS SAY SO IN THE FIRST LIMITATION: ESTIMATES “SHOULD BE TREATED AS PRELIMINARY”¹

THE SCEPTIC • READ THE LIMITATIONS, WHICH THE AUTHORS WROTE

Three things, all from the paper. First, this is a **pilot**: “a modest sample that... limits precision for platform-specific and subgroup analyses. Estimates should be treated as preliminary.” Second, the differences “may partly reflect differences in the user populations rather than platform design alone” — their own example is that Pinterest scored worse than Discord, which may say more about who uses Pinterest than about Pinterest. Third, and most damaging to the table above: “availability heuristics may also lead adolescents to attribute experiences to platforms they use most rather than where the harm occurred, which may inflate estimates for high-use platforms.” Note also what the paper does not contain: there are no odds ratios with confidence intervals anywhere in the main text or tables, only unstandardised coefficients. Anyone printing “X times more likely” from this study has invented it.¹

One more thing is worth reading, and it is unusually candid: the disclosure list is longer than the methods section. Among eleven authors, one joined YouTube’s family advisory board in January and reports fees from it; three worked at Meta or Facebook between 2006 and 2025, one of them on youth safety and age assurance, and one still holds the stock; one reports paid expert-witness work in platform trials; one has consulted for Google and YouTube; and one participates in the group serving as lead academic partner evaluating Australia’s social-media delay law.²

General guidance, not medical advice. The practical reading for a household with a twelve-year-old is on the next page, where a study forty times the size lands much closer to her actual age.

SOURCES 1 Lee AY, Lim CHJ, Liu SX, et al., “Adolescents’ Experiences of Online Harms by Platform,” *JAMA Pediatrics*, online 18 September 2026, doi:10.1001/jamapediatrics.2026.4692. Free full text read at source, including the prevalence figures, the platform table and the limitations, quoted verbatim. Fielded February to December 2025 on a probability-based national panel. 2 Conflict-of-interest disclosures as published with the paper, summarised by role; all are on the record there. 3 Iyer via USC’s own write-up, 24 September; Lee via PsyPost, 26 September, which independently flags the design limits. Stanford’s write-up refused this desk.

THE TWEEN YEARS • THE AGE BAND THAT MATTERS

Forty thousand children, and the effect is strongest at eleven

Published online on 24 September, in the *Scandinavian Journal of Public Health*: a population study of **40,748 Swedish adolescents aged 11, 13 and 15**, drawn from the Swedish arm of the World Health Organization's school-aged children survey. It asks what travels with trouble getting to sleep on more than one night in a week, adjusted for gender and how the child rates the family's finances.¹

Four things travelled with sleeplessness in every age group: problematic social-media use, cyberbullying victimisation, poorer mental wellbeing, and a household the child judged to be worse off. Lower wellbeing was the strongest, at roughly **threefold** higher odds at eleven and **fourfold** at thirteen and fifteen. Cyberbullying victimisation ran at about twofold at all three ages. Being a cyberbullying perpetrator was associated with sleeplessness mainly at eleven and thirteen.¹

11

THE AGE AT WHICH PROBLEMATIC SOCIAL-MEDIA USE WAS MOST STRONGLY ASSOCIATED WITH TROUBLE FALLING ASLEEP. THE ASSOCIATION WAS WEAKER, NOT STRONGER, IN THE OLDER ADOLESCENTS — THE OPPOSITE OF THE INTUITION MOST PARENTS CARRY¹

That is the finding worth carrying around. The instinct is that risk rises with age, as autonomy and screen time rise together; in this sample it did not. The authors report the odds in words rather than figures — “approximately threefold”, “roughly twofold” — and the publisher refuses this desk, so no confidence intervals are printed. And it is cross-sectional: a child who cannot sleep has more hours in which to scroll.¹

THE CHATTER • ONE OF THE AUTHORS DEFLATES IT HIMSELF

The most quotable line came from a co-author, and it points away from the headline. **Ravi Iyer** of USC Marshall: “Contact harms were often experienced more in person (e.g. more teens reported being bullied in person than on any platform).” The lead author, **Angela Y. Lee**, frames the contribution narrowly and correctly — the hope is that this “can help add specificity to conversations about social media and well-being effects.” Specificity, not causation.⁴

THE SCEPTIC • DON'T BE ALARMIST, FROM THE PEOPLE WHO ARGUE THIS FOR A LIVING

The previous page's paper carries Jonathan Haidt as an author, which makes his standing critics the right sceptics. **Andrew Przybylski** of the Oxford Internet Institute puts the epistemics plainly: “Extraordinary claims require extraordinary evidence. Right now, I'd argue he doesn't have that.” His sharpest line is about the shape of the argument: “It's a collision between a knowledge-discovery process, which is necessarily messy, and a narrative-generation process, which is unencumbered by critique.” **Christopher Ferguson** of Stetson supplies the counter-data: “Look at kids today — they're less violent, they smoke less, they use drugs less. All that happened during the social media smartphone age.” And **Sonia Livingstone** of the LSE aims at bans: “the argument for a ban is an admission of failure that we cannot regulate companies, so we can only restrict children.” None was commenting on either paper here.²

What this means in a Singapore household.

Secondary students may not use phones or smartwatches during school hours at all — recess and co-curricular activities included — under measures effective from January. School-issued devices now lock by default at **10.30pm**, brought forward from eleven, with sleep mode until 6.30am. Put the Swedish finding beside that policy and the conclusion is unglamorous: the curfew is doing the one thing the largest study here supports, and it matters most for the youngest child in the house, not the oldest. General guidance, not medical advice.³

SOURCES 1 Ellertsson AS, Ekstrand J, Hedin G, Sandlund C, Garmy P, *Scandinavian Journal of Public Health*, online 24 September 2026, doi:10.1177/14034948261485372. Design, sample and findings from the PubMed record. The publisher refuses this desk, so the exact odds ratios and intervals could not be read and only the abstract's wording is printed. 2 Przybylski, Ferguson and Livingstone are quoted from a June 2026 compilation rather than the original interviews, which this desk could not reach; none was commenting on either paper here. Flagged rather than dressed up. 3 September 2026, including Robinson's quotation and the stated next step. Outside this issue's window; carried as context. 3 Przybylski, Ferguson and Livingstone are quoted from a June 2026 compilation rather than the original interviews, which this desk could not reach; none was commenting on either paper here. Flagged rather than dressed up. 4 Iyer via USC's own write-up of the platform study, 24 September 2026; Lee via PsyPost, 26 September.


THE BIG PICTURE · THE PAR CURVE, READ COLUMN BY COLUMN

The long end did not lead. The two-year did

Monday's selloff is being written up as a term-premium shock. The Treasury's own curve says something narrower and more useful: this was a bet on the next meeting.

Take the numbers from the source and line them up. Between Friday's close and Monday's, on the Treasury's own par yield curve, the two-year note rose **eleven basis points**, to 4.92 per cent. The seven-year rose nine, the ten-year seven, to **5.24**, the twenty-year six. The two-to-ten spread narrowed from 36 basis points to 32.¹

That is a flattening, and a flattening is a statement about the central bank rather than the national debt. A long end that leads means worry about supply and the far future; a front end that leads means the next few meetings. Monday was the second — in a market already given one hike, the Federal Reserve having raised its range to **3.75–4.00 per cent** on 16 September, twelve votes to nil.³¹



The cause is not in dispute. Saturday's rejection of Iran's seven-day plan (p3) removed the only near-term route to cheaper crude, and Brent closed near **\$105.28**. An oil shock landing on an economy whose central bank has already restarted hiking is not a scare about growth. It is an inflation problem, and those price at the front of the curve.²

Equities, global bonds, gold and bitcoin all fell on the same day. In a genuine flight to safety, gold does not do that.

The levels are on p2. The tell is what fell together. The S&P 500 gave up 0.8 per cent, the Nasdaq 0.9.

SOURCES 1 US Treasury, Daily Par Yield Curve Rates, September 2026, read at source and mapped against the table's verbatim header row after this desk's tenor error in No. 115. 25 September: 2-year 4.81, 7-year 5.06, 10-year 5.17, 20-year 5.54, 30-year 5.49. 28 September: 4.92, 5.15, 5.24, 5.60, 5.56. The flattening and the figure are this desk's arithmetic on those rows. 2 AFP market wrap, 28 September, c.20:15 GMT, for Brent, WTI and the equity closes. 3 FOMC statement, 16 September 2026: the range raised a quarter point to 3%–4 per cent on a 12–0 vote. 4 Levels from stockanalysis.com's London lines and its BNDW page; TradingEconomics for DXY, USD/SGD and the STI; CoinDesk for bitcoin. VWRA's page mislabels the currency; the level was confirmed against the sterling line at 1.32539. Gold's exact Monday close could not be separated from a later Asia-hours print and is held. 5 This desk's reading of the four instruments. Explanation, not investment advice. 6 Zwemmer and Sharma via Fortune, 10 and 14 September; Klein via NBC News, 1 September — all predate Monday. 7 TradingEconomics; MAS's advance release calendar. A survey expectation circulating this week rests on a newsletter rather than MAS's own, and is not printed.

VWRA closed at \$192.30; **AGGU** slipped to \$5.71, and BNDW sits eleven cents above its fifty-two-week low, so the equity–bond correlation offers no shelter. Gold broke through \$4,200 to a seven-week low; bitcoin fell to \$83,588. A risk-off day bids gold. A monetary-tightening day sells it, and everything priced off a discount rate with it.⁴

WHAT IT MEANS FOR VWRA, GLOBAL BONDS, GOLD & BTC

One force is setting all four prices, and it is the expected path of policy rather than the growth outlook. For a global all-world equity line the drag is the discount rate, not earnings — which is why the index fell less than a per cent on a day the oil story worsened. For global bonds there is no diversification while the front end reprices. Gold is trading its real-yield and dollar sensitivity and losing on both; bitcoin is the liquidity proxy it usually is. All four readings turn on a number that has not printed yet.⁵

THE SCEPTIC · THE REPRICING MAY BE WRONG BY WEDNESDAY

Grace Zwemmer of Oxford Economics reads the producer-price detail as implying a monthly core PCE increase of roughly **0.15 per cent** — too soft to justify another hike. That print lands on **30 September**: Monday's conviction has a day of life in it. **Matthew Klein** of The Overshoot attacks from the other side — “today's rates are obviously too high only if inflation and growth are both poised to slow sharply from here” — reading rising yields as health, not distress. And **Ruchir Sharma** supplies the threshold nobody wants tested: the AI trade breaks when the ten-year “decisively breaches” five per cent. It closed at 5.24.⁶

Singapore is the quiet one. The Straits Times Index rose 0.31 per cent to **5,729**, the dollar held at 1.2772, and the Monetary Authority's statement is due no later than **14 October**. It has tightened twice this year; August core inflation ran at 2.2 per cent, a two-year high.⁷

DOWN THE RABBIT HOLE · ORIENTEERING

The sport where the map is the opponent

Two days from now, ninety-eight teams start a five-day race across Corsica with no satellite navigation at all. They will be handed the maps at the line.

The rules are almost embarrassingly simple. You are given a map and a list of control points, and you visit them in order, as fast as you can, through terrain you have never seen. A compass is allowed. A GPS is not. At each control you punch — these days by holding a card near a beacon — and at the finish you download the card, which proves the order. Get one wrong and you are disqualified. There is no appeal.¹

What makes it a sport rather than a walk is that the straight line is almost never the right answer. The map shows a printed red line between controls; the whole game is deciding which of the wrong routes around it is least wrong. Hence the vocabulary: you pick an **attackpoint**, a large unmistakable feature near the control, run to that confidently, and only then navigate precisely. You **aim off** deliberately to one side of a control sitting on a stream, so you know which way to turn when you hit it. You keep a **catching feature** beyond the control in mind, so you know when you have overrun.²

1897

THE FIRST PUBLIC CIVILIAN COMPETITION, IN NORWAY. THE INTERNATIONAL ORIENTEERING FEDERATION WAS FOUNDED IN COPENHAGEN IN 1961 BY TEN NATIONS, HAS BEEN RECOGNISED BY THE IOC SINCE 1977, AND NOW HAS 78 MEMBER FEDERATIONS³

The map itself is a specification, not a drawing. Jukola, the Finnish relay, prints at **1:10,000 with a 2.5-metre contour interval** to the ISOM 2017 standard; urban sprint races use a separate spec at much larger scales. Colour is a claim about how fast you can run: yellow is open land, **white is forest you can run through**, and progressively darker greens mean progressively thicker undergrowth. Blue is water. A newcomer's first shock is that white does not mean empty.⁴

It is also, genuinely, a lifelong sport. The **World Masters** championships in Poland in August ran classes from **M35 and W35 to M95 and W95**. That is not a courtesy category; it is a race, with a start list. At the other end the British course ladder is colour-coded so a beginner can choose honestly: white, yellow and orange are short and follow visible paths; green through black get longer and more committing.⁵

Singapore's scene is small and real. The **Orienteering Federation Singapore** is an IOF member, restored by former members of Singapore Polytechnic's outdoor club, and runs a league, a National Orienteering Race — this year at Republic Polytechnic on 18 July — plus Urban Wayfinder events and one on Pulau Ubin. A schools race has run with Northbrooks Secondary since 2016. What Singapore lacks is Nordic forest; what it has is sprint terrain, which since 2019 is a world-championship discipline in its own right.⁶

THE SCEPTIC · THE RECRUITMENT PROBLEM, STATED BY RECRUITERS

The people who run the sport are unusually candid about its central failure. On Attackpoint in March, **Gordhun** noted that Florida alone puts “some 2500 high school age people” through events every year — then asked the question nobody has answered: “Then they stop. Why do they stop?” **yurets** gave the blunt version: “the brutal truth: the sport only works for a small, rare type of participant.” **TheInvisibleLog** supplied the diagnosis — “contemporary culture is interested in low social commitment outdoor activity,” which a sport needing a club, a map, a start time and a forest is not. In Singapore, **Eugene Chow** of the federation puts it locally: “people have the stigma that the sport is difficult for them to complete.”⁷

What follows goes down: where the sport argues with itself, which of the gear is theatre, what a map costs to commission, and an honest ninety-day path in.

SOURCES 1 British Orienteering's newcomers' guide; SPORTident's product documentation. 2 Terms as used in the sport's own writing, including an Attackpoint glossary thread; the definitions are this desk's. 3 The federation's founding, IOC recognition and the 78-federation count from its encyclopaedia entry — orienteering.sport refuses this desk on every path, so the count is not from the IOF directly. 4 Jukola 2026's official invitation for the scale, contour interval and specification; British Orienteering for the colours. 5 The World Masters 2026 site for the class structure; British Orienteering for the course ladder. 6 Orienteering Federation Singapore's own listings and Hypeandstuff's report. No Singapore entry fee is printed: pricing is published only per event, on a login-walled page. 7 Attackpoint, March 2026, read at source; Chow via Hypeandstuff.

THE DEEP END · WHERE IT ARGUES, AND WHAT IT COSTS

Two timing vendors quietly obsoleted their own installed base in eighteen months

The sport's conversation lives in three places, and one is dying in public. **Attackpoint** is the anglophone spine, a training log with a forum attached where the same handles have argued for twenty years. **World of O**, Jan Kocbach's site, is the analysis brain; its ritual is **Route to Christmas**, in its nineteenth edition, posting one leg a day through December with the routes hidden so readers draw their own first.¹

Then there is **Nopesport**, the British board, whose own numbers tell the story: 7,451 topics and 125,031 posts, with Banter last active in 2021 and Reviews in 2019. In a May thread titled "What is the point of nopesport?", **DJM** did the sums: "When nopesport started there would be typically one new post each day. So far in 2026, there have been 14 new posts or approximately one post every 10 days."²

995,000

ROUTES UPLOADED TO LIVELOX IN THE PAST YEAR, ACROSS MORE THAN 28,000 EVENTS IN 83 COUNTRIES. IT HAS BECOME THE DEFAULT PLACE A RACE IS RE-RUN AFTERWARDS; PREMIUM IS €52.90 A YEAR, AND THE FREE TIER SHOWS ONE ROUTE AT A TIME³

The timing card is the first real purchase, and it just became compulsory. SPORTident discontinued the SI-Card5 and 6 in December 2024 "after more than 27 years of reliable service," killing the borrowed-card economy at a stroke. Its replacement, the contactless **SIAC**, runs **€69.90**: 128 punches, readable at three metres and at 80km/h, four years of battery — and it still works as a passive card when that dies. Jukola charged **€20 to hire one before May, €40 after.**⁴

Emit, the Norwegian rival, did it harder. Every emiTag with a serial below **4,252,870** must be replaced before this summer's season; the batteries are cast in and cannot be changed. Both vendors, in eighteen months, obsoleted the hardware their own communities already owned.⁵

Then the compass, and the mistake a Singapore buyer will make. Compasses are balanced for a magnetic zone — middle north, magnetic equator or

middle south — because the earth's field dips at different angles by latitude, and the wrong one tilts and drags. It matters more for a racing thumb compass than a hiking baseplate: **gruver** notes it is "worse than just a needle tilt for low-inertia compasses." **blairtrewin** gives the tolerance — "one zone away should be fine, even two is reasonably OK" — and New Zealand's GODZone race legislates it.⁶

A **Silva Arc Jet** thumb compass is €68.99 at the Lithuanian specialist all4o; the Finnish cult brand **STR8** runs €59 to €89. The zone-proof answers are Suunto's **M-3 Global at €72.30** and the mirrored **MC-2 Global at €87.80**. The connoisseur's choice, Moscompass — lately rebranded, identical product — is A\$93.90 in Australia with a quoted **one-and-a-half-second settling time**; a British retailer lists it as pre-order with no delivery date, which is its own comment on 2026. One useful mechanic: all4o's prices include VAT, deducted outside the EU.⁷

THE SCEPTIC · THE FLAGSHIP SHOE IS A CONSUMABLE

Orienteering shoes are studded, and the best-regarded ones do not last. On Attackpoint's long ownership thread, **Jagge** — who wrote the mapping software half this sport runs on — delivered the verdict on VJ's race flagship: "vj bold race is expensive and will not last more than 50 km." That shoe is €199 direct. **AndrewM**: "VJs are the comfiest and grippiest, although I find they don't last quite as long as inov8s," and on the rival, "icebugs do not last much more than 150km." Nor are metal studs a clean win — **jjcote**: "Slippy wet wood, too. There's no one right answer for everyone." For a buy-once reader the inversion matters: here the premium buys lightness, and lightness costs life.⁸

Which points at the durable middle. The **VJ Integrator 20** is €139.90 at all4o against €179 from VJ direct; **Icebug's Sisu OLX**, on steel studs, is €122.90. The same shoes cost €40 to €60 more from the maker than from the specialists — a rare case where going direct is the expensive mistake. Sizing is the hazard for anyone ordering into Singapore: VJ's last runs snug, and owners report needing two EU sizes above their street shoe.⁹

SOURCES 1 Attackpoint and World of O. The former's index pages refuse automated requests but individual threads read normally, and every quotation here was read that way. 2 Nopesport's own board index for the counts and activity dates; the May 2026 thread read at source. 3 Livelox's homepage and subscriptions page. 4 SPORTident's news page for the card end-of-life; its SIAC documentation for the capacity, range, speed and passive fallback; all4o for the price; Jukola's invitation for the hire charges. 5 Emit's own product pages. 6 Silva's and Suunto's explanations of balancing zones; four orienteers quoted from an Attackpoint thread; GODZone's published kit list. 7 all4o and Silva's Swedish store, read 29 September 2026; Wildfire Sports for the O-Compass specification; Kong Running for the pre-order status. "Race Jet" is no longer in Silva's catalogue. 8 Attackpoint, July–August 2026. 9 all4o and vjshoes.com, both read 29 September; sizing from the same thread.

THE DEEP END · THE CRAFT, THE MONEY, THE WAY IN

Sixteen to sixty hours per square kilometre, and the laser did not change it

The hidden craft here is mapping, and its economics are brutal. British Orienteering's guidance puts field survey at **16 to 60 hours per square kilometre**, with cartography adding 50 to 100 per cent of that. The worked example is quietly devastating: four square kilometres at forty hours each "could easily take sixteen months" for a volunteer with ten free hours a month. The same document defines a professional mapper, for insurance purposes, as anyone earning more than **£5,000 a year**.¹

Commissioned prices, from mappers talking to each other rather than a rate card: **US\$1,500** for a 150-acre sprint map, **US\$3,600** for a 750-acre forest map, travel excluded. When a mapper asked the same forum last November what modern sprint mapping costs per square kilometre, nobody answered.²

THE CUTTING EDGE · WHAT LIDAR ACTUALLY DID

Karttapullautin, the free Finnish engine written by the Attackpoint regular **Jagge**, turns airborne laser data into a usable basemap — rather than computing mathematical contours it lets elevations drift onto features it has already detected, giving 2.5-metre contours with less jagged noise. **Cristina** reports that "plenty of mappers prefer the KP-generated contours over the OCAD ones." Jagge volunteers its limit unprompted: it "does not do well with really thin lidar data." And the honest verdict came from a beginner, **arya**: "I used the DEM lidar of the area to produce 2.5m contours, and heavily edited them by hand." The laser makes the basemap. It does not touch the fieldwork.³

Genuinely new this year, honestly graded: **Purple Pen 4.0**, the free course-setting standard, reached cross-platform beta in February — real, because the tool had been Windows-only. **OCAD** shipped LiDAR contour and vegetation tooling for field tablets; incremental, but it is how mapping is now done. Against that, **SPORTident's** 2026 news was a card reader and two new colourways. There is no AI map generation in this sport, whatever you have read.⁴

The first ninety days — and the honest budget line. Everything in this sport is cheap except getting to it:

a compass, a card, gaiters and a durable shoe come to about three hundred and thirty euros all in, and the return fare from Changi is the rest of the budget.⁵

Spend a fortnight spending nothing: read *World of O*, open a free Livelox account and watch other people's routes on races you will never run, and turn up to a federation event in borrowed shoes. Then buy in this order: the compass first, the cheapest permanent upgrade in the sport; your own SIAC, now that borrowing is dead; gaiters at **€33–43** and a control-description holder at **€8.90**, which everyone forgets. Leave the shoes until last, and buy the durable model. The club one-piece cannot be bought at all: Trimtex needs five pieces over five to six weeks, so it is a club decision.⁶

THE SCEPTIC · WHAT THE GEAR IS WORTH, GRADED HONESTLY

The compass is real. A needle that settles in a second and a half changes how often you can take a bearing on the move, and four orienteers independently document the zone-balancing failure. **The shoe is real but inverted** — the premium buys lightness, and lightness costs durability. **The headlamp is the weakest case.** The engineering is genuine, from Petzl's NAO RL at US\$199.95 to Lupine's Blika at €469 — but when a British orienteer asked the forum last September what the best head torches were, he got **zero replies**. And the Olympic dream is pronounced dead by its own advocates: **robplow**, who started the thread arguing for inclusion, concluded the quest "has little or no chance of success." The honest counter-data on decline: the US championships drew 323 in 1985 and **356 in 2026**. Not collapse. Stagnation.⁷

NEXT DESCENTS — YOU PICK TOMORROW'S HOBBY

Home cheesemaking — cultures, rennet, and a cave in a wine fridge.

Bookbinding & book repair — Coptic to case binding, and the archival-materials rabbit hole.

Slot car racing — braid, magnets and a garage banked oval.

SOURCES 1 British Orienteering Map Group, "Procurement of Maps". Dated November 2001: the survey hours are craft constants and printed as such; its money figures are not used except the £5,000 threshold, which is definitional. 2 Attackpoint's mapping-cost thread, September 2022, with the November 2025 follow-up that went unanswered. A 2022 baseline. 3 Attackpoint's Karttapullautin threads, December 2025 and January 2026. 4 purple-pen.org for the February 2026 release and the cross-platform beta; OCAD's blog for the 2026 LiDAR posts; SPORTident's news page. No AI map-generation announcement was found in twelve months. 5 Prices as sourced on the preceding page. 6 Trimtex's custom arm; components as on **p16**. 7 Petzl's and Lupine's own pages; the unanswered head-torch thread, 19 September 2025; Attackpoint's Olympic thread and its 1985-against-2026 comparison. A separate thread on format reform is not quoted: it opens with an April Fools hoax.

THE LONG GAME • CADENCE

Ten thousand is the wrong variable. The right one is your best thirty minutes

A hundred and three thousand wrist accelerometers say that how fast you walk during your briskest half hour does more work than how far you walk all day.

The paper reached print on 24 September, built from exactly the data a wrist device already generates. **103,684** people wore a triaxial accelerometer for a week; after exclusions for existing disease and for deaths in the first two years, the analysis covers **64,743** people, mean age 61.5, and 1,697 deaths over about eight years.¹

The exposure is the interesting part. Rather than daily steps, the authors use **PK30** — the average cadence over the thirty highest-stepping minutes of the day. Against a reference of about 45 steps a minute and under 5,000 steps a day, the two variables come apart. Under 5,000 steps at **80 a minute**: hazard ratio **0.72** (0.54–0.97). Between 5,000 and 7,500 at 60 a minute: **0.70** (0.58–0.86). The lowest risk sits at **7,500–10,000 steps with a peak cadence near 100**: **0.57** (0.47–0.69).¹

A hundred steps a minute is the conventional threshold for moderate-intensity walking, and where the curve bottoms out; above ten thousand steps a day it flattens and more steps buy nothing. Cardiovascular mortality mirrors it. The corresponding author's translation is characteristically unheroic. **Emmanuel Stamatakis**: “Both step counts and stepping intensity matter, and the combination that works is the one a person can sustain.”²

THE CHATTER • AND THE NUMBER THAT WILL BE MISUSED

A larger analysis published on 15 September pools German and British cohorts: **147,414 adults**, 5,545 deaths. Against no activity, 150 minutes a week of moderate activity gives a hazard ratio of **0.70** (0.67–0.72) — and the same association is reached by roughly 510 minutes of light activity or about **twelve minutes a week of vigorous**. Twelve minutes is the most shareable number in exercise science this month and the most abusable: accelerometers cannot tell resistance work from aerobic, and people already dying stop doing vigorous things first.³

THE SCEPTIC • THE PAPER'S OWN METHODS SECTION, AND A CARDIOLOGIST

The load-bearing assumption is stated openly and is an assumption: the cohort's response rate was **5.5 per cent**, and the authors argue, citing prior work, that this does “not materially influence” the association. They concede that reverse causation “cannot be excluded,” that step metrics miss cycling and swimming, and that wrist devices misread other movements as steps.

Nothing here excludes the possibility that a hundred-step peak at sixty-one is a *marker* of an already-healthy cardiovascular system rather than a cause.

Kevin Shah, a cardiologist not involved, deflates the optimisation instinct: “The biggest health gains often come from going from no activity to some activity.”

What would settle it is a trial randomising a cadence target against a volume target. Nobody has run one.⁴

What to do with this. Program the wearable to a peak-cadence target rather than a step target: one deliberate thirty-minute block a day at a hundred steps a minute or better — three steps every two seconds, genuinely brisk. Then convert one existing hour a week of easy movement into something properly vigorous. The honest catch is that no mainstream wearable surfaces PK30 natively; it has to be approximated as the steps taken during your best continuous half hour. The metric that predicts the outcome is not the one the product shows you. Not medical advice.⁵

SOURCES 1 Wei L, Ahmadi MN, Koemel NA, Biswas RK, Del Pozo Cruz B, Stamatakis E, *Br J Sports Med* 2026;60(18):1401–1408, doi:10.1136/bjsports-2025-111173, print issue 24 September. Read in full from the journal's own press copy, the publisher's site being closed to this desk; all hazard ratios and intervals are the paper's. Competing interests: none declared. 2 Stamatakis via Medical News Today, 9 September. 3 Stein MJ, Lendt C, et al., *Int J Behav Nutr Phys Act* 2026;23(1):85, published 15 September; free full text read at source, including the limitations quoted. 4 Limitations verbatim from the first paper; Shah via Medical News Today. 5 This desk's reading of the two together. Neither device's support for a peak-cadence metric was verified, and the limitation is stated rather than assumed away. General guidance, not medical advice.

THE DIARY • THE DATED AGENDA

The race is in ten days and the cheap seats are the ones still on sale

Every row re-checked this week against the venue's own page or the official seller. Two of the biggest names on it cannot be bought at any price.

SINGAPORE

BOOK NOW • PLAN AHEAD

- 2 OCT

Octoburst! 2026
 ESPLANADE • THE CHILDREN'S FESTIVAL, TIMED TO CHILDREN'S DAY • PER-SHOW PRICING NOT PUBLISHED ON THE LANDING PAGE¹
- 2-3 OCT

The Weeknd
 NATIONAL STADIUM • TWO NIGHTS²
- 7 OCT

Hans Zimmer
 SINGAPORE INDOOR STADIUM²
- 9-11 OCT

★ Formula 1 Singapore Grand Prix
 MARINA BAY STREET CIRCUIT • A SPRINT-FORMAT WEEKEND. **THE KILLERS** ON THE PADANG AFTER SATURDAY QUALIFYING, **LANA DEL REY** SUNDAY, ACROSS TEN STAGES. **GONE:** PADDOCK CLUB, SUPER PIT GRANDSTAND, SEVERAL SUITES. **STILL ON SALE:** STAMFORD GRANDSTAND FROM **S\$228/DAY**, PIT FROM S\$428, ZONE 4 WALKABOUT FROM S\$198, EMPRESS PLATFORM FROM S\$38. CHECK THE ZONE — ZONE 4 DOES NOT REACH THE WHARF OR BARGE STAGES³
- 17-18 OCT

BIGBANG
 NATIONAL STADIUM • THE 17TH IS REPORTED SOLD OUT; THE 18TH WAS STILL LISTED²
- 10-11 NOV

My Chemical Romance
 SINGAPORE INDOOR STADIUM • TWO NIGHTS²
- 17-22 DEC

★ BTS — four nights
 NATIONAL STADIUM, 17, 19, 20 AND 22 DEC • **SOLD OUT ON PRIMARY.** THE LARGEST BOOKING ANY ACT HAS TAKEN HERE THIS CYCLE, AND THE STORY IS THE RESALE MARKET RATHER THAN THE SHOW²
- 22-24 JAN

★ ART SG 2027, fifth edition
 SANDS EXPO • VIP PREVIEW AND VERNISSAGE 21 JAN, INVITATION-LED; PUBLIC DAYS THE 22ND TO 24TH. PUBLIC TICKET PRICES NOT YET PUBLISHED⁴

JAPAN & THE REGION

THE ONES WITH A TICKET CLOCK

- to 8 OCT

teamLab Borderless — special exhibition
 AZABUDAI HILLS, TOKYO • ADULTS **¥3,600**, 13-17 ¥2,800, 4-12 ¥1,500 • ADULT PRICING IS **DYNAMIC** AND RISES ¥200 ON SITE • TIMED ENTRY, NO DOOR SALES • CLOSED TUE 20 OCT⁵
- 8-22 NOV

★ Kyushu Basho — Fukuoka
 FUKUOKA KOKUSAI CENTER • ON SALE SINCE 19 SEP • BOX **¥34,000-60,000**, FIRST-FLOOR CHAIR ¥8,500-10,000, SECOND-FLOOR ¥5,000-7,500, SAME-DAY GENERAL ADMISSION **¥3,500** • ENGLISH AUDIO GUIDE AND SIGNAGE • WEEKENDS AND THE LAST THREE DAYS GO FIRST⁶
- 28 NOV-5 MAY

★ Niseko United opens
 ALL FOUR AREAS FROM 28 NOV • NIGHT SKIING 12 DEC TO 22 MAR • ANNUPURI AND GRAND HIRAFU RUN TO 5 MAY; NISEKO VILLAGE CLOSURES 4 APR, HANAZONO 11 APR • THE ONLINE LIFT PASS NEEDS A ¥500 CARD⁷

PLANNING NOTES • WHAT THE CLOCK IS DOING

Two of the four marquee picks are already unbuyable at face value, which makes the other two the decisions. At the Grand Prix the prestige end has gone and the cheap end has not: a Stamford grandstand day at S\$228 or a Zone 4 walkabout at S\$198 still exist ten days out — and the walkabout means standing for hours, worth knowing before buying it for a twelve-year-old. Kyushu is the best-value marquee by some distance: a chair seat from ¥5,000 in a ten-thousand-seat hall with an English audio guide, over a fortnight, so a first-week weekday is gettable where a Saturday is not. And set reminders for two things with no ticket yet: Hatsu Basho sales in early December, and O-Ringen's early discount before 1 November (p15).⁸

SOURCES All entries re-verified 28-29 September against a venue, organiser or official ticketing page rather than an aggregator, except where noted. 1 Esplanade's own festival pages, neither of which publishes per-show pricing. 2 Concert dates from the one listings aggregator this desk could reach: ticketmaster.sg and sistic.com.sg both refused automated requests, so no individual concert price is printed and the sold-out notes are reported as reported. 3 singaporepgp.sg's tickets page and entertainment announcement, for the line-up, the zone rule, the sold-out categories and every price. 4 artsg.com. 5 teamLab's own English pages, including the dynamic-pricing note and the October closure. 6 Japan-based sumo ticketing guides; sumo.or.jp rejected the schedule URL, so January is given as unpublished rather than guessed. 7 Japan ski-season listings; Hakuba's dates were genuinely unannounced. 8 This desk's reading. Christmas Wonderland is not listed: the organiser's site still showed last year's programme.

THE TABLE · SINGAPORE'S NEWEST OPENINGS

Ten seats, no reservations, and nobody is arguing about the food

One kitchen returns to this page for its verdict run, seven issues after it opened. The other three are new.

Start with the one finally judged. **MUKAI**, in the basement of Plaza Singapura, is the first overseas outlet of a Shibuya seafood-rice-bowl specialist. The format is a two-stage meal: eat the bowl as a conventional *kaisendon*, then pour over a dashi — katsuobushi or asari clam — turning what remains into *ochazuke*. The Kaiho Don Nami is **S\$28.90**; the Tokujo, adding otoro and uni, **S\$40.90**; the Goku Goku **S\$79.90**, capped at thirty bowls a day.¹

The verdict, when it came, split exactly where it matters. Middleclass gave it three and a half out of five and singled out the second stage — “the ochazuke experience stood out to us” — while praising the portions and the solo-friendly counter. Then the criticism, which goes at the core proposition: the fish “lacked strong, sweet flavour” and needed soy sauce to carry it, and at these prices “we wished the fish quality could be better.” Worth going once; they doubted a return. The theatre lands, the fish is contested. Walk-ins only, and a basement mall unit, so two o'clock or half past five will get you in.¹

THE BUZZ · THE QUEUE IS THE REVIEW

Nice Noodle, in Oxley Tower on Robinson Road, is a ten-seat bar from the team behind Nice Rice, with chef **Mandel Ban** out of one-starred Vianney Massot. Three proteins against six bases, brothy or dry; bowls run **S\$10 to S\$16**. The photographed dish is The Cool at **S\$12** — thin noodles in an iced tare broth.

HungryGoWhere headlined its review “Nice, or not?” and answered itself: “Yes, we think it is,” with the chashu “tender with a silky streak of fat” placed “right up there” against specialist shops in Singapore and Japan, and the Indoramén judged to “stack up against many popular Indomie goreng spots.” The only consistent criticism is structural: ten seats and a queue. Nobody is arguing about the food; they are arguing about the door. Hold the note that at these prices the value case is doing some work in the praise.²

Lulu's Lounge has reopened after six years, moved from a Pan Pacific ballroom to a Chinatown hotel and

reskinned as 1930s Shanghai — jade, lacquer, red neon, a jukebox beside the DJ booth. The cocktail programme was built by the outfit behind Manhattan and ATLAS, the most credible line on its CV. Small plates run **S\$12 to S\$45**: caviar siu mai at **S\$28**, short-rib skewers at **S\$26**; cocktails **S\$23 to S\$25**. Two independent reviews were warm — the beef “melt-in-the-mouth tender,” the cod “pristine.”³

THE SCEPTIC · TWO REVIEWS, NO RESERVATIONS, AND THAT IS THE PROBLEM

Neither published review of **Lulu's Lounge** contains a single stated criticism, which for a venue with a hotel behind it, a famous cocktail lineage and a burlesque programme usually means the coverage is running ahead of the cooking. The fair framing is the venue's own: small plates “meant to accompany the cocktail programme.” Twenty-eight dollars for siu mai is supper-club pricing for supper-club reasons. Judge it as a night out, not a dinner. Credit where due: no cover charge, and it runs to three in the morning.³

Furthest from town and cleverest for it: **Joylion Buffet Hotpot**, in the basement of Lentor Modern, brings conveyor-belt hotpot — still rare here — at **S\$22.90++** for weekday lunch and **S\$32.90++** otherwise. Each diner gets a pot built into the table, so there is no shared broth and no trips to a station — good for a solo diner, or a table where one wants mala and another does not. The catch is in the format, and the coverage names it: you wait for what you want to come past, so “patience helps.” At a buffet, wanting is the whole psychology. It has been reviewed only lightly, and this desk will not invent a verdict on the seafood.⁴

Two names this desk killed. Both appeared on reputable September opening lists and neither is new: one has been reviewed at the same address since 2022, the other since 2024. The lists had caught a relaunch and a menu refresh — the commonest way a restaurant page goes wrong.⁵

SOURCES 1 Middleclass.sg's review for the rating and both quoted criticisms; Eatbook's opening coverage and the venue's own site for the format, prices and the thirty-bowl cap. This is the venue's second and final appearance here: it opened in No. 110, and the Table Ledger allows a verdict run no sooner than seven issues later. 2 HungryGoWhere's dedicated review for every quotation, the format and the prices. Two sources disagree on both the opening date and the closing time, so neither is printed as settled. 3 Middleclass.sg and HungryGoWhere for the dishes, prices and verdicts; the venue's own pages for the programme, hours and split booking. The absence of published criticism is stated as an absence, not filled in. 4 Eatbook for the pricing, the format, the halal position and the “patience helps” caveat. 5 Both discarded venues were checked against earlier reviews at the same addresses, from 2022 and 2024; they are not named, because the error was the listing's.

THE UNDERCURRENT · TAIPEI

The brittle place

A country of twenty-three million supplies about a fifth of one social network's traffic on earth — more than the United States. Two months before an island-wide election, somebody has built two hundred and ninety-four accounts on it and has not yet said anything political.

Taiwanese people do not call it Threads. They call it 脆 — *cùi* — which means crisp, crunchy, brittle, and of a person or a thing, fragile. It is not a translation. Mandarin has no “th” sound, so users needed something sayable, and this was the rough phonetic stab that won. A rival candidate, 串, *chuàn*, meaning a string of beads, was closer in sense and lost. A technology magazine documented the contest in April 2024 and noted that the winner was “too far-fetched” as a transliteration but had “already caught on among the users.”¹

You can tell a nickname has taken when it starts producing children. Taiwanese users are 脆友, *cùi-yǒu*, crisp friends. The feed is the 河道, the river channel. The platform's native register — low-stakes, unglamorous, personal — is 廢文, trash posts. The local trade press now puts the character in headlines without glossing it.¹

How a country ends up here is worth retracing, because nobody designed it. Threads launched globally in July 2023 and was, by common consent, abandoned within weeks. Taiwanese adoption began quietly around November 2023, among users leaning one way politically, and then accelerated through the presidential election of January 2024 until the app sat at the top of the local download charts for months. In June this year, announcing five hundred million monthly users worldwide, the platform's own product head named Taiwan as central to that success and named elections as the reason it took off there. A network built expressly to avoid politics had become, in one country, the place politics happens.²

28.8%

OF TAIWAN'S ENTIRE POPULATION INSIDE THE PLATFORM'S ADVERTISING REACH — 6.65 MILLION PEOPLE, OR 31.9 PER CENT OF THE ELIGIBLE AUDIENCE AGED THIRTEEN AND OVER³

The intensity shows up in the texture as much as the totals. An analytics firm read **36,465,742** Taiwanese posts across 2025 and found the weekly engagement

peak on Thursday evening, between eight and nine, with Sunday the trough — an inversion of the global weekend pattern, and the signature of a workplace-adjacent public rather than a leisure one. The average Taiwanese post drew **177.89 interactions** against a local Facebook benchmark of 110.06.⁴

What Taiwan built, then, is a public square nobody built: privately owned, foreign, with no local corporate entity, publishing no per-country numbers, and answerable to no Taiwanese regulator. In 2026 the bill for that arrived three times in six months, and each instalment is worth taking separately, because the lesson of each is different.

A network built expressly to avoid politics had become, in one country, the place politics happens.

The first instalment came in April, and it was an arithmetic error. A figure circulated saying Taiwan held **94.86 million** Threads accounts — four times its population, and therefore proof, to anyone who wanted proof, that the square was mostly fake. On 5 May the Taiwan FactCheck Center ruled it false. The firm behind it had multiplied a global active-user count by Taiwan's share of web traffic, which are two incompatible units; the platform publishes no per-country account numbers at all. The firm acknowledged the error and withdrew it. A companion oddity — that seventy per cent of Taiwanese usage was on desktop — turned out to be an artefact of counting only the web version and excluding the app.⁵

Hold onto that, because it sets the standard for everything that follows. The single most-shared piece of evidence that Taiwan's public square was hollow was a spreadsheet mistake, and it travelled because people wanted it to be true.⁵

SOURCES 1 MIT Technology Review, 2 April 2024, for the nickname contest and the quoted assessment; the derived vocabulary from Taiwanese platform-culture writing. 2 The same report for the adoption timeline; the product head's remarks of 17 June 2026 as carried by a Taiwanese outlet and reproduced on a public forum, the form in which this desk could read them. 3 DataReportal's Digital 2026 figures for Taiwan, citing the platform's own advertising tools, underlying data dated October 2025. Advertising reach is not an account count, and no reliable Taiwanese account total exists. 4 QSearch's analysis of 36,465,742 posts across 2025, as reported in the Taiwanese business press. 5 Taiwan FactCheck Center ruling of 5 May 2026. The much-quoted 21.08 per cent traffic-share figure comes from the same firm whose account figure was withdrawn and is not printed, though Taiwan's outlier status is independently supported above.

The second instalment came in June, and it was an accident in California. A Meta age-verification mechanism, deployed in May partly for European compliance, began misclassifying accounts as belonging to under-thirteens and suspending them. Taiwan was hit disproportionately. Within two days the casualty list included a legislator, a Taoyuan city councillor with eighty-eight thousand followers and thirty million weekly impressions, a former Taipei mayor, a television news service, a party youth organisation — and, overnight, the Ministry of Foreign Affairs’ own account.¹

Taiwan’s Ministry of Digital Affairs issued a statement on 15 June demanding that Meta actively assist restoration, and opened an inquiry. On the 17th it announced that Meta had committed to proactive detection and restoration rather than making every user appeal individually, and that roughly **ninety per cent** of wrongly suspended accounts were back by that afternoon. More than two hundred media organisations and public figures had been affected, across every party, plus thousands of ordinary users. Meta’s position throughout was that this was a technical fault and not targeting, and there is no reason to doubt it.¹

Which is precisely what makes it instructive. No Taiwanese statute compels an offshore platform to verify accounts or to disclose per-country numbers; the digital intermediary services bill never passed. The ministry’s entire leverage, in the week its own foreign ministry was deplatformed, was persuasion. Ninety per cent came back within three days. That also means ten per cent did not.¹

The state’s only instrument, in the week its own foreign ministry vanished, was a press release asking nicely.

The third instalment is the one this desk thinks is genuinely new, and it arrived in two parts from two separate investigators who were not looking for the same thing.

In April, Doublethink Lab published an audit of a content farm it names as the **Borderless Group**,

headquartered in Qinhuangdao, Hebei. The operation runs **1,217 identified domains**, of which 475 remain live — 134 in Japanese, seventy in Portuguese, plus English, German, Finnish and Czech editions — each dressed as an ordinary small news site and filled with plagiarised copy and machine-written summaries. Against Taiwan it ran **549 inauthentic Threads accounts**, created between July 2023 and February 2026, which produced at least **467,547 posts and replies**. Eighty-three per cent carried a link back to a farm domain. The best-performing single post reached 1.2 million views and 90,600 likes.²

467,547 POSTS AND REPLIES — WHAT THEY WERE ABOUT

Gossip, lifestyle and curiosity-bait, with a link attached — 99.62%

1,779 political posts — 0.38%, and ideologically mixed

The vermilion sliver is drawn to scale. The infrastructure is real; the propaganda is almost absent.

A Meridian original, from Doublethink Lab’s published counts, April 2026.

The finding that makes this worth writing about is the one nobody expected.²

Here is the uncomfortable part. Of those 467,547 posts, exactly **1,779** — nought point three eight per cent — were political, and the political content was ideologically mixed rather than consistently hostile to any party. The automation is not subtle: 6,438 posts from 171 accounts share an identical template, exactly six lines with the hashtags confined to the last. But the money is not political. It is programmatic advertising. This is a business.²

The method, reconstructed from the audit and from an engineer’s account given to a Chinese-language outlet, is patient in a way propaganda rarely is. Register domains in bulk. Fill them with stolen copy. Create accounts in bulk with stolen profile photographs. Then 養號 — “raise the account” — feeding it entertainment news and life hacks for weeks until the algorithm and the human reader both file it as ordinary. Cross-like inside the network to manufacture early velocity. Only then attach the link. Keep political content to a trickle so the accounts never trip the coordinated-behaviour heuristics platforms actually enforce. And reap the laundering: **566 genuine Taiwanese users** shared the farm’s domains of their own accord, at which point the content is inside the authentic graph.³

SOURCES 1 Taiwan’s Ministry of Digital Affairs, statements of 15 and 17 June 2026, for the demand, the inquiry, the Meta commitment and the ninety-per-cent figure; Taiwanese public broadcasting and a Taiwanese wire for the named casualties and the foreign ministry’s account. The absence of an enabling statute is a matter of record. 2 Doublethink Lab, “Assessing the Borderless Group’s activity on Threads,” 28 April 2026. The publisher errors for this desk, so the report was read through a text proxy and every figure was cross-checked against a second outlet’s independent write-up of 22 July and the lab’s own index. 3 The same report for the sequence and the template count; the engineer’s account of the method comes from an outlet with a declared editorial stake, is given under a pseudonym, and is used only where it matches the audit’s technical findings.

The second part arrived in June from NewsGuard, which was looking at something else. It documents **294 coordinated accounts** on the same platform, launched in **May 2026**, presenting as women seeking Taiwanese men. A hundred and eighteen share a handle pattern, an animal name followed by digits; seventy posted identical profile text; ninety instances of near-identical simultaneous posting were found. A researcher who examined 457 posts across four days found the network active from nine in the morning to seven at night **Beijing time**, with a break between twelve and one. The profiles say “currently in Taiwan”; most that disclose are in China.¹

And as of 24 June, the date of the analysis, the number of political posts across all 294 accounts was **zero**. They are pre-positioned. **Jerry Yu** of Doublethink Lab supplies the timetable: fake accounts “usually start emerging roughly six months before the elections,” pivoting once they have an audience. May plus six is November. Taiwan votes island-wide on **28 November**.²

O

POLITICAL POSTS FROM 294 COORDINATED ACCOUNTS, SIX MONTHS BEFORE A NATIONAL ELECTION. NEWSGUARD COULD NOT DEFINITELY LINK THE CAMPAIGN TO THE CHINESE GOVERNMENT¹

This is why the correct frame is not “China is flooding Taiwan with propaganda.” On the published evidence it is not. What has been built, over thirty-one months, is *rentable distribution* inside the platform where Taiwan argues with itself — a worse sentence, because an ad-arbitrage business has no ideology to abandon and no reason to refuse a better offer.²

The guns exist. They are pointed at a slimming advertisement. Nobody has decided what to do with them.

The sceptics deserve the floor, and the strongest is the lab that wrote the report: anyone reading 549 accounts as an election attack is reading past a finding of 0.38 per cent. Nor does this country have a clean record of telling an operation from a bug. In 2024, when a viral trick claimed that blocking two

SOURCES 1 NewsGuard, Charlene Lin, 25 June 2026, read at source, for the account count, the launch month, the handle and profile-text patterns, the batch-posting instances, the working-hours analysis of 457 posts, the location data, the zero political posts as at 24 June, Jerry Yu’s quotation, and the statement that the campaign could not be definitively linked to the Chinese government. 2 This desk’s reading of the two investigations together. 3 Chiu Wei-chieh’s tests as reported by a Taiwanese daily in 2024; the April error as ruled by the FactCheck Center (p21). 4 Liberty Times, 28 September 2026, for Shen’s follower count and its growth, the ranking, the incumbent’s 88,000, and the slogan’s origin at Ningxia Night Market on 18 August. His co-founding of Doublethink Lab is public record. 5 Taipei Times, 21 September 2026, for the Ministry of Justice measures and the seven-year maximum; that report does not itself mention this platform. 6 This desk could not search the major Western outlets, which refuse it; no feature surfaced in nine searches, and the only Western prior is a 2024 technology piece arguing the opposite thesis.

accounts would collapse a politician’s 350,000 followers to a few dozen, the former city councillor **Chiu Wei-chieh** tested it across several politicians, found no pattern, and called it a privacy-setting artefact or an unfixed bug. Taiwan has already mistaken a display error for an information operation once, and an arithmetic slip for a hollow public square in April.³

Meanwhile the election is being fought on 脆 in a way that has no precedent. The governing party’s Taipei mayoral candidate, **Shen Po-yang**, passed **640,000 followers** yesterday, up two hundred thousand in three months, second among Taiwanese politicians only to a former president. The sitting mayor, of the opposition and holding the city, has **88,000** — better than seven to one, in a seat his party does not hold. His slogan, 洋流 — *yángliú*, ocean current, a pun on his name — was not devised by the campaign: supporters made orange placards at Ningxia Night Market in August while he canvassed, and the campaign adopted it within weeks.⁴

An irony worth stating without making too much of it: Shen co-founded Doublethink Lab, whose audit is the spine of this piece. The biggest beneficiary of Taiwan’s migration onto 脆 is the man who built the institution now counting its fake accounts.⁴

The state has begun to move, narrowly: on 21 September the Ministry of Justice promised harder enforcement against election deepfakes, with penalties for platforms that ignore removal orders and up to **seven years** for individuals. A real instrument, aimed at synthetic content. Nothing in it touches an account that has spent six months posting gossip.⁵

So watch three things before 28 November. Whether the political share of that network’s output moves off 0.38 per cent. Whether the digital ministry converts June’s embarrassment into a statute rather than a press release. And whether *yángliú* proves that a text feed can originate a political brand faster than a party can. Taiwan chose this square in 2024 partly because it felt clean. It will vote in it in November knowing it is not. That is not a worse position. It is an honest one, and more than most democracies can say.⁶

MERIDIAN

THE LAST WORD

Loaded, and not yet fired

Two hundred and ninety-four accounts were opened in May on the network where Taiwan now does its arguing. They post about romance. Between them they have made, as of the last published count, not one political statement. The election is on the twenty-eighth of November, and the researcher who has watched this before says six months is the usual lead time.

It was that kind of issue. A forty-two-fold speed-up sits finished and benchmarked in a personal fork, because a maintainer was annoyed. A pull request has one of the two approving reviews it needs. A pair of headphones went on sale here this morning with the feature that justifies the price shipped in early access. A chip company signed eight point two billion dollars of its own stock for a laboratory, subject to a regulator who has not spoken.

Capacity is cheaper than intent, and it arrives first. The question is never whether the thing can be done. It is who decides, and when, and whether anyone is watching the wall.

Twelve Desks · One World MERIDIAN · No. 117 · 29 September 2026

A note on method. This edition was researched with working network access and its load-bearing figures were read on the pages that carry them — the US Treasury's own par yield curve, mapped cell by cell against its header row; the Federal Reserve's statement of 16 September; AMD's investor release and Nvidia's filing; the JAMA Pediatrics paper in full, with its limitations and disclosures; the sports-medicine cohort study from the journal's own press copy; Doublethink Lab's audit and NewsGuard's investigation; the llama.cpp pull requests as HTML; Albin's and C. Tessile's corporate pages; and every Diary entry against a venue or official page. Where a host refused this desk — the Singapore broadsheets, Reddit everywhere, Permanent Style, the orienteering federation, Variety, Axios, the Guardian and CNBC, and the major wires to our search tools — it is named on the page where the gap falls, and the claim is flagged or dropped rather than dressed up. Two figures are printed as irreconcilable rather than averaged: the Hormuz transit counts, and the price of a Louis Vuitton watch. The Macro Desk is explanation, not investment advice; the Family and Fitness desks are general guidance, not medical advice.