

THE CURRENT BRIEFING OF

MERIDIAN

Twelve desks. One world.

WEDNESDAY 7 OCTOBER 2026



A DEFENCE PACT WAS ACTIVATED ON MONDAY. NO SOLDIER HAS YET BEEN SEEN (P3)

SINGAPORE

Four Unhealthy regions became three overnight — and the hourly reading says the plume has swung west

Twelve Desks · One World

THE ATELIER

The American binocular is built in Japan, the Japanese one in China, and the house that makes them names no client

Source: US Navy photograph by PH1 Kevin H. Tierney, public domain, via Wikimedia Commons — Prince Khalid bin Salman Al Saud, Saudi Minister of Defence, at the Pentagon, 24 February 2025. Specific subject; not the Riyadh meeting of 5 October 2026. Cropped; no other alteration.

THE LONG READ

Thailand rents five million citizens an AI subscription — then the company being paid per user stops publishing the count

THIS ISSUE · NO. 125

Contents

- p3–4** The World — a mutual-defence clause invoked, a secretary-general appointed, and a troop number its own government calls invented

p5 Singapore — the haze eases by one region, and the Act has four files open from 2015

p6 The Atelier — binoculars: two arcseconds, and the ten-second test you can do without taking the caps off

p7 Technology — two open-weight models, no weights

p8 The Home Lab — a cache that is 2.2×, and 0.35× if you size it wrong

p9 The Kit — forty-six decibels, measured by nobody

p10 The Good Life — \$26,000 for a worse movement

p11 Screen & Sound — twenty-one Metascore points, one thousand seven hundred and seventy-one dollars
- p12–13** The Family Desk — 5,515 step counters, an effect size of 0.08, and a panel that refused a ban

p14 The Macro Desk — oil down a quarter, the long bond up anyway

p15–17 The Rabbit Hole — home cheesemaking: the hobby whose forum is being spammed to death

p18 Fitness — the optimal dose, and the authors who say there isn't one

p19–20 The Diary — a marathon that sold out under a new name; and two “October openings” from 2025

p21–23 The Long Read — Thailand rents cognition by the head

p24 The Last Word

The spine: announced, not delivered — the gap between a declaration and the thing it describes, in eight desks

FROM THE DESK

“For a successful technology, reality must take precedence over public relations, for nature cannot be fooled.”

Richard P. Feynman, closing line of his personal appendix to the Rogers Commission report on the Challenger accident (1986)¹

He meant O-rings. It applies this morning to a pact, two model releases, a decibel figure and a usage count.

THE STRIP · AT MONDAY’S CLOSE

US 2-year	4.84%	▼ 8bp w/w
US 10-year	5.31%	▲ 7bp w/w
10-year real	2.95%	▲ 5bp w/w
VWRA	\$195.58	—
DXY	101.84	▼ 0.32%
USD/SGD	1.2769	▼ 0.21%
Brent	\$101.15	▲ 0.83%
Gold	\$4,165.30	—
Bitcoin (intraday)	\$85,541	▼ 0.21%

Treasury series at Monday’s close; spot and crypto are 6 October prints. Detail on **p14**. Not advice.²

THE ARCHIVE · WEDNESDAY — FOUR THINGS THIS MAGAZINE SAID IT WOULD COME BACK TO

- No. 124 reported a film fifth at eight million dollars** off Sunday estimates. The actuals put it **fourth, by \$1,771** (**p11**).

No. 124 printed that 91% of the ten-year’s rise was real yield. Three weeks on it is **106%** — breakevens are now pulling the other way (**p14**).
- The Deepavali date held.** It is **Sunday 8 November**, with Monday the 9th in lieu only where the rest day falls on the 8th (**p19**).
- And the Hormuz executive has added a fourth number** to the three that did not reconcile in No. 124: **nearly three billion barrels** of supply lost since February (**p4**).³

SOURCES ¹ Feynman’s sentence closes Appendix F, “Personal Observations on the Reliability of the Shuttle”, in the report of the Presidential Commission on the Space Shuttle *Challenger* Accident, 1986. ² Levels at the close of Monday 5 October 2026 and, for the spot and crypto rows, 6 October prints; every figure is sourced on **p14**, where the real-yield share is shown as this desk’s own derivation from three published series. VWRA’s daily change is disputed between two sources and is not printed. ³ Each item is checked against the primary source named on the page cross-referenced.

THE WORLD · RIYADH, MONDAY

A mutual-defence pact two months old has been activated for the first time, and given a Pakistani general as its first secretary-general. What it has not yet been given is a single observed soldier

Islamabad has publicly called the troop numbers “invented”. Read the denial closely: it disputes the figures, and not the deployment.

On Monday 5 October the Strategic Political and Defence Committee of the **Makkah Alliance for Defence** met in emergency session in Riyadh and resolved “**to move immediately to the practical implementation of the collective defence commitments... and ensure their rapid deployment in Saudi Arabia.**” The pact — Saudi Arabia, Türkiye and Pakistan, signed on 7 August — holds that an armed attack on any one of the three is an attack on all. It had never been invoked.¹

The institutional detail matters more than the language. The same meeting appointed **Lieutenant General (Ret.) Nauman Mahmood of Pakistan** as the alliance’s **first Secretary-General**, for three years from that day. A treaty acquired a standing command on the morning it was first used.¹

What triggered it was the holy cities.²

1.4m · 3,400 · 6,000

COMBINED ACTIVE TROOPS, MILITARY AIRCRAFT AND TANKS ACROSS THE THREE SIGNATORIES — ONE OUTLET’S OWN TALLY, NOT AN OFFICIAL FIGURE, AND PRINTED AS SUCH²

THE SCEPTIC · ISLAMABAD REJECTS THE NUMBER AND NOT THE THING

A wire service reported on 2 October that Pakistan had sent **30,000 to 40,000 troops** to Saudi Arabia for border defence and missile interception, short of combat in Yemen.³

THE CHATTER · THE REALIGNMENT NOBODY DISCUSSED

The largest security realignment in the Muslim world this year registered almost nowhere on the English-language technical internet. The wire story on Yemeni forces taking the strait drew **three points and no comments** on Hacker News; the original August signing of the pact, **ten points and no comments**. Searches for Hormuz, Brazil and the election in the same window returned **no stories at all**. The front page that week went to a French language model and a data-frame library.⁴

“Filming, publishing, or sharing information related to countering missiles... exposes you to legal liability.”

And one small notice that says more than the communiqué. On 6 October the Saudi Interior Ministry warned that filming or sharing anything about missile interceptions and their impact sites carries legal liability. No penalty was specified. A state does not prohibit the filming of interceptions that nobody is seeing.⁵

ARTICLE 5

UNOBSERVED

SINGLE-SOURCED TALLY

SOURCES 1 The Turkish Ministry of Foreign Affairs’ joint statement on the emergency meeting of 5 October 2026, read at source, for the quoted implementation language, the appointment and the three-year term. The Saudi Press Agency’s item rendered its headline but not its body to this desk, and Pakistan’s foreign ministry returned a service error twice; the full paragraph-by-paragraph text is therefore held and only phrase-level quotation is printed. 2 Al Jazeera, 5 October, read at source, for Prince Khalid bin Salman’s quoted remarks and for the combined force tally, which is that outlet’s own aggregate and not an official figure. Arab News, 1 October, read at source, for the interception counts, the debris at Khamis Mushait and the Madinah power-station strike; The National, 6 October, for the two airports and the injuries. 3 The troop figures were reported by a wire service this desk could not reach; Pakistan’s rejection is quoted from The Express Tribune, 6 October, read at source, so the denial is sourced and the claim it denies is not. The Houthi denial and the fact-checked video are from the same publication’s check of 16 September. Allison Minor quoted from an Atlantic Council dispatch of 7 August, read at source and dated here because it precedes the activation. 4 Story points and comment counts read through the Hacker News search API for 3–8 October. Reddit returned nothing on any route. 5 The National, 6 October, read at source, for the Interior Ministry notice and the absence of a stated penalty.

THE WORLD · FIVE CAPITALS, ONE MEETING

Doha reads it as the Gulf growing up, Beijing as a message rather than a war, Lyon as a nuclear power joining a Gulf army outside Western command — and New Delhi as a bill

The most useful framing for a Singapore reader comes from the capital with the same problem: two chokepoints, and no navy at either.

DOHA — AL JAZEERA

Frames the activation through the violation of the holy sites rather than through Iran, and lets Riyadh set the terms. Its analyst **Abdulaziz Alghashian** calls the moment **“a critical juncture in the history of the region”** and argues the Houthi strikes **helped** the alliance acquire its structure — an adversary doing an institution’s organisational work for it.¹

BEIJING — XINHUA

Datelines its analysis **Sanaa** and sources it to Yemenis rather than Western institutes.²

LYON — EURONEWS

The European read is arithmetic and uneasy: this bolts **NATO’s second-largest army** and **the Islamic world’s only nuclear-armed state** to the largest Gulf power, outside any Western command chain.³

NEW DELHI — THE WEEK

India reads it as a chokepoint problem, and the sentence travels: **“the two chokepoints hurt differently. A Hormuz closure cuts supply. A Bab el-Mandeb closure inflates cost.”** Its figures: **35% of India’s foreign trade** via Suez and the Red Sea, and crossings down from **30 to 15 ships a day** in September.⁴

HORMUZ, DAY TWO HUNDRED AND TWENTY · A FOURTH NUMBER FROM THE SAME EXECUTIVE

No. 124 set out Amin Nasser’s three irreconcilable inventory figures. He has now added a fourth, in London: **nearly three billion barrels of supply lost** since the conflict began at the end of February, with rebuilding stocks taking **as long as two years**. Meanwhile the measurement problem this desk flagged has a name.⁵

AROUND THE DESKS

Brazil goes to a run-off on 25 October. First round: **Flávio Bolsonaro 47.03%** to **Lula’s 45.16%**, turnout **78.92%**. The held thread from No. 123 closes.⁶

Federal forces retook Mekelle on Saturday, the TPLF saying it had **“temporarily moved out”**. The next day Egypt, Eritrea, Somalia and Sudan jointly declared that **control of the Red Sea should remain with the coastal states** — four heads of state answering one prime minister’s demand for sea access.⁶

France has arrested over 5,000 people in a week of school protests, 85% of them minors, with 400 to 500 schools closed. The budget behind it: **€65.5bn** proposed for education in 2027 against **€63.4bn** for defence.⁶

Germany has arrested its former intelligence chief. August Hanning, 80, once president of the BND, was detained on 6 October on suspicion of obtaining state secrets and treason-related espionage. No foreign state is named in the source this desk could read.⁶

Two Nobels. Medicine to **Deisseroth, Hegemann and Nagel** for light-gated ion channels and optogenetics; physics to **Francis Halzen** alone, for IceCube and the discovery of high-energy astrophysical neutrinos. He proposed the South Pole idea in **1988**; the cubic kilometre of instrumented ice was finished in **2011**.⁷

RUN-OFF 25 OCT

RED SEA

OPTOGENETICS

SOURCES 1 Al Jazeera, 5 October, read at source. 2 Xinhua’s news analysis datelined Sanaa, 6 October, read at source, for all three named analysts. 3 Euronews, 6 October, read at source. 4 The Week (New Delhi), 6 October, read at source; the trade, shipping and insurance figures and the Brent and import-bill projections are that publication’s own and are attributed rather than adopted. 5 Al Jazeera, 6 October, read at source, for Nasser’s remarks as reported, the bypass estimate, the strait’s recovery percentages and the tanker incidents; the 14.2m figure is from a commentary piece of 5 October, read at source, and is single-sourced. The reconciliation between the two sets of figures is this desk’s. 6 The Brazilian result from the electoral authority’s count as reported by Xinhua, 5 October, read at source; Ethiopia and the four-leader declaration from Al Jazeera, 4 and 5 October; the French figures are the interior ministry’s as reported by Euronews, 6 October, and differ from a same-day digest’s summary; the Hanning arrest from Euronews, 6 October. 7 Both citations read on the awarding academy’s own press releases.

SINGAPORE · THE HAZE, DAY ELEVEN

Four Unhealthy regions became three overnight — and the hour-by-hour reading says the plume has swung west. The law written for exactly this has four files open from 2015 and no charges

Eleven years of open investigation, and a professor of environmental law explains why: the Act was built to need Jakarta's cooperation, and Jakarta has never given it.

At five o'clock this morning the environment agency's own feed read **Central 129, East 118, West 114, North 99, South 96**. Three regions Unhealthy, against five at five yesterday afternoon and four at dawn. But the twenty-four-hour index is a trailing average, and the one-hour particulate reading underneath it has already moved: **West is now the dirtiest air in Singapore at 87 micrograms**, against Central's 74. The plume has swung.¹

The agency names the source without hedging: **"moderate to dense smoke plumes from southern Sumatra and West Kalimantan... drifting towards Singapore"**, with dry weather forecast to persist and south-easterlies to keep it coming.¹

4 files, 0 charges

SINGAPORE'S INVESTIGATIONS INTO FOUR INDONESIAN COMPANIES UNDER THE TRANSBOUNDARY HAZE POLLUTION ACT, OPENED AFTER THE **2015** EPISODE, REMAIN OPEN AND UNCHARGED **ELEVEN YEARS ON**²

THE SCEPTIC · THE ACT IS NOT UNUSED BY ACCIDENT

No. 124 printed that the Act has never been used. The better question is why, and it has an answer on the record.²

What the state is doing instead runs through **28 agencies across 13 ministries**. Every classroom from primary to junior college already has air purifiers. Community centres open as air-conditioned respite **only above PSI 200** — seventy points above this morning.³

SOURCES 1 The National Environment Agency's real-time PSI and PM2.5 feeds on data.gov.sg, read at source at the 05:00 reading of 7 October, timestamped 05:15; the quoted cause and the forecast from NEA's own haze situation updates of 5 and 6 October, read at source. The 7 October advisory had not been published at press. 2 An AFP report of 1 October, read at source, for the four open investigations and for both professors' quoted remarks and the structural obstacles. Single-wire-sourced; this desk did not obtain confirmation from the agency or the Attorney-General's Chambers. 3 NEA's whole-of-government haze measures release, read at source, for the agency and ministry counts, the school and workplace measures, the PSI 200 respite trigger and the mask guidance. 4 HardwareZone's haze discussion thread, read at source on 6 October; sensor readings are hobbyist devices reported by their owners and are labelled as such, not as measurements. Reddit returned nothing on any route. 5 Race dates and the sprint format from the championship calendar; the contingency-planning line and the no-cancellation sentence from a motorsport outlet of 6 October whose published index figures reconcile exactly with NEA's, read at source. The organiser's quote is single-sourced. 6 The health ministry's release of 3 October on the National Population Health and National Nutrition Surveys 2025, read at source, for every figure in both directions. 7 The finance ministry's release of 30 September, read at source, for the household count, the rebate bands and the quoted purpose.

THE CHATTER · THE READINGS PEOPLE TRUST ARE THEIR OWN

On HardwareZone's haze thread — four thousand replies, half a million views, moving by the minute — Singaporeans have built a parallel air-quality network and it disagrees with the official one.⁴

"There is currently no indication that this weekend's Grand Prix will be cancelled."

And the deadline nobody can move: **the Grand Prix runs 9–11 October**, a sprint format here for the first time, track action from Friday. Organisers say haze is in the contingency plan, with masks and medical support at the circuit. Marina Bay is a street circuit without a roof.⁵

ALSO THIS WEEK

The national health surveys say we are healthier, and the release does not dwell on the half that got worse. Activity is up to 87.3%, daily sugar down to 48g from 60g in 2018. But smoking rose to 8.8% in a year, sodium hit 4,000mg a day with 95% over the limit, and eating out daily went 60% to 70%.⁶

Over a million HDB households receive U-Save and service-and-conservancy rebates this month, **\$55 to \$95** by flat type, including a top-up announced as part of the second support package for the Middle East shock. Brent was \$130.80 on 15 September. It is \$101 now — see **p14**.⁷

PSI 129

RACE WEEK

ACT UNUSED

THE ATELIER • BINOCULARS

The American Razor is built in Japan and the Japanese Monarch is built in China. The one Japanese house that makes binoculars for the famous names has a website, a seventy-five-year history, and not one client's name on it

Two numbers decide a pair of binoculars and neither is on the box. Here is how to read an unfamiliar one in ten seconds, at arm's length, without taking the caps off.

Start with the fact that makes the others make sense. A **Porro** prism must be aligned to within **ten arcminutes**. A **roof** prism's critical angles must be right to within **two arcseconds** — three hundred times tighter. That ratio explains why a cheap Porro can be excellent and a cheap roof pair never is, and why the slim modern shape costs what it costs.¹

A roof prism also needs something a Porro does not. Its internal reflections impose a phase lag that degrades contrast and resolution *in one axis only*.¹

52% → 78.5% → 90%+

LIGHT THROUGH SIXTEEN AIR-GLASS SURFACES, UNCOATED, SINGLE-COATED, MODERN MULTI-COATED. ZEISS PUBLISHES **92%** FOR THE VICTORY SF 42MM; AN INDEPENDENT LAB MEASURED **91.2 ± 1%**²

Mirrors are blunter: aluminium reflects **87–93%**, silver **95–98%**, dielectric **over 99%**. An **Abbe–König** prism skips the mirror entirely, so in principle it is brighter.²

THE TEN-SECOND TEST, AND WHY “BAK-4” ON A BOX MEANS LITTLE

Hold any pair at arm's length against a bright sky and look at the two discs of light in the eyepieces. **Round is right. Square-shouldered means light is being thrown away** — clipped by undersized prisms, not by the glass type.²

Kamakura Koki of Warabi, Saitama has made binoculars since **1950** and calls itself an OEM

SOURCES 1 Prism alignment tolerances from the standard reference article on binoculars and, independently, from Holger Merlitz's own page on Porro prisms, both read at source; the two agree. The phase-lag mechanism, its 1988 invention at Zeiss and Merlitz's quoted sentence are from the same two sources. 2 Coating transmission arithmetic and the aluminium, silver and dielectric reflectivity bands from the same reference, read at source. Zeiss's published 92% from its own Victory SF page; the measured 91.2 ± 1%, the Vortex Razor UHD's 86.5 ± 1%, the Monarch 7's 89.1 ± 1.5% and every truncated-exit-pupil and country-of-manufacture finding from allbinos.com's measured review tables, each read at source. That publisher does not describe its transmission method on the page, which is noted rather than glossed. 3 Kamakura Koki's own English company and custom-manufacture pages, read at source, for both quoted sentences and the 1950 founding; the circulating client lists trace to BirdForum, which returned 403 on every route, and are not printed. German Precision Optics' quoted sentence and its Japanese manufacture from binocular.ch, read at source. Meopta's divisions and headcounts from its own site. 4 Singapore prices from Cathay Photo and Alan Photo, read at source on 6 October. 5 Pupil-diameter figures by age from allbinos' article on exit pupil, read at source. 6 The allbinos 10x42 ranking and the one-word cons list; Merlitz's Fujinon-versus-Sky Rover comparison and his United Optics review, both read at source; the CL Companion cons from its measured review. 7 Purging and gas-permeability from an optics dealer's own technical blog, read at source — vendor-published and single-sourced.

“providing products that do not damage the value of our customers' brands” for “famous camera manufacturers and major optical equipment manufacturers.” It names no client anywhere.³

THE TIERS, AND WHAT EACH STEP BUYS

The entry point of quality, S\$490–740. Nikon Monarch M5 8x42 at **S\$488**, M7 at **S\$738**, Zeiss Terra ED 8x42 at **S\$660**. The Monarch 7, made in China, measures **89.1%** against the grail's 91.2. Two points.⁴

And the reader is forty-five. Maximum pupil diameter is about 7mm to age 40–45 and after 50–55 **rarely exceeds 5mm**. An 8x42 already gives 5.25mm. Buying an 8x56 for low light in your forties is buying weight.⁵

THE SCEPTIC • THE GRAIL'S ONLY LISTED FAULT IS ITS PRICE

The measured ranking puts **Swarovski's NL Pure 10x42** first, and the entire cons list on its review is one word: “price?”⁶

NEXT ON THE BENCH

Espresso machines (the boiler, not the badge) — heat exchanger against dual boiler against lever, and where the curve goes flat.

Mattresses & slat bases — latex against pocket springs, and what tropical humidity does to both.

Cutlery & flatware (forged against stamped) — eighteen-ten against eighteen-zero, Sheffield against Tsubame-Sanjo, and the weight in the handle.

TECHNOLOGY · OPEN, BUT NOT DOWNLOADABLE

Mistral trained a trillion-parameter model on its own European GPUs and published eleven benchmark wins. An independent harness has already measured one of them five points lower

Both of this week's open-weight announcements ship no weights. The index that tracks frontier models currently files one of them under "proprietary".

The sovereignty claim in **Mistral Large 4**, published Monday, is not the parameter count. It is a sentence about geography: the model was **"trained from scratch on 3,800 NVIDIA Grace Blackwell GPUs in Mistral's own datacenters in Europe"**, and ships with **"a European deployment that Mistral operates end-to-end... under European law."** European lab, European silicon, European jurisdiction, and **160 languages** including every official tongue of the Union.¹

The machine is **one trillion parameters, 49 billion active**, priced at **\$1.36 per million input tokens and \$4.18 per million output**. The community nicknamed it *Le Chonk* within hours and Mistral adopted the name. The weights, the post says, "drop end of this month". No licence is named.¹

28.3 vs 22.73

TERMINAL BENCH 4.0 AS THE MAKER PUBLISHED IT, AGAINST WHAT **VALS AI** MEASURED ON ITS OWN HARNESS — **5.6** POINTS ON THE SAME NAMED BENCHMARK²

THE SCEPTIC · THE INDEPENDENT HARNESS DISAGREES ALMOST EVERYWHERE

Vals runs its own evaluations rather than reprinting a maker's table, and where the two overlap they diverge. It puts the model **32nd of 44 on its overall index, last of twenty-four** on MysteryMechanism, and **6th of 75** on Harvey's legal benchmark, which Mistral says it leads among open-source models.²

THE CHATTER · 1,388 POINTS, AND AN ARGUMENT ABOUT WHAT EUROPE IS FOR

The release took the top of Hacker News with **1,388 points and 871 comments**. **saberience**: "Looks like it's about a year behind still." **esafak**, answering the obvious question: "Mistral is one of the few European AI labs. Look up 'sovereign AI'." Per-comment scores are not exposed by the archive, so none is claimed.³

NEW YORK — REFLECTION AI

The week's other open-weight announcement, **Beam**, is **501 billion parameters, 23 billion active, Apache 2.0**, pretrained on **6,144 GB300 GPUs over 23.8 trillion tokens**, then reinforcement-trained through **"over 100 million rollouts"**. Read down its own table and the Chinese open-weight models take most agentic rows: **SWE-Bench Pro v2-Hard 77.2 against Kimi K3's 88.2**.⁴

"Couldn't find any download option. Strange way to distribute an open-weights model."

SINGAPORE SILICON · TWO NUMBERS, ONE SOURCE

VSMC, the VIS-NXP joint venture, has finished its first 12-inch wafer fab here in **22 months**. Mass production starts **Q1 2027**, targeting **44,000 wafers a month by 2029**.⁵

ASE Technology has bought a **449,340 sq ft** plant here for about **US\$68.44m** for AI-chip testing and packaging. Both figures reach this desk through one trade outlet and neither company's newsroom, and are printed as single-sourced.⁵

SOVEREIGN AI

WEIGHTS PENDING

SINGLE-SOURCED

SOURCES 1 Mistral AI's own announcement post, read at source on 6 October, for the parameter counts, the training-fleet and European-deployment sentences, the language count, the prices, the weights timing and every benchmark figure attributed to the maker. 2 Vals AI's model page for Mistral Large 4, read at source, for the Terminal Bench, index, MysteryMechanism and Harvey figures; Artificial Analysis's model page, read at source, for the Intelligence Index rank and the output-token comparison; the Surge AI row is in Mistral's own table. 3 Hacker News story 49977979; points and comment count read through the site's search API, comments attributed by username. The API returns null for per-comment scores, so none is given. 4 Reflection AI's own Beam post, read at source, for the parameter counts, licence, training fleet, token and rollout counts and the comparison table; the quoted line is from Hacker News story 49969183 (531 points, 166 comments). 5 A DigiTimes weekly roundup of 5 October, read at source. **Single-sourced**: this desk did not reach VSMC's or ASE's own announcements and says so rather than implying a second check. **Refused this desk**: GitHub's API and datacenterdynamics.com.

THE HOME LAB · THE CACHE THAT IS 2.2× AND ALSO 0.35×

A pull request keeps the experts in host RAM and caches them on the card. On an RTX 5090 it runs 2.20 times faster — and a third as fast as baseline if you size the cache wrong

It is not merged. The headline multiple is the author's own figure, and the one tester who looked closely reports the output diverging from the uncached path.

The machine this matters to is the one with a 32-gigabyte card and a 93.7-gigabyte model. Pull request 29887, opened on 3 October by the contributor **aml17an**, puts a mixture-of-experts model's experts in host memory and keeps a least-recently-used cache of them in VRAM, so only a cache miss has to cross the PCIe bus. Its limits are stated plainly: batches of **32 tokens or fewer**, one device only.¹

On **Qwen3.8-Flash-Next** at Q4_0 — 93.7 GiB, of which 65.4 is experts — the author's table reads: RTX 5090 baseline **30.8 tokens a second**; with every expert offloaded behind a **19GB cache**, **67.8**. That is **2.20×** at an **89% hit rate**. A 4090 goes 25.0 to 40.7, and a Radeon AI PRO R9700 owner reported **11.0 to 27.6** independently.¹

20.9 → 7.30

TOKENS A SECOND ON A 4090 WHEN THE CACHE IS SET TO 1,024 MiB INSTEAD OF NINETEEN THOUSAND — THE SAME PATCH, **A THIRD OF BASELINE SPEED**, REPORTED IN THE THREAD BY A TESTER THE AUTHOR HAS SINCE ANSWERED¹

THE SCEPTIC · NOT MERGED, AND NOT BYTE-IDENTICAL

The tester **antomicblitz** found the cliff above, and something worse: **"greedy output also diverges from the non-cache path at the same position in both runs, so it is not a byte-identical cache."** A cache that changes the answer is not a cache.¹

The same thread explains why half of llama.cpp looks the way it does.¹

"The spills are real — I just missed the kernel in the list."

A second patch makes the point without the ambiguity. **PR 29910** claims to eliminate register spills on Q2_K and reports **+168.72%** on an older AMD card. A maintainer measured **0.75×** on an MI100 and **0.68×** on a P40; a third tester got **0.71–0.80×** on MI100, **3.79×** on a 7900 XTX at one batch size and **0.88×** at another.²

What did merge on Monday is **PR 26610**, tensor-split inference over RPC — **619 tokens a second of prefill** across two DGX Sparks on an RDMA fabric. On gigabit Ethernet a tester measured **8 tokens a second of prefill against layer split's 100**.³

THE CHATTER · THERE WASN'T ONE

All of this happened in public, on a repository with tens of thousands of watchers, and none of it reached the wider developer internet. **llama.cpp has not had a Hacker News thread above six points since 26 September**. The week's front page went to **Polars 2.0** (379 points) and a **Gleam compiler change** (265).⁴

WORTH INSTALLING THIS WEEK

EmbeddingGemma 2 — 740M parameters, modular down to 270M for text alone, **Apache 2.0, 8K context**. Quantised it wants about **191MB of active RAM** text-only, **567MB** with vision and audio, takes **5.5 minutes** of audio in context and runs offline.⁵

UNMERGED

RTX 5090

AUTHOR-REPORTED

SOURCES 1 Pull request 29887 on the llama.cpp repository, read at source on 6 October: the description's own benchmark table for every tokens-a-second figure, hit rate and multiple; the third-party AMD figures, the 1,024 MiB cliff and the divergence report from named testers in the thread; the force-push timestamp and the author's reply from the thread. The patch was OPEN and unreviewed at press, with the maintainer self-assigned the same day. 2 Pull request 29910, read at source: the author's claimed percentages and spill counts, and the contradicting measurements posted by a maintainer and one further tester on 5 and 6 October. 3 Pull request 26610, read at source: merged 6 October; the author's RDMA figures and the Metal and gigabit-Ethernet counter-measurements are all from that thread. 4 Story points and comment counts read through the Hacker News search API, including a dated search for llama.cpp back to 26 September. Reddit returned nothing on any route from this environment. 5 Google's own EmbeddingGemma 2 announcement, read at source, 6 October, for the parameter counts, licence, context length, memory footprints, the MTEB Code figures and the Matryoshka storage claim; the llama.cpp support PR is 30054, merged the same day. **Note:** GitHub's API is blocked from this environment, so every pull request was read as a web page.

THE KIT · THE NUMBER NOBODY HAS MEASURED

Nothing's flagship headphone has a Singapore price at last: \$549. It also has a claim of “up to 46 dB” of noise cancellation that no laboratory on earth has yet checked

The maker's own Singapore page says Pre-order and names no dispatch date. Four issues after launch, three of the review labs that would settle the question return 404.

The peg this desk named in No. 123 has landed. The **Nothing Headphone (1) Pro** now carries a Singapore price on the maker's own storefront: **S\$549**, against US\$399 and £349 at launch. The button beside it reads **Pre-order**, and the page states no dispatch date, no GST note and no delivery term — absences verified on the page, not inferred.¹

The specification is genuinely ambitious: **327 grams**, a triple-driver array pairing a 40mm bass driver with an xMEMS treble unit, **Bluetooth 6.1**, IP52, **68 hours** of playback with cancellation off and **30** with it on, from ten microphones rated at “**up to 46 dB**”.¹

0 of 3

OF THE MEASUREMENT LABS A BUYER WOULD CONSULT — RTINGS, WHAT HI-FI, SOUNDGUYS — HAVE PUBLISHED A REVIEW. ALL THREE RETURN **404**. THE 46 DB AND THE 68 HOURS REMAIN THE MAKER'S OWN FIGURES²

THE SCEPTIC · THE CHEAPER ONE MAY BE THE ANSWER

Adam Doud of SlashGear recommends it and then takes it apart. The new pads are softer, and “**the isolation they provide isn't as good — it's a sacrifice of function over comfort.**” On the tuning: “**with the flat EQ enabled, I found the midrange to be emphasized more than I'd normally like.**”³

THE CHATTER · THERE IS NONE, AND THAT IS THE FINDING

This desk went looking for owner sentiment and found a wall. Hacker News returns **zero results** for the product. **Head-Fi** served a security challenge rather than a thread. **Reddit** refused the request outright. So there is no attributable community verdict on a flagship headphone eight days after sale, and this page prints that rather than a sentence of invented consensus.²

“A parametric equalizer is a lot more complicated than something like a ten-band EQ.”

The honest position, then. The Singapore price is real and single-sourced to the maker; no local retailer confirms it. The design claims are plausible and unverified. A buyer who waits a fortnight loses nothing but the queue, and gains the first independent measurement of the one number the product is sold on.¹

STILL ON THE SHELF

Ray-Ban Meta Audio, the camera-free pair, ships **13 October**. Two outlets give two prices — \$329 and \$349 — and this desk reached neither the maker's newsroom nor a confirming page, so no price is printed.⁴

Garmin's Enduro 4 and Approach S72 both launched on 6 October. The reader asked this desk to drop one Garmin product permanently and it has; the brand is not covered here by default, and these two wait for a reason to run.⁴

And the held thread from No. 124 stays held: no Singapore price has appeared for the Huawei FreeBuds Neo. Fifth issue of asking.

S\$549

PRE-ORDER

UNMEASURED

SOURCES 1 Nothing's own Singapore storefront page for the Headphone (1) Pro, read at source on 6 October, for the price, the pre-order state, every published specification and the absence of a dispatch date. The S\$549 figure appears nowhere else this desk could reach and is printed as single-sourced. The US and UK launch prices are from No. 123. 2 The review pages at RTINGS, What Hi-Fi and SoundGuys were each requested and each returned 404 — a genuine absence, not a block. Hacker News was queried through its search API and returned no matching story. Head-Fi served a security interstitial and Reddit refused every route; no sentiment is claimed from either. 3 Adam Doud's review for SlashGear, 28 September, read at source; every quoted sentence is his. 4 The Ray-Ban Meta Audio ship date and the two conflicting prices from two roundups, neither confirmed at the maker; the Garmin launch dates from a specification outlet, read at source. The standing instruction to drop the Garmin CIRQA is the reader's, recorded 3 August.

THE GOOD LIFE · MOUNT JOY, PENNSYLVANIA

RGM's first perpetual calendar since the nineteen-nineties costs US\$39,900. The same house will sell you the same case, with a better movement and a longer reserve, for thirteen thousand nine hundred

The \$26,000 between them buys fired enamel and a calendar module. It also buys a base movement that is a step down from the cheaper watch's.

The Model 501-PC arrived on Tuesday from America's most stubborn independent watchmaker, and it is a handsome thing: **39mm by 11.5mm** of American-made 316L steel, **48mm lug to lug**, under a raised box crystal, with a **white Grand Feu fired enamel** dial, blued steel pear hands, Roman numerals and a red five-minute track. Day, date, month and leap year on three sub-dials, a moonphase at six. Price **US\$39,900**, confirmed on the maker's own page and in the one outlet that covered it.¹

Inside is an **ETA 2892-A2** — a commodity Swiss automatic, rhodium-finished and fitted with a solid-gold rotor — carrying a bought-in **Dubois Dépraz** perpetual-calendar module. Twenty-one jewels, 28,800 beats, **52 to 56 hours** of reserve.¹

\$26,000

THE GAP TO RGM'S OWN **501 TEDDY BALDASSARRE EDITION** — SAME 39MM STEEL CASE, SAME LUG SPAN, AN IN-HOUSE-FINISHED SCHWARZ ETIENNE MICRO-ROTOR AND AN **86-HOUR** RESERVE, AT **US\$13,900**²

THE SCEPTIC · THE ARGUMENT IS IN THE MAKER'S OWN CATALOGUE

Set the two beside each other. For **\$13,900** RGM sells a 39mm steel watch with a **micro-rotor movement and eighty-six hours** of reserve under a hand-cut guilloché dial.²

One oddity deserves its line. The maker's own product page lists the functions as "**hour, minute,**

second, and date (on select models)" — boilerplate that does not mention the perpetual calendar at all. The specification block for the headline complication does not describe the headline complication.¹

THE CHATTER · THREE DOORS, ALL SHUT

Hacker News has **no result** for the watch. **Watchuseek** redirected to a metering host; **Hodinkee** refused the request. No attributable collector sentiment exists for this reference yet, which is unsurprising at forty-eight hours and worth saying rather than filling.³

"Contact RGM." — the whole of the availability note

ELSEWHERE ON THE BENCH

Depancel's Serie-R Chronograph Skeleton opens for pre-order **8 October at 16:00 Paris**, 100 numbered pieces, delivery mid-December: a fully skeletonised column-wheel, vertical-clutch automatic chronograph in 40mm steel, **€1,595 / US\$1,595**. The identical euro and dollar figures are their own small tell about how these prices are set. Single-sourced to one outlet.⁴

Three salerooms open this week — Phillips in New York closing Friday, Sotheby's Paris from today, Bonhams New York from Thursday. Every auction-house page this desk tried refused it, and no result exists before this issue goes out, so no estimate and no lot is printed. Results in a later edition.⁴

And still held: no Singapore stockist for RGM could be found, and the Omega thread from No. 124 is unchanged — no Singapore price published.

MOUNT JOY

GRAND FEU

NO SECOND OPINION

SOURCES 1 RGM's own product page for the Model 501-PC, read at source on 6 October, for the price, case and movement specification, the power reserve, the dial description and the quoted functions line; the complication layout is confirmed against the maker's own press photograph. 2 Monochrome Watches, 6 October, read at source, independently carrying the US\$39,900 price and naming the module maker; Monochrome's January piece on the 501 Teddy Baldassarre Edition, read at source, for that watch's case dimensions, movement, reserve and US\$13,900 price. The comparison is this desk's, drawn from both published specifications. 3 Hacker News queried through its search API with no result; Watchuseek redirected to a metering host and Hodinkee refused the request. No sentiment is claimed from either. 4 Monochrome, 6 October, for the Depancel specification, pre-order time and price — single-sourced, the maker's own site not reached. The three sale dates are from trade coverage; every auction-house page returned 403 or 404 to this desk, so no estimate is printed.

SCREEN & SOUND · THE MARGIN

A film critics scored 46 beat a film they scored 67 by one thousand seven hundred and seventy-one dollars — and the estimate this desk printed on Tuesday had the result the other way round

No. 124 reported fifth place off Sunday estimates. The actuals moved it to fourth. Twenty-one points of critical opinion, two hundredths of one per cent of money.

The weekend of 2–4 October closed with the two films this desk paired on Tuesday separated by less than the cost of a week's electricity in one cinema. **Digger**, Iñárritu's, took **\$8,001,856**. **Primetime**, A24's, in its second weekend, took **\$8,000,085**. The gap is **\$1,771** — **0.022%**. Two independent charts agree to the dollar.¹

And the correction. No. 124 printed **Digger** fifth at \$8.0m off Sunday estimates; on actuals it finished **fourth**. The desk also printed **Verity** at **\$32.6m**; the real figure is **\$32,031,011**, about **\$569,000** lower. Sunday estimates are a forecast wearing a result's clothes.¹

427 screens

HOW MANY MORE CINEMAS **DIGGER** NEEDED TO WIN BY \$1,771 — **3,321 AGAINST 2,894**. ITS PER-SCREEN AVERAGE WAS **LOWER**: \$2,409 AGAINST \$2,764¹

On Tuesday the two films were one Metascore point apart. They are not any more. Current scores: **Digger 46** from 58 critics, **Primetime 67** from 48. The critical gap widened to **twenty-one points** in the same days the money gap collapsed to nothing.²

THE SCEPTIC · A 46 IS NOT A MEDIOCRE FILM, IT IS TWO FILMS

Digger's 46 is built from **16 positive**, **27 mixed** and **15 negative** reviews — a bimodal distribution, not a consensus. **David Rooney** in *The Hollywood Reporter* scored it **20**: it is “so meta... it disappears up its own ass”.²

THE CHATTER · WHERE AUDIENCES AND CRITICS PART COMPANY

The only attributable audience signal is the user data on the same pages, and it inverts the critics twice. **Verity** holds a Metascore of **45** and a user score of **4.5** from 43 ratings, **40% of them negative** — critics and audiences agreeing for once.²

“A grandiose folly of a film, a failure on its own perhaps impossible-to-satisfy terms.”

The weekend's quietest number belongs to neither. **Drishyam: The Conclusion** opened ninth on **541 screens** for **\$2,715,000** — a per-screen average of **\$5,018**, the second-highest in the top ten and more than double **Digger's**.¹

THE ALBUM CHART · A FIFTY-ONE-WEEK-OLD RECORD HOLDS OFF TWO DEBUTS

Taylor Swift stayed at No. 1 in Britain for the chart week to 8 October — **week fifty-one** on chart. What is new is what she held off: **two straight-in top-three debuts in the same week**, **Keo** at 2 and **Nothing But Thieves** at 3.³

And the catalogue number that should not be possible: three **Oasis** albums sit in the top sixteen, two of them past **660 weeks** on chart — *Time Flies* at six in week **682** and *Morning Glory* at nine in week **769**. **Abba's Gold** is at twenty in week **1,273**.³

Held: no units figure exists — the chart body publishes none. The US **Billboard 200** is behind a paywall this desk could not pass, and is not reconstructed from memory.³

\$1,771

BIMODAL

CORRECTION

SOURCES 1 Box Office Mojo's domestic weekend chart for the weekend of 2–4 October 2026 and *The Numbers*' chart for the same weekend, both read at source on 6 October; every gross, theatre count and per-screen average above appears identically in both. The estimates corrected here were printed in No. 124. 2 Metacritic's pages for **Digger**, **Primetime** and **Verity**, read at source on 6 October, for the Metascores, the critic counts, the positive/mixed/negative splits, the user scores and every quoted review with its published grade. *Hacker News* was queried through its search API and returned no in-window story; *Reddit* refused every route. 3 The Official Albums Chart for the week of 2–8 October 2026, read at source, for positions, debut status and weeks-on-chart. The chart body publishes no sales or streaming figure; *billboard.com* redirected to a metering host that returned a payment-required response, so no US chart is printed.

THE FAMILY DESK • FIVE THOUSAND FIVE HUNDRED TWELVE-YEAR-OLDS

A step counter strapped to 5,515 children found more walking, better mental health and measurably thicker cortex in two places. The effect size is 0.08, and that is the whole argument

The mean age in the study is 11.93 — the exact band that matters here. What the finding cannot do is tell you which way the arrow points.

Most research about children at this age measures screens. This one measured legs. Published in *JAMA Network Open* on **1 October**, it took **5,515** adolescents from the long-running American brain-development cohort, put a wrist tracker on each for about three weeks, and followed their mental health, their school grades and their brain scans. **Mean age 11.93 years. Mean daily steps 9,595.**¹

More steps went with fewer mental-health problems: **B = -0.080**, confidence interval -0.109 to -0.052. Children taking **12,000 steps a day or more** scored **0.102 points** better than those under 5,000. The authors found no threshold — their phrase is that **every step counts.**¹

B = 0.057

THE STANDARDISED ASSOCIATION BETWEEN DAILY STEPS AND **PARACENTRAL CORTICAL THICKNESS**; THE ROSTRAL ANTERIOR CINGULATE CAME IN AT 0.052. THE MEDIATION CARRYING STEPS THROUGH CORTEX TO MENTAL HEALTH: AN INDIRECT EFFECT OF **-0.002**¹

DON'T BE ALARMIST • WHAT AN INDIRECT EFFECT OF 0.002 CAN CARRY

Read those numbers again before acting on them. The mediation — the part that says walking changes the brain and the brain changes the mood — is an indirect effect of **-0.002**, and the authors themselves warn that selecting the mediators in the same sample **may have inflated it**. The design is observational, so the arrow could run backwards: a child who is already low moves less. Steps capture ambulatory movement and miss the hard running and the organised sport.¹

Two days later a different literature went the other way with far more confidence.²

THE CHATTER • SIX RESEARCHERS ON THE MORNING THE SERIES LANDED

Britain's Science Media Centre put the Series to independent experts on 4 October and the responses are sharper than the papers.³

“Calling for precautionary action is a legitimate position, but it is a value judgement, not a scientific finding.”

There is a serious reply to that, and it should be printed too. Writing in *Paediatrics & Child Health* last month, **Vaillancourt and Brittain** invoke Geoffrey Rose's prevention paradox: **“a large number of people exposed to a low risk is likely to produce more cases than a small number of people exposed to a high risk.”** When over **90%** of a country's twelve-to-seventeens are on social platforms, a tiny effect size is still a public-health problem.⁴

MEAN AGE 11.93

OBSERVATIONAL

NOT MEDICAL ADVICE

SOURCES 1 Li, Paluch and Rakesh, “Physical Activity and Adolescent Brain Structure, Mental Health, and Academic Outcomes”, *JAMA Network Open*, 1 October 2026; full text read at source, for the sample size, mean age, mean step count, every coefficient and interval, the categorical comparison, the mediation estimate and all of the limitations quoted in the sceptic box, which are the authors' own. Point estimates for the intermediate step bands sit in supplementary tables this desk could not open and are not printed. 2 Holly and colleagues, “The public health challenges of children's and adolescents' digital media use”, *The Lancet*, 4 October 2026; the journal's own site refused this desk, so the abstract was read through Europe PMC and only sentences present in it are quoted. 3 The Science Media Centre's expert reaction of 4 October, read at source, for all three researchers' quoted remarks and their stated interests. These experts were responding to the Series, not to the step-count study, and are not presented as commenting on it. 4 Vaillancourt and Brittain, *Paediatrics & Child Health*, 16 September 2026, read at source for the quoted prevention-paradox argument and the participation figure. The population-attributable fractions were not obtainable and are not quoted. **General guidance, not medical advice.**

THE FAMILY DESK • THE SINGAPORE LINE

Three of the six authors of the week's global paper on children and screens work at a school of public health in Kent Ridge. Singapore's own expert panel reached the opposite conclusion to Australia's, and said so first

The panel rejected a blanket age ban on the ground that children would simply get round it — and told the platforms to change instead.

The *Lancet* Series paper on children and digital media carries six named authors on its children's paper, and **three of them are at the Saw Swee Hock School of Public Health** at the National University of Singapore, one jointly at Duke-NUS. The global framework for thinking about this was partly written here, which is worth knowing before reading what Singapore then did with it.¹

In July the Ministry of Health's **Expert Panel on the Mental Health Impact of Social Media Use**, convened the previous September and co-chaired from the Institute of Mental Health and NUS, delivered five recommendations. The headline one is a refusal.²

1 in 6

YOUNG PEOPLE AGED **10 TO 24** HERE SHOW SIGNS OF PROBLEMATIC SOCIAL-MEDIA USE, ON THE PANEL'S OWN FIGURES; SEPARATELY **ONE IN THREE AGED 15 TO 17** SPEND MORE THAN THREE HOURS A DAY ON IT — A BAND THAT DOES NOT CONTAIN A TWELVE-YEAR-OLD²

THE SCEPTIC • A POLICY WITH NO PUBLISHED OUTCOME MEASURE

No. 124 noted that the schools' phone rule is three terms old with no defined outcome. Nothing has changed, and the same gap now runs through the panel's own work: there is no published metric by which any of the five recommendations will be judged to have worked. A precautionary approach that cannot be falsified is a posture, not a policy — which is Professor Woodward's objection on **p12** arriving from a different direction.²

“Whether online time helps or harms depends more on content, context, and user characteristics than screentime.”

TWO FINDINGS A PARENT CAN USE

Evening movement did not wreck sleep. A study of 1,175 children and adolescents using accelerometers and the American youth fitness survey found that higher movement intensity on a given day went with earlier sleep onset and longer sleep that night — and that the association close to bedtime was stronger in adolescents.³

It depends on the child, measurably. In 150 young people aged 7 to 17, longer screen time went with weaker inhibitory control — but **only in those with lower white-matter integrity** in one tract. Cross-sectional, small, and the screen time was reported by caregivers. Still: the same exposure, different children, different result.⁴

General guidance, not medical advice. Persistent low mood, lost sleep or school refusal is a reason to see a doctor, not to adjust a step target.

AGE 12

NO BAN

NO OUTCOME MEASURE

SOURCES 1 Author affiliations for the *Lancet* Series children's paper read through the Europe PMC record for that article; the journal's own site refused this desk. 2 The Ministry of Health's release on the Expert Panel on the Mental Health Impact of Social Media Use, 30 July 2026, read at source, for the co-chairs, the convening date, both prevalence figures and all five recommendations as quoted. The 15-to-17 figure is noted as covering a band that does not include a twelve-year-old. 3 Bourke, Verswijveren and Hesketh, *Sleep Medicine*, 26 September 2026, abstract read at source; the sample's mean age is nine with a wide spread, which is why only the direction is printed. 4 Jang and colleagues, *Developmental Cognitive Neuroscience*, 16 September 2026, abstract read at source; no effect sizes are published in it and none is printed. **General guidance, not medical advice.**

THE MACRO DESK • THE REGIME

Oil has fallen twenty-three per cent in three weeks and the long bond sold off anyway. The real-yield share of the move is no longer ninety-one per cent — it is a hundred and six

Inflation expectations have stopped explaining the sell-off and started working against it. What is left is supply — and it is repricing equities and bonds in the same direction at once.

No. 124 printed that **91%** of the ten-year's rise was real yield. Three weeks on, the decomposition has gone past its own limit. Since **15 September** the ten-year has risen **31bp** — and the real yield **33bp**, to **2.95%**, while the breakeven fell 2bp to **2.36%**. The real share is **106%**: inflation expectations are not failing to explain this move, they are pulling the other way.¹

THE STRIP • AT MONDAY'S CLOSE

US 2-year	4.84%	▼ 8bp w/w
US 10-year	5.31%	▲ 7bp w/w
10-year real	2.95%	▲ 5bp w/w
VWRA	\$195.58	—
Bitcoin (intraday)	\$85,541	▼ 0.21%

Treasury and breakeven series at Monday's close; spot and crypto are 6 October prints. VWRA's daily change is disputed between two sources and is not printed. Market explanation, not advice.¹


The proof sits in the oil price. **Brent fell from \$130.80 on 15 September to \$101 on Tuesday — down 22.7%** as Gulf exports recovered and reserves were released. The largest disinflationary shock in a year landed, and the thirty-year went up **37 basis points** through it. A bond market trading inflation does not do that.²

+38bp

THE RISE IN THE NEW YORK FED'S TEN-YEAR TERM PREMIUM OVER THE THIRD QUARTER, **0.5062 TO 0.8850**. THE EXPECTED-RATES LEG ROSE 40BP. OF Q3'S 78BP, ROUGHLY HALF IS THE PRICE OF ISSUANCE ITSELF¹

SOURCES 1 The St Louis Fed's own data service, downloaded as a series file on 6 October, for the two-, ten- and thirty-year Treasury yields, the ten-year inflation-indexed yield, the ten-year breakeven, the federal funds target range and the Singapore and Japan exchange rates. The 106% real share is this desk's own division of the published real-yield change by the published nominal change since 15 September. The New York Fed's term-premium workbook, downloaded and converted locally, for the quarter-end values; the correction issued in No. 124 against a circulating figure of 1.45 stands. VWRA's level is confirmed on two sources and its daily change is disputed between them, so only the level is printed. The global bond proxy's net asset value, weekly range, year-to-date return and yield to maturity are from the issuer's own product page. 2 Brent and West Texas prints from two independent commodity services for 6 October; the 15 September high from the St Louis Fed series. 3 Nicky Shiels quoted in The Globe and Mail, 1 October, read at source. Central-bank and fund flows from the World Gold Council's Q2 2026 demand report, read at source, which also carries the 29 January record of US\$5,405; one outlet gives US\$5,595 for the same record and the Council's own figure is used. 4 Gold spot from a price service timestamped 6 October; bitcoin from a market data service's spot endpoint, same day. Explanation and education, not investment advice; forecasts are flagged as forecasts. 5 Exchange-rate moves derived from the St Louis Fed series; the index level and moves from one market data service, single-sourced; the policy dates and both 2026 decisions from the central bank's own advance release calendar and decisions page, read at source. **Held:** the weekly dollar-index and index changes, the daily term premium, and the August core inflation print, whose publisher refused this desk.

THE MOVE SINCE 15 SEPTEMBER, IN BASIS POINTS



Real yield accounts for 106% of the nominal move. Inflation expectations now subtract from it.

Meridian original, computed from the St Louis Fed's published ten-year nominal, inflation-indexed and breakeven series, 15 September to 5 October 2026.

THE SCEPTIC • A STRATEGIST WHO SAYS REAL YIELDS NO LONGER RULE

Nicky Shiels, metals strategist at MKS PAMP, argues the opposite of this page: **"the real-yield framework still anchors fair value, but since 2022 a persistent, structurally higher premium driven by reserve diversification and geopolitical hedging has become the dominant driver that macro factors no longer capture, and it is not fading."** Her mechanism is real — central banks bought a net **289 tonnes** of gold last quarter, up **62%** year on year, while funds sold **44.8 tonnes**.³

WHAT IT MEANS FOR VWRA, GLOBAL BONDS, GOLD AND BITCOIN

Diversification is not working, and the reason is that both legs face the same force.⁴

BEAR STEEPENER

REAL YIELDS

SUPPLY, NOT PRICES

THE RABBIT HOLE · HOME CHEESEMAKING

Making curd takes thirty minutes and a lemon. The hobby begins at the other end, in a cold damp box, where you are not cooking but farming a microbial community and losing arguments with it

The most-followed voice in the hobby spent last August telling three hundred and fifty-five thousand subscribers they do not need the expensive thing.

Milk plus acid, or milk plus rennet, makes curd. That part is genuinely easy and genuinely fast — a same-day ricotta, a paneer, a thirty-minute mozzarella. Everything that makes this a hobby rather than a recipe sits on the other side of a word: **affinage**, the controlled rotting of a wheel at **10 to 13°C and 80 to 95% humidity** for three months to three years.¹

The ladder the community itself teaches runs: acid-set cheeses in week one; rennet-set with a short age, which **“introduces curd cutting and stirring”**; then the technical cheddars; then washed-curd Goudas and Havartis, which **“require curd-washing techniques”**; and at the far end an eighteen-month Alpine wheel.¹

£5 → gold

MAT LLOYD WAS GIVEN A FIVE-POUND CHEESEMAKING KIT BY HIS SISTER-IN-LAW ONE CHRISTMAS. FOUR YEARS LATER HE TOOK A GOLD MEDAL AT THE **2023 INTERNATIONAL CHEESE AND DAIRY AWARDS**. SINGLE-SOURCED, AND THE TRADE PRESS'S OWN FAVOURITE GATEWAY STORY¹

The moment people describe is always small. **Paula Butler** visited a Virginia creamery in 1993: **“it was magical”**; she went professional twelve years later. Hunter, asked what she bought first: **“my first big investment was a thermometer.”** On the forum the hook is cruder and better — a member three batches in reported going from two pounds of cheese per two gallons of milk to **“basically three pounds”** after one change in technique.²

THE SCALE OF IT, COUNTED RATHER THAN ESTIMATED

CheeseForum.org holds **133,754 posts across 14,538 topics** going back to January 2008. **Gavin Webber**, the Australian who is the hobby's closest thing to a famous person, has **355,000 YouTube subscribers and 984 videos**. Two private Facebook groups carry **53,500 and 46,100 members**.³

“Busting the cheese cave myth: seven beginner cheeses you can make without a cheese fridge.”

THE SCEPTIC, EARLY · THE BIGGEST CHANNEL ARGUES AGAINST THE UPGRADE

In **August 2025** Webber published a post whose entire thesis is that the expensive thing is unnecessary: **“there’s a whole world of cheeses that don’t require a dedicated cheese cave... your kitchen fridge is more than enough to get started.”** His seven are ricotta, paneer, quick mozzarella, queso fresco, goat’s-milk feta, Bel Paese and halloumi. He also sells kits, courses and controllers.⁴

What follows over two pages: where these people actually gather and what they argue about; the gear, from a ten-dollar sachet of mesophilic culture to a **four-thousand-dollar pneumatic press**; the vocabulary and the real money; ninety honest days; and the question of doing any of it at thirty degrees and eighty-two per cent humidity.

AFFINAGE

355,000 SUBSCRIBERS

BUY LESS

SOURCES ¹ *Culture* magazine’s survey of home cheesemaking, 5 January 2026, read at source, for the difficulty ladder, the quoted technique notes and the quoted remarks of Caitlin Hunter, Paula Butler and Mat Lloyd; the Lloyd gold medal is single-sourced to that piece, the awards’ own site having returned an error. The temperature and humidity band is the board’s stated consensus. ² CheeseForum thread 15203, read at source, for the yield report and its date. ³ CheeseForum’s own statistics page, read at source on 6 October, for the post, topic and date figures; the subscriber and video counts read from the channel’s own page the same day; both Facebook group sizes are reported by *Culture* and could not be checked at the platform, which served a login wall; New England Cheesemaking Supply’s own site for its founding and its claims. Reddit refused every route. ⁴ Gavin Webber’s own post of 13 August 2025, read at source, for the title, the quoted sentence and the list of seven.

THE RABBIT HOLE · THE COMMUNITY, AND THE GEAR

The best free technical library this hobby has is being posted into, every six minutes, by a bot advertising a dating site — inside a board called “New website owner”

The conversation has moved to Facebook and YouTube comments. What is left behind is fourteen thousand threads of genuine chemistry and an intermittent server error.

CheeseForum.org has run since January 2008 and its structure says who is there: nine and a half thousand posts in Introductions alone, then boards for every cheese family and for equipment, forming, ageing and ingredients. Its culture is long-form and chemistry-forward, carried by a handful of very deep accounts — one member has 5,820 posts, another 5,016. The oracle is a user called **linuxboy**, whose posts from 2010 are still quoted back at newcomers sixteen years later.¹

And it is dying in a specific, visible way. The recent-posts feed in October 2026 is almost entirely spam: one thread, *This is NOT a Dating Site!*, posted at six-to-ten-minute intervals through 19 September, and another called *Test Your Luck Today!*, both inside a board named “**New website owner**”. Board pages intermittently return **HTTP 508, Loop Detected**, with a four-hour retry — a shared host throttling an under-resourced site. The archive is the asset; the conversation has left.¹

THE RUNNING ARGUMENTS, IN THE PARTICIPANTS’ OWN WORDS

On **UHT milk**, which is the whole question in a tropical supermarket. **linuxboy**, 2010: “**the heat treatment damages caseins, making it more difficult for them to bond... your curds will simply be too weak to hold together**” — but, separately, “**UHT is fine for lactic-coagulated cheeses. In fact, it will make for a thicker set.**” Settled position: fine for yoghurt, quark and mascarpone, useless for anything rennet-set. On **calcium chloride**.²

The gear ladder, priced. The press is where the money starts. A spring press is US\$15.97. New England’s Dutch-style lever is \$199.95 — “a 5 lb.

weight hung on the furthest notch will equal 20 lbs. of pressure”, to a hundred pounds. Lehman’s all-stainless American-made press is \$239. And at the top, Glengarry’s two-piston pneumatic Dutch press at CAD \$3,995. A spring loses force as the wheel compresses and a weight stack must be re-hung; an air cylinder holds the set pressure for six hours unattended.³

CAD \$436.89

ONE **KADOVA 10KG MOULD** — MORE THAN THE ENTIRE COMMITTED-YEAR KIT LIST BELOW IT. THE DUTCH MOULDS ARE WHERE QUIET MONEY GOES³

THE CAVE · WHAT THE MOST-FOLLOWED PRACTITIONER ACTUALLY RUNS, AFTER ABANDONING THE OBVIOUS ANSWER

Everyone assumes a wine fridge. **Webber** tried one and left it: “**the temperature would drift during our hot summer here in Australia.**” What he runs instead is a conventional under-counter fridge governed by an **Inkbird ITC-308**, from \$35.99 — plug the fridge into the controller’s cooling socket, set the temperature, feed the probe through the door seal.⁴

Then the instruments, which is where the serious and the casual diverge. A **pH meter** runs from an Exttech at US\$129.95 to a Hanna Foodcare at CAD \$710.54 with a CAD \$315 consumable probe. Whether a hobbyist needs one at all is the board’s most revealing thread: three experienced makers, in one day, all declining. **Aris**: “**a decent pH meter is expensive and a lot of hassle to take care of.**” **broombank** uses four-to-seven range strips instead and finds them “**very accurate and easy to use.**”⁵

508 LOOP DETECTED

\$15.97 TO \$3,995

WAX IT

SOURCES 1 CheeseForum’s board index, statistics page and recent-posts feed, all read at source on 6 October 2026, for the founding date, the board structure, the per-member post counts, the spam threads with their posting interval and board, and the 508 responses with their retry interval. 2 CheeseForum threads 4006, 15203 and 19144, read at source; every quoted sentence is attributed to the member and date carried there. 3 Prices read at source on 6 October from New England Cheesemaking Supply, The Cheesemaker, Lehman’s and Glengarry Cheesemaking, including the quoted lever-arithmetic sentence from the maker’s own page and the Glengarry mould and press prices in Canadian dollars. The explanation of why a pneumatic press holds pressure is this desk’s reasoning from the mechanism, not an owner’s testimony, and is flagged as such. 4 Gavin Webber’s own cheese-fridge and cheese-cave posts, read at source, for the quoted drift sentence, the controller method and the humidity position; Inkbird’s own site for the controller price; the condensation failure from a build published by New England Cheesemaking Supply, read at source. 5 Instrument prices from New England Cheesemaking Supply and Glengarry, read at source; the quoted remarks from CheeseForum thread 19554, all posted on 7 August 2021.

THE RABBIT HOLE · THE MONEY, AND THE TROPICS

A cave at twelve degrees in a room at thirty is an eighteen-degree lift with no winter — and the one thing Singapore makes easier is the half everyone else struggles with

Milk alone costs about thirty-three dollars a kilo here. Nobody serious claims the hobby saves money. What it does have, surprisingly, is a better law than California's.

The weather service puts Singapore's daily range at **23 to 33°C**, the coolest month at a twenty-four-hour mean of **26.8°C**, and **“mean annual relative humidity... around 82%”**, over 90% before dawn. So a cave at twelve degrees is a sustained **eighteen-degree lift, all year**. A thermoelectric cooler cannot hold that — its achievable drop below ambient is limited, which is the community's stated physics rather than a measurement this desk has seen. **Buy a compressor unit or do not build a cave.**¹

And then the surprise: **the humidity half, which defeats temperate-country makers, is the half Singapore gives away**. At 82% ambient a cave here runs too wet rather than too dry.¹

≈ S\$33 / kg

MILK COST ALONE, COMPUTED HERE FROM MEIJI FRESH AT **S\$6.70 FOR 2 LITRES** AND THE HOBBY'S OWN ROUGHLY TEN PER CENT YIELD — BEFORE CULTURE, RENNET, SALT, WAX, ELECTRICITY OR TEN MONTHS OF ATTENTION. MERIDIAN'S ARITHMETIC FROM PUBLISHED PRICES²

THE SINGAPORE SHOPPING PROBLEM, IN THREE PARTS

The milk. Buy from the chiller, not the shelf, and read the carton for “pasteurised” rather than “UHT” — the chilled two-litre lines are Meiji at S\$6.70, Farm Fresh at S\$6.00, Greenfields at S\$6.05, against shelf-stable UHT at S\$3.07 that will not set.³

THE VOCABULARY THAT EARNS ITS PLACE

Flocculation multiplier — time from adding rennet to first visible flocculation, times a number per cheese, gives the

moment to cut. One member's table: **2 to 2.5** for Alpine and Grana, **3** for Gouda and Parmesan, **4** for feta and blue, **5 to 6** for Brie and Camembert. Higher multiplier, wetter curd, softer cheese, shorter age.⁴

Mite — *Acarus siro*. Not optional: **“if you have natural rind cheese, you will have mites”**, and the same maker adds that they **“do contribute something positive in terms of flavor and aging.”** Compressed air removes about ninety-five per cent.⁴

And the strains people obsess over — *Geotrichum candidum* for wrinkled rinds at \$11.95 to \$23.95 a sachet, *Penicillium roqueforti* at \$43.95 to \$51.95. Strain choice is a genuine faction.⁵

THE SCEPTIC · THE METER NOBODY CALIBRATES, AND THE CHEESE YOU COULD JUST BUY

Three things the hobby admits when pressed.⁶

One last thing, and it cuts the other way. A Californian cannot legally sell a wheel made at home: cheese is not a cottage food anywhere in that state's list.⁷

NEXT DESCENTS

Ham-free satellite tracking & SatNOGS — a receive-only space hobby with a real open-hardware ground-station ladder.

Leathercraft — veg-tan against chrome, the saddle stitch, and a tool roll that outlives you.

Competitive pistol & air-rifle target shooting — the most heavily regulated hobby a Singaporean can legally hold a licence for.

SOURCES 1 The Meteorological Service's own climate page, read at source, for the temperature range, the monthly means and the quoted humidity figure. The limit on a thermoelectric cooler's temperature drop is the community's stated physics; the forum thread carrying it refused this desk and it is flagged rather than cited. 2 Milk prices read at source from the retailer's own listings on 6 October; the per-kilo figure is this desk's own arithmetic from those prices and the yield reported on CheeseForum, and is labelled as computed. 3 Retail listings as above; one bakery supplier's own catalogue, searched at source, carries neither cultures nor rennet; the shipping sentences are quoted from The Cheesemaker's own terms, read at source; the personal-import allowance from the Singapore Food Agency's own guidance, read at source, which does not name dairy cultures as a category. 4 CheeseForum threads 16171 and 14588, read at source, for the flocculation table and both quoted members on mites. 5 Culture prices from New England Cheesemaking Supply, read at source. 6 Instrument prices as above; the quoted calibration complaints from CheeseForum threads 6917 and 19554, read at source. No published figure for discarded wheels exists in anything this desk could open, and none is printed. 7 California's cottage-food list read at source; the Singapore position from the Food Agency's own page on home-based food businesses, read at source.

FITNESS · THE DOSE

Two hundred and thirty-nine trials were pooled to find the optimal exercise dose for back pain. The answer is about a hundred minutes a week — and the authors' own abstract says the evidence does not support an optimal dose at all

Eighty-four per cent of the trials were at high risk of bias. The people who ran the numbers put the demolition in their own summary, which is rarer than the finding.

Chronic low back pain is the signature complaint of a mid-forties body, and the literature on it has for years said only that exercise helps. A dose-response network meta-analysis published in the *British Journal of Sports Medicine* on **24 September** tried to say how much. It pooled **239 trials** across six databases and two registries, pre-defining a **five-point change on a hundred-point scale** as the threshold for clinically meaningful.¹

The headline doses are modest and specific. For **disability** in the short term: **four sessions a week, twenty-five minutes each, over thirteen weeks**. For **pain intensity**: **three sessions a week, thirty minutes, over eight weeks**. Call it a hundred minutes a week for roughly three months.¹

84%

OF THE 239 POOLED TRIALS WERE AT **HIGH RISK OF BIAS**; THE CERTAINTY OF EVIDENCE IS GRADED **VERY LOW TO LOW**, AND ADVERSE EVENTS WERE REPORTED IN 28.5% OF TRIALS, PREDOMINANTLY NON-SERIOUS¹

THE SCEPTIC · AND IT IS THE AUTHORS' OWN

This desk looked for an external critic and could not find one — no accompanying editorial, no expert roundup. It did not need to.¹

The useful consequence for a reader in his forties is deflationary and welcome: **more bought little extra**. The curve flattens early. Whatever is driving the

benefit, it is not volume, and three short sessions beat a heroic one.¹

A smaller trial published three days earlier shows the same shape from the other end. **EduSwim** randomised **76 adults** with chronic low back pain to an individualised swimming programme plus education, or education alone, over four telehealth physiotherapy sessions in eight weeks.²

THE CHATTER · THE LONGEVITY FORUMS HAVE STOPPED ARGUING ABOUT TRAINING

This desk ran the technical internet for exercise discourse and found an empty room. Queries for VO2 max, strength training, creatine, zone two, rucking and protein returned **zero stories** in the window.³

“The available evidence does not support a single optimal exercise dose.”

ALSO ON THE BENCH

What you are training for in twenty years. A second review in the same journal pooled 227 studies, with 105 in its network analysis, on exercise and cognition in older adults without dementia.⁴

General guidance, not medical advice. New pain running down a leg, or pain with weakness or numbness, is a reason to see a doctor before starting anything — the lead paper says as much about its own evidence.

100 MIN/WEEK

VERY LOW CERTAINTY

NOT MEDICAL ADVICE

SOURCES 1 Arora and colleagues, “Optimal exercise dose in adults with chronic low back pain disorders: systematic review with dose-response network meta-analysis”, *British Journal of Sports Medicine*, 24 September 2026; the complete abstract was read through Europe PMC after the journal's own site refused this desk, and every figure, dose and quoted sentence above is from it. The total participant count across the 239 trials is not in the abstract and is not printed. 2 Wareham and colleagues, “Effectiveness of an individualised swimming and education programme for chronic low back pain (EduSwim): a randomised controlled trial”, same journal, 21 September 2026, abstract read through Europe PMC for the randomisation, both effect sizes with their intervals, the 52-week result and the adverse-event note. 3 Story points and comment counts read through the Hacker News search API on 6 October, including dated searches on each training term named. Reddit returned nothing on any route. 4 Hernández-Gamboa and colleagues, same journal, 21 September 2026, abstract read in summary form through Europe PMC; the standardised mean differences and their intervals were not obtainable and are held rather than approximated.

THE DIARY • THE DATED AGENDA

The marathon changed its name, closed its entries and sold out its local places while its own old website still showed last December

Every row here was put back to a promoter or a venue this morning. Two listings circulating elsewhere are wrong, and one is a festival whose organiser still says “stay tuned” about last year.

SINGAPORE

BOOK NOW • PLAN AHEAD

- 7-8 NOV

★ Pokémon RUN 30 — Singapore Edition
 STADIUM ROAR, THE KALLANG. A **4KM NON-COMPETITIVE** RUN, **S\$138**, LATE ENTRY S\$158. PRE-REGISTRATION ONLY — NO ENTRY ON THE DAY — BUT THE **FESTIVAL ZONE IS FREE**, 8AM TO 2PM BOTH DAYS¹
- 8-10 JAN

★ Jay Chou — Carnival II World Tour
 NATIONAL STADIUM, THREE NIGHTS. ON SALE SINCE MAY, FULL LADDER PUBLISHED: **S\$238 / 278 / 318 / 358 / 398**¹
- 18 OCT

NCT 127 — Neo City: The Redline
 INDOOR STADIUM, DOORS 5.30PM. ON SALE SINCE 21 AUGUST, **S\$228-368**¹
- 31 OCT

SING浪 2026 — Mandopop meets EDM
 INDOOR STADIUM, THE VENUE'S FIRST 360° STAGE. **S\$18 / 28 / 35**, THE CHEAPEST BIG-ROOM NIGHT IN THE WINDOW. **NO ADMISSION UNDER 13**¹
- 10 OCT-8 NOV

Deepavali — Little India Festival Village
 CAMPBELL LANE AND SERANGOON ROAD, FREE TO WALK. **DEEPAVALI FALLS ON SUNDAY 8 NOVEMBER**; MONDAY THE 9TH IS A HOLIDAY ONLY WHERE THE REST DAY FALLS ON THE 8TH²
- 13-22 NOV

Singapore Writers Festival, 29th edition
 “HIDDEN IN PLAIN SIGHT”, 300+ WRITERS. **FESTIVAL PASS S\$24 EARLY BIRD, YOUTH PASS S\$12**, COVERING 100+ PROGRAMMES¹
- 4-6 DEC

BYD Singapore International Marathon
 THREE DAYS FOR THE FIRST TIME. **REGISTRATION CLOSED 30 SEPTEMBER; LOCAL MARATHON AND HALF PLACES ARE SOLD OUT**. SEARCH THE OLD NAME AND YOU REACH A SITE STILL SHOWING DECEMBER 2025¹

JAPAN & THE REGION

PLAN THE FLIGHTS

- 11-13 DEC

★ Tokyo Comic Con, tenth anniversary
 MAKUHARI MESSE. **¥5,500 A DAY** IN ADVANCE AND **¥500 FOR UNDER-15S**, WITH PRIMARY-SCHOOL CHILDREN FREE BESIDE A GUARDIAN. NICOLAS CAGE MAKES HIS TOKYO DEBUT ON THE 12TH AND 13TH³
- 12 NOV-25 DEC

★ Marunouchi Bright Holiday — Harry Potter
 TOKYO. A HOGWARTS SNOW TREE IN THE MARUBIRU ATRIUM AND **1.2KM OF CHAMPAGNE-GOLD LIGHTS** ALONG NAKA-DORI; THE ILLUMINATION ITSELF RUNS TO **14 FEBRUARY**. FREE, NO TICKET³
- 2-3 DEC

Chichibu Yomatsuri
 EIGHTY MINUTES FROM IKEBUKURO. A UNESCO-LISTED NIGHT FESTIVAL, FIREWORKS OVER THE FLOATS FROM ABOUT THREE ON THE 3RD. FREE TO STAND IN THE STREET; THE PAID GRANDSTAND IS NOT YET ON SALE³
- 3 NOV

Lexus Melbourne Cup Day
 FLEMINGTON, FIRST RACE 10.45; THE CUP IS 3,200M FOR **A\$8.66M**. GENERAL ADMISSION **A\$120**⁴
- 15-17 JAN

Tokyo Auto Salon 2027
 MAKUHARI MESSE. **¥4,000 FRIDAY, ¥3,000 AT THE WEEKEND**. **DIGITAL TICKETS ONLY, ADVANCE ONLY** — NO PAPER, NO DOOR SALES, RESALE PROHIBITED³
- 29 JAN-7 FEB

Kobe Luminarie
 STREETS FREE; THE MERIKEN PARK VIEWING AREA IS **¥500 WEEKDAY AND ¥750 WEEKEND** IN ADVANCE, DOUBLING ON THE DAY³

PLANNING NOTES • WHAT HAS GONE, AND TWO LISTINGS THAT ARE WRONG

Almost everything large here went on sale months ago — Jay Chou in May, NCT in August. The pattern is fan-club presale, then a bank or promoter presale, then general sale a day or two later, so hold a free Live Nation membership *before* you want something, and note which seller: SISTIC for some promoters, Ticketmaster for others. **Already gone:** local marathon and half places, and the Pokémon run has no entry on the day. **And two circulating listings this desk will not carry.** Christmas Wonderland is being given an opening of 1 November; its own site still serves a 2025 closure page saying “stay tuned”.⁵

SOURCES 1 The Kallang’s own event pages, read at source on 7 October, for every Singapore row’s date, venue, on-sale date and published price ladder; the Pokémon run additionally against the franchise’s own Singapore portal, and the writers’ festival against its own site. Where a venue publishes no price — LANY and wave to earth among them — no price is printed and no row is run. 2 Festival Village dates from the festival’s own site; the public-holiday date and the rest-day condition from the Ministry of Manpower’s own table, read at source. 3 Tokyo Comic Con’s own ticket page, whose three-day exchange window carries stale dates, so only the session dates are printed; the developer’s dated press release for the Marunouchi illumination; and the Chichibu, Tokyo Auto Salon and Kobe Luminarie official sites, each read at source. 4 The Victoria Racing Club’s own race-day page. 5 Checked against christmaswonderland.sg and the Setouchi Triennale’s own site, neither of which supports the circulating listing. **Refused this desk:** Ticketmaster Singapore, Marina Bay Sands and SISTIC’s listings. The Hakone Ekiden’s January dates are convention rather than verification and are not listed.

THE TABLE · SINGAPORE'S NEWEST OPENINGS

A Netflix cooking-show contestant has opened in a Hougang mall and will not take a booking. And two restaurants on this month's "new openings" lists opened in May 2025 and June 2026

Five genuinely new rooms, three of them too new to judge — which this desk says rather than inventing a verdict for them.

Begin with the one that is reviewed. **Sansok** opened on **7 September** at Hougang Mall, and its chef, **Min Yohan**, was a first-season contestant on Netflix's *Culinary Class Wars* — which is why a suburban mall opening drew four write-ups. The format is Korean *hansang*: a main, a soup, free-flow banchan from the chef's own recipes, and *sothap* rice cooked to order in individual pots so it arrives with a scorched golden crust. Whole salt-cured mackerel **S\$21.90++**, pork galbi marinated forty-eight hours **S\$18.90++**.¹

The consensus is warm and specific on value — one reviewer calls it worth it for "affordable price points, generous portions and food quality", singling out the braised mackerel stew. The dissent is equally specific: the smoked beef bulgogi came out "**a tad tough and chewy**", which is a consistency flag rather than a verdict, and the fish is bone-in.¹

Dolpan Hwajoo, open since **14 September** on Tanjong Pagar Road, is the first outpost of a Busan original and is built to be drunk at rather than dined in: large hot-stone sharing plates, several finished with a tableside flame, and a flavoured makgeolli list at S\$22–23.¹

THE BUZZ · TWENTY-FOUR THOUSAND AFFOGATOS, AND A ROOM TO PUT THEM IN

AiFOKATO began in April as a hole-in-the-wall kiosk in Telok Ayer selling build-your-own affogato and sold more than 24,000 of them across March and April alone.²

Noci Bakehouse has taken a second site at Orchard Gateway this month and done something odd: it has **dropped the shio pan it is known for**. In its place, thirteen rotating Korean bagel flavours with house cream cheeses, and a made-to-order *panuozzo* menu — the wood-fired Neapolitan sandwich bread — from **S\$18 to S\$23**.²

WALK-INS ONLY

TOO NEW TO CALL

HOUGANG

LUNCH IS THE VALUE

The most consequential opening is not a restaurant. **Food Opera** reopened under ION Orchard on **28 September** after a full refit as **eighteen stalls across 600 seats**, curated as heritage plus new talent rather than the usual mall default: **Ah Er Herbal Soup** and **MP Thai**, both Michelin-recognised, beside **Berempah Bros** from a MasterChef Singapore winner.²

THE SCEPTIC · TWO OF THIS MONTH'S "NEW OPENINGS" ARE YEARS OLD

The honesty ledger, and it is mostly about the lists. Of the candidates this desk's research surfaced as October openings, two had to be thrown out on their actual dates.³

STILL ON THE SHELF

The exact opening day for the Orchard bakery is not published anywhere this desk could read. Held until it is.²

SOURCES 1 Eatbook's reviews of Sansok and Dolpan Hwajoo, read at source, for the opening dates, addresses, chef attribution, formats and every price; the quoted assessments and the quoted dissent are from middleclass.sg's reviews of both, read at source. 2 Eatbook and littlebigreddot on AiFOKATO, read at source, for the April kiosk, the volume figure, the 24 September opening, the room and both prices; Eatbook on Noci Bakehouse and on the ION food hall, read at source, for the formats, the price ranges, the stall count and the named tenants. Where no assessment of the food has been published, this desk prints "too new to call" rather than borrowing the opening coverage's tone. 3 Opening dates for both excluded restaurants verified against 2025 and 2026 coverage of their original launches. Burpple returned 403, Instagram rate-limited, Reddit served a consent gate and HardwareZone's search required a login, so no social sentiment is claimed anywhere on this page.

THE UNDERCURRENT · THAILAND

Thailand has bought five million of its citizens a year of frontier artificial intelligence, and it is paying by the head, by the month, at twenty-four baht and six-nine-five-one satang

Everywhere else, AI sovereignty means chips and datacentres. Thailand went to the shop. The contract is a procurement, and the invoice line is a human being who logged in.

There is a new line in a national budget, and it was written in Bangkok. Under a programme called **AiPASS** — the TH-AI Passport — the Thai state has bought **five million** of its citizens twelve months of access to more than thirty frontier models from fourteen providers, for **1,621 million baht**, roughly fifty million US dollars. Registration opened on **19 August 2026**. The platform went live on the 31st. Within two days a million people had signed up.¹

Set beside what other states have done, the strangeness is the shape rather than the size. The Emirates built a gigawatt of compute. India is subsidising GPUs and indigenous models. Vietnam legislated. Singapore made direct arrangements with the model developers and attached a laboratory and a research commitment to them. Thailand did none of that. It went to the shop and bought retail seats.²

฿24.6951

THE CONTRACTUAL PRICE, PER **ACTIVE USER**, PER MONTH — ABOUT ONE SINGAPORE DOLLAR. THE PAYMENT BASIS WAS CHANGED FROM A LUMP SUM TO A HEADCOUNT IN JUNE, AFTER THE FIRST ROW¹

And it did not contract OpenAI, Google or Anthropic. It contracted a consortium: a listed Thai systems integrator and a small firm registered in 2017, forty-nine per cent Hong Kong-owned.³

WHAT A CITIZEN ACTUALLY GETS, AND WHAT THEY MUST DO TO KEEP IT

A Thai national aged fifteen or over registers in person — no proxies — and verifies through the state identity app built out during the pandemic and the welfare-transfer years. Behind one login sit more than thirty models, with a lightweight one as the only uncapped default; everything above it is metered in credits, and video generation is capped at four-to-six-second clips at modest resolution. Access is tiered, and **the tier must be earned**: users take courses from a library of well over ninety curricula and must hold **at least a hundred points a month** to stay at the top tier.⁴

“Remaining funds of fifteen hundred to sixteen hundred million baht, divided by three hundred baht per user.”

Where did five million come from? On the opposition’s account of what officials told parliament, from division. The target was reached by taking the **unspent balance of a digital-economy fund** and dividing it by an assumed **300 baht a head a year**. The population to be served was derived from the money left over, not from any estimate of who wanted it. That claim is single-sourced to one hearing, and it is the single most damaging sentence anyone has said about the programme.⁵

FIVE MILLION SEATS

TWELVE MONTHS

EXPIRES 31 AUG 2027

SOURCES 1 Contract value from Thai PBS, 21 September 2026, and MGR Online, 1 September, both read at source; the tender’s reference price of 1,645.69m baht, against which the winning bid came in 1.5% lower, from Nation Thailand, 29 May. The per-active-user rate from The Bangkok Insight, 5 September, read at source. Registration and launch dates and the five-million cap from the programme’s own site. Early sign-up counts from InfoQuest, 20 August, and Bangkokbiznews, 20 August, both read at source. 2 The international comparison is drawn from a consultancy’s published comparison, read at source, and is a single secondary source; the individual national programmes were not separately verified by this desk. 3 Consortium composition and ownership from Nation Thailand, 29 May 2026, read at source; the quoted sentences from the same publication’s interview of 25 June, read at source. 4 Registration mechanics, model list, metering, tiering, point threshold and lapse rule from the programme’s own site and from MGR Online, 1 September, both read at source; course-completion and certificate figures from The Bangkok Insight, 5 September. 5 Thai PBS’s report of the parliamentary session of 28 May 2026, read at source. The division is the opposition’s characterisation of officials’ own explanation and is single-sourced; no official has confirmed it in those terms.

THE UNDERCURRENT · THE METER

One and a half million people had registered. Somewhere between a hundred thousand and four hundred and fifty thousand were using it. Then the company being paid by the user asked the ministry to stop publishing the number


Asked why, the agency said the contractor gave no reason. Cumulative sign-ups are still released daily. Engagement is not.

Begin with the thing that did work. 1,455,289 people had registered by midnight on 2 September, from a standing start a fortnight earlier. For a state that struggles to reach its informal workforce with anything, that is a real response, and the minister's defence of the price — that it is “**twenty times cheaper than Singapore**” — is not easily dismissed on the published per-capita figures.¹

Now the number that matters, which is how many of them came back.¹

“did not provide reasons”

WHAT THE AGENCY SAID WHEN ASKED WHY THE CONTRACTOR HAD REQUESTED THAT ACTIVE-USER COUNTS NO LONGER BE PUBLISHED. REGISTRATION TOTALS STILL GO OUT DAILY¹



THE SCEPTIC · THE METRIC IS THE PRODUCT

The payment basis changed in June, after the first parliamentary row, from a lump sum to **pay per monthly active user** — which was presented, reasonably, as a taxpayer protection. It created something else at the same time.²

The objections did not begin there. **Pawoot Pongvitayapanu**, an e-commerce founder turned opposition MP, asked in May the question nobody has answered: “**the winning price was only 1.5% below the median price.**”³

THE CHATTER · AND THE HANDS-ON VERDICT

A social-listening report put the programme among Thailand's top ten national conversations for June, at **37.7 million engagements across 3,532 posts**, centred on transparency and feasibility.⁴

“The real gap lies not in the number of AI users, but in the depth of people's ability to use the technology.”

The government's answer on the data is specific and worth recording.⁵

PAY PER ACTIVE USER

METRIC WITHHELD

1,481M BAHT

SOURCES 1 Thai PBS, 3 September 2026, read at source, for the 1,455,289 registrations at midnight on 2 September, for the 450,000 launch-day active figure, and for the agency's quoted answers about the contractor's request; the 110,000 figure from Ken Lohatepanont's walkthrough of 7 September and the 100,000 figure from MGR Online, 1 September, a partisan outlet, both read at source. The minister's comparison quoted from Thai PBS's report of the 28 May parliamentary session. 2 The June contract amendment and the definition of an active user from Bangkokbiznews and The Bangkok Insight, both read at source. Dr Karndee's filing, the quoted sentences, the Thai phrase and the 1,481m baht figure from Thai Post and Dailynews, both 21 September, read at source; this desk could not resolve the romanisation of the filer's surname across outlets and gives none. 3 Nation Thailand, 29 May 2026, read at source, for the quoted question, the tender length, the billboard requirement and the bulk-token point. 4 The engagement figures from a social-listening report published 17 July, read at source; the campaign site read at source, which publishes no verifiable signature count; Lohatepanont's quoted question and findings from his own piece of 7 September, read at source. The analyst's remark in the pull quote is from an economist cited by The Nation, 3 June, read at source. 5 The permanent secretary's quoted assurances from Nation TV, 21 August, read at source; the registration terms as described in a regulatory wire's summary of 25 August and in local coverage of 19 August. **Held:** no reporting after 22 September was found; whether either investigating body has acted, and whether any instalment has been paid, is unknown to this desk.

THE UNDERCURRENT · MONTH THIRTEEN

Thailand will know something by late 2027 that Singapore will not: what five million ordinary people actually ask a frontier model when it costs them nothing. Nobody in the Thai parliament has put a price on that

The argument here is not that Thailand did it badly. It is that Thailand did it first, at scale, to ordinary people, for a dollar a head.

Fourteen years ago a Thai government handed out tablets. The One Tablet Per Child scheme put a device in the hands of **860,000** first-year primary pupils, with a plan for two and a half million more; the auditor later found roughly **thirty per cent** of them broken, and the programme was scrapped after the 2014 coup. Those figures reach this desk through summaries rather than primary sources and should be treated as approximate. The shape, though, is exact: a mass per-child entitlement, procured fast, delivered to a headline number, audited afterwards.¹

What is genuinely new is not the haste. It is the **meter**. The tablets were bought once. AiPASS is rented monthly, per head, against a usage definition — and that turns a capital purchase into a subscription whose value depends on behaviour the vendor can observe and the state cannot easily audit.²

31 Aug 2027

THE DAY THE PROGRAMME ENDS. FIVE MILLION PEOPLE WILL BY THEN HAVE TWELVE MONTHS OF HABIT BUILT ON A SUBSIDY, AND THE STATE WILL HAVE A RENEWAL QUOTE AND NO LEVERAGE IT DID NOT BUY²

Three things follow, and they are why this is three pages rather than a brief.

First, the metric becomes the product. The moment a state agrees to pay per active user, it manufactures a financial instrument out of human

attention — and the counterparty holding the meter stopped publishing the reading within days.²

Second, it is the welfare state with a churn metric. Thai fiscal politics has run for five years on transfers — co-pay schemes, digital wallets, welfare cards. This is the first transfer a recipient forfeits by not consuming it. Lapse ten days, lose the seat. Fall below a hundred points, drop a tier.²

THE CASE AGAINST THIS READING — AND IT IS A SERIOUS ONE

Two defences deserve printing.³

“Thailand must turn Thais from users into creators.”

Third, and this is the part for a Singapore reader: Singapore bought a counterparty and Thailand bought a bill. Singapore’s arrangements are direct, with the model developers, bundled with courses, and came with a laboratory and a research commitment attached.⁴

The uncomfortable part is not that Thailand did it badly. It is the asset nobody has valued. By late 2027 Thailand will hold a year of records of **what five million ordinary people ask a frontier model when it is free** — in Thai, from provinces, by people the surveys do not reach.⁴

A PROCUREMENT, NOT A POLICY

THAILAND

SOURCES 1 Figures for the 2012 tablet programme — the pupil count, the planned expansion and the auditor’s breakage finding — reach this desk through search summaries only; the newspaper archive carrying them refused the request, and no standalone encyclopaedic entry exists. They are printed as approximate and flagged here rather than presented as verified. 2 The payment mechanism, the active-user definition, the ten-day lapse rule, the points threshold and the programme’s end date are as sourced on **p21** and **p22**, each read at source. The characterisation of what a usage-metered entitlement is, and the comparison with Thailand’s earlier cash transfers, is this desk’s argument rather than a sourced claim. 3 The comptroller’s letter, its date and the subcommittee’s decision not to select either project for the integrity-pact programme, reported by Thai Tabloid, 22 September 2026, read at source — a partisan outlet, single-sourced, and printed here because it cuts against this desk’s own reading. 4 The quoted ministerial line is from Thanettakij, 28 August 2026, read at source. The comparison with Singapore’s arrangements rests on the consultancy comparison cited on **p21**, a single secondary source. The closing argument, including the valuation of the query corpus and the eighteen-month forecast, is *Meridian*’s own and is offered as a forecast, not a finding.

MERIDIAN

THE LAST WORD

The reading

A pact was activated and no soldier has been seen. Two laboratories released open-weight models that cannot be downloaded. A headphone is sold on forty-six decibels that no laboratory has measured. A watchmaker's own specification sheet does not mention the complication on the dial. And a government that agreed to pay by the active user was asked, by the company it pays, to stop publishing how many there are.

None of these is a lie. Every one of them is a declaration that has outrun its evidence, and they are not rare — they are most of what a newspaper is handed in a day. The announcement is cheap and arrives on schedule. The thing it describes arrives later, if at all, and nobody issues a second press release to say so.

Which is why the most useful habit this magazine has is also its least impressive: going back a fortnight later to see whether the number held. It did not, twice, in this issue alone — a film finished fourth rather than fifth, and the long bond rose through the largest fall in oil in a year. Reality is slower than public relations, and it does not send a correction.

Twelve Desks · One World MERIDIAN · No. 125 · 7 October 2026

A note on method. This edition's load-bearing figures were read where they are published: the environment agency's own five-o'clock feed this morning (p5); the St Louis Fed's yield series and the New York Fed's term-premium workbook, both downloaded as data (p14); three llama.cpp pull requests and their review threads (p8); a foreign ministry's own joint statement (p3); two box-office charts cross-checked against each other (p11); a measured optics laboratory's test tables (p6); and a venue's own event pages for every dated row (p19). Computed here rather than quoted: the real-yield share on p14, the milk cost per kilo on p17, and the screen-count comparison on p11. Refused this desk today: Reddit on every route, which is why six desks print a null result rather than a sentiment; BirdForum, Cloudy Nights, Head-Fi, Watchuseek and Hodinkee; Billboard, SISTIC, Ticketmaster Singapore and Marina Bay Sands; Bangkok Post, GitHub's API, and the journals whose abstracts had to be read through a repository instead.