

THE CURRENT BRIEFING OF
MERIDIAN

Twelve desks. One world.

SATURDAY 10 OCTOBER 2026 · THE WEEKEND EDITION



PARANAL, WHERE THE PUBLIC WALKS IN FREE ON SATURDAYS. ALMA, WHERE IT NO LONGER CAN (P21)

THE TRAVEL DESK

An eight-metre telescope is free to walk into. The observatory you imagined visiting is shut

THE RABBIT HOLE

One of the four top-shelf stitching-iron makers on earth works out of Singapore

THE ATELIER

The only rule that put a number on LED flicker expired in March. Nothing replaced it

THIS ISSUE · NO. 128

Contents

p3 The World — a prize for international law, a sanction on its court
p4 The World — an eviction, a death, and a dissolved parliament
p5 Singapore — nothing free opens below a PSI of 200
p6 The Atelier — the flicker rule expired in March and nobody noticed
p7 Technology — thirty million a quarter, two days before the hearing
p8 Technology — the cache that doubles decoding and taxes your prompt
p9 The Kit — a 390-gram constant f/2.8 that no bench has measured
p10 The Good Life — a million for a steel watch; no bid for a tourbillon
p11 Screen & Sound — a weekend chart with no studio label on it

p12–13 The Family Desk — the policy worked where it was done
p14 The Macro Desk — the model that stopped explaining gold
p15–17 The Rabbit Hole — leathercraft, two needles and one thread
p18 Fitness — the week’s biggest number, and its interval
p19–20 The Travel Desk — a launch put back a third time
p21–22 The Travel Desk — the Atacama, where the air thins first
p23–24 The Diary — what is on, and four rooms that just opened
p25–27 The Long Read — 848,000 drivers, 11,106 rides a day
THE SPINE: the threshold — nine desks, and the line below which nothing happens

FROM THE DESK

“The improver of natural knowledge absolutely refuses to acknowledge authority, as such. For him, scepticism is the highest of duties; blind faith the one unpardonable sin.”
Thomas Henry Huxley, “On the Advisableness of Improving Natural Knowledge” (1866)¹
A lapsed standard still cited. A gold model that failed on the day its data is missing. Relief that opens above a number nobody reached.

THE STRIP · ON THE WEEK

US 2-year	4.75%	▼ 8bp
US 10-year	5.22%	▼ 6bp
10-year real	2.87%	▼ 5bp
VWRA	\$193.50	▲ 0.39%
DXY	102.23	▲ 0.30%
USD/SGD	1.2804	▲ 0.13%
Brent	\$104.43	▲ 2.13%
Gold, futures	\$4,220.30	▲ 1.39%
Bitcoin	\$82,509	▼ 2.35%

Friday to Friday; Treasury series are Thursday’s. The Straits Times Index closed 5,401.89, down 4.13% (**p14**). Not advice.²

THE SCOREBOARD · SATURDAY

Thirty-four calls came due. Nineteen landed, thirteen missed, two cannot be settled. Misses first.³
Micron, one number, opposite fates. No. 114 said revenue would clear fifty billion; No. 118 said fifty plus or minus one. It printed **\$54.23bn** — the loose call won, the precise one missed by **\$3.23bn**.

On core PCE the timid calls won. Right at **3.0%**: Nos. 115 and 118. Wrong: No. 110 at 3.2, No. 114 at 3.3.
Wrong by a week, and by six. No. 114 had a country single at No. 1 for a twenty-fifth week; it stopped at **24**. No. 109 put the Arctic agreement in November; it was signed **22 September**.

Contradicted by a banner. No. 110: entry closes 30 September “without extension”. The organiser’s homepage now reads **extended till 31 October**.
The tariff cluster, five of six right: the Q4 rate fell to **28.59 cents**, down 10.4%. No. 113 said it would not fall; No. 108 said it would rise.

SOURCES 1 Huxley’s 1866 essay, read at source in the Clark University transcription of *Collected Essays* i, where the sentences stand on page 40, and checked again in the 1907 *Aphorisms and Reflections*, aphorism III. 2 Levels at Friday’s close, each sourced on **p14**; the Treasury and real-yield series carry Thursday’s values because the Federal Reserve’s daily release had not published Friday’s, and gold is quoted on futures because spot and futures differed by about twenty-six dollars. 3 Every claim in this magazine’s own ledger whose check-by date had passed was graded against a figure read where it is published: a fourth-quarter release filed with the Securities and Exchange Commission, the Bureau of Economic Analysis’s August release, the chart-history records for the 10 October charts (Billboard’s own site refused us, redirecting to a host returning a payment error), wire reporting of the New York signing, the marathon organiser’s homepage, a watchmaker’s own collection pages, a reviewer’s full hands-on of 10 September, and the utility’s own tariff page, whose release carries a mislabelled address naming 2027.

THE WORLD · INTERNATIONAL LAW

The prize and the designation, hours apart

Navi Pillay won it in Oslo on Friday morning. By Friday evening Washington had listed the International Criminal Court itself — the first time the court, and not a judge, has been designated.

The citation is one sentence. The Norwegian Nobel Committee gave the 2026 Peace Prize to Navanethem “Navi” Pillay, 85, “for her efforts to promote peace and international law” — a judge of the International Criminal Court’s Appeals Chamber from 2003, UN human-rights commissioner from 2008, now on the International Court of Justice. **12.0 million kronor, from 287 nominations.**¹

The committee was not subtle about why now. International law, it said, “**is no longer just a supplement to peace and security. It is an absolute necessity.**” Pillay chaired the UN inquiry that found Israel responsible for genocide in Gaza.²

Hours later the Office of Foreign Assets Control published a new entry. Not a name: an institution, at an address. **International Criminal Court, Oude Waalsdorperweg 10, The Hague**, established 1 July 2002. Secretary of State **Marco Rubio** left nothing between the lines: “**Either the ICC will end its threats, or we will end the ICC.**”³

“Judges guard the thin red line between order and chaos.”

Then read the small print. The same day OFAC issued **four general licences** — transactions, telecoms and software, pensions, detainees — and set **180 days** before the listing bites. The Court employs about **900 people** and has **125 states parties**.⁴

The reaction came as a bloc: eight foreign ministers — **Britain, Denmark, France, Germany, Italy, the Netherlands, Canada, Japan** — jointly “**regret and strongly disagree**”. The Court answered with the

sentence that matters: “**The Court cannot stand alone.**” One caution, because the ordering invites it: no source this desk could read asserts the designation answered the prize. Hours after is not because of.⁵

DOHA — AL JAZEERA

The law is the target: a dismantling at the moment the thing was honoured. It leads on Rubio’s threat, hands the floor to the Court, and puts first that it was Pillay’s own commission that found genocide in Gaza.⁶

ANKARA — TRT WORLD / ANADOLU

Not a story about America but about European capability. Belgium’s **Maxime Prévot** wants the EU’s blocking statute switched on, and that goes in the subheading; seven condemnations are stacked against the Commission’s note that the sanctions “**would not take effect for another six months**”.⁷

THE SCEPTIC · THE COURT HAS BEEN PACKING FOR A YEAR

William Schabas of Middlesex University London will not call this fatal. “**They have been developing a plan B on how to get around this**”, he told Al Jazeera. His detail is the useful part: the Court has spent over a year migrating off American software and has **ended its relationship with Axa**. Its deputy prosecutor said in September, before the listing, that “**we have done a lot of work to ensure that the court is buffered**”. A designation with four carve-outs and six months of grace is a bargaining instrument. What would settle it: whether a US-exposed bank, insurer or cloud provider drops the Court inside the window anyway.⁸

180

DAYS OF GRACE IN THE DESIGNATION, WITH FOUR GENERAL LICENCES ISSUED THE SAME DAY⁴

SOURCES 1 The Norwegian Nobel Committee’s own citation, read at its prize site; the sum and the nomination count are the committee’s, as carried in reporting of 9 October 2026. “287 candidates” and “208 individuals and 79 organisations” are the same figure decomposed, not a disagreement; a 9.0m-krona figure in search summaries is wrong on both sources read. nobelprize.org refused this desk on every path. 2 The committee’s statement, its chair’s remarks and Pillay’s words on the day, as reported 9 October 2026; her commission’s finding is its own, quoted from the same coverage. 3 The Treasury’s Office of Foreign Assets Control recent-actions listing for 9 October 2026, read at source: the designation notice, the specially-designated-nationals entry with the address, establishment date and tax number, and general licences 13 to 16. The Secretary of State’s lines are from his announcement, carried by two independent outlets. 4 Licences and the 180-day period from note 3. The seventeen prior designations are the State Department’s figure, the looser formulation a US network’s; staff and states-parties counts as reported on the day. 5 The eight ministers’ joint statement, the Court’s own response, and the Belgian and Slovenian positions, as reported by two independent outlets on 9 October 2026. This desk searched for any assertion of causation between prize and designation and found none. 6 Al Jazeera, “US sanctions International Criminal Court itself in latest escalation”, 9 October 2026, with its companion prize report. 7 TRT World, carrying Anadolu Agency, “EU vows to protect ICC as Europe condemns US sanctions plan”, 9 October 2026. 8 Professor Schabas as quoted by Al Jazeera, 9 October 2026. The deputy prosecutor’s September remark was given to the *Financial Times*, which refused an automated request, and is quoted as a US title reported that interview on 25 September. Reuters, the FT and CNBC were unreachable for this page; the Associated Press chain, Al Jazeera, TRT, Xinhua and the Australian broadcaster were not.

MADRID · SEVEN DECADES, ONE STAIRCASE

She died the day before she could sign

María del Carmen Abascal Martín was carried out of her flat of seventy years on 23 September. A new tenancy restoring her was due to be signed on Wednesday.

She was 87, and everyone called her Maricarmen. An international investment fund bought her rent-controlled building and raised the rent **more than 200%**, to about **€1,750** a month. She fought it for **seven years**; on 23 September she was carried out on a stretcher. She died in a Madrid hospital on Wednesday **7 October**, hours before she was due to sign the agreement that would have put her back.¹

From her hospital bed on 27 September she asked for one thing: **“I want what happened to me to help ensure that it doesn’t happen to other people.”** She wanted a law with her name on it, and told the tent camps to continue **“as long as your bodies allow”**.²

The numbers behind her are national. Spain’s housing deficit is about **750,000 homes**, which the Bank of Spain expects to pass **one million by 2028**. Average rent is **€1,176** — **98.7%** of a young person’s average monthly salary, by the Youth Council of Spain’s arithmetic. Parliament was dissolved on Monday 5 October after lawmakers rejected an emergency housing bill carrying an eviction moratorium. A general election has been called for **29 November**, and the two largest unions have jointly called a housing strike for **11 November** — their first joint call in **fourteen years**.³

“She literally put her body on the line for seven years.”

What is left is a threshold of a different kind: a protection that arrived one day late. The tenancy existed. The signature did not.²

THE CHATTER · ONE WORD, SHOUTED AT A PARLIAMENT

Outside the Madrid parliament on Wednesday the crowd chanted one word: **“¡Asesinos!”** — murderers. **Diana**, a protester, told *Democracy Now!* in translation: **“She’s a woman who died of heartbreak. They killed her through heartbreak”**, describing neighbours stunned and cooking for one another in the street. From Barcelona: **“Maricarmen has died. But she didn’t just die. She was killed.”**⁴

THE SCEPTIC · A HURRICANE, AND THE PRESSURE IT MUST BEAT

Mexico holds a **\$175m** World Bank catastrophe bond on its Pacific coast, and this weekend it may take a major hurricane and collect nothing. At 2100 UTC Friday the hurricane centre had **Simon at 105 mph and 966 millibars**, 195 miles south of Manzanillo. The parametric trigger needs **937 mb or lower** inside the box for any payout, and market sources cited by the insurance trade put **932** as the bar for even a quarter. Forecasts have Simon bottoming near **940**. A Category 4 landfall and a zero payout are not a contradiction; they are the contract working as written.⁵

STILL RUNNING · THREE DELTAS, NO RE-EXPLAINING

The Makkah pact. The delta is a troop number. Pakistan’s foreign office said Friday it **“already has a substantial number of forces... in Riyadh”**, while its information ministry called a 30,000-to-40,000 figure **“speculative”**. Stepping back in words, stepping in on the ground.⁶

The tanker campaign has left the strait. Two ships hit outside Hormuz inside 24 hours — one **51 nautical miles** north of Madinat ash Shamal, the first there in six months; a second off Fujairah the maritime authority had **not recorded at all**, confirmed only by satellite fire detection.⁷

France’s lycées. Cumulative custody is now **over 6,800, 86% minors; 18** misconduct inquiries are open and **stun grenades are suspended** after a 15-year-old lost a hand. *Setting the record straight:* this desk could not re-verify the 71,500 turnout or the 6,100 arrests it printed yesterday, and does not restate them.⁸

SOURCES 1 Spanish and wire reporting of 7 October 2026 and a Spanish title’s own account of the death and the eviction, both read at source; the rent increase and the resulting figure are as those reports give them. CNN’s account, which carries a different pair of rent figures, returned a geographic block (451) and is therefore not relied on or quoted. 2 Her remarks of 27 September as carried in the same reporting and in the 8 October broadcast at note 4. 3 Agence France-Presse and Reuters copy carried 7 October 2026, read at source; the Bank of Spain’s deficit projection, the Youth Council of Spain’s rent-to-salary figure, the dissolution, the election date and the joint strike call. 4 *Democracy Now!*, 8 October 2026, read at source, including its interview with the journalist María Carrión and its translated protester interviews. 5 The National Hurricane Center’s Advisory 10 for Hurricane Simon, 300 PM CST Friday 9 October 2026 (2100 UTC), read at source for position, winds, pressure, warnings and rainfall. The bond’s structure and the 937 and 932 millibar thresholds are reported by the insurance-linked-securities trade press, which attributes them to unnamed market sources; they are presented here as that attribution and not as the World Bank’s own published terms. 6 Pakistan’s foreign-office briefing of 9 October 2026 and the defence minister’s remarks, both as carried by two Pakistani titles read at source. 7 European broadcast reporting of 8 October 2026, read at source, which names the maritime authority’s confirmation for the first vessel and its absence for the second. 8 French and English-language reporting from Paris of 9 October 2026, read at source. The United Nations human-rights office entered the story the same day. No source read gives a turnout figure for 8 or 9 October.

SINGAPORE • THE HAZE THRESHOLDS

Nothing free opens below two hundred

The island's worst region peaked at 192 on Friday evening and held there four hours. Every piece of free relief announced that afternoon begins above 200.

It turned, and then it stopped turning. The 24-hour index in the Central region reached **192** at five on Friday afternoon and sat there **four consecutive hours** — the season's high. By **five this morning it was 177**, every region still above a hundred.¹

The hourly reading tells a different story. The West's one-hour fine-particulate figure went from **87** at dawn on Thursday to **171** at two on Friday morning, the episode's highest single hour, then fell all day. Between four and five this morning Central went **118 to 123**. It has flattened, not cleared.¹

On Friday the Haze Task Force was convened: the environment agency and **28 government agencies**. The relief it described is real and specific. **N95 masks** from **vending machines** at selected transport hubs, against an identity card. **Air-conditioned shelters** at community clubs. Preschools to stop strenuous indoor activity. All of it begins **above 200**.²

A reader at 123 and a reader at 60 are given the same label and the same advice.

So this weekend a very large crowd stands outdoors at Marina Bay in Unhealthy air, no masks free and no shelters open, because the island has not reached the number at which they open. The forecast is “Unhealthy to Low Very Unhealthy” — possibly just across it.²

THE CHATTER • THE THRESHOLDS, AGENCY BY AGENCY

Worth knowing before Monday. **Schools stay open**, indoors, with purifiers, and the ministry watches the one-hour figure rather than the headline; **national examinations are being sat in enclosed rooms**, closure considered only on a forecast above **300** — the 2015 line. The armed forces published the clearest pair: outdoor activity adjusted if **the value** passes 200, non-essential activity cancelled if **the forecast** passes 300. One keys off what happened, the next off what is expected.³

THE SCEPTIC • TWO OF THE FOUR BANDS HAVE NO NUMBERS AT ALL

Yesterday this desk reported that the agency quantifies only two of its four hourly bands. Read again today it is worse: **Band 1 is 55 and below, Band 2 is 56 to 150, and Bands 3 and 4 — High and Very High — carry no published range whatever**. The same page adds that in significant haze about **10%** of hourly readings land in those two unquantified bands. One automated read returned ranges for the upper two; this desk could not re-verify them on any agency page and does not print them.⁴

8.6×

CENTRAL SINGAPORE'S 24-HOUR PARTICULATE AT FIVE THIS MORNING, AGAINST THE WHO GUIDELINE OF 15¹

IN BRIEF • TWO NUMBERS

The road first promised for 2020 now opens in 2031. The transport minister told Parliament on Tuesday that the North-South Corridor's viaduct slips to **2029** and its tunnel to **2031** — the fifth published completion date. Eleven MPs filed questions. His concession: “**Soil investigations are not exhaustive**.” And no compensation: the Government “**does not compensate stakeholders for project delays**”.⁵

Every bank and insurer here has twelve months. The monetary authority's AI risk-management guidelines of 7 October cover **all** financial institutions and **all** forms of AI; the institution, not the vendor, stays liable for a third-party model. Sections 3 and 4 bite from **7 October 2027**.⁶

SOURCES 1 The environment agency's own hourly index and fine-particulate readings, read as machine-readable data through the national data service for 8, 9 and 10 October 2026; the agency's own index page returned a not-found error, so the data service was used. Sub-indices for five this morning show fine particulate binding in all five regions. The World Health Organization's 24-hour guideline is its own published figure. 2 The agencies' joint release of 9 October 2026 as reported by a Singapore title read at source, and the environment agency's own haze situation update of 9 October, read at source, for the readings, the forecast wording and the quoted ranges. The preschool, mask and shelter thresholds are as that release sets them out. 3 The education ministry's and early-childhood agency's positions as set out in the same release. The armed forces' two thresholds are quoted from the defence ministry's own release of 9 October 2026, read at source. Reddit and the principal local forum both refused automated requests; no comment-thread sentiment is quoted anywhere on this page. 4 The environment agency's own air-pollution answers page, read at source, which publishes numeric ranges for the first two bands only and gives the 10% figure. The haze portal's health-advisory page and the index guide both returned not-found errors. 5 The transport ministry's own published oral reply of 6 October 2026, read at source, for both dates, the eleven questioners, the quoted concession, the absence of a revised cost and the compensation principle. 6 The monetary authority's own media release of 7 October 2026 and the guidelines it links, read at source.

THE ATELIER · READING LAMPS AND BEDSIDE LIGHT

The one rule that measured flicker expired in March

It was inactivated on 26 March 2026 and nothing replaced it. At the only maker publishing a flicker figure, the flicker-free bulb is the one you cannot dim.

Start with the document, because it is gone. **IEEE 1789-2015**, the recommended practice for modulating current in high-brightness LEDs, is marked **Inactive-Reserved** as of **26 March 2026**. Its rule was usable: flicker not to exceed **frequency × 0.08 — 10%** at 120 hertz, or **4%** for the flicker-sensitive. Nothing replaced it.¹

The mechanism is in one maker's three sheets. Its **non-dimmable** bulb measures **under 1% at 0 hertz**; its **dimmable** one, **under 6% at 120 hertz** — inside the general limit, outside the sensitive one. Its filament bulb publishes **no number**, only "completely flicker-free".²

"At anything under several kilohertz there is a risk of someone not liking it."

So the one behaviour you want at a bedside — dimmed — is the one nobody certifies. And it is the cheap end that publishes: an **\$18 bulb gives R9 of 91** and a **Duv of 0.0000**. A £950 lamp publishes neither.³

Two more. **Warm, not dim**: the evening target is **10 lux melanopic** three hours before bed, which a 2,700K source meets at about **30 lux**. And **the socket test**: an E27 lamp is refillable; an integrated-LED lamp dies with its module.⁴

THE CHATTER · HERTZ, NOT COLOUR

The measurement scene does not argue about CRI. Reviewing a bias light, **maukka** on BudgetLightForum recorded Ra 98, R9 98 — then killed it, because the dimmer runs at **155 hertz**, "too slow... to be used on any other mode than max". The maker's representative argued back, "**should have been ~210Hz**".⁵

THE SCEPTIC · THE SHOP YOU COMPARE ON

Read the retailer's own history page. The design shop **Lumens** was **acquired in April 2021** by the group owning **Flos, Louis Poulsen and Audo Copenhagen**, a group "**jointly owned by... Investindustrial and Carlyle**". Three "independent" houses and the shop that ranks them are one balance sheet.⁶

18%

OF A CLAIMED 25,000-HOUR LAMP LIFE ACTUALLY COVERED BY ITS 36-MONTH WARRANTY, AT FOUR HOURS A NIGHT⁷

NEXT ON THE BENCH

Hand-knit & boot socks — gauge, nylon heels, testable guarantees.

Enamelware & camp mugs — glass fused to steel, and why the rim goes first.

Carbon-steel frying pans — the thin forged pan, and who still makes it.

The Singapore reality is legal. A table lamp here is a **controlled good**: tested, registered, Safety Marked. A circular of **18 March 2026** fined a supplier **S\$15,000** for three unregistered goods, and applies the ban "**regardless of whether the product has been labelled 'for export only'**".⁸

SOURCES 1 The standards body's own record for IEEE 1789-2015, read at source, giving the status and the inactivation date of 26 March 2026. The numeric rule, the two population tiers and the 120-hertz example are quoted from the United States energy department's own published presentation on the practice, co-authored by the chair of the committee that wrote it; the standard's text is paywalled. 2 Three specification sheets from one bulb maker, read as published PDFs — the non-dimmable A19, the dimmable BR30 and the filament lamp — for flicker, index, frequency, dimmability and the incandescent reference. 3 The same A19 sheet for the colour figures, which it publishes per sample from R1 to R15, and for the deviation figure. The heritage lamp's data block was read on a retailer's page carrying the full manufacturer specification; the maker's own portable-lamp category page lists nine lamps with no light-quality specification at all. 4 Brown and others, *PLOS Biology* 20(3):e3001571, 17 March 2022, read at source, for the evening target and the warm-white conversion; sockets and integrated-LED details from each maker's or its retailer's published data block. 5 A measurement thread on BudgetLightForum of 8 October 2018 with a follow-up of 3 November, read at source; posters are quoted as they sign themselves and the maker's representative posts under a named account. 6 The retailer's own history page, read through its commerce host after its main domain refused us, for the acquisition and the ownership sentence; the brand list, revenue and renaming corroborated from an entry carrying its own citations to the Italian financial press. The Danish brand's and the British maker's own pages, read at source. 7 The maker's published L90 figure and warranty term, with the division done here; the energy department's own reliability fact sheet supplies the mechanism — the driver is the weak link, its electrolytic capacitors are temperature-limited, and "lifetime and reliability are not synonymous". 8 The consumer product safety office's own list of controlled goods and its circular of 18 March 2026, read at source. Reddit refused this desk entirely; one maker's spares pages, a distributor and two local retailers also refused, and nothing from them is printed.

TECHNOLOGY · THE MEMORY PATENTS

Six hundred million, two days before the hearing

Micron licensed Netlist's portfolio on 5 October. On the 6th the trade commission terminated the parallel case against Samsung, Google and Supermicro. The evidentiary hearing was set for 23 November.

The filing is unusually specific. On 5 October Micron took a five-year, worldwide, non-exclusive licence to Netlist's entire patent portfolio, covering **server memory modules and high-bandwidth memory: \$30 million a quarter for twenty quarters**, to the third quarter of 2031 — **“a total value of \$600 million”**, in the document's own words. A second, stranger line: Micron also bought **ten million shares** of Netlist for an **aggregate \$1 million**, locked up in fifths over five years.¹

Then the other front, 48 hours later. On 6 October the International Trade Commission **terminated investigation 337-TA-1472** on settlement: Netlist against **Samsung, Google and Super Micro**, six patents, its hearing set for 23 November.²

The patents at issue sit between your memory and your accelerator. This is who collects the rent.

One number should not be adopted. The trade press put the **Samsung** side at “up to **\$898 million**”; Netlist's own quarterly filing anticipates only **“the upfront payment of \$200 million”**.³

THE CHATTER · THE WEEK'S OTHER THRESHOLD

TSMC filed September revenue on Wednesday: **NT\$511,857m**, up **54.6%**; the quarter sums from its own filings to **NT\$1,494,243m**, a record. But it guides in **US dollars** — \$44.6–45.8bn, on an assumed **NT\$32**. Work backwards and the top of guidance is reached only if the Taiwan dollar is **stronger than NT\$32.63**. One wire converted at NT\$31.99 and called it a beat. At NT\$33 it is a miss. **TSMC has published no dollar figure of its own.**⁴

THE SCEPTIC · THE COLLATERAL IS NOT THE CHIPS

The sixty-billion-dollar AI financing is moving, and the “secured on TPUs” framing does not survive its sourcing. As reported this week: **\$42bn senior-secured backed in part by residual-value support from Broadcom**, plus **\$18bn junior** led by Blackstone. The collateral is a **chipmaker's residual-value support**, not the processors.⁵

68%

OF A NINE-HUNDRED-DOLLAR NOTEBOOK'S BILL OF MATERIALS TAKEN BY PROCESSOR, MEMORY AND STORAGE IN THE THIRD QUARTER⁶

IN BRIEF · TWO DELTAS HELD OVER

Mistral Large 4 has a size and a deadline, and still no licence. Its maker's post of 6 October gives **~1 trillion total, 52 billion active** parameters and **“We will release the weights by the end of the month”** — no licence, no context length, nothing on Hugging Face. Meanwhile a 30 September release of the Apple inference stack is titled **“Fix gradients through Mistral4...”**: the toolchain already carries architecture code for weights that do not exist.⁷

Samsung's segment split now has a date: 29 October, its own investor calendar fixing the third-quarter call. The guidance figures this desk printed stay previously-reported: the newsroom timed out.⁸

SOURCES 1 Netlist's current report filed with the Securities and Exchange Commission on 6 October 2026 for the period 5 October, items 1.01, 3.02, 8.01 and 9.01, read at source, together with its exhibit press release: the four agreements, the licence scope and term, the quarterly fee, the twenty-quarter count, the quoted total, and the share purchase with its lock-up schedule. The filing states no per-share price. 2 The Federal Register notice of 9 October 2026 for investigation 337-TA-1472, read as its own full text after the register's HTML pages refused us: the commission's determination not to review, the termination on settlement, the respondents, the asserted patents, the institution date and the scheduled hearing date. 3 The \$898m figure is a patent-trade publication's of 5 August 2026 and is attributed, not adopted; the \$200m anticipation is Netlist's own language in its quarterly report for the period ended 27 June 2026, read at source. 4 TSMC's own monthly revenue filing of 8 October 2026, filed with the Commission and published on its investor site, and its own third-quarter guidance page giving the dollar range and the assumed exchange rate; the two break-even rates are this desk's arithmetic from those figures. The wire conversion implying NT\$31.99 is Reuters', as carried by a market research service. 5 A market-research service's round-up of 8 October 2026, which names Bloomberg as its source for the structure and the participants. The two unprinted figures appeared in no readable source. 6 The same service's cost analysis of 8 October 2026, read at source. 7 The maker's own announcement post dated 6 October 2026, read at source, and the release title of the Apple inference library dated 30 September 2026. Secondary reporting naming 27 October specifically, and disagreeing on licence, is not adopted. 8 The company's own investor-relations calendar, read at source.

THE HOME LAB • THE EXPERT CACHE

It doubles decoding and taxes your prompt

The MoE expert cache merged on Tuesday and went multi-GPU on Wednesday. In the author's own table, the setting that gives 1.85× decoding costs 27 per cent of prompt processing.

Three of the pull requests this desk has been tracking landed at once. 29887, the GPU cache for experts held in host memory, merged 7 October; 29928, multi-token prediction, the same day; 30112, removing the single-device restriction, on 8 October. The maintainer silence reported on Friday is over.¹

The headline is real. On two RTX 4090s the author measures decoding at **34.78 tokens a second** on master and **64.39 — 1.85×** — with a fifteen-gigabyte cache. On a single 5090, **30.8 to 67.8**, a 2.20×, at an 89 per cent hit rate.¹

Read the next column of the same table. Prompt processing on that 1.85× setting goes **440.7 to 321.5** tokens a second — **down 27 per cent**. For long contexts, the actual home-lab workload, that is a net loss.¹

Two independent failures to reproduce the headline, inside the merge threads themselves.

And it does not generalise. One tester ran GLM 5.3 Flash and watched generation fall from **33 to 16 tokens a second**, the hit rate **6.9 per cent** at larger batches. Another measured **+1 to 5 per cent on one GPU and a 25 to 34 per cent slowdown on two**. Both are in the merge threads.²

THE CHATTER • THE FLAG THAT BRICKS IT, AND A FIX A USER FOUND

Tying the cache to the existing split flag has a sharp edge: with `-ts 1,0` a secondary card gets **zero** cache, and one owner of a 16- and an 8-gigabyte pair reports the model **will not load at all**. Per-device sizes are now their own open issue. A 5090 owner also found **repetitive output inside the reasoning block** with the cache on, and fixed it by raising top-k to **2048**.³

THE SCEPTIC • THE BEST SPEED-UP IS AN ARTEFACT

An M5 Max owner filed careful numbers on 7 October: at a 179,000-token prompt, **free-form summarisation decodes at 26.9 tokens a second**. Ask the same machine at the same depth to **quote the last six paragraphs**, with n-gram speculation on, and it hits **73.9**. The 2.7× is real and useless: the output is already in the input, so the drafter is copying. His comparator is the uncomfortable part — a different runtime doing **53.9 tokens a second at 252,000 tokens**. No maintainer has replied.⁴

14%

PROMPT-PROCESSING REGRESSION ON AN INTEL ARC CARD, BISECTED TO TWO BUILDS, FILED 9 OCTOBER — OPEN AND UNANSWERED⁵

STILL ON THE SHELF • TWO UPGRADES TO READ BEFORE YOU RUN THEM

Ollama turns on MLX by default on Apple silicon in v0.40.0 (6 October). The warning is in v0.40.2: models pulled with earlier versions are **converted on disk in the background on first run, the originals kept as backups** — roughly double the disk until you run the supplied cleanup, and no way back below 0.40 without re-pulling.⁶

Home Assistant 2026.10 makes your house an MCP server — one click, administrators only — and adds frequency and power-factor reporting over Matter. Against that: **23 published breaking changes**, including state conditions combining for with attributes that now fail validation outright.⁷

SOURCES 1 The three pull requests on the inference project's own repository, read at source for merge dates, merge commits, the flag added, the device restriction and its removal, and the author's own benchmark tables; the maintainer exchange on per-layer against global sizing is in the third thread, dated 8 October 2026. The project's own web pages were read directly and its release feeds through their published syndication, as the repository interface is outside this desk's tooling scope. 2 Two commenters' figures in the same merge threads, dated 7 and 8 October 2026, read at source; the second had locally removed the single-device check to test it. 3 The per-device sizing complaints and the loading failure are in the 8 and 9 October comments on the multi-GPU thread, now reopened as a separate feature request; the repetition finding and the top-k remedy are in the 8 October comments on the first thread. The downstream fork and the removal merge are both dated 8 October 2026. 4 Issue 30074, filed 7 October 2026, read at source: machine, build and commit, quantisation, prompt depth, both decode rates, the draft-acceptance figures and the comparator. It remains open with no maintainer reply. 5 Issue 30248, filed 9 October 2026, read at source: the backend, the model, and the benchmark figures of 991.74 against 852.47 tokens a second between two named builds, with generation unchanged inside the noise. 6 The runtime's release notes for v0.40.0 to v0.40.2, read at source, with dates taken from the published feed where the rendered page disagrees. 7 The home-automation project's own release post of 7 October 2026, read at source, for the protocol server, the Matter additions and the breaking-change list; the absence of any Lutron change was checked against the same post. Reddit refused this desk entirely, so every community voice on this page is from the project's own threads.

THE KIT · A CONSTANT APERTURE AT 390 GRAMS

Everything measured here was measured by its maker

Sigma's 50-120mm F2.8 was announced on Wednesday at \$799 and ships on the 22nd. No independent bench has touched it, and nobody has published a Singapore price.

The specification is the reason to care. A **constant F2.8** from 50 to 120mm on a cropped sensor — about **75-180mm** in full-frame terms — in **390 grams** and 99 millimetres, with 18 elements in 13 groups, eleven rounded blades, a 55mm filter and a 23-centimetre close focus at the wide end. **US\$799** in all three mounts, shipping **22 October**.¹

The comparison that makes it interesting is one a reviewer made physically rather than from a spec sheet: Canon's RF 70-200mm F4 is about **25% wider and 20% longer retracted, and roughly twice the price**. The Sigma grows a quarter of its length at full zoom.¹

Now the part this desk exists for. The stabilisation claim — **6.5 stops at 50mm, 4.5 at 120** — is a **CIPA self-test by Sigma on a body of Sigma's choosing**. As of today there is **no independent modulation-transfer measurement, no optical-bench number and no lab entry** for this lens anywhere. Sigma publishes its own curves; nothing exists to check them against.²

A lens the specialist press is calling a legend has drawn one reader comment.

What does exist is a week of handling, and it is not uniformly kind. At 50mm near minimum focus on a flat subject, the reviewer reports, **"corners get extremely soft and swirly"** — and says plainly that he overlooked **"some pretty major image quality issues"** for the size and price. He has not tested sports autofocus at all.²

THE CHATTER · THE SILENCE IS THE FINDING

There isn't any. On the enthusiast site that reproduced the press release in full, the post had drawn **exactly one reader comment** when this desk read it, and its pre-order live blog **none**. Hacker News has no story on the lens at all. The live blog lists **24 pre-order retailers** across America and Europe — **not one in Asia**. The photography forums that might have broken the quiet all refused this desk.³

THE SCEPTIC · A MEASUREMENT FORUM THAT HAS NEVER HEARD OF IT

Worth holding beside the Sigma: the same failure in a different category. Schiit's Bifrost 3 publishes **distortion below 0.00015 per cent and a signal-to-noise ratio above 125 decibels**. Exact-phrase searches of the most measurement-obsessed audio forum for "Bifrost 3" return **"No results found"** — not one post, let alone a bench run. It has measured the *old* Bifrost; the new architecture has never been on an instrument. Its ship date has slipped to **"around Halloween"**, at **\$699**.⁴

\$699

SONOS ACE ULTRA, NOW IN STOCK AT THE OFFICIAL SINGAPORE STORE, AGAINST US\$449 ON THE MAKER'S OWN AMERICAN PAGE⁵

STILL ON THE SHELF · AND ONE PRICE THAT FINALLY READ

Ray-Ban Meta Audio: the maker's own page says \$349.

Asked for a sixth issue, and now answered — but not by the press, which printed **\$329** and a 13 October ship date. Meta's own store says **"starting at \$349"**, with no ship date and no specification beyond "Bluetooth AI glasses".⁶

The Fire tablet is dead. Amazon's Alexa Tablet 12 Pro runs **Android with Google Play**, not Fire OS, from **US\$499.99**. No Singapore launch, no camera specification, and the "up to 20 hours" applies only to reading mode — video is rated 16.⁷

SOURCES 1 Sigma's own product page for the 50-120mm F2.8 DC OS Contemporary, read at source, for every optical and physical figure, the mounts, the equivalent focal lengths and the supplied accessories; the US price and the 22 October date are from the American subsidiary's release as reproduced in full by an enthusiast title, and the price was confirmed live on a major retailer's pre-order listing. The physical comparison with the Canon zoom is the reviewer's own, in the review at note 2. 2 A camera title's hands-on of 8 October 2026 after about a week with the lens, read at source. The stabilisation figure is Sigma's own CIPA measurement as published on the page at note 1. This desk searched for an independent optical measurement and found none; the maker publishes its own transfer curves. 3 Comment counts read on the enthusiast title's own two posts on 10 October 2026, and its pre-order live blog, which lists its retailers by country. Hacker News was checked through its search interface. The two large photography forums and the camera title's own forums all refused automated requests. 4 Schiit's own product page, read at source, for the measured specification, the price and the quoted shipping expectation; the 13 October date is a consumer title's of 1 October 2026. The absence of forum results is this desk's own exact-phrase search of that forum, run twice. 5 The official Singapore flagship store's own product page, read at source, with the full specification; the American price is from the maker's own page. The maker's own Singapore product address returns a not-found error, so the local price exists only at the distributor. 6 Meta's own product page, read at source; the differing price and ship date are a consumer title's of 1 October 2026. 7 Amazon's own newsroom announcement of 8 October 2026, read at source, for prices, specifications, operating system, launch markets and the battery wording; the comment is quoted from the Hacker News item, with its score, read at source. No Singapore price exists for this product to read.

THE GOOD LIFE · ONE OF EIGHT, IN STEEL

A million dollars for a thirty-nine-millimetre watch

The Voutilainen 28Sport closed on Friday at 8.47 times its low estimate. Forty-eight hours earlier, a tourbillon at a fifteenth of the price found no bid at all.

Lot 22 was a 39mm stainless-steel Voutilainen 28Sport of about 2021, a navy guilloché hour track around a black onyx centre, from its original owner. The movement is the point: **in-house calibre 28, manually wound, 21 jewels** — a manufacture movement, not a bought base. **One of eight in steel**, and the catalogue calls this configuration “possibly one of only two made”.¹

Estimate **US\$120,000 to 240,000**. It sold for **US\$1,016,000** with premium: **8.47 times the low estimate, 4.23 times the high**. Every one of the sale’s 88 lots sold, and summing the house’s own printed results gives about **\$5.71m** — so this single lot was roughly **17.8 per cent of the entire auction**.¹

The press was behind it the whole way. While the lot was open one title reported the top bid at **\$700,000**, another **\$650,000**. It closed **45 per cent above the highest figure in print**.²

Eight-piece provenance soared. Modern limited-edition horology could not find a bid.

Because forty-eight hours earlier the same market said no. At another house on **8 October**, the headline lot — an **H. Moser Pioneer Tourbillon “Miami”**, estimate **\$50,000 to 70,000** — **did not sell**, nor did a 2026 Streamliner. What cleared was mid-market and mispriced: a 1980s **Breguet Classique** at **\$53,760** against an **\$8,000-to-12,000** estimate.³

THE CHATTER · THE FORUMS WERE TALKING ABOUT SOMETHING ELSE ENTIRELY

This desk went looking for the collector reaction and found none. On the morning of 10 October the principal collector board’s front page carried **eleven live threads** — Fortis Cosmonauts, a Panerai “Cali” dial, **“A mini FPJ GTG at the Wynn in Las Vegas last night”** — and **not one on a million-dollar Voutilainen result**.⁴

THE SCEPTIC · WHAT A MILLION DOLLARS DID NOT BUY

It bought scarcity and a story, not watchmaking anyone can check. **Nobody has published a rate-and-amplitude test, a chronometric certification or an independent teardown** of this piece. Read against the unsold tourbillon, the week does not say “independents are soaring”. It says **eight-piece provenance with an in-house calibre is soaring, and a modern limited edition at a fifty-to-seventy-thousand estimate cannot find a bid**. Different markets wearing one word.⁵

\$66.5m

PAID FOR 95% OF EBEL. THE ASSET LIST IN THE FILING IS TRADEMARKS, INVENTORY AND A LE CORBUSIER HOUSE — NO FACTORY, NO CALIBRE⁵

STILL ON THE SHELF · A CLAIM CLOSED, A RATING ADDED

Bremont now prints 250 itself — the story closes. Three titles printed a 250-piece limitation and the maker printed none. Read today, its own product page says **“limited to 250 pieces worldwide”**, at **£6,450**. A hands-on names what the “BB64AH” designation conceals: a **Sellita SW300-1** base.⁶

And a fourth water-resistance figure for the Seiko. A hands-on of 7 October gives the Power Design Project “Panda” at **30 metres** — the lowest yet printed — with Hardlex rather than sapphire, **US\$540**, from November. One maker, four ratings.⁷

SOURCES 1 The auction house’s own lot page for lot 22 of its New York online sale, read at source, for the description, the movement, the edition size, the provenance, the estimate and the realised price including premium; its own sale page gives the 88 lots and the concluded date. The sale total is this desk’s own sum of the 88 printed results — the house publishes no official total — and the percentage follows from it. 2 Two watch titles’ in-sale reports, read at source, for the \$700,000 and \$650,000 figures. The discontinuation claim appears in one of them and in no other source, and is not in the house’s catalogue. 3 A watch-industry results round-up of October 2026, read at source, for the unsold tourbillon, the unsold Streamliner and the Breguet result against its estimate. 4 The collector board’s own front page, read on 10 October 2026; thread titles are as it displayed them. Two other watch titles refused automated requests, as did a headphone forum. 5 The seller’s current report filed with the Securities and Exchange Commission on 8 October 2026 for the period 5 October, and its exhibit press release, both read at source: the price, the 95 per cent, the retained 5 per cent, the buyer group led by the independent watchmaker’s company with the fashion house participating, the incoming chief executive, and the asset list including the Le Corbusier house. How the 95 per cent splits is disclosed by nobody. 6 The maker’s own product page and its two regional stores, read at source today, for the limitation and both prices; the movement base is named in a watch title’s hands-on of 8 October 2026. The maker’s Singapore store returned a not-found error on re-test. 7 The same title’s hands-on of 7 October 2026, read at source. No yen or Singapore price is published.

CULTURE DESK · TWO WIDE OPENINGS

A nineteen-million horror beat Sorkin five to one

Other Mommy took \$3.6m on Thursday against *The Social Reckoning*'s \$650,000 in previews. Neither film has an audience score yet; both are still counting critics.

Thursday's dailies are the only published numbers for this weekend, and they are lopsided. **Other Mommy** — Rob Savage's Josh Malerman adaptation, Jessica Chastain, a \$19m Blumhouse budget — took **\$3,600,000** on opening day from **2,850** sites. **The Social Reckoning**, Aaron Sorkin's sequel at a reported **\$30–42m**, took **\$650,000** in previews. That round \$3.6m reads as a studio estimate.¹

The critics have gone the other way and still not settled. *Other Mommy* sits at **18% from 38 critics**, its consensus conceding a “valiant performance” inside a film “too silly”. *The Social Reckoning* is at **56% from 179**. Both were moving hour by hour as this issue closed, so both are timestamped to Friday evening.²

That 56% is this desk's own continuity delta, and it is a fall. Yesterday we printed **61% from 120**. Fifty-nine reviews later it is five points lower — what a score does when the early, invited critics stop being the whole sample.²

Neither film has an audience percentage. One has “250+” ratings, the other “50+”.

Last weekend's chart, for the record: **Verity** opened at **\$32,031,011** from 3,510 sites and **Digger** at **\$8,001,856** — the hard number behind this desk's correction two issues ago. Sixty-eight titles reported **\$98.4m** between them.³

THE CHATTER · THE ARGUMENT IS ABOUT AUTHORSHIP, NOT THE FILM

With the forums shut to this desk, the letterboxing site is the record — and it is too early for either film to carry a weighted average, which is itself the finding. On *The Social Reckoning* the fight is not whether it is good but **who made the first one: jonathan fujii** (two stars) — **“I’ve come to the conclusion that David Fincher actually deserves MORE credit”**; and, drily, a l e x: **“imagine being such a villain, you get two movies before you actually die”**.⁴

THE SCEPTIC · NAMED CRITICS, ON BOTH

On the horror: **Peter Travers** (one and a half of four) says it **“crashes and burns as human drama”**; **Jacob Oller** of *The A.V. Club* finds **“barely any actual harm or threat”**. On the sequel: **Annie Berke** in *The New Republic* — **“a journalism movie that feels like a play, and a didactic one at that”**. For balance, **Sonny Bunch** at three of four: it **“shows how he filled that hole: with ragebait”**.⁵

1,540m

MINUTES OF *TOY STORY 5* STREAMED IN FOUR DAYS OF ITS DEBUT WEEK, AFTER A \$1.148BN THEATRICAL RUN — 3.7 TIMES THE NEXT FILM⁶

IN BRIEF · WHAT A NUMBER ONE COSTS NOW

Sixty-one thousand units, and 47,000 of them were purchases. Miley Cyrus's *Bass Persuades* entered at No. 1 on the chart dated 3 October with **61,000 equivalent units — 47,000 pure sales** — from an artist whose lead single peaked at **No. 35**. Seventy-seven per cent of a number-one week was people buying a physical object.⁷

And no tour to report. This desk could not verify a single major tour announcement dated inside the window at source; everything that surfaced was announced in August or September. Rather than print a stale date, it prints none.

SOURCES 1 The Numbers' own daily domestic chart for Thursday 8 October 2026, read at source, for both films' figures, theatre counts and the previews flag; the Friday chart had not published when this issue closed, and the specialist trade's Friday read returned a payment-required error. Budgets: \$19m from the horror film's own encyclopaedic entry citing the trade press; \$30–42m for the sequel from the same source citing a named trade reporter on 7 October 2026, against a different title's “around \$50 million” — the lower, better-sourced range is printed and the dispute noted. 2 Both films' aggregator pages, read 9 and 10 October 2026, for scores, review counts, consensus wording and the verified-rating counts; the second aggregator's figures are as the films' encyclopaedic entries cite them. Scores on these pages change hourly. 3 The Numbers' own weekend chart for 2 October 2026, read at source; the page carries no estimate-or-final label and no date range, and is treated as its reported chart rather than a studio final. 4 The letterboxing site's own film pages, read at source; reviewers are quoted as they sign themselves, with their star ratings as displayed. Reddit refused this desk entirely. Neither film yet carries a weighted average there. 5 Critic excerpts and scores as published on the aggregator pages at note 2, each attributed to the named critic and outlet. One further negative review cited elsewhere could not be read and is not quoted. 6 Nielsen's own streaming top ten for the week of 21–27 September 2026, the most recent published, read at source; the platform debut date of 23 September is from the film's encyclopaedic entry citing the trade press, and the theatrical gross from the same entry citing two box-office compilers. Nielsen measures United States television-set viewing only. 7 The album's own encyclopaedic entry and the chart-history record for 2026, read at source. The 3 October chart date falls a few days before this issue's window, and is printed as context for the following week's figure.

THE FAMILY DESK • TWELVE THOUSAND CHILDREN

Twenty-six minutes a day buys 0.04 of a standard deviation

The largest device-measured study of its kind reports small, unstable effects — and finds more sitting associated with slightly lower blood pressure. General guidance, not medical advice.

Published on Tuesday: a harmonised meta-analysis of **12,874 young people**, mean age **10.4**, followed a mean **3.1 years**, with activity measured by accelerometer across 14 cohorts rather than asked about. The units matter. One standard deviation of moderate-to-vigorous activity is **26 minutes a day**; one of sedentary time is **106 minutes**.¹

What 26 extra minutes a day was associated with: insulin **-0.24** microunits per millilitre, LDL **-1.54** milligrams per decilitre, and a clustered risk score **-0.04 of a standard deviation**. Real, consistently signed, and small.¹

Two rows go the other way, and the authors print them rather than bury them. More activity was associated with **higher** systolic pressure, by **0.25 mm Hg**; more **sedentary time** with **lower**, by **0.90**. Their own verdict: the associations are “**weak, sensitive to the modeling approach, and heterogeneous across subgroups**”, and sitting time is **unlikely to be an important determinant** of cardiometabolic risk in the young.¹

The rules worked where they were adopted. About half of the services did not adopt them.

A second trial the same week shows what that costs. Across **197 after-school care services** and about **10,000 children** in three Australian states, a voluntary accreditation scheme returned, on intention to treat, **an odds ratio of 1.63, interval 0.53 to 5.00** — nothing. Among services that actually accredited, **7.35, interval 1.39 to 38.86**. Both are true. The width of the second is the honest part.²

SOURCES 1 Husoy and others, “Device-Measured Physical Activity, Sedentary Time, and Cardiometabolic Risk Markers in Youth: A Harmonized Individual Participant Data Meta-Analysis”, *JAMA Pediatrics*, published online 6 October 2026, doi:10.1001/jamapediatrics.2026.5176, read at source: the design, the eight contributing datasets and fourteen cohorts, the sample and its mean age, the follow-up, the standard-deviation units for each exposure, every coefficient and interval quoted, and the authors’ own conclusion. One row of that table — the body-mass-index result — is reported with one interval in the abstract and a different one in the table; this desk does not print that row. 2 Maher and others, “A Policy-Based Intervention for Physical Activity and Screen Time in Childcare Outside School Hours: A Cluster Randomized Clinical Trial”, *JAMA Network Open*, published online 5 October 2026, doi:10.1001/jamanetworkopen.2026.36741, read at source: the cluster-randomised design, the 197 services and their split, the approximate child count, the accreditation requirement, and both the intention-to-treat and per-protocol odds ratios with their intervals. Children were not individually enrolled. 3 Li, Paluch and Rakesh, “Physical Activity and Adolescent Brain Structure, Mental Health, and Academic Outcomes”, *JAMA Network Open*, published online 1 October 2026, doi:10.1001/jamanetworkopen.2026.36973, read in full at source. The timing discrepancy between its Methods and its Abstract and Results is stated here as this desk found it; the paper itself does not resolve it. 4 Staiano, “Boosting Children’s Physical Activity in Outside School Hours Care”, invited commentary, *JAMA Network Open*, published online 5 October 2026, doi:10.1001/jamanetworkopen.2026.36701, read at source; all three quotations are from that page. This desk could not verify any fresh publication in the window from the researchers most often cited against the screen-time consensus, and does not attribute anything to them here.

THE CHATTER • A CAVEAT IN A PAPER WE CITED

A note on our own sourcing. The study this desk reported two issues ago — mean age **11.93**, mental-health coefficient **-0.080** — carries an internal discrepancy worth stating: **its Methods describe the outcome at three-year follow-up, its Abstract and Results one year after the step data**. The paper does not resolve it. Its other figures stand: **5,515** participants, a mean **9,595 steps a day**, and **no significant association with academic performance** at all — the part the coverage mostly skipped.³

THE SCEPTIC • DON’T BE ALARMIST, FROM INSIDE THE SAME JOURNAL

Amanda Staiano of the Pennington Biomedical Research Center, in an invited commentary alongside the trial, declines the easy reading in both directions. “**The effect sizes on children’s physical activity were modest**”, she writes — but “**this is an implementation challenge, rather than a critique of the intervention itself**”, and “**systems-level strategies are needed beyond voluntary uptake**”. The programme works through adults, so organisational readiness decides who benefits. Which is the quieter finding: a voluntary scheme sorts by capacity, and the average then tells you nothing about either group.⁴

-0.90

MILLIMETRES OF MERCURY: THE SYSTOLIC BLOOD PRESSURE ASSOCIATED WITH ONE MORE STANDARD DEVIATION OF SITTING — 106 MINUTES A DAY — IN 12,874 CHILDREN¹

Which leaves a parent roughly where the evidence does. Twenty-six minutes a day is worth having and is not a transformation; the sitting panic is, on this evidence, unsupported; and a policy that half the institutions decline is a policy with no average effect. *General guidance, not medical advice.*

THE TWEEN YEARS • WHAT THE TIME WAS WORTH

Homework fell by a third. Enrichment spending rose fivefold

A repeat of a 2016 study of 1,500 Singapore parents: \$260 a month on enrichment against \$48 a decade ago, and 4.3 hours of homework against 6.46.

Read the two numbers together, because separately each is ordinary. Weekly homework for a Singapore primary-school child has fallen to **4.3 hours**, from **6.46** a decade ago. Monthly enrichment spending has risen to **\$260** from **\$48** inflation-adjusted, private tuition to **\$500** from **\$242**. The school gave the time back, and the market took it.¹

The study repeats a 2016 survey, fielded May to August across **1,500** Singaporeans and permanent residents with a child in a local primary school. Its sharpest finding is a contradiction people hold without discomfort: **89.8%** agreed children spend too much time on tuition — and **seven in ten said they would spend more**. About half named peer pressure.¹

And it is not remedial. The researcher who presented it put it plainly: for nearly **60 per cent** of parents, **“tuition is not about their children needing some additional help in school... It’s just simply because they want to stretch their potential.”** And **84.9%** agreed it confers an advantage in the leaving examination.¹

Paid tuition in private housing does not replace the coaching parents do themselves.

For a household at the secondary transition the sorting is the interesting part. Weekly tuition hours run **2.53** in one-to-three-room flats, **4.29** in condominiums, **5.01** in landed homes — but the hours parents spend coaching their own children barely move, **4.5** to **4.7**. And **50.4%** in the smallest flats used enrichment to prepare for direct admission, against a quarter to a third of larger households: the

discretionary route is worked hardest by those with least.¹

THE CHATTER • WHAT THE CALENDAR IS SELLING, IN THE SAME FORTNIGHT

With both parent forums closed to this desk, the proxy is what a Singapore parenting site published, dated on its own page. In three weeks: **“Secondary School Open House 2026”** (7 October), **“Screen Time After PSLE: Do Something About It Now”** (28 September), **“PSLE Marking Days 2026”**. On the same day as the open-house piece, a sponsored slot: a **guide to year-end holiday enrichment classes**. The conversation is sequenced to the exam calendar, not the research, and the moment a child gets time back is the moment it is sold.²

THE SCEPTIC • THE CHATBOT THAT WORKED HAD NO AI IN IT

A randomised trial published on Wednesday put a decision-support tool in front of **106** young people aged 12 to 24 with hereditary cancer-risk syndromes. Accurate knowledge of their own risk by age 30 went from **37%** in the control arm to **78%** — odds ratio **3.50**, interval **1.45** to **9.19**. Every secondary outcome was null. The detail worth keeping is in the methods: **“The research team wrote all content without generative AI.”** In 2026, the chatbot that beat a clinic visit on knowledge was a hand-written decision tree.³

75.8%

OF PARENTS WHO WOULD STILL PREFER SECONDARY ENTRY DECIDED ON EXAMINATION RESULTS — WHILE 57.5% CALLED DIRECT ADMISSION THE BETTER ROUTE FOR THEIR CHILD¹

One standing rule, before the holidays: since January, phones and smartwatches have been barred across the whole secondary school day, and the default sleep time on school learning devices moved to **half past ten**. Not medical advice.⁴

SOURCES 1 A study of parents’ attitudes to the education system by a Singapore public-policy institute, a repeat of its 2016 work, presented at a media briefing on 25 September 2026; the sample, fieldwork window, spending figures, homework hours, agreement percentages, the quoted remarks of the principal research fellow who presented it, the housing breakdowns and the secondary-transition findings are all as reported from that briefing by *The Straits Times* and read by this desk in a regional reprint. **The institute’s own page for the study returned an empty body to this desk, so the primary document is held** and every figure here is attributed to the briefing as reported. 2 A Singapore parenting site’s own published article list, read at source on 10 October 2026, with each item’s own displayed date. The two parent forums this desk would normally read both refused it, and nothing here is attributed to any parent. 3 Kamihara and others, “Chatbot Intervention and Cancer Risk Knowledge in Adolescents and Young Adults With Cancer Predisposition: A Randomized Clinical Trial”, *JAMA Network Open*, published online 7 October 2026, doi:10.1001/jamanetworkopen.2026.35794, read at source: the design, the randomised count, the age range, the primary outcome with its odds ratio and interval, the null secondaries, the authors’ own statement that the trial was underpowered for distress, and the quoted sentence from its methods. 4 The education ministry’s own press release of 30 November 2025, read at source, which remains the operative rule; this desk found no ministry or agency announcement dated inside this issue’s window and prints none.

THE MACRO DESK • A MODEL THAT STOPPED WORKING

Real yields fell, and gold fell with them

For three issues the real-yield share explained nearly all of gold's move. On the only fully published window this week it explains none, and the sign is inverted.

This desk has printed the real-yield share of gold's move three issues running: **91%, 106%**, then about 98. Take the only window this week with both endpoints published — 2 to 8 October — and the ten-year real yield fell **five basis points, 2.92% to 2.87%**. On the quarter's own regression that should have lifted gold **0.55%**. Gold fell **0.13%**. The share is not small. It is **sign-inverted**.¹

A naive ratio prints -436%, an artefact of dividing by near-zero, so this desk prints the failure instead. And gold's whole weekly gain — futures **+1.39%** — landed on Friday, **the one day the real yield has not been published**.¹

The regime is otherwise legible. The Fed is **hiking**, at 3.75–4.00%, and Wednesday's minutes said **“all participants viewed a higher target range... as appropriate”**. Ten-year breakevens sat **dead flat at 2.36%**: the entire move was **real**.²

The twenty-one-day equity-bond correlation is above the two-hundred-and-fifty-two-day on every pair.

Which leaves diversification where nobody wants it. On two years of daily closes, global equities against global bonds run **+0.65 at twenty-one days** against **+0.49 at two hundred and fifty-two**. Bitcoin, **-2.35%** on the week, behaved as a liquidity asset: **+0.55** to US equities.³

SOURCES 1 The Federal Reserve's own daily yield series, read as published data: the ten-year inflation-indexed yield at 2.92% on 2 October and 2.87% on 8 October, with no value published for 9 October when this issue closed. Gold is quoted on futures throughout, because spot and futures differed by about twenty-six dollars on Friday and this desk does not mix them in one series. The expected move is from a regression of daily gold-futures returns on daily changes in that real yield over the trailing quarter, run here; the coefficient and the working are this magazine's own, not a published statistic. 2 The same series for the nominal ten-year, including the 5.31% reading of 5 October, with the full series scanned to establish the 14 May 2002 comparison; the breakeven series for the flat 2.36%; and the Federal Open Market Committee's own published minutes of its 15–16 September meeting, released 7 October 2026 and read at source, for the quoted sentence and the year-end expectation. 3 Correlations computed here from two years of daily closes to 9 October 2026, on publicly quoted series; they are this magazine's own arithmetic. Weekly changes for bitcoin and the other instruments are close-to-close on the same series. 4 Analyst remarks as published by a wire on 9 October 2026 and read in syndication, each attributed to the named individual and firm. Auction results are from the Treasury's own auction data service, read as published data for the 7 and 8 October reopenings. 5 The New York Fed's own term-premium workbook, downloaded and read row by row: the values for 2, 5, 6, 7 and 8 October, the absence of a 9 October row, and a scan of the full daily series establishing 11 June 2014 as the last higher reading. The decomposition of the September-to-October move into term-premium and expectations components is arithmetic from that file. 6 Index and bank closes from quoted market data; the broker's downgrade and the two comparative bank-index falls are as published by a Singapore investment research site on 9 October 2026, read at source. 7 Currency and yield levels as above; the policy date is a market calendar's and is **single-sourced** — the authority's own site is rendered by script and listed nothing this desk could read. The July analyst split is as reported at the time. Reuters, the *Financial Times*, Bloomberg directly, CNBC, Forbes and the oil producers' group all refused automated requests; every figure above was routed to a statistical or official source instead.

THE CHATTER • REGIME OR SPIKE

The week's consensus is that the term premium is a regime. **Frank Rybinski** of Aegon: **“There's multiple ways to calculate it, but they're all going higher”**. **Florian Ielpo** of Lombard Odier: **“The bond market is no longer waiting for the Fed to tighten financial conditions.”** Against them: **\$39bn of ten-years went at 5.300%, dealers taking 2.5%**. If this were a buyers' strike, that is the wrong number. *Explanation, not advice*.⁴

THE SCEPTIC • CONTRADICTED BY ITS OWN SOURCE

This desk printed **0.9848** on the ten-year term premium yesterday and called it a post-June-2014 high. Both halves check out: the last higher reading was **11 June 2014, at 0.9951**. But the New York Fed's own workbook now has **8 October at 0.9198**. Week on week the term premium moved **+0.6bp**; **expectations** fell by 4.3. A major wire ran the regime story on Friday citing roughly 0.98, a figure its own source had already revised away from.⁵

5.31%

THE US TEN-YEAR ON 5 OCTOBER, HIGHEST SINCE 14 MAY 2002 — 104 BASIS POINTS ABOVE THE 2025 CLOSE. BY THURSDAY IT WAS 5.22²

THE SINGAPORE READ • A 279-POINT GAP, AND WEDNESDAY

The index fell 4.13%, it was banks, and it was not local. The Straits Times Index closed **5,401.89**, its lowest since 8 July. OCBC lost **8.40%** on the week, UOB 7.31%, DBS 4.99%. The wider frame: a US bank index fell **12.3%** from mid-August.⁶

And the currency did nothing at all. The Singapore dollar absorbed a 24-year high in US yields with a **0.13%** move, at **1.2804**. Ten-year stock here yields **2.45%** against 5.24% there — **279 basis points**. The policy statement is due **Wednesday 14 October**; in July **13 of 18 analysts** wrongly expected no change.⁷

THE RABBIT HOLE · LEATHERCRAFT

Two needles, one thread, and each pass locks the last

The saddle stitch cannot be made by machine, which is why a hobby of 888,000 people still turns on a hand tool. One of the four makers everyone wants is in Singapore.

Start with the stitch, because everything else is downstream of it. A sewing machine makes a lockstitch: one thread from above, one from below, knotted between the layers. Cut it anywhere and it unzips. A **saddle stitch** uses **two needles on one thread**, passing each through the same hole from opposite sides, so every stitch locks the one before. Cut it and it stops. No machine does this. That one mechanical fact is why the craft still turns on a piece of toothed steel you hit with a mallet.¹

The loop is short and learnable: cut, **skive** (thin the leather where it folds), glue, mark, pierce, saddle stitch, then bevel and **burnish** the edge until it is glass. The skill that separates levels is not the stitching. It is the knife work.¹

The scale is real and lopsided. The biggest subreddit runs **888,000 members** but grew only **3.5%** this year; its beginner-facing sibling has **68,000** and grew **31.8%**. The old guard — saddlery, holsters, tack — lives on a forum with **98,052 members and 740,900 posts** since 2006. The centre of gravity has moved to video: a Korean atelier channel has **1.2 million subscribers**.²

“We have a lot of traffic, but most of it is window shopping.”

Which is the tension this desk came away with. The forum that holds the craft’s technical memory is openly discussing its own decline. In a thread begun in February and still running in September, one moderator counted **602 users online and six logged in**; another gave the diagnosis: **“YouTube also takes its share of folks who would rather just watch than participate.”**³

SOURCES 1 The mechanism of the saddle stitch and the order of operations are as the craft’s principal forum and the teaching channels at note 2 describe them; the forum’s own board structure, read at source on 9 October 2026, is the best published map of the craft’s sub-disciplines. 2 Subreddit membership and annual growth are from a third-party analytics service, dated 30 September 2026 and labelled as such, because Reddit itself refused this desk on every route tried, including two mirrors. The forum’s member, topic and post counts are displayed on its own index, read at source. Channel subscriber counts and last-upload dates were read on each channel’s own page on 9 October 2026. 3 A thread in the forum’s own feedback board, opened 20 February 2026 and last answered 21 September 2026, read at source; posters are quoted as they sign themselves, and the online-user counts are one moderator’s own observation in that thread. Companion threads on hosting costs and server response, dated June and August 2026, are in the same board. 4 A forum thread of June 2022 on whether a pricking iron can be used as a stitching chisel, read at source; both positions are quoted from named posters. 5 The United States producer price index for hides and skins including cattle, downloaded as published data; its last observation is December 2025. The 22 per cent projection and the bootmaker’s figure are from a San Antonio television report of 16 January 2026, read at source. The retailer’s figures and the quoted explanation are from its own second-quarter release of 11 August 2026, filed as an exhibit with the Securities and Exchange Commission and read at source. 6 The host guild’s own public fundraising page, created 4 May 2026 and read on 9 October 2026, for the amount raised, the goal and the donor count; the show’s dates and venue are from the forum’s own events calendar.

THE CHATTER · THE ARGUMENT THAT DEFINES IT

Ask what the difference is between a **pricking iron** and a **stitching chisel** and you have started it. The traditional position, from a long-running forum thread: a pricking iron **“just does what it says, makes a small prick marker to give the equal spacing for the awl to follow”**, with blades **“maybe 5mm at max”**. The pragmatic reply, same thread: they **“will penetrate thin leather, especially if they are sharp”**. One marks; the other punches through and skips the awl entirely. Which you own decides which tradition you joined without being told.⁴

THE SCEPTIC · HIDES GOT CHEAPER; LEATHER GOT DEARER

The story everyone tells is hide-price inflation, and the producer-price index says the opposite. US prices for hides and skins went from **110.8 in January 2023 to 70.6 in December 2025 — down 36 per cent**, with nothing published since. What rose is the finished object: a 111-year-old Texas bootmaker said prices had risen **by more than \$200** a pair. The mechanism is in a retailer’s own accounts: the largest hobby chain reported sales up **1.5%** and gross margin up **5.3 points to 64.8%**, **“primarily due to better pricing of our products”**. The raw material got cheaper and you did not.⁵

\$300

RAISED OF A \$4,500 GOAL, FROM THREE DONORS, BY THE GUILD HOSTING THE SIXTIETH INTERNATIONAL LEATHER-GUILD SHOW, RUNNING THIS WEEK⁶

The next two pages go down: where the craft gathers, the gear from the tannery up with real prices, the custom tier and what is actually waiting for you there, and the ninety days that get you from nothing to a wallet you would carry.

The leather first, because it decides what you can do. **Vegetable tanning** fixes the hide's collagen with plant tannins over weeks to months; the result holds a crease, takes tooling, burnishes to a polish and patinas. **Chrome tanning** uses chromium salts and takes hours; softer, more uniform, water-resistant, cheaper — and it will not burnish, will not tool, and must never be put under a laser. At one British supplier the gap is plain: undyed veg-tan shoulder **£6.75** a square foot against undyed chrome side at **£2.50**.¹

Above that sits the material the obsessives talk about. **Shell cordovan** is a dense membrane from a horse's hindquarter, vegetable-tanned at least 60 days then hot-stuffed with grease; the full process runs a **minimum of six months**, each hide yields **two ovals**, and the raw material runs about **\$100 a square foot**. A hobbyist buys it by the piece: Italian shells **\$89.99 to \$298.99**. Or a bag of offcuts at **\$80**, which is how most people meet it.²

For everything else the default is Tuscan veg-tan. A half-hide of Buttero — 2.8mm, firm, drum-dyed through — is **\$119.99** for six to eight square feet; a panel is **\$64.99**; a single square foot **\$24.99**. American bridle and harness leather is sold differently again: at one British house a **whole pair of bridle butts is £400**, not priced by the foot at all.¹

“There is only so much you can do with a pricking iron, and mainly it's the finish that costs the big bucks.”

Then the status tool, where the money actually goes. French irons from the traditional house run **\$64.99** for two teeth to **\$399.99** for fifteen — and on the day this desk looked, **only 10 of 39 variants were in stock**. A Korean maker starts at **\$45**. A Dutch retailer lists a two-piece Kevin Lee set at **€130.95**, backordered. One celebrated brand has been **retired outright**, its seller writing: **“this classic has been retired, but we keep moving forward.”**³

And here is the Singapore fact. **Crimson Hides**, which the community's own tier-lists put on the top shelf alongside the Korean and Japanese makers, is **based here** — its prices are in Singapore dollars, it says, “where we are based”. Its Japanese-style irons are **S\$155** in DC53 steel, 2.7 to 5mm, with 14-tooth sizes in ready stock and about **three weeks** otherwise, on **265 reviews**. One of the four names the hobby covets ships from this island.⁴

THE CHATTER · THE CLONES, AND WHAT A FILE CAN DO

The tool-tier argument is the hobby's class system, and its own members keep puncturing it. On the Chinese clones, from a 2018 thread still cited: **“Spend half a hour with a fine file between the teeth and they will do a very reasonable job like most others”** — then the line that undercuts the whole ladder: **“There is only so much you can do with a pricking iron and mainly it's the finish that costs the big bucks.”** The tier-list that circulates puts Sinabroks, KS Blade, Crimson and Amy Roke on the top shelf and Kevin Lee in the middle — and advises spending the money on **knives, edge bevellers and irons**, in that order.⁵

THE SCEPTIC · WHAT A LASER WILL AND WILL NOT DO

The hobby's newest tool got a mixed verdict in its own forum. A thread opened in October 2025 and still running in June 2026 collects real tests. One owner of a 60-watt carbon-dioxide laser: **“great success with cutting and engraving leather”**. One with a diode: **“It's STINKY!!”**, cutting slowly, needing the char sanded off, with 8-to-10-ounce veg-tan taking nearly as long as eighth-inch rosewood. A third, with air assist and a vented enclosure: **“the edge is clean... no soot... no stink.”** And the safety line that matters: **“I would not cut Chrome Tanned leather with the laser; the chemicals can be nasty when burnt!”** The settled use is not cutting parts at all — it is acrylic templates and engraved tooling patterns.⁶

29 of 39

FRENCH PRICKING-IRON VARIANTS OUT OF STOCK AT ONE MAJOR SUPPLIER ON 9 OCTOBER — WITH KEVIN LEE BACKORDERED AND AMY ROKE RETIRED. AVAILABILITY IS THE TAX, NOT PRICE³

SOURCES 1 Tanning mechanisms as described by the suppliers' own product copy and the forum sources at note 6; the vegetable-tanning certification named on one Italian tannery's listing. British per-square-foot prices, and the bridle-butt pair price, from a Newcastle supplier's own hides page, read at source on 9 October 2026, first figure excluding value-added tax. The Buttero half-hide, panel and square-foot prices, with the thickness, temper and dyeing description, are from an American supplier's own product page, read at source the same day. 2 The cordovan process, the minimum tanning periods, the two-ovals-per-hide yield and the approximate raw cost per square foot are from a footwear title's explainer updated 3 May 2026, read at source. The Italian and Japanese shell price ranges are from one supplier's own variant-level product data; the offcut bag from another supplier's own listing. 3 Iron prices, tooth counts and stock status read at variant level on the supplier's own pages on 9 October 2026; the Korean maker's own site; the Dutch retailer's own listing for the backordered set and its quoted retirement notice for the other brand. The retirement date itself is not published and is not printed. 4 The Singapore maker's own product page and its own frequently-asked-questions page, read at source, for the quoted basing statement, the steel, the sizes, the stock and lead times, the price and the review count. The tier-list placing it on the top shelf is a Reddit comment surfaced through a search index, not read at source, and is attributed as such. 5 A forum thread of September 2018 on one clone brand, read at source; all three quotations are from named posters there. 6 A forum thread opened 4 October 2025 and last active 8 June 2026, read at source; the four quoted verdicts are from named posters, with their machines as they describe them. Reddit refused this desk on every route, including two mirror services; one tool maker's site reset the connection and one supplier blocked it outright, and no figure from either is printed.

The custom tier is not where you think it is. In most crafts the flex is a commissioned object; here it is a **tool**. One American knifemaker prices round knives from **\$45 to \$275** and currently lists every priced item **out of stock**. A Wyoming maul maker warns of **two to ten business days** and takes orders by telephone. The Singapore iron maker runs an open pre-order list for a creasing iron. The waiting list is for irons and knives — not for wallets.¹

Where the money really goes at the top is machinery, and the prices are published. A 441-class stitcher runs **\$3,397 to \$4,357** from one house and **\$4,295 to \$6,495** from another; a five-ton clicker press is **\$2,499** and a ten-ton **\$4,987.50**. One of the best-known machine brands publishes no prices online at all and asks you to telephone — in this corner of the craft, the headline price is still a phone call.²

So the economics, built only from prices read at source. A **credible quality first kit is about \$380**: a nine-tooth French-style iron set at \$99.99, a diamond awl at \$25.99, harness needles at \$6.99, braided polyester thread at \$13.99, a table stitching pony at \$34.99, a skiving knife, a maul, burnishing compound and a bottle of dye — plus a \$64.99 panel of Italian veg-tan. A **serious hobbyist lands near \$1,500 to \$2,000**. The **grail tier is \$12,000 to \$18,000**, and almost all of the difference is the machine and the clicker press.³

“Saddle stitching, saddle stitching and more saddle stitching.”

Ninety days, honestly. Spend **two weeks buying nothing** and reading the getting-started and sharpening boards. In week three **take one class** — in Singapore a card-holder workshop is **S\$75 to S\$95**, and a three-hour session using Hermès leather with a six-guest cap is **S\$185**. Then buy the \$380 list and make the same wallet four times. What not to buy before day ninety: a \$400 iron, a round knife, a maul over \$90, a clicker press, a sewing machine.⁴

THE CHATTER · THE TROPICS ARE THE REAL CONSTRAINT HERE

This is the part a Singapore buyer has to plan for, and the hobby is quiet about it. Conservation guidance is not: mould on leather needs **relative humidity above 65 per cent** to grow, and the recommended storage range is a stable **45 to 55 per cent at 18 to 20°C**. The trade's own figure for inhibition is about **70 per cent**, with the warning that goods “**should not be close wrapped in polythene**” but in something that lets the leather breathe. A national standards study of vegetable-tanned strap leather in humid tropical conditions found mildew growth “**from moderate to very heavy**”, stiffening the leather and weakening it. Singapore sits above 65 per cent most of the year. Veg-tan here lives in a dry box, not on a shelf. No local hobbyist thread on this could be found.⁵

THE SCEPTIC · THE CRAFT'S OWN ANSWER TO ITS TOOL LADDER

Asked what a beginner should not skimp on, the forum answered mostly by naming what not to buy. “**‘Professional’ tools won’t make you a professional, but good quality tools will make it easier to learn the craft**” — followed immediately by “**BUT make sure the tool is worth the money and you aren’t paying for the hype around the name.**” Another, on the collecting trap: “**Skimping can also mean not buying tools you don’t need.**” Asked the same question, a third gave the whole answer in four words: “**Saddle stitching, saddle stitching and more saddle stitching**” — and told the beginner to go and buy cheap irons today. On selling the results, the pricing threads converge on doubling materials and adding an hourly rate of about **\$50**, with the honest coda that until you are known, you cannot charge as though you were.⁶

NEXT DESCENTS

Lockpicking & physical-security sport — a craft whose whole point is that the lock is a claim.

Bookbinding & book repair — sewn signatures, paste and board, and a trade that nearly vanished.

Longboard & distance-skate commuting — the deck as transport, and the push that goes twenty miles.

SOURCES 1 The knifemaker's own store, read at source 9 October 2026, for prices and stock status; the maul maker's own product page for the price and the quoted delay notice; the Singapore maker's own listings for the pre-order and lead times. This desk could not verify any leather-specific commission terms — deposits, waitlists or lead times for finished goods — at source, and prints none; one celebrated knifemaker's site reset every connection. 2 Machine prices from two suppliers' own current listings, read at source; the third maker's own site lists models without prices and directs buyers to telephone. 3 Every component price is read at variant level on the suppliers' own pages on 9 October 2026; the three totals are this magazine's own arithmetic from those prices and are labelled as such. 4 Singapore workshop prices from the two providers' own booking pages, read at source; the ninety-day shape follows the forum's own getting-started advice at note 6. 5 The Canadian Conservation Institute's published note on the care of vegetable-tanned leather for the 65 per cent threshold and the recommended storage range; a leather-trade publication for the roughly 70 per cent figure and the quoted packaging warning; and a United States national standards study of vegetable-tanned strap leather under humid tropical conditions for the quoted mildew finding. No Singapore or Malaysian hobbyist discussion of this could be found, and none is invented. 6 A forum thread of 25 December 2020 on what a beginner should not skimp on, read at source, for the first three quotations and the fourth; the pricing consensus is drawn from four named pricing threads on the same forum, surfaced through a search index rather than opened individually, and is characterised rather than quoted at length. Reddit was unreachable throughout.

FITNESS · GRIP PLATEAUS, FITNESS DOES NOT

Your watch is about three millilitres optimistic

A validation study measured twenty runners before and after twelve weeks. The watch read high both times, then promised a half-marathon seven and a half minutes faster than they ran.

Twenty people trained twelve weeks for a half-marathon wearing the same running watch, tested against a laboratory criterion before and after. At baseline the watch read **43.95** against a measured **41.10** millilitres per kilogram per minute — **2.85 high**, and significant. After twelve weeks, still **2.57 high**. The bias did not wash out with training.¹

Then race day. The watch predicted **129.97 minutes**; the official times averaged **137.55**. The authors are fair — mean absolute error stayed under 10 per cent, and they conclude the device “**may suit tracking general fitness trends**”, which is not the claim most people act on.¹

Why the optimism matters is on the next page of the literature. In **562,234** adults followed a mean **10.2 years**, each one-MET increase in measured fitness carried a hazard ratio of **0.84**, and adjusted ten-year mortality ran **32.9% in the least fit against 13.4% in the fit**.²

Grip strength tops out at thirty-two kilograms. Aerobic fitness does not top out at all.

And it argues against the fashionable metric. In **10,462** participants followed a median **12.5 years**, aerobic fitness was inversely associated with mortality **with no plateau**. Handgrip was weaker and **did plateau, at 32 kilograms**. Fifteen-year survival ran **39.3 to 77.4 per cent** across fitness, only **49.8 to 71.7** across grip.³

THE CHATTER · SAVING THE MUSCLE, AND WHAT IT BOUGHT

The argument has moved on from whether the new weight-loss drugs cost you muscle to whether protecting it achieves anything. On **3 October**, Peter Attia's team worked through a phase-two trial of a myostatin inhibitor: **102 adults**, both arms on tirzepatide for 24 weeks, **both losing about 12 kilograms**. Lean mass lost: **1.6 kg against 3.5 on placebo**. Grip: **+0.9 kg against -0.3, not significant**. The verdict: “**they don't deliver what's most important about muscle: Strength**”.⁴

THE SCEPTIC · A NULL PRIMARY, IN EXACTLY THIS AGE BAND

Published on Friday, and worth reading before the next supplement cycle. A double-blind trial gave **1,000mg a day of urolithin A for four weeks to 50 adults aged 40 to 64** with obesity. The primary outcome, flow-mediated dilation: **$\beta = 1.74$, interval -1.33 to 4.81 , $P = 0.33$. Null**. A secondary cerebrovascular measure did favour the active arm, with no interval given.⁵

0.662 → 0.682

WHAT MEASURED FITNESS ADDED TO A CONVENTIONAL RISK MODEL'S DISCRIMINATION, IN 562,234 ADULTS. THE MORTALITY GAP IS HUGE; THE ADDED INFORMATION IS TWO HUNDREDTHS²

The protocol worth trying. Eight weeks, 45 trained men, squats three days a week: **three single sets of 20 spread across the day** against **one session of three sets of 20**, volume matched. Both beat the control — **+20.59 kg** and **+28.93 kg** on one-rep maximum — the difference not significant, the fragmented arm rated **more enjoyable**. Permission, not equivalence. Not medical advice.⁶

SOURCES 1 Howard and others, “Validity of the Garmin Forerunner 265 to predict VO2max and race performance times”, *Journal of Medical Engineering & Technology*, 1 October 2026, doi:10.1080/03091902.2026.2741912, read through the national library's own record service after the web interface imposed a cookie wall: the sample, the twelve-week design, both pairs of means with their significance, the race-time comparison, the error figure and the authors' own conclusion. 2 Myers and others, “Cardiorespiratory Fitness Reclassifies Mortality Risk Beyond Traditional Risk Factors: The ETHOS Study”, *Medicine & Science in Sports & Exercise*, 6 October 2026, doi:10.1249/MSS.0000000000004175, read the same way: the cohort, the follow-up, the hazard ratio per unit of fitness, the adjusted ten-year mortality figures and the change in the concordance statistic. The published interval around that hazard ratio is implausibly tight and is reproduced here as the abstract reports it rather than smoothed. 3 Tur-Boned and others, “Associations of Pre-Diagnostic Cardiorespiratory Fitness and Handgrip Strength With All-Cause Mortality and 15-Year Survival After Cancer Diagnosis”, *Mayo Clinic Proceedings*, 8 October 2026, doi:10.1016/j.mayocp.2026.09.023: the sample, the median follow-up, both hazard ratios against the cohort median, the plateau finding and both survival ranges. 4 Rae, Fritsch and Attia, “Myostatin inhibitors and the paradox of larger but weaker muscles”, published 3 October 2026, read at source, including the trial figures it reports and both quoted sentences. 5 Pinaffi-Langley and others, “Urolithin A and Endothelial and Cerebrovascular Function in Middle-Aged Adults With Obesity: A Randomized Clinical Trial”, *JAMA Network Open*, published online 9 October 2026, doi:10.1001/jamanetworkopen.2026.38047, read at source: the design, the dose and duration, the sample and its mean age, the primary outcome with its coefficient, interval and probability, the secondary finding and the absence of an interval for it, and the reason for stopping at fifty. 6 Zhai and others, *Journal of Exercise Science & Fitness*, 2026;24(4):200521, read at source: the eight-week design, the three arms, the exact squat protocol and load progression, both improvements against control with their intervals, the three non-significant between-arm comparisons and the enjoyment result. Rest intervals are not reported in the abstract. This desk could verify no output in the window from the other commentators usually cited here, and attributes nothing to them.

DEPARTURES · THE LINK AND THE CERTIFICATE

Built, tested, and still not allowed to carry anyone

The cross-border rail link slips to February 2027. The infrastructure is “substantially complete”; the testing is not, and after that comes an independent assessor's certificate.

A joint release on **2 October** moved the Johor–Singapore rail link's opening to **February 2027**. The wording repays attention. Infrastructure on both sides is “**substantially complete**”. But final system tests and trial runs are not expected to finish until **31 December 2026** — and only then can the operator obtain certification from an **independent safety assessor**, which the authorities of both countries must review and accept before anyone rides. Three gates, in series, after the thing is built.¹

The previous public target was end-2026, and Malaysia's transport ministry had repeatedly said **1 January**. Its minister said on **8 October** that authorities “**did not want to take unnecessary risks**”, and does not expect further slippage. A Malaysian daily ran commuter reaction under the headline “**No more RTS Link delays, please**”.¹

The other date to put in the diary is **1 November**, and it changes how shopping in Japan works. Until **31 October** the consumption tax is deducted at the register. From 1 November you **pay the full tax-inclusive price in the shop** and claim the **10 per cent** back at departure, after customs confirms the goods are leaving. There is **no parallel-running transition**.²

If one item on a receipt is missing, nothing on that receipt qualifies.

The mechanics matter more than the headline. The **¥5,000** minimum per shop per day survives. Sealed packaging for consumables is gone; **bullion is now ineligible**. Within **90 days** you scan your passport at a terminal in the departure lobby **before checking**

bags — green means done, red means carry the goods to inspection. Each receipt is assessed separately, at your **final** point of departure from Japan.²

THE CHATTER · TWO LAUNCHES THAT KEEP NOT LAUNCHING

Singapore Airlines put **Riyadh** back a third time on 7 October, to **28 March 2027**; the original start was 2 June 2026 and the last target 1 December. The step up to four weekly in summer 2027 is dropped, taking weekly capacity from 160 business and 1,052 economy seats to **120 and 789**. Tickets held for December through March **will not operate**. **Dubai** has moved again too — now **2 December**, the sixth date set since February, returning two-class with **no First or Premium Economy** until late March at the earliest.³

THE SCEPTIC · THE HAZE HAS NOT TOUCHED A SINGLE FLIGHT

Worth stating plainly, because the inference is tempting. Central Singapore's index hit its season high on Friday and Johor closed **all state schools for two days** from 7 October after a reading of 203 the previous night. The motorsport authority said on the 8th that **outdoor entertainment** at this weekend's race could be restricted if conditions worsened, though track activity was not expected to be affected. But **no flight or airport disruption has been reported at Changi**. This desk went looking for it and found none, and will not imply otherwise. Haze degrades visibility long before it closes a runway, and nothing published this week says it has come close.⁴

31 Dec

WHEN THE OPERATOR'S FINAL SYSTEM TESTS AND TRIAL RUNS ARE NOW EXPECTED TO FINISH — BEFORE CERTIFICATION, BEFORE REVIEW, BEFORE ANYONE BOARDS¹

SOURCES 1 The joint release of the Singapore land transport authority and Malaysia's transport ministry, dated 2 October 2026, read at source on the authority's own newswire: the February 2027 target, the “substantially complete” wording, the 31 December date for tests and trial runs, the independent safety assessor and the two-country acceptance step. The release names no officials and publishes no journey-time or capacity figures. The minister's remarks of 8 October and the commuter reaction are from two reports in a Malaysian daily, read at source. 2 Japan's national tourism organisation's own explainer on the new tax-free system, read at source: the 1 November switch, the refund-at-departure mechanism, the absence of a transition period, the ¥5,000 threshold, the removal of sealed packaging, the bullion and used-consumables exclusions, the 90-day window, the terminal process and its colour outcomes, the per-receipt assessment and the final-departure rule. A companion claim that the departure tax tripled on 1 July 2026 is reported by two travel titles but this desk could reach no tax-authority or customs primary, and it is not printed as fact. 3 A Singapore aviation specialist's reports of 7 and 1 October 2026, read at source, which reproduce the airline's advisories; the airline's own newswire index returned a not-found error on every path this desk tried, and its budget subsidiary's media centre refused outright. Capacity figures, cabin configuration and the refund eligibility are as those advisories state them. 4 The environment agency's own haze portal, statement timestamped 17:24 on 9 October 2026, read at source; the Johor school closures and the motorsport authority's statement are as reported on 7 and 8 October. The absence of airport disruption is this desk's own negative finding after searching for it.

The winter schedule starts on **25 October**, and it is the most consequential document in Singapore aviation this month. **Madrid returns after 22 years**, five weekly on an A350 routed through Barcelona — the fifteenth European destination, and the end of the Singapore–Milan–Barcelona routing. The Barcelona–Madrid leg is a Spanish domestic sector and **cannot be sold or redeemed on its own**.¹

What is being taken away is less advertised. **Colombo halves**, from fourteen weekly to seven, withdrawing a daily narrowbody service that began in May. **Johannesburg halves** too. **Bali** drops to 35 weekly off-peak — seat capacity down over **20 per cent** year on year. Against that, **Western Sydney** launches daily on **23 November**.¹

And a genuinely new short-haul option: **Gulf Air** opens a fifth-freedom **Singapore–Kuala Lumpur** sector from **26 October**, daily, on a 787-9, loaded to September 2027. Round trips from **S\$246** in economy and **S\$1,025** in business for November. A widebody on a seventy-minute hop, because it is continuing to Bahrain anyway.²

Eight overseas lounges, down from thirteen before the pandemic.

The reveal to watch is **early November**: a **S\$1.1 billion** long-haul cabin programme covering **41 A350-900s**, adding First Class to the ultra-long-range aircraft, which go from 161 seats to **132**. Existing 777s will **not** be retrofitted. The first aircraft enters service in **the first quarter of 2027**, a year late.³

THE CHATTER · A LOUNGE CLOSES, AND THE COMMUNITY COUNTS THE COST

The Taipei lounge shuts after **31 October** — the network's smallest and oldest, in a terminal being redeveloped, with no reopening date. Eligible passengers go to a contract lounge, unnamed. That takes the overseas lounge network to **eight**, from thirteen before the pandemic. With the airline forums shut to this desk, the argument is in the comments under a Singapore miles site's report: **emercycrite** asked why the airline would pay a contract operator rather than use a Star Alliance partner's lounges; the site's owner, **Aaron Wong**, replied drily, **"maybe plaza premium charges less than EVA."**⁴

THE SCEPTIC · A TRANSFER PARTNER VANISHED

This is the one to act on. **Qantas quietly disappeared** from a major Singapore bank's rewards conversion list in late September; revised terms with the airline removed went up at about **seven in the evening on 1 October**, and the last references were gone by **3 October**. No notice was given. The lost ratio was **5,000 bank points to 10,000 Qantas points**; KrisFlyer and Asia Miles remain at the same ratio. Routes to Qantas from Singapore now: a card issuer at 20,000 or 22,000 for 10,000, no fee; two others with a **S\$27.25** fee; one at 25,000 with none. Qantas points expire after **18 months** of inactivity, so an orphaned balance is on a clock.⁵

8 a day

FIRST CLASS SEATS EACH WAY TO JAKARTA THIS WINTER, ON TWO OF NINE DAILY ROTATIONS — SIXTEEN LAST WINTER, FORTY BEFORE THE PANDEMIC⁶

IN BRIEF · THREE DATED CHANGES

Scoot adds Guiyang from 26 October, three weekly, one-way economy from **S\$158** including taxes — taking it to 81 weekly flights to sixteen mainland Chinese cities.⁷

The 3D seat-map rollout finished on 7 October, adding the long-haul A350 and the A380. One caveat the airline does not print: once retrofits begin, the render you see at booking will not match the seat you fly.⁸

And a nominee rule worth knowing before you book for someone. A redemption nominee added today cannot be booked on flights departing **within 72 hours**; five slots per account, each locked for six months.⁹

SOURCES 1 A Singapore aviation specialist's report of 20 September 2026 on the airline's northern-winter schedule, read at source, for the Madrid return and its routing and restriction, the Milan change, the Colombo, Johannesburg and Bali reductions and the Western Sydney launch; the airline's own release on the Madrid launch could not be opened and its exact launch-date wording is therefore not quoted. 2 The same specialist's report of 6 October 2026, read at source, for the new fifth-freedom sector, its aircraft, frequency, loaded season and the quoted fares. 3 A long-haul aviation title's report of 1 September 2026, read at source, for the programme value, the aircraft count, the seat-count changes, the exclusion of existing 777s and the revised entry-into-service quarter. The reveal event itself is reported rather than confirmed by the airline, and no date is stated here as official. 4 The specialist's and a Singapore miles site's reports of 5 October 2026, both read at source, for the closure date, the terminal redevelopment, the contract-lounge arrangement, the lounge count and the Melbourne slippage. The quoted exchange is from the comment thread under the second of those, read at source; commenters are named as they sign themselves. The two airline forums this desk would normally read both refused it. 5 The specialist's report of 8 October 2026, read at source, for the removal, its timing, the lost ratio, the remaining conversions and the alternative routes with their fees. The bank gave no notice, which is the finding. 6 The specialist's report of 6 October 2026 on the winter First Class allocation, read at source. 7 Reported 20 August 2026 and read in two regional titles, the airline's own media centre having refused this desk; the service begins inside this fortnight. 8 The specialist's report of 7 October 2026, read at source, including the mixed-fleet caveat. 9 The same site's guide of 4 October 2026, read at source, which records the rule added to the airline's own answers in August 2026.

THE GRAND TOUR · THE ATACAMA

You can walk into an eight-metre telescope. Free


Paranal runs free guided tours every Saturday of the year, booked a quarter ahead. ALMA, the one everyone pictures, is shut.

Begin with the correction, because most planning starts from a false premise. **ALMA's public visit programme is suspended**, as its own page says and Chile's tourism board confirmed in February. And the antennas people picture, on the **5,000-metre** Chajnantor plateau, **were never open**: the visits that existed went to a support facility at **2,900 metres**.¹

What **is** open is better. **Paranal**, at **2,635 metres**, runs guided tours **every Saturday of the year**, in two slots, **free of charge**, booked online only, and you walk into the dome of one of the **8.2-metre** Unit Telescopes. It has run these since 1996, for **more than 7,000 people a year**. Two changes: booking now opens in the **first week of the month before each quarter**, and **night tours exist** — the last of this year is **tonight**, announced on Instagram, not the website.²

Paranal is not near San Pedro. It is a separate trip altogether.

Which is the geography nobody tells you: Paranal sits about 120 kilometres south of Antofagasta, on the coastal range — a separate leg, not a San Pedro day trip. From Singapore the routing is **one stop through Australia**, about **23 and a half hours**, then two hours to Calama and **100 kilometres** of road: **thirty to thirty-four hours** door to door.³



THE CHATTER · THE ALTITUDE IS THE ITINERARY

San Pedro sits at **2,407 metres**. The public-health threshold for altitude illness begins at a sleeping altitude of **2,450 metres “and sometimes lower”**, so the town sits on the line while the excursions go to **4,320**. The saving grace is that these are day trips: you touch the height and sleep back down, so the **order** matters more than the ceiling. One account of getting it wrong: **“Booked Tatío for day two and spent the entire morning at 4,300 m with a headache so bad we couldn’t focus”**.⁴

THE SCEPTIC · IS THE SKY ACTUALLY THE BEST PART?

The informed dissent is not about darkness but **transparency**. On an astrophotography board, an experienced visitor writes that the sky at San Pedro **“is not as clear in terms of transparency as it is closer to the coast, where the optical observatories are”**. The professionals' own siting says the same: Paranal and La Silla sit on the coastal range, **not** on the altiplano. And nobody publishes a sky-brightness number for the town: the standard meter is at its limit beyond 21.5 magnitudes per square arcsecond, so a Bortle 1 for the Atacama is an instrument that cannot resolve it.⁵

2,407 m

SAN PEDRO'S ALTITUDE, AGAINST THE 2,450-METRE SLEEPING HEIGHT AT WHICH RISK BEGINS⁴

SOURCES 1 The observatory's own visits page, read at source, which states the public programme is suspended; the European observatory's own visitors page for the same site, which states flatly that public visits are not currently possible — note those two pages contradict each other, the second still carrying the legacy tour description; and Chile's official tourism page, updated 10 February 2026, which records the suspension and work towards reactivation. The support-facility and plateau altitudes and the never-open status of the high site are from the observatories' own pages. 2 The European observatory's own Paranal visitor page, read at source, for the altitude, the Saturday schedule and both arrival slots, the free admission, the online-only booking, the telescope aperture, the age rules and the annual visitor figure; and its own announcement of 2 September 2026 for the new quarterly booking windows, the free transport arrangements and the night-tour programme, including the 10 October date. 3 Flight times and the one-stop routing are from a Singapore fare-comparison page carrying a 9 October 2026 data stamp and from a published schedule analysis of the Sydney–Santiago service, with the block time computed here from its departure and arrival times; the Calama transfer price is a registered local operator's own published rate, read 9 October 2026. The door-to-door range is this desk's own arithmetic and is labelled as such. 4 The town's altitude from its own municipal and encyclopaedic record; the risk threshold, the quoted phrase and the single-day ascent guidance from the United States public-health travellers' handbook, read at source. The quoted traveller account is a 2026 travel guide, read at source and treated as lived texture rather than authority — its own cost figures contradict each other elsewhere. 5 A quotation from an astrophotography forum thread, obtained through that page's search index because the page itself renders only by script and returned an empty shell; it is attributed as such. The siting of the professional observatories is from the European observatory's own pages, and the meter's limit from the United States national park service's own published note on sky-quality metrics. Tripadvisor's destination forum, Reddit, one astronomy board and the web archive all refused this desk.

Everything here is a separately ticketed concession, most run by Atacameño communities, most with a foreigner tier. The clearest proof is the forestry service's own table for the Altiplanic Lagoons: a foreign adult pays CLP 8,000 and a Chilean CLP 4,000 — **exactly double** — with holders of an Indigenous certificate at 1,000. The Atacama is not expensive because Chile is; it is expensive because every sight is its own gate.¹

The prices, converted at Friday's rate of **CLP 762.8 to the Singapore dollar**. **El Tatio**: entry CLP 15,000 (S\$19.70), shared tour CLP 35,000 (S\$45.90), leaving at five, dawn around 0°C, entry paid **in cash the day before**. **Valle de la Luna**: CLP 10,800 (S\$14.20), booked a day ahead. **Altiplanic Lagoons** tour CLP 76,000 (S\$99.60). A family of three doing the five signature excursions pays roughly **S\$790 to 920** in tour fees and **S\$130 to 170** in entries before a single night's bed.²

Two corrections to the guidebooks. You **can no longer swim in Laguna Cejar**: one ticket covers two lagoons, Cejar is viewed from walkways, and bathing is only in adjacent **Laguna Piedra** — CLP 15,000, **not suitable for under-sixes**, about an hour on site, no sunscreen beforehand, nothing sold there. And **Valle de la Luna is closed on the Tuesdays of this October** for staff training.³

“A tourist trap, but the scenery more than makes up for it.”

Where to sleep splits cleanly. The all-inclusive lodges buy guiding, a private vehicle and no five-a.m. minibus — one runs its own **40-centimetre telescope** — at roughly **US\$910 to 2,090 a person a night**. Two people doing it from town for a week land nearer **US\$2,000 to 3,000 all in**, and the lodges buy access to nothing that is closed to others. One caveat: a third-party listing puts one celebrated lodge at **14 and over only**, which its own pages do not state — confirm before booking with a child.⁴

THE CHATTER · THREE DAYS, OR SEVEN

The classic forum line is that three days is enough. The most detailed recent account disagrees flatly: **“Fewer than 4 and the altitude bites back”, “if you can spend 6 or 7 days in Atacama, spend them”,** and on the compressed version, **“we'd seen a highlights reel instead of the actual place”**. On the tours, the useful concession comes from a traveller who normally dislikes them: **“San Pedro is the exception”,** because **“it would be extremely hard to do San Pedro on your own because of the terrain, weather”**. His summary is the fairest sentence anyone has written about the place: **“A tourist trap, but the scenery more than makes up for it.”**⁵

THE SCEPTIC · WHAT UNDERWHELMS, NAMED

Four honest disappointments, each from a source rather than a mood. The **Salar de Atacama** is **“not as impressive as those found in Uyuni, Bolivia”** — and the Uyuni fix is a **three-day** trip with two border crossings and hostel rooms above 4,000 metres, the worst possible fit inside an altitude-adjustment week. Agencies have **“shortened some of the interior walking circuits”** at Valle de la Luna, which is where the payoff is, and the visit is **“completely doable in your own rental car”**. The classic Salar de Tara route has quietly been swapped by many operators for a broader circuit. And on restaurants: the town's best-known gastronomic address carries **167 one-star and 152 two-star reviews** against 2,479 five-star, while the highest-rated place by volume is a **French bakery**.⁶

2×

WHAT A FOREIGN ADULT PAYS AGAINST A CHILEAN AT THE ALTIPLANIC LAGOONS, ON THE FORESTRY SERVICE'S OWN PUBLISHED TABLE: CLP 8,000 AGAINST 4,000¹

And the timing, which for this reader is an astronomy problem wearing a school calendar. From latitude -23 the **Magellanic Clouds ride high from November to February**; the galactic centre does not — that is June to August, when El Tatio dawns are below freezing. Commercial star tours run **CLP 25,000 to 45,000 (S\$33 to 59)**, and none run in the seven nights around a full moon. The new-moon calendar and the tour calendar are the same calendar.⁷

SOURCES 1 The national forestry service's own reserve page, read at source, for the two-tier entry table in full. 2 The same page and the operators' own published list prices, read 9 October 2026; the exchange rate is a public rate feed stamped 9 October 2026, and every Singapore-dollar figure is this desk's own conversion at it. The family totals are this desk's arithmetic from those list prices and are labelled as such. The cash-the-day-before condition is the operator's own stated term. 3 The lagoon concession's own site, read at source, for the bathing restriction, the prices, the age limits, the timings and the on-site rules; the forestry service's own page for the October Tuesday closures and the booking requirement. 4 Lodge rates are third-party editorial and booking-platform figures, clearly secondary: neither of the two best-known lodges publishes rates on its own site, and one of their rate pages returned a not-found error. The telescope is named on that lodge's own page. The à-la-carte comparison is this desk's arithmetic from the list prices at note 2. The age restriction is a specialist agency's listing and is **not** stated on the lodge's own site or in its own answers page — it is printed here as the unresolved item it is. 5 A 2026 travel guide read at source for the first three quotations, and a personal travel account of 1 March 2011 for the last three; the second is old, and is used only for sentiments about terrain and character that have not changed. 6 The same 2026 guide for the salt-flat comparison, the shortened circuits, the rental-car point and the operator substitution; the Bolivia itinerary from a registered operator's own multi-day listings. The restaurant review distributions are from an aggregation of two public review platforms, read at source, and are quoted as counts rather than as sentiment. 7 Object coordinates are the published positions of the two galaxies and of the galactic centre; the seasons follow from them at this latitude, and the computation is this desk's own. Tour prices and the full-moon suspension are the operators' own published terms.

THE DIARY · THE DATED AGENDA

Japan is already selling January. Singapore has not priced it

Every row below was re-checked against a venue, promoter or official seller. Nine candidates failed verification and are named at the foot.

SINGAPORE

BOOK NOW · PLAN AHEAD

- 21 OCT–1 NOV** ★ **37th Singapore International Film Festival**
PUBLIC SALE OPENED 28 SEP — **LIVE NOW**. UNDERCURRENT PASS **\$50**; CLOSING-NIGHT SEATS FOR *VINTAGE VIOLENCE* RELEASE **11 OCT**¹
- 10–11 NOV** **My Chemical Romance — The Black Parade**
INDOOR STADIUM · FROM **\$148** · MOVED FROM APRIL; APRIL TICKETS TRANSFERRED TO THE 10TH. THE SELLER'S PAGE RENDERS NO PRICE THIS DESK COULD READ²
- 26–29 NOV** **Affordable Art Fair**
F1 PIT BUILDING · PRIVATE VIEW THU 2–10PM, THEN 12–10 AND 10–7 · BOOKING LIVE; **ADMISSION PRICE NOT PUBLISHED**³
- 28–29 NOV** **BABYMONSTER — 춤 (CHOOM)**
ON THE PROMOTER'S OWN LIST · FROM **\$168** · THE 28TH REPORTED SOLD OUT, THE 29TH AVAILABLE — NEITHER VERIFIABLE AT A PRIMARY²
- 3–6 DEC** **Voices — A Festival of Song**
ESPLANADE'S OWN LISTING · VENUE AND PRICES NOT YET PUBLISHED³
- 22–24 JAN** ★ **ART SG, fifth edition**
SANDS EXPO · VIP PREVIEW 21 JAN · **NO PUBLIC TICKET PRICE PUBLISHED AT ALL**, AND IT FALLS INSIDE THE AUSTRALIAN OPEN FORTNIGHT⁴
- to 14 MAR** **Disney Garden of Wonder, Floral Fantasy**
GARDENS BY THE BAY · RESIDENT **\$12 / \$8** AGAINST **\$24 / \$16** · 10AM–9PM · **CLOSED 13 OCT** FOR MAINTENANCE⁵

JAPAN & THE REGION

THE ONES WITH A TICKET CLOCK

- 8–22 NOV** ★ **Kyushu Basho — and it has gone**
FUKUOKA · ON SALE 19 SEP; THE ASSOCIATION'S OWN ENGLISH SITE NOW READS **"TICKETS ARE SOLD OUT ALL DATES"**. CHAIR SEATS HAD RUN **¥2,000–6,000**⁶
- 14 NOV–28 MAR** ★ **Masterworks from the Musée d'Orsay**
TOBIKAN, UENO · THE CENTENARY SHOW · TWO FREE COMPANION EXHIBITIONS FROM 19 NOV · **ADMISSION NOT YET PUBLISHED**⁷
- 28 OCT–8 FEB** **Shojo Manga Infinity**
NATIONAL ART CENTER, ROPPONGI · HAGIO, YAMAGISHI AND YAMATO · THE CROSSOVER PICK FOR A TWELVE-YEAR-OLD⁸
- to 13 DEC** **Renaissance Masterpieces from the Louvre**
NATIONAL ART CENTER · NOW ON VIEW, CLOSES MID-DECEMBER — AN AUTUMN TRIP OR NOT AT ALL⁸
- 8 JAN–4 APR** ★ **Final Fantasy, 40th anniversary**
HYOKEIKAN, TOKYO NATIONAL MUSEUM · THE MUSEUM'S OWN LISTING ALREADY READS **"NOW ON SALE"** — A QUARTER AHEAD OF OPENING⁹
- 7 MAR** **Tokyo Marathon 2027 — entry closed**
THE TWENTIETH RUNNING · GENERAL ENTRY RAN 14–28 AUG AND WENT TO A BALLOT; RESULTS 11 SEP. **NO FURTHER ROUND ANNOUNCED**¹⁰
- 9–11 APR** **Cathay/HSBC Hong Kong Sevens 2027**
KAI TAK STADIUM · REGISTRATION OPEN ON THE UNION'S OWN PAGE; **NO PRICE OR ON-SALE DATE YET**. FOR SCALE, 2026'S THREE-DAY PASS WAS HK\$2,250¹¹

PLANNING NOTES · THE ASYMMETRY, AND WHAT FAILED VERIFICATION

Buy Japan now; watch Singapore. That is the week's lesson, and it is in two rows above: a Tokyo museum is already selling an **8 January** exhibition, while a Marina Bay art fair ten weeks out has **published no ticket price at all**. The November sumo went on sale on 19 September and was gone across all fifteen days before most people read about it — the habit is not watching for the tournament, it is **setting an alarm for the on-sale**. **For the twelve-year-old:** the Disney display is the cheapest substantial family booking in the window, at a 50 per cent resident rate, and it runs to March. **Dropped for failing verification:** four arena shows whose only sources are secondary and whose two ticketing platforms both refused us; a December garden event rendering unfilled template placeholders; a marathon still carrying last year's branding; a February parade with no 2027 dates; a club festival with no edition announced; and the January Tokyo tournament, whose schedule page errors on every variant.¹²

SOURCES 1 The festival's own site, read at source, for the dates, the opened public sale, the pass price and the closing-film ticket release. 2 The promoter's own Singapore event list, read at source, confirms the name and dates of the arena shows; venue, price, the transfer of April tickets and the reported sell-out are from a Singapore concert guide last updated 30 September 2026 and are flagged as secondary. The ticketing platform's event page exists and is live but renders client-side, yielding only a title. 3 The fair's own page and the Esplanade's own festivals index, both read at source; neither publishes admission or venue detail, which is stated rather than filled in. 4 The fair's own site, read at source; the clash with the Melbourne tournament fortnight is this desk's observation from the two published date ranges. 5 The venue's own attraction page, read at source, for the run, both price tiers, the hours and the single-day closure. The venue's own page also warns that its non-resident ticket link currently points at an unrelated event, so confirm the price at checkout. 6 The sumo association's official English ticket service and its own English site, both read at source, for the dates, the on-sale date, the price bands and the quoted sold-out statement. 7 The museum's own exhibition index, read at source. 8 The gallery's own exhibition page, read at source. 9 The museum's own events listing, read at source, where the show appears under "Now on Sale". 10 The race's own site, read at source, for the date, the entry window, the ballot and the notification date. 11 The union's own event page, read at source; the 2026 price is given for scale only. 12 Each dropped item was checked against the promoter, venue, organiser or museum named, and is omitted rather than printed from a secondary source. The children's-space closure is on the Esplanade's own listings page. Two Singapore ticketing platforms and the national stadium's own site all refused this desk.

THE TABLE · FIVE NEW ROOMS

One of these has been reviewed. The rest are too new

Four of this week's five openings have no rated review and no settled verdict. Rather than invent one, this page prints what is known.

The one with a verdict first. **Blue Label Tavern**, opened this month at 61 Stamford Road, is the gastropub spin-off of a pizza-and-wine group — stone-fired pizzas, a dry-rubbed rotisserie line and tavern-only dishes, opposite the university. Buttermilk chicken dunks **\$15**, a smash burger **\$18**, the outlet-exclusive chicken plate **\$19** for a half. Closed Sundays, shuts at nine.¹

The single full review scores it **four out of five** — dunks “addictive”, the chicken tender, good for a group — and lands two criticisms worth more than the score. The pizzas here are **ten inches against twelve** at the sibling outlets; and the reviewer says plainly it “**might not be very clear what’s the distinction between Blue Label Tavern and Blue Label Pizza outlets**”.¹

Then two on the same Joo Chiat stretch, both too new to call. **Base Gelato** at 198 is a takeaway counter from a homegrown gelato brand: six flavours, several local — Joo Chiat Heritage, coconut with ondeh ondeh, **\$10**; kaya butter toast **\$9**. The binding constraint is the week: **Friday to Sunday only, noon to seven**. **Accompany Coffee** at 292, in the Grafunkt building, uses Uji ceremonial-grade matcha; an einspanner is **\$8.50**. It shuts at **half past four** on weekdays.²

Four of five have an announcement and a photograph. That is not a verdict.

Two more. **Din Tai Fung** has opened a hundred-seat room at ION Orchard built as a replica of its original Taipei shop, with two dishes found nowhere else: wantons in peanut sesame sauce at **\$14.80** and in spicy sauce at **\$12.80**. The one published assessment

recommends it “**mainly for the two exclusive wonton dishes**” — a verdict about two plates, not a restaurant. And **Pull by TPG** at 60 Joo Chiat is a Muslim-owned bagel-sandwich cafe, seven till nine daily, sandwiches from **\$11.90**. The source says Muslim-owned; it does **not** say halal-certified.³

THE CHATTER · WHAT THERE IS, AND IS NOT

On the tavern, one rated review and no aggregate — and a trap for anyone checking: the listings that surface under its name belong to the **pizza** outlets, not to this room. On the gelato counter, the write-up predicts a flavour “will probably be a hit”, which is forecast, not verdict. There are **no queue reports, no social consensus and no dissent on record** for any of the four. That is the honest state of a Singapore opening in its first fortnight, and printing enthusiasm into the gap is how a restaurant page stops being worth reading.⁴

THE SCEPTIC · ONE THAT READS NEW AND IS NOT

A Japanese tea and wagashi counter on Joo Chiat is appearing in this month's round-ups and reads like an October opening. It is not: it was reviewed on **10 July** and again on **4 August**, so it opened around July and one round-up has filed it late. It is left off this page for that reason, not on its merits. The caution follows: an aggregator's “new openings” list is a publication schedule, not an opening schedule, and at least one entry in most months is a re-run.⁵

TOO NEW TO CALL ×4

ONE RATED REVIEW

TWO ON ONE STREET

WEEKEND-ONLY HOURS

The planning, then. None of the five takes reservations. The tavern empties at nine; the gelato counter opens three days a week; the matcha cafe shuts before most people leave work. The dumpling room is the safe bet with a twelve-year-old.²

SOURCES 1 A Singapore food blog's full review of 7 October 2026, read at source, for the scores, the quoted criticism, the pizza-size and selection comparison and the dish prices; the address, hours and closed day are from the same review. The group's other outlets are named there. 2 A Singapore food title's write-up of 9 October 2026 and a lifestyle title's new-cafes round-up of 7 October 2026, both read at source, for the two Joo Chiat venues' concepts, operators, flavours, prices and opening hours. The walking distance is from the two published addresses. 3 A Singapore food title's monthly round-up of 1 September 2026, read at source, for the dumpling outlet's seat count, the Taipei replica conceit, the two exclusive dishes and their prices, and the quoted recommendation; and the lifestyle round-up at note 2 for the bagel cafe's prices and hours. That source describes the cafe as Muslim-owned and does not state halal certification. 4 Each verdict above is characterised from the pieces named in notes 1 to 3; where a piece is an announcement rather than a review this page says so. No forum or social consensus could be found for any of the four, and none is invented. 5 Two Singapore reviews, dated 10 July and 4 August 2026, establish the earlier opening of the venue left off this page; the hours disagreement between two round-ups was resolved in favour of the reading corroborated twice. This desk also checked every candidate against its own standing record of venues already covered, and set aside three that had run too recently to earn a second appearance.

THE UNDERCURRENT · INDIA'S SARATHIS

Eight hundred thousand drivers, eleven thousand rides

India built the largest driver-owned ride-hailing cooperative on earth and it is doing 11,106 trips a day. Read as a business that is a catastrophe. It is not a business.

At Terminal 3 of Delhi airport a driver named Mukesh Kumar explained how he is paid. For the first couple of months the platform gave him **₹100 a ride** as a launch incentive. Then it stopped, and replaced the cash with **free wheat flour, rice and lentils, redeemable at Amul outlets**. He would prefer the cash. A ride-hailing company paying its drivers in flour from a dairy cooperative is not a metaphor; it is the operating model, and it is the first clue that this is not a ride-hailing company.¹

The company is **Bharat Taxi**, run by Sahakar Taxi Cooperative Limited, registered on **6 June 2025** as a multi-state cooperative with authorised capital of **₹300 crore** and promoted by eight national cooperatives — the fertiliser giants, the dairy federation behind Amul, the agricultural marketing and rural-development bodies. Drivers are **Sarathis**, from the Sanskrit for charioteer: Krishna was Arjuna's sarathi. They buy shares at **₹100**, minimum five, and sit on the board. The slogan is *Saarathi Hi Maalik* — the driver is the owner.²

Now the numbers, and they are the government's own. Answering in the upper house on **12 August**, the home and cooperation minister gave **848,000 registered Sarathis** and **4,427,000 registered customers**. In the same reply: **11,106 rides a day**. Delhi and its satellites account for 5,574 of them, Gujarat 3,817, Pune 632, Jaipur 372, Lucknow 324, Mumbai 197, Chandigarh 190. The city figures sum exactly to the total, which is how you know the number is real rather than rounded.³

One ride, per registered driver, roughly every seventy-six days.

Do the division. **0.013 rides per registered driver per day** — one ride each, about every **seventy-six days**. Cumulative Sarathi earnings to **4 August** were **₹83.70 crore**, which across 848,000 drivers is roughly **₹987 each** for six months of existence. And the enrolment has not slowed: between the July and August parliamentary answers, **21,000 drivers joined in a week** while daily rides stayed near eleven thousand.³

Jaipur is the whole thing in one city. More than **30,000 Sarathis** and nearly **700,000 registered customers** — against **372 rides a day**.³

The fee structure is the part everyone leads with, and it is the weakest part. Bharat Taxi charges **no commission**: either **₹50 a day** for unlimited rides, or **₹10 a ride** in a car and **₹5** in an autorickshaw. Against an aggregator's **₹139 plus tax** daily pass, that is real money. But the private platforms had already abandoned commissions for subscriptions, so zero commission is the same model with a cheaper sticker — which is why a rival's spokesman **welcomed** the launch, saying it supported his own company's view that subscription-led models suit the future of ride-hailing. Being endorsed by the incumbent you came to disrupt is a particular kind of review.⁴

A reasonable person reads that and concludes the thing has failed. The drivers mostly agree about the rides. At Terminal 1, Rajneesh Kumar gets about **fifteen Uber requests a day and two or three from Bharat Taxi**; Uber pays his rent. At Terminal 3, Kapil Pardal points out there is still **no digital wallet**, so incentives, payments and parking are squared up by hand — and says something more damaging than any ride count: that **nearly 90 per cent** of drivers are not educated enough to understand what owning a share in a cooperative means.¹


SOURCES 1 An Indian digital title's reported piece of 12 July 2026, read at source, which interviewed named drivers at both Delhi airport terminals; the incentive history, the grocery substitution, the ride-request comparison, the absent wallet and the quoted estimate about share ownership are all from those named drivers in that report. 2 The cooperative's registration date, legal form, authorised capital, promoting cooperatives, share price and minimum holding, and the slogan, are from the same report and from a second Indian title's profile of 20 August 2026, read at source, and corroborated against the company's own published answers page. 3 The driver, customer and daily-ride figures, the city breakdown and the cumulative earnings are from written answers given by the minister in the Rajya Sabha on 12 August 2026 and in the Lok Sabha on 28 July 2026, as reported by a cooperative-sector news service on 13 August 2026 and a second title on 6 August, both read at source. **The government's own press service refused this desk** (403 on the relevant release), so the parliamentary figures are cited to those reports of the answers rather than to the primary release. The per-driver and per-day arithmetic in this essay is this magazine's own, computed from those published totals, and is labelled as such wherever it appears. 4 Fee structures for both the cooperative and the private aggregator, and the rival spokesman's quoted welcome, are from the report at note 1, read at source.

Why now, and why a ministry. The Ministry of Cooperation is the home minister's own project, and its doctrine — *sahkar se samriddhi*, prosperity through cooperation — needed a flagship in a market everyone understands. What makes the structure unusual is what the state did **not** do: the minister has told Parliament that the government has given **no financial assistance, no equity support and no incentives**. The capital comes from the eight cooperatives, which is to say from the balance sheets of Indian agriculture. The state is lending its distribution, not its money.¹

The anger it is riding was already there. Drivers have spent two years in open revolt over commissions of twenty to twenty-five per cent. In Ahmedabad they struck for **four days from 7 September** under the banner *Rozgar Bachao* — save our livelihood — and a fare discussion on the 10th went nowhere. Bharat Taxi did not create that anger. It gave it an app.²

Then the per-kilometre argument, where the two sides disagree and both are on the record. Sarathis put their net take at **₹10 to 12 a kilometre** on the cooperative against **₹14 to 18** on the private apps; a senior Bharat Taxi official counters with **₹14 to 16**, and notes that a four-kilometre ride bills about **₹100** on every platform. What the cooperative cannot match is the cash on top: a rival pays **₹250 for eight weekend rides and ₹500 for twelve**, and a platform that has promised not to burn money cannot answer that.⁵

Nor is it reliably cheaper to ride. Three passengers interviewed the same day gave three answers: **15 per cent cheaper** than the incumbent in Delhi, **10 to 15 per cent dearer** according to another at the same terminal, and in Mumbai **20 to 30 per cent higher**. A price that depends on which city you stand in is not yet a price.⁵



SOURCES 1 The ministry's doctrine and the absence of state funding, equity or incentives are as the minister stated in a written answer to Parliament of 12 August 2026, reported by two Indian titles and read at source; the framing of the launch is from the state broadcaster's own report of 5 February 2026. 2 The four-day Ahmedabad strike, its banner, its grievances and the failed fare discussion of 10 September are reported by two Gujarati and Indian titles, read at source; the twenty-to-twenty-five per cent commission range is as the launch coverage states it. 3 An Indian technology title's report of 26 September 2026, read at source: the boycott date, the passenger's account, the described disruption, the honeypot bookings, the reported fines, and the title's own statements that it could not verify the accounts, that sources feared going to the police, and that the platform and both police chiefs did not respond. The two contradictory videos are described in the same report. Nothing in this passage is presented as established fact. 4 The insurance, health, portal and e-governance arrangements and the profit-sharing rule are from an Indian business title's profile of the venture, read at source. The figure above is drawn by this magazine from the parliamentary answer cited on **p25**; both bars share one scale and the ratio stated beneath them is arithmetic. 5 The per-kilometre figures, the official's counter-figures, the rival's weekend incentives and all four passenger price comparisons are from the reported piece of 12 July 2026 cited on **p25**, read at source, except the Ahmedabad boycott fares, which an Indian financial daily published on 22 September 2026 and which this desk read in a summary that names and timestamps it, the paper's own page being unreadable to us.

And then September turned ugly. On the **21st**, Ahmedabad's drivers began boycotting the private apps to force riders onto the cooperative. A woman arriving from Pune told an Indian technology title that minutes after her cab pulled away, **ten to fifteen people surrounded the vehicle**, asked which app she had booked on, took the driver's phone and told her to cancel and **delete** the rival apps. She did, then booked off-app. Drivers described about **fifteen days** of disruption, threats, theft, and **"honeypot"** bookings — people ordering a rival's car specifically to threaten whoever accepted.³

Eight hundred and forty-eight thousand people signed up. Eleven thousand rides a day found them.

Treat those as allegations, because the title that published them does. It could not verify the accounts; most of its sources would not go to the police; the platform did not respond and neither did either city's police chief; and two contradictory videos of uniformed officers circulated, one promoting the cooperative's insurance cover and one warning against coercion. Union tactics elsewhere are better documented: free tea for passengers who switch, and reports of **₹1,000 fines** levied on drivers who refuse to join the boycott.³

What the drivers actually get is not the rides. It is a bundle: **₹5 lakh** of personal-accident cover through a cooperative insurer, **₹5 lakh** of family medical cover through the national health scheme once they register on the unorganised-workers portal, group health through a payments company, and a tie-up wiring the platform into the state's document locker and services apps. Profit-sharing does not begin for three years, and is then **80 per cent distributed by kilometres driven**.⁴

So set the threshold honestly, because nobody involved has. At **11,106** rides a day the platform is not competing with anything; it is a pilot with a press office. The private apps in India move rides by the million. To matter in one city — to be the app a driver opens first in Ahmedabad — the cooperative needs something like **a hundred thousand rides a day** in that city alone, which is **nine times** its present national volume. That is the line. Below it, the **848,000** registrations are not traction. They are a mailing list.¹

Which is why the Jaipur number is the one to keep. **Thirty thousand** drivers signed on there; **372** rides a day came through. Do that to enough drivers and you have taught eight hundred thousand people that the cooperative is where you go to be on a list — hard to unteach, and the one real risk to a venture whose only genuine asset is the drivers' belief that it belongs to them.¹

The sceptic's case is not that the model is wrong. The insurance bundle is worth something even to a driver who never gets a booking. The case is that rider-side demand is the hard half of a two-sided market, and the cooperative has spent all its effort on the easy half. Drivers are cheap to recruit because they are angry; riders are expensive because they are comfortable. **₹83.70 crore** of cumulative driver earnings is a real number and also, across 848,000 registrations, about **₹1,000** each — eight months of a platform, one tank of fuel.²

The drivers were cheap to recruit because they are angry. The riders are expensive because they are comfortable.

And the coercion is the tell. A platform that is winning does not need a cordon at an airport rank. If the only way to move riders onto the cooperative is to make the alternative frightening, it has not solved the rider problem; it has outsourced it to men standing outside arrivals. Nobody has shown the

venture asked for any of it, and the title that reported it could not verify the accounts. But the pattern — drivers organised, riders pressured, the platform silent — is a market being forced rather than won.³

Which brings the essay to the thing nobody has named. What Bharat Taxi delivers at scale is not rides. It is **848,000 self-employed people on a single register**, through which accident cover, a national health entitlement, group insurance and a bank's credit reach workers who own a vehicle, carry a loan and have no employer to reach them. The app is not the product; the app is the **enrolment mechanism**, and the welfare bundle is the product. At 0.013 rides a driver a day this is a catastrophe as a business and a triumph as a census — India has built, more or less by accident, **a union with an interface**.⁴

And a register of that size is leverage whether or not a single car moves. The minister's own claim at the Gujarat launch is the admission: that the cooperative's arrival had **already forced the private operators to revise their fares**. A body holding the loyalty of 848,000 drivers is a standing threat to any aggregator's take rate without ever carrying a passenger. September in Ahmedabad was the first demonstration of that, and the ugliest: drivers worked out that the state's app could be used as a strike fund. Which raises the question the next twelve months will answer. What happens when a government becomes the counterparty in a labour dispute it also arbitrates? India has not nationalised ride-hailing. It has nationalised the drivers' bargaining position, and is now discovering it cannot control what they do with it.⁴

What to watch, then. First, a city-level daily-rides figure published by the venture rather than extracted from Parliament: its absence is the number. Second, whether profit-sharing survives to its third year. Third, the registrations-to-rides ratio, which on the minister's own figures stands at **0.013**. A cooperative that gets that to **one** has built something. One that leaves it where it is has built a membership roll with a logo, and the drivers who joined will remember who asked them to.

SOURCES 1 The Jaipur figures and the national daily-rides total are from the written parliamentary answer of 12 August 2026 cited on **p25**, read at source. The hundred-thousand-rides threshold is this magazine's own stated test, not a figure anyone has published, and the nine-times multiple is arithmetic against the 11,106 national total. 2 The commission range, the zero-commission and subscription terms and the five-lakh accident cover are as sourced on **p25** and **p26**. The ₹83.70 crore cumulative-earnings figure is from the same parliamentary answer; the division across 848,000 registrations is ours and is offered as a scale check, not as any individual driver's income. 3 The reported coercion is as described on **p26**, including that title's own inability to verify it; the inference drawn here is this magazine's. The venture has not been shown to have sought or sanctioned any of it. The three things to watch are editorial judgements; the 0.013 ratio is arithmetic from the parliamentary figures. 4 The register size and the welfare arrangements are as sourced on **p25** and **p26**. The minister's claim that the cooperative's arrival had already made private operators revise fares was made at the Gujarat launch of 28 June 2026 and is reported by two Indian titles read at source. The reading of the platform as an enrolment mechanism, and of the register as bargaining power, is this magazine's own argument; no source read advances it.

MERIDIAN

THE LAST WORD

Who drew the line

Relief for haze-hit households opens above a reading of two hundred, and two of the four bands the public is given have no published range at all (p5). A catastrophe bond pays nothing until the pressure in the eye falls below nine hundred and thirty-seven millibars (p3). A flicker standard half the lighting trade still cites was inactivated in March (p6). A schooling effect of four hundredths of a standard deviation was reported as a finding (p12).

None of those thresholds is wrong. Each had to be put somewhere, and somebody did the work. The trouble is what happens on either side of the line, where the arithmetic stops and the consequence does not: the household at a hundred and ninety-nine, the island at nine hundred and forty, the driver who registered and waited (p25).

So the question was never whether a threshold is accurate, but whether it was drawn to describe the world or to be defended in it. A line drawn to describe will tell you its own range. A line drawn to be defended will not.

TWELVE DESKS · ONE WORLD

MERIDIAN · No. 128 · 10 October 2026

A note on method. This edition's load-bearing figures were read where they are published: the National Environment Agency's own haze advisory, air-quality band definitions and hourly readings, with the 24-hour concentrations converted here against the World Health Organisation guideline; Indonesia's forestry ministry hotspot count; the Treasury's sanctions listing and the Nobel committee's own citation; a catastrophe-bond trigger table; TSMC's monthly revenue filing and guidance assumptions, with the exchange-rate crossover computed here; Micron's own quarterly release, against the claim this magazine printed in August; the Federal Reserve's daily yield series and the New York Fed's term-premium workbook, taken as data; the IEEE standards register, which records 1789-2015 as inactivated on 26 March 2026; three bulb specification sheets; Sigma's product page, with the absence of any independent bench result stated as the finding; four journal papers read at source; the European Southern Observatory's and ALMA's own visitor pages; and — computed here from a written parliamentary answer rather than quoted — every ratio on p25 to p27. **Corrections:** two, and one withdrawal. No. 127 said one scheduled American release landed all week; at least six did — services on the 5th, trade on the 6th, the committee minutes and consumer credit on the 7th, claims and wholesale inventories on the 8th, the Michigan preliminary on the 9th. The conclusion it drew about the long end survives; the premise does not (p14). No. 127 reported that one of Singapore's upper air-quality bands carries no published range; in fact both do not (p5). And No. 127's turnout and arrest figures for the French demonstrations could not be re-found at any source this week, so they are withdrawn rather than repeated. **Refused this desk today:** Reddit on every route and as a search domain, which is why six desks print a null result rather than a reader sentiment; the Indian government's own press bureau, on the day its minister's answer was the document that mattered; Reuters, the *Financial Times*, Bloomberg direct, CNBC, Forbes and TheStreet at article level; the Nobel Foundation; Billboard; Box Office Mojo; Singapore Airlines; Head-Fi, Hodiinkee and WatchPro; an American leather supplier that reset every connection; the regional haze monitoring centre; and the Singapore food titles Seth Lui and Burpple.