


THE CURRENT BRIEFING OF

MERIDIAN

Twelve desks. One world.

MONDAY 21 SEPTEMBER 2026



THREE GOVERNMENTS CONFIRMED THE SAME AGREEMENT ON FRIDAY, AND NOT ONE OF THEM HAS RELEASED IT (P3)

THE WORLD

Permanent control of Greenland, announced from a text nobody outside the room has read — and a Danish lecturer's verdict: the same deal as 1951

THE LONG READ

Vietnam wired the shop tills to the tax office, so five million shopkeepers peeled the QR codes off the counters

THE RABBIT HOLE

Knife making, all the way down: the step that decides whether a blade works is the one step nobody can see

THIS ISSUE · NO. 109

Contents

p3–4 The World — permanent control of Greenland, from a document nobody has published

p5 Singapore — a haze law twelve years old and never once used

p6 The Atelier — desk lamps: the CRI on the box averages eight pastel swatches and leaves out red

p7 Technology — Firefox's assistant runs on Mistral, and Mozilla won't say where the thinking happens

p8 The Home Lab — transparent output, ten reference images, and a licence that isn't Apache

p9 The Kit — a watch that samples every five seconds and shows you one reading in five

p10 The Good Life — a fourteen-storey waterfall on Orchard Road, and no rate published

p11 Screen & Sound — fourteen Emmys on the lowest critical score of the night

p12–13 The Family Desk — 3,486 children who had no phone at twelve; half of them got one

p14 The Macro Desk — every leg of a global portfolio lost money in the same week, gold included

p15–17 The Rabbit Hole — knife making: two steel mills closed, and the replacement is better

p18 Fitness — the split session lost on every measure, and won on the one that lasts

p19–20 The Diary — eighteen days to the street circuit; a hawker legend nobody queues for

p21–23 The Long Read — Vietnam: the QR codes came off the counters

FROM THE DESK

“Absence of evidence is not evidence of absence.”

Carl Sagan, *The Demon-Haunted World* (1995), popularising a formulation of the astronomer Martin Rees

An unpublished treaty text. A statute never invoked.

THE STRIP · AT FRIDAY'S CLOSE

US 10-year (par)	4.93%	▲ +6bp
Gold (spot /oz)	\$4,383.45	▲ +0.97% day
Bitcoin	\$80,799	± Mon 05:22
Brent crude	\$103.87	▼ -1.2% wk
USD/SGD	1.2757	▲ +0.68% wk
VWRA (All-World)	\$191.66	▼ -1.03% wk
AGGU (global agg)	\$5.758	▼ -0.10% wk
STI	5,656.11	▼ -0.70% wk
USD/JPY	156.86	▲ +2.15% wk

Bitcoin is this morning's print, not Friday's. Gold is spot and is **not** a record — it sits 21.8% below January's high. Not advice.¹

A CLARIFICATION TO OUR OWN STRIP

The ten-year, twice. No. 108 printed the US ten-year at 4.997 per cent under a note saying the yields were on the par curve. They were not: 4.997 was Friday's market close; the Treasury's own constant-maturity par yield that day was **4.93**. Both series are real, run six to ten basis points apart, and give the same weekly change. Above is the par figure, matching the note.

— p14

SOURCES 1 Levels at the close of Friday 18 September 2026, each read at source and set out with its origin on p14; the bitcoin print is 05:22 SGT on Monday 21 September, read as this edition closed. Treasury yields are the official constant-maturity par series. Gold is spot; the all-time high of US\$5,608.35 was set in January 2026. The all-world and global-aggregate tracker prices are single-sourced — the fund issuer and the exchange both refused this desk's requests.

And a straight book. After two themed issues, this one runs as ordinary news: no spine forced across the desks. The pattern on the back cover turned up on its own.

THE ARCTIC · AN AGREEMENT ANNOUNCED BEFORE IT EXISTS

Three governments confirmed the same Greenland deal. None has published it.

Washington calls it permanent control; Copenhagen calls it recognition of sovereignty. It is due to be signed in New York this week — and a Danish lecturer's verdict is that the Americans have held these rights since 1951.

Three governments confirmed the same agreement on Friday and not one of them released it. The Danish prime minister's office put out a release headed "Expected agreement" — expected, four days before a signing — saying the three would sign in New York during the General Assembly, entry into force "subject to the necessary national parliamentary procedures." It ends: "There are no further comments at this time."¹

Donald Trump was less procedural. The United States, he wrote, now has "permanent control over security, and all other needs, in Greenland," and no adversary can "EVER" hold a base there.²

Every specific term in circulation traces to United States officials briefing one American broadcaster: permanent basing and overflight with no expiry; new installations on consultation rather than consent; China, Russia and non-NATO states barred from bases *and* from investing in telecoms, cables, ports and critical minerals; the agreement surviving Greenlandic independence.² Al Jazeera, Xinhua and the AP wire note the same absence — no published text.³ They are printed as single-sourced, because that is what they are.

56,000

THE POPULATION OF GREENLAND. IN JANUARY 2025, **85 PER CENT** OPPOSED BECOMING AMERICAN³

COPENHAGEN & NUUK **STATSMINISTERIET, 18 SEP**

A crisis ended on Danish terms: an agreement "which at the same time recognizes the sovereignty and territorial integrity of the Kingdom and the Greenlandic people's right to self-determination."¹

WASHINGTON

VIA FOX NEWS, 18 SEP

A conquest completed and, pointedly, a cheap one. "Greenland will forever be part of the strategic defense area of North America and exclude any adversary from it." A senior official on the price: "we're not spending any money for the right."²

"Congratulations Trump. You now have what you've had since the 1950s, when the base agreement was signed."

THE SCEPTIC · A RERUN OF 1951, DRESSED AS A RUPTURE

Peter Viggo Jakobsen of the Royal Danish Defence College told DR that Washington has been handed the rights it already held under the 1951 defence treaty, and that the annexation campaign was "more theatre than genuine military concern" — but that the administration had to sell the outcome as a victory, or Trump "would have accomplished nothing."⁴ The footprint supports him: seventeen installations and more than ten thousand troops in Greenland at the Cold War's height, against one base now — Pituffik, a hundred-odd personnel.⁵ Permanent control of what, exactly, is what the unpublished text would answer.

THE CHATTER · WHAT THE ANALYSTS ACTUALLY SAID

Richard Weitz, NATO Defense College, to Al Jazeera: Trump "backed down and settled for less" than he sought, but "got pretty much everything he wanted, except formal annexation."

Mikkel Olesen of the Danish Institute for International Studies on the headline concession: barring Russian and Chinese bases is "unlikely to be a major concession because the Danes and Greenlanders never wished" for them.³

SOURCES 1 Statsministeriet (Danish Prime Minister's Office) press release, "Expected agreement on strengthened security in the Arctic and the North Atlantic Area," 18 September 2026, read in full at source; both quotations are from that release. 2 Fox News, 18 September 2026, sourced to United States officials including Secretary Rubio — the only published account of the agreement's specific terms, and printed here as single-sourced. Trump's own wording is quoted identically by NBC News, Al Jazeera, Xinhua, Euronews and Fox from his Truth Social post; the post itself was not readable by this desk. 3 Al Jazeera explainer, 19 September 2026, read at source: the unreleased text, the January 2025 poll, and the Weitz, Bayoumi and Olesen comments. 4 DR (Danmarks Radio), read 20 September; the page carries a relative timestamp ("yesterday, 14:55") rather than an absolute date, so the date is inferred and the quotation is verbatim. 5 Euronews, 19 September 2026, for the Cold War footprint and the present Pituffik establishment. Reuters, the BBC, AP, the Guardian, CNN and the Washington Post all refused this desk's requests; wire copy was read through SCMP, Arab News and Euronews instead.

GERMANY · A GOVERNING PARTY POLLS 4.9 AND LEAVES WITH NOTHING

The AfD wins a state outright. In Berlin, Die Linke comes first.

Two Sunday elections, and the Christian Democrats fell below the five-per-cent threshold in Mecklenburg-Vorpommern — out of the parliament altogether, on a night the chancellor called a disaster.

The five-per-cent hurdle is a blunt instrument, and on Sunday it did what it was built to do, to the wrong party. In Mecklenburg-Vorpommern the AfD took **38.3 per cent and 32 of 71 seats**; the SPD 35.5 and 29. Friedrich Merz's CDU took **4.9 per cent and zero seats** — its worst state result anywhere since the war, and an exit from a Landtag rather than a defeat inside one.¹

Berlin voted the same day and went the other way. Die Linke came first on **25.8 per cent**, up thirteen points; the CDU fell nine to 18.8; the AfD took 15.9; the SPD, which has run Berlin for most of two decades, finished fifth.¹

Merz, who cancelled his General Assembly trip under domestic pressure, promised more of the same: "The reforms have to come." The AfD's Tino Chrupalla read that as a gift — the chancellor "announced that he'd carry on as before."¹

4.9%

THE CDU'S MECKLENBURG-VORPOMMERN VOTE — A TENTH OF A POINT BELOW THE HURDLE, AND THEREFORE **ZERO SEATS**¹

THE SCEPTIC · WHAT A PRELIMINARY COUNT IS WORTH

These are preliminary numbers, compiled off the returning officers' feeds rather than the officers' own pages, which this desk could not reach; the exit polls corroborate the order, not the last digit. The 4.9 is the figure to watch — a tenth either way is nought seats or five.¹

RUSSIA VOTES; UKRAINE SENDS SIXTEEN HUNDRED DRONES

The Duma elections ran Friday to Sunday. On the final voting day Ukraine mounted what Moscow's mayor Sergei Sobyanin called the largest drone attack on the capital of the war: **more than 1,600**

shot down across Russia in a day, 450 aimed at Moscow. Several reached the Moscow refinery. Two were killed in the Moscow region, two in Russian-held Kherson including an election official, four in the Kyiv region.²

Moscow called the strikes an attempt to "disrupt the democratic will of Russian citizens." The number worth keeping is domestic: the exit poll put United Russia at **49.4 per cent**, the preliminary count at **57.5**. Eight points is a wide gap to open between a poll and a tally in one evening.²

BEIJING

XINHUA ENGLISH, 19 SEP

On Greenland (p3) the framing is carried by one verb: the headline reads "Trump *claims* deal." The foreign ministry did not raise it at its 17 or 18 September briefings — Beijing's official line is, so far, no line at all.³

THE CHATTER · RIYADH, AND IRAN'S PRICE FOR THE STRAIT

The first Houthi missile aimed at the Saudi capital since July was intercepted on Friday; on Sunday black smoke was seen near fuel storage at King Khalid International Airport, no casualties reported. Tehran has put a number on itself: its own Statistical Center reports **GDP down 10.1 per cent year on year** for the quarter to late June. Speaker Mohammad Bagher Ghalibaf says Iran's conditions for reopening Hormuz have been "clearly conveyed through mediators" — an end to the war on all fronts, frozen funds released, the naval blockade lifted.⁴

IN BRIEF

17 Sep · Riyadh's shopping list. Washington notified Congress of a proposed sale of **48 F-35s to Saudi Arabia, worth US\$24.3 billion.** Approval still to come.⁵

This week. The Greenland accord is signed in New York. Xi Jinping goes to Washington on Thursday for a summit on trade, tariffs, Taiwan and technology restrictions — the week's longest read-through to this book's other desks (p8, p14).³

SOURCES 1 Mecklenburg-Vorpommern and Berlin results, turnout and seat allocations from preliminary counts compiled on 20–21 September 2026 — the state returning officers' own pages were not reachable by this desk, so the decimals are tertiary and provisional and are labelled as such above; ARD exit polls, Al Jazeera and Euronews (both 20 September) corroborate the order and the margins. Merz and Chrupalla quotations from Al Jazeera, Euronews and AP via SCMP, 20–21 September. 2 Drone totals, casualty counts and the Sobyanin quotation from AFP and Reuters copy carried by SCMP, 20 and 21 September 2026, and Euronews, 20 September; the exit-poll figure is AFP's, the preliminary tally a partial central-commission count at about 37 per cent of precincts reporting. 3 Xinhua English, published 19 September 2026; China's position checked against the foreign ministry's own 17 and 18 September briefing transcripts, the latest posted. The week's diary from SCMP, 21 September. 4 Al Jazeera, 19 September 2026, and SCMP, 20 September, on the Riyadh interceptions; Arab News (AFP), 20 September, for the Iranian GDP print, the Ghalibaf statement and the Rezaei conditions. 5 Al Jazeera, 17 September 2026, for the F-35 notification; Arab News (AFP), 20 September, for Ethiopia.

THE HAZE · A STATUTE THAT HAS NEVER ONCE BEEN USED

Singapore has a law for this. In twelve years it has opened no investigation.

In February the environment minister explained why: the haze had not reached levels that would warrant one. The peak then was a PSI of 69; on 15 September, Central read 154.

At three this morning the twenty-four-hour PSI read 66 north and south, 70 east, 86 west and centre; the one-hour PM2.5 figures ran from 13 micrograms in the south to 34 in the west.¹ That is the moderate band, and better than the week has been.

Six days ago it was not. On the morning of 15 September all five regions were unhealthy at once for the first time since 2019: Central **154**, East and West 128, South 111, North 104. The environment agency's advice contained the sentence that tells you how a city calibrates alarm: N95 masks "are not required for short exposure, like commuting from home to school or work."²

Singapore passed the Transboundary Haze Pollution Act in 2014: an offence to cause or contribute to haze pollution here, wherever the fire is lit. On 12 February 2026, answering Christopher de Souza, Joan Pereira and Fadli Fawzi, Grace Fu told Parliament that no investigation had ever been initiated under it, because "the haze has not reached such levels which would warrant investigation." The highest reading she had was a PSI of 69. She ruled out cloud seeding as ineffective here, and noted that the Haze Task Force spans twenty-eight agencies.³

69 → 154

THE PEAK READING WHEN THE MINISTER LAST EXPLAINED THE THRESHOLD, AGAINST THE PEAK SIX DAYS AGO. THE MINISTRY'S SEPTEMBER LISTING CARRIES A COOLING SPEECH AND A BATCH OF WRITTEN REPLIES — NOTHING ON THE ACT³

That answer is seven months old and has not been revisited. The question is neither rhetorical nor hostile. At 154, does it now warrant investigation? Nobody has been asked, and nobody has said.

THE CHATTER · THE FORECAST THAT CALLED THIS IN JUNE

The Singapore Institute of International Affairs rated the risk "**Red**" on **24 June** — its highest band, only the second since the Outlook began in 2019 — and named **August–September 2026** as the window. It has landed almost to the month.

Chairman **Simon Tay**: "we must not be fatalistic. There is much that can be done." Senior assistant director **Aaron Choo**: the worst episodes come "in years with both a strong El Niño and 'positive' IOD, like in 1997 and 2015." The Outlook puts 2015 at **US\$16.1 billion** in Indonesia alone.⁴

THE SCEPTIC · THE FIRE MAY BE NOBODY'S NEGLIGENCE

The same Outlook argues against reaching for the statute first. **Khor Yu-Leng**, its associate director, points at demand rather than carelessness: "The fuel crisis is also driving demand for biofuels, with Southeast Asian countries increasing the percentage of vegetable oil in diesel."⁴ A blending mandate is a price signal on palm oil and therefore on land, which makes the burning a regional policy output rather than a local crime.

What would settle it: a ministerial answer this sitting, or an NEA notice served.

ALSO THIS WEEK

467 addresses, one wrong field. The National Cancer Centre invited patients to a hereditary breast-and-ovarian-cancer event using cc rather than bcc, exposing every address to all of them. Chief operating officer Chong Pang Boon: "we are deeply sorry."⁵

Singapore declines to follow America. The Health Ministry kept the colorectal-screening start age at **50** on 17 September, against America's 45, citing local registry data.⁵

Sixteen minutes, four tenths. Gan Ching Hwee took silver in the 1,500m freestyle at the Asian Games on Saturday in **16:00.45** — the first Singaporean woman to medal in the pool at the Games since 2018.⁵

SOURCES 1 National Environment Agency 24-hour PSI and 1-hour PM2.5, read directly from the data.gov.sg real-time endpoints at 03:00 and 03:30 SGT on 21 September 2026 — the readings are this edition's own, taken as it went to press. 2 Peak regional readings of 15 September 2026 and the NEA mask advisory from IBTimes Singapore, 15 September, read at source. The Straits Times and CNA both refused this desk's requests, so no reading is sourced to them. 3 Ministry of Sustainability and the Environment, oral reply to a parliamentary question on transboundary haze, 12 February 2026, read at source — the quotation, the PSI 69 reference, the cloud-seeding answer and the twenty-eight-agency figure are all from it. The absence of any September statement is this desk's own reading of the ministry's latest-news listing on 20 September. 4 SIIA Haze Outlook 2026 media release, 24 June 2026, read at source, for the Red rating, the August–September window, the Tay, Choo and Khor quotations and the 1997–98 and 2015 loss figures. 5 AsiaOne, 20 September 2026, for the National Cancer Centre disclosure — note the outlet's own account places the send on Friday evening and the statement on Saturday 19 September. Ministry of Health written reply to Dr Desmond Wai, 17 September 2026, read at source, for the screening-age decision, which rests on the Screening Test Review Committee's February 2026 report. Race result and historical note from AsiaOne carrying Straits Times copy, 20 September.

BUY IT ONCE • THE LAMP ON YOUR DESK

A lamp with an unreplaceable LED has a service life, not a lifetime.

The specification that decides whether you still own the thing in twenty years is the socket — and it is never on the front of the box.

Begin with the number makers print and the one they don't. **Ra — the CRI on the box — is the mean of only the first eight test colours,** all pastels; R9, deep red, is calculated and left out. A blue-pump white LED has its trough exactly where R9 lives, so a lamp can advertise **Ra 90 with an R9 near zero** and be telling the truth.¹ Ask for R9 in writing.

Then flicker. PWM dimming chops the current on and off — a square wave near **100 per cent depth**; constant-current dimming lowers it continuously and stays shallow. On 50 Hz mains the ripple is 100 Hz, and IEEE 1789 puts the no-effect level at **8 per cent modulation, the low-risk level at 3.3**. A halogen filament sits near 10.² And judge the output where the work is, not at the bulb: the Luxo Ninety emits **300 lumens**, less than a forty-watt bulb, and puts **1,100 lux** on the page at four hundred millimetres — double the office standard — because it has an optic.¹

The arm is older engineering and better. **George Carwardine**, a Bath engineer who specialised in **vehicle suspension**, filed the spring-and-lever patent in 1932 and licensed it to Herbert Terry & Sons, a **spring factory**.³ A near-zero-free-length spring across a parallelogram gives a restoring moment varying exactly as gravity's; they cancel at every angle. Jieldé does the opposite and is equally sound: no springs, only friction joints. A fatigued spring is not adjustable; a loose joint is.

THE CHATTER • AND THE SHOP TEST

On Hacker News's "LED Strain" thread, **posterboy** corrected the room in one line: "No, the diodes don't, the drivers do."² That is the tier line. **vitaflo** gives the shop test: a fast phone-camera shutter at the lit page — "if there are black lines across the screen there is a PWM issue."²

SOURCES 1 CRI construction (Ra as the mean of R1–R8; R9 excluded from the headline figure) and the TM-30-20 fidelity/gamut pair from the standards literature, read at source; EN 12464-1:2021 task values (500 lux maintained, UGR ≤ 19, Ra ≥ 80) from three independent renderings of the table. Luxo Ninety output, lux at 400 mm and CRI from the dealer specification sheet, read at source. The IEC 62471 / IEC TR 62778 versus CIE S 026 distinction from the standards' own scope statements and from BenQ's own claim page. 2 IEEE 1789–2015 thresholds as published, worked here for the 100 Hz ripple of a 50 Hz supply; the PWM-versus-halogen modulation comparison from DIAL's treatment of the standard. Hacker News quotations from the "LED Strain" thread (10 December 2018) and the "PWM flicker" thread (18 June 2025), read at source, handles as posted. 3 Carwardine's patent, the February 1934 licence to Herbert Terry & Sons and the model chronology from the Design Museum and the encyclopaedic record, which differ by three days on the filing date; manufacturing location, the Giant exception, and the Luxo and Wästberg ownership facts from the same record and the makers' own statements where reachable. anglepoise.com refused this desk throughout, so every Anglepoise price here is an authorised dealer's. 4 Prices read at makers' and authorised dealers' own stores in September 2026: A+R and 2Modern for Anglepoise in USD; jielde.com for the Loft range in EUR, carrying the five-week lead time and the fifty-kilometre sourcing claim; ambientdirect and smow for Tolomeo; artemide.net for the LED and Tunable White specifications. 5 L70, LM-80 and TM-21 including the six-times projection limit from the lumen-maintenance literature; driver failure as the leading cause of premature LED system failure from the lighting-engineering sources cited there. Singapore prices from EOMM's Anglepoise listing and a Singapore retail report of the ScreenBar Halo 2, both read at source, with BenQ's lux, coverage and Rf figures from its own specification pages.

THE WHITE-LABEL TRUTHS

Anglepoise is a British design made in China for every model but the **1227 Giant**, hand-built in Portsmouth; the **Paul Smith Type 75 at US\$430** is the **US\$320** lamp in different paint. **Luxo** is a Glamox subsidiary, and Jacobsen's 1937 L-1 was itself "a modification of the earlier Anglepoise lamp." Artemide's fixed Tolomeo LED is 10 W at **CRI >90**; its Tunable White upgrade is 12 W at **CRI >80** for €250 more.³

THE SCEPTIC • TWO NUMBERS THAT MEAN LESS THAN THEY LOOK

"**No blue light hazard**", which the eye-care lamps advertise, is a class every consumer desk LED passes: the lamp is certifying that it is not a welding arc.¹ "**Fifty thousand hours**" is not a measurement either. It is an L70 projection — output down to seventy per cent, a threshold chosen because people fail to notice the loss until well past it. What kills the lamp is the driver.⁵

Entry of quality: Anglepoise **Type 75, S\$339.30** at EOMM, or a BenQ **ScreenBar Halo 2 at S\$289**, which publishes lux at distance and an Rf. **Sweet spot:** 1227 at US\$320–360, Tolomeo Tavolo at €323–342, Luxo Ninety at US\$395. **Grail:** a Jieldé Loft at **€936–1,428**, handmade and numbered at Saint-Priest.⁴

NEXT ON THE BENCH — PICK THE NEXT CATEGORY

Wooden cutting boards & butcher blocks — end grain against edge grain.

Barware & the home bar tools — Yukiwa, Koriko, the Japanese mixing glass.

Hard-shell luggage & carry-ons — polycarbonate against aluminium.

SOVEREIGNTY · THE BROWSER BECOMES THE FRONT LINE

Firefox's assistant now runs on Mistral. Mozilla will not say where the thinking happens.

The announcement is four paragraphs long and contains a zero-data-retention promise, two launch territories and no model name. For a reader who owns the hardware to run this himself, the omission is the whole product.

Mozilla published it on 16 September, without a benchmark or a version number: Firefox's **Smart Window**, its beta browsing assistant, is "now powered by Mistral models." Conversations "aren't saved on Mozilla's servers by default," and partners "like Mistral agree to zero data retention." It is live in **France and North America**.¹

Read what is absent. The post does not name which Mistral models, or say whether inference runs on the machine or in a data centre, or whose. No Firefox version, no pricing, no word on Singapore.¹ Zero retention is a promise about what is kept, not about where the text goes first — and for a European-sovereignty pitch arriving in the last independent browser, the location of the compute is not a detail. It is the claim.

2 of 4

TERRITORIES LIVE AT LAUNCH. BRITAIN AND GERMANY, BOTH LARGE FIREFOX MARKETS, ARE "EXPECTED LATER THIS YEAR," WITH NO DATE¹

The practical answer is the one the announcement makes obvious by not addressing it: a browser assistant is the clearest case yet for pointing an application at an endpoint you control. The home lab runs the models already (p8).

THE SCEPTIC · THE CRACKED INTERN

Jay Krueer's essay of 15 September is the sharpest bear case in the window. His frame for a frontier model: "a cracked intern: quick and effective in the hands of an adult but not given run of the place."²

The analogy is from silicon: a CPU project typically employs **three times as many specification and validation engineers as design engineers**, reaching five to one. Rigorous specification is the expensive part of building anything correct, and a model that writes the design does not remove it. His conclusion lands on the next page.²

AND THE WEEK'S REAL SILICON STORY IS A MEETING

Thursday, Washington. Xi Jinping arrives for a summit on trade, tariffs, Taiwan and technology restrictions — the week's one event with a read-through to every other desk here, from the dollar (p14) to the export-control regime that prices every accelerator in a home lab.³

Still on the shelf. Nvidia's two Rust tracks for GPU kernels — **cuda-oxide**, a custom rustc backend emitting PTX, and **cutile-rs**, tile-based on stable Rust — hit the front page on 16 September on a post dated the eighth. Nvidia calls cuda-oxide "early alpha" and warns "coverage is incomplete and APIs will move," aiming at 2027, with no performance numbers. cutile-rs is already inside mistral.rs, which is how the two halves of this page meet.⁴

THE CHATTER · WHAT THIS DESK COULD NOT READ, AND SAYS SO

Three stories in the window resisted verification and are not run as fact. **Fujitsu's MONAKA** processor — a made-in-Japan sovereignty story with a press release dated 14 September — returned HTTP 429 to every attempt, so no specification is printed. **Z.ai's** account of its own inference infrastructure served an empty script shell. And a pair of widely-circulated figures for a frontier lab's raise came back 403 from the lab's own page.⁵

SOURCES 1 Mistral AI announcement, "Mistral x Mozilla," dated 16 September 2026 and read in full at source; every quotation is from that post, and every omission noted above is an omission in it. 2 Jay Krueer, "Why I'm still bearish on LLMs after Navier-Stokes," read at source; the page carries no visible date beyond its URL slug, so the date given is its Hacker News submission of 15 September 2026, with the score and comment count as read on 20 September. 3 Summit date and agenda from SCMP, 21 September 2026. 4 Nvidia developer blog, "Introducing CUDA-Rust: two tracks for writing GPU kernels," dated 8 September 2026, read at source, including the toolchain requirements, the "early alpha" characterisation and the absence of performance figures; the Hacker News submission of 16 September and its score read on 20 September. 5 Access failures recorded by this desk on 20 September 2026: global.fujitsu returned HTTP 429 to repeated requests; z.ai returned a two-kilobyte script shell; techmeme.com returned 403; the frontier-lab funding figures surfaced in search results but their source page returned 403 and the numbers are therefore not printed.

THE HOME LAB · OPEN WEIGHTS, CLOSED LICENCE

Qwen-Image 2.1 does transparent output and ten reference images. It is not Apache, and it does not tell you what it needs.

The model card went up on 14 September and reached the front page on the twentieth. It publishes an architecture, an aspect-ratio ceiling and a licence — and not one benchmark score or a single figure for video memory.

The interesting part of the feature list is the alpha channel. **Qwen-Image 2.1** carries a **seven-billion-parameter** visual generation component across **32 single-stream DiT layers**, does text-to-image and editing, accepts **up to ten reference images**, runs to **2752×1536**, and outputs **transparent RGBA** — the difference between a picture and an asset, and the reason this one matters to anyone with a node graph open.¹

It wants transformers 5.17 or newer and diffusers from source, and offers CPU offload for the memory-constrained. What it does not offer is a number: the card publishes **no VRAM requirement and no benchmark**, and this desk found no independent evaluation.¹

Not Apache

THE LICENCE IS THE **QWEN RESEARCH LICENCE AGREEMENT**, NOT THE PERMISSIVE ONE THE LAST FEW OPEN-WEIGHT RELEASES TRAINED READERS TO EXPECT. READ IT BEFORE IT GOES ANYWHERE NEAR CLIENT WORK¹

That licence line is the buy-once question in software form. An open-weight model under a research licence is a loan, not a purchase: run it, study it, tear it apart locally — but do not rely on it commercially without reading the terms. Treating “open weights” and “Apache” as synonyms is a habit the last two years built, and this is where it costs something.

THE CHATTER · THE REPRODUCTIONS THAT DIDN'T

The ternary 27-billion-parameter release this desk covered on Sunday has now met its first hands-on reports, and they are not the table.² On the launch thread, **redox99** found the WebGPU build “immediately started looping” and called the near-lossless framing false: “reasoning that it looped on was clearly wrong.” **raylad** ran a memorisation check — recite Jabberwocky — and reports the full-precision version nearly passing while the ternary and fp8 builds “fail badly.” **Aurornis**: it falls apart “spectacularly” on longer tasks. The other direction, same thread, same day: **jedbrooke** got **7–8 tokens a second** generating on a sixteen-gigabyte M2 Mac mini; **jakswa** got **89 with speculative decoding** on an RX 7900 XTX. Both things are true. A retention percentage is an average over benchmarks, and benchmarks do not ask a model to remember a poem.²

THE SCEPTIC · WHAT THE COMPARISON LEAVES OUT

The sharpest published critique of that release asks the question its own table cannot answer: every comparison is against other quantisations of the same 27-billion model. At roughly six gigabytes on disk, is a crushed 27B actually better than a native eight- or nine-billion model at four bits in the same memory? Nobody has published that test.² Two of three Apple throughput figures were measured on an older build and carry a pending-re-measurement note, and the energy claim cites “an 8B model” without naming it. The verdict to carry: do not build on it yet, and note that the common desktop runtimes will not load the packings at all until the ternary kernels land upstream.

STILL ON THE SHELF

No Matter, no Thread, no Home Assistant. Nothing shipped in the window. The project's blog index still tops out at the **2 September** release — the Matter network map, the connectivity settings page, the activity view. Three weeks stale is not news, and this desk is not going to dress it as such.³

SOURCES 1 Qwen-Image 2.1 model card, read at source on 20 September 2026, for the parameter count, layer count, reference-image ceiling, maximum resolution, transparent-output support, dependency versions and licence; repository creation date, last-modified stamp and like count read the same day from the host's own API. The absence of a VRAM figure and of any benchmark score is an absence on the card itself. The maker's blog post is JavaScript-rendered and returned no readable content. 2 Hands-on reports, handles and figures from the launch discussion thread of 17 September 2026, read at source through the site's API; the published critique and its evaluation conditions from DataNorth AI, 18 September 2026, read at source. That release was covered in No. 108 and is carried here only as a delta — the first independent hands-on reports and the failed reproductions — not as a fresh story. A third outlet's coverage of the same release returned HTTP 202 with no body and was not used. 3 Home Assistant's own blog index, read 20 September 2026; newest post dated 2 September 2026.

WEARABLES · IT SHIPS, AND THE PRICE LANDS

The Ultra 4 samples your heart every five seconds. It shows you one reading in five.

Singapore price S\$1,199, on shelves since Friday. The reviewer who measures these things for a living calls the sensor the best Apple has made — then documents twenty-second holes in the data during hard efforts.

It went on sale on **18 September** at **US\$799**, and at **S\$1,199** here — read off Apple's own Singapore store, which is the peg this desk was waiting for.¹ The case is recycled titanium, natural or black. The battery is rated **50 hours** of normal use, 84 in low power, 25 hours of GPS workout and 45 in the extended mode. Satellite now carries messages as well as emergency calls.

The spec that changed is the sensing: the new health system samples heart rate **every five seconds all day** and measures variability **up to twenty-four times more often**, feeding a **Readiness score** out of ten.¹

DC Rainmaker, whose in-depth review landed on the seventeenth, calls the optical sensor “the best Apple has ever made” and the watch “a solid upgrade for earlier Apple Watch Ultra series owners.” His own measurements broadly hold: roughly 18–20 hours of standard GPS, 25 extended, 45 maximum, and the claimed two days of daily life confirmed.²

THE SCEPTIC · THE ACCURACY STUDY GRADES ITS OWN HOMEWORK

The same review takes Apple's heart-rate accuracy methodology apart. Apple records only points that meet its internal “quality confidence levels” — which produces **gaps of twenty seconds and more during hard efforts**, precisely when the reading matters, and precisely the interval a comparison against a chest strap would otherwise expose. In workouts the watch displays a figure **every five seconds rather than every second**, so the sampling improvement does not reach the screen.²

The rest of the honest list: the touchscreen “falls apart in colder conditions with gloves”; navigation and mapping remain the biggest gap against the dedicated outdoor brands; GPS accuracy is “nearly identical” to the previous generation; and Readiness shipped with speed and display bugs. For a buyer replacing an Ultra 2, the sensor is the reason. There is no other.

SOURCES ¹ Apple newsroom announcement of the Apple Watch Ultra 4, read at source, for the availability date, case materials, battery ratings, satellite messaging and the health-sensing specification including the five-second sampling, the heart-rate-variability multiple and the Readiness scale. The Singapore price of S\$1,199 was read from Apple's own Singapore store on 20 September 2026; the US price is Apple's own. ² DC Rainmaker in-depth review, dated 17 September 2026, read at source: the sensor verdict, the measured battery figures, the methodology critique, the twenty-second data gaps, the five-second workout display, the cold-weather touchscreen, the navigation gap, the GPS comparison and the Readiness bugs are all his. ³ Kit prices, accessory prices and the reservation and purchase-email timeline from Road to VR, 14 September 2026; the silicon, memory, battery, panel and radio specification from BGR, 14 September 2026, with the panel resolution and the “120 Hz (144 Hz experimental)” wording independently confirmed in Valve's own store markup, read 20 September. Valve's store exposes no prices and no Singapore availability, so none is printed. The “eye-watering” characterisation is BGR's. ⁴ Apple newsroom announcement of AirPods 5, read at source, for the availability date and the noise-reduction claim; Singapore prices read from Apple's Singapore store, 20 September. Review verdicts from MacRumors' hands-on of 15 September 2026 and Engadget, 9 September 2026, both read at source.

ALSO SHIPPING

Valve's **Steam Frame** closed reservations on the seventeenth and sent its first purchase emails on the **eighteenth: US\$1,059** for the 256-gigabyte kit, **US\$1,299** for the terabyte, each including controllers, the Wi-Fi 6E adapter for PC streaming and a copy of *Half-Life: Alyx*.³ Inside: a Snapdragon 8 Gen 3, sixteen gigabytes of LPDDR5X, a 21.6 watt-hour battery and dual **2160×2160** panels.

120 Hz

THE HONEST REFRESH NUMBER. VALVE'S OWN STORE COPY LISTS **144 HZ AS “EXPERIMENTAL”** — WHICH IS THE FIGURE THE WRITE-UPS LEAD WITH. NO SINGAPORE PRICE OR DATE HAS BEEN PUBLISHED³

THE CHATTER · THE PRICE, AND THE COMPANY IT KEEPS

The trade reaction is not about the optics. BGR calls the price “eye-watering” and “a significant leap” over the Meta alternatives, landing in a season of broad hardware price rises.³ The counter-argument is the bundle: a standalone headset that is also a wireless PC display, shipped by the company that owns the library, is a different proposition from a headset that needs a store to grow around it. Whether that is worth roughly double is a judgement, and nobody outside Valve has spent a week with one yet.

And the cheap one that is mostly fine. AirPods 5 also landed on the eighteenth — **S\$199**, or **S\$219** with the wireless charging case — claiming up to fifty per cent more noise removed than the ANC AirPods 4 via a new multiport acoustic architecture, at five hours a charge.⁴ MacRumors' hands-on calls them “an excellent value” and then says the sound is “muddier than on the AirPods Pro 3, especially in the low mids,” the fit is unchanged from the AirPods 4, and existing Pro or ANC owners should not upgrade. Engadget's word for the generation is “modest.”

ESCAPES · THE CAIRNHILL CORNER, FILLED AT LAST

Hilton is putting NoMad on Orchard Road, with a fourteen-storey waterfall and 1,860 artworks.

A hundred and seventy-three rooms over twelve floors on the old Faber House site, WOHA outside, Kevin Klein inside — and roughly nine in ten of the pictures by artists based here. Opening is “late 2026”; no rate has been published.

The announcement went out on **17 September** and is short on the thing a reader wants, which is a number with a dollar sign. What it does give is a building. NoMad — the Hilton luxury-lifestyle brand, here with **UOL Group** — takes the Cairnhill intersection of Orchard Road, on the site of the former Faber House: **173 rooms and suites over twelve levels**, of which 25 are suites, on a 10,225 square-metre plot.¹

The room ladder is unusually flat: 71 Atelier rooms at 32 square metres, 42 Parlour Suites at 42, 43 NoMad Suites at 43. Architecture is **WOHA**, interiors **Kevin Klein Design**, and the building carries a cascading **fourteen-storey waterfall** behind a vertical green façade.¹

~90%

SHARE OF THE HOTEL'S **1,860 ARTWORKS** MADE BY ARTISTS BASED IN SINGAPORE. THAT IS THE DETAIL THAT SEPARATES THIS FROM THE BRAND'S NEW YORK AND LONDON HOUSES — AND THE ONLY SOURCED CLAIM ON THE PAGE A GUEST CAN CHECK BY WALKING THE CORRIDORS¹

Two things to hold honestly. The opening date is “late 2026” and **no rate has been published**, so anyone planning around it is planning around a press release. And this desk found no comparable Japanese opening inside the window: the 2026 wave landed in spring, and the autumn Sapporo property does not open bookings until December.¹

OBJECTS · THE FOUR-THOUSAND-POUND QUESTION

Two headphone amplifiers were listed as new in the same weekly round-up on **18 September**, and the gap between them is the whole buy-once argument. iFi's **iCAN PHANTOM 1.5** at **£3,749** is a fully differential design with three switchable input stages — solid-

state J-FET, tube and tube-plus, on a matched pair of General Electric 5670 valves — delivering more than 15,000 milliwatts into 16 ohms balanced and a signal-to-noise ratio above 145 decibels, and, crucially, driving electrostatics at up to 640 volts of bias through interchangeable cards. One box for the Stax and the Susvara.² FiiO's **FT15**, announced the same day, is a diamond-coated planar at **£499**.²

THE SCEPTIC · NOBODY HAS REVIEWED THE THING YOU WOULD BE BUYING

The **1.5** is a revision, not a new platform, and iFi has published no dated statement saying what changed — its own news feed stops in August. Every named specialist verdict in existence is for the *original* iCAN Phantom, at the same price: Hi-Fi's Simon Lucas called that machine “an absolute masterclass of headphone amplification” and “always an absorbing, revealing, vaultingly dynamic and forensically insightful listen” — in a review dated **October 2024**, with a reservation about the control app failing to initialise.²

So the honest position: what the 1.5 changes is unknown to this desk, and a four-thousand-pound purchase against a two-year-old review of a different revision is a leap. Wait for the first dated hands-on. Darko refused this desk's requests and Stereophile has not covered it.

THE CHATTER · THE WEEK'S OTHER LISTINGS, AT A GLANCE

From the same round-up, prices as published: **KLH Model Four** floorstanders at US\$1,999 the pair, described as the small footprint of the Model Three with the architecture of the Model Five; **Schiit Bifrost 3** at US\$699; **Clearaudio CompactPhono** at US\$639 on the seventeenth; and Devialet's **Phantom Ultimate in Pulse Blue** at US\$2,200 — which is a finish, not a product, and is listed here as one.²

SOURCES 1 NoMad Hilton Singapore announcement carried 17 September 2026, read at source: partner, site, room and suite counts and sizes, floor count, plot area, architect and interior designer, the waterfall and green façade, the artwork count and its local share, and the food-and-drink line-up. The opening is stated as “late 2026” with no date, and no room rate is published anywhere in the release — both absences are the release's. The Japanese comparison rests on this desk's own search of the window: the 2026 Kyoto and Nara openings fall in March and June, and the Sapporo property's bookings are trade-reported to open on 1 December 2026, a pipeline date its own source warns is subject to slippage. 2 iFi iCAN PHANTOM 1.5 price and specification from iFi's own product page, read at source, and its listing in Gear Patrol's hi-fi round-up of 18 September 2026, which is also the source for the KLH, Schiit, Clearaudio and Devialet prices and for the FiiO FT15 announcement of the same date. The Hi-Fi+ verdict is Simon Lucas's review of the original iCAN Phantom, dated 30 October 2024, read at source and cited here as a review of the earlier revision, which is what it is. iFi has published no dated announcement of what the 1.5 changes; its own news feed ends in August 2026. Darko.Audio returned HTTP 403 to this desk.

TELEVISION · THE RECORD, AND THE REVIEWS THAT DON'T MATCH IT

A freshman comedy won fourteen Emmys. Its Metascore is 78, the lowest of the night's winners.

Matthew Rhys became the first person to take two lead acting Emmys in one year. The show he did it for has a critical temperature that runs a long way below the industry's enthusiasm for it — and one review at 40.

The ceremony ran on **Monday 14 September** at the Peacock Theater, hosted by Mariska Hargitay — the first woman in the chair since 2011. *Widow's Bay*, a first-season horror comedy created by Katie Dippold, took **six awards on the night and fourteen in total**, an all-time record for a comedy series.¹

The individual record is stranger. **Matthew Rhys** won Lead Actor in a Comedy and Lead Actor in a Limited or Anthology Series — the first person ever to take two lead acting Emmys in the same year. Stephen Root and Kate O'Flynn took the supporting awards, Hiro Murai directing, Dippold writing. In drama, *The Pitt* repeated, as did Noah Wyle; **Rhea Seehorn** took a first Emmy for *Pluribus*, **Jean Smart** an eighth for *Hacks*.¹

78 / 8.1

METAScore FROM 26 CRITICS AGAINST THE USER SCORE FOR *WIDOW'S BAY* — THE LOWEST AGGREGATE AMONG THE NIGHT'S MAJOR WINNERS, ATTACHED TO THE BIGGEST HAUL IN THE CATEGORY'S HISTORY²

THE SCEPTIC · THE BOARD UNDERNEATH THE AVERAGE

The spread is the story. **The Guardian's Lucy Mangan** filed a hundred: "rich and wonderful. Grownup, funny, scary, true." **Rolling Stone's David Fear**, at ninety, says it "plays like a mixtape of the Master of the Macabre's work." And **Slate's Sam Adams** filed **forty**, on a specific charge rather than a mood: the supporting cast is never developed, so the stakes never land and there is nobody in the town a viewer actually fears for.² A hundred and a forty average to something in the high seventies, and the mean is the one number that describes none of them. What the Academy rewarded and what the critics measured are not the same quantity — which is worth remembering the next time a score is offered as a verdict.

SOUND · TWENTY-FOUR TRACKS, AND A SAG IN EACH HALF

Carly Rae Jepsen released *Day and Night* on Friday: **twenty-four tracks** across a double album, issued whole rather than as a later B-side, with Cole M.G.N. among the producers.³ **Pitchfork's Marissa Lorusso** gives it **7.5** and does not hide the problem: "most songs don't contain the razor-sharp hooks that usually make Jepsen's work so indelible. In their absence, each side becomes a little unwieldy, and both *Day and Night* sag in their second halves." She lands warm anyway — a decade ago the concern was hammering perfect synth-pop from a dozen angles, and this is "a broader, bolder palette." The detail worth keeping: a bossa nova hidden in one closer, a sitar-like solo in another, and what Lorusso calls the cleanest ABBA homage Jepsen has managed.³

THE CHATTER · THE MAJORS OPEN A SECOND FRONT, TWICE IN ONE WEEK

Two suits, five days apart. On **15 September** Universal sued **DistroKid** in Delaware — its first ever action against the distributor — alleging it is "flooding platforms with AI-generated 'slop' that siphons revenue and listeners from legitimate artists." A thousand tracks are named at US\$150,000 each.⁴ On **18 September** Universal and Sony sued **Suno** a second time, in Boston, over **60,202 recordings** said to have been copied to build its models including the new v6. The line the filing will be remembered for: "v6 is not a fresh start; it is the fruit of the same poisoned tree... Training a 'new' model on the outputs of an infringing model does not eliminate the infringement; it launders it." Statutory exposure runs past **US\$9 billion**. Suno's answer to this filing is not yet on the record.⁴

Still running. The video-game adaptation this desk covered on Sunday opened to **US\$60 million** domestically and **US\$108.3 million** worldwide against a US\$75 million budget — the best opening for both the franchise and its director. The audience grade that came in with it is a **B-plus**, which from a horror crowd is the tell the aggregate scores were smoothing.²

SOURCES 1 78th Primetime Emmy Awards: date, venue, host and the full winners list read at source from the encyclopaedic record of the ceremony and an independent trade list of winners, cross-checked against each other. 2 *Widow's Bay* Metascore, critic count and user score read at source on Metacritic, as were the individual critic scores and the quoted lines from the Guardian, Rolling Stone and Slate reviews. Opening-weekend grosses, budget and CinemaScore for the video-game adaptation from the sourced record of the film, read at source; the trade outlets that publish the preview split refused this desk (307 to a paywall gateway returning HTTP 402). 3 Track count, release date, format and production credits for *Day and Night* from the album's sourced record; the Pitchfork review and its 7.5 score read directly at source, with all quotations Marissa Lorusso's. Metacritic's music pages are behind a challenge that refused this desk, so no aggregate score is printed for the album. 4 Both complaints as reported by Music Business Worldwide, 15 and 18 September 2026, read at source; all quoted language is from the filings as that outlet renders them. DistroKid had not commented at publication, and Suno's response to the second suit was not in the piece.

TWELVE • THE YEAR THE FIRST PHONE ARRIVES

A study of 3,486 children who had no phone at twelve. Half of them got one.

It was accepted thirteen days ago and it is not a screen-time paper. It is a paper about one particular year — and its own mediation analysis says the phone is not what does the damage.

The design is the reason to read it. Researchers took the American ABCD cohort and kept only the children who **did not have a smartphone at age twelve** — 3,486 of them. Between that assessment and the next, **1,546 acquired a first phone and 1,940 did not**. Outcomes were measured at thirteen, and — this is the part that matters — the models adjusted for each child's own score on the same measure *before the phone arrived*, along with the usual family and developmental covariates.¹

So it is not simply that sadder children get phones. These children were matched against themselves a year earlier.

What moved: internalizing symptoms, attention and total problem scores were worse in the acquiring group, all at $p \leq 0.003$; screen time rose; the odds of holding a social-media account rose; **sleep fell**; and bedtime use of social media, games and streaming rose, all below $p = 0.001$.¹

What did not move, and this is printed as prominently as the rest: **externalizing symptoms, academic performance, exercise, and bullying or cyberbullying — no association, all $p > 0.05$** .¹

$p \leq 0.003$

AND NO MORE THAN THAT. THE PAPER WAS ACCEPTED ON **8 SEPTEMBER** AND IS NOT YET TYPESET: THE ABSTRACT PUBLISHES **P-VALUES ONLY — NO EFFECT SIZES, NO CONFIDENCE INTERVALS, NO MINUTES OF SLEEP**. IN A COHORT THIS LARGE, THAT IS FULLY COMPATIBLE WITH A REAL EFFECT THAT IS SMALL¹

THE SCEPTIC • WHAT THIS DESK WILL NOT PRINT

No number describing *how big* any of this is. The full text is not public and the paper is not yet indexed, and a p-value in a sample of three and a half thousand says an effect is unlikely to be noise — not that it is large enough to matter in a house.¹

It is observational, so it cannot establish cause: nobody assigned a child a phone. The outcome instrument is a brief self-reported screening questionnaire. And the cohort is American, which matters more than usual here, because a Singaporean twelve-year-old's year is structured by an examination and a device policy no American child has.

THE CHATTER • THE MEDIATION RESULT, WHICH CUTS AGAINST THE HEADLINE

The paper's own path analysis is the most useful thing in it, and it is not a finding about phones. The indirect effects ran through **total screen time, holding a social-media account, and bedtime media use** — all at $p < 0.001$.¹

Read plainly: the handset is the vehicle, and the account and the hour are the cargo. A phone that goes into a drawer at half past nine and carries no social platform is a different object from the one the study measured — and that is the one thing both camps in this argument agree on.

AND THE SCALE OF IT, DURING SCHOOL HOURS

A separate study of **640 American adolescents aged 13 to 18**, using objective on-device data, measured **1.16 hours a day of phone use during school hours**, social media the largest share. Two caveats printed because they matter: the paper was **corrected nine days after publication**, and this desk could read only the institutional release rather than the results section.²

General guidance, not medical advice. If sleep, mood or attention are genuinely deteriorating, that is a conversation with a doctor who knows the child, not with a router.

SOURCES 1 "Smartphone Acquisition as an Inflection Point in Early Adolescence Digital Exposome: Relevance to Mental Health," Bren, Tran, Clinton, Visoki, Tran, Michel, Moore, Pimentel and Barzilay, *Frontiers in Child and Adolescent Psychiatry*, vol. 5, DOI 10.3389/frcha.2026.1940646 — received 17 July 2026, **accepted 8 September 2026**, open access. Abstract read at source on 20 September 2026. Sample sizes, exposure window, outcome measures, the covariate adjustment including the pre-exposure score, the significance levels and the mediation pathways are all as the abstract states them. The full text is not yet typeset and the article is not yet indexed in either of the two major biomedical databases; the publisher's own file endpoint returns no document. **No effect size, odds ratio or confidence interval is published anywhere in the available record, and none is printed here.** The characterisation of the outcome instrument as a brief self-reported screening measure is from the instrument's own published description. 2 Nagata, Kim, Huang, Sportsman, Hale, Baker and Christakis, "Smartphone Use During School Hours by US Youth in the Adolescent Brain Cognitive Development Study," *JAMA* 2026;335(5):453–456, published online 5 January 2026 and **corrected 9 January 2026**. The public archive holds only front matter for this paper, so the sample size, the daily figure and the quoted author come from the university's own newsroom release of 5 January 2026, read at source, and are attributed to it rather than to the paper. The authors' declared litigation consultancy is disclosed in the paper's own record.

THE OTHER SIDE · THE BAN THAT WAS MEASURED

England put restrictive phone policies in twenty schools and costed the result. The wellbeing difference was statistically nothing.

Eight hundred and fifteen pupils, twelve to fifteen. The quality-of-life gain was 0.009 years, on an interval straddling zero — and the saving was in staff time.

The SMART Schools health-economics analysis, published in *BMJ Mental Health* on 10 February, compared thirteen schools with restrictive phone policies against seven permissive ones. Incremental quality-adjusted life years: **0.009, with a 95 per cent interval from -0.014 to 0.032.** Incremental wellbeing-adjusted life years: **-0.004, from -0.044 to 0.036.** Both cross zero.¹

What the policies achieved was administrative: 3.1 full-time-equivalent staff enforcing against 3.3 in the permissive schools, a per-pupil difference of **-£94** a year. The authors' own summary: restrictive policies came with "minimal differences in quality of life or mental well-being of pupils" but "may offer some cost savings to schools."¹

Put the two pages together. The phone arriving at twelve is associated with less sleep and more bedtime use (p12); removing it during school hours does not appear to fix that. The action is at home and at bedtime, not at the gate.

THE SCEPTIC · CANDICE ODGERS, WHO HAS BEEN ARGUING THIS FOR YEARS

Candice L. Odgers, professor of psychological science and informatics at UC Irvine, is the most credible voice arguing that the panic outruns the evidence, and she is arguing it this year, not in 2019. From her talk transcript published in July: "social media is one of the least influential factors in predicting teen mental health," and "there's not one single study to date that has actually tested whether shutting it off impacts their mental health."²

Her line for a parent is the uncomfortable one: "caregiver mental health is the most important predictor of teen mental health by far," and the stressors adolescents report most often are "conflict in their home and pressure to do well at school."²

Note where she and the paper opposite agree: not on the handset. On the account and the hour.

WHAT ACTUALLY CHANGES IN THIS HOUSE, AND WHEN

January. Since the start of 2026, secondary students here may not use smartphones or smartwatches outside lesson time — including **recess, CCAs and supplementary lessons.** Primary schools have had the equivalent since January 2025.³

Half past ten. The nightly lock on school-issued Personal Learning Devices moved earlier, from 11pm to **10.30pm**, from January 2026, with sleep mode running to 6.30am. The ministry's stated reason is to "nudge them to sleep earlier."³ That is the study opposite, turned into policy, arriving in a Primary 6 household in about fifteen weeks.

THE CHATTER · AND WHAT THE REGULATOR IS CIRCLING

Two parliamentary replies on **10 September** show where this is going. Asked by **Dr Charlene Chen** how the effectiveness of age-assurance requirements will be evaluated, and whether data will be published on how far they reduce under-age access, the ministry deferred to a written answer given two days earlier.⁴ Answering **Mr David Hoe** on emerging apps, it named **Setlog and BeReal**, citing "time-sensitive prompts, camera-first design and closed-group sharing," and set out its test: effort goes to services "that have become widely used in Singapore and where safety risks have been clearly flagged."⁴ An earlier reply named the features rather than the platforms — direct messaging and autoplay. Naming mechanisms rather than brands is the more durable approach, and the one a parent can copy at home.

General guidance, not medical advice.

SOURCES 1 Perry, Goodyear, Pallan, Adab, Fenton, Michail, Patterson, Randhawa, Sitch, Wade and Al-Janabi, "Health economics analysis of restrictive school smartphone policies in secondary schools in England (SMART Schools)," *BMJ Mental Health* 2026;29(1):e301892, published 10 February 2026, open access; abstract read at source. Sample, school counts, both incremental estimates with their intervals, the staffing figures, the per-pupil cost difference and the authors' conclusion are as published. 2 Candice L. Odgers, University of California, Irvine: quotations from the published transcript of her talk on social-media bans, 6 July 2026, read at source; the ABCD observation and the correlational critique from her article in *Nature* 628, 29–30, 29 March 2024, read at source. 3 Ministry of Education secondary-school device rules announced 30 November 2025 and in force from January 2026, and the change to the Personal Learning Device lock time, reported 30 November 2025 — read via a Singapore news outlet rather than the ministry's own newsroom, whose release index returned an empty body to this desk; the substance is consistent across three independent outlets. 4 Ministry of Digital Development and Information replies to parliamentary questions from Dr Charlene Chen and Mr David Hoe, both dated 10 September 2026, read at source on the ministry's own newsroom, as is the earlier reply of 6 May 2026 naming direct messaging and autoplay. The substantive written answer the first reply defers to was not located by this desk and is not characterised here.


THE REGIME · DIVERSIFICATION STOPPED WORKING IN ALL FOUR DIRECTIONS

Equities fell, bonds fell, and gold sat twenty-two per cent under its own record.

Stocks and bonds have moved together for a month — daily-return correlation +0.57 over twenty-one sessions, +0.67 last week. That is not a bad week for a sixty-forty; it is the signature of an inflation-shock regime.

Nothing has printed since Friday. Asia does not open for four more hours and Tokyo is shut for a holiday; the only markets that traded over the weekend are the ones that never stop. Bitcoin was **US\$80,799** at twenty past five this morning against **US\$80,901** at Friday's close — a weekend that did nothing, after a Friday that added 5.9 per cent.¹

Which leaves the week itself. **Every leg lost money.** The all-world tracker fell 1.03 per cent, the global-aggregate bond proxy 0.10, the Straits Times Index 0.70; gold ended **21.8 per cent** below the high it set in January.²



The cause is one variable. Three central banks raised rates inside eight days — the ECB, the Federal Reserve **25 basis points to 3.75–4.00 per cent, unanimously**, and the Bank of Japan to 1.25 — all reacting to energy.⁴ August headline inflation in the United States ran at **3.4 per cent**, energy up 16.3 on the year; core was **2.4**.⁵

From there it is mechanical: the energy print sets the reaction function, that sets real yields, real yields set the dollar. The ten-year touched **5.04 per cent** intraweek, its highest since 2007. The yen fell 2.15 per

cent despite its own central bank hiking, because 1.25 against a four-per-cent Fed is not a contest.²

WHAT IT MEANS FOR VWRA, GLOBAL BONDS, GOLD AND BTC

Name the regime: **energy-shock re-tightening** — dollar-bid, real yields up, liquidity tightening. While oil is the marginal driver of the headline print, **bonds will not hedge equities; they will amplify them**, both repriced by the same discount rate rather than by growth. **Gold** trades as a real-rates asset, not an inflation hedge. The thing to watch is not the Fed. It is Brent. Explanation, not advice.

THE SCEPTIC · A BEAR CALL, AND ITS SCORECARD SO FAR

Dean Curnutt of Macro Risk Advisors argued on 15 September, the day before the decision, that a hike would trigger a pullback of **8 to 10 per cent** in the S&P 500, invoking 2018.⁶ The hike came, the index closed the week down 0.08 per cent, and volatility fell to 14.81. His call is so far wrong, and printed because the analogue he cites took three months: the question is whether the rally priced a credible central bank or an oil bounce that unwinds the next time Hormuz is in the headlines (p4).

THE WEEK AHEAD, AND ONE CORRECTION TO THE CONSENSUS CALENDAR

Today Japan is closed and China sets its loan prime rates. **Wednesday** brings **Singapore's August core inflation**, the last price print before the October policy statement. **Thursday** the Bank of Japan's 1.25 takes effect and Xi Jinping meets the American president (p7).⁷

And the one that isn't. Several calendars put the US personal-income-and-outlays print this week; the agency's own schedule says **30 September**. Central-bank speakers are the week's event risk.⁷

SOURCES 1 Bitcoin at 21:22 UTC on 20 September 2026 — 05:22 SGT on the 21st — read by this desk directly from the exchange chart endpoint as the edition closed; the Friday daily close and the Friday move from the same series. 2 Levels at the close of Friday 18 September 2026, each read at source: the all-world and global-aggregate trackers and the Straits Times Index from exchange chart data carrying explicit 18 September timestamps, with the index cross-checked against a second service; US Treasury yields from the Treasury's own constant-maturity series, with the 5.04 per cent intraweek print and its 2007 comparison from a market-data service; the dollar index, the yen, the S&P 500 and the volatility index from exchange data. Gold spot and its January 2026 all-time high of US\$5,608.35 from a market-data service, the high independently corroborated; the level is spot and is not a record. The all-world and global-aggregate prints are single-sourced — the fund issuer and the exchange both refused this desk. 3 Correlations are this desk's own, computed from daily closes of the S&P 500 and the US aggregate bond ETF over the windows stated; the 10 September move is from the same series. 4 Federal Open Market Committee statement of 16 September 2026 and the accompanying projections, both read at source on the central bank's own site; Bank of Japan policy statement of 18 September 2026, read at source, effective 24 September. The European Central Bank's 10 September increase is well corroborated in direction but its own statement pages returned errors to this desk, so no rate triplet is printed. 5 August consumer price index, released 11 September 2026, read at source on the statistics bureau's own release. 6 Dean Curnutt's 15 September 2026 comments as published, read at source; the index levels against which they are scored are from exchange data as above. 7 Week-ahead calendar from published central-bank and statistical-agency schedules; the personal-income-and-outlays date of 30 September 2026 read at source on the issuing agency's own release schedule. Singapore's August non-oil domestic exports print of 17 September was reported at +46.2 per cent year on year by several outlets, but the statistics department refused this desk's requests and the figure is therefore **not verified at source and is not carried as a level on this page**. This desk explains markets; it does not advise on them.

DOWN THE RABBIT HOLE • KNIFE MAKING

Everything they argue about is one trade-off: how hard, against how tough.

Steel names, quench recipes, twenty-year forum feuds — all of it a position on one axis. And the step that decides whether a blade works is the one you cannot see.

A knife is a wedge of hardened steel with a handle on it, and making one is four problems stacked. **Shape:** cut the outline from a flat bar, or hammer it hot. **Geometry:** grind the flats into a bevel, symmetrical, with a crisp transition where it stops. This is the part that takes years. **Heat treatment:** heat past the critical temperature, quench fast enough to trap the steel hard and brittle, then temper it back to buy toughness. **Handle:** scales, pins, epoxy.

Step three is where the knife is actually made or ruined, and it leaves no mark you can see. That is the fact the whole hobby is organised around, and the reason its economy runs on trust (p17).

The doors in are worn smooth. The first is a **railroad spike** — free steel that looks like a knife; a Georgia supplier runs a spike-forging class at **US\$200**.¹ The joke is structural: spike steel is too low in carbon to harden properly, so falling in love with one and then finding it will not hold an edge is the on-ramp to metallurgy.

The second is a weekend. American one-day classes run **US\$325–399**, a Sydney workshop **A\$559–599**. The serious version is two weeks at the Bill Moran School for **US\$1,100**, where the college's own copy says most students finish with a blade “capable of passing the American Bladesmith Society Journeyman test” — and which knocks a year off the eligibility clock (p17).²

The third is television: *Forged in Fire*'s own channel carries **1.04 million subscribers**, and is at once the hobby's biggest recruiter and its most reliable irritant.³

1.67M vs 44.9K

SUBSCRIBERS FOR **ALEC STEELE**, WHO HITS HOT STEEL ON CAMERA, AGAINST **KNIFE STEEL NERDS** — THE CHANNEL OF THE METALLURGIST WHO DESIGNED THE ALLOY EVERYONE ARGUES ABOUT³

That ratio is the hobby in miniature, and it is not a complaint. Spectacle recruits; the metallurgy keeps the people spectacle brought in.

THE CHATTER • THE SCALE OF THE WORLD

The institutions are older than the internet. The **American Bladesmith Society** was founded in **1976** by William F. Moran and adopted its Apprentice–Journeyman–Master ladder in 1988; there are **114 living Master Bladesmiths**. The **Knifemakers' Guild** dates to **1970** and admits by peer review: four recommendations from voting members, and four handmade knives displayed at two successive annual shows.⁴ The gathering is **Blade Show**, Atlanta, **5–7 June 2026**. Online, the forum that let this desk in is **KnifeDogs**: **23,761 members, 361,903 messages**.⁵

THE SCEPTIC • WHAT THIS DESK COULD NOT GET INTO

BladeForums — the old-growth forest of this hobby, where the technical archive lives — refused every request on every path. No member count, and not one post quoted at source in these three pages.⁵ Reddit was unreachable as usual. That matters, because the missing forum holds the institutional memory. Everything attributed in the next two pages comes from KnifeDogs, the societies' own documents, suppliers' catalogues, or a named metallurgist's published testing.

WHAT THE NEXT TWO PAGES DESCEND INTO

Page 16 — the gear, and a supply shock nobody outside the hobby noticed. Two steel mills closed within six months, and the shortage is visible on a product page today. The replacement is measurably better than what it replaced.

Page 17 — the test, the money and the trust. An examination that ends with the candidate bending his own blade ninety degrees in a vice; what a knifemaker earns an hour; and the brand that closed in March after its steel was sent for spectrometry.

SOURCES 1 Pop's Knife Supply catalogue, read 20 September 2026 through the shop's own product feed: the railroad-spike class, its price and the quoted description. 2 Class prices from the providers' own listings; the Texarkana College “Introduction to Bladesmithing” price, duration, instructors and the quoted outcome from the college's own community-education page, read at source. 3 Subscriber counts read directly from the channel pages on 20 September 2026. 4 Founding dates, the rank structure, the count of living Master Bladesmiths and the Knifemakers' Guild membership requirements from the societies' own pages and the encyclopaedic record; the American Bladesmith Society's own HTML pages returned 403 to this desk, though its rule PDFs served (see p17). 5 Blade Show dates and self-description from the show's own site, read 20 September 2026; KnifeDogs statistics read from its own front page the same day. BladeForums returned a challenge page on every path attempted, and Reddit returned 403 on every endpoint; subreddit sizes of roughly 182,000 and 170,000 members come from a third-party statistics mirror and are marked secondary here and throughout.

THE DEEP END · THE GEAR, AND THE YEAR THE SUPPLY BROKE

Two mills closed six months apart. You can see it on a product page today.

Every size of W2 and of 52100 is out of stock at the biggest beginner supplier in America — a German furnace shutting in October 2025, arriving on a shopping cart in September 2026.

Steel comes in four families. **Simple carbon** — 1084, 80CrV2, W2, 52100 — is iron, carbon and little else: cheap, forgiving, and the only route to a *hamon*. **Non-powder stainless** is razor stock, sharp but needing a real oven. **Powder** steels — 3V, S35VN, MagnaCut — are the money tier. A bar of 1084 is **US\$33.75**; the same length of MagnaCut is **US\$155**.¹

Then the shock. **Crucible Industries ceased operations in 2025**, ending American powder production of the CPM line; **Buderus Edelstahl shut on 31 October 2025**, taking W2, 80CrV2, Nitro-V and 15N20 with it. The American supplier that had bought from Buderus for fifteen years published the warning: **“Once supplies run out some alloys may not be returning for a long time. If at all.”**²

Out of stock

EVERY SIZE OF **W2** AND EVERY SIZE OF **52100** AT POP’S, READ 20 SEPTEMBER 2026. IF YOUR READER WANTS A HAMON BLADE THIS YEAR: BUY IT WHEN YOU SEE IT.¹

And then the twist. MagnaCut is now made by Niagara Specialty Metals and Erasteel, pressed in Sweden. When **Larrin Thomas** tested the new supply against the original in January, the Erasteel version showed **significantly higher toughness** — matching AEB-L at 62 Rockwell, beating it at 64 — and took a hotter austenitising temperature for roughly **2.5 more points of hardness at the same toughness**.³ The American mill died and the European replacement is better, and almost nobody knows it.

THE CHATTER · THE ARGUMENT OF SEPTEMBER 2026

The Damascus fight is no longer about whether pattern-welding performs. It is about whether the Damascus is real. **Black Mamba Knives** on KnifeDogs, **9 September 2026**: “Acid etched pattern that sits on the surface... At a glance it fools a lot of buyers, and honestly it fools some sellers too.”

His tells: the pattern must continue over the spine and the choil; and the two faces of a forged billet are never identical, so “if the pattern is mirror perfect... it came from one file, not from the anvil.”⁴

THE GRINDER — THE ONE BIG PURCHASE

Why 2×72 and not 1×30: ecosystem and heat. It is the industrial belt size, so every abrasive exists in it, a longer loop runs cooler, and the machines are platforms rather than sealed appliances.

The ladder, priced 20 September. Torus 2×72 assembled **US\$845** — but the motor and drive is **US\$1,495** on top, landing a complete one near US\$2,340, which is KMG territory. Beaumont’s KMG-PL is **US\$2,840**; Wuertz’s TW-90 **US\$4,100–4,550**.⁵

And the consumable is not steel. Ceramic belts run **US\$5.25–20.50** each and a knife eats six to ten. One knife a month is **US\$500–800 a year**. A forty-dollar bar of 80CrV2 makes four knives; the belts to grind them cost more than the bar.¹

THE SCEPTIC · THE STEP THAT SHOULD STAY OUTSOURCED

An Evenheat KO 18 — the standard heat-treating oven — is **US\$2,600**, draws thirteen amps at 240 volts continuously and radiates 2,400°F; a Paragon professional reaches **US\$6,190**.⁶ Commercial heat treating has long been quoted in the tens of dollars a blade.

Print the flashpoints, because this is where people get hurt: Parks 50 quench oil flashes at **275°F**, and you are lowering 1,500°F steel into it. The right kit costs less than a belt.¹ One caution on outsourcing: the best-known heat treater’s pricing page is behind a challenge this desk could not pass, and a second operation’s domain was **disabled** on 20 September.⁷

SOURCES 1 Steel, belt and quench-oil prices and the stock status of W2 and 52100 read on 20 September 2026 from Pop’s Knife Supply’s own product feed, cross-checked against Jantz Supply’s catalogue for the same alloys. Flashpoint and quench speed as the oil’s own specification states them. 2 New Jersey Steel Baron, “Steel Signals #1 — On the Closures of Crucible and Buderus,” 15 October 2025, read at source, for both closures, the affected alloys, the new Niagara–Erasteel production arrangement and the quoted warning. 3 Larrin Thomas, “Testing Erasteel MagnaCut (vs Crucible Original),” Knife Steel Nerds, 12 January 2026, read at source, for the toughness comparison, the austenitising temperatures, the hardness differential and the corrosion result. 4 Black Mamba Knives, KnifeDogs thread of 9 September 2026, read at source; all three quotations are verbatim from that post. 5 Grinder prices read on 20 September 2026 at each maker’s own site. Two further recommended American makers publish no readable prices and are therefore not quoted. 6 Oven prices read the same day from the two suppliers’ own catalogues. 7 The commercial heat treater’s pricing page is behind a captcha this desk could not pass, so no 2026 price is printed; a second operation’s domain returned a disabled-domain error on 20 September 2026.

THE DEEP END · THE TEST, THE MONEY, THE TRUST

The examination ends with the candidate bending his own blade ninety degrees in a vice.

It may crack at the edge, but not more than a third of the way through. If anything breaks off, he fails — and waits six months.

The American Bladesmith Society's Journeyman test is the most interesting ritual in the hobby, because every part of it is a proxy for something invisible. A Master Smith administers it, on a blade the candidate forged alone — “no other person physically assisting” — from **homogeneous carbon steel**. Once it begins, “**no work, not even light stropping, may be done to the test blade.**”¹

Four tests. **Cut** a free-hanging one-inch sisal rope through in one stroke — edge geometry. **Chop** a two-by-four through twice, after which “any nicks, chips, flat spots, rolled edges, or other deformations” fail it. Then **shave hair with the part you chopped with**. Then the vice.¹

The bend demonstrates a soft spine and a hard edge in one blade. Fail and you wait six months; pass and you still lay five knives before a judging panel, where “the failure of any one knife may cause the applicant to fail.” And note the reversal: Damascus is **forbidden** for the Journeyman test and **mandatory**, three hundred layers on a hidden tang, for the Master. Prove the knife works before you make it beautiful.¹

THE SCEPTIC · MARCH 2026, AND WHAT A BUYER CANNOT CHECK

Bark River Knives closed on 20 March 2026, after model runs were made with **Chinese steel stamped as American CPM-154**. The admission is the founder's own: **Mike Stewart** “had chosen to use alternate Chinese steel in lieu of the U.S.-made CPM-154/154CM indicated on the blades.” Dealers sent pre-controversy stock for **X-ray fluorescence testing**.²

This is the keystone: every claim in this hobby — every steel name, every Rockwell number — is unverifiable without a spectrometer and a hardness tester, and the instrument gap runs from a **US\$120** set of files to a bench tester costing thousands.³

THE CHATTER · THE METALLURGIST ANSWERS BOTH SIDES

Larrin Thomas, February 2026, on the forum line that premium steels are marketing: his verdict is “**They are different. And even if you didn't notice that doesn't mean they weren't different.**”

Then the concession that guts the marketing: **AUS-6, a budget steel, matched S110V** when the edge angles differed. Geometry beats alloy — **US\$34** of 1084 ground thin out-cuts **US\$155** of MagnaCut ground thick.⁴

WHAT IT PAYS, AND WHAT IT COSTS

The wage. American knifemaker pay averages **US\$17.79 an hour**. A hand-ground blade is twelve to twenty hours; at a mid-tier US\$500, with US\$200 left, that is ten to seventeen an hour before overhead and belts.⁵ And the top is not a ladder: **Kramer's knives went from US\$475 in 2008 to US\$30,000 in 2015.**⁶

The ninety days, from here. Singapore rules out most of page 16: no garage, no open flame, no 220-volt circuit. Book a class at **Tombalek** in Sungei Kadut, the only forge here open to the public. Then **stock removal, heat treatment sent out** — 1084, Micarta, pins, epoxy, files and sandpaper, about **S\$330**, no fire at any point. Only if you are still in at day ninety does a grinder make sense, and the expensive mistake there is electrical: **get it in writing that the drive is rated for 50 Hz single-phase input.**⁷

NEXT DESCENTS — PICK TOMORROW'S RABBIT HOLE

Leathercraft — the hand saddle-stitch and a tool economy that never depreciates.

Falconry & birds of prey — a licensed craft with a mandatory apprenticeship.

Home cheesemaking — cultures, rennet, and a cave in a spare fridge.

SOURCES 1 American Bladesmith Society, “Journeyman Smith Applicant,” revised 4 November 2019, and the corresponding Master Smith document — both read at source as the society's own PDFs; all quoted conditions, the four tests, the failure criteria, the six-month wait, the five-knife judging and the Damascus rules are verbatim from them. The society's HTML pages returned 403 to this desk; its PDFs served. 2 KnifeNews, 30 March 2026, read at source, for the closure date, the mislabelled steels, the founder's quoted admission, the dealers' spectrometry testing and their response. 3 Hardness-tester price points are drawn from forum discussion surfaced in search and are **secondary**: the host was unreadable to this desk, and the figures are printed as indicative only. 4 Larrin Thomas, “Knife Steel — Can You Tell the Difference?”, 23 February 2026, read at source, for the cutting-test range, the quoted comparison, the conclusion and the AUS-6 against S110V geometry result. 5 Hourly and annual pay figures as published by a recruitment data service as of 1 January 2026 and cited in an abrasives-industry article of 6 May 2026, read at source; the per-knife arithmetic is this desk's own, from the hours and margin given there. 6 Kramer price history and the 2019 auction result from the sourced record and the auction-house reporting, read at source. 7 Tombalek's location, hours and public-forge status from its own materials and Singapore press coverage of October 2024; its current class price is behind a challenge this desk could not pass and is **not printed**. The quoted local maker posted on a Singapore forum on 28 March 2019, read at source. The electrical caution is this desk's own, from the makers' published motor and drive specifications. Nothing on these three pages is attributed to Reddit or to BladeForums, neither of which was reachable.

THE LONG GAME · A TRIAL THAT SPLIT THE SESSION UP

Three short bouts across a day beat nothing, and lost to one session on every measure.

The differences were not statistically significant — which is not the same as equal. The finding worth having is that the fragmented group enjoyed it more.

The trial, published on 1 September, randomised **45 recreationally trained men aged 18 to 25** into three arms for eight weeks: fragmented training — three short daily bouts of twenty squats — against one continuous session of three sets of twenty, against a time-matched control. Same exercise, same machine, same volume.¹

Both arms beat the control clearly. One-rep maximum rose **20.59 kg** fragmented, interval 8.69 to 32.49, and **28.93 kg** traditional, 17.05 to 40.82. Between them, no difference reached significance.¹

p = 0.034

THE ONE PLACE THE FRAGMENTED GROUP WON OUTRIGHT: **ENJOYMENT**. OVER THIRTY YEARS, NOT A TRIVIAL VARIABLE¹

Read the point estimates rather than the p-values and the picture is less friendly. The continuous session beat the fragmented one on **every single measure**; failing to detect a difference is not demonstrating equivalence. This is a permission slip to split a session when life is the constraint, not evidence that splitting it is as good.

THE SCEPTIC · WHO WAS ACTUALLY IN THE ROOM

Forty-five young men, eighteen to twenty-five, at one Chinese university, one exercise on one machine for eight weeks. Not middle-aged, not women, not free weights.¹

The transfer to a reader in his forties is narrow, and the claim that survives is modest: **better than the session you skip**.

AND A META-ANALYSIS THAT FLATTERS NOBODY'S RING

A systematic review published on 22 May pooled sixteen studies — ten in the meta-analysis — of

SOURCES 1 Zhai, Yang, Dai, Wang, Jiang, Wu and Wang, "Fragmented versus traditional resistance training: a randomized controlled trial on muscular strength, endurance, and enjoyment in recreationally trained young adults," *Journal of Exercise Science & Fitness* 2026;24(4):200521, published online 1 September 2026, trial registered; abstract read at source. Sample, arms, protocol, all three effect estimates with their intervals, the between-arm p-values and the enjoyment result are as published, as is the authors' own note on adherence. 2 Agostinho, Borges, Pereira, Borges and Soares, "Are Wearable Sleep-Tracking Devices Reliable Alternatives to Polysomnography? A Systematic Review and Meta-Analysis," *Behavioral Sleep Medicine*, published online 22 May 2026; abstract read at source. Study counts, the direction of the biases, the device-specific findings and both quoted conclusions are as published. 3 The absence of independent validation for proprietary readiness and recovery composites is this desk's own search result: what it found were vendor and affiliate pages, which are excluded, so no figure is printed. 4 Jack McNamara, senior lecturer in clinical exercise physiology, University of East London, writing in *The Conversation*, 22 June 2026, read at source. The cohort study itself was covered in an earlier issue and is carried here only as a one-line delta. 5 Peter Attia and Kathryn Birkenbach, "With exercise, results matter more than 'time served'," 4 May 2024, read at source. This desk found no published comment by any of the named longevity voices on the fragmented-training trial above, and has invented none. *All of the above is general guidance, not medical advice.*

consumer sleep wearables against laboratory polysomnography. The result: **wearables systematically overestimate total sleep time and sleep efficiency and underestimate wake after sleep onset**, and **"no device demonstrated consistently superior performance."**²

Within individual studies the third-generation Oura ring came closest on latency and efficiency, certain Fitbit models on deep and REM. The authors conclude these devices "provide reasonable estimates of global sleep metrics" while their variability "currently limit[s] their use as stand-alone diagnostic tools."²

THE CHATTER · TRUST THE TREND, NOT THE STAGE

The error is not random: these devices err toward **flattering** you — more sleep, better efficiency, less waking. A ring is decent at telling you whether last night beat your own last month. It is not measuring your REM.²

One distinction this desk will not blur: a validated sensor is not a validated score. The composite readiness and recovery numbers are proprietary and were **not** measured here.³

Still on the shelf. The 147,374-person cohort on resistance-training dose, covered here in June, has had no replication since; its ceiling — benefit flattening past roughly two hours a week, aerobic work doing more than lifting — stands. **Jack McNamara's** caution holds: it was observational, "so although it can show a strong link... it can't prove that one directly causes the other."⁴

The unit problem, from the other camp. Peter Attia's objection is that minutes a week is the wrong measure: "we should only concern ourselves with duration insofar as it influences what we really care about: results."⁵

Epidemiology measures what its questionnaire asked, and it asked about time. *General guidance, not medical advice.*

THE DIARY · EIGHTEEN DAYS TO THE STREET CIRCUIT

What is worth booking, and two dates this desk will not print as settled.

Every row below was re-checked against a promoter, a venue or an official ticketing page. Two shows carry conflicting dates across sources and are flagged rather than dropped; Singapore's three biggest ticketing hosts refused this desk outright, so several famous names are absent for that reason alone.

SINGAPORE

BOOK NOW · PLAN AHEAD

- 24
SEP

The Radicalisation of Mrs. Mary Lim-Rodrigues — opens
 WILD RICE @ FUNAN · RUNS TO 10 OCT · A PRINCIPAL'S JUDGEMENT TESTED BY AN INCIDENT INVOLVING A PUPIL.
 ADULT¹
- 1
OCT

Satay by the Bay — last day
 GARDENS BY THE BAY · CLOSING FOR REDEVELOPMENT. GO THIS WEEK OR NOT AT ALL²
- 2
OCT

Octoburst! — the Children's Day festival
 ESPLANADE · PROGRAMMES TO 4 OCT · **EVERY EVENT FREE** · PITCHED AT FOUR-TO-TEN, SO A TWELVE-YEAR-OLD IS AT THE TOP OF THE BAND¹
- 9-11
OCT

★ Formula 1 Singapore Grand Prix
 MARINA BAY STREET CIRCUIT · **MOST THREE-DAY GRANDSTANDS AND HOSPITALITY ALREADY SOLD OUT**; STILL OPEN WHEN READ — STAMFORD DAILY FROM \$228, ZONE 4 WALKABOUT FROM \$198, WHEELCHAIR-ACCESSIBLE EMPRESS PLATFORM FROM \$38³
- 16-25
OCT

Fat Kids Are Harder To Kidnap 2026
 WILD RICE @ FUNAN · 16-18 AND 23-25 OCT · THIRTY-ONE PLAYS IN AN HOUR. PRICES NOT PUBLISHED ON THE LISTING¹
- 12
NOV

★ Monkey Goes West — opens
 WILD RICE @ FUNAN · THE YEAR-END MUSICAL: SONGS, GAGS, WUSHU. WILD RICE YEAR-ENDERS SELL OUT; THE CLOSING DATE IS NOT YET POSTED¹
- 13-22
NOV

Singapore Writers Festival — "Hidden in Plain Sight"
 ARTS HOUSE AND FESTIVAL VENUES · **YOUTH PASS \$15 → \$10.50** AND FESTIVAL PASS \$30 → \$21 ON THE SUPER-EARLY-BIRD, LIMITED. THE BEST VALUE ON THIS PAGE¹
- 25
NOV

Guns N' Roses — World Tour 2026
 NATIONAL STADIUM · 8PM, ONE NIGHT · GENERAL SALE OPEN; THREE VIP TIERS. PRICES ARE NOT SHOWN ON THE PROMOTER'S EVENT PAGE¹
- 12
DEC

Slowdive — date disputed
 FORT CANNING GREEN · THE PROMOTER'S OWN SITE SAYS 12 DEC; A WIDELY-READ LOCAL ROUND-UP SAYS THE 21ST. **GOING WITH THE PROMOTER**, AND SAYING SO¹

JAPAN & THE REGION

PLAN THE FLIGHTS

- 23
SEP

Ron Mueck — closes
 MORI ART MUSEUM, ROPPONGI HILLS · THIS WEEK. MARIKO MORI, "ALL THAT SHINES," OPENS 31 OCT AND RUNS TO 28 MAR⁴
- 19
SEP-4
OCT

★ Asian Games — Aichi-Nagoya
 ~53 VENUES ACROSS AICHI, WITH SWIMMING AT THE TOKYO AQUATICS CENTRE · RUNNING NOW · OVERSEAS TICKETS NEED A SUPPORTER REGISTRATION; THE TICKET PAGE DID NOT RESOLVE FOR THIS DESK⁵
- 26
OCT-4
NOV

★ Tokyo International Film Festival
 HIBIYA, YURAKUCHO, MARUNOUCHI, GINZA · CENTREPIECE *GODZILLA MINUS ZERO* · **THE PRE-ORDER LOTTERY HAS ALREADY OPENED**⁴
- 29
OCT

MTV VMAJ 2026
 TOKYO DOME · 16:00-18:00⁴
- 3
NOV

Lexus Melbourne Cup Day
 FLEMINGTON, MELBOURNE · GENERAL ADMISSION A\$120, TWO-DAY PASS A\$150; HILL STAND RESERVED FROM A\$575 AND FLAGGED SELLING FAST. SEVERAL MARQUEES ALREADY GONE⁶
- 6-8
NOV

Stray Kids — Run It Japan
 TOKYO DOME⁴
- 5-6
DEC

YOASOBI — Super Planet
 TOKYO DOME · THE TEN-CITY DOME AND STADIUM RUN⁴
- 23-26
DEC

Snow Man — All Suite
 TOKYO DOME · FOUR NIGHTS OVER CHRISTMAS⁴
- 4
JAN

★ Wrestle Kingdom 21
 TOKYO DOME · THE TRADITIONAL FOURTH-OF-JANUARY SHOW. THE SINGLE BEST THING ON THIS PAGE TO DO WITH A TWELVE-YEAR-OLD⁴

SOURCES 1 Singapore theatre and concert rows read at the venue's or promoter's own site on 20 September 2026: Wild Rice, Singapore Repertory Theatre, Esplanade, the Singapore Writers Festival and the promoters' own event pages. The Slowdive conflict is between the promoter's own listing and a local aggregator last updated 2 September; the promoter's date is printed. Singapore's major ticketing hosts and the national stadium's own site all refused this desk, which is why several widely-advertised shows — and a postponed one whose replacement dates conflict across sources — are absent rather than listed. 2 Gardens by the Bay's own what's-on page, read at source. 3 Formula 1's own 2026 calendar for all race dates; ticket availability and prices from the Singapore Grand Prix's own ticketing page, read 20 September 2026 — availability moves daily. 4 Tokyo Dome's own English events page, the Mori Art Museum's own exhibitions index and the film festival's own site, all read at source. 5 Games dates and venue spread from the organising committee's own site; the swimming venue is confirmed by a result reported from it on 20 September (p5). The official ticket path returned a 404 to this desk, so no price is printed. 6 Victoria Racing Club's own race-day page, read at source, for admission prices and the sold-out marquees. Grand Sumo's November tournament is confirmed by its own site as sold out on every date, but no schedule path resolved, so it is not listed as a row.

THE TABLE · FOUR OPENINGS, TWO OF THEM UNREVIEWED

An eighteen-seat basement doing a thirty-six-hour black curry, and a hawker legend that nobody is queueing for.

The Jewel flagship has had its verdict, and it is not a rave. The two most expensive rooms on this page have had none at all, said here plainly rather than dressed as anticipation.

Kuro Kare, Havelock II, is the one to book: a Japanese curry specialist built around a **thirty-six-hour slow-cooked black curry**, run by **Aidan Wee** and **Ho Bing Rru**, both formerly of Fleurette with time at one-starred Meta behind them. It is the relocation of the cult outlet that closed near the university, and an **eighteen-seater**.¹

Prices sit where the pedigree does not: fried chicken curry **\$12.90**, Iberico pork katsu **\$16.90**, charcoal-grilled striploin don **\$18.90**, and **GST and service absorbed at lunch**. Eatbook told readers in July to go “before the crowds start to flock there”; by September Time Out reported it **packed daily**. Not halal-certified. **Dinner by reservation, lunch walk-in only**: arrive at 11.30 sharp, or book a week out.¹

The one with a verdict is **Tian Tian**’s first proper restaurant, at Jewel — ninety-eight air-conditioned seats, soft-opened 11 August, grand-opened on the 22nd. The set is **\$11** against **\$6** for a medium plate at the Maxwell stall, with outlet exclusives including a steamed chicken curry at \$13.80 and a heritage platter at \$25.80.²

Time Out’s **Adira Chow** found the profile intact — “soft rice, silky chicken and sweet sauce” — at “significantly higher prices,” then criticised rice “soft, slightly broken and a touch starchy” and a thicker, sweeter gravy. Her verdict is the hedge itself: more popular with “international travellers looking to get a taste of Tian Tian between flights.” An invited tasting scored the chicken and the curry 4.2 out of 5, the platter 3.²

THE CHATTER · THE QUEUE THAT ISN’T

The detail worth the trip, or rather worth not making it: **neither reviewer reported a queue**. For a Tian Tian outlet that is unusual, and it suggests the flagship trades on airport footfall rather than pilgrimage — a good pre-flight proposition and a poor reason to cross the island.²

AND TWO ROOMS NOBODY HAS JUDGED YET

Assaggi, Neil Road, opened in August: Italian, set menus only, the name meaning “tasting,” the house line “imperfect by choice.” Head chef **Dong Kai Ti** comes from years alongside **Denis Lucchi** at one-starred Buona Terra, who still guides it. Express lunch **\$58++** for three courses; the Signature and Pasta journeys **\$88++**.³

Le Ju Xuan opened this month in the Marina Bay Sands Tower 2 lobby: **188 seats**, regional Chinese, the accessible sibling to its Cantonese fine-dining counterpart, with marble goby and a Sichuan mapo tofu with sea cucumber among the signatures and a claimed **seventy-plus label baijiu collection**. Explicitly family-friendly — on this page, the one built for a table with a twelve-year-old at it.³

THE SCEPTIC · TOO NEW TO CALL, SAID PLAINLY

Neither Assaggi nor Le Ju Xuan has been reviewed by any of the Singapore food outlets this desk reads, or by the Michelin guide. What exists is listing coverage — a seat count, a baijiu wall, a lineage — and listing coverage is not praise. No chef is named for Le Ju Xuan anywhere this desk could read, and no price is published for it, which is why none appears above.³ The Buona Terra lineage is real, and an eighty-eight-dollar tasting menu is cheap for the genre in this city. Whether either room is good is not yet knowable. Ask again in six weeks.

On the bench. A Shibuya kaisendon brand’s first outlet abroad at Plaza Singapura, capping its top bowl at thirty a day; a Korean fried-chicken chain back four years after closing every Singapore store, and **not halal-certified**, which readers will assume wrongly; and the century-old kaya-toast institution’s second outlet, at Nex.³

SOURCES 1 Kuro Kare’s own menu and booking page, read at source on 20 September 2026, for the chefs, the address, the service hours, the seat count, every price quoted and the reservation rule. Reception from Eatbook, 20 July 2026, and Time Out Singapore, 3 September 2026, both read at source; the halal status is Eatbook’s. 2 Tian Tian Jewel flagship: address, seat count, opening dates and prices as reported at source; the comparison price at the original stall as published. The quoted review and its verdict are Adira Chow’s in Time Out Singapore, 21 August 2026, read at source. The numerical scores come from an invited media tasting published 20 September 2026 and are labelled as such, since an invited tasting is not an independent one. 3 Assaggi’s menu, prices and the chef’s lineage from the restaurant’s own site, read at source — note its own spelling of the guiding chef’s name, which one listing renders differently. Le Ju Xuan’s seat count, location, signatures and baijiu claim from Time Out Singapore’s listing of 3 September 2026; **no review of either restaurant exists at any outlet this desk reads, and no price is published for Le Ju Xuan**. Bench items from Eatbook’s and HungryGoWhere’s September listings, read at source, which are listings rather than reviews. SETHLUI and several local outlets refused this desk’s requests.

THE UNDERCURRENT · VIETNAM

Không nhận chuyển khoản

“We don’t accept bank transfers.” Handwritten, taped to a till, in the country that grew QR payment value eighty-five per cent in a year. Vietnam has just run the world’s first large experiment on whether digital-payment adoption is a ratchet. It is not.

The signs went up in June 2025, and then the QR codes came off the counters. Both are small acts and both are legible: a shopkeeper with no lobbyist, no association and no column has exactly one lever, and it is the decision about how money enters the shop.

What they were refusing was not a tax. It was a record. On 1 June 2025, Decree 70 obliged household businesses above a revenue threshold to issue electronic invoices from cash registers **wired directly to the tax authority**. Seven months later, on 1 January 2026, Vietnam abolished **thuế khoán** — the lump-sum tax that had stood for more than thirty years — for some five million household businesses that together supply **39 per cent of the country’s jobs and about a quarter of its GDP**.¹

Under the old system a tax officer and a local council estimated your annual revenue, and you paid a flat monthly sum against the estimate. You kept no books, issued no invoices and proved nothing. The estimate was negotiable, which is to say it was social. Under the new one you declare actual revenue, monthly or quarterly, and you file **even in a month when you sold nothing** — five million nil returns a year being a category of paperwork that did not previously exist.²

7×

THE FINANCE MINISTRY’S OWN ARITHMETIC: THE AVERAGE LUMP-SUM PAYMENT WAS ABOUT **₫672,000–700,000 A MONTH**; THE AVERAGE UNDER DECLARATION IS ABOUT **₫4.6 MILLION** — “NEARLY SEVEN TIMES HIGHER.” SAME SHOP, SAME STREET, SAME BOWL OF PHỞ³

The ministry published that comparison as the *justification* for the reform — proof that the old system let the larger operators pay least. Read from the counter, it is a notice of a sevenfold increase arriving with a bookkeeping requirement attached.

WHAT THE OLD SYSTEM ACTUALLY COST

The flat monthly figure was never the whole payment. Radio Free Asia quotes a shop owner paying **₫2.5 million a month to local officials** — roughly four times the official average assessment — as the unofficial layer that sat on top of a negotiable estimate.⁴ A system that is social is also a system that is discretionary, which is the strongest case the reformers have, and they made it.

Adoption was bought on an unstated premise: that nobody at the counter thought of the rail as a ledger.


The response was immediate and physical. Radio Free Asia reported on 11 June 2025 that stalls were shutting across the country, with **eighty per cent of shops closed in Nghệ An’s largest market**; traders at An Đông, Bình Tây, Tân Bình and Bến Thành in Ho Chi Minh City pulled their shutters down.⁴ By July, VOV was describing Hàng Đào street in Hanoi’s Old Quarter with shutters half-lowered through the day and trade conducted behind them — what the traders themselves call **buôn bán du kích**, guerrilla trading.⁴

The anxiety was never mainly about the rate. It is about the input side. This sector buys from farmers, backyard workshops and informal importers who issue nothing at all — the model traders call **ba không**, three nothings: no invoice, no origin papers, no declaration. With no documented cost you are taxed on turnover rather than margin, and a trader at An Đông explained the rest to a Saigon paper in plain terms: the goods come from small independent suppliers, so the invoices do not exist to be produced.⁴ A fifty-year veteran at Bình Tây added the other half — most stallholders simply cannot work the software, training sessions notwithstanding.

SOURCES 1 Decree 70/2025/NĐ-CP, in force 1 June 2025, and the abolition of the lump-sum regime from 1 January 2026, from Vietnamese government and government-press accounts read at source; the sector’s share of employment and output as reported by Radio Free Asia, 11 June 2025. The count of household businesses varies by definition across official sources — roughly 2.2 million operating regularly, 3.6 million registered, and “over five million” in the policy documents, which include individual businesses — and that range is printed as a range here rather than resolved. 2 The declaration mechanics, the nil-return obligation, the sector VAT rates and the two personal-income-tax routes from the government portal’s own explanations of Decree 68/2026/NĐ-CP and Circular 18/2026/TT-BTC, both dated 5 March 2026, read at source. 3 Finance ministry figures as reported by Dân Trí, 11 June 2025, read at source by this desk: “Mức thuế khoán bình quân của mỗi hộ kinh doanh hiện khoảng 672.000-700.000 đồng/tháng... cao gấp gần 7 lần.” 4 Market closures and the Nghệ An figure from Radio Free Asia, 11 June 2025; the Hàng Đào description and the term buôn bán du kích from VOV, 19 July 2025; the An Đông and Bình Tây traders from Sài Gòn Giải Phóng, 6 June 2025. All read at source; all rendered in this desk’s own words except where quoted.

Now the measurable part. Between September 2025 and the end of January 2026, cash in circulation as a share of Vietnam's total means of payment rose from **9.48 per cent to 11.52** — its highest in three years. In money, roughly **₫360 trillion** walked out of the banking system in four months.¹

A second series says the same thing from the other side. In the first quarter of 2026, cashless payments rose **37.98 per cent by volume and only 14.22 per cent by value** — a twenty-four-point divergence. Many more transactions; far less money moving through them. That is precisely the shape you would expect if small purchases stayed digital and larger ones went back to notes.¹



It shows up in the banks. **Twenty-four of twenty-seven listed Vietnamese banks** reported falling current-account balances in the first quarter of 2026; by the end of June, MB's ratio was down 3.25 points on the year-end and MSB's down 6.23.² Current-account money is the cheapest funding a bank has. It is also the money that leaves first when a counter stops taking transfers.

A nation's payments strategy now turns on which of two identical black-and-white squares is stuck to the till.

The state saw it and pressed harder. In November 2025 the central bank convened a seminar on **separating merchant QR codes from personal**

transfer QR codes — because a personal QR is how you receive money without recording revenue. On 10 March 2026 the finance ministry issued a dispatch warning that “some organizations and individuals are exploiting cash payments or using personal accounts to conceal revenue,” and instructing provinces to **promote cashless payment in order to improve tax management**.³ That is the quiet part, in an official document, over a minister's signature. And from **20 April 2026**, every household business had to hand the tax office a list of every bank account and e-wallet it used to take money.³

Nine days later the state retreated. **Decree 141/2026**, dated 29 April, raised the taxable threshold from **₫500 million to ₫1 billion** a year, backdated to 1 January, taking **2,556,042 household businesses** back out of the net at a self-costed revenue loss of **₫16.65 trillion**.⁴ It is rare to see a government put a precise price on a retreat and then publish it.

84.4%

SHARE OF 1,001 HOUSEHOLD BUSINESSES SURVEYED BY THE CHAMBER OF COMMERCE IN FEBRUARY–APRIL 2026 WITH **NO PLAN TO BECOME A COMPANY**. THE REFORM'S STATED GOAL WAS TWO MILLION ENTERPRISES BY 2030, HARVESTED FROM THIS SECTOR⁵

226,543 → 93,988

HẢI PHÒNG'S REGISTER AGAINST THE NUMBER ACTUALLY TRADING, IN ONE CITY, IN AUGUST 2026. IN THE FIRST HALF OF THE YEAR THE PROVINCE LOGGED 5,149 NEW REGISTRATIONS AGAINST **3,653 SUSPENSIONS** AND MORE THAN THREE THOUSAND DISSOLUTIONS⁶

The same survey found **81 per cent** reporting falling revenue in 2025, a third planning to shrink, and 4.4 per cent expecting to close within two years. Its most quietly devastating line came from the chamber's own deputy chief of legal affairs, **Phạm Ngọc Thạch**: treating household businesses as feedstock for an enterprise target will fail, because most of them do not regard becoming a company as a step up at all.⁵

SOURCES 1 Cash-in-circulation share for September 2025 and end-January 2026, the implied ₫360 trillion and the first-quarter divergence between transaction volume and value, from central-bank data as analysed by SHS Securities and reported by Vietnamese financial press on 27 and 31 May 2026, read at source. Two published observations, not a continuous series — the chart above plots exactly those two points and says so. 2 Bank current-account ratios from Vietnamese banking press, 15 May and 12 August 2026, read at source, reporting first-quarter and end-June figures respectively. 3 The central bank's November 2025 seminar on separating merchant from personal QR codes, and finance ministry Dispatch 01/CĐ-BTC of 10 March 2026 including the quoted warning and the instruction to provinces, read at source on the government portal. The 20 April 2026 account-disclosure deadline from the government portal's own account of Decree 68/2026/NĐ-CP, 7 March 2026. 4 Decree 141/2026/NĐ-CP, 29 April 2026, as published by the government portal: the new threshold, the backdating, the exempted count and the costed revenue loss are all official figures. 5 Chamber of commerce survey of 1,001 household businesses, fieldwork February to April 2026 across 34 provinces, published 23 April 2026 and read at source; the revenue, intention and closure figures and the quoted official are from it. 6 Hải Phòng registration, suspension and dissolution figures from the provincial newspaper, 4 August 2026, read at source.

So what is it that actually happened here, beyond a tax reform going badly?

The cashless miracle was bought on an unstated premise. VietQR on every counter; **ninety million mobile-banking accounts**; payment value in 2025 at **28.17 times GDP** — all adopted because it was free and frictionless.¹ Decree 70 and the death of thuế khoán told five million people what the rail was for, and within twelve months a measurable slice of the economy unplugged.

That is the first-order story, and it is already the interesting one. The second order is better, and nobody in Hanoi has published it.

The bill is not landing on the tax office. It is landing on the banks, and through them on everybody's cost of credit. What left was current-account money — the cheapest deposits a Vietnamese bank holds. When it migrates into cash, gold and term deposits, banks replace it with dearer funding and price loans accordingly.

The arithmetic, plainly. The sector paid **₫32.84 trillion** in 2025.² In April 2026 the ministry handed **₫16.65 trillion** of it back. Against a funding base that shed something like **₫360 trillion** in four months, the net is not obviously positive — and it is paid, at the margin, by the formal enterprises the reform meant to give a level field.

A state can spend the trust in its payment infrastructure exactly once. Vietnam has just published the price.

Third order, and the part that travels. Financial inclusion and fiscal legibility are *the same infrastructure*. The contradiction is not inferred — it is written down, in a dispatch instructing provinces to promote a payment technology *in order to* improve revenue collection (p22). Every state that built a QR rail on an inclusion story holds the same unexercised option: India's UPI, Indonesia's QRIS, Brazil's Pix. Vietnam exercised it first.

SOURCES 1 Active mobile-banking accounts and the ten-month QR growth figures from the national payment switch's own reporting, 13 December 2025; the 2025 ratio of cashless payment value to GDP from Vietnamese financial press, 15 May 2026. Both read at source. 2 The sector's 2025 tax take and its year-on-year increase as reported by Vietnamese press, 17 February 2026, read at source. 3 Vietnamese banking press, 12 August 2026, read at source, including the regulated demand-deposit rate and the quoted chief executive; the absence of any reference to tax in that article is this desk's own observation of it. 4 The record deposit total, the securities house's own caveat about gold, foreign-exchange and property transactions, and the electronic-filing and on-time percentages from Vietnamese press of May and July 2026 and the finance ministry's reply to citizen petitions of 28 July 2026. The filing percentages are official and are reported here as what they are — a measure of paperwork. 5 Conversion counts and the 84.4 per cent intention figure from the chamber of commerce survey of 23 April 2026 and Vietnamese business press of early 2026, the latter surfaced in search rather than read at source and flagged accordingly; Decree 141/2026 as on p22; the June 2026 National Assembly proposal to reinstate a fixed or assessed regime for the ₫500 million to ₫5 billion band from Vietnamese press of 20 June 2026, read at source. Reddit was unreachable; no Vietnamese social platform was read directly, and every community sentiment above is summarised in this desk's own words from named, dated local reporting.

THE SCEPTIC · THE BANKS BLAME INTEREST RATES, AND THEY WOULD KNOW

The strongest objection is that the current-account collapse has nothing to do with tax. Vietnamese banking press attributed it in August entirely to the regulated demand-deposit rate of **0.1 to 0.2 per cent** against far higher term rates; MB's chief executive **Phạm Như Anh** said customers are moving idle balances into fixed-term deposits for the yield. **The article does not mention tax at all.**³

Three more. System-wide deposits hit a record in July 2026. The securities house behind the 11.52 per cent figure notes that large-value cash sits in gold, foreign-exchange and property deals, where it always did. And compliance shows no revolt: **99.6 per cent** of filings electronic, **98.1 per cent** on time.⁴

That last lands least: compliance measures paperwork, and eighty per cent of the sector was exempted in April precisely so that its compliance would become trivial. But the interest-rate objection is not answerable from published data, and the chain from reform to funding cost remains an *inference*.

Which leaves the deepest cut, already conceded. Resolution 68 wanted two million enterprises by 2030, grown out of five million household businesses; about a thousand converted in the first two months of 2026, and the chamber found 84.4 per cent never will. So on 29 April the finance ministry spent **₫16.65 trillion** to make two and a half million of them invisible again — voluntarily, retroactively, to save the rest of the project.⁵

That is not a tax reform completing. It is a state discovering that the informal sector is not a transitional stage on the way to being a company, but a stable, chosen, rational form of economic life with its own technology — and when you point a ledger at it, it has somewhere to go.

WHAT WOULD FALSIFY THIS

The thesis dies if the cash ratio falls back below about ten per cent through 2027 while household-business receipts keep climbing. It is confirmed if the threshold is raised again, or the lump-sum tax reinstated for the middle band — a live proposal in June 2026.⁵

And one tell above all: **if the state ever moves to compel merchants to accept digital payment.** That is the government conceding the rail is no longer a convenience it offers but infrastructure it requires.

MERIDIAN

THE LAST WORD

The things you are asked to take on trust

Nothing was meant to connect this morning's desks, and something did anyway. Three governments announced an agreement over Greenland and none of them published it; every specific term in circulation comes from officials briefing one broadcaster. Singapore has had a law against transboundary haze since 2014 and has never once opened an investigation under it, the last explanation being that a reading of sixty-nine did not warrant one. Six days ago the number was a hundred and fifty-four.

A lamp can print a colour-rendering index of ninety and hide a deep-red score of nothing, because the headline number averages only eight pastel swatches. A comedy took fourteen Emmys on a critical score of seventy-eight, the lowest of the night. An American knife brand closed in March because somebody finally sent its blades for spectrometry, and the steel stamped on them was not the steel inside them. A paper on twelve-year-olds and first phones reports its significance levels and, not yet being typeset, no effect size at all — so this book printed the direction and refused the magnitude.

And in Vietnam, five million shopkeepers worked out what the free, instant, frictionless payment rail on their counter was actually for, and took the sticker off. The state priced its own retreat at sixteen and a half trillion dong and published the figure. Everywhere in this issue the same shape: the number you are shown, and the thing it was standing in for.

Twelve Desks · One World MERIDIAN · No. 109 · 21 September 2026

A note on method. This edition was researched with working network access, and the load-bearing figures were read on the pages that carry them: the Danish prime minister's own press release; the National Environment Agency's live air-quality endpoints, read at three o'clock on the morning of publication; the Federal Reserve's and the Bank of Japan's own decision documents; the American Bladesmith Society's own test rules; the Vietnamese government portal's own decrees and the finance ministry's own dispatch; the journals' own abstracts; and each maker's or dealer's own catalogue for every price printed. Hosts that refused this desk are named on the page where their absence matters, and nothing has been substituted for them. Reuters, the BBC, the Associated Press, the Guardian, the Straits Times, CNA, BladeForums and Reddit were all unreachable; no claim anywhere in this issue is attributed to any of them. Market levels are Friday's close except the bitcoin print, which is this morning's. The Macro Desk explains markets and does not advise on them; the Family and Fitness desks carry general guidance, not medical advice. Where a figure could not be verified at its source it is flagged in place or omitted — including a widely-reported Singapore export figure and an effect size that has not yet been published. *Twelve Desks · One World.*