

THE CURRENT BRIEFING OF

MERIDIAN

Twelve desks. One world.

SATURDAY 19 SEPTEMBER 2026 · THE WEEKEND EDITION



A MANIFESTO WRITTEN IN 2007 SAID ONLY THIS: DRAW WHAT IS IN FRONT OF YOU, AND BE THERE WHILE YOU DO IT. EIGHTEEN YEARS LATER IT READS LIKE A PROOF OF HUMANITY (P15)

THE RABBIT HOLE

Urban sketching, all the way down — and the blind test a watercolour teacher ran on her own readers, and lost

THE TRAVEL DESK

Taitung: the bridge is shut until 2027, the balloons never leave the ground, and it is still worth five days and a hire car

THE ATELIER

The pepper mill is a cutting tool wearing a decorative body — and two test kitchens both picked the sixty-dollar one

Twelve Desks · One World

Source: Art Toolkit — a “Hike and Draw” outing at Discovery Park, Seattle. Sketchers working on location at sunset, which is what the movement on p15 actually looks like. Honestly labelled; not a Singapore sketchwalk.

THIS ISSUE • THE WEEKEND EDITION • NO. 107

Contents

p3–4 The World — the Houthis take both shoulders of the Bab el-Mandeb, and Riyadh’s escape route runs through it
p5 Singapore — the haze index read worst at midnight and calmest when the air was dirtiest; Waymo picks the west
p6 The Atelier — pepper mills: the burr is the product, and the two-hundred-dollar one has the shortest warranty
p7 Technology — blocked at the lithography node, Huawei went round it with light
p8 The Home Lab — a twenty-seven-billion-parameter model in five point nine five gigabytes
p9 The Kit — half a terabyte in one address space, five weeks late, and not yet measured by anyone
p10 The Good Life — a turntable that never touches its own platter, beaten on its one published number
p11 Screen & Sound — the audience score arrived and agreed with the

critics, which helps nobody
p12–13 The Family Desk — it was the new school, not the birthday; and the programme that lost to a relaxation class
p14 The Macro Desk — Tokyo raised rates and named the cause. Then oil fell and yields rose anyway
p15–17 The Rabbit Hole — urban sketching: a 2007 manifesto that reads as proof of humanity, and a blind test its own teacher lost
p18 Fitness — twelve minutes a week, if they are the right twelve
p19–22 The Travel Desk — the world’s best airline posts a loss; and Taitung, where three postcards are false
p23–24 The Diary — sumo on sale at nine this morning; and twenty-six seats nobody has eaten in yet
p25–27 The Long Read — the children learned the word “rate”

FROM THE DESK

“For want of a naile the shoe is lost, for want of a shoe the horse is lost, for want of a horse the rider is lost.”

George Herbert, *Outlandish Proverbs*, 1640; the familiar longer version is Franklin’s, in *Poor Richard’s Almanack*, 1758

A whole issue about narrow parts: a strait, a burr, a sizing agent, twelve minutes, a published rate.

THE STRIP • AT FRIDAY’S CLOSE

US 10-year	5.01%	▲ +7bp
Gold (spot /oz)	\$4,377.80	▲ +0.84%
Bitcoin	\$81,274	▲ +6.50%
Brent crude (Nov)	~\$103	▼ -1.5%
USD/SGD	1.2757	▼ SGD softer
VWRA (All-World)	\$191.66	▼ -0.40%

Treasury yields on the par curve. Full caveats, including a missing 17 September close, on **p14**. Not advice.¹

THE SCOREBOARD • THIRTY-NINE CLAIMS CAME DUE. TWENTY-FOUR RIGHT, THIRTEEN WRONG, TWO PUSH

The turn is datable. Nos. 89 and 90 called it a coin-toss (*push*); from No. 92 the desk said hike, **nine times running, all right.** What changed our mind was Warsh at Jackson Hole; we should have moved sooner.

Two corrections we owe you regardless of grade. Nos. 104–106 called gold “record” territory at \$4,307–4,340. **Gold’s record is about \$5,589–5,608, set in late January** — we printed a drawdown as a record, three times. And No. 81’s “19-year high” of 5.33 per cent on the thirty-year was stale: it closed at 5.37 on 10 September. **Venice:** three right, two wrong. **Zambia and Garmin:** both wrong.

— **p14**

SOURCES ¹ Levels at Friday’s close, 18 Sep 2026; Treasury yields from the US Treasury par curve, the rest from exchange data with gold and the STI cross-checked. Two caveats carried rather than smoothed: neither VWRA nor the global-aggregate proxy has a 17 September close on any venue we could read, and Friday’s Brent is a range because intraday prints could not be reconciled to a settlement notice. Every Scoreboard verdict was re-established at source this week. The ledger is `state/predictions.json`.

THE MERIDIAN DISPATCH · THE GULF, DAY 202

The second gate closes, and the first one stops mattering quite so much

Last weekend this desk reported Saudi crude escaping the Strait of Hormuz by changing ships off Oman. The escape route ran through the Bab el-Mandeb — and over four days in September the Houthis took the whole Yemeni coastline that commands it. Two narrow gates, one hand on both. Oil fell anyway.

this story  

The islands went in four days. Zuqar and the port of Mocha on 10 September; Mayyun — the island the charts still call Perim — with the Greater and Lesser Hanish on the 11th.¹ The Houthis now hold the length of Yemen's Red Sea coast and both shoulders of the gateway beneath it.

The war it has started is no longer proxy. On Thursday a Houthi drone was intercepted over Taif and the debris killed a man on Saudi soil — the first death Riyadh has announced inside the kingdom this month; a second, aimed at Mecca, was brought down.² Saudi aircraft struck telecoms towers in Taiz, killing a civilian, and three children died elsewhere in the same exchange.² The Houthis claim strikes on Aramco's Yanbu facility and a downed Saudi F-15. Those are claims, unconfirmed, and printed here as claims.²

What makes this the week's story is the geometry. The East-West pipeline carries Saudi crude away from the Gulf and out to the Red Sea — which is to say out through the Bab el-Mandeb. Drones hit a pumping station on 10 September; the kingdom shut the whole 745-mile line the next day, and Aramco is restoring about half of its seven-million-barrel capacity within days, the rest in roughly six weeks.³ The bypass and the blockade are now the same piece of water.

TEHRAN AS RELAYED IN THE ISRAELI PRESS

"Now the Iranians don't only control Hormuz and the Strait of Hormuz, but they control the other most strategic chokepoint in the Middle East." — Mohammad Mokhber, adviser to the supreme leader.⁴ An academic reading, from Professor Alam Saleh of the Australian National University: it means Iran holds "the Strait of Hormuz and the Strait of Hormuz plus."⁴

SANA'A

HOUTHİ POLITBURO

"Freedom of navigation and international trade in the Red Sea and Bab el-Mandeb are safe and orderly." — Hazem al-Assad.⁴ Twelve per cent of world trade now passes a strait held by a movement making that promise.

And the counter-framing acquired evidence on Wednesday. Reuters reports, on three Iranian sources, that Beijing has privately pressed Tehran to rein the Houthis in — and that the request came to Beijing from Riyadh.⁴ China took more than eighty per cent of Iran's seaborne oil exports last year and draws about half its own imports from the region. Alex Vatanka of the Middle East Institute puts the limit plainly: "Iran can disrupt Hormuz. But at some point that leverage begins hurting the very power Tehran increasingly depends on."⁴ Tehran's answer, per those sources, was that peace "depend[s] on ending the US-Israeli war on Iran."

THE SCEPTIC · THE FORECASTERS ARE LOOKING STRAIGHT PAST IT

Every house that publishes a number thinks this ends. Commerzbank raised its year-end Brent forecast on 14 September — to US\$85, still eighteen dollars below Friday's print — on the view that "transit volumes will recover to around 8 million barrels per day by the end of the year."⁵ HSBC has \$90 then \$85, calling the market "neither fully closed nor fully open, but persistently impaired"; the US Energy Information Administration forecasts Brent near \$90 for the second half of this year and \$67 by late 2027.⁵ Andreas Krieg of King's College London adds the military version: the Houthis are "terrible" at governance and unlikely to hold a blockade for long.⁵ And the market agrees — December Brent closed Friday at \$98.77, below the front month (p14). The curve is pricing the disruption out.

SOURCES 1 Island seizures of 10–11 September, the twelve-per-cent trade share and the casualty and displacement figures: *The Times of Israel* and Al Jazeera, 11 September 2026; IOM via Reuters, 18 September. 2 The Taif death, the Mecca interception and the Taiz strikes: Al Jazeera, 17 September. Houthi claims collated by GlobalSecurity.org, 18 September — unconfirmed, and flagged in the text. 3 The pumping-station strike and Aramco's restoration timetable: Bloomberg, 16 September, corroborated by gCaptain, 17 September. 4 Mokhber and al-Assad: *The Times of Israel*, 11 September. Saleh and Krieg: Al Jazeera, 11 September. Vatanka, the Beijing approach and the oil-dependency figures: Reuters (Parisa Hafezi), 17 September 2026. 5 Commerzbank, HSBC and Bank of America forecasts as reported 14 September 2026, single-sourced on the Commerzbank quotation; US EIA *Short-Term Energy Outlook*, 9 September, read at eia.gov; December Brent from the exchange close, 18 September.

AROUND THE DESKS · NEW YORK, VILNIUS, GAZA CITY

The inspectors are voted out of existence, and the mandate runs out a week on Sunday

Three stories about the moment a system stops being watched: a sanctions panel that dies on 27 September, a drone carrying explosives shot down over a NATO member, and a city where the buildings have begun killing people on their own.

On Wednesday the Security Council failed to renew the panel of experts that monitors sanctions on Iran. Eleven members voted in favour. China and the Russian Federation voted against, which was enough; Pakistan and Somalia abstained.¹ The panel's mandate expires on **27 September**.

What dies on the 27th is therefore not surveillance but the *record* of surveillance — the thing that would have told anyone, later, what had moved and where. The deputy American ambassador, Jennifer Locetta, put it in the one sentence worth keeping: without an independent panel “this council loses its principal source of impartial, evidence-based reporting on sanctions violations.”¹

THE CHATTER · A SMALL CORRECTION TO OUR OWN ARITHMETIC

We flagged the vote as pending on Friday and gave the expiry as the 27th; at least one outlet runs the 26th. The date above is the Security Council Report's own September forecast, which is the document the Council works from.¹

BALTIC · THE DRONE THAT WAS CARRYING SOMETHING

An Italian NATO fighter shot down a drone over Lithuania, and the thing it was carrying was an explosive device. It came down near Pratkunai, about a hundred kilometres west of Vilnius, and had “most likely entered the country's airspace from Belarus,” according to Lithuania's national crisis management centre; President Gitanas Nausėda confirmed the shoot-down.²

Four days earlier, at sea. A Russian frigate fired two emergency flares at a Danish air force Fennec helicopter in international waters off southern Denmark. They passed within “a matter of metres,” said defence minister Jeppe Bruus, with no warning. Russia's ambassador in Copenhagen, Vladimir Barbin, called the helicopter's conduct “provocative”; Denmark summoned him.²

The framing from Brussels. Ursula von der Leyen called it “part of a broader pattern of Russian aggression and provocation against Europe testing our readiness and our

resolve.” Denmark's prime minister Mette Frederiksen: “Russian actions are meant to intimidate and divide. They will not succeed.”²

The escalation is worth naming precisely, because a year of incursion stories has blurred the category. Previous drones were unarmed and the response was tracking. This one was armed and the response was a missile. A different rung.

THE SCEPTIC · AGAINST OUR OWN PATTERN-MAKING

Two incidents four days apart in adjacent seas invite a single narrative, and the narrative may be doing work the facts do not support. A flare from a warship is an established if aggressive signal; a drone crossing from Belarus may reflect a war economy's chaotic periphery rather than a decision in Moscow. A correction, too: the Copenhagen airport closure in our continuity file belongs to **September 2025**, not this month, and we have re-dated it rather than let it thicken a pattern it was never part of.

GAZA · A YEAR AFTER THE CEASEFIRE, THE BUILDINGS

Twenty-one people were killed when a multistorey building collapsed in Gaza City on Wednesday, according to the Gaza health ministry. On Friday a second roof caved in, lightly injuring four children.³

The scale underneath it. Seventy-six per cent of residential units are damaged or destroyed; roughly 320,000 homes are damaged, affecting 1.7 million of Gaza's 2.1 million people; more than 2,000 further buildings are judged at risk of collapse, and in one Gaza City neighbourhood 22 per cent face imminent risk. More than a million people lack adequate shelter.³

Shaina Low of the Norwegian Refugee Council supplies the sentence that explains the death toll: “People are choosing to stay in dangerous buildings because the idea of living in a tent is worse.”³ Israel's COGAT says it regularly facilitates shelter supplies, citing the risk that construction material is repurposed by Hamas.³ A year after the shooting mostly stopped, the mechanism killing people in Gaza City is masonry.

SOURCES 1 The vote, the tally and the Locetta quotation: UN Meetings Coverage SC/16455, 17 September 2026; Al Jazeera, same date. The 27 September expiry, the blocked programmes of work and the vacant committee chair: Security Council Report, *September 2026 Monthly Forecast*. 2 The Lithuanian shoot-down, the Danish helicopter incident and the von der Leyen and Frederiksen quotations: *The Irish Times*, 15 September 2026, corroborated by *Foreign Policy*. The Copenhagen airport closure re-dated to September 2025 on the same check. 3 The two collapses, the casualties and the damage assessment: Associated Press, 18 September 2026; Low and COGAT quoted in the same report.

THE LITTLE RED DOT • SATURDAY MORNING

The air cleared overnight, and Waymo picked the west

Two things landed on Friday while the haze was quietly falling out of the Unhealthy band: Alphabet's robotaxi arm became Singapore's third licensed autonomous-vehicle operator, and it is starting at Labrador and HarbourFront. Commercial rides in 2028.

this story  

At four this morning the twenty-four-hour PSI read 79 in Central and 76 in the West — Moderate in all five bands, for the first time in a week.¹ Friday's had peaked at 102 at midnight and fallen all day.

Keep the shape of that in mind. On Friday the twenty-four-hour index read its worst at midnight, when the one-hour PM2.5 was around twenty micrograms — and sat inside Moderate at five in the afternoon, the precise hour the one-hour reading peaked at sixty-one.¹ The number that triggers purifiers in classrooms lags the air in your lungs by most of a day — not a failure of the instrument, which does what a rolling average does, but of how the rest of us read it.

The other Friday story will outlast the weather: the Land Transport Authority confirmed **Waymo** as the country's newest autonomous-vehicle operator, alongside ComfortDelGro and Grab.² Vehicles arrive within months, a manual-driving phase runs through 2027, and commercial ride-hailing opens in 2028, on all-electric Jaguar I-PACEs adapted to "local road geometry, monsoon weather patterns."²

The first operating domain is reported as Labrador and HarbourFront, with Tampines North and Tengah flagged next, against sixteen passenger vehicles running in Punggol today. The government is subsidising up to ninety per cent of training salaries for drivers moving into fleet management and remote operations.³

THE CHATTER • THE LOCAL FORUMS GOT THERE YEARS AGO

HardwareZone has been running this argument since June. **radish**: "not so soon. sg gov need to protect all the taxi, bus, phv drivers. a lot of ppl will go jobless." **fulaien**: "San fran already has it for a few years and much safer than human taxi." **cheesybites1**: "pap can delay but inevitable it will be here. this is the future." And the question an engineer asks, from **buffybuffy**: "Can this autonomous vehicle react to traffic marshaller and sinkhole?"³ On the haze index the same forum is blunter — a thread titled "NEA air quality index is misleading" has been running since August, which is the paragraph above, arrived at independently.¹

ALSO THIS WEEK

Twenty-seven teenagers, one street, five statutes. Police released the outcome of a 9 September operation near Petir Road in Bukit Panjang: twenty-seven young people aged **12 to 19** stopped, thirteen under investigation, seven non-compliant bicycles seized by LTA.⁴ The offences run from bicycles without working handbrakes to vaporiser possession (three of them) to nuisance calls made to the 999 line. The youngest person investigated is twelve.

THE SCEPTIC • NINETY-FOUR PER CENT OF WHAT, EXACTLY

Waymo's headline figure is a 94 per cent reduction in injury-causing crashes, drawn from 220.6 million rider-only miles.² The useful thing is that Waymo says the quiet part itself, on its own safety page: "There is no perfect 'apples-to-apples' comparison between human and AV data available today," and the human benchmark rests on police crash reports carrying "limited information about the specifics of each individual human crash" beside Waymo's own sensor-rich record.² Add the geography: those miles are surface streets in Phoenix, the Bay Area, Los Angeles and Austin. Not expressways, not a monsoon squall, and not a lane of motorcycles filtering at a Bukit Merah junction at six o'clock. The 2027 familiarisation year is not a formality; it is the whole experiment.

SOURCES 1 PSI and PM2.5 readings for 17–19 September pulled from the NEA / data.gov.sg real-time API, read 19 Sep 2026 at 04:30 SGT; NEA, "Haze Situation Update (18 September 2026)." The HardwareZone thread "NEA air quality index is misleading," opened 8 August 2026. 2 LTA news release, 18 September 2026; Waymo, "Singapore, Next Stop," 17 September, and its safety-impact page, including the 220.6 million-mile basis and the "no perfect apples-to-apples" caveat. 3 The first operating domain, Punggol fleet size and training subsidy: Mothership, 18 September 2026, single-sourced. 4 Singapore Police Force, 18 September 2026; the operation itself ran on 9 September.

BUY IT ONCE · PEPPER MILLS & SALT GRINDERS

The burr is the product. Everything else is a handle

A pepper mill is a cutting tool wearing a decorative body, and almost nobody sells you the cutting tool. Two test kitchens have now independently named a sixty-dollar mill best overall — and the two-hundred-dollar one carries the shortest warranty in the category.

this story  

A peppercorn is a dried drupe: hard, brittle, full of oil held in cells. You want to fracture it, not smear it — which makes a pepper mill a cutting tool, and the metal at its centre the only part that decides whether it works.

Peugeot's mechanism is the reference, and the order of operations is the lesson. The grooves and teeth are *cut into the steel first*, and the steel is **case-hardened afterwards** — carbon driven into the surface so the tooth faces are glass-hard while the core stays tough enough not to shatter.¹ A supermarket burr is pressed from soft sheet and never hardened: its teeth are folded metal, not cut edges. Salt needs ceramic instead, for a chemical reason and not a marketing one — dissolved in moisture it is an electrolyte, and it pits steel, stainless included.¹ Note what that means for “stainless”: ordinary 18/10 cannot be heat-hardened at all. The phrase that tells you something is “hardened high-carbon”.

THE CHATTER · WHO MAKES YOUR MILL, AND THE SINGAPORE CLAUSE

CrushGrind invented the ceramic mechanism in 1994 and names its customers — Tupperware, WMF — with over eighty million mills carrying it. **Holar**, Taipei, 1966, advertises “OEM, private label, logo printing” and names none — roughly where a mill comes from when a brand won't say where it is made.² And Peugeot is *not* a badge licensed from the carmaker: the mill company is family-held, and the business began in 1810 as a steel foundry.¹

The Singapore clause. Among Peugeot's warranty exclusions sits “**extended humid storage**” — here the specification, not fine print: buy acrylic, stainless or cast iron, not beech.¹ Cook N Serve carries twenty-six Peugeot mills at Peugeot's own prices; ToTT stocks none, shelving Cole & Mason and AdHoc — one company, both made in China.²

THE SCEPTIC · TWO TEST KITCHENS, AND NEITHER PICKS THE EXPENSIVE ONE

The US\$199.99 Pepper Cannon carries a five-year warranty; the US\$68.95 Peugeot u>Select carries an unlimited-duration guarantee on the mechanism. On buy-it-once logic the expensive one loses.³ America's Test Kitchen crowned the Unicorn Magnum in 1996 and 2007 — then *replaced* it after a 2013 test with the Cole & Mason Derwent, around US\$60. Food Network's August 2026 test reached the same Best Overall, relegating the Cannon to “Best Splurge.”³ Where the logo coasts: **Le Creuset** sells a plastic-bodied, China-made mill at US\$50 on a cast-iron reputation with no connection to mills.³

And the honest limit. The quality delta lives between about US\$15 and US\$50; above US\$70 you are buying throughput and presence. Nobody has run a blind taste test, because there is nothing to blind: piperine, the heat, survives months, and what you lose to pre-ground is the volatile top note — defensible on a two-hour braise, the whole point over a fried egg. **And skip the salt mill:** salt does not go stale, grinding Maldon destroys the flake you paid for, and CrushGrind's own FAQ says not to grind wet salt.²

NEXT ON THE BENCH

Coffee grinders — the same essay with real measurement behind it: burr alignment, particle distribution, and a maker landscape that makes mills look sleepy.

Kitchen scales & thermometers — the only kitchen tools with a specification you can actually verify.

Wooden cutting boards & butcher blocks — end-grain versus edge-grain, and what a tropical kitchen does to a four-hundred-dollar walnut block.

SOURCES 1 Mechanism manufacture, the cut-then-harden sequence and the grain limits from Peugeot's factory description via a dealer and from Peugeot Saveurs' warranty page, which also carries the “extended humid storage” exclusion; company history and the 1840 coffee mill from its history page, correcting the widely-repeated 1842 date. Ceramic-versus-steel from Cole & Mason and Zassenhaus product pages; MÄNNKITCHEN's “hardened high-carbon stainless” specification from its own page. 2 CrushGrind's founding, named customers and wet-salt advice from crushgrind.com; Holar's OEM offer from holar.com.tw, which names no clients and none is asserted here; DK Household Brands' ownership from coleandmason.com; Singapore prices read from Cook N Serve's and ToTT's own catalogues, 18 September 2026. 3 Warranties and prices from each maker's own page. America's Test Kitchen's 1996 and 2007 verdicts and the post-2013 switch via the archived Cook's Illustrated review; Food Network's six-mill test, updated 24 August 2026. Community sentiment is drawn from dated, attributable Hacker News and Virtual Weber threads; Reddit was unreachable this session and nothing is attributed to it.

SILICON & SOVEREIGNTY • SHANGHAI, THURSDAY

Blocked at the lithography node, Huawei went around it with light

At Huawei Connect on Thursday the company pulled its next AI accelerator forward by three quarters and launched a supercluster that replaces forty-eight thousand optical modules with five and a half thousand optical engines. The chip it cannot buy is still the chip it cannot buy. So it moved the bottleneck to the wires.

this story 

The interesting announcement was not the chip. It was the cabling. Huawei's Atlas 960E SuperPoD is billed as the industry's first supercluster built on near-packaged optics: **5,500 Hi-ONE optical interconnect units in place of 48,000 conventional 800G modules**, which the company says removes more than 550 kilowatts of power draw and lifts system availability to 99.8 per cent.¹

Read that against the constraint. Huawei cannot buy the leading lithography node, and no amount of ambition changes that. What it can do is attack every part of the system that is not the transistor: if each individual accelerator is behind, make the interconnect so good that more of them behave like one machine. Against the current Atlas 950 the company claims 2.3 times the training throughput, 2.5 times the inference throughput, and 70 per cent lower inference latency.¹

The roadmap moved too, which is the part the market will read. Deputy chairman David Wang brought the **Ascend 960DT forward to the first quarter of 2027 — three quarters earlier** than the schedule published a year ago — with the 960PR at Q3 2027, the 970 in 2028 and the 980 in 2029.¹ UnifiedBus, the protocol beneath all of it, is pitched as spanning clusters up to a million accelerators.

48,000 → 5,500

OPTICAL MODULES REPLACED BY OPTICAL ENGINES IN A SINGLE ATLAS 960E POD. THE SAVING HUAWEI CLAIMS IS 550 KILOWATTS — REMOVED FROM THE INTERCONNECT ALONE¹

THE CHATTER • NOBODY TURNED UP

The honest community read is that there isn't one. Two Hacker News submissions of the Atlas 960E news drew **seven points and no comments**, and **one point and no comments**.² The live discussion was one thread over, on how Zhipu built its inference stack, where the useful observation was about supply rather than silicon — user **HarHarVeryFunny**: the GLM team “use Huawei Ascend processors made by SMIC. Huawei use a combination of domestic memory from CXMT and leftover (pre-sanctions) memory from Samsung.”² That is the constraint arriving from the software side, unprompted. *Reddit was unreachable from this desk this week and nothing here is attributed to it.*

THE SCEPTIC • THE POD GOT SMALLER WHILE THE CHIP GOT EARLIER

The China analyst **Rui Ma** caught the trade in one line: “The chip itself is coming WAY earlier, but the SuperPod they announced is much smaller.”³ Huawei's earlier roadmap described an Atlas 960 scaling to 15,488 chips; the pod actually launched holds 4,096, and the million-accelerator configuration is described as theoretical rather than demonstrated.³ On raw silicon the gap is unglamorous: the 960DT's roughly 288GB and 4 petaflops at FP4 is about half an Nvidia B300 at FP8 and around a third of it at FP4.³ And then the part no announcement fixes. Huawei cannot make enough of them. It raised Ascend 950DT prices 60 per cent citing tight component supply; DeepSeek alone plans more than 160,000 chips, which is more than current capacity; and rotating chairman Eric Xu has publicly projected parity only with domestic demand, and only by 2030.³ An independent estimate puts Huawei near 1.5 million AI chips this year against Nvidia's 5.9 million — reported second-hand here, and flagged as such. A faster roadmap into a fixed packaging ceiling is a schedule, not a supply.

SOURCES 1 Huawei newsroom, David Wang's HUAWEI CONNECT 2026 opening keynote “Advancing the Agentic World, Building a Solid Silicon Foundation,” 17 September 2026, read in full at huawei.com — the pod specification, the near-packaged-optics claim, the Hi-ONE figures, the roadmap dates and the Atlas 950 comparisons are all from it. Corroborated by The Next Web, 18 Sep 2026, and Wccftech, 17 Sep 2026; the per-chip specification table is single-sourced to Wccftech and is not relied on above. 2 Hacker News items 49747675 and 49742250, engagement counts read via the site's search API, 18 Sep 2026; the HarHarVeryFunny comment from the “How GLM built its own inference infrastructure” thread, 17 Sep 2026. 3 Rui Ma quoted in The Next Web, 18 Sep 2026, which also carries the 15,488-versus-4,096 comparison, the “theoretical rather than proven” characterisation, the B300 comparisons, the 60 per cent price rise, the DeepSeek order and Eric Xu's 2030 projection. The Epoch AI capacity estimate reached this desk through a search summary rather than Epoch's own page and is flagged second-hand in the text.

THE HOME LAB · A TWENTY-SEVEN-BILLION-PARAMETER MODEL IN SIX GIGABYTES

Someone squeezed a flagship into a phone-sized file, and kept ninety-eight per cent of it

Ternary weights are an old idea that has mostly produced broken models. Thursday's release is the first this desk has seen where the benchmark gap against ordinary two-bit quantisation is large enough to be the story.

PrismML released Ternary-Bonsai-2-27B on Thursday under Apache 2.0: a ternary compression of Qwen3.8-27B, in which every weight is forced to minus one, zero or plus one. The smallest pack is **1.75 bits per weight and 5.95 gigabytes**, against 53.8 gigabytes for the sixteen-bit reference — a factor of nine and a bit.¹

Read the category breakdown rather than the average, because that is where the honesty is. Maths comes through nearly intact at 96.57 against 97.06, and coding at 89.42 actually edges above the full-precision baseline, which is noise rather than magic. The losses are concentrated and predictable: reasoning falls from 85.55 to 79.86, vision from 71.36 to 66.19, and document OCR from 60.99 to 56.88.¹ Ternary weights cost you the things that need fine gradations.

84.78 : 72.59

BONSAI 2 AGAINST A CONVENTIONAL TWO-BIT QUANTISATION OF THE SAME MODEL, ON THE MAKER'S OWN FOURTEEN-BENCHMARK SUITE. TWO SIMILARLY-SIZED FILES, FOURTEEN POINTS APART — THAT GAP IS THE ENTIRE CLAIM¹

THE CHATTER · READ OFF THE MODEL'S OWN DISCUSSION TAB

The repository took roughly 406,000 downloads inside a day, and the complaints arrived at the same speed.¹ **phil111**: “This breaks the model.” **DemetriusMichael** reports it getting “stuck on reasoning loops” on harder tasks. Two hardware findings are more useful than either: **hedomn** measured an RTX 5060 Ti at 37 tokens a second against an RTX 3060 at 34 — almost no scaling on mid-range cards — and **shenghedadi** reports AMD inference collapsing to **0.1 tokens a second** on the two-bit tensors.² A format this new is a software problem before it is a hardware one.

THE SCEPTIC · BENCHMARKS ARE NOT A HARNESS

Every figure above is the maker's own, run on the maker's own suite, and nobody has published independent agentic-harness numbers for this release. The one independent data point available concerns the previous generation: a documented run of July's Bonsai 1 scored **7.9 per cent on Terminal-Bench 2.0** on an eight-gigabyte laptop GPU — below a nine-billion-parameter model at ordinary four-bit on the same card — while the one-bit variant “could not finish a task at all, rambling past 14,000 tokens without emitting a stop.”² That reached this desk second-hand, through a write-up whose author discloses he has not run the new model either; the original thread was on a platform unreachable this week, and it is flagged rather than laundered.

The fair counterweight: Bonsai 1 retained 94.6 per cent and Bonsai 2 retains 98.2, so the method is improving fast enough that last quarter's failure is weak evidence about this quarter's file. The correct posture is the cheap one — it is a six-gigabyte download, and your own eval is the only number that settles it.

STILL ON THE BENCH

Home Assistant 2026.9 shipped on 2 September with the thing a Lutron-to-Home-Assistant household has been missing: a **Matter network map**. The Matter panel now draws the fabric, separates Thread devices from Wi-Fi, marks routers and border routers, and shows per-link signal strength by line colour and thickness — parity, at last, with the Zigbee and Z-Wave maps. The other quiet win is an Activity view that shows a full trigger chain down to millisecond timestamps with links to the automation and its trace, which turns “why did the hall light do that” from an evening into a minute.³

SOURCES 1 PrismML, Ternary-Bonsai-2-27B model card, read in full on Hugging Face, 18 September 2026 — pack sizes, bits per weight, the fourteen-benchmark thinking-mode averages, the IQ2_XXS comparison, the per-category scores, the RTX 5090 throughput figures and the download count are all from it or from the repository listing. All are the maker's own measurements. 2 Comments read first-hand on the model's Hugging Face discussions tab, 18 September 2026, quoted with the handles as posted. The Bonsai 1 Terminal-Bench figure comes from a dev.to write-up of 18 Sep 2026 relaying a thread this desk could not open, and whose author states he has not run the model himself — flagged second-hand in the text. 3 Home Assistant release notes, “2026.9: There's room on this bus,” 2 September 2026, read at home-assistant.io.

GEAR WE'RE WATCHING · ON SALE TUESDAY

Half a terabyte of memory in one address space — and nobody has measured it yet

Apple's new Mac Studio reaches Singapore on Tuesday from S\$3,499. The configuration that makes it interesting to a home lab — 512GB of unified memory — does not ship until late October, and no independent review of any version exists.

this story  

Apple Singapore's store page currently reads: available starting 22 September, limit of two per customer. The M5 Max version starts at S\$3,499, the M5 Ultra at S\$7,999.¹

The specification that matters here is not the core count. It is that the M5 Ultra addresses up to **512GB of unified memory at 1.2 terabytes a second** — roughly one and a half times the M3 Ultra's bandwidth, in a pool the GPU sees directly.¹ For anyone running models locally, memory capacity, not compute, has been the binding constraint for two years.

Thunderbolt 5 arrives at 120 Gb/s with RDMA clustering, and Apple claims four machines lashed together run inference three times faster than one.¹ Wi-Fi 7 and Bluetooth 6 appear on a Mac Studio for the first time.

Late October

WHEN THE 512GB CONFIGURATION ACTUALLY SHIPS — APPLE'S OWN STATEMENT. TUESDAY'S LAUNCH IS THE 128GB M5 MAX AND THE LOWER ULTRA BUILDS; THE MACHINE THIS DESK IS DESCRIBING IS FIVE WEEKS OUT¹

Now the part this desk has learned to lead with rather than bury. **There is no independent measured review of this computer.** None existed as of Friday; the configuration that justifies it is not in reviewers' hands until late October. The only measured M5 Max silicon in public comes from the sixteen-inch MacBook Pro, where Macworld recorded a Geekbench multi-core score of **29,233** — beating the thirty-two-core M3 Ultra Mac Studio's 27,726, on which their verdict was that “it would be hard to feel the difference.”² That is a fact about a chip in a laptop, not about this chassis.

No. 106 demoted a pair of headphones for exactly this reason. The rule holds when the product is one the desk would like to recommend.

SOURCES ¹ Apple Newsroom, “Apple introduces new Mac Studio with M5 Max and M5 Ultra,” 25 August 2026, read in full — memory ceilings, bandwidth, Thunderbolt 5 and RDMA clustering, the clustering claim and the late-October availability of the 512GB configuration are all Apple's own statements, and the performance multiples are Apple's with its own baselines. Singapore availability and the two-per-customer limit read on apple.com/sg; the S\$3,499 and S\$7,999 entry prices from HardwareZone Singapore, 25 August 2026. ² Macworld's M5 Max sixteen-inch MacBook Pro review for the Geekbench and Cinebench figures — a laptop, not this machine, and labelled as such in the text. The memory-pricing arithmetic and the configuration objection are relayed from MacRumors' own reporting of 25 August 2026; the underlying forum threads were unreachable from this desk and no quotation from them is printed as verified.

THE CHATTER · THE ARITHMETIC OWNERS ARE DOING

The forum objection is structural rather than grumbling, and it is checkable: since going from 96GB to 256GB costs US\$4,000, a 512GB machine lands, in the publication's own wording, “well above the \$10,000 mark.”² The sharper version is about where the value sits — the Ultra only separates itself from the Max at 256GB and above, and the 512GB option is restricted to the eighty-core GPU model, so the configurations that justify the machine sit far above its advertised starting price. *Those threads were unreachable from this desk this week; the reasoning is relayed from the publication that quoted them, and is not presented as verified quotation.*

THE SCEPTIC · YOU ARE BUYING CAPACITY, NOT SPEED — AND CAPACITY GOT CHEAPER ON THURSDAY

1.2 terabytes a second sounds enormous until you put it beside a single RTX 5090, which does roughly 1.8 to 32GB. The Mac Studio's trade is explicit: it buys you room, at a bandwidth penalty, for five figures. Which would be a clean trade, except that on Thursday the same week produced the opposite argument. A twenty-seven-billion-parameter model that keeps 98.2 per cent of its full-precision benchmark score now fits in **5.95 gigabytes (p8)**. If the frontier is learning to fit into the memory people already own, a machine whose entire proposition is memory capacity is buying insurance against a constraint that is actively loosening. Both things can be true — but only one of them costs eight thousand dollars, and only one of them can be tested tonight for the price of a download.

ALSO ON THE SHELF

Mac mini with M6 and M5 Pro — same Tuesday, from US\$899 and US\$1,699. The detail for a home lab is unglamorous: 2.5 gigabit Ethernet as standard with a 10-gigabit option, which is the difference between a NAS that feeds a model and one that starves it. *Figures from a secondary reading of Apple's release rather than the page itself; indicative.*¹

OBJECTS & ESCAPES · SHIPPING AT THE END OF THIS MONTH

Esoteric took the motor away from the record and left a magnetic field behind

The T-01 has no belt, no idler and no spindle motor under the platter — the platter is dragged round by induction, with nothing touching it. It is fifty thousand US dollars. A Technics at a tenth of that publishes a better wow-and-flutter figure.

this story  

Every turntable is an argument about one problem: the motor that turns the record is also a machine that vibrates, and the stylus cannot tell the difference between vibration and music. Belts, plinths, suspensions and heavy platters are all attempts to put distance between the two.

Esoteric's answer, unveiled in Tokyo on 21 August and shipping internationally at the end of this month, is to remove the contact entirely. A magnetic driver with alternating poles spins *beside* a soft-iron induction wheel set into the rim of a non-magnetic aluminium platter; the platter is pulled round by induced magnetic force **with no physical contact at all**.¹ The company spent more than five years getting from the idea to the patent.

The second trick is better. Magnetic levitation cuts the platter's *effective* mass from a real thirteen kilogrammes to roughly four — the flywheel inertia of a heavy platter without the bearing load that normally comes with it.¹ Speed is adjustable by twelve per cent in tenth-of-a-per-cent steps, and there is a ten-megahertz BNC clock input so the deck can be slaved to an external reference.

US\$50,000

WITH THE TA-9S ARM; US\$46,000 ARMLESS. THE GRANDIOSO T1, WHICH INTRODUCED THIS DRIVE, WAS US\$80,000 WITH AN ARM IN 2023 — SO THE TECHNOLOGY HAS COME DOWN ROUGHLY THIRTY THOUSAND DOLLARS WHILE KEEPING THE LEVITATION AND THE CLOCK INPUT¹

0.06% vs 0.025%

PUBLISHED WOW AND FLUTTER: THE US\$50,000 ESOTERIC AGAINST THE US\$4,999 TECHNICS. ON THE ONE NUMBER BOTH MAKERS CHOOSE TO PRINT, THE CHEAP ONE CLAIMS BETTER THAN TWICE AS GOOD³

THE CHATTER · THERE ISN'T ONE, AND THAT IS WORTH SAYING

This desk opened both enthusiast threads on the T-01 and found announcement posts with no replies: one AudioShark thread containing a single line from user **brodrickj** restating the press release, and an AV NIRVANA thread with none at all.² Rather than manufacture a community verdict, here is the trade reaction, attributed. **Michael Lavorgna** of Twittering Machines calls the drive “sci-fi” and the package “lust-worthy to me” — and then catches the incoherence nobody else did: the T-01 already carries what Esoteric bills as a “High-Quality ‘Master Sound Discrete Clock’” motor driver, *and* offers a socket for an external one. His question is the right one: “What’s better than exceptional?”²

THE SCEPTIC · ONE NUMBER, PUBLISHED BY BOTH MAKERS

Esoteric specifies the T-01's wow and flutter at **0.06 per cent or less**, W.R.M.S. Technics specifies the SL-1200G-S, which costs **US\$4,999.99**, at **0.025 per cent** W.R.M.S., with rumble at 78 dB.³ On the single measurement both companies choose to publish, the turntable at a tenth of the price claims better than twice as good.

The caveat is real and this desk will give it: the two figures are almost certainly not measured to the same standard or under the same conditions, and Esoteric's is a ceiling rather than a typical value. But that is precisely the point. A fifty-thousand-dollar object that declines to publish a number beating a five-thousand-dollar one, on the only axis the industry has agreed to measure, is asking to be taken on faith. The burden of explaining the gap sits with the maker, and so far nobody has made them carry it.

STILL ON THE SHELF

Schiit Bifrost 3 — US\$699, shipping around Halloween, and a hundred dollars less than the model it replaces.⁴

SOURCES 1 Esoteric (TEAC Corporation) news release and T-01 product page, esoteric.jp, read at source — the MagneDrive and Magne-Float descriptions, the effective-mass figure, speed adjustment, the clock input, dimensions and the end-September international shipping date are the maker's own. US pricing from the launch press release of 21 August 2026. The Grandioso T1's 2023 price as reported at the time. 2 AudioShark thread 27254 and AV NIRVANA thread 17261, both opened and read 18 September 2026 and both carrying no discussion. Michael Lavorgna, Twittering Machines, 21 August 2026. 3 Esoteric's own specification for the T-01; Technics' own specification for the SL-1200G-S, read at us.technics.com, including the price. The measurement-standard caveat is this desk's and is stated in the text. 4 Schiit Audio's own Bifrost 3 product page, including the quoted comparison with Yggdrasil MIB; iFi Audio's own iCAN Phantom 1.5 page. No independent measurements of either exist yet, and the Head-Fi discussion of the Bifrost was unreachable from this desk.

CULTURE DESK · YESTERDAY'S RELEASES

The audience score landed, and it agrees with the critics

Friday's edition reported *Resident Evil* opening at three different percentages with no audience number at all. That number now exists: ninety-two, against ninety-six from a hundred and sixty-nine critics. The gap this desk expected did not appear.

this story 

Read off Rotten Tomatoes on Friday: **Tomatometer 96 per cent from 169 critic reviews, Certified Fresh; Popcornmeter 92 per cent** from what the site describes only as a thousand-plus verified ratings.¹ Metacritic has it at 81 from 43 critics. Zach Cregger directs; the budget was US\$75 million.

Two cautions about those numbers, both of which matter more than the numbers. First, the critic count moved from 151 to 169 *within the same day* this desk was reading, so any percentage of this kind is only meaningful with the count and the hour attached.¹ Second, this page is being printed on Saturday morning. The opening weekend does not exist yet. What exists is **US\$8.8 million in Thursday previews** — a franchise record — and tracking in the forty-to-fifty-million range against the series' previous best of \$26.65 million, set by *Afterlife* in 2010. Those box-office figures reached this desk second-hand, through a paywall, and are printed as tracking rather than as result.¹

96 : 92

CRITICS AND AUDIENCE ON *RESIDENT EVIL*. THE FRANCHISE'S PREVIOUS BEST AUDIENCE SCORE WAS 67 PER CENT, FOR THE 2002 ORIGINAL. IT IS NOW RECORDED AS THE HIGHEST-RATED VIDEO-GAME ADAPTATION ON BOTH MAJOR AGGREGATORS¹

151 → 169

THE *RESIDENT EVIL* CRITIC COUNT, AS READ OFF THE SAME PAGE ON THE SAME DAY. ANY PERCENTAGE OF THIS KIND MEANS NOTHING WITHOUT THE COUNT AND THE HOUR ATTACHED¹

ALSO OUT

Shaun the Sheep: The Beast of Mossy Bottom — Aardman, distributed by GKIDS, eighty minutes. Tomatometer **98 per cent from 54 reviews** as read on Friday; the audience score does not exist yet, the site reporting fewer than fifty

verified ratings. A note for anyone checking: search results widely report this film at 100 per cent. The live page said 98.²

Beck, "Ride Lonesome" — released Friday on Capitol, self-produced. Metacritic **87 from 11 reviews**, which is its "universal acclaim" band; Any Decent Music has 8.0 from sixteen scores. Five out of five from both *The Guardian* and *Mojo*; four and a half from *NME*, whose Jordan Bassett calls it "a drop-dead gorgeous meditation on heartache and solitude" and "a record that should offer succour to those who need it."³

THE CHATTER · WHERE THE DISSENT ACTUALLY IS

On the Beck record it is not in the Metascore, which is eleven reviews deep and therefore fragile. It is in the spread underneath: **AllMusic and Slant both at three and a half out of five**, sitting under an 87 and two perfect scores.³ That is not a consensus, it is a small sample with a long tail, and eleven reviews is roughly a quarter of what *Resident Evil* gathered from critics in a day. No user or audience score for the album was available anywhere this desk could read, so none is reported.

THE SCEPTIC · AGAINST OUR OWN REFLEX

This desk has spent several issues on the critic-versus-audience gap as a diagnostic — *Practical Magic 2* at roughly 35 against 92 was the fingerprint. The temptation now is to treat the *absence* of a gap as evidence of quality. It is not. Ninety-six and ninety-two, four days after a premiere, mostly measures who turns up first: the people who buy a ticket on opening Thursday for a video-game adaptation are the people most disposed to like it, and Rotten Tomatoes' verified-rating pool is drawn from ticket purchasers. The number that will mean something is the one three weeks out, after the audience stops being self-selected.

And a second caution on the aggregate itself: eleven critics for a major album, against a hundred and sixty-nine for a horror film, tells you something about what gets reviewed rather than what is good. Print the count beside the score, always, or the score is doing work it cannot support.

SOURCES ¹ *Resident Evil* (2026) Tomatometer, Popcornmeter, critic count, release date, runtime, director and cast read off rottentomatoes.com on 18 September 2026; the 151-to-169 movement observed against a Wikipedia snapshot the same day. Metacritic score and budget from the film's Wikipedia entry, accessed 18 September. Thursday preview gross and opening-weekend tracking reported by *Deadline* and the *Hollywood Reporter* and reaching this desk through search summaries only — the publication was paywalled, and the figures are flagged as tracking in the text. ² *Shaun the Sheep: The Beast of Mossy Bottom* figures read off rottentomatoes.com, 18 September 2026, including the fewer-than-fifty-ratings notice. ³ *Ride Lonesome* aggregate scores and the individual publication ratings from the album's Wikipedia entry, accessed 18 September 2026, with the Metacritic figure independently corroborated against Metacritic's own release-date browse listing; the publisher's album page itself was unreachable. Jordan Bassett's review, *NME*, 16 September 2026, read at source.

RAISING THEM WELL • SECONDARY ONE

It wasn't the birthday. It was the new school — and it lasts two years

A South Australian reform accidentally ran the experiment everyone needed: two different age cohorts started secondary school in the same month. Twenty thousand students, seven years of census data, and the decline tracks the transition rather than the age.

this story  

In 2022 a South Australian education reform separated them. Two age cohorts — one at Year 7, one at Year 8 — started secondary school at the same moment. Researchers at the University of South Australia used the state's annual Wellbeing and Engagement Collection census to compare them: **20,910 students, 104,800 observations, seven waves from 2019 to 2025.**¹

In the first two years after transition, all eight measured domains moved adversely. The positively-worded ones fell between 0.18 and 0.44 of a point; the negatively-worded ones rose between 0.08 and 0.13. The two largest falls were **cognitive engagement, -0.44**, and **perseverance, -0.31**.¹ The younger and older cohorts declined similarly — which is the finding. It is the transition, not the year of birth.

Set that against how transition support is actually delivered here. Singapore's secondary orientation is front-loaded into the first week — one school's 2026 programme ran across three days in January.² Three days of welcome against an effect the data measures in years is not a scandal; it is a mismatch worth knowing about, because it tells a parent where the gap is and that the gap is theirs to cover.

DON'T BE ALARMIST • READ WHICH DOMAINS MOVED

A shift of 0.44 on a five-point scale is under half a point, on one domain, in a school census that is not a clinical instrument. And look at the pattern rather than the headline: the *school-effort* constructs — engagement and perseverance — fell three to five times as far as sadness and worry rose.¹ That is the signature of a motivation and belonging dip, not a mental-health emergency. The authors' practical advice follows from it and is refreshingly unfashionable: teach secondary-level *learning strategies*, rather than running more feelings sessions.

THE SCEPTIC • THE PAPER'S OWN BIGGEST PROBLEM, IN ITS OWN WORDS

The transition studied “occurred in 2022, only 2 years after the major disruptions caused by the pandemic in 2020... The pre-transition period was also affected by the pandemic and therefore may not provide an ideal reference.”¹ That is the authors flagging that their own baseline may be distorted, and it is the single largest reason to hold this finding loosely. Two more: the census runs once a year, so “short-term fluctuations in well-being immediately before or after the transition may not have been captured”; and it is all self-report, “which may be influenced by transient mood states or social desirability.” One state, mostly government schools.

What survives all of that is the comparison itself, because the two cohorts sat inside the same pandemic, the same census and the same state. Whatever distorts one distorts both — which is exactly why a natural experiment is worth more than a bigger sample.

AND A GAP WORTH KNOWING ABOUT

The national health scaffolding mostly stops before she gets there. Grow Well SG health plans reached around 114,000 Primary 1 to 3 pupils in 2025 and extend to roughly 80,000 Primary 4 and 5 pupils from January 2026, with K1 and K2 added; Primary 6 gets classroom resources only, from 2027.² The founding statement said the strategy “will focus on children up to 12 years old for a start.” A Secondary 1 student is standing at the far edge of it — in the year the evidence above says the support should be thickest.

SOURCES 1 Zhou M, Maher C, Brinkman S, Cools J, Dumuid D, “Well-being decline during adolescence: school transition as a predominant driver beyond age progression,” *Journal of Child Psychology and Psychiatry* 2026;67(9):1513–1524, doi 10.1111/jcpp.70154 — open access, full text via PubMed Central; sample, design, effect sizes and all quoted limitations are the paper's own. 2 The three-day January 2026 orientation is one school's published schedule, cited as illustration, not a national standard. Grow Well SG figures and the “up to 12 years old for a start” statement from Ministry of Health releases of January 2026 and January 2025; MOE's own site was unreachable, so that material is at one remove. **General guidance, not medical advice.**

THE TWEEN YEARS · WHAT ACTUALLY WORKED

The branded anxiety programme did not beat a relaxation class

Six hundred and seventy-six adolescents, ten sessions each, a proper active control — and both arms improved about the same. The authors' explanation for why is the most useful paragraph a parent will read this month.

A multicentre randomised trial across Spanish schools took 676 adolescents with below-threshold anxiety and depressive symptoms, mean age 13.8, and split them between the Unified Protocol for Adolescents — ten fifty-minute small-group sessions of structured transdiagnostic therapy — and ten sessions of protocolised progressive relaxation.¹

Both arms improved significantly at post-intervention, six months and twelve months. On the primary outcomes there was **no significant difference between them**. The branded protocol came out ahead only on hyperactivity, conduct problems and total difficulties, plus marginally on depressive symptoms at twelve months.¹

The crucial design point is that the comparison was an *active* control, not a waiting list. Most school-programme trials compare doing something with doing nothing, and then report that something works. This one compared a specific something with a generic something, and the specificity did not earn its keep.

No difference

BETWEEN TEN SESSIONS OF A BRANDED TRANSDIAGNOSTIC PROTOCOL AND TEN SESSIONS OF PROGRESSIVE RELAXATION, ON THE PRIMARY OUTCOMES, AT POST-INTERVENTION, SIX MONTHS AND TWELVE¹

DON'T BE ALARMIST · THIS IS A GOOD-NEWS FINDING WEARING A NULL RESULT

The authors attribute the parallel improvement to “common therapeutic factors — such as structured adult attention, emotional validation, the experience of being heard, and supportive group interactions — rather than from protocol-specific techniques alone.”¹ For a parent deciding whether the school's wellbeing programme is worth a timetable slot, the answer is probably yes — and the brand on the folder matters less than whether the same teacher turns up for ten weeks running. That is a cheaper and more replicable finding than a proprietary protocol would have been.

THE SCEPTIC · AND IT MAY SIMPLY BE UNDERPOWERED

The authors say it themselves: “the study may be underpowered, as the power calculation assumed a large effect size ($d = 0.80$), whereas school-based psychosocial interventions typically show small to moderate effects.”¹ A trial built to detect a large difference will miss a real small one. Add what they call the standout problem — **dropout exceeded thirty per cent in both arms** — plus no blinded assessment and reliance on self-report, and the honest reading is “no difference detected,” not “no difference exists.” Worth holding, because the null is the more comfortable conclusion and comfort is not evidence.

AND ON EXERCISE, A NUMBER TO DISTRUST

The believable numbers are the small ones, and one of them is genuinely useful: the effect was larger in **high-autonomy settings, -0.62** — when the young person had a say in what they were doing. Group-based settings scored -0.30 , aerobic -0.19 . The adolescent subgroup estimate of -2.97 is roughly three standard deviations and is not a credible real-world effect; the authors call their subgroup findings exploratory, and this desk takes them at their word.²

SOURCES 1 Fonseca-Pedrero E, Pérez-Esteban A, Díez-Gómez A, et al., “A multicenter randomized controlled effectiveness trial of transdiagnostic treatment for adolescents with subthreshold anxiety and depressive symptoms,” *Psychological Medicine* 2026;56:e264, published online 4 September 2026, doi 10.1017/S0033291726105443 — open access, full text read at source; ClinicalTrials.gov NCT05322642. Sample, arms, outcomes, adherence and all quoted limitations are the paper's own. The abstract gives an age range of 12–18 while the results text says 11–18; the discrepancy is noted and neither figure is relied on above. 2 “Acute exercise for anxiety reduction in children and adolescents: a meta-analysis of developmental, contextual, and modality-specific effects,” *European Journal of Pediatrics*, published online 2 September 2026, doi 10.1007/s00431-026-07371-x (confirmed via Crossref — note the PubMed record carries an incorrect DOI). The publisher's page redirects to an authentication gateway, so figures are from the National Library of Medicine abstract record. **General guidance, not medical advice.**

THE BIG PICTURE · TWO CENTRAL BANKS, ONE INPUT

Tokyo raised rates, and the governor said out loud what the week was about

The Bank of Japan went to 1.25 per cent, its highest since 1995, two days after the Fed moved. Asked why so many central banks were tightening at once, Kazuo Ueda answered: the Middle East. Then oil fell on Friday and yields rose anyway.

this story  

US 10-YEAR	BRENT, NOV	GOLD / OZ	BITCOIN	USD/JPY
5.01%	~\$103	\$4,378	\$81.3k	156.9

The Bank of Japan raised its policy rate to around 1.25 per cent on Friday morning, a quarter-point rise and the highest Japanese policy rate since 1995.¹ A detail almost nobody printed: neither the new guideline nor the deposit and loan rates take effect until **24 September**.

THE CHATTER · THE SECOND SPLIT, WHICH THE WIRES MOSTLY MISSED


Seven to two on the rate — Asada and Sato against, both citing inflation running below two per cent — is the headline. Read the attachment and there is a second seven-to-two pointing the other way: Takata and Tamura opposed the statement's language on the price outlook, on the ground that underlying inflation has *already* reached the target.¹ Four of nine members dissented from something. A board that divided is not a board with a confident next move.

“The central banks all seem to be responding to similar factors. For one, it’s inflationary pressure from the Middle East conflict.” — Kazuo Ueda, 18 September

Last week yields rose because oil rose. On Friday Brent fell a third straight session and the whole US curve rose anyway, the two-year leading at nine basis points. The tightening is now in the system independently of the shock that caused it.

THE WEEK THE DRIVER INVERTED

Brent front month (US\$/bbl, left) against the US 10-year par yield (%), 14–18 September 2026



A Meridian figure. Brent settled lower for a third session on Friday while the whole US curve rose. Brent from exchange data; yields from the US Treasury par curve.

THE SCEPTIC · AND A CORRECTION THIS DESK OWES THE READER

The correction first. Nos. 104, 105 and 106 described gold at \$4,340, \$4,338 and \$4,307 as “record” territory. That was wrong, and not marginally: gold’s all-time high is about **\$5,589–5,608, set in late January 2026.**² Friday’s \$4,377.80 is roughly twenty-two per cent below the record — a one-week high and the first weekly gain in four. We printed a drawdown as a record three times.

Then the scepticism. Yimin Xu of YX Insights notes that rate rises “do not directly affect any single line item” in an energy-driven inflation, so with unemployment at 4.1 per cent the Committee is accepting job losses to fight a shock it cannot reach: “the Fed thinks that even if it is making a policy mistake, it can afford to do so now.”² And against reading Tuesday as an ending — Chair Warsh called the move the removal of “a dose of accommodation,” which puts policy at 3.75–4.00 still on the stimulative side of neutral.²

HIGHER-FOR-LONGER

SYNCHRONISED TIGHTENING

CORRELATION HURTING

SOURCES 1 Bank of Japan, “Change in the Guideline for Money Market Operations,” 18 September 2026 — the decision document itself; the vote, the dissents, the quoted language and the 24 September effective date are from it. Market levels from exchange data, 18 September, with Treasury yields from the US Treasury daily par curve and gold cross-checked against TradingEconomics. Two caveats printed rather than smoothed: **VWRA and the global-aggregate proxy have no 17 September close** on any of three venues, and Friday’s Brent is a range because prints between \$102.41 and \$103.21 could not be reconciled to a settlement notice. No. 106 gave the 10-year as 4.951% on a different convention; JGB yields are held, the ministry’s page being unreachable. 2 Gold’s record independently verified against two sources before this correction was printed. Yimin Xu, YX Insights, 17 September 2026, single-sourced and attributed. Warsh’s “dose of accommodation” from the 16 September press conference as widely reported; the FOMC statement and the 12–0 vote read at [federalreserve.gov](https://www.federalreserve.gov). Ueda’s press-conference remarks from the Reuters quote file.

DOWN THE RABBIT HOLE • URBAN SKETCHING & WATERCOLOUR

The proof is that you were there

A movement that began as a Flickr group in 2007 has passed five hundred and ninety-nine chapters. Its founding document is eight sentences long, and the first of them — draw from direct observation, on location — has turned out to be a rule about authenticity that nobody in 2007 knew they were writing.

this story  

In November 2007 Gabriel Campanario, a Spanish-born journalist and illustrator in Seattle, started a Flickr group for people who drew the cities they lived in and visited — “always on location, not from photos or memory.”¹ A year later he invited correspondents to a blog under the line “See the World, One Drawing at a Time.” Urban Sketchers became a Washington State nonprofit in 2009.

What makes it a movement rather than a hashtag is the manifesto, and it is worth reading as a constitution because that is how the community treats it. Eight points, in full: *We draw on location, indoors or out, capturing what we see from direct observation. Our drawings tell the story of our surroundings, the places we live, and where we travel. Our drawings are a record of time and place. We are truthful to the scenes we witness. We use any kind of media and cherish our individual styles. We support each other and draw together. We share our drawings online. We show the world, one drawing at a time.*¹

Note what points one, three and four do together. They locate the truth of the drawing not in how accurately it resembles the thing, but in the fact that a person sat in front of the thing for forty minutes. It is a claim about *presence*, and it is the opposite of the photographic contract, where the guarantee runs the other way: the camera was pointed at something real, and the operator need not have looked at it at all.

599

WORLDWIDE CHAPTERS, AS THE ORGANISATION’S OWN PAGE PUT IT THIS WEEK. IT ANNOUNCED ITS FIVE HUNDREDTH IN MAY 2025 — ROUGHLY A HUNDRED NEW CHAPTERS IN SIXTEEN MONTHS, EACH VOLUNTEER-RUN AND FREE TO WALK INTO¹

That definitional clause is now doing work nobody designed it for. In April this year Hasselblad unveiled seventy Masters finalists; within a day a Street-

category frame was flagged as generative AI — a Coca-Cola bottle with impossible typography — and by 18 May the entry was disqualified and the company was reaffirming “the authenticity of photography.”² A sketchbook cannot be prompted. The manifesto is, functionally, this hobby’s proof-of-humanity protocol, written eighteen years before it was needed.

Singapore’s own best-known sketcher, Teoh Yi Chie, presses the purist case further and his reasons are explicitly anti-camera: “It’s not urban sketching if you’ve sketched an urban or rural scene from a photo because you were not on location.” Among his eleven differences: a camera distorts perspective where the eye does not, and binding yourself to a photograph forfeits the compositional choice that being there hands you for free.²

THE CHATTER • WHERE PEOPLE ACTUALLY FALL IN

Not the gear. The walk. A Vienna sketcher who joined Urban Sketchers Singapore at the Sentosa Boardwalk found “over 50 sketchers” against her home chapter’s usual ten to twenty, and stayed three hours in thirty-three degrees and humidity above ninety per cent.³ The counter-route is Liz Steel’s, and it is deliberately unequipped: “simply get a pen or pencil, a small notebook (with plain unlined paper) and go outside,” then do three sketches before judging anything.³ Both are the same advice. Start before you shop.

THE SCEPTIC • THE DENOMINATOR

Point seven of the manifesto — “We share our drawings online” — is the mechanism that built the movement and the mechanism that makes it discouraging. There are eleven million reels under #urbansketchers and 7.8 million under #urbansketching.¹ That is the denominator against which a beginner grades their third drawing, and it is the single best argument for doing the first ninety days privately (p17).

SOURCES 1 Urban Sketchers, “Who We Are” — the manifesto quoted verbatim, the founding account and the 599-chapter figure, read 18 September 2026; the 500-chapter milestone announced 24 May 2025. Instagram tag counts read from the tag pages the same day (the platform counts reels, not posts). 2 The Hasselblad Masters 2026 disqualification: finalists unveiled 28 April, the entry flagged the next day and disqualified 18 May 2026, per PetaPixel. Teoh Yi Chie, parkablogs.com. 3 The Sentosa Boardwalk account from whereismrwong.com, posted 2019 describing a 2018 walk — historical, and flagged. Liz Steel, lizsteel.com.

THE DEEP END · THE GEAR, AND THE ARGUMENT UNDERNEATH IT

You don't paint on the paper. You paint with it

The most expensive thing in a watercolour kit is usually the paint. The thing that decides what the paint can do is a treatment a few microns thick on the surface of the paper, and almost nobody selling you paint will mention it.

Start with what separates artist grade from student grade, because it is four specific things and not a vibe. **Pigment load:** artist grade carries more pigment and less filler, so a brush-load delivers more colour and you fight it less. **Single-pigment versus convenience mixes:** single pigments mix cleanly and predictably; convenience mixes go muddy the moment a third colour arrives. **Lightfastness,** published per colour. And the **binder** — gum arabic plus a humectant, which is where honey changes how a dried pan rewets, and which matters enormously in the tropics (p17).¹

The only honest label is the colour-index code on the tube. **PB29** is ultramarine: granulating, non-staining, warm. **PY150** is nickel azo yellow: transparent, staining, huge tinting strength. Two tubes called “Hooker’s Green” are almost certainly different mixes. Learn to read codes and you can spot a downgrade: in May one maker’s professional line was found to have quietly changed Viridian from PG18 to a PW6 mix and Cerulean from PB35 to PW6 and phthalo.¹ PW6 is white filler. That is what a downgrade looks like on paper.

THE PALETTE, WHICH IS THE OBJECT PEOPLE COVET

Art Toolkit (Seattle) is the default cult object because it solved a real problem: a flat anodised case with a magnetic tray taking loose pans you fill yourself. Demi US\$30, Pocket US\$45, Folio US\$55 — or US\$36 empty, holding up to sixty mini pans.² **Portable Painter** is the sleeper at US\$29.95, because it has integrated water wells and nobody else does.

And the grail, honestly reported. Craig Young’s Paint Box Company made brass boxes by hand in Ludlow, each numbered, hinge pins included — “around 15 boxes a month,” with owners recording five-to-twelve-month waits in 2008 and later listings citing eighteen months.² There is now no live website, and multiple resale listings state he has retired. *That last point rests on sellers rather than on the maker, and is flagged rather than asserted* — but the practical position is not in doubt: the grail of this hobby is a secondary-market object you cannot order.

THE CHATTER · THE PAPER ARGUMENT, WHICH IS THE REAL ONE

Sizing is the whole game. Arches’ own explanation: “The aim of sizing is to make the paper more resistant to liquid absorption so as to limit the penetration of water and inks” — done internally, where the agent coats fibres inside the pulp, and at the surface, applied to the finished sheet.³ Sizing is why you can lift and scrub on Arches and cannot on cheap paper. Cold press is the forgiving default; hot press is for pen; 300gsm is the floor for washes in a bound book. The disruption is **Baohong**, whose Academy grade is 100 per cent cotton at US\$14.95 a block. Singapore’s Teoh reviewed the artist grade at S\$29.50 against US\$36 for Arches, finding the layering “very similar to Arches and Fabriano.” But the story has a second half: the Swedish reviewer Erik Lundgren notes that Baohong now prices *above* comparable Arches at some European retailers, and that the cheaper Academy grade “fluff[s] up when you paint wet on wet” — a fault he does not find in Master.³ The Arches-killer has become an Arches competitor.

THE CUTTING EDGE · TWELVE MONTHS, DATED

The tell that this hobby has arrived commercially: in March, **Winsor & Newton and Shinhan both released six-tube themed sets, and both named one “Urban Sketching.”**⁴ Two majors putting the hobby on the box in one quarter.

The losses are the more interesting news. **Winsor & Newton discontinued genuine Viridian, PG18, at the end of 2025;** Blockx dropped Cobalt Violet and Cobalt Blue Dark; Schmincke reformulated Indian Red after its pigment was discontinued.⁴ Meanwhile Daniel Smith marked fifty years by shipping an inventory registry filtering by granulation, staining and single-pigment status — which tells you the bottleneck is now cataloguing what you own.

And the genuine price shock is tariffs. Kremer Pigments: “We expect prices to rise or severely limit the colors we can offer artists.”⁴ A Toronto retailer’s banner, read on Friday: “Due to recent changes in the de minimis rules, we’re currently pausing all shipments to the United States.” A cross-border art-supply market breaking in real time.

SOURCES 1 Artist- versus student-grade criteria and the maker’s own concession from Greenleaf & Blueberry’s explainer; the Phoenix reformulations from Watercolor Scientist, 23 May 2026; paint prices read at rileystreet.com, gwartzmans.com and the makers’ own shops, 18 September 2026. 2 Art Toolkit and Portable Painter prices read at source; Craig Young’s production rate from The Watercolour Log, 10 March 2011, and wait times from WetCanvas owner threads of 2008–11 — the retirement is reported by resale listings only, and flagged. 3 Sizing quoted from Arches’ own explainer; Teoh Yi Chie on Baohong artist grade, 11 December 2021; Erik Lundgren on Baohong Master, 22 April 2025; paper prices as read at artdiscount.co.uk and rileystreet.com. 4 New releases, discontinuations and reformulations from Watercolor Scientist’s monthly reports of January, March and May 2026; tariff quotations from Hyperallergic, 4 March 2025; the Toronto retailer’s shipping notice read 18 September 2026.

THE DEEP END · THE CLIMATE, THE MONEY, AND THE NINETY DAYS

In Singapore, wet-in-wet does not work — and that is a technique lesson, not a complaint

The tropical version of this hobby is a different craft: carry less water, commit to fewer washes, lean on the pen. Which is exactly why line-and-wash dominates here, and why the chapter meets at half past nine and stops at half past twelve.

Gail Wong, writing from the 2015 Symposium in Singapore, recorded the failure precisely: “Wet on wet did not work well for me and my palette was a muddy mess with my paints oozing out of the wells.” Everyone struggled on the first day, and the local fix she was given was one sentence — “you have to adjust how you approach painting in the humidity, a little less wetness to control the paints.”¹

Three consequences. Honey-bound paints stay permanently tacky here, so hard-poured pans travel better, and a closed box that never dries is a mould risk. Drying times invert: a wash can sit wet for ten minutes then bloom when you touch it, so use fewer, more decisive washes. And anything under 270gsm buckles.

THE GLOSSARY, FOR READING ANY OF THIS

Sizing — the size that stops paper drinking like blotting paper. The narrow part of the whole craft.

Fude — a bent fountain-pen nib that runs hairline to broad on angle alone.

THE FIRST NINETY DAYS — ABOUT US\$250, AND MOST OF IT OPTIONAL

Practise, in this order. Weeks one to three: pen only, no colour, no eraser, three sketches an outing, judged as a set. Weeks four to six: value studies, three greys, before any colour. Weeks seven to nine: three pigments at a time, learning which granulate, which stain, which lift. Weeks ten to twelve: **go to a sketchwalk**. That last step, not the gear, is what determines whether anyone keeps going.²

THE SCEPTIC · A TEACHER RAN THE BLIND TEST, AND LOST

In September 2018 Liz Steel — who sells watercolour courses for a living — published four of her own sketches, some in student-grade Cotman and some in Winsor & Newton Professional, and asked readers which was which. **Two got it right.**³ Her conclusion concedes and then salvages: good results are possible with student grade, but “it requires a lot of work and care when mixing.” Read honestly: the audience could not see the difference; the painter could feel it. That is the truest sentence available about this entire gear ladder.

The paper threshold is real and low. Brian T Meyer, on a long thread about the cost of paper: “You don’t paint on the paper, you paint **WITH** the paper” — and immediately after, Calvin_0: “as long as the paper is 100% cotton, it’s a good paper.”³ Both are right, and Baohong Academy is cotton at US\$14.95.

Even the top of the market says so: Greenleaf & Blueberry, who will sell you a US\$945 set, write that if your work “is not meant to be sold or preserved” you “really don’t need to invest a lot of time or money into your supplies.”³ Gear acquisition syndrome is a term this hobby borrowed from photography. Anyone with a camera shelf will recognise it first.

NEXT DESCENTS

Leathercraft — a cheap entry, a bottomless tool ladder, and a genuine atelier tier with year-long waits. What this desk did for palettes, it does for stitching ponies.

Knife making / bladesmithing — heat treatment is the whole craft; everything before it is shaping. With an honest Singapore constraint: no forge in an HDB flat.

Hand-tool woodworking — buy-once-cry-once at its purest, and the vintage-Stanley-versus-new-premium argument that never ends.

SOURCES 1 Gail Wong, writing on urbansketchers.org from the 2015 International Symposium in Singapore, 26 August 2015. The humidity consequences that follow are this desk’s reasoning from the binder chemistry set out on **p16**, not a sourced quotation. 2 Prices as sourced on **p16**; the practice sequence follows Liz Steel’s published beginner method, and the common first-purchase regrets are drawn from Steel, from drawingwithpri.art (29 July 2023) and from urbansketchingworld.com’s beginner-mistakes list. 3 Steel’s blind test, lizsteel.com, 24 September 2018, including the count of correct answers and her own conclusion. Meyer and Calvin_0 quoted from the WetCanvas thread “Discouraged by the cost of watercolor paper,” 7–9 November 2019, with handles and dates as posted. Greenleaf & Blueberry’s own explainer.

THE LONG GAME · INTENSITY, MEASURED RATHER THAN REMEMBERED

Twelve minutes a week, if they are the right twelve minutes

A pooled analysis of 147,414 people wearing accelerometers puts about twelve minutes of vigorous activity alongside a hundred and fifty minutes of moderate. Then, in the same month, the same field's own cardiology body published a statement saying mortality is the wrong thing to be measuring.

this story  

Published on Monday: the largest pooled accelerometer dataset yet assembled, joining the German National Cohort to UK Biobank — **147,414 adults aged 40 to 74**, seven days of device-measured movement each, a mean 7.9 years of follow-up and 5,545 deaths.¹

Against no activity, a hundred and fifty minutes a week of moderate activity carried a hazard ratio of **0.70**. The same association appeared at roughly **510 minutes of light** activity — or about **twelve minutes of vigorous**.¹

THE CHATTER · THE LONGEVITY CROWD IS BEING TALKED DOWN OFF A NUMBER

The discourse this autumn is about VO_2 max, and the mood is relief. Eric Topol's February piece argues that the studies carrying the longevity claim mostly *did not measure VO_2 max at all* — one 2024 meta-analysis found 234 times more participants assessed by exercise capacity than by actual VO_2 max, or over 99 per cent of the data. He puts smartwatch estimates at “7-16% mean absolute percentage error,” consistently underestimating in fit people and overestimating in unfit ones.² Steve Magness is blunter: “The majority of the studies cited did NOT even use Vo2max as the main variable. They used performance!”² In Topol's own comments, **Naomi B Bishop, MD**: “Although it is not quite the same, it reminds me of the 10,000-step myth, which also lacked a firm scientific basis.” A retired statistician and seventy-year-old runner, **Chris Stidley**: “Thanks for providing a clearer way to view and not view my fitness.”²

A PROTOCOL, IF YOU WANT ONE

Take the paper at its own word, not at the headline's. What it actually models is substitution: minutes moved from one intensity band into another, at the same total. On that arithmetic the exchange rate is roughly **forty to one** — 510 minutes of light against twelve of vigorous for the same association — and it is the only number here that survives the objections on this page. The usable version is two hard

quarter-hours a week that you already have, rather than a seventh day you do not, and no test to tell you whether they worked.¹

THE SCEPTIC · THE FIELD SAYS THE OUTCOME IS WRONG, AND ONE AUTHOR SIGNED BOTH

In the same month, the European Association of Preventive Cardiology published a scientific statement whose opening is a demolition: “observational data suggest that regular physical activity lowers all-cause mortality. **However, this was not confirmed by Mendelian randomization studies and randomized controlled trials.**”³ Their diagnosis covers recall and recruitment bias, Hawthorne effects, reverse causation, healthy-volunteer selection and short trial follow-up — and the awkward modern point that medicine now keeps people alive despite cardiovascular disease, compressing any mortality signal exercise could show. Their proposal is to stop measuring lifespan and measure healthspan: function, cognition, quality of life.

Hansjörg Baurecht is an author on both documents.

The same research community that produced this month's large mortality number is on record, this month, saying mortality is the wrong outcome. Two more things the lead paper concedes about itself, in its own words: it “precludes causal inference,” and accelerometers cannot distinguish “between resistance and aerobic exercise.” And the recruitment: **NAKO's baseline response was 15.6 per cent and UK Biobank's was 5.5 per cent.**¹ These are not populations; they are the fifth and the twentieth of people who said yes. The two halves also wore different devices in different places on the body, at different intensity thresholds.

The honest synthesis is narrow and still useful: the direction is consistent across designs, the absolute numbers are not to be trusted, and the one finding robust to all of it is that intensity does more per minute than duration. Which is a reason to go harder for a quarter of an hour, and not a reason to buy a lab test.

SOURCES 1 Stein MJ, Lendt C, Bamberg F, et al., “Intensity-specific physical activity and mortality in 147,414 adults,” *IJBNPA*, 15 September 2026, doi 10.1186/s12966-026-01980-3 — full text via PubMed Central. Cohorts are the German National Cohort and UK Biobank; hazard ratios, the reallocation model, the response proportions and quoted limitations are the paper's own. This is a different study from the similarly-sized strength-training analysis in No. 104. 2 Eric Topol, “The Flawed VO_2 Max Craze,” 23 Feb 2026, and comments from the same page; Steve Magness, 12 Nov 2025. 3 Baahs M, et al., EAPC Scientific Statement, *Eur J Prev Cardiol*, Sep 2026, doi 10.1093/eurjpc/zwaf578 — abstract via the NLM; the publisher's page was unreachable.

DEPARTURES · WHAT CHANGED THIS FORTNIGHT

The trophy and the loss arrived in the same quarter

Singapore Airlines was named the world's best airline again on Friday, by 24.8 million votes. It also posted a S\$76 million net loss on record revenue, because its fuel bill rose seventy-eight per cent. The two facts are the same fact.

this story  

At the Skytrax ceremony in London on Friday, Singapore Airlines took World's Best Airline for the sixth time, along with Best Airline in Asia and the two economy-class awards; Scoot took Best Long Haul Low-Cost Airline for the sixth year running. The survey drew on more than 24.8 million travellers rating over three hundred airlines between September 2025 and August 2026.¹

Set that against the accounts. For the quarter to 30 June the group posted **record revenue of S\$5,714 million, up 19.3 per cent**, with passenger yields up twelve — and a **net loss of S\$76 million**. The reason is one line: net fuel cost of S\$2,253 million, **up 78.5 per cent**, a S\$991 million increase the airline attributes directly to the Middle East conflict.¹ Operating profit fell by S\$299 million to S\$106 million.

+78.5%

SIA'S NET FUEL BILL, YEAR ON YEAR, IN THE JUNE QUARTER. IT IS THE WHOLE DIFFERENCE BETWEEN A RECORD REVENUE QUARTER AND A LOSS — AND IT IS THE HORMUZ STORY (P3) ARRIVING IN A SINGAPORE COMPANY'S ACCOUNTS¹

ROUTES AND METAL

Four daily to Ho Chi Minh City, and fewer seats. From 25 October SIA goes to four daily on Saigon and three on Da Nang. But one existing 787-10 rotation is downgauged to a 737-8 MAX, so weekly seats to Saigon *fall about three per cent* and Business Class capacity falls about fourteen. Da Nang, by contrast, is a clean fifty per cent lift on consistent metal. Both are seasonal and revert on 28 March.² *Schedule-derived: SIA has published no release.*

EgyptAir returns to Changi after about twenty-five years, announced with Changi Airport Group and Sydney Airport at the El Alamein Air Show, 9–12 September. Cairo–Singapore–Sydney on an A350-900, thirty Business and 310 Economy, from early 2027. Frequencies and on-sale unannounced, and subject to approvals — including whether it can sell Singapore–Sydney standalone.²

The RTS Link has slipped, quietly. LTA announced completion of Woodlands North station on 4 September — a five-minute, four-kilometre crossing, up to 10,000 passengers an hour each way, immigration cleared once at the departure station. But the reporting around it puts passenger service in **early 2027**, while LTA's own project page still reads "end of 2026."² Nobody has announced a delay; the date simply moved in the coverage and not on the page.

RULES, MONEY AND THE CLOCK

Thailand's tourist fee is now a flat 450 baht, endorsed on 14 August, covering multiple entries within thirty days. The public hearing closes on **28 September**; collection is targeted from the first quarter of 2027, air arrivals first. It supersedes the earlier split 300-by-air, 150-by-land plan.³ The TDAC arrival card remains mandatory and free — any site charging for it is a scam.

THE SCEPTIC · A BORDER SYSTEM THAT QUIETLY STOPPED APPLYING

The EU's Entry/Exit System became fully operational on 10 April, and the flexibility window for member states expired on **6 September**. Reporting on the 14th says at least **nine Schengen states — France, Germany, Italy, Belgium, the Netherlands, Greece, Malta, Portugal and Switzerland — have been given informal approval to keep suspending biometric capture**, to maintain border fluidity, with no new deadline set. Malfunctioning kiosks, connectivity faults, queues, staffing. The Commission declined to comment.⁴

Two honest caveats. The European Commission's own pages were unreachable from this desk, so the primary confirmation is held and this rests on one specialist outlet. And ETIAS — the separate pre-travel authorisation — remains scheduled for the last quarter of this year with no confirmed date and no applications open.⁴ For a traveller the practical read is unchanged and slightly absurd: plan for the checks, expect not to get them, and do not assume the queue you met in June is the queue you will meet in December.

SOURCES 1 Singapore Airlines' Skytrax release, 18 September 2026, and the group's Q1 FY2026/27 results of 28 July 2026, read as the company's own PDF. 2 SIA's Vietnam frequencies are derived from loaded schedules and reported by Mainly Miles, 16 September 2026; the airline has issued no release, and this is flagged in the text. EgyptAir's return announced 9–12 September 2026. RTS Link: LTA's completion announcement of 4 September as reported, against LTA's own project page read 18 September, which still carries the end-2026 target — the discrepancy is printed rather than resolved. 3 Thailand's 450-baht fee endorsed 14 August 2026, hearing to 28 September. 4 EES grace-period expiry and the nine states: Biometric Update, 14 September 2026, single-sourced; the Commission's own pages were unreachable, so primary confirmation is held. ETIAS status per ABTA.

DEPARTURES • THE AIRPORT, THE POINTS, THE SMALL PRINT

Changi has stopped growing, and freight is doing the work

Forty point seven million passengers through July, up one-tenth of one per cent. Five of the first seven months were down year on year. Meanwhile the cargo is up eight and a half.

Read off Changi's own traffic statistics on Friday: **40.70 million passenger movements from January to July, up 0.1 per cent** on the same period last year. July alone was 5.88 million, **down 1.4 per cent** year on year and down half a per cent against July 2019.¹ Aircraft movements are flat at 219,000.

Airfreight, over the same seven months, is **up 8.5 per cent** at 1.28 million tonnes. An airport whose passenger number is flat while its cargo number climbs is an airport in a fuel-price economy: freight prices through, and discretionary long-haul does not.

40.70m

CHANGI PASSENGERS, JANUARY TO JULY 2026, AGAINST 40.66 MILLION A YEAR EARLIER. LAST YEAR FINISHED AT A RECORD 69.98 MILLION; THIS YEAR IS RUNNING LEVEL, AND AUGUST WAS STILL UNPUBLISHED ON FRIDAY¹

THE CHATTER • FOUR FLIGHTS, FEWER SEATS

Singapore's most-read miles blog broke the Vietnam expansion on Tuesday and led not with the extra flight but with what was buried inside it: a 787-10 rotation swapped for a 737-8 MAX, leaving Ho Chi Minh City with three per cent fewer weekly seats and fourteen per cent less Business Class *despite going to four daily*. Its verdict on Da Nang, by contrast, was that it is a clean lift on consistent metal.² That is the Singapore frequent-flyer read of 2026 in miniature — more flights on the timetable, smaller aeroplanes underneath them, and the premium cabin quietly thinning. *This desk could not reach the specialist flyer forums this week, so nothing here is attributed to them.*

IN BRIEF

KrisFlyer, this fortnight: a redemption sale of ten to fifteen per cent off Economy and Premium Economy awards from 8 September; a fifty per cent transfer bonus to and from Accor; and a **ten per cent bonus on Marriott Bonvoy transfers to 30 September** — 3,000 miles becoming 2,475 points instead of 2,250, capped at 250,000 miles a year, and still one of the weaker things you can do with a KrisFlyer mile.²

Changi's new private terminal. The former JetQuay CIP terminal is being rebuilt by the airport group with Plaza Premium, opening **mid-2027**. Terminal 5 remains mid-2030s, though 1.7 kilometres of tunnel now link Terminal 2 to the T5 site — one for baggage, two for the people-mover.¹

Still cancelled: SIA's Dubai service through 24 October, with the earliest restart on the 25th; the Riyadh launch has been pushed a second time, from 1 September to **1 December**, and cut by a quarter in frequency.²

Cathay's Aria Suite reaches Perth from 26 October on selected services to 27 March — roughly seventeen of thirty-five 777-300ERs retrofitted so far. *Reported by a specialist tracker; Cathay's own newsroom carried no Perth-specific release this desk could reach.*²

Korean Air received final regulatory approval for the Asiana mileage integration plan on 17 September — relevant to anyone still holding Asiana miles and connecting through Incheon.²

Japan's departure tax has been ¥3,000 since 1 July, up from ¥1,000, and is normally bundled into the ticket. **Bali's levy** stays at IDR 150,000; what changed is enforcement, with spot checks at Tanah Lot, Uluwatu, Besakih and the airport, and hotels now permitted to collect it. The proposed rise to 300,000 still has no date.³

And for the second weekend of October: Grand Prix road closures run from one minute past midnight on Monday 7 October to 5.30am on Tuesday the 13th, with extended MRT hours on the North-South and East-West lines. If you are flying out that weekend, the transfer needs planning that the ticket does not mention.³

THE SCEPTIC • WHAT A FLAT AIRPORT ACTUALLY MEANS

It is tempting to read Changi's flat line as a Singapore problem. It is more likely a fuel problem with a regional shape. Asian jet fuel is reported to have gone from about US\$88 a barrel in mid-February to around US\$209 by late March after the strait closed; Lufthansa cancelled twenty thousand flights through October over cost and supply.¹ An airline that must pass through a seventy-eight per cent fuel increase either raises fares or shrinks the aeroplane, and both suppress the marginal leisure passenger. Changi's flat number and SIA's downgauged Saigon rotation are the same decision seen from two ends.

The barrel figures above are reported rather than read from a price assessment, and are flagged accordingly.

SOURCES ¹Changi Airport Group traffic statistics, read at changiairport.com on 18 September 2026; August figures were not yet published. Terminal 5 and the JetQuay rebuild from Changi's own future-developments page and reporting of 18 April 2026. The jet-fuel price path is reported by a travel-industry blog and is flagged as secondary; the Lufthansa cancellations from Al Jazeera, 23 April 2026. ²Mainly Miles, 16 and 17 September 2026, for the Vietnam schedule analysis and the KrisFlyer promotions; the same publication of 16 July 2026 for the Dubai and Riyadh changes. Cathay's Aria Suite rollout from a specialist tracker, flagged in the text as lacking an airline release. Korean Air's mileage approval reported 17 September 2026. ³Japan's departure-tax increase effective 1 July 2026. Bali levy status from a specialist aggregator; the provincial government's own page was unreachable and the figure is held at one remove. Grand Prix road closures and transport arrangements from the Land Transport Authority news release of 4 September 2026.

THE GRAND TOUR · TAITUNG AND TAIWAN'S EAST COAST

Three of the postcards are currently false. Go anyway — and go in six weeks

The eight-arch bridge at Sanxiantai is closed until the end of 2027. The balloons mostly never leave the ground. Dulan's sugar factory is dark six nights a week. What is left is the best of it: a coast road with nothing built on one side, a valley that turns gold twice a year, and the world's only shore-side seawater hot spring.

this story  

Start with the fact that will otherwise ruin a day. The **Sanxiantai eight-arch bridge — the image every article about Taitung uses — closed to the public on 4 May 2026 and will not reopen until the end of 2027**: corroded steel box girders, cracked piers, salt.¹ You cannot walk to the island. The coastal park, the pebble beaches and the erosion landscape are open, and the bridge photographs beautifully from the shore. The English-language page on the official site still says “always open.” The Chinese one carries the closure notice. Anyone trusting the English will drive two hours and be turned round.

The Puyuma tilting express does Taipei to Taitung in **three and a half hours for NT\$936**, about forty Singapore dollars, with twenty-two departures a day. **Book it the day it opens**. Reservations open twenty-eight days ahead at midnight, and east-coast seats are the scarce commodity on the network — for Lunar New Year the railway now releases the east-coast lines a full day *before* the west coast, which tells you everything about demand.² Taitung County's own advice is to book at least fourteen days out and, for peaks, to log in at half past eleven and submit at midnight.

WHAT THE FORUMS ACTUALLY SAY · AND ONE MAN WHO CHANGED HIS MIND

The sharpest line about Taitung City on Taiwan's oldest English forum is a slur: “**the Zhongli of the east coast.**” It traces to a 2013 post quoting a Lonely Planet author, and was sharpened in 2015 by **afterspivak**, who called it “a dispersed and dysfunctional and somewhat dumpy town with a train station that's been (re-)located to a ridiculously remote location.”³ The rebuttal is blunter — **hannes**: “total BS. The two places have next to nothing in common.” And then the thing that settles it: **fanglangzhe**, the original sceptic, came back in December 2015 after four days and recanted — “A lovely town to be based in and pretty central to explore surrounding areas... Relative to other Taiwan townships, not dumpy or dysfunctional at all.”³ The station complaint is true; the dumpy complaint dissolves after four days. The consensus underneath both is unanimous — **Hokwongwei**: “you'll definitely need a motor vehicle, preferably a car.” Without one you can do Taitung City and Chihshang by train and bicycle. You cannot do the valley and the coast on one trip, and you cannot be at Luye Highland at quarter past five in the morning.

THE SCEPTIC · AND THE TWO CONDITIONS

Does it reward the effort? Yes, with two conditions and a concession. The conditions: **rent a car, and give it five days**. This is a hundred and seventy kilometres of linear destination with two parallel highways and two offshore islands; it punishes the three-night traveller. The concession is the honest one. Come in July, base yourself in Dulan without a car, and build the trip around Sanxiantai and the balloons, and *it will feel exactly like a long drive to a nice beach*. The bridge is shut, the balloons stay tethered, and Dulan's sugar factory — the arts complex the town is sold on — is, in the words of a Taiwan-resident writer, “almost totally closed” outside events.³ Time Dulan for a Saturday night, when the café has live music, or accept a quiet village with overpriced rooms.

SOURCES 1 The Sanxiantai closure notice read in Chinese on the East Coast National Scenic Area's own attraction page, works running to end-2027; reported by Public Television Service and China Times, 29 April 2026. The English version of the same official page carried no notice when read on 18 September 2026. Typhoon seasonality from Taiwan climate guides; the Central Weather Administration's own tables were unreachable and are held. Balloon Festival dates and the two August typhoon cancellations from the festival's own site. 2 Train times, fares and frequencies from timetables.tw and Taitung County's rail page; the 28-day window and the Lunar New Year east-coast-first release from Taiwan Railway Corporation's published 2026 booking calendar. 3 Forumosa threads read directly and quoted with handles and dates as posted, 8 September and 26 December 2015; Dulan assessments from Nick Kembel's 2025-updated guide.

THE GRAND TOUR · WHAT ACTUALLY EARNS THE TRIP

A seawater hot spring, a prison island, and a valley that turns gold twice a year

The things worth planning around, with what they cost — and an honest verdict on the dark skies, which are marketing rather than certification.

Zhiben, fifteen kilometres from town, is the hot-spring valley: weakly alkaline carbonate, pH seven to eight, thirty-nine to forty-five degrees, odourless. The value here is genuine — Toyugi, run by the county farmers' association and therefore unusually transparent about rates, publishes one night with two meals **from NT\$3,399**, about S\$136. Hotel Royal Chihpen, Taiwan's first international-class hot-spring resort, has 183 rooms including 41-square-metre family suites, but publishes no rack rate.¹

One of three

ZHAORI, ON GREEN ISLAND, IS ONE OF ONLY THREE SHORE-SIDE SEAWATER HOT SPRINGS IN THE WORLD — THE OTHERS IN KYUSHU AND SICILY. SOURCE TEMPERATURE SIXTY TO SEVENTY DEGREES; THE PATTERN IS TO SOAK FROM BEFORE DAWN, BECAUSE THE POINT IS THE SUNRISE²

Green Island is twenty minutes by fast ferry from Fugang Harbour, or fifty on the standard boat. The crossing is routinely rough — transparent sick-bags are handed out — and cancellations inside two days are typically non-refundable. Two nights minimum, for one full day.²

And the island carries the other thing. The **White Terror Memorial Park**, free, nine to five, closed Mondays from October, with free guided tours at ten and three, holds Oasis Villa — the political prison — and a wall inscribed with the names of those jailed and executed under the Kuomintang.² Green Island still has a working prison. Some visitors find the prison-themed restaurants nearby hard to take beside the memorial, and they are right to.

DOING THE INDIGENOUS CULTURE PROPERLY

Seven peoples live in Taitung County — Amis, Paiwan, Bunun, Puyuma, Tao, Rukai and Kavalan. The answer to “where can you meet it without it being a show” is unusually concrete: the county runs a **day-of-the-week programme of village food experiences at NT\$1,980 a head**, minimum six people, with pickup from Taitung City — Bunun spinning-top culture and red quinoa on Tuesday at Yanping;

a smoked-fish and flying-fish meal on Saturday; Rukai hunter's barbecue and wild-vegetable porridge on Sunday at Beinan.³ A booked meal with a family that cooks, rather than a stage.

MONEY · FOR TWO ADULTS AND A TWELVE-YEAR-OLD, ON THE GROUND

At **S\$1 to TWD 24.99**, read on Friday: hire car from about S\$60 a day for an economy, S\$110 for an SUV, plus S\$20 for full insurance; a Zhiben hot-spring night with two meals about S\$136; a Taitung City four-star around S\$127–149; the Puyuma from Taipei about S\$37. **Comfortable comes to roughly TWD 10,000 a day, about S\$400 for the family**; lean, on a guesthouse and night-market food, about S\$225. Which makes it materially cheaper than Japan or Korea for the same quality of room — and the hot-spring hotels are the standout. Bring cash; the east coast is far more cash-based than Taipei. And carry both your Singapore licence *and* an international permit, S\$20 from the AA: driving on the Singapore licence alone is treated as driving unlicensed, and Taiwan drives on the right.⁴

THE SCEPTIC · THE DARK SKIES ARE MARKETING, NOT CERTIFICATION

Taitung sells itself as Taiwan's starry county and has run a free night-sky programme since 2018. Here is the honest position. **Taiwan has exactly one certified International Dark Sky Park, and it is not on the east coast** — Hehuanshan, in Nantou, in the Central Mountain Range, certified in August 2019, a full day's drive away across the island.⁴ And when the Taiwan Dark-Sky Association's own chairman listed the country's six best stargazing sites, Taitung, Lanyu and Ludao were all absent; the one east-coast site he rated was a fishing harbour near Hualien.⁴ What is true: Taitung is thinly built, the mountains wall off the west-coast glow, and there is nothing but two hundred kilometres of Pacific to the east, so the skies are genuinely dark by the standards of a crowded island — and Lanyu and Green Island, offshore and barely lit, are properly dark. Soaking at Zhaori under stars before sunrise is real. But go for the coast, the valley and the hot springs. Treat the stars as a bonus, not a reason.

SOURCES 1 Chihshang cycling costs and the parallel-road advice from Nick Kembel's guide; Zhiben water chemistry and Toyugi's package rates from the resort's own site, updated 4 Sep 2026; Hotel Royal Chihpen's room count from the group's own page, which publishes no rates. Taitung City hotel figures are aggregator bands, not quotes. 2 Zhaori, the ferries and the whale-watching season and rates from the East Coast National Scenic Area's official pages; a 2026 guide reports the lower Zhaori pools still closed after typhoon damage, which we could not confirm and have not asserted. Memorial park hours and free admission from the National Human Rights Museum. 3 The seven peoples, the day-of-the-week food programme and the harvest-festival visitor rules from Taitung County Government's own pages; Lanyu costs and respectful-visiting guidance from a 2026 trip report. Museum and music-village hours cross-checked across two secondary sources after the venues' own sites proved unreachable. 4 Exchange rate read 18 Sep 2026; car hire is one operator's published rates. IDP price and conditions from the Automobile Association of Singapore; Taiwan's 30-day rule from the MOTC. Hehuanshan's 2019 certification from Taipei Times and Taiwan News; the Dark-Sky Association chairman's six recommendations from Taipei Times, 22 Jul 2022.

THE DIARY • THE DATED AGENDA

Sumo goes on sale at nine this morning, and the marathon shuts in eleven days

Every item below was re-checked against a promoter, a venue or an official ticketing page before it was printed. Two things that were live on Friday are dead this morning, and they are named rather than quietly dropped.

this story  

SINGAPORE

BOOK NOW • PLAN AHEAD

25
SEP

Halloween Horror Nights 14 — opens

UNIVERSAL STUDIOS SINGAPORE • 18 SELECT NIGHTS TO 1 NOV • \$68 NON-PEAK / \$78 PEAK • **RWS STATES IT IS NOT RECOMMENDED UNDER 13**¹

9–11
OCT

★ F1 Singapore Grand Prix — the first Sprint weekend

MARINA BAY • HOSPITALITY **SOLD OUT**; GRANDSTANDS \$1,698–1,798, SINGLE-DAY FROM \$228 • SAT: THE KILLERS. SUN: LANA DEL REY, JANET JACKSON¹

11
DEC–
4 APR

★ Pop after Pop: Art Around Us

NATIONAL GALLERY SINGAPORE • 200+ WORKS, THE REGION'S LARGEST POP ART SHOW, FIRST COLLABORATION WITH THE GUGGENHEIM • \$15 CITIZENS AND PRS; MEMBERSHIP FROM \$20 BEATS TWO VISITS²

17–22
DEC

★ BTS — World Tour 'Arirang'

NATIONAL STADIUM, 7PM • FOUR NIGHTS: 17, 19, 20 AND 22 DEC • NOW **SOLD OUT ON EVERY OFFICIAL CHANNEL**, NO FIFTH NIGHT • THE PROMOTER PROHIBITS RESALE, SO TREAT SECONDARY LISTINGS AS A RISK, NOT A BARGAIN²

13–18
APR

Bruno Mars — The Romantic Tour

NATIONAL STADIUM, 2027 • FOUR NIGHTS WITH ANDERSON .PAAK • GENERAL SALE OPENED 14 SEP AND HAS **NOT SOLD OUT**²

JAPAN & THE REGION

PLAN THE FLIGHTS

8–22
NOV

★ Kyushu Basho — general sale opens TODAY

FUKUOKA KOKUSAI CENTER • **10AM JST, WHICH IS 9AM HERE** • TICKET OOSUMO, TICKET PIA OR SEVEN-ELEVEN COUNTERS • ALL FOUR LOTTERIES HAVE RUN, SO SEAT TYPES THAT CLEARED WILL NOT APPEAR — BEST ODDS ARE FIRST-WEEK WEEKDAYS. **PRICES ARE NOT PUBLISHED OFFICIALLY**³

9
SEP–
13
DEC

★ Masterpieces of the Renaissance from the Louvre

NATIONAL ART CENTER, TOKYO, ROPPONGI • CLOSED TUE • 50+ WORKS; LEONARDO'S *LA BELLE FERRONNIÈRE* IN JAPAN FOR THE FIRST TIME³

22
OCT

Jidai Matsuri, and the Kurama Fire Festival the same night

KYOTO • ~2,000 COSTUMED PARTICIPANTS FROM MIDDAY; TORCHES AT KURAMA FROM 6PM • **POSTPONED A FULL DAY FOR RAIN, DECIDED AT DAWN** — DO NOT BOOK A SAME-NIGHT FLIGHT OUT OF KANSAI⁴

4–11
FEB

★ 77th Sapporo Snow Festival

ODORI PARK, SUSUKINO AND TSUDOME • FREE — THE CONSTRAINT IS SAPPORO HOTEL ROOMS, WHICH GO SIX TO NINE MONTHS AHEAD. TSUDOME IS THE SITE THAT WORKS WITH CHILDREN⁴

PLANNING NOTES • AND TWO THINGS THAT DIED THIS WEEK

Dead, and named rather than dropped. The F1 Thursday Pit Lane Experience ballot closed at 11.59pm on 18 September — the night before this page. If you were going to apply, the answer is no, not “soon.” And **The Art of the Brick** at Singapore Expo has closed; the official site now says so plainly.²

Japanese ticketing is lottery-first, not queue-first. The Fan Club and Pia draws are applications, not races; only the general sale afterwards is a race. Sequence matters more than speed, and joining the sumo association's fan club before the first draw is the only real edge. Expect the January 2027 Tokyo calendar to post in November — the dates circulating now conflict and this desk will not print them.³

Two booking traps. Kyoto's maples are forecast about **six days later than average** this year, peaking around 11 December, while Tokyo is roughly on time at 29 November — a late-November Kyoto trip booked on habit will arrive early.⁴ And on the Grand Prix: the promoter's haze contingency is masks, aid stations and real-time PSI tracking, not refunds. Read that as the show going on regardless (p5).¹

SOURCES 1 Singapore Grand Prix ticket and entertainment pages, read 18–19 Sep 2026; the haze contingency is quoted from the promoter via press coverage of 19 Aug and is not on its own newsroom. Halloween Horror Nights from Resorts World Sentosa's official guide, including its under-13 advisory. SGIFF dates from sgiff.com. 2 Marathon dates and sponsor change from singaporeinternationalmarathon.com; Pop after Pop from the National Gallery's media-centre release; BTS and Bruno Mars from Live Nation Singapore; The Art of the Brick closure from the exhibition's own site. 3 Kyushu Basho sale time, channels and completed lotteries from Ticket Oosumo; seat prices are not published there and no third-party figures are printed. Asian Games from the organising committee's ticket site; the Louvre exhibition and Shōjo Manga Infinity from the National Art Center, Tokyo. 4 Jidai Matsuri from Kyoto City's tourism site; Sapporo Snow Festival from sapporo.travel; Australian Open from ausopen.com; the foliage forecast from the Japan Meteorological Corporation, 2 Sep 2026.

THE TABLE • SINGAPORE'S NEWEST OPENINGS

Twenty-six seats on Mohamed Sultan Road, and nobody has eaten there yet

The most consequential opening of the season is a restaurant this desk cannot review, because it has not served a paying meal. That is worth printing plainly — and worth booking before it becomes impossible.

Arō is Johanne Siy's own place at last. She was head chef at Lolla and was named Asia's Best Female Chef in 2023, off a curriculum vitae that runs Le Bernardin, Café Boulud, Restaurant André, Noma, Relæ and Fäviken. She has taken the Mohamed Sultan Road shophouse vacated by the one-starred Esora. The name does double duty: arō is Latin, to plough or cultivate; áro is love in the dialect of her hometown in Pangasinan.¹

The format is a single ingredient-led tasting menu she calls Creator's Cuisine, built around under-recognised ingredients and the growers behind them, seafood-forward, in a room by Keiji Ashizawa Design hung with work from her own collection. **Twenty-six seats** — ten at the counter, two four-tops, and a private room seating eight as four pairs — five days a week, Thursday to Monday, bookings through SevenRooms.¹

And the buzz is nothing, because there cannot be any. The Straits Times previewed it in July; opening trackers have it; there are no reviews because nobody has eaten a paid meal. What the trade is reacting to is a curriculum vitae and an address. This desk will say that rather than dress it up: too new to call. The menu price is not published and is not guessed at here, and the opening date is reported as 25 September by one tracker while the restaurant itself publishes none — so call it late September.¹ **The pick, and the reason is structural:** it is the only genuinely new thing on this list rather than an import, a relocation or a second outlet, it is the debut of the most decorated female chef working in Singapore, and at twenty-six seats the window between opening and impossible is usually two or three weeks. Book the counter now, eat in October, and judge it then.

OPEN, AND REVIEWABLE

Decker Barbecue, Old Bukit Timah Fire Station — Decker's second site and new flagship, in the conserved fire station, smoking up to twelve hours in an imported offset smoker sited where diners can see it. Converted fire doors are now floor-to-ceiling windows. Brisket, pulled pork, kaffir-lime hot honey wings \$18, platters \$40–288. **The warmest reception of the five and no meaningful dissent found;** one reviewer found the room packed on a Tuesday evening five weeks after opening.² The only caution worth printing is that it is a second outlet of an eleven-year-old brand, not a new idea. Book the weekend brunch, which nobody thinks of, sit outside, and ask them to hold brisket when you reserve — it sells out.

THE CHATTER • AND AN HONEST NOTE ON HOW WE READ IT

The reception above is drawn from published reviewers and food media — Eatbook, SETHLUI, Time Out, HungryGoWhere, Honeycombers and the local blogs — read this week. **The social layer is missing and this desk will not pretend otherwise:** Reddit and Burpple both refused every request from this build. So where the word “consensus” appears above it means a consensus of critics, not of diners, and that is a narrower thing. The gap matters most on Mukai, where a divisive verdict from three publications is exactly the case where a hundred ordinary diners would be worth more than three professionals.

THE SCEPTIC • THE PICK IS THE ONE WE CANNOT VOUCH FOR

There is an obvious objection to leading a restaurant column with a restaurant nobody has eaten in, and it is a good objection. A twenty-six-seat debut from a decorated chef is exactly the kind of opening that gets a rapturous first month on pedigree and a quieter second year on cooking; this desk has watched it happen. Against that: the only service a magazine can perform for a restaurant this size is to tell you before the booking window shuts, because afterwards the information is worthless. So we are telling you early and telling you the risk in the same breath, which is the honest version of a tip. If it disappoints, that will be reported here too.

SOURCES 1 Árō: the restaurant's own site for concept, seat count, design credit and hours; Johanne Siy's record as widely reported; *The Straits Times* preview, July 2026. The 25 September date comes from one opening tracker and is not published by the restaurant, so “late September” is printed instead; no menu price is published and none is estimated. 2 Tembusu, Mukai and Decker Barbecue: concept, addresses, hours and prices from the venues' own listings and from HungryGoWhere, Eatbook and Time Out Singapore, read 18 Sep 2026; quoted verdicts are attributed to the publications that made them. Reddit and Burpple were unreachable and no social sentiment is reported. Not repeated here, having run recently: Hikiniku to Come, Kikukawa, Modu High, Foura, Kuro Kare, Casa Lola, Katsu by Kyu, Le Ju Xuan, Toriei, Noor and Sushidan.

THE UNDERCURRENT · JAPAN

The children learned the word “rate”

A five-hundred-yen sticker sheet for seven-year-olds became, in about twenty months, a speculative asset with a published price index, a counterfeiting problem, criminal prosecutions, school bans and a crash. What it destroyed on the way was a playground gift economy older than anyone now trading in it.

this story  

The sentence that should stop you is not about money. It is a thing Japanese primary-school children now say to each other, recorded by a writer watching them trade in February: **“that sticker’s rate doesn’t match, so I can’t trade.”**¹

The product is **ボンボンドロップシール** — Bonbon Drop stickers, universally shortened to **ボンドロ**, *Bondoro*. A printed base under a domed resin cap, so each motif sits under a glassy, sweet-like blister. Five hundred to five hundred and fifty yen a sheet. It is made by Q-Lia, a ninety-seven-person company in Osaka with eighty-four women among its staff, and distributed by Sun-Star Stationery. It launched in March 2024.¹

By the end of last December, cumulative shipments had reached fifteen million sheets, running at two million a month.² By late June the figure reported was thirty-four million. In June alone the company shipped around four million units, planned five for July, and held orders for some thirty million more across the rest of the year. Its president told an awards ceremony in Tokyo that production was running at three and four times previous levels and added: **“Some people wonder when the Bonbon boom ends — it’s nowhere near over.”**²

That last number is the whole essay. Forty-six per cent of children reporting a sense of having *won or lost* is not a description of swapping. It is a description of a market, and it is being reported by nine-year-olds.

The first was purchasing power. The demand did not come from the stated target of preschool and lower-primary girls. It came from **平成女児**, the

“Heisei girls” — women now in their twenties and thirties who traded stickers in the early 2000s and are now solvent. Q-Lia’s own head of the president’s office noted that adults who once traded stickers are “now purchasing with greater economic capacity,” and the company confirmed it would not reposition away from children.² A receipt-data company’s release quotes one such buyer: “Back then I agonised over spending my meagre pocket money. Now I can buy as much as I like with an adult’s means. It’s like granting a childhood wish.”³ Pocket-money prices met salaried demand, and the pocket-money side lost.

“That sticker’s rate doesn’t match, so I can’t trade.”

The second was liquidity. Mercari and Rakuten supplied continuous, public price quotation. Five-hundred-yen sheets were reported moving at two to five thousand, popular designs above ten thousand, with one reported outlier at a hundred thousand.⁴

The third, and the one that changes the story from a mania into a system, was data. A site called Bonbon Navi — explicitly unofficial — publishes a sortable rate ranking by series, derived from Mercari and Rakuten trading data, alongside a buy-back price board, a lottery-sale tracker, a first-come-sale tracker and stock listings for all forty-seven prefectures. It was last updated on 14 September.⁴ A rival publishes a guide to finding and behaving at exchange meets. This is a functioning financial data vendor for a children’s toy, and it is what the children are reading when they say a rate does not match.

SOURCES 1 The children’s exchange language recorded on note.com, 18 February 2026. Product description, maker, distributor and the March 2024 launch from Q-Lia and Sun-Star Stationery’s own statements. 2 Shipments: fifteen million cumulative at end-December 2025, ITmedia Business Online, 4 March 2026. The thirty-four-million figure for late June 2026 is a Kyodo figure reaching this desk through two independent renderings but no page we could open, and is reported as such; note two other outlets were still carrying “over ten million” in February, apparently a stale figure recycled. The president’s remarks and the order book from Advertimes, 18 June 2026. 3 Survey of 2,484 Nifty Kids users, December 2025 to January 2026; receipt-data findings and the quoted buyer from Research and Innovation Inc., 6 February 2026. 4 Resale prices as reported in February 2026; Bonbon Navi’s rate ranking and prefecture stock listings read at source, 18 September 2026.


The writer who recorded the “rate doesn’t match” line noticed something worse in the same sessions: children gifting stickers to newcomers not out of kindness but as a booked receivable against future trades.¹ Generosity had not disappeared. It had been re-priced as credit.

Then the machinery arrived, in the order such machinery always arrives. **Retail broke first.** On 27 January a chain opened online pre-orders for a mini range; the servers fell over within seconds and sales were cancelled. On 1 February a Shibuya store suspended three-dimensional sticker sales “for customer safety and trouble prevention.” On **4 February the chain suspended three whole series across every store and its online shop**, citing nationwide shortages and crowd control. The next day Q-Lia announced twenty-four-hour production. Around three hundred people queued from morning at one discount store’s new branch.¹

¥500 → ¥100

THE MARKDOWN A REPORTER FOUND ON SHELVES IN SAITAMA AND TOKYO THIS MONTH — AN EIGHTY PER CENT CUT, WITH TOY SHOPS PILING STOCK AT A HUNDRED TO TWO HUNDRED YEN, AND CHILDREN WALKING PAST IT TO CROWD AROUND OTHER MERCHANDISE²

Then the counterfeiters, and then the police. In May two Osaka men were arrested in Nagoya holding 1,016 counterfeit sheets for sale, and were fined by the Nagoya Summary Court in June; fakes were priced at 280 and 580 yen against roughly 550 for the real thing.² **And then the schools, because no regulator came.** Public primary schools sent emergency notices banning three-dimensional stickers, citing money disputes and bullying.²



SOURCES 1 The gifting-as-receivable observation from the same note.com account of 18 February 2026. The January and February retail suspensions, including the chain-wide withdrawal of 4 February 2026 and its stated reason, reported by ITmedia News, 4 February 2026; the queue and the twenty-four-hour production announcement from ITmedia Business Online, 4 March 2026. 2 Daily Shincho, reporting by Miyahara Takashi, 5 September 2026, read in full via its mirror — the markdowns, the shop-floor observations, the retailer’s verdict and both quoted lines are from it. 3 Spending accounts and the children’s survey reported by Shueisha Online and by Research and Innovation Inc.’s release of 6 February 2026; the social posts are reported as reported, not verified individually. Earlier: 15 February 2026, which collates social-media posts from adults describing compulsive collecting; those posts are reported as reported, not verified individually, and no individual is named.

That is the crash, reported on 5 September. The shop staff quoted in it sound relieved rather than sorry: “We used to get stock enquiries by phone daily; they’ve stopped — the boom has calmed down.” The reporter’s own comparison was to rice reappearing on supermarket shelves after a shortage — *where on earth was it all this time?*²

70.2%

OF 2,484 JAPANESE CHILDREN SURVEYED OVER DECEMBER AND JANUARY WHO WERE COLLECTING STICKERS — 77.5 PER CENT OF PRIMARY PUPILS. NINETY PER CENT OF ALBUM OWNERS TRADE, AND **46.1 PER CENT SAID THEY HAD FELT THEY GAINED OR LOST ON A TRADE**³

And then the line that is the reason this piece exists. Asked what the boom had been worth, a mass-merchant employee said the category is thin-margin, high-volume: the profits were small against the units moved, and staff had to be pulled onto queue management. “儲かったのは転売屋だけ” — only the resellers made money. And then, from shopkeepers: “Please, God, don’t let anything we stock become a boom.”²

“Please, God, don’t let anything we stock become a boom.”

When retailers begin praying against demand, the resale layer has stopped being a market and started being a tax.

Here is the argument. Japan has built a resale and price-discovery apparatus so complete, so cheap and so fast that it can convert almost any consumer good into a speculative asset within weeks — and in doing so it destroys whatever social function the good had. Bondoro is the proof precisely *because* the good was trivial: a five-hundred-yen sticker for seven-year-olds, and the thing it destroyed was a children’s gift economy.

Count who lost. The chains took the reputational damage and diverted staff to crowd control for thin-margin revenue. Q-Lia spent capital expenditure tripling capacity into a market that turned. Parents spent — one reported over fifty thousand yen in six weeks.³ Children learned that their treasures had a bid, and then that the bid had gone. The only party that cleared a profit produced nothing at all.

Japan has done this before, and the way it ended the last time is the most useful fact in the story.

In August 1985 Lotte put a sticker inside a thirty-yen chocolate wafer and called the series Devil versus Angel. **Bikkuriman** ran at thirteen to thirty-three million units a month through the late 1980s, around four hundred million a year at peak. The parallels are exact, item for item. Children bought the chocolate for the sticker and threw the chocolate away in bulk, which became a national controversy. Shops imposed limits — three per customer. And a rate economy emerged with a rarity hierarchy, holographic over angel over charm over devil, in which children traded not only stickers but *information*: which shop, what time, which position in the carton.¹

Bikkuriman proves that the only thing which has ever killed a Japanese children's-collectible bubble is the deliberate destruction of the scarcity premium — and that in 1988 a state regulator did it, at the manufacturer, in three months.

In 2026 the Fair Trade Commission is absent. Japan's only anti-scalping statute covers event tickets; the constitutional presumption in favour of freedom of contract means there is no general law against reselling goods. So the enforcement burden has migrated to the marketplace — and **Mercari has already accepted the role**. It abandoned its 2021 position that anything legal is freely tradable in favour of pre-emptive prohibition: temporary listing bans on the Switch 2, then on the McDonald's Chiikawa toys in May, then, from midnight on **16 September 2026**, on unopened Pokémon Card thirtieth-anniversary products. The Pokémon Company has gone further and introduced identity verification for its own lottery sales.²

Restricting resale turned out to be good for the resale business.

And here is the part a Western paper will notice three weeks late. Mercari told its results call on 6 August that these restrictions *helped drive growth* — marketplace gross merchandise value of ¥1,285.6 billion, up 14.7 per cent, its first double-digit year in

three.² The marketplace has discovered it is cheaper to be a regulator than to be regulated, and is monetising the role. Bondoro is not yet on that list. On current trajectory the next children's craze will be — and a listed company, not the Fair Trade Commission, will decide which toys children may trade at a price.

One last detail, because it is the one that stays with you. The market research found that purchases peaked twice a day — at opening, and at restock — and that the buyers in their thirties and forties were mothers and, visibly, fathers.³ A published circuit even existed: eight in the morning at the convenience store, the variety chains at opening, two o'clock at the mall, four at the toy shop.

THE SCEPTIC · FOUR ARGUMENTS AGAINST THIS PIECE, AND WHAT SURVIVES THEM

One: the bubble did not burst, the copycats did. July reporting is explicit that Bondoro itself stayed scarce while the me-too sheets glutted and collapsed, and the crash article concedes on its own second page that rare designs still trade at a premium. **Two: the timing is against us.** A major chain ran a Bondoro lottery on 4 September, and Bonbon Navi lists another opening on **17 September — the day before this issue went to press.**³ A product that needs a ballot two days ago is not a crashed asset. **Three: the maker has an order book** — some thirty million units for the rest of the year. **Four, and the strongest:** Takafumi Horie's position that resale "is just filling a gap between supply and demand," under which the resellers were the clearing mechanism for a maker who under-produced by a factor of eight.³ **Three of those four land.** This desk is not claiming Bondoro is finished, and has printed the hundred-yen markdown and the 17 September ballot in the same breath deliberately, because the contradiction is the story. What none of them touches is the actual claim, which is not about one sticker. Whether Bondoro trades at a hundred yen next spring or ten thousand, Bonbon Navi will still publish the rate, the chat groups will still relay stock, Mercari will still decide what may be listed — and children will still be saying that a rate does not match. **The mechanism outlives the product.**

Somewhere in that apparatus is a parent who set out to buy their daughter a sticker, and came home having learned the word *rate* from her.

SOURCES 1 Bikkuriman sales volumes, the discarded-chocolate controversy, the three-per-customer limits and the children's rarity hierarchy as documented in Japanese reference sources and in Nikkei xTrend's account of 30 June 2023. The Japan Fair Trade Commission recommendation of August 1988 and Lotte's November 1988 compliance — the cost cut, the increased print run and the equalised assortment rate — from the same reference record. 2 Japan's anti-scalping statute covering event tickets only; Mercari's shift from its 2021 principle to pre-emptive listing bans, the Switch 2, Chiikawa and Pokémon restrictions including the 16 September 2026 start, and the Pokémon Company's identity verification, as reported by Japanese trade press. Mercari's FY2026 results call of 6 August 2026 for the gross merchandise value and growth rate. 3 The copycat-versus-Bondoro distinction from MoneyPost WEB, 1 July 2026, and from the second page of the Daily Shincho report of 5 September 2026. The 4 and 17 September lottery sales listed by Bonbon Navi, read 18 September 2026. The order book and tie-ups from Advertimes' report of 18 June 2026. Takafumi Horie's remark posted 29 October 2025 and quoted in Advertimes, 10 August 2026. The twice-daily purchase peak, the buyer split and the shopping circuit from Research and Innovation Inc.'s release of 6 February 2026.

MERIDIAN

THE LAST WORD

The narrow part

A week of small things deciding large ones. Two stretches of water off Yemen and Iran are now an input to the policy rate in Tokyo, and the governor of the Bank of Japan said so out loud. Huawei, unable to buy a lithography node, attacked the wires instead. Twelve minutes a week of hard effort carried the same association as a hundred and fifty of moderate.

And the two we liked best, because nobody sells them to you. In a pepper mill the entire product is a burr you never see; the body is a handle. In watercolour the paint gets the shelf space and the argument, and what actually decides whether you can lift a wash is a sizing agent a few microns thick.

Then Japan, where a five-hundred-yen sticker acquired a published rate, and a playground economy that had run on asymmetric generosity for two generations was repriced by the smallest possible piece of infrastructure: a number, on a website, that the children could read.

Twelve Desks · One World MERIDIAN · No. 107 · 19 September 2026

A note on method. This edition was researched with working network access, and the load-bearing figures in it were read on the pages that carry them — the Bank of Japan's decision document, the US Treasury's yield curve, Huawei's keynote, Apple's and Esoteric's product pages, the journals themselves, and each maker's own price list. Where a host refused this desk — Reddit throughout, several paywalled publishers, two Singapore retailers and the European Commission's own border pages — that is said on the page in question, and nothing has been attributed to a source we could not open. Every hero image was fetched and looked at before it was captioned. Figures we could not stand behind are flagged in place or absent. Full source lists appear on the desk pages.