

THE CURRENT BRIEFING OF

MERIDIAN

Twelve desks. One world.

WEDNESDAY 23 SEPTEMBER 2026



THE PIECE IS FINISHED WHEN IT IS COOL, NOT WHEN YOU STOP WORKING IT (P15)

THE WORLD

A treaty signed on Tuesday has no end date — and Article XI binds a country that does not yet exist

THE MACRO DESK

Two prices for the same barrel of oil, twenty-two dollars apart. One of them is a cargo

THE LONG READ

China is naming hospital clinics after what the patient would say. Someone has worked out what the signboard rents for

THIS ISSUE • NO. 111

Contents

p3–4 The World — a treaty with no end date, and a clause that binds a country not yet founded

p5 Singapore — a tribunal decides a resignation was a dismissal, and prices the harm separately

p6 The Atelier — hard-shell luggage: the shell is not what breaks

p7 Technology — 216 gigabytes, and nobody will say what kind of memory it is

p8 The Home Lab — MIT-licensed, and the nine-billion clears a number we spent a page on

p9 The Kit — two earcups that are separate radios, and no way to plug them in

p10 The Good Life — a numbered watch that is not a limited edition

p11 Screen & Sound — no critic filed a negative review; the Metascore is eighty-one

p12–13 The Family Desk — puberty at ten, read against executive function at sixty

p14 The Macro Desk — the futures board says the crisis is over; the cargoes disagree

p15–17 The Rabbit Hole — glass: the piece is finished when it is cool

p18 Fitness — the metric your ring displays predicted nothing at all

p19–20 The Diary — two things close before October does; four openings, two judgeable

p21–23 The Long Read — China: the clinic name has become the unit of demand

FROM THE DESK

“Glass, china, and reputation are easily crack’d, and never well mended.”

Benjamin Franklin, *Poor Richard’s Almanack*

Everything in this issue is about the slow cool: what a thing is like after it has set, and how little of that is decided while you are still working it.

THE STRIP • AT TUESDAY’S CLOSE

US 10-year (par)	4.96%	unchanged
US 2-year (par)	4.71%	▼ -5bp
Brent (Nov)	\$98.65	▼ -1.68%
Gold (spot /oz)	\$4,358	▲ +0.34%
Bitcoin	\$86,253	▼ -0.41% 24h
VWRA (All-World)	\$194.96	▲ +0.28%
AGGU (global agg)	\$5.76	▼ -0.10%
STI	5,723.76	▲ +0.86%
USD/SGD	1.2747	▼ -0.02%

Tuesday 22 September closes. Treasury figures are the par yield curve, not the market close; both run on **p14**.¹ Bitcoin and the currency are spot at 05:15 SGT. Not advice.

WHAT THIS ISSUE KEEPS FINDING

The thing is finished, and then it sets. A glassblower’s piece looks perfect leaving the flame and shatters on a shelf a week later (**p17**). An agreement signed on Tuesday has no end date and binds a state that does not yet exist (**p3**).

And what you measure it with. Two oil prices for the same barrel, twenty-two dollars apart (**p14**). Two film scoreboards fourteen points apart with no negative review under either (**p11**). A clinic named for what the patient would say, not what the organ is (**p21**).

SOURCES 1 Par yields read from the US Treasury’s own daily file for 22 September 2026; other levels as set out with their origins on **p14**. Gold is spot, not a fix, and is not a record — it sits roughly 22 per cent below the January high. Brent is the November contract. A note on method: the two-year par figure here is 4.71 per cent, and an earlier reading of the Treasury’s rendered table gave 4.43 — that table carries an extra column, and a misaligned read shifts every maturity by one. The raw file is the authority, and it agrees to the basis point with the Federal Reserve’s published series.

THE ARCTIC · SIGNED IN NEW YORK ON TUESDAY, PUBLISHED THE SAME AFTERNOON

Article XI binds a country that does not yet exist.

Two new American defence areas in Greenland, a veto over who may invest there — and a clause obliging Denmark and Nuuk to guarantee the choices of a state neither of them governs.

They signed at half past ten on Tuesday morning at United Nations headquarters, and by the afternoon the text was on a White House server.¹ That is what changed. For four days the only public account had been the President's, and his word for it was *control*: the deal “gives the United States permanent control over security, and all other needs, in Greenland.”² Copenhagen and Nuuk said nothing had moved, and No. 109 reported the accord as confirmed by three governments and published by none — so neither claim could be tested.

The document never uses the word *control*. Article VI lets American aircraft “fly over and land in any territory in Greenland” and American vessels move undersea in its waters, “without restriction except as mutually agreed.” Article IX bars non-NATO states from installations there; Article X bars non-NATO, non-EU investors from control in sensitive sectors. Article IV reopens **Narsarsuaq**, abandoned in 1958, and **Mestersvig**.¹

“Permanent Nature”

THE HEADING ARTICLE XI GIVES ITSELF. “THIS AGREEMENT DOES NOT HAVE AN END DATE.”¹

Then the clause with no 1951 analogue. Should Greenland become independent, Denmark and Greenland “shall together ensure that the independent Greenlandic state will (i) agree to remain in NATO ... and (ii) affirmatively assume ... all of the rights and obligations of the Kingdom of Denmark specified in this Agreement.”¹ The preamble recognises Greenlanders as “a people ... with the right to self-determination.” Eleven articles later the text disposes of what they might decide with it.

Nothing here obliges Greenland to stay Danish. Everything here obliges whatever it becomes.

COPENHAGEN — STATSMINISTERIET

Frederiksen: an agreement “which strengthens our common security in the Arctic,” and one “that can last forever.” Sovereignty intact. Unmentioned: that Article XI outlives the Kingdom.³

BEIJING — CHINA DAILY

Leads on process: the provisions “were not revealed in advance.” Nothing new, and done in the dark. Unmentioned: Articles IX and X, which are pointed at Chinese money.⁴

THE CHATTER · THE SPECIALISTS, AND THE HOLE IN THE COVERAGE

Reddit was unreachable all session; nothing here comes from it. **Marc Jacobsen**, of the Royal Danish Defence College, reads the ceremony itself as the product: “It’s important for him to show that he’s a good businessman, that he kind of forced the pressure on Denmark and Greenland into something new.”⁶ The absence is louder: no Greenlandic-language outlet was reachable, and **Naleraq**, the opposition and the party of independence, has said nothing about the clause that lands on them.

THE SCEPTIC · THE CASE THAT TUESDAY CHANGED NOTHING

Take the French reading seriously, because Beijing shares it. **Mikaa Blugeon-Mered** told France 24 the agreement “doesn’t really give the US many more rights”: the 1951 text already let the Americans operate across Greenland.⁵ Three things resist that. Article XI has no 1951 analogue; neither does Article X’s screening duty; neither does the missing end date. And none of it binds yet: it takes effect only after “the necessary national parliamentary procedures”¹. The vote in Nuuk is the one to watch.

SOURCES 1 The agreement itself, read in full on [whitehouse.gov](https://www.whitehouse.gov), 22 September 2026: “Agreement ... to Amend and Supplement the Agreement of 27 April 1951.” Every article number, heading and quoted clause here is taken from that text, Article XI in full. 2 Truth Social, 18 September, as reported by Arab News, 19 September, read at source. 3 Statsministeriet release, dated 18 September, updated 20 September, read at stm.dk; “last forever” from Al Jazeera’s account of the signing, 22 September. 4 China Daily/Xinhua, 22 September, read at source. 5 France 24, *Perspective*, 22 September, read at source; Blugeon-Mered holds the Canada Research Chair in Clean Hydrogen at the Université du Québec à Trois-Rivières. 6 PBS NewsHour, 19 September, read at source, for Jacobsen, Gad and the Joelsen post. Reuters and DW refused this desk on every path; nothing here is attributed to either.

BRUSSELS · A MIDNIGHT DEADLINE, AND WHAT WAS TRADED TO BEAT IT

The EU bought three years of sanctions with two names.

The whole listings regime would have lapsed at midnight on Tuesday. It was renewed for thirty-six months instead of the usual six — and Alisher Usmanov and Mikhail Fridman came off the list.

The regime governing EU sanctions on Russia renews every six months, which means twice a year any one of twenty-seven governments can hold the whole list hostage. On Tuesday the Council removed five of the next six hostage moments, rolling the listings over for **thirty-six months**. The count locked in is itself disputed — Euronews says “over 2,600” individuals and entities, other accounts roughly three thousand — and this desk prints both rather than pick.¹

6 months → 36

THE RENEWAL, BEFORE AND AFTER — FIVE FEWER OCCASIONS ON WHICH ONE MEMBER STATE CAN THREATEN THE LOT.

Latvia blocked it on Monday, then abstained on Tuesday. Its prime minister **Andris Kulbergs** had put the objection plainly: “The price for that should not simply be removing two individuals from the sanctions list.”² France championed Usmanov; Luxembourg backed Fridman, who has a sixteen-billion-dollar asset-freeze claim against it. Zelensky, who gains from the lock and objected to the price, said the world needs “not a political delisting, but timely and unconditional rollover.”¹

A procedure that could be taken hostage twice a year has been made hostage-proof by paying the ransom once.

MOSCOW'S REFINERY IS NOT ALIGHT. IT IS SHUT.

No. 110 reported the drone wave that hit the Gazprom Neft plant at Kapotnya. The delta is what followed. Two units were struck: the AVT-6 crude column, **53 per cent** of plant capacity, and the Euro+ complex, **47 per cent**. With both down there is nothing left running.³

The cleanest evidence is commercial, not photographic. From Monday 21 September the refinery's fuel stopped being offered on the St

Petersburg exchange.³ A plant selling nothing is not a plant with a damaged outbuilding, which is what the mayor described. Repairs are put at several weeks. Kyiv has meanwhile offered to stop striking Russian oil if Russia stops striking Ukrainian energy — an offer Macron carried to Trump on Monday.⁴

THE CHATTER · WHAT THE ANALYSTS ARE COUNTING

The Institute for the Study of War's 21 September assessment has Russia expanding production of composites for Geran-type strike drones and building underground at Alabuga; it logs 172 Russian drones against Ukraine overnight and no territorial change either way.⁵ One claim needs care: **Fire Point**, the Ukrainian manufacturer, says its FP-1 drones “completely shut down the refinery” — a vendor taking credit for its own product.

THE SCEPTIC · AGAINST THE FORTY-FIVE PER CENT

Ukraine's General Staff says strikes have removed **45 per cent** of Russian refining capacity.⁶ It is a belligerent's figure, and **Sergey Vakulenko** of Carnegie Politika has taken the method apart. Nameplate capacity is 327m tonnes a year against roughly 270m of actual output, so “at least 22% of the country's total capacity is always idle”. And damage is not proportional: losing one of two identical columns yields “a 30 percent drop — not 50 percent.”⁷

SETTING THE RECORD STRAIGHT

No. 110 put Russian H1 2026 refining earnings at about **US\$12.3bn**. The underlying Rosstat figure is **938bn roubles**, which the source reporting it converts to **US\$11.1bn** — a difference of rouble rate, not of fact. Hold the rouble number. The growth rate, +16.2 per cent, is unchanged.⁷

IN BRIEF

Yemen. UNHCR warned on 22 September that displacement could rise from over 130,000 to more than **230,000**. Its Yemen operation is **18 per cent funded**.⁸

Washington. Xi Jinping's state visit begins today, his first in nearly three years. Global Times previewed it on “strategic stability” and mentioned neither Taiwan, tariffs nor the Arctic.⁹

SOURCES 1 Euronews, 22 September, for the 36-month renewal, the listing count, the French and Luxembourgish positions and the Zelensky quotation; the Council's press server returned 403 on every path, so the decision text was not read here and a second outlet confirmed the duration. 2 Meduza, 22 September. 3 Kyiv Post, 20–21 September, for the AVT-6 and Euro+ shares and the withdrawal of the plant's fuel from the St Petersburg exchange; accounts of the plant's distance from the Kremlin differ and none is printed. 4 Kyiv Post, 21 September, for the Macron–Trump meeting and the conditional moratorium. 5 ISW's 21 September assessment via Kyiv Post. 6 Kyiv Independent, 21 September; single-sourced to Ukraine's General Staff and flagged in the text. 7 Carnegie Politika; the Rosstat half-year figures via Euromaidan Press, 17 September. 8 Al Jazeera, 22 September. 9 Global Times, 22 September. All read at source.

THE BENCH · A TRIBUNAL DECIDES WHAT A RESIGNATION WAS

Nine thousand nine hundred dollars for the lost income, and nine thousand nine hundred for the harm.

The Employment Claims Tribunal ruled on Tuesday that a woman who resigned had in fact been dismissed — then doubled the award to price what the dismissal did to her.

The claimant joined in April 2021 on S\$3,300 a month. A new director arrived in March 2023. She was hospitalised that September at thirty-six weeks, gave birth in October, and resigned that December. On Tuesday the Employment Claims Tribunal held that she had not resigned at all.¹

The mechanism is in the Employment Act's own definition, which counts as a dismissal "the resignation of an employee where the employee can show ... that the resignation was not voluntarily but was compelled by the conduct ... on the part of the employer." Tribunal Magistrate **Lee Li Choon** found she "did not resign voluntarily but was compelled to do so", and that dismissing her for exercising her maternity rights "amounted to discrimination ... on account of her pregnancy and maternity leave."¹

S\$19,800

THE AWARD, AND HOW IT WAS BUILT — S\$9,900 — THREE MONTHS' GROSS — FOR LOST INCOME, AND THE SAME AGAIN FOR HARM: A FIFTY PER CENT UPLIFT, STATED AS SUCH.¹

The uplift is the part to read twice. Lost earnings are arithmetic. The second S\$9,900 is a judgement about conduct, explained in terms courts rarely use: the dismissal came "at perhaps the most profound and yet vulnerable moment in a woman's life" and was "both degrading and deeply insensitive."¹ The employer is anonymised. The reasoning is not.

A payslip can price three months. It cannot price the month a child is born.

THE CHATTER · WHERE THE NUMBER LANDS

Read against No. 110's labour page the scale is honest rather than triumphant: second-quarter retrenchments of 4,620, the highest since the end of 2020, and re-entry within six months down from 60.7 to 54.9 per cent. Against that, one award of S\$19,800. Reddit was unreachable and HardwareZone's board sits behind an age wall, so no local forum voice is quoted.

THE SCEPTIC · WHAT A FIFTY PER CENT UPLIFT IS WORTH

Three months' salary deters nobody, and doubling it still leaves a figure a mid-sized employer treats as a cost rather than a risk. The claimant also carried the burden of proving her own resignation involuntary — hard to evidence. The counter is that the uplift is precedent, not penalty: a published decision naming a fifty per cent addition gives the next magistrate a figure to reach for.

THE DATA UNIT THAT WOULD HAVE SEEN IT

The Ministry of Social and Family Development briefed media on Tuesday on **SAFER**, stood up quietly in June, which cross-reads data that previously sat in separate agencies. In two months it drove intervention in **70 cases**. The counterfactual is the reason it exists: had it run in 2020, Megan Khung's case would have been flagged **at least three times, up to five months before she died**.² One briefing write-up only.

ALSO THIS WEEK

The haze broke, and the Act had nothing to do with it.

Tuesday's update put the 24-hour PSI at **59–77** and one-hour PM2.5 at **11–21** micrograms — Band 1, against No. 108's PSI of 108. What changed was rain.³

And the bus rules bite today. Up to **S\$500** for feet on seats, loud audio or eating; up to **S\$5,000** for obstruction. The mechanism is the interesting part: enforcement is devolved to the operators, with inspectors empowered to issue notices while **bus captains expressly cannot**. A driver may ask you to leave. He may not book you.⁴

SOURCES 1 *JKP v JKQ* [2026] SGECT 13, Employment Claims Tribunal, Tribunal Magistrate Lee Li Choon, decision dated 22 September 2026, read in full at elitigation.sg; the statutory definition is from the Employment Act 1968, and every quoted phrase is from the grounds of decision at the paragraphs cited there. 2 Mothership, 22 September, read at source, reporting the ministry's media briefing; msf.gov.sg's media room served no content on any path tried, so this desk could find no primary release and flags the item as single-sourced in the text. 3 National Environment Agency Haze Situation Update, 22 September 2026, as of 6pm, read at nea.gov.sg. 4 Land Transport Authority news release, 21 September 2026, read at lta.gov.sg, for both sets of regulations, the penalty ceilings and the enforcement provisions. Reuters, the Straits Times, CNA and SingStat's web pages all refused this desk; nothing on this page is attributed to any of them.

BUY IT ONCE · HARD-SHELL LUGGAGE AND CARRY-ONS

The shell is not what breaks. Turn the case over and look at the wheels.

Screws or rivets is the whole test — and in this climate the part that dies first is a polyurethane tyre quietly hydrolysing in a cupboard between trips.

Start with the material, because the marketing does. **Polycarbonate** flexes under impact and returns to shape — provided it is virgin resin. Regrind has shorter chains, goes brittle, and cracks at a corner or a wheel mount instead. The tell is transparency: a maker who names **Makrolon**, Covestro's polycarbonate, is telling you the grade. One who says “premium polycarbonate” is not.¹

Aluminium does the opposite — it deforms permanently. Not worse, different: aluminium can be hammered true, polycarbonate cannot be uncracked. **Curv** is the interesting one, self-reinforced polypropylene woven from stretched tape, so fibre and matrix are the same polymer and nothing can delaminate. It is also not Samsonite's: Curv is **Propex**'s trademark, pressed at Oudenaarde.²

1.7kg → 5.1kg

EMPTY CABIN CASE AGAINST A 7KG ALLOWANCE — A CURV COSMOLITE AGAINST AN ALUMINIUM CARRY-ON³

None of which is the failure point. **YKK** — the zip maker, testifying against its own industry's marketing — states that “replacing broken zippers is by far the most common task at luggage repair facilities.”⁴ Then wheels, then handles. One owner's catalogue across three aluminium cases: hinges misaligning, seals failing, metal torn near the wheel assemblies.⁵ The shell is not on the list.

Nobody here publishes drop-test data, wheel-cycle data, or how long they will stock a wheel.

So buy the wheels. The connoisseur's name is **Hinomoto** of Nagoya, whose SILENT RUN casters use a **Lisof** compound developed with Mitsubishi Chemical and sold on *resistance to hydrolysis*.⁶ Polyurethane tyres hydrolyse — tacky, then crumbling — and warm wet storage speeds it up. A case sitting eight months in a Singapore flat between

trips is in exactly the wrong place. At the counter the test is cruder: turn it over. Riveted wheels mean a repair shop has to drill it.

THE CHATTER · WHAT THE PEOPLE WHO FLY A LOT KEEP

Reddit and FlyerTalk were both unreachable; nothing here comes from either. **Pack Hacker** calls Hinomoto casters “some of the smoothest we've tried” — the one component the gear press names by its maker rather than the brand on the lid. On the Rick Steves forum, unprompted: “I'm sticking with my Briggs & Rileys,” and “I adore my Muji small hardside bag — countless flights and it's still going strong, 7+ years in.”⁷

On origin, received wisdom is wrong. LVMH bought 80 per cent of **Rimowa** in 2016, but the cases are still made in Rimowa's own four plants, aluminium still riveted by hand in Cologne.⁸ The logo tax is on the badge, not the bag's origin. **Titan** publishes the cleanest claim going — every case made in Germany since 2020, ten-year cover, “except Mervus, Prior and X-Ray.”⁹

THE SCEPTIC · WHERE THE MONEY STOPS WORKING

Above roughly **US\$700** a hard case stops getting functionally better. Hinomoto casters turn up on cases costing a third of a Rimowa, and Briggs & Riley's lifetime cover at US\$695 beats anything above it at any price. Past that line you buy material, provenance and appearance — and not durability. Read the warranties for the exclusions: Rimowa's excludes cosmetic wear, and on an aluminium case a dent is cosmetic.¹⁰

NEXT ON THE BENCH — PICK THE NEXT CATEGORY

Coffee grinders — burr geometry is most of the cup.

Tea ware (teapots, kettles & gaiwan) — zisha, tetsubin, and what the clay does to the brew.

Merino base layers & travel tees — micron grade against the tropics.

SOURCES 1 market.com's polycarbonate guide for the shell bands and the virgin-versus-regrind mechanism; Makrolon is Covestro's polycarbonate. 2 Propex's own composites site for Curv's ownership and the Samsonite relationship; the Oudenaarde line from the engineering contractor's own project page. 3 Zero Halliburton's own product page for 5.1kg; the 1.7kg Cosmolite figure is from retail listings, not Samsonite, and is flagged. 4 YKK Americas. 5 Live and Let's Fly, 28 July 2024. 6 Hinomoto Jomae's own site, including the hydrolysis claim. 7 Pack Hacker; the owner verdicts from the Rick Steves forum, quoted as posted. 8 The four-plant list from the ERP vendor's account of its Rimowa installation; rimowa.com returned 403 on every path, so no Rimowa price or warranty clause here comes from Rimowa. 9 titan-bags.com. 10 The briggs-riley.com, uk.tumi.com and lojel.com warranty pages, quoted verbatim. All read at source.

HANGZHOU · AN ACCELERATOR ANNOUNCED AT THE APSARA CONFERENCE

Two hundred and sixteen gigabytes, and nobody will say what kind of memory it is.

Alibaba's chip arm put the Zhenwu V900 on a stand on Tuesday. The number on the placard is the one the American export rule was written to catch — if it is the memory the rule means.

T-Head, the silicon arm Alibaba spun up in 2018, unveiled the **Zhenwu V900** on 22 September: a training-and-inference accelerator with **216GB** of on-package memory, **1,200GB/s** chip-to-chip interconnect, and native FP8 and FP4. Mass production is claimed for the first quarter of 2027 — fifteen months out. More than a thousand can be lashed into one system, and a cluster scales to half a million parts, which Alibaba says supports models of five to ten trillion parameters.¹

2 GB/s per mm²

THE CLAUSE THE NUMBER RUNS INTO — ECCN 3A090.C CONTROLS HBM ABOVE THAT MEMORY-BANDWIDTH DENSITY — “REGARDLESS OF WHETHER THEY CONFORM TO THE JEDEC STANDARDS.”²

Here is the mechanism worth learning. The controls added in December 2024 do not catch accelerators alone; they catch the *memory*. The rule defines memory-bandwidth density as package bandwidth divided by package area in square millimetres, and controls high-bandwidth memory above two. Stack enough fast memory close enough to the die and you are inside the rule by arithmetic, not by intent.²

So: 216GB of what, made by whom? T-Head's own placard says 显存 — “video memory” — not HBM. One outlet wrote it up as HBM; that is an inference, and neither trade outlet names the foundry, the node or the memory supplier.¹

A control written around a machine is being tested by a package. It measures bandwidth per square millimetre; it does not ask who stacked it.

THE CHATTER · AND AN HONEST BLANK

There is no community reaction to report, and inventing one would be worse than the silence: a search of the Hacker News index for Alibaba chip stories after 21 September returns nothing. Vendor claims are all that is on the table, and the outlet most sympathetic to the “export controls backfired” reading concedes in its own text that “these are vendor-reported figures, not independently benchmarked.”³ Even the predecessor's installed base is contested — 400-plus external customers by one account this week, 650-plus by another.¹

MEANWHILE, THE RULE CONGRESS WANTS NEXT

Representative **Michael Baumgartner** said on 21 September that Congress should pass the MATCH Act despite Dutch and Japanese objections, conceding the bill “has some notoriety” and that foreign firms “would rather not have the Match Act be passed.”⁴ The bill text names five Chinese firms — CXMT, Hua Hong, Huawei, SMIC and YMTC — gives Commerce **150 days** from enactment to issue regulations, and defines a chokepoint tool by a market test: equipment no country of concern makes in sufficient volume at comparable capability to the global leader.⁵

THE SCEPTIC · INDUSTRIAL POLICY IN A SECURITY JACKET

Mark Dalton of the R Street Institute argues the case against with numbers rather than adjectives: affected US suppliers have seen “a \$130 billion aggregate decline in market capitalization” and “an 8.6 percent drop in revenue”; Chinese domestic equipment adoption rose from 25 per cent in 2024 to 35 in 2025; and since semiconductor firms put 17.7 per cent of revenue into research, every lost revenue dollar costs about eighteen cents of R&D. His political reading is that Micron is the bill's primary advocate and is losing share to ordinary competition.⁶

Set that beside Baumgartner's own remarks, which named Micron's New York operations. The lobbyist and the critic agree on who is asking; they disagree about whether that settles it.

SOURCES 1 TechNode, 22 September 2026, for the V900's specifications, the cluster claims and the predecessor's installed base; the placard wording was read from the photograph accompanying that report. TrendForce, same day, corroborates the specifications and gives the differing customer count. CNBC also covered the launch but returned 403, so nothing here rests on it. 2 Bureau of Industry and Security interim final rule, published 5 December 2024, read in full text at federalregister.gov, for ECCN 3A090.c, the memory-bandwidth-density definition and the licence-exception threshold. 3 Forkast, including its own caveats quoted here. 4 Export Compliance Daily, 21 September 2026. 5 S.4281, the MATCH Act, read in full at govinfo.gov; congress.gov returned 403. 6 Mark Dalton, R Street Institute, 30 April 2026. All read at source.

THE HOME LAB · A TRILLION-PARAMETER RELEASE, AND THE NINE BILLION THAT MATTERS

MIT-licensed, and the small one clears the number No. 108 spent a page on.

Xiaomi shipped MiMo-V2.6 on Monday. The distilled nine-billion scores 61.1 on SWE-bench Verified — and llama.cpp merged support the same week.

Start with the licence, because No. 109 ended on one. That issue reported Qwen-Image 2.1 arriving under a research licence — weights you could download and not use. All three MiMo-V2.6 repositories carry **MIT** on the model card: commercial use, modification and redistribution, with attribution, and no acceptable-use rider.¹ That is the delta, and for a reader who builds things it is the whole story.

The family is three. **Pro** is a sparse mixture of experts, 1.02 trillion parameters total with 42 billion active. **Flash** is 309 billion total, 15 billion active. Both take a million tokens of context and handle text, image, video and audio. **Distill-Qwen-9B** is nine billion dense, trained on a mixture the card describes as “77.4B total tokens, including 27.2B loss-bearing tokens.”¹

61.1

THE NINE-BILLION, ON SWE-BENCH VERIFIED — NO. 108
DISSECTED AN AGENT SCORING 60.8 ON THE SAME
BENCHMARK. THAT ONE DID NOT FIT ON A DESK.¹

What the reader can actually run is unusually concrete. A llama.cpp pull request adding MiMo-V2.6 conversion support merged on **22 September**, and the project is publishing its own quantisations: Flash at **167GB** in MXFP4, **126GB** in Q2_K, plus a vision-and-audio projector and a speculative-decode draft. The 9B fits a single card comfortably — **9.53GB** at Q8_0.²

A hundred and twenty-six gigabytes is six over the line the community drew for itself.

THE CHATTER · THE THREAD, BY HANDLE

The Hacker News thread ran to 469 comments. **alfiedotwtf** wants a Q3 that fits 128GB of shared memory and calls “95–115GB the sweet spot” — which the 126GB Q2_K misses by six. **embedding-shape** reports the Flash variant “reasons more and for longer than Pro”. **wren6991** explains why the listings show wrong parameter counts: packed four-bit weights are “often identified as u8 byte arrays”, so the platform reports file bytes, not parameters.³

On Apple silicon there is a measured number rather than a claim: a four-bit conversion of Flash sustains **59.4 tokens a second** decoding on an M3 Ultra, peaking near 167GB of unified memory — so a 256GB machine in practice. Nobody has published the same run on the M5 Ultra of No. 110.⁴

THE SCEPTIC · READ THE VENDOR’S OWN CHART CAREFULLY

Xiaomi publishes a comparison chart that shows its model losing — on DeepSWE the Pro sits at 71.9, behind three rivals — and the community has praised it for the candour. Read the footnote and the candour gets complicated: “Hollow bars are placeholders; evaluation pending.” Several competitor bars in that chart are not measured results at all.⁵

A second wrinkle turned up on cross-reading. The chart gives ExploitBench at 43.8 for the Pro; the model card gives 47.9. Same vendor, same week, two numbers. And most of the suite is Xiaomi’s own — one benchmark is labelled “internal” on the chart itself — with a single independent index as the outside check. On that index the Pro ranks first of 114 models. It is also an index whose own weighting, one commenter noted, “had to be radically changed when Astra came out.”⁶

STILL ON THE SHELF

Frontier inference on one card — claimed, not shipped.

Tim Dettmers posted on 21 September promising quantised inference at “1.5 bits per weight”, a 35B model at 450 tokens a second, and a 125-billion-parameter model on a single 24GB GPU. The repositories are not up: the lab’s public organisation shows nothing newer than 17 September, and he says himself he is “not going to give away everything before the open-source week starts.” Watch it; do not plan around it.⁷

SOURCES 1 The three MiMo-V2.6 model cards on Hugging Face, read at source, for the MIT licence line, the architecture, the context length and every benchmark quoted; the vendor’s own site is a JavaScript shell and served nothing readable. 2 The llama.cpp pull-request list and the ggml-org GGUF repositories for the merge date, the quantisation sizes and the projector and draft files. 3 Hacker News item 49792730, read through the site’s API; handles as posted. 4 The MLX conversion’s model card for the sustained-decode figure and memory ceiling — measured on an M3 Ultra, stated because it is not the current part. 5 Xiaomi’s own comparison chart, downloaded and read, including the placeholder footnote. 6 The independent index’s model page; the commenter quoted is from Hacker News item 49796660. 7 Tim Dettmers’ post of 21 September and the laboratory’s public repository listing. All read at source.

THE KIT · ANNOUNCED TUESDAY, IN SHOPS ON THURSDAY

Both earcups are separate wireless devices. The headband carries one power cable.

The Beats 360 is built so the parts that always die — pads, headband foam — are consumables you can buy. It is also the first flagship here with no wired path at all.

Orders opened on 22 September; the headphone reaches shops on the 24th in more than a hundred countries, at **US\$349.99** and **S\$429**.¹ The architecture is the interesting part and it is not the noise cancelling. Each earcup carries its own custom chip, its own battery, its own antenna and its own wireless circuitry; the slim headband runs a single power cable between them.² They are two radios that happen to be joined.

What that buys is repairability of the right components. Ear cushions and headband cushions are user-swappable, sold as kits at **US\$49.99** — with an introductory price of **US\$9.99** to 31 October.¹ The two things that always fail first on a headphone are foam and fabric, and they have been made into stock.

294.5g · 32h

WEIGHT, BATTERY, MICROPHONES — NINE MEMS MICROPHONES, EIGHT SERVING NOISE CANCELLATION; BLUETOOTH 5.3; IPX4; FIVE MINUTES TO AN HOUR OF PLAY.³

The noise-cancelling claim is stated precisely enough to quote precisely: “up to 1.75x more total noise cancellation than Beats Studio Pro,” measured to IEC 60268-24.¹ That is a vendor measuring against its own previous model by a named standard — better than most, and still not an independent laboratory.

The most repairable flagship they have made is also the one you cannot plug into anything.

THE CHATTER · THREE REVIEWERS, PUBLISHED THE SAME DAY, CONVERGING

No owner sentiment is quoted: the forum thread returned 403 and Reddit was unreachable, so what follows is named reviewers only. **Billy Steele** at Engadget scored it 8.2 and found the battery beat its rating — 47 per cent left after twenty-four hours — but flagged no wear detection, no auto-pause, and clamping force that “becomes uncomfortable” on larger heads. **Malcolm Owen** at AppleInsider gave it 4.5 of 5, objecting that USB-C is charge-only and that bass “creeps in noticeably at high volume”. **José Adorno** at BGR calls it “the best sound in any pair of Beats headphones I’ve ever used” and says thirty-two hours lags rivals offering fifty.⁴ All three arrive at the same complaint from different directions.

THE SCEPTIC · THE FALSIFIABLE BIT

Take the repairability story seriously and it runs into its own catalogue. A headphone with no 3.5mm jack, no wired listening and no lossless is a Bluetooth-only device for its whole life; the swappable-cushion longevity case ends wherever the radio does. A competitor at US\$499 sells a jack, lossless and thirty-five hours.⁵

And the customisation headline deflates on contact: the “250-plus combinations” figure circulating is four frames times eight cushions times eight headbands — an arithmetic product, not a catalogue. The honest test is scheduled and dated: **whether cushion kits are still stocked and listed after the US\$9.99 promotion ends on 31 October**. If they quietly leave the catalogue, the consumable-parts bargain was a launch price, not a design.

STILL ON THE SHELF

A firmware update on announcement day. Apple pushed build 4A262 to the Beats 360 on 22 September, up from 4A260, contents undisclosed.⁶

Sonos Ace Ultra — US\$499, announced 1 September, pre-orders live, nothing since. It answers the gap above: USB-C and a 3.5mm jack, lossless, thirty-five hours. One outlet only; the maker’s press room refused this desk.⁵

Meta’s Connect keynote runs today and tomorrow, with new glasses and a watch expected. Nothing is announced, so nothing is printed.

SOURCES 1 The maker’s own press release for the launch and ship dates, the US price, the cushion-kit pricing and promotion end date, and the noise-cancellation claim with its IEC standard. 2 MacRumors, 22 September, for the per-earcup architecture; the specification page confirms a custom chip in each earcup. 3 The Singapore product page for S\$429, and the specification page for weight, battery, microphones, Bluetooth and ingress rating. The release says nine microphones and one review says eight; the specification page reconciles them as nine total, eight serving the noise-cancelling system. 4 Engadget (Billy Steele), AppleInsider (Malcolm Owen) and BGR (José Adorno), all 22 September. The MacRumors forum returned 403 and Reddit was unreachable, so no owner sentiment appears here. 5 9to5Toys, 1 September; the maker’s press room returned 403, so this item is single-sourced. 6 MacRumors, 22 September. All read at source.

THE GOOD LIFE • TWO PANERAI RELEASES, PUBLISHED HOURS APART

A numbered watch that is not a limited edition, on a movement from 2016.

The caseback is engraved “Annual Production” over a blank plaque. On the same day the same house launched a 42mm three-hander at a twenty-fifth of the price.

The Luminor Tourbillon GMT Titano, reference **PAM01784**, was announced on 22 September at **€160,000**, available from December. Forty-seven millimetres of brushed grade-5 titanium, a hundred metres of water resistance, and — the figure that matters for a case that size — about **105 grams** all in.¹

The movement is the calibre P.2005/T: hand-wound, skeletonised, 277 components, 31 jewels, and **three barrels in series** holding a combined mainspring 1.36 metres long, which is where the **six-day** power reserve comes from. The tourbillon is horizontal and turns once every thirty seconds rather than the conventional minute. The GMT is a true one — the local hour hand jumps in whole hours without disturbing home time. Indices sit directly on the blue-treated movement plate, and the power-reserve indicator is lumed in green while the GMT hand is blue.¹

“Annual Production”

ENGRAVED ON THE CASEBACK, UNDER THE SAPPHIRE —
BESIDE A PLAQUE READING 00/00 AA. SERIALIZED, THEN —
AND NO LIMITATION NUMBER IS STATED ANYWHERE.¹

That phrase does the work. A limited edition tells you how many exist. An annual-production number tells you only that yours was counted. It is a watch that carries the furniture of scarcity without the commitment, and Panerai publishes no figure. This desk looked for one and could not find it.

One house, one Tuesday, two opposite arguments about what it is for.

THE CHATTER • THE OTHER RELEASE, AND THE SILENCE

Hours before the tourbillon, the same specialist title reviewed Panerai’s new **Luminor 42** collection — two steel references at **€6,500** and a coated one at **€7,500**, all permanent. **Brice Goulard**’s verdict runs the other way entirely: “it’s often a case of less is more... simple, clean, ultra-robust, and true to its origins.”² One house, one day, two opposite arguments about what it is for. On the tourbillon there is nothing to report from collectors: this desk searched the watch forums and found no thread on the reference. Timing, probably — but worth recording rather than filling.

THE SCEPTIC • WHAT EXACTLY IS NEW HERE

The calibre P.2005/T launched in **2016**. The case is the 47mm Luminor that has been in the catalogue for years. A version in Goldtech with the same blue-finished bridges already exists. On the evidence published on Tuesday, the novelty is an execution in titanium and a colour — a case-and-dial exercise on a ten-year-old movement, priced at a hundred and sixty thousand euros.¹

What would settle it in the watch’s favour: a stated limitation, a revised escapement or finishing specification, or an independent hands-on confirming the 105-gram figure and how 47 millimetres actually wears. None of those exists yet. Note also what this desk could not read — Panerai’s own product page returned a server error, so no dollar or Singapore price is printed, and the sterling figure circulating is unverified.

ELSEWHERE ON THE SHELF

The most valuable Cartier ever sold at auction is a quartz watch. The unique white-gold **Cheich** made **HK\$26.5m** (about US\$3.38m) with fees at Sotheby’s Hong Kong, against an estimate of HK\$8–40m — a record for the house and, as far as anyone has counted, for any quartz watch. It is a one-off made roughly twenty-five years after the four yellow-gold Dakar Rally originals; the first of those made a little over a million euros with fees in 2022. The buyer is anonymous, and the sale date is not stated in the account read here.³

SOURCES 1 Monochrome Watches, 22 September 2026, by Denis Peshkov, read at source, for the reference, the euro price, the December availability, the case and movement specifications and the lume layout; the caseback engraving and the 00/00 AA plaque were read from the photographs accompanying that piece. Panerai’s own product page returned HTTP 503 to this desk, so no price other than the euro figure is printed and no production number is asserted. 2 Monochrome Watches, same day, by Brice Goulard, read at source, for the Luminor 42 collection, its three references, their euro prices and both quotations. 3 Watches by SJX, 22 September 2026, by JX Su, read at source. The sale date is not given there and a secondary account placing it on 18 September could not be verified, so no date is printed.

SCREEN & SOUND · SIXTY MILLION, AND TWO AGGREGATORS FOURTEEN POINTS APART

Not one critic filed a negative review. The Metascore is eighty-one.

Zach Cregger's *Resident Evil* opened at \$60.2m on ninety minutes of runtime. The gap between the two scoreboards is the more interesting number.

It opened on Friday 18 September to \$60,152,878 across 3,684 screens — a per-screen average of \$16,328 — and by Tuesday had taken \$65.5m domestically and \$111.9m worldwide against a \$75m production budget.¹ The rest of that weekend's top five did not come close: the second film managed \$12.2m.

Then the scoreboards. Rotten Tomatoes has it at **95 per cent from 235 critics**. Metacritic has a Metascore of **81 from 47** — and, unusually, a breakdown of forty-five positive reviews, two mixed and **zero negative**.² The lowest score visible on that page is 88.

95% vs 81

SAME FILM, TWO SCOREBOARDS — 235 CRITICS AGAINST 47. A BINARY THUMB AGAINST A WEIGHTED AVERAGE — AND NOBODY VOTED IT DOWN.²

Fourteen points with no negative reviews underneath is not disagreement about quality. It is two instruments measuring different things: one counts how many critics liked it, the other asks how much. The answer the second returns — delighted, not awed — is the honest verdict, and it is the one the enthusiastic notices say too when you read past the headline.

The opening weekend was ninety-two per cent of everything it had taken by Tuesday.

THE CHATTER · WHAT THE HUNDREDS ACTUALLY PRAISE

Read the perfect scores and they agree on a register nobody expected. **Dominic Baez** in the *Seattle Times*: “It’s honestly shocking how funny this movie is — even more so that that humor doesn’t distract.” **Brandon Yu** in the *New York Times* calls it “a thriller that’s as funny as it is jumpy, with a remarkable number of jokes.” **Bob Strauss** in the *San Francisco Chronicle* credits “formal mastery... through what is essentially a creature spectacular.”³

Rotten Tomatoes’ own consensus lands in the same place — scares in a sleek runtime, comic timing providing “crucial levity”. A video-game adaptation is being reviewed as a comedy that happens to frighten people. The box office is not telling you that.

THE SCEPTIC · THE DISSENT IS TONAL, AND IT IS CONSISTENT

Searching deliberately for pans turns up almost nothing in the aggregated pools, which is itself a finding: with zero negatives on one scoreboard and 95 per cent on the other, perhaps a dozen rotten notices exist and the filtered page would not load, so this desk cannot name a single negative critic. That is a gap, printed as one.

The substantive objection is **Phil de Semlyen**’s at *Time Out*, three stars, filed on 16 September before the weekend: the film is “more goofy and good-looking than actually frightening”, “a handsome-looking and occasionally inspired addition to the funhouse zombie genre” that is “still a fun Friday night ride”.⁴ The structural version of the same complaint recurred across reviews: Cregger’s earlier films worked by withholding information and fracturing structure, and this one trades that for a linear shape — one level after another, which is what the source material is.

THE NUMBER TO WATCH

By Tuesday the opening weekend accounted for **91.9 per cent** of the domestic total — the front-loaded pattern horror always shows.¹ Everything hangs on the second-weekend hold, which prints this Sunday. A drop under sixty per cent would make it a genuine crossover; over seventy and the aggregator story was the story.

SOURCES 1 Box Office Mojo’s weekend chart for 2026W38 and The Numbers’ title page as of 22 September 2026, both read at source, for the opening gross, the theatre count, the per-screen average, the running totals and the production budget. Trade accounts giving a \$108.3m worldwide *opening* are a different measure from the \$111.9m cumulative printed here, and all three trade titles were paywalled to this desk, so nothing is taken from them. 2 Metacritic and Rotten Tomatoes, both read at source on 23 September; scores move as counts rise and are printed with their sample sizes for that reason. 3 The three quoted notices as published on Metacritic’s critic page, read at source. 4 *Time Out*, Phil de Semlyen, three stars, 16 September 2026, read at source.

THE FAMILY DESK · A PAPER PUBLISHED MONDAY, AND THE BAND IT ACTUALLY COVERS

Puberty timing at ten, read against executive function at sixty.

Ten thousand seven hundred and twenty-two children, plus a genetic causal test on six hundred thousand adults. The authors wrote their own headline: “limited clinical significance.”

The *Journal of Child Psychology and Psychiatry* published Dinkelbach and colleagues on 21 September, and it does something unusual: it runs the question twice, by two methods that fail in different ways. First an observational cohort — 10,722 children in the American ABCD study, measured at three waves. Then a two-sample Mendelian randomisation, which uses genetic variants as a natural experiment, drawing on genome-wide data for age at menarche in 632,955 women and executive functioning in 427,037 adults.¹

The age band matters more than the finding, so take it first. The cohort’s second wave is **12.0 ± 0.7 years** — a twelve-year-old cohort almost exactly. But puberty itself was scored at *baseline*, around age ten, by a parent-reported maturation scale. The exposure is measured two years before the age this desk is written for.¹

$\beta = 0.04$

THE ASSOCIATION, IN STANDARD DEVIATIONS — PICTURE-SEQUENCE MEMORY, GIRLS, 95% CI 0.02–0.06. WORD-LIST LEARNING: 0.05 IN GIRLS, 0.03 IN BOYS.¹

Earlier puberty was associated with slightly poorer performance on memory and word-list learning in both sexes, and on list sorting in boys only. The genetic arm agreed in direction: better executive function with later menarche, **0.005 per year** — with a confidence interval of 0.000 to 0.011, whose lower bound sits on zero.¹

A beta of 0.04 standard deviations is roughly half a point on a hundred-point standardised score. It is a population signal. It is not a parenting signal, and the authors say so in a discussion section they titled, in the paper itself, “**Small Effect Sizes – Limited Clinical Significance**”: the effect is “likely negligible for everyday functioning for puberty timing within normal ranges.”¹

A study that pre-empts its own moral panic is rare enough to be the story.

THE CHATTER · WHERE THIS LANDS AFTER THREE WEEKS OF SCREENS

This desk has spent three issues on phones — No. 108 on a “fifty per cent more effective” claim that was *d* –0.30 in 205 children and –0.07 in all 1,793; No. 109 on a first-phone paper against a null trial; No. 110 on bans that coincided with in-class use rising from 28 to 42 per cent. The through-line is not phones. It is that the published effect sizes in this literature are small and the public conversation is loud.

Worth noting what this paper is *not*. The causal arm’s outcome was built from tests given to adults aged roughly forty to sixty-nine. So the defensible claim is “earlier menarche, very slightly lower executive function in middle age” — not “worse at school at twelve.” The authors name that mismatch themselves.¹

THE SCEPTIC · FOUR THINGS WRONG WITH IT, THREE NAMED BY THE AUTHORS

They flag the first three. Puberty was measured indirectly — parent-reported maturation in the cohort, recalled menarche and voice-break in the genetic arm — and “neither approach employed the gold standard of Tanner staging by a clinician.” Scoring at age ten “may be too early to fully capture gonadal maturation in boys”, since the scale may be reading adrenal changes that start at six to eight. And the causal outcome is not the cohort outcome: the adult tests “are not directly comparable.”¹

The fourth is ours. A causal estimate whose interval touches zero, drawn from European-ancestry genetic data, generalised to a multi-ethnic American cohort, is a fragile instrument — and the published paper and its preprint do not even agree on the sample, 10,722 against 10,174. The published figure is the one printed here.

SOURCES ¹ Dinkelbach L, Serio B, Valk S, Weisner FE, Bläsche LM, Peters T, Hinney A, Schmidt B, Hirtz R, “Puberty timing and cognitive functioning: insights from the Adolescent Brain Cognitive Development (ABCD) study and Mendelian randomization,” *Journal of Child Psychology and Psychiatry*, published online 21 September 2026, DOI 10.1111/jcpp.70220. The published article is paywalled; the abstract was read at source on PubMed for the sample sizes, betas and confidence intervals, and the authors’ own methods, discussion and limitations were read in the full preprint PDF, where the section heading and the quoted limitation sentences appear verbatim. Note that the published abstract and the preprint report the same associations with opposite signs, because one scales per standard deviation of *earlier* puberty; the finding is identical and the published convention is used here. General guidance, not medical advice.

THE TWEEN YEARS • A THIRTY-DAY DIARY STUDY, THREE REPORTERS PER FAMILY

It is not the arguing that costs everyone sleep. It is disagreeing about whether there was any.

A hundred and seven families kept nightly records. The worst sleep in the house belonged to those who could not agree on what the parents' relationship was like.

Child Development published this on 22 September, and its design is the reason to read it. 107 American families kept daily diaries for thirty days, with three reporters in each — adolescent, mother and father — each rating how they perceived the parents' interactions that day, and each recording their own sleep.¹

Stable, high three-way agreement on positivity predicted better sleep for everyone. Minimal perceived negativity predicted the best sleep in the house. And the poorest sleep, for every family member, went with **mixed daily perceptions and family-level fluctuation** — days on which the three of them did not see the same marriage.¹

Age 14.75

THE SAMPLE, AND THE AGE IT ACTUALLY STUDIED — ADOLESCENTS, NOT TWEENS. MOTHERS 41.8, FATHERS 44.0. THE ABSTRACT GIVES NO RANGE, SO A TWELVE-YEAR-OLD SITS IN THE TAIL AT BEST.¹

Print that caveat rather than bury it: this is a study of mid-adolescents. Applying it to a twelve-year-old is an extrapolation, and the desk that spent last week complaining about age bands ought to say so about its own secondary.

What survives the caveat is the mechanism, because it is unusual. Most of this literature measures conflict — how much shouting, how often. This measured *perception*, three times over, and found the cost in the gap between the three accounts. Mothers' sleep tracked the family's collective reading. Adolescents' and fathers' sleep tracked days that came in mixed, or lower than they had expected.¹

Three people in one house, filling in the same form about the same evening, and disagreeing.

THE CHATTER • THE PRACTICAL READING

There is no forum voice here and it would be dishonest to manufacture one for a paper five days old. What can be said is what the design implies. A household that is managing conflict by containing it — the argument had quietly, the atmosphere left unexplained — is producing exactly the condition the study associates with everyone sleeping worse: an adolescent reading a room differently from the adults in it. That is a hypothesis this study supports and does not prove. It is a thirty-day observational diary with no intervention, so nobody has shown that explaining things helps.

THE SCEPTIC • WHAT A HUNDRED AND SEVEN FAMILIES CANNOT TELL YOU

No effect sizes and no confidence intervals appear in the abstract, and the full text sits behind a paywall this desk could not pass — so the size of any of this is unknown here, and unknown is what gets printed. Multilevel latent profile analysis on 107 families generates a lot of parameters from a small sample; profiles found this way have a habit of not replicating. And the direction is assumed, not tested. A tired adolescent misreads a room. A badly-slept father is shorter at dinner. Every arrow in this paper could run backwards and the data would look the same, which is why the authors make no causal claim and neither does this page.¹

ALSO ON THE DESK

And the band problem, one more time. JAMA Pediatrics published a nationally-representative survey of online harms on 18 September — a weighted sample of 472, of whom 44.7 per cent had witnessed harassment and 29.8 per cent experienced it directly in the past four weeks, with online exceeding offline for violent imagery, 10.7 against 1.6 per cent. Its band is **13 to 17**, mean 14.9.² After three issues of screen research, the recurring finding is not what the studies show. It is how rarely any of them contain a twelve-year-old.

SOURCES ¹ "Daily and family profiles of adolescent-mother-father perceived interparental interactions and family members' sleep: a multi-reporter multilevel study," *Child Development*, 22 September 2026, DOI 10.1093/chidev/aacag168; abstract read at source on PubMed for the design, the sample, the reporter structure, the mean ages and every finding quoted. The full text is paywalled and was not read, so no effect size or confidence interval is printed and the absence is stated in the sceptic box. ² "Adolescents' Experiences of Online Harms by Platform," *JAMA Pediatrics*, 18 September 2026, DOI 10.1001/jamapediatrics.2026.4692; abstract read at source, cross-sectional, weighted n = 472, ages 13–17, fielded February to December 2025 across twenty platforms. General guidance, not medical advice.

THE BIG PICTURE · THREE MARKETS, THREE DIFFERENT STORIES

The futures board says the oil crisis is over. The last physical print says it had not started.

On 15 September, dated Brent was assessed twenty-two dollars above the front contract. The paper price has fallen nine per cent since. The agency publishes the next physical print today.


US 2Y PAR	US 10Y PAR	BRENT NOV	NASDAQ
4.71 ▼	4.96	98.65 ▼	27,244 ▲

Full strip on **p2**; par curve, not the market close.¹ Not advice.

The Energy Information Administration publishes a physical price for dated Brent — actual North Sea cargoes, not contracts — about a week in arrears. Its last print, for **15 September**, was **\$130.80** a barrel, up from \$96.02 on 1 September. That same day the front futures contract settled at **\$108.75**. Twenty-two dollars, and widening every week of the month.²

\$1.19 vs \$22.05

THE CONTROL THAT RULES OUT A DATA FAULT — ON 15 SEPTEMBER THE SAME AGENCY'S WTI SPOT SAT \$1.19 OVER WTI FUTURES — ENTIRELY NORMAL. ONLY THE SEABORNE BENCHMARK BLEW OUT.²



That control is why this is printable. Were the feed or the method broken, both series would misbehave; American landlocked crude behaved perfectly and only the waterborne grade detached. Paper Brent has since fallen five straight sessions, down **9.3 per cent**. Whether the cargoes followed is not yet knowable — the next physical release is today — but Murban, loading outside the Strait of Hormuz, rose 3.8 per cent on Tuesday as the Brent board fell 1.8.³

Several outlets now print the futures price under the words “dated Brent” — which is how you know one of them is wrong.

The other force is a central bank moving the opposite way. The Federal Reserve **raised** its target range a quarter point on 16 September, twelve to nothing. The projections matter more: the median for end-**2027** moved from 3.6 per cent in June to **4.1**. Three months ago the committee expected to be cutting next year; it now expects no net easing at all through 2027.⁴

THE CHATTER · WHAT IT MEANS FOR VWRA, GLOBAL BONDS, GOLD AND BTC

Hold the four together and they disagree. The front end is welded near its ceiling and global aggregate bonds are two cents off a fifty-two-week low: a bond market pricing stagflation. Equities read the same tape and priced the opposite — a record Nasdaq close. The hedges did neither: gold sits **22 per cent below** its January record in the middle of an inflation scare.

HIGHER-FOR-LONGER

DOLLAR BID

DISPUTED

THE SCEPTIC · AND THEY CANNOT BOTH BE RIGHT

The dissent runs both ways this week, which is rarer than a single contrarian. MUFG's George Goncalves and Agron Nicaj wrote on 11 September that “the market has priced-in a lot of the Fed path already and there is a risk it has overshot”. Tuesday closed 38 and 21 basis points above their quarter-end targets.⁵ The physical oil price implies the opposite.⁶ One is wrong, and what adjudicates is a spreadsheet posted today.

AND THE SINGAPORE READ-THROUGH

August exports rose 46.2 per cent, and almost none of it is momentum. S\$19.414bn against S\$13.280bn a year earlier, confirmed from the statistics department's own table.⁷ In level terms August was **up 0.2 per cent on July and flat since April**. Selena Ling of OCBC: one “should not extrapolate the 46%” indefinitely.⁸

SOURCES 1 Par yields read from the US Treasury's own daily CSV for 18, 21 and 22 September; market closes from Investing.com, read at source. The two differ by a few basis points because the par curve is a mid-afternoon snapshot and the close is not; both are printed rather than one implied. 2 EIA Europe Brent and Cushing WTI spot series, downloaded from the agency's own spreadsheets and parsed here; matching futures settlements from Investing.com's historical data. The EIA states its next release as 23 September. 3 Oilprice.com's live grade table, read at source, 22 September. 4 The FOMC statement of 16 September and the Summary of Economic Projections PDF, both read at federalreserve.gov. 5 MUFG, 11 September, read at source. 6 Graham Summers, 16 September, read at source; his claim about pipeline volumes is his own and is not carried. 7 Singapore Department of Statistics table M451301, read through the department's own data API, datasource Enterprise Singapore. 8 IBTimes SG, read at source. The Macro Desk explains markets and does not advise on them.

DOWN THE RABBIT HOLE · GLASSBLOWING AND LAMPWORKING

The piece is finished when it is cool, not when you stop working it.

Two crafts share one word and almost no equipment. One needs a furnace that never switches off. The other is a torch bolted to a bench — and it is the one you can have.

Say “glassblowing” and most people picture the hot shop: a furnace holding a pool of molten glass at around 1,150°C, a gathering iron, a glory hole to reheat in, and a team, because the piece is heavy and must never stop turning. That furnace runs twenty-four hours a day, every day — on Murano they shut down once a year, in August, for maintenance.¹ You cannot have one. A free-standing pot furnace from the reference American maker starts at **US\$36,231** before the glory hole and before the annealer.²

Lampworking — flameworking, to the Americans — is the other craft, and the name is a fossil: glassworkers once worked over an oil lamp with bellows. Today it is a torch on a bench. You start with cold solid rod and feed it into the flame, a few grams at a time. One person, sitting down, in a room.

\$11 vs \$33–\$75.50

AN HOUR OF STUDIO ACCESS, NEW YORK — A FLAMEWORKING BENCH AGAINST A HOT SHOP BY GLORY-HOLE SIZE, THREE-HOUR MINIMUM. THAT RATIO IS THE WHOLE ECONOMIC ARGUMENT FOR THE TORCH.³

WHY GLASS HAS NO MELTING POINT

Water becomes ice at exactly 0°C; glass does not work that way. Metals, as the craft’s teaching text puts it, “change from a solid to a liquid at a specific temperature (a melting point). Glass, conversely, transforms gradually from behaving like a solid to behaving like a liquid.” Even cold it has “the molecular structure of a stiff liquid” — a supercooled liquid.⁴

What exists instead is a **working range**: soft enough to move, stiff enough to hold shape. Too cold and it shatters; too hot and it sags off the rod. The number governing everything else is **COE**, the coefficient of expansion. Soda-lime “soft glass” runs about 104 parts per million per degree; borosilicate a third of that, 33.

Two consequences follow. Put a boro dot on a soft-glass bead and the two shrink by different amounts as they cool, loading the join with permanent tension; it may crack in the kiln, or on a shelf three weeks later. And boro’s low expansion is why it tolerates the torch, and why it became the glass of laboratories. Lampworkers name glass by COE before colour.

The skin sets hard while the core is still soft. Nothing shows.

THE CHATTER · HOW PEOPLE FALL IN, AND HOW BADLY IT GOES

Three doors. The forty-minute make-your-own session, where “experienced gaffers handle the most complex work” and you pick colours. The weekend intensive — and the truest sentence written about this craft comes from someone two days into one, who managed flowers, a paperweight and a bud vase, then tried a drinking glass: “**I completely and utterly failed.**”⁵ And the first bead class, the door actually open here (p17).

Philip Bailey came to glass by collecting paperweights, took a lampworking class in 1994, and describes his progress as going “from creating a collection of drunken shot glasses, to a few items I am pleased with.”⁶ Collector to maker to auction floor.

THE SCEPTIC · THE WORLD IS SHRINKING, WHATEVER THE MARKET REPORTS SAY

The honest headline in the numbers is contraction. The American Scientific Glassblowers Society lists “approximately 650 members.”⁷ Murano is down to about 300 glassblowers from thousands in the 1960s and 70s.¹ And in 2026 the field’s umbrella body shared its annual conference with the beadmakers’ society and the stained-glass association, billed as “One Conference — One Community”⁸ — three organisations that used to fill their own halls now filling one. Consultancy reports put the market on cheerful growth; trust the institutional counts.

SOURCES 1 PBS NewsHour’s Murano reporting, 11 October 2021, read at source, for the continuous furnace operation, the annual August shutdown and the count of remaining glassblowers. 2 Wet Dog Glass’s own product pages, read at source, for free-standing pot furnace pricing. 3 UrbanGlass’s published studio rental card, read at source; the page notes rates change on 1 January 2027. 4 Bullseye Glass, “TechNotes 4: Heat & Glass,” the manufacturer’s own technical PDF, downloaded and read here; both quoted sentences are verbatim from it. 5 A first-timer’s account of a two-day workshop at Sightdoing, 11 July 2018, read at source. 6 Philip H. Bailey’s own glass pages at MIT, read at source. 7 The American Scientific Glassblowers Society’s own membership page, read at source. 8 The Glass Art Society’s conference page and the beadmakers’ society’s own 2026 page, both read at source. Reddit and the main lampworking forum were both unreachable to this desk on every path tried — see p16 — and nothing in this desk is attributed to either.

THE DEEP END · THE GEAR, AND THE DECISION NOBODY TELLS YOU ABOUT

The torch is the cheap part. The argument is about oxygen.

A concentrator delivers its rated flow or its rated purity, never both — which is why you buy a ten-litre machine to run a torch that needs seven.

The mechanical distinction to learn first is **surface mix against premix**. In a premix torch, fuel and oxygen blend inside the body before the jets; in a surface-mix torch they meet only at the face. Surface mix cannot flash back into the body and dominates the art-glass market. The second distinction is circuits: one flame, or independently valved centre, middle and outer fires. What the money buys going up the ladder is not heat but **range** — a needle-fine point and a large melt from one torch — and control of flame chemistry, which is the mechanism behind silver glasses and striking colours.

THE LADDER, AT REAL PRICES

The line, not a rung. A Hot Head at \$50 screws onto a barbecue canister, runs no oxygen and cannot work borosilicate. A buy-once buyer who starts here buys twice. The quality entry point is a Nortel Minor at \$224 or a Bethlehem Alpha at \$350.¹

The sweet spot, and the one to buy. The GTT Lynx at \$525: seven jets, triple-mix surface mix, a pinpoint flame for detail and solid borosilicate up to two inches. It does both glass families, runs at pressures a concentrator can feed, and is the centre fire the whole upper range is built around.²

*Every rung a real improvement.
There is never a moment when the
next torch is obviously unjustified.*

Now oxygen, where a home bench is won or lost. Tanked oxygen gives high flow and purity, a permanent bill, and a cylinder to store away from any trace of oil or grease. The alternative is an **oxygen concentrator**: a compressor and a molecular sieve stripping nitrogen out of room air, with nothing to refill. Reconditioned ten-litre machines run \$850–\$995.³

85% at 5 LPM

THE SPEC SHEET'S QUIET LIE — ...90% AT 4.5, AND 95% AT 4, FOR A UNIT RATED FIVE LITRES. THE FLOW ON THE BOX IS THE FLOW AT WHICH IT PERFORMS WORST.⁴

SOURCES 1 Frantz Art Glass's own live product pages, read at source 22–23 September, for the Hot Head, Nortel, Bethlehem and Paragon prices and the Effetre and Double Helix lists. Note for the record that the entire Nortel entry and mid ladder was showing sold out when those pages were read. 2 ABR Imagery's own product pages, read at source, for the Lynx, Bobcat, Phantom, Delta Elite and Kobuki prices and the Lynx's published specifications. 3 Bradshaw Oxygen's ten-litre concentrator listings, read at source. 4 De Clarke's technical write-up on medical oxygen concentrators, read at source, for the purity-against-flow figures, the sieve life and the two-unit experiment. 5 A working beadmaker's account of living with a concentrator, read at source; quotations as posted. 6 The DragonJools lampwork blog, read at source; an older post, dated as such. 7 Michel Sun's ventilation guide, read at source, including his own caveats.

THE CHATTER · NAMED MAKERS, SINCE THE FORUMS WERE SHUT

The central English-language lampworking forum refused this desk on five separate routes and Reddit refused it outright, so there is no forum consensus here — only people writing under their own names. **De Clarke**, whose write-up on medical concentrators is the useful document here, supplies the purity curve above, puts sieve life at “20,000 to 40,000 hours” — which is why second-hand is reasonable — and records the ceiling: two units feeding one big torch and “the cannon barely whispered out a flame.”⁴ A beadmaker of twenty years supplies the warning nobody budgets for: “**Dust and moisture is the number one killer of oxycons.**” In this climate that is not a footnote.⁵

THE KILN IS NOT AN ACCESSORY

A bead annealer that merely holds a temperature cannot run a controlled descent; a programmable controller can. A Paragon BlueBird runs \$2,200 new.¹ Below that you are buying a box that gets hot.

The glass splits the same way. Italian soft-glass transparents run \$3.51–\$7.74 a quarter-pound. Then the silver-glass tier, where one rod costs more than a quarter-pound of the ordinary stuff: Aurae \$20, Skiron \$25, **Zephyr \$50 a rod.**¹ These only bloom if the flame chemistry is right — one maker, on a Double Helix colour: “after some significant practice — I am starting to get reproducible results.”⁶

THE SCEPTIC · THE VENTILATION NOBODY BUDGETS FOR

What comes off the bench is not smoke. It is nitrogen oxides from the flame, carbon monoxide, and metal oxides boiling off the coloured glass — the very metals that make the colours. The accepted design is capture at source: a hood, a fan, ducting that terminates outdoors, and make-up air entering at the rate the fan exhausts. That last clause is the one people skip, and it is why their system does not work. The lampworker who built the best-documented amateur system is candid that his own return air comes from under doors, leaving the room at negative pressure, and advises seeking professional help first.⁷ An open window is not a system, and neither is a kitchen range hood.

THE DESCENT · THE SLOW COOL, THE REAL COST, AND WHERE TO START HERE

Six millimetres anneals in three hours. Two hundred takes six hundred and fifty-four.

The tax on this craft is thickness, and it is brutal — which is why the kiln is the purchase, and why in Singapore the studio is not a compromise but the route.

Annealing is the mechanism the whole craft turns on. “As glass heats, it expands; as it cools, it contracts,” the field’s teaching text says, and those movements “set up stresses within glass, especially between the interior and the surface.”¹ The outside of a hot bead chills and sets rigid while the core is still soft, and when the core finally contracts it pulls against a shell that can no longer move with it. Nothing shows. The piece is beautiful, and it is a loaded spring.

Annealing unloads it — holding the work where the stress can still flow out, then cooling so slowly that no part of the piece is ever more than 5°C warmer than any other.¹ That tolerance is why you cannot open the kiln to look. Below the strain point whatever remains is permanent; above it, there is still time.

3 hrs → 654 hrs

ANNEAL CYCLE BY THICKNESS, FLAT SLAB — 6MM AGAINST 200MM — TWENTY-SEVEN DAYS, COOLING AT 0.23°C AN HOUR. DOUBLE THE THICKNESS AND THE CYCLE ROUGHLY TRIPLES.²

That chart, the maker notes, is “derived from Corning’s method” via McLellan and Shand’s *Glass Engineering Handbook* — a lineage running from a glassworks laboratory into every hobbyist kiln controller sold. The caveat beneath it is the argument against a cheap kiln, made by a glass maker rather than a kiln salesman: “**Even the most conservative annealing cycle may fail if the kiln cannot cool the work uniformly.**”²

The piece that shatters a week later was never unlucky. It was un-annealed.

THE VOCABULARY, TRANSLATED

Gather — the blob of molten glass collected on a blowpipe. Also a verb.

Glory hole — the open reheating furnace a hot-shop worker returns to, over and over.

THE CHATTER · AND THERE IS A SCENE HERE

This is the finding the desk did not expect. Singapore has a small, young lampworking scene — the discipline this page recommends. **Gurasu House** in Tai Seng calls itself “Singapore’s only studio fully dedicated to the specialised art of lampworking”: two-hour beginner workshops at **S\$158**.³ **Art Glass Solutions** in Joo Chiat runs bead-making for one to three. A hot shop exists too — a Mandai studio reported to hold the country’s only glass-melting furnace.⁴

THE SCEPTIC · THE PLACES YOU WOULD GO TO USE THE EQUIPMENT ARE CLOSING

Britain’s **National Glass Centre** closed to the public on 31 July this year, despite a campaign gathering some **34,000 signatures** which argued the department had been “profit-making for at least the last five years”. The university would not carry a subsidy of about £800,000 a year.⁵ In the same stretch North Lands Creative was liquidated and Nazeing Glass Works closed. The demand side has been shrinking in public for a decade too. The largest consumer show for jewellery-makers shut for good in 2020, and an attendee’s account of the decline is the truest line in the file: “**the aisles were spaced wider and wider apart.**”⁶

THE FIRST NINETY DAYS

Take a class; take a second with a different teacher. Then buy the eyewear before the torch — ordinary safety glasses do not block the sodium flare that hides your work, and didymium runs **\$120**. Bought once the bench is roughly **US\$4,240**. Here, the honest advice is the studio.

NEXT DESCENTS — PICK TOMORROW’S RABBIT HOLE

Letterpress printing — polymer against lead, and a Vandercook in a garage.

Lapidary & gem cutting — faceting to a meetpoint, and buying rough blind.

Sourdough & home bread baking — a starter that behaves differently in the tropics.

SOURCES 1 Bullseye Glass, “TechNotes 4: Heat & Glass,” the manufacturer’s own technical PDF, downloaded and read here; the expansion sentence and the 5°C tolerance are verbatim. 2 Bullseye’s “Annealing Thick Slabs” chart, downloaded and read here, for every cycle quoted, the attribution to Corning’s method via McLellan and Shand, and the quoted caveat. 3 Gurasu House’s own site, read at source, including the quoted claim; its street address is given differently by two sources and is not printed. 4 Art Glass Solutions’ own education page, read at source; the Mandai studio’s prices and the “only furnace” description are Home & Decor Singapore’s, that studio’s own site being behind a bot challenge. 5 UrbanGlass’s *Glass Quarterly* report on the National Glass Centre, read at source; the closure date and the subsidy figure are second-hand from the university and local press. 6 Craft Industry Alliance, read at source; both quotations as published.

THE LONG GAME · A POOLED ANALYSIS OF WHAT YOUR WRIST ALREADY RECORDS

The metric that looks most like a consistency score predicted nothing at all.

Seven studies, 109,365 people. A flattened daily rhythm carried a hazard ratio of 1.46 for death from any cause. Day-to-day regularity carried none.

A wrist accelerometer does not really measure sleep. It measures a twenty-four-hour curve, and researchers read three things off it: **interdaily stability**, how alike your days are; **intradaily variability**, how broken up each one is; and **relative amplitude**, the gap between your most active ten hours and your least active five. The *Journal of the American Heart Association* pooled seven cohort and cross-sectional studies — **109,365** participants — on 18 September.¹

HR 1.46

LOW RELATIVE AMPLITUDE, ALL-CAUSE MORTALITY — 95% CI 1.27–1.68; CARDIOVASCULAR MORTALITY 1.50 (1.13–1.98).
FRAGMENTATION: 1.05 (1.02–1.09).¹

A flatter curve — less difference between your active hours and your still ones — and a more fragmented day both tracked worse outcomes. And then the result nobody has led with: **interdaily stability showed no significant association with any outcome.**¹ That is the parameter closest to “did you keep the same schedule?”, which is the thing every wearable scores you on and every protocol tells you to fix. In this pooling it did not predict anything.

Amplitude predicted death. Regularity did not. Your ring grades you on the second one.

The practical half comes from a different paper, and it lands on exactly this reader’s cohort. *Diabetes, Obesity and Metabolism*, 17 September: **39,197** UK adults, mean age **45.8**. Each additional millilitre per kilogram per minute of cardiorespiratory fitness went with a **2.9 per cent** lower count of metabolic-syndrome traits and **6.0 per cent** lower prevalence. Lowest fitness quintile against highest: **2.4 times** the trait count, **5.4 times** the prevalence.²

The shape matters more than the headline. The splines were graded and non-linear, **steepest at the**

low-fitness end — the first few points of aerobic capacity buy the most. An already-trained mid-forties reader is on the flat part of that curve, which is an argument for spending the next hour on the rhythm rather than the interval.

THE CHATTER · THE ARGUMENT THIS LANDS IN

The longevity discourse has spent two years on *regularity over duration*, and it has real papers behind it: a 2023 *Sleep* analysis of more than ten million hours of accelerometry from 60,977 people found the four most regular quintiles carried 20 to 48 per cent lower all-cause mortality than the least regular, with regularity beating duration in equivalent models. A 2024 companion in *PNAS*, 88,905 people wearing light sensors, found the brightest daytime decile at hazard ratios of 0.66 to 0.83 and the brightest *night* decile at 1.21 to 1.34 — the paper behind every “morning sun, dark bedroom” protocol in circulation.³

Which makes this week’s null awkward rather than fatal. Those studies measured a regularity index; this pooling measured interdaily stability. They are cousins, not twins — and the cousin that failed is the one most wearables actually display.

THE SCEPTIC · A FLATTENED RHYTHM IS WHAT ILLNESS LOOKS LIKE

The strongest objection is on the record and it is simple. Writing in *The Lancet Healthy Longevity*, **Leandro Brito** and **Saurabh Thosar** note that “the observational nature of the retrospective analysis does not allow cause–effect conclusions to be drawn” and that the measure “is still in the early stages.”⁴ State the mechanism plainly: a flattened amplitude is what undiagnosed illness looks like on a wrist. Sick people move less by day and sleep worse by night, months before diagnosis.

There is also a cost to watching. A 2024 study of 523 adults found **35.8 per cent** using sleep trackers and orthosomnia — anxiety driven by the tracking itself — at between 3.0 and 14.0 per cent depending on threshold.⁵ General guidance, not medical advice.

SOURCES ¹ *Journal of the American Heart Association*, 18 September 2026, DOI 10.1161/JAHA.125.048656; structured abstract read at source on PubMed for the design, the seven studies, the participant total and every hazard ratio and interval, including the interdaily–stability null. The journal returned 403, so the paper’s own limitations and heterogeneity statistics were not read and are not characterised. Observational only. ² *Diabetes, Obesity and Metabolism*, 17 September 2026, DOI 10.1111/dom.71367; abstract read at source. Cross-sectional, capacity estimated from submaximal protocols, sample 69 per cent male and not population-representative. ³ Windred et al., *Sleep*, DOI 10.1093/sleep/zsad253, and *PNAS*, DOI 10.1073/pnas.2405924121; both abstracts read at source. ⁴ Brito and Thosar, *The Lancet Healthy Longevity*, read at source. ⁵ Jahrami et al., *Brain Sciences*, 6 November 2024, read at source including the authors’ own limitations.

THE DIARY · TWO THINGS ON THIS PAGE CLOSE BEFORE OCTOBER DOES

A show that quietly added a week, and an exhibition that leaves for good on the eighteenth.

Every row was re-checked at a promoter, a venue or an official ticketing page. Where no price is published, this page says so rather than repeat an aggregator's guess.

SINGAPORE

BOOK NOW · PLAN AHEAD

→ 11
OCT

★ Hiroshi Sugimoto — Form Is Emptiness

SINGAPORE ART MUSEUM, TANJONG PAGAR DISTRI PARK · **EXTENDED** FROM 4 OCTOBER · \$15 CITIZENS AND PRS, \$20 OTHERS · 63 WORKS ACROSS ELEVEN SERIES, PLUS FOURTEEN FOSSILS¹

→ 18
OCT

★ The F1 Exhibition

GARDENS BY THE BAY, EASTERN GREENWAY · **\$49** ADULTS, \$40 AGES 13–17, \$29 AGES 5–12; RESIDENTS 20% OFF · HOURS EXTEND TO 7PM FROM 5 OCTOBER · THE ASIAN DEBUT, AND IT CLOSES A WEEK AFTER THE RACE RATHER THAN TOURING ON²

9–11
OCT

GrandBite @ Bayfront

BAYFRONT EVENT SPACE · FREE ENTRY, REGISTER ON THE ORGANISER'S PAGE · 70+ STREET-FOOD VENDORS OVER THE THREE MOST EXPENSIVE NIGHTS IN THE CITY, REPORTED FULLY HALAL — A CLAIM MADE IN THE ORGANISER'S PUBLICITY, NOT ON ITS LISTING³

24–25
OCT

★ In The Mood For Love in Concert

SANDS THEATRE · FOUR PERFORMANCES ONLY · THE RESTORED FILM AGAINST UMEBAYASHI'S SCORE PLAYED LIVE · ON SALE; **NO PRICE IS PUBLISHED ON ANY PRIMARY PAGE**⁴

29
OCT–15
NOV

Cats — the fifth Singapore run

SANDS THEATRE · THE PROMOTER CONFIRMS OCTOBER 2026 AND A FIFTH RUN; THE EXACT DATES ARE FROM SECONDARY WRITE-UPS AND THE TICKETING PAGE RENDERS NO PRICES. THE SAFEST BIG FAMILY NIGHT IN THE WINDOW, ONCE A PRICE EXISTS⁴

26–29
NOV

The Nutcracker — State Ballet of Georgia

SANDS THEATRE · SIX PERFORMANCES, ~115 MINUTES WITH AN INTERVAL · THE COMPANY'S SWAN LAKE SEASON SOLD OUT, WHICH IS THE REASON TO BOOK EARLY RATHER THAN TO SKIP IT · PRICES SECONDARY ONLY⁴

25
NOV

Guns N' Roses — World Tour 2026

NATIONAL STADIUM, 8PM · ON GENERAL SALE, NO SOLD-OUT NOTICE · THREE VIP TIERS ARE NAMED ON THE PROMOTER'S PAGE WITH **NO FACE VALUES PUBLISHED**; THE FIGURES CIRCULATING ARE RESALE⁵

JAPAN & THE REGION

WORTH THE FLIGHT

31
OCT–1
NOV

★ Roppongi Art Night 2026

ROPPONGI HILLS, MORI ART MUSEUM, TOKYO MIDTOWN, 21_21 AND THE NATIONAL ART CENTER · SUNSET TO DAWN · **FREE** · THE FIRST TRUE ALL-NIGHT EDITION IN THREE YEARS, WITH THIS YEAR'S FOCUS ON FRANCE⁶

22
OCT

Jidai Matsuri, then Kurama

KYOTO · THE PROCESSION LEAVES THE GYOEN AROUND NOON; PAID SEATS **¥5,000–¥7,500** THERE AND **¥13,500–¥15,500** AT HEIAN JINGU, WITH ENGLISH COMMENTARY · THE KURAMA FIRE FESTIVAL IS FREE FROM 6PM THE SAME EVENING. TWO OF KYOTO'S THREE GREAT AUTUMN SPECTACLES, BACK TO BACK⁷

10
OCT–20
DEC

Tokyo Atlas

DAIBA, AOMI AND TENNOZU · **FREE ADMISSION**, CLOSED MONDAYS · A NEW CITY-SCALE ARTS FESTIVAL USING A PORT PARK AND AN UNDERGROUND CAR PARK⁸

5–27
NOV

Autumn Night Museum — Tokyo's first

SIX MUSEUMS OPEN TO 8PM ON FRIDAYS · **FREE AFTER 5PM FOR STUDENTS** WITH ID, GROUP RATES OTHERWISE — THOUGH NOT FOR THE TAIGA DRAMA OR MUSÉE D'ORSAY SHOWS⁹

2–4
OCT

MotoGP — Grand Prix of Japan

MOBILITY RESORT MOTE GI, TOCHIGI · A DAY TRIP FROM TOKYO · ON SALE THROUGH THE OFFICIAL STORE; GRANDSTAND AND GENERAL-ADMISSION TIERS ARE LISTED WITH **NO PRICES PRINTED**¹⁰

4–6
DEC

Clockenflap — Hong Kong

CENTRAL HARBOURFRONT · "BLIND BIRD" PASSES ON SALE WITH **NO ARTISTS ANNOUNCED AT ALL** — CHEAPEST TICKETS, MAXIMUM RISK. THE REPORTED PRICE IS SECONDARY AND THE DATES COME FROM THE TICKETING LISTING'S TITLE¹¹

THE SCEPTIC · THREE THINGS WIDELY LISTED THAT THIS PAGE WILL NOT LIST

Post Malone stays barred, and our own wording needed correcting: the reader reported it cancelled, but the official position is *postponed* — the 25 September date was postponed on 27 August with automatic refunds and **no new date**.

Christmas Wonderland has no 2026 dates and a homepage still titled for last year's closure. **ZoukOut** cannot be verified at all.¹²

SOURCES 1 Singapore Art Museum's own page, carrying the "extended until 11 October" banner. 2 The exhibition's own Singapore page for dates, prices and extended October hours; the ticketing agent returned 503. 3 The organiser's own listing; its start time differs from the organiser's publicity and the halal claim appears only in the publicity — both flagged. 4 The promoter's own shows list and the official ticketing listings; the venue refused this desk and those pages render prices client-side, so no price is printed for these three. 5 Live Nation Singapore. 6 Roppongi Art Night's English page. 7 Kyoto City Tourism Association's seat-sale pages, sale announced 8 September. 8 Tokyo Atlas's own site. 9 The Tokyo Metropolitan Foundation for History and Culture, release dated 18 August. 10 The official MotoGP store; the circuit's English site returned 404. 11 The festival's own site; its ticketing host served a truncated response. 12 Ticketing and press accounts of the 27 August postponement. All read at source.

THE TABLE • FOUR OPENINGS, AND ONLY TWO OF THEM CAN BE JUDGED

A viral hamburg brand's second attempt at Singapore soft-launched yesterday.

Nobody has reviewed it, so this page says so. The three-tonne Texan smoker out west, by contrast, has a verdict — and a parking problem.

Yamamoto's Hamburg opened for soft launch at Lau Pa Sat yesterday, grand opening 5 October. The pedigree is the story: founder Shohei Yamamoto's other brand is the hamburger operator that queued out of VivoCity, and this one is deliberately not that. "Yamamoto Hamburg was born from the kind of food I grew up eating at home," he says. Founded 2025, twenty outlets in Japan, and Hong Kong and Taipei before here.¹

Everything is under twenty dollars: the signature at **S\$19.90**, a beef patty stuffed with gorgonzola and mushroom under a miso-dashi demi-glace. **Too new to call.** All five pieces published so far are previews written before the doors opened; nobody has eaten there and filed. Go during the soft-launch fortnight, while the queue is only the people who read the previews.¹

Decker Barbecue is the one with a real verdict. Central Texas barbecue in the conserved 1956 Bukit Timah fire station, which worked until 2005 — a hundred and twenty seats and a fourteen-foot, three-tonne offset smoker shipped from Texas, reported the only one in Southeast Asia, running meats up to twelve hours. An independent smokehouse rebuilt with a large group's money: the PS.Café operators acquired it in 2025. Platters run **S\$40** to **S\$288**, and the Singaporean overlay is deliberate — a rendang brisket rice bowl, kailan for collard greens.²

THE CHATTER • THE CONSENSUS, AND THE ONE LOUD DISSENT

On the meat the local critics agree. *Time Out* called it "one barbecue joint that's worth the trip west," crediting the setting, the smoker and a menu that goes past the standard formula.²

The dissent is not culinary, and it is what a Bukit Timah readership will actually be caught by: **"parking sucks, with limited lots and no nearby parking available."**

Take the train to Beauty World. And note the booking policy, which tells you the demand: **walk-ins only at weekends.**

Two smaller ones. **Jeju Haenyeo** has put a second concept in the courtyard of The Arts House, modelled on the *bulteok* — the stone fire pits where Jeju's free-diving women warmed up after a dive. The food is built to be finished at the table, including a S\$25++ pancake that arrives as raw batter for a hot griddle. The one write-up that exists is enthusiastic on value and stops short of a verdict. **Too new to call on the cooking; the set is the draw.** A drinks-led pub, five till midnight — not for a twelve-year-old.³

Miura is the opposite in every way: a Japanese tea and wagashi counter hidden at the back of a clothing shop's Joo Chiat flagship, from the team behind Paaru. Matcha whisked to order from **S\$7.50**, a fruit daifuku at S\$7 — made fresh in limited daily batches. The coverage is consistently positive and genuinely thin: no negative review, and nobody pushing hard either way.⁴

THE SCEPTIC • WHAT "TOO NEW TO CALL" IS PROTECTING YOU FROM

Three of these four opened inside six weeks, and two have no published criticism whatever. The temptation in a column like this is to launder a press preview into a verdict by changing the tense — "will suit the lunch crowd" becomes "suits the lunch crowd" — and the reader never sees the join. Note too what inherited anticipation is worth: a queue is evidence of a reputation, not of a meal. The scoreboard entry is simple — whether anyone has filed an actual review by the end of October.

THE BASICS

Yamamoto's Hamburg — Lau Pa Sat; two sources give different street numbers. Walk-ins; halal status not stated, so do not assume it. **Decker** — 260 Upper Bukit Timah Road, closed Mondays; pork and beef. **Jeju Haenyeo** — The Arts House, Wed–Sun 5pm–midnight, not halal-certified. **Miura** — 316 Joo Chiat Road; two sources disagree on the opening days, so ring ahead.

SOURCES 1 HungryGoWhere, 21 September 2026, read at source, for the concept, the founder's quotation, the outlet counts and the prices, with Eatbook and Time Out Singapore's previews of 17 September corroborating the pricing. All five pieces published to date are previews written before opening, which is why this entry carries no verdict. 2 Decker Barbecue's own menu page for hours, platter pricing and the weekend walk-in policy; Time Out Singapore, 7 August 2026, for the opening and the quoted assessment. The brisket line and the parking complaint come from a local review site via search extract and are flagged as second-hand. 3 Eatbook, 25 August 2026. 4 Eatbook, 2 September 2026, with Honeycombers and Time Out Singapore consulted for the room and format. Reddit was unreachable and no social chatter is quoted anywhere on this page. All read at source.

THE UNDERCURRENT · CHINA

The Fishbone Clinic

A photograph of a doorway went round Chinese social media this month. The sign translates itself, into one run-on English word. Behind it is the quiet re-cutting of a health system along the grain of what patients say, rather than what organs do.

The sign is a dark maroon lightbox over a consulting-room door. In large white characters it reads 夜间鱼刺门诊 — night fish-bone clinic — and underneath, in smaller Roman type, someone has set the English as a single word: **Fishboneclinic**. Through the open door there is a pastel-green room, a fluorescent panel, a vinyl bench with a rucksack and a coat on it, and the elbow of somebody waiting.¹

It circulated in early September. The first reaction quoted in the coverage is the right one: 现在看病分得这么细了吗? — are they slicing medicine this thin now? Reporters went looking and found the same clinic in Nanjing, Chongqing, Wuhan and Shanghai.¹

The obvious read is that this is a joke about specialisation, and the obvious read is wrong twice over. It is wrong because the clinic is not new — the Nanjing one opened in May 2020, and the national press audited the phenomenon in 2022. And it is wrong because the clinic is not frivolous. What is new is the word that has grown up to describe the whole class of them, and who has started using it.

The word is 网红门诊, *wǎnghóng ménzhěn*: *wǎnghóng* is the ordinary Chinese for an influencer or a viral hit, *ménzhěn* the formal administrative term for a hospital outpatient clinic. Weld them together and you get something English has no phrase for — a hospital department named after a colloquial complaint rather than an organ or a specialty, which spreads because the name goes viral. It is not slang from the comment section. It appears in state press as an established label: in some cities, a party newspaper noted last year, the school-refusal clinic “has become a *wǎnghóng ménzhěn*.”²

The class is bigger than one doorway. Chinese hospitals now run a **school-refusal clinic**, a **learning-difficulty clinic**, weight-management, sleep and memory clinics, and — the purest specimen — a **director’s clinic**, named after the

seniority of the person sitting in it rather than the complaint it treats.²

You had to know which organ was failing you before you could buy a ticket.

Consider what a clinic name used to be. 耳鼻喉科 — ear, nose and throat. 内分泌科 — endocrinology. 消化内科 — gastroenterology. These are not brands; they are budget lines and legal statuses, drawn along the boundaries of what doctors are trained in. They ask the patient to perform a small act of self-diagnosis before the system will engage: work out which system is failing, then queue at the right window.

What is happening now is that hospitals — and increasingly the state — are naming clinics after *what the patient would say out loud*. There is a fish bone in my throat. My child will not go to school. I cannot sleep. I am too heavy. I cannot remember things. The deputy chief physician who runs the Nanjing clinic is admirably plain about why. They called it the fish-bone clinic, he says, so ordinary people would remember it — and in practice it also deals with chicken bones, meat bones, date pits and fragments of steel-wool scourer.³

500 yuan

THE SPECIAL CONSULTATION FEE AT THE BEIJING SCHOOL-REFUSAL CLINIC, OPENED JUNE 2024 UNDER PSYCHOSOMATIC MEDICINE. ROUGHLY A FIFTH TO A THIRD OF THAT DEPARTMENT’S GENERAL CASELOAD WAS CHILDREN AVOIDING SCHOOL, WHICH IS WHY IT WAS SPLIT OUT²

The naming also solves a real navigational problem. A parent whose twelve-year-old will not leave the bedroom does not know the department they want is called psychosomatic medicine. They know what they would say to a friend.

Which is the whole thesis in one sentence. The name is not a description. It is a search term.

SOURCES 1 The Shangyou News report as carried on Tencent News, 10 September 2026, read at source in the original Chinese, for the photograph and its caption, the netizen reaction quoted, and the confirmation of clinics in Nanjing, Chongqing, Wuhan and Shanghai. The photograph is itself credited there as an image from an online post. 2 A signed commentary in the Xinhua Daily group’s title, 15 April 2025, read at source, which is where the phrase *wǎnghóng ménzhěn* appears as an established label applied to the school-refusal clinic. 3 Dr Liu Honggang of Nanjing BenQ Hospital, quoted in a Health Times investigation republished by The Paper in November 2022, read at source in Chinese; the list of objects the clinic actually removes is his. The hospital’s own 2020 press coverage, also read at source, dates the clinic to May 2020; a 2026 account says 2019, and the earlier, contemporaneous date is used here.

The reason the fish-bone clinic works is not marketing. It is a hole. When Health Times audited more than twenty top-tier hospitals across Beijing, Shanghai and Guangzhou, it found that **more than half could not remove a deeply lodged fish bone at night** — many emergency departments simply do not stock a laryngoscope.¹ And the demand arrives on a schedule: patients turn up concentrated between six and nine in the evening, which is precisely when the scope room has gone home.¹

So a hospital bought a scope for the emergency area, put a doctor beside it, and gave the arrangement a name a person in distress could remember at eight o'clock at night. One evening a single physician removed twelve bones.² A Nanjing saying has grown up around it — for fish bones, go to BenQ. The clinic reports handling tens of thousands of throat-foreign-body cases since it opened, which is the hospital's own figure and is printed as such.

BY ORGAN	BY COMPLAINT
Otolaryngology	Fish bone, at night
Endocrinology	Won't go to school
Gastroenterology	Can't sleep
Psychosomatic medicine	Weight management

The patient must self-diagnose first. The name is the search term.

The same hospital, cut two ways. *Meridian original.*

Then the state joined in, and that is what makes 2026 the year rather than 2020. The National Health Commission ran a national **Weight Management Year**; by this year every tertiary public general hospital must offer the service, and more than five thousand five hundred institutions had set one up by February.³ Sleep clinics were mandated at least one to a prefecture-level city; memory clinics passed six hundred and sixty. Chinese health media now runs an annual *chart* of the year's ten new clinic names.⁴

60%

OF ADULT OESOPHAGEAL FOREIGN BODIES LODGE IN THE UPPER, CERVICAL OESOPHAGUS — WITHIN REACH OF A LARYNGOSCOPE, IF THE HOSPITAL OWNS ONE AND SOMEBODY IS AWAKE. CHINESE RESEARCHERS HAVE REPORTED OVER TWENTY THOUSAND ADULT CASES IN FIVE YEARS¹

The Nanjing physician is blunt about the folk remedy too. Swallowing a ball of rice to force a fish bone down, he says, is like hammering a nail in deeper. He describes one patient whose bone was eventually found lodged between the thyroid and the internal jugular vein, and another who believed he had swallowed his and returned two years later with a liver abscess and a four-and-a-half-centimetre bone.²

When a statute instructs a public clinic to have a brand, the franchise pitch is not a scandal. It is a forecast.

In June, Shenzhen put the logic into a draft regulation. Community health centres are to be supported in setting up weight-management, smoking-cessation and sleep-disorder clinics under a development model the text calls “**one centre, one specialty, one brand**”, inside a fifteen-minute health-service circle.⁵ A municipal legislature has written branding into health law.

Two weeks before that, the market had already found it. In May a post appeared on Xiaohongshu offering a franchise of a Shenzhen community health centre's *sleep clinic*: fifty-nine thousand eight hundred yuan, three years rent-free, equipment supplied, revenue split down the middle. Reporters found a “cell-optimisation sleep pod” from a biotech firm installed in the centre, and staff conceding they could only use the equipment to help improve sleep, not diagnose, because they were not credentialled physicians. The centre denied placing the advertisement. The district health bureau said no franchising policy exists and opened an investigation.⁶

That is the second-order consequence nobody has drawn out. A clinic name has become an asset, and assets get rented. What was leased in Shenzhen was not a service or a building. It was a signboard with the state's name on it.

SOURCES 1 The Health Times audit of more than twenty tertiary hospitals in Beijing, Shanghai and Guangzhou, republished by The Paper, 10 November 2022, read at source in Chinese, for the night-time capability finding, the six-to-nine arrival pattern and the foreign-body statistics it cites from a 2020 expert consensus. 2 Nanjing Morning Post coverage of 17 September 2020 as reprinted by the hospital, read at source, for the twelve-bones evening, the rice-ball warning and both case histories; the cumulative caseload is the hospital's own claim and is attributed as such. 3 China News Service, 12 February 2026, for the count of institutions with weight-management clinics; read in summary only and flagged. 4 A Chinese health-media round-up of the ten new clinics of 2024, read at source. 5 Shenzhen press coverage of the draft community health regulations, comment closing 27 June 2026, for the quoted development model. 6 Shenzhen Special Zone Daily, 20 May 2026, for the franchise offer, the sleep pod, the staff's concession, the denial and the bureau's response.

Britain has run this experiment, and the results are in. Between 2000 and 2010 the NHS opened **more than 230 walk-in centres** in England — nurse-led at first, no registration, no appointment, longer hours than a general practice. They were, the regulator’s own review records, “popular with the public,” and “attendances at many centres have exceeded expected levels.”¹

Then the sentence that could stand as this essay’s epigraph, written by a regulator summarising its critics: “Others believe that walk-in centres **create demand** for care for self-limiting, minor conditions. They say that the resources used to provide walk-in centres would be better spent on other priorities.”¹

By 2013, of the 238 estimated to have opened, **53 had closed in three years**. And the review’s patient survey found the diversion was real but partial: asked what they would do if the centre were not there, 34 per cent said they would go to a general practice and 21 per cent to accident and emergency — but 16 per cent would go to a different walk-in centre, and 8 per cent would have stayed at home.¹ A sixth of the traffic was circulating inside the new system, and nearly a tenth was traffic the system had invented.

Popularity was mistaken for relief; a fifth were shut by commissioners who decided they were duplicating rather than diverting.

China’s clinics are at the “extremely popular” stage. The school-refusal clinic at Beijing Children’s Hospital, a five-hundred-yuan special consultation opened in June 2024, took nearly ten thousand visits in ten months and appointments became almost impossible to get.²

And there is a sharp Chinese sceptic already on the record, which is worth more than a foreign one. Writing in a party newspaper, **Lu Yan** argued that the clinics are a genuinely positive intervention *and* that one should ask coldly how much of the fervour is real need and how much is 虛火, *xūhuǒ* — a term from traditional medicine meaning “false fire”, heat without substance. She reported doctors saying not

every child who arrives has a problem: some parents are simply over-anxious, some children have decent grades and low motivation, some parents believe one consultation will send a child back to school. And the clinic label itself, she warned, risks fixing a “school-refuser” tag on a child. Her closing line is the standard the whole system should be held to: the clinic is a medical institution, it does not sell smart pills, and it is not in the business of manufacturing top students.³ A second title put the same charge in its headline — scientific navigation, or the outsourcing of anxiety?

Hold the strongest version of the case against. This is not new: the Nanjing clinic is six years old and Chinese media have audited it twice before, so what trended in September was a photograph of a signboard, not a phenomenon. And lumping the fish-bone clinic in with weight and sleep clinics conflates two different things — one a single hospital’s bottom-up differentiation play, the others top-down public-health programmes with coverage targets.

Both are right, and neither touches the argument. The claim is not that a clinic with a funny name is new. It is that in the space of eighteen months the naming has been taken up by the state as policy, written into a draft municipal statute as a development model, and independently discovered by a private firm willing to pay fifty-nine thousand eight hundred yuan to stand behind one. The signboard that went viral this month is not the story. It is the moment the story became visible.

What is being emptied out is the word 门诊 itself. In a system where a department name carried legal status and a budget line, it is being refilled with a retail meaning. That is not cosmetic. It is the moment a health service stops organising itself around what its doctors know and starts organising itself around what its patients can say — which is more humane, and more legible, and generates its own traffic, and cannot easily be undone once every door has a name on it and a queue behind it.

Britain took thirteen years to find out. It is worth knowing, in September 2026, that somebody in Shenzhen has already worked out what the door is worth to rent.

SOURCES 1 Monitor, *Walk-in centre review: preliminary report*, November 2013 — the English health-sector regulator’s own review. The PDF was downloaded and read here; the count of more than 230 centres opened between 2000 and 2010, the 238 estimated openings against 53 closures in three years, the survey percentages and all three quoted sentences are verbatim. 2 Reported by CCTV News and relayed in the commentary below; the clinic, its June 2024 opening and its fee are from a New Beijing Post interview read in Chinese republication rather than the original, and flagged. 3 Lu Yan, signed commentary, *Xinhua Daily*, 15 April 2025, read at source in Chinese. The second headline quoted is from *Yangcheng Evening News*, whose body this desk could not read. Reuters, the BBC, the *Financial Times* and the *New York Times* all refuse this desk’s crawler, so the under-radar test rests on repeated open-web searches finding no English-language feature — stated here rather than implied.

MERIDIAN

THE LAST WORD

The slow cool

A glassblower can do everything right and lose the piece anyway. The outside chills and sets rigid while the core is still soft; when the core finally contracts it pulls against a shell that can no longer move with it. Nothing shows. The object leaves the kiln beautiful and loaded, and three weeks later it comes apart on a shelf. The craft's word for the remedy — annealing — is not a technique so much as a refusal to hurry. Six millimetres of glass needs three hours. Two hundred needs six hundred and fifty-four.

This issue kept finding the same shape. An agreement signed at the United Nations on Tuesday is headed “Permanent Nature”: no end date, and an obligation laid on a state nobody has founded. A suitcase is sold on a shell that does not break, while the wheel quietly hydrolyses in a cupboard between trips. A barrel of oil is ninety-eight dollars on a screen and a hundred and thirty on a quay. And somewhere in Nanjing a hospital put a laryngoscope in the emergency room at eight in the evening and gave it a name a frightened person could remember — and the name has outrun the thing. A door creates its own traffic. Britain learned that over thirteen years and closed fifty-three of them. Things do not set badly. They set later than you think, and somewhere you are not looking.

Twelve Desks · One World MERIDIAN · No. 111 · 23 September 2026

A note on method. Load-bearing figures were read on the pages that carry them: the Greenland agreement in full on the White House server, Article XI included; the US Treasury's own daily yield-curve file; the Energy Information Administration's Brent and West Texas spreadsheets, downloaded and parsed here; the Federal Reserve's statement and projection tables; Singapore's judiciary and its manpower, environment and transport agencies; and each maker's own catalogue for every price printed. The glass physics comes from a manufacturer's technical notes and annealing chart, read as published PDFs. Hosts that refused this desk are named where their absence matters — Reuters, the BBC, CNBC, the *Financial Times*, the *New York Times*, the *Straits Times*, CNA, Rimowa, the Council of the EU, the main lampworking forum, FlyerTalk and Reddit were all unreachable, and nothing here is attributed to any of them. What could not be verified is flagged in place or left out. Market levels are Tuesday's closes. The Macro Desk explains markets and does not advise on them; the Family and Fitness desks carry general guidance, not medical advice. Twelve Desks · One World.