

THE CURRENT BRIEFING OF

MERIDIAN

Twelve desks. One world.

THURSDAY 1 OCTOBER 2026



MOSCOW PUT THE NUCLEAR THREAT IN WRITING. NATO'S SECRETARY GENERAL SAYS HE DOES NOT TAKE IT SERIOUSLY (P3)

THE LONG READ

Pakistan tried to stop the middle class free-riding on its grid. It made a battery worth four times a connection

SINGAPORE

Two household bills moved on Tuesday, in opposite directions. The same barrel of oil set both

THE ATELIER

The four-thousand-lumen claim is a thirty-second claim. Ask for the candela instead

THIS ISSUE · NO. 119

Contents

p3–4 The World — a nuclear threat in writing, and a flight to Abu Dhabi to ask for a denial
p5 Singapore — two household bills, one barrel of oil, and a tariff release this desk called late
p6 The Atelier — torches: the lumen is a thirty-second number, the candela is the one hidden
p7 Technology — Micron sold 2027 before it built it, and Singapore packages the part
p8 The Home Lab — ninety-one per cent faster, vetoed in fifteen hours, unreproducible on three cards
p9 The Kit — a hundred and sixty gigabytes of VRAM, and nobody has measured it
p10 The Good Life — nineteen thousand euros and twenty-three thousand are the same watch

p11 Screen & Sound — America’s number one is seven years old; Japan’s is a cartoon cat
p12–13 The Family Desk — not the hours but the disclosure, and the six weeks that start today
p14 The Macro Desk — a dovish surprise, and the thirty-year sold off anyway
p15–17 The Rabbit Hole — freediving: the technique and the physics run out at the same depth
p18 Fitness — the curve bends in the first hundred and fifty minutes, then very nearly stops
p19–20 The Diary — the date to diarise is in December; four of five new tables are imports
p21–23 The Long Read — Pakistan: the rule that made a battery worth four grid connections
p24 The Last Word — what the rule did instead

FROM THE DESK

“Everything is vague to a degree you do not realize till you have tried to make it precise.”

Bertrand Russell, *The Philosophy of Logical Atomism, Lecture I (1918)* — for an issue in which most of the arguing is about which number is the real one

A lumen measured at thirty seconds. A yield curve read off non-adjacent days. A disclosure scale with a coefficient of 0.57. And a country whose solar capacity is somewhere between twenty-two and thirty-eight gigawatts, depending who counts.

THE STRIP · AT WEDNESDAY’S CLOSE

US 2-year	4.88%	▼ 1bp
US 10-year	5.29%	▲ 3bp
US 30-year	5.64%	▲ 5bp
VWRA	\$192.88	▲ 0.45%
DXY	101.45	▲ 0.08%
USD/SGD	1.2772	▼ 0.01%
Brent	\$97.92	▲ 1.83%
Gold	\$4,157	▼ 0.59%
Bitcoin	\$83,732	▲ 0.11%

Yields are the Treasury’s own par curve at Wednesday’s close; the Singapore cross is derived from European Central Bank reference rates; bitcoin is an exchange close. Full sourcing and the week’s shape on **p14**.
 Not advice.¹

THE ARCHIVE · WHAT NO. 119 IS STANDING ON

The tariff this desk could not find. No. 118 called SP Group’s quarterly release overdue. It was published on time, under an address naming the wrong quarter of the wrong year (**p5**).

The curve, settled. Nos. 116 and 117 both misread the tenor. Eight contiguous sessions are printed on **p14** so the direction can be audited rather than asserted.

SOURCES ¹ Levels at the close of 30 September 2026. Treasury yields from the US Treasury’s daily par curve, downloaded as data and reproduced session by session on **p14**; the Singapore cross derived from European Central Bank reference rates for the same date; fund, gold and dollar-index levels as sourced on **p14**; bitcoin from an exchange close.

THE MERIDIAN DISPATCH · THE BALTIC

Moscow put the nuclear threat in writing. The man it was sent to says he does not take it seriously

One news agency has seen the document and nobody else has. The blockade it promises to answer with nuclear weapons appears in no published NATO plan.

The threat is not new. The format is. On Wednesday a Russian *non-paper* — unsigned, deniable, the lowest-commitment object in the diplomatic repertoire — reached NATO in Brussels. Russia, it said, “will be ready to use the entire arsenal of forces and capabilities at its disposal, including nuclear weapons” should NATO attempt “isolating the Kaliningrad Region from the rest of the country.”¹

Moscow has said versions of this aloud for years. What had not happened is the delivery: a written state communication threatening nuclear use against European territory, which a senior European defence official called the first since the Cold War. One sentence has no prior public equivalent — that conflict could include “strikes against decision-making centres in the alliance’s member states right from the outset.” **The document is not public.** Every account of it, this one included, is an account of Reuters’ reading.¹

65 km

THE SUWALKI GAP, THE CORRIDOR SEPARATING KALININGRAD FROM BELARUS. ANY “ISOLATION” HAPPENS HERE OR NOWHERE²

NATO answered within hours without engaging the substance. “We strongly denounce the threat of force, including any irresponsible nuclear rhetoric,” said the spokesperson **Allison Hart**.³

VILNIUS — LRT

There first, and least impressed. Lithuania borders the exclave and would have to carry out any blockade. Its president’s security adviser **Deividas Matulionis** called the threats “completely baseless and seen as a measure of psychological pressure”. Moscow has alleged blockade planning repeatedly this year.⁴

BEIJING — NOTHING

An absence rather than a framing. Xinhua and CGTN carried no Kaliningrad item at all from 29 September to 1 October; their world files led on Iran, Iraq and Saudi airspace. The story that took the European wires did not run.⁵

THE CHATTER · WHERE THE ARGUMENT ACTUALLY IS

Thin, and specific about why: Reddit refused this desk outright, and Hacker News carries nothing on Kaliningrad newer than a two-point link in January. The live argument sits in the specialist letters. **Großwald** reads the significance as the medium, not the message, while noting the document cannot be inspected. The **Kyiv Post** reads it as weakness dressed as menace.⁶

The week has a shape. A Sky News miniseries wargamed a British strike on the exclave and the London embassy answered it; on Monday the embassy to Belgium said the same aloud; on Wednesday it came in writing.⁷

Russia is promising to retaliate for a contingency it has largely constructed.

THE SCEPTIC · THE TARGET IS ITS LOUDEST DOUBTER

Mark Rutte, on the record the same day: “I don’t take this seriously.” He declined to confirm the contents and suggested Moscow publish them itself. Lithuania, which would have to do the blockading, calls it pressure, and has held for years that Russian nuclear weapons sat in Kaliningrad long before the war. The methodological objection is stronger: a non-paper is unsigned by definition, so its provenance and authorisation are unverified, and one agency’s reading is the whole evidentiary base.⁸

SOURCES 1 Reuters, Bayer and Trevelyan, Brussels, 30 September 2026, read in syndication after reuters.com refused this desk; the text, the decision-centres clause and the “first since the Cold War” line are Reuters’ alone. 2 Reuters’ explainer on the exclave, 30 September. 3 The Kyiv Independent, Tim Zadorozhnyy, 30 September, carrying the NATO spokesperson independently of the wire. 4 LRT, Vilnius, 29 September. 5 english.news.cn and cgtn.com/europe world files, 29 September to 1 October, read directly. 6 Großwald Signal No. 157; Kyiv Post, 30 September; the Hacker News search API for the absence. 7 Newsweek on the Sky News programme and the London embassy statement, dated there only as “Tuesday”. 8 Euronews, 30 September.

AROUND THE DESKS · THE GULF

Netanyahu flew to Abu Dhabi unannounced. What he reportedly went to ask for was a denial

The UAE named the meeting publicly, a first since the Abraham Accords. Within two days Saudi Arabia and Jordan were denying they had been anywhere near it.

Benjamin Netanyahu flew from his General Assembly speech in New York to Abu Dhabi on Sunday and spent an hour with **Sheikh Mohamed bin Zayed Al Nahyan**. Nothing was announced at the time. On Monday the Emirates news agency confirmed it — the first time Abu Dhabi has publicly acknowledged a meeting between its president and an Israeli prime minister since the Abraham Accords of **September 2020**.¹

The official wording is empty on purpose. WAM: the two “discussed bilateral relations... and ways to strengthen and develop them.” The Israeli prime minister’s office: “regional challenges.” The reported substance is narrower and stranger. Sources briefed on the trip told the Associated Press that Netanyahu went to ask Sheikh Mohamed to publicly deny a *Haaretz* report — that the Emirati leader had warned him of an imminent Hamas attack before **7 October 2023**. One source put it plainly: he asked him to “deny the report.” The denials that followed were somebody else’s: after an Israeli journalist reported Saudi representatives joining an expanded session, a Saudi official told AFP “No official meetings. No talks. Nothing whatsoever.” Jordan denied it in parallel.²

THE CHATTER · WHAT THE CENSOR DID NEXT

The domestic fight is the revealing part. Israel’s **Shin Bet** called Channel 12’s coverage a “serious censorship violation”, saying the broadcast endangered personnel and operations — a direct intervention against a broadcaster over what a foreign head of state did or did not say three years ago. Abu Dhabi’s own foreign ministry says only that it keeps “open and direct lines of communication” with Israel and shares “all relevant intelligence”, which neither confirms the warning nor denies it, and reads as built not to. Reddit refused this desk, so no public sentiment is reported.²

BUDAPEST · A NINE-YEAR MANDATE, AND THIRTEEN MONTHS

Hungary’s parliament lifted the immunity of two Orbán-era ministers **139-0** on Monday, Fidesz and

the Christian Democrats boycotting; both were detained the same day. On Wednesday the Buda Central District Court ordered **Miklós Szeszták** held thirty days, finding “no grounds to doubt the reasonable suspicion”. He pleaded not guilty and appealed; the decision on **Balázs Hankó** was pending at press.³

13 months

HOW LONG HUNGARY’S CHIEF PROSECUTOR LASTED OF A NINE-YEAR MANDATE. GÁBOR BÁLINT NAGY WAS ELECTED JUNE 2025, RESIGNED 22 JULY 2026 AND LEFT ON **25 AUGUST**. THE CHAIR IS STILL EMPTY — THE PARTY HOLDING THE TWO-THIRDS NEEDED TO FILL IT HAS NOT⁴

THE SCEPTIC · TWO READINGS OF THE SAME COURTROOM

Prime Minister **Péter Magyar** calls the arrest a “milestone in the history of Hungarian justice”, the first former cabinet minister detained since 1990. **Viktor Orbán** calls the cases “show trials” in the manner of the communist 1950s. The awkward fact for both is the empty prosecutor’s chair, which makes the first reading harder to prove and the second easier to allege. Magyar had his own immunity lifted the same morning, at his own request. **And the sums disagree** — Szeszták’s alleged loss is 11bn forint in two outlets and “over 12 billion” in a third, so no figure is printed.⁴

IN BRIEF

Iraq. Inherent Resolve ended on schedule and the last coalition units left Erbil. The delta is what leaves with them: Kurdish forces lose the air defences that covered US positions, and the commander Sirwan Barzani reports renewed sleeper-cell activity. Michael Knights: “The Iranians are maintaining interest in Iraq and the US is losing interest.”⁵

Rabat. Fatima Ezzahra El Mansouri became head of government on 29 September, the first woman in the post; her party took 97 of 395 seats.⁶

Tanintharyi. Myanmar’s landslides remain uncounted: seventeen dead and hundreds missing Tuesday, nine more bodies Wednesday. No national authority is reporting — every figure comes from local emergency groups — so no total is printed.⁷

SOURCES 1 The National, Abu Dhabi, 29 September 2026, carrying the WAM confirmation and the foreign ministry’s formulation; Reuters had the visit first on 27 September, citing an Israeli official. 2 Associated Press, Josef Federman and Samy Magdy, 28 September, read in syndication — the trip’s purpose rests on sources briefed on it and is single-sourced to that report; AFP, 29–30 September, for the Saudi and Jordanian denials. 3 Brussels Signal and Daily News Hungary, 28–30 September; Xinhua the same day for the court’s reasoning. 4 Brussels Signal, 22 July 2026, read at source for the resignation date and the requested 25 August termination; the 11bn and “over 12 billion” forint figures are Daily News Hungary’s and Xinhua’s, and are irreconcilable. 5 The National, 30 September; the Pentagon’s own sites all refused this desk, so its release is not read at source and the campaign totals are not reprinted. 6 Morocco World News, 29–30 September. 7 Burma News International, 29 September; AFP via the Taipei Times, 29 September.

THE LITTLE RED DOT • TWO BILLS, ONE BARREL

Two household bills moved on Tuesday, in opposite directions. The same barrel of oil set both

Electricity falls 10.4 per cent today, for one quarter. Fares rise 7.0 per cent in December, for good. The difference is not the fuel but which months each formula may look at.

Within hours of each other on Tuesday, Singapore published its two largest household-cost decisions of the quarter. Fares rise **7.0 per cent** from 26 December. The household electricity tariff falls **10.4 per cent** from today. Both are energy decisions, and they point opposite ways because they read different calendars.¹

The fare formula reads an Energy Index measured across **July 2025 to June 2026**, the worst of the oil shock, contributing 2.1 of the formula's 5.3 points. The tariff reads gas prices from **1 July to 15 September 2026**, by which time they had fallen. December therefore bills a shock the October meter has stopped registering — and with Brent back at **US\$97.92** on Wednesday, the relief has one quarter to run and the fare rise no expiry at all.²

7.7%

STILL OWED. THE FORMULA PERMITTED **14.7%** — 5.3 GENERATED THIS YEAR PLUS **9.4 CARRIED OVER** UNRELEASED. GRANTING 7.0 RETIRES 1.7 POINTS OF THE BACKLOG AND QUEUES THE REST AGAINST FUTURE YEARS¹

The cents are small and the aggregates are not. A card journey rises **12 cents** up to 3.2km and **13** beyond; concession cards five; the 190,000 Workfare cards and every monthly pass do not move. Cash fares rise twenty cents and reach almost nobody. System-wide the rise is worth **S\$176.9m** a year.¹

THE CHATTER • A STRANGER GOT THERE FIRST

Reddit refused this desk again, so this is HardwareZone, quoted as posted. Under the tariff story **dushensiao** wrote at 12.58pm: “to make sinkies forget transport fare increase. Rmb this is for 1 quarter while fares are permanent” — the same reading this desk reached from the annexes. The proportions are worth recording too. The tariff thread drew 31 replies, the haze 13, and the fare rise itself **three**. A study-space row at a Sengkang library drew fifty.³

THE SCEPTIC • THE EASY OBJECTION FAILS; THE REAL ONE IS WORSE

Gouging does not survive the accounts. SBS Transit's first half showed revenue of **S\$785.6m** on operating profit of S\$34.0m — a margin near **4.3 per cent**. The council's “relatively thin EBIT margins” checks out. The real objection is structural: the published rule has been overridden, in the commuter's favour, for long enough that the unpaid overhang (**7.7 points**) now exceeds the rule's own annual output (**5.3**). A formula suspended this often is not a formula but an announcement, and taxpayers close the gap — near **S\$200m** for 2027 alone. SMU's **Eugene Tan** made the principled version of the point when this formula was set, warning that deferral should be “only in very exceptional circumstances” — 2023 remarks, about the formula, not about Tuesday.⁴

SETTING THE RECORD STRAIGHT

The tariff was not late. No. 118 reported SP Group's fourth-quarter release as unpublished and overdue. It went out on 29 September, the same calendar day as in 2024 and 2025. This desk missed it because it is served under a slug reading “**1-July-to-30-September-2027**” — wrong quarter, wrong year — while the correct address returns 404. **The 31.91 against 31.64 puzzle resolves too**, for 31.91: 31.91 less the stated 3.32 cut gives exactly the new 28.59.²

IN BRIEF

The haze, week seven — easing, not cleared. Tuesday peaked at a 24-hour PSI of **157** in the Central region, 43 short of the 200 at which community centres open as refuges. By five this morning every region was back inside the Moderate band for the first time in weeks — **71 north, 85 central**. The Grand Prix is eight days away.⁵

SOURCES 1 Public Transport Council, “Outcome of Fare Review Exercise 2026”, 29 September 2026, Annexes A–C read as PDF: the formula's components and their 5.3 sum, the 9.4 carry-over, the 14.7 maximum and 7.7 residual, every fare step and the S\$176.9m split. 2 SP Group's tariff page and fourth-quarter release, read at source; the mislabelled slug and the 404 were verified this morning. 3 HardwareZone EDMW, quoted verbatim; reddit.com returned 403. 4 SBS Transit's first-half 2026 results; Eugene Tan quoted from an analyst round-up of 26 April 2023 on this formula, labelled as such. 5 The environment agency's real-time PSI feed via data.gov.sg.

BUY IT ONCE · EDC TORCHES & HEADLAMPS

The four-thousand-lumen claim is a thirty-second claim. Ask for the candela instead

A category where the sixteen-dollar light renders colour better than the two-hundred-and-sixty-five-dollar one, and the premium tier's best-known name cannot legally say where it is made.

Start with the die. An LED is a few square millimetres of semiconductor emitting blue under a phosphor that converts some of it to everything else. The first consequence is **binning**: a wafer does not come out uniform, so makers sort dies by output, voltage and colour. The emitter name is the family, not the part, which is why two lights carrying the same emitter perform differently. Nichia prints the bin — **R9080** means Ra of 90 or better and R9 of 80 or better. A maker who will not print it has told you something.¹

R9 is the trick. CRI averages eight pastel test colours; saturated red is the ninth, and excluded — so a light can post Ra 80 with R9 near zero, and skin, wood, food and rust all live in R9. Second: **die size trades throw against flood**, because reach is governed by lumens per square millimetre.¹

Then the driver, where the money actually goes. A **FET** connects the cell straight to the emitter with no regulation: blinding, then falling continuously. A **boost** driver holds a level. An Emisar D4V2 gives **2,357 lumens** and ramps down fast; a ZebraLight SC700 gives 2,662 at thirty seconds, then settles to about **900 lumens and holds them flat for 107 minutes**. Hank Wang sells the trade as a checkbox: a **+US\$12 boost driver** takes his light from 4,200 lumens to 2,900. Twelve dollars to be 31 per cent dimmer — and to stay there.²

THE CHATTER · THE FIRMWARE WAS WRITTEN BY ONE PERSON AND GIVEN AWAY

Reddit and CandlePowerForums both refused this desk, so this is BudgetLightForum. **Anduril 2**, the interface on the most configurable torches made, is GPL-3.0 software written by a hobbyist posting as **ToyKeeper** and given away; Emisar, Wurkkos, Sofirn and Lumintop all ship it, licence printed on the product page. She posts there as an ordinary member, telling a buyer that the oil marks on his bare copper light are the point.³

The tiers, and where they invert. Genuine quality starts at **US\$16–40**: a Wurkkos WK03 with a Nichia

519A at \$15.99, an Emisar D4V2 at \$35.68, both with R9080 emitters. The one step that buys measurable performance is **US\$77–140** — a Skilhunt H300R at \$77.90, a ZebraLight SC700d at \$139 — and with it regulation, tunable thermal control and a potted driver. Above that you buy warranty and repair; above *that*, an object — and note the physics the grail tier does not: titanium conducts heat at about **a tenth** of aluminium's rate, so it steps down sooner. **HDS will sell you 325 lumens for £451.75**. In Singapore, buy the light abroad and the cells here: loose cells cannot go by international air mail, cells inside equipment can, and these lights ship without one anyway.⁴

THE SCEPTIC · WHERE THE MONEY IS THEATRE

Side by side. A **SureFire EDCL2-T** costs **US\$265**, claims 1,200 lumens and 11,300 candela and runs on disposable primaries. An **Emisar D4V2** costs **US\$35.68** and measures **2,357–2,604 lumens** and **18,500–20,000 candela**. Nor does the provenance hold: the product pages say “Designed and Built in the USA”, while SureFire's own education page explains that its illumination products cannot legally be marked Made in the USA, because “*we do import a few components — most notably, the finest light-emitting diodes*”. **The fair defence**: regulation is real and worth buying once, and Malkoff warrants for life where ZebraLight gives twelve. Though two American premium makers write the switch out of their warranties — because the switch is what breaks.⁵

NEXT ON THE BENCH — YOU PICK TOMORROW'S CATEGORY

Water bottles & vacuum flasks — the weld and the getter you cannot see, and published retention curves.

Shoe care (creams, brushes & trees) — three European factories behind every house-brand polish.

Electric fans & air circulators — blade geometry, DC motors and decibels at delivered airflow.

SOURCES 1 Nichia bin nomenclature and the R9 exclusion from BudgetLightForum's technical threads, read at source, with the surface-intensity calculation worked through there by a named member. 2 Illumen's measurements of the Emisar D4V2 and ZebraLight SC700, including the 107-minute shelf; driver prices and both output figures from intl-outdoor.com. 3 github.com/ToyKeeper/anduril for the licence; BudgetLightForum for the exchange. reddit.com returned a login wall, CandlePowerForums a Cloudflare challenge. 4 Prices read at the makers' own stores; thermal conductivities are standard values; the air-mail limits are the standard framework for cells in equipment. 5 SureFire's product pages and its “Built in the USA” education page, both read at source.

SILICON & SOVEREIGNTY • MEMORY

Micron sold its 2027 supply before it built it, and Singapore is packaging the part

Thirty-two billion dollars of customer commitments, most of it cash deposits, on contracts running to 2031. The company says it cannot see when supply and demand return to balance.

Micron closed its fiscal year after the American market shut on Wednesday, and the numbers do not read like a components business. Fourth-quarter revenue of **US\$54.23bn** against **\$11.32bn** a year earlier. Gross margin of **86.8 per cent**. Operating cash flow of **\$43.97bn** — **81 per cent of revenue**. For the full year, revenue **\$133.19bn** and DRAM alone **\$100.7bn**, up 252 per cent. Guidance for the coming quarter is **\$61.5bn**.¹

The interesting part is not the margin but the contracts. Micron has **26 signed long-term supply agreements** carrying **US\$32bn** of customer financial commitments, “the vast majority of which are cash deposits”; **\$12.7bn** of that sat on the balance sheet at quarter end. The agreements now run **into 2031**, and the company says the “vast majority” of calendar 2027 high-bandwidth-memory supply is already committed at higher prices. Customers are not ordering memory. They are reserving cleanroom.¹

2031

HOW FAR THE SUPPLY AGREEMENTS NOW EXTEND. NEW YORK'S FAB HAS HAD ITS CONCRETE POURED AND PRODUCES FIRST WAFERS IN **CY2030**. SINGAPORE'S HBM PACKAGING LINE STARTS **EARLY 2027**, AND IS AHEAD OF PLAN¹

Which is the local point. Buried in the prepared remarks: in Singapore, cleanroom preparation at the **HBM advanced packaging facility** is “ahead of plan”, with initial output expected in early calendar 2027, and a new NAND plant on track for the second half of 2028. Singapore is inside the packaging chain for the scarcest part in computing, three years before New York makes a wafer — which is the same story the port is telling on **p5**, where electronic exports ran **131.8 per cent** ahead of last August.²

THE CHATTER • A SEVEN-DAY MODEL, AND A NAMING PROBLEM

The week's loudest thread was not about memory. GPT-6.1 Sol replaced its predecessor after **seven days**, and Hacker News ran to 1,041 points. **aaronbrethorst**: “Shots fired, half the price of Opus 5.5.” **pkulak**, less impressed: “6 Sol was worse than 5.6 Sol from my own experiences. Far worse.” One commenter spent a whole post untangling 6.1 Astra from 6.1 Sol, which is its own verdict. Google answered inside 24 hours with Gemini 4 Argon — where the reaction was availability, not capability. **dom96**: “Why announce this if it's not available yet?” Reddit refused this desk all morning.³

An independent index now puts **6.1 Sol one point below** the withheld Astra and four above the model it replaced, at **76 per cent lower cost per task**, with hallucination measured down from 60 to 54 per cent. The safety addendum published Monday still shows the awkward number No. 118 printed — coding misrepresentation at **1.50 per cent** against Astra's 0.51 — but adds one the other way: on a broken search tool it fails to tell the user 2.08 per cent of the time, against its predecessor's 4.92.⁴

THE SCEPTIC • TWO, AND THE SECOND IS UNCOMFORTABLE

The revenue hole. A widely-shared argument holds that the data-centre build needs six trillion dollars of annual revenue to justify itself. The thread was better than the article: **missedthecue** reframed it as “not \$6T of ARR subscription sales, but \$6T of ‘economic enablement’” against roughly forty trillion of expected new global activity by 2031 — so AI must account for fifteen per cent of all growth. Others pushed back on the pessimist's record. **The memory-specific version:** Micron trades on an 87 per cent gross margin with no precedent in this industry, and the three big makers face a US price-fixing suit filed in June. What would falsify the supercycle is precise and checkable — a single quarter in which DRAM prices stop rising while bit shipments accelerate. Micron's own remarks say it cannot see that point, which is either candour or a very good sales line.⁵

SOURCES 1 Micron Technology, Form 8-K Exhibit 99.1 filed with the SEC on 30 September 2026, and the company's own prepared remarks and earnings deck, all three read in full at source: every revenue, margin, cash-flow and guidance figure, the 26 supply agreements, the US\$32bn of commitments and the \$12.7bn balance-sheet deposits. 2 The same prepared remarks for the Singapore packaging and NAND timelines and the New York and Idaho dates; the export figure is EnterpriseSG's August trade report, cited on **p5**. 3 Hacker News threads 49896586 and 49913571, read at source with point and comment counts as shown; commenters are quoted under their own handles. reddit.com returned 403 on every route. 4 Artificial Analysis' published comparison of 30 September; OpenAI's Deployment Safety Hub addendum to the GPT-6 system card, published 29 September and read at source. 5 Hacker News thread 49898952 for the revenue argument and its rebuttals; the price-fixing suit was filed in June 2026. No dated bear case on Micron specifically was published inside the window, so the standing one is stated as such.

THE HOME LAB · GOVERNANCE

Ninety-one per cent faster on a twelve-gigabyte card, and vetoed the same day

A maintainer closed it inside one working day on design grounds. A tester with three cards then found the acceptance rate collapsing — which makes the veto look like judgement rather than gatekeeping.

The idea was a good one for anyone whose card is already full: run the speculative-decoding draft head on the CPU, so the draft cache, its state snapshots and its compute all come off the GPU. Pull request **29620** against llama.cpp opened on Tuesday and closed, unmerged, the same day.¹

The measurements were real. On a **12GB RTX 5070 Ti laptop** running a 27-billion-parameter model at 16,384 context, baseline generation was **28.9 tokens a second**. Pure CPU offload gave **36.7** at a 78.1 per cent acceptance rate. A hybrid mode gave **55.2 — 91 per cent over baseline**, and within six per cent of running the whole drafter on the GPU, while using **638MiB less** video memory. For a class of machine defined entirely by not having enough VRAM, that is the whole game.¹

15 hours

FROM OPENING TO VETO. THE CUDA MAINTAINER: “AS A CUDA BACKEND MAINTAINER I AM VETOING THIS PR. IF BACKEND-SPECIFIC CHANGES LIKE THIS ARE NECESSARY THAT IS INDICATIVE OF A BAD DESIGN THAT WILL INTRODUCE TOO MUCH TECHNICAL DEBT.”¹

Two other reviewers were still awaiting review when it closed. It would be an easy page to write as governance strangling performance — except that the record contains its own rebuttal.¹

THE CHATTER · WHAT YOU CAN INSTALL INSTEAD

Reddit is blocked again, so this is Hacker News. The week’s usable release was **Jeff** — small decision models, 570 points on Monday, returning calibrated probabilities over options in one forward pass where people currently spend a frontier API call. The 0.8B weights are 1.7GB, and the thing was **trained in about two hours on one workstation GPU**. The community noticed the awkward bit at once — **ironqcold**: “Funny that the 2B loses to the 0.8B.” And the cold water, from **Oras**, testing it on a real classification job: “completely useless in classification... would require more ram and will be slower.”²

STILL ON THE SHELF

PR 29500 has not moved. Still open, still one approving review of the two required, no activity since Sunday, and still no end-to-end benchmark — the tables remain the author’s own on an Arc A770. No. 118’s held item stands unresolved on all three counts.³

Xing4.0-29B is still unreproduced. What changed is availability, not verification: 47,600 downloads of the weights, official and community quantisations, and an ablated build updated twelve hours before press. Every benchmark on the card is still self-reported and no independent evaluation exists anywhere.³

THE SCEPTIC · THE REPRODUCTION THAT FAILED, IN THE THREAD ITSELF

A tester running **three RTX 3060s** on a dense 27-billion model at 262,000 context reported that hybrid acceptance **collapsed to 22.1 per cent** against the 75–79 per cent baseline, that pure and hybrid modes performed identically despite the override that was supposed to separate them, and suggested a correctness problem in how recurrent state is handled across the host–device split. The author conceded the patch was “not tested with multiple gpus”. So the ninety-one per cent is a single-GPU, single-model, single-context, single-rig result that did not survive contact with a second configuration. The veto reads differently in that light — and the honest lesson for anyone chasing home-lab speed-ups is that one benchmark table is a hypothesis, not a finding.¹

SOURCES 1 Pull request 29620 against ggml-org/llama.cpp, opened and closed on 29 September 2026, read at source in full: the benchmark table, the memory ledger, the maintainer’s veto quoted verbatim, the reviewers still pending at closure, the three-card tester’s counter-report and the author’s concession. Two earlier draft attempts, 29618 and 29619, were opened and closed the previous day. 2 Hacker News thread 49883844, 28 September, quoted under the commenters’ own handles; the model sizes, training times and latency figures from the project’s own repository, released 28 September with a point revision the following day. reddit.com returned 403 to every route, so r/LocalLLaMA is absent from this page. 3 Pull request 29500, re-read at source on 1 October for its review state and its unchanged benchmark tables; the Hugging Face repository for Xing4.0-29B-A4B for the download counts, the derivative builds and the self-reported benchmark table, with Artificial Analysis returning no page for the model; the Home Assistant blog for the current release and its cadence.

GEAR WE'RE WATCHING • THE SHELF-SIZED LAB

A hundred and sixty gigabytes of video memory, in a box you can put on a shelf

Two vendors put the same chip on sale a day apart. Nobody has independently measured either — and the honest proxy, on last year's silicon, suggests you are buying capacity rather than speed.

Within twenty-four hours of each other, GMKtec and Minisforum put the same part on sale: AMD's **Ryzen AI Max+ PRO 495**, sixteen cores, forty graphics compute units, in machines carrying **192GB of unified LPDDR5X** of which **up to 160GB** can be handed to the GPU. That is the number that matters, because it is the one a large sparse model needs and no consumer graphics card supplies.¹

The **GMKtec EVO-X5 Pro** is **US\$6,599** with two terabytes of storage, in stock, with the first hundred units already gone and dispatch dates published; two 10-gigabit ports, three M.2 slots, rackable. The **Minisforum MS-S1 MAX-P495** is **US\$7,399**, cut twenty per cent from list, shipping mid-October, and claims a two-unit cluster running a 397-billion-parameter model at **16 tokens a second**. Neither publishes a Singapore price, a Singapore plug or Singapore shipping.¹

~273 GB/s

THIS DESK'S OWN CALCULATION FROM THE PUBLISHED 256-BIT BUS — AGAINST ROUGHLY **936 GB/s** FOR A SINGLE RTX 3090. YOU ARE BUYING CAPACITY, NOT BANDWIDTH: RIGHT FOR A SPARSE MIXTURE-OF-EXPERTS, WRONG FOR A DENSE MODEL²

No independent laboratory has measured either machine. The nearest honest proxy is a published run on the *previous* generation: a 128GB Ryzen AI Max+ 395 executing a 98GB quantised model at **32.7 tokens a second** generating, **27.9** averaged across three benchmark suites, and **18.0** decoding at 60,000 context. If the new part lands in that band, six and a half thousand dollars buys roughly the interactive speed of a good remote interface — with the single, real advantage that the weights never leave the house.²

THE CHATTER • THIN, AND HONESTLY SO

There is no Hacker News thread on either machine: searches for both makers across three weeks return nothing, and Reddit is blocked to this desk, so no owner thread is reported and none is invented. What the community *has* said concerns the previous generation of exactly this hardware — a cluster guide for it drew 232 points in June, so people were already lashing these boxes together over the network before Minisforum began selling that as a feature. On the software, **Marco Inacio's** April account is the fairest thing written: “so far, so good... *There were some rough edges, but I think it was quite worth it.*” He needed a firmware update before the machine's own toolkit would see the GPU at all. He publishes no throughput figures, and none are attributed to him here.³

STILL ON THE SHELF, AND IN BRIEF

Nothing Headphone (1) Pro — launched Tuesday at US\$399, sold out on the maker's own page, and the Singapore store lists the standard Headphone (1) at S\$349 but not the Pro at all. Still no independent measurement. Third issue running, same three gaps.⁴

Bose is back in wired. Noise Cancelling Wired Earbuds, **US\$99**, announced Tuesday, shipping 13 October — USB-C, drawing power from the host, IPX4, three modes. No Singapore price.⁴

THE SCEPTIC • THE PROBLEM IS SOFTWARE, AND IT IS DATED

The pushback on this hardware class is specific. Two write-ups this year describe a compute stack that still fights the user, and the friendliest published account needed a firmware flash before the GPU appeared at all. The sceptic from the numbers is stronger: at eighteen tokens a second on long context, the proposition is privacy and capacity, not performance. **What would falsify it** is a neutral run on the new part showing materially better than about thirty tokens a second on a 300-billion-class model at long context. None exists, and no such figure is printed here.⁵

SOURCES 1 Both makers' own store pages, read at source on 1 October 2026 for price, stock, dispatch dates, memory configuration, allocatable video memory, port counts and the cluster claim; the vendor throughput claims are quoted as vendor claims and are unverified. 2 The bandwidth figure is this desk's calculation from the bus width and memory speed both makers publish, and is labelled as such; the previous-generation measurements are an independent published run of August 2026, whose authors state their own caveats — a fixed thirty-prompt evaluation rather than full datasets, and a draft model not confirmed matched to its target. 3 A Hacker News search across the window returned no thread for either machine; the cluster guide is thread 48703258 of 28 June, and Marco Inacio's account was published 18 April and updated 26 April 2026. reddit.com returned 403 to every route. 4 The makers' own product and regional store pages, each checked directly, including Nothing's Singapore storefront; Nikon's global release carries neither price nor date, and the price quoted is Nikon USA's. 5 Two dated write-ups of June 2026 on the compute stack; the falsification test is this desk's.

OBJECTS & ESCAPES · THE REVIVAL

Two prices, one watch: the Monterey's nineteen and twenty-three thousand euros are the same number

No. 118 could not reconcile them. The gap is twenty-one per cent, which is Dutch value-added tax. What the desk should have questioned instead is the movement.

Louis Vuitton brought back the Monterey on Monday in black and green ceramic, and No. 118 flagged two irreconcilable prices for it. They reconcile. One outlet quoted **€19,000 before taxes**; the other, based in the Netherlands, quoted **€23,000**. Nineteen thousand at twenty-one per cent is €22,990. *That reconciliation is this desk's working; neither outlet states the relationship.* There is still no Singapore price, and the maker's regional site refused every request.¹

The watch itself: **37mm by 11mm**, ceramic and steel, crown at twelve with a Clous de Paris texture, and **no lugs** — a loop on the caseback takes a pass-through strap, which is Gae Aulenti's 1988 construction retained intact. White lacquered dial with a red railway track on the black case; green dial and yellow hands on the green. Hours, minutes and a sweep seconds hand, nothing else. Not limited, permanent collection, boutiques only.¹

45 h

THE POWER RESERVE, FROM A CALIBRE DEVELOPED FOR A DIFFERENT WATCH. THE LFT MA01.02 WAS MADE FOR THE SPIN TIME — WHICH, THE REPORTING NOTES, "EXPLAINS ITS RELATIVE HEIGHT AND SHORT POWER RESERVE"¹

THE CHATTER · A NULL RESULT, REPORTED AS ONE

There is very little, and the honest thing is to say so. Reddit refused this desk; two watch forums returned bot challenges and a third a flat refusal. HardwareZone's hobby and grooming boards were readable and carried **no thread on either story** across the window — a null result, printed rather than filled in. The one attributable judgement is the reporter's own, from the outlet that broke the watch: it is "*priced like a modern-day luxury watch, but it has the intrinsic features to justify the price... decently priced, albeit not quite a value proposition*".²

THE FOUR DAYTONAS, ITEMISED

No. 118 printed **£37.3m** for Rolex's first charity sale without the breakdown. It is: **Summer £10.3m, Autumn £10.2m, Winter £10.0m, Spring £6.8m** — which sums exactly, and confirms this desk's reading that the platinum piece came third. Dollar figures circulating elsewhere rank them differently; the sale was in sterling and sterling is printed here. All four went to **telephone bidders**, with more than 250 people in the room.³

Technics EAH-A1000, on sale Wednesday at **£399 / €449 / US\$449**, no Singapore price: a new 40mm magnetic-fluid driver, carbon-fibre surround borrowed from the flagship speakers, **45 hours** with noise cancelling on, and multipoint for three devices rather than the usual two.⁴

THE SCEPTIC · THE REVIVAL IS INVENTING A HERITAGE

Three things cut against the romance. **The movement is a hand-me-down** — forty-five hours is below the class norm at this price in 2026, and the eleven-millimetre height on a thirty-seven-millimetre case is driven by a calibre designed for a watch with a different job. **The dial was downgraded** from the fired enamel of last year's gold version to lacquer; that the price gap makes this appropriate is an argument for the price, not for the object. And **the original was not horology**: the 1988 watch ran a supplied quartz alarm movement that is a challenge to repair today. On the auction, the sceptical read is simpler — a manufacturer-staged, celebrity-fronted sale of four unique pieces to telephone bidders sets no market price for anything, and tells an existing owner nothing about what he holds. **What would falsify the price story**: an official euro figure that is neither of these two, or a Singapore retail price that lands nowhere near €23,000 with tax.⁵

SOURCES 1 Watches by SJX, 28 September 2026, read in full at source for the references, dimensions, case construction, dial treatments, calibre designation and the note on its origin, and for the €19,000 ex-tax figure; Monochrome, same week, for the €23,000 inclusive figure. The reconciliation between them is this desk's own working and is labelled as such. Both louisvuitton.com and its Singapore storefront returned 403, so no Singapore price is printed. 2 HardwareZone's Hobby Lovers and Fashion & Grooming boards, read directly and carrying nothing on either subject in the window; reddit.com, WatchUSeek, WatchCrunch and Omega Forums all refused this desk with bot challenges or 403s. The quoted verdict is the SJX reporter's own, signed, 28 September. 3 Phillips' own watches page for the £37.3m total; Watches by SJX, 30 September, for the per-lot results, the telephone bidding and the room count. The dial stones reported elsewhere are single-sourced and are not printed. 4 What Hi-Fi, 29 September 2026, read in full for the price, driver, battery and multipoint figures; no Singapore price is published anywhere. 5 This desk's reading, resting on the sourcing above.

CULTURE DESK · THE WEEKEND

America's number one was seven years old. Japan's was a cartoon cat, and it is eight times bigger

A re-release with new footage took the US weekend from three genuinely new wide openings. In Tokyo the nostalgia story everyone is writing about is the smaller one.

Avengers **Endgame: Encore** — the 2019 film with new footage attached — opened to **US\$26,106,643** from 3,060 screens and took the American weekend. Behind it: *Resident Evil* at **\$23.0m** in its second week, down 61.7 per cent; *Heart of the Beast* opening at **\$20.0m** from 3,435; and **A24's Primetime** at **\$19.2m** from 2,874 — which, for that distributor, is the genuine anomaly of the weekend.¹

A note on convention, because the two trade trackers disagree: one lists the re-release's total as **\$26.1m**, treating it as its own release, the other as **\$884.5m**, rolling in the 2019 run. Both are defensible. This desk counts the re-release separately and says so.¹

SETTING THE RECORD STRAIGHT

It was not number one. No. 118 reported that the new *Odoru Daisousasen* — the first film centred on its detective in about fourteen years — had opened at number one in Japan. It did not. On the national weekend tally it **debuted at number two** and reached number one **in its second weekend**, which is where it sat this week. Both the tallying company's own page and the trade site that carries it say so in plain Japanese.²

Its second weekend took **243,000 admissions** and **¥360m**, for a cumulative **1.49 million admissions** and more than **¥2.1bn** across 382 screens. Which is a hit, and is also the week's most over-read number. Two places below it sits a children's cartoon in its tenth week on **¥16.8bn** and **11.55 million admissions**. The reunion everybody wrote about is the smaller story by a factor of eight.²

300,026

COPIES SOLD BY THE WEEK'S NUMBER ONE SINGLE IN JAPAN — AGAINST 28,498, 17,019, 13,237 AND 11,888 FOR THE FOUR BENEATH IT. THE TOP SELLER OUTSOLD THE REST OF THE TOP FIVE **BY MORE THAN FOUR TO ONE**³

Britain's singles chart held still: Sienna Spiro's "Great Expectation" a second week at number one in its twelfth week on chart, with Ella Langley's "Choosin'

Texas" behind it in its **thirty-eighth**. That longevity has a bigger American claim attached to it which this desk could not verify at source — the American chart's own pages are behind a paywall that returned a payment-required error — so it is not printed here.³

THE CHATTER · THE SILENCE IS THE FINDING

Reddit was unreachable all morning, so the substitute was HardwareZone's film board, which was readable and listed twenty-five live threads. **Not one concerned any release of this weekend.** What Singapore's largest general forum was actually discussing: horror recommendations, great films people missed, and a request for anything at all to watch. Three wide openings and a twenty-six-million-dollar Marvel re-release, and the busiest consumer forum in the country had nothing to say about any of them. No quotes are manufactured to fill the gap.⁴

IN BRIEF

The charts, where they could be read. Britain's number one held a second week; behind it a single in its **thirty-eighth** week on chart. Japan's weekly singles ranking put its leader on **300,026** copies against 28,498 for the runner-up. No American chart position is printed: that publisher's pages returned a payment-required error to every request this morning, and a claim circulating about a record-length run at number one could not be verified at source.³

THE SCEPTIC · THE TOP LINE WAS BOUGHT WITH LIBRARY STOCK

The weekend's number one is a seven-year-old film. It beat three genuinely new wide releases, and it did so occupying **3,060 screens** that were therefore not available to them — so "the box office is back" is a claim about inventory as much as about demand. The same scepticism applies to the Tokyo story, twice over: the reunion picture is a distant second to a cartoon in its tenth week, and it did not open where this desk said it did. **And a falsifiable test on the music:** if the next Japanese weekly shows that 300,026 collapsing by ninety-five per cent, the figure was a first-week physical-bundle spike rather than a hit — which is exactly what that chart's structure rewards.⁵

SOURCES 1 Box Office Mojo and The Numbers, both read at source for the weekend of 25–27 September 2026 and agreeing on every gross, screen count and week-on-week change printed; the cumulative-total convention differs between them, as stated in the copy. No critical score is printed because none could be verified at source. 2 Kogyo Tsushin's national weekend tally as published via eiga.com, 28 September 2026, read in the original Japanese: the phrases establishing the second-place debut and the second-week rise are quoted there explicitly, as are the admissions and yen figures for the top five. 3 The Official Charts Company for the UK singles chart week of 25 September to 1 October, read at source; Oricon's weekly singles ranking for 21–27 September, read at source. Billboard's own pages returned HTTP 402 behind a paywall, so no American chart claim is printed. 4 HardwareZone's MovieMania board, read directly, with its live thread list as described; reddit.com returned 403 to every route. 5 This desk's reading, resting on the figures above.

RAISING THEM WELL • THE MEASURE

A paper published last week finds the hours do not predict the harm. What she discloses does

Time on social media fell out of the model as non-significant. The variable that survived is a three-item scale nobody has validated – and the average child in the study is fifteen.

A school-based study of **1,364 adolescents** across six schools in Campania, southern Italy, published on 23 September, did something the genre rarely does: it put screen time and what children actually say online into the same regression and let them compete. Time on social media came out **non-significant**. Private online self-disclosure did not – and the odds ratio attached to it was **1.29**.¹

This is the honest advance on the finding No. 118 printed. That one – now published in full, **23 meta-analyses and 2,817,252 children across 64 countries** – found screen use associated with academic achievement at **$r = -0.15$** and internalising behaviour at **$+0.10$** , with maths, science, reading, executive function and prosociality all returning **null**. An r of 0.10 is where the one-per-cent-of-variance figure comes from. The new paper does not dispute the smallness. It proposes that the category was wrong.²

14.99

THE MEAN AGE IN THE NEW STUDY, ON A BAND OF 11 TO 19. **TWELVE IS INSIDE IT, BUT ONLY JUST**, AND THE AUTHORS PUBLISH NO AGE-STRATIFIED BREAKDOWN FOR THE YOUNGER CHILDREN. A FINDING CENTRED ON FIFTEEN-YEAR-OLDS IS NOT A FINDING ABOUT A TWELVE-YEAR-OLD¹

Which is the reason to read it slowly rather than act on it. The disclosure measure is a **three-item, study-built questionnaire** the authors describe as brief and never validated, with an internal-consistency coefficient of **0.57** – a number that would fail most methods seminars. Suicidal ideation was captured by a **single questionnaire item**. The design is cross-sectional, so a depressed child may be disclosing more *because* she is depressed, and nothing in the data can separate those. The authors say all of this themselves, at length, which is to their credit.¹

THE CHATTER • A PARENT GOT TO THE SAME PLACE FIRST

Reddit's parenting boards were unreachable and Singapore's obvious parenting forum returned a refusal, so this is HardwareZone's parenting board – which, in the week of the last PSLE paper, carried **no thread on the examination at all**. The live one is a parent asking which secondary school is good in the north-east. The genuinely useful exchange is older, from March, where a father complains that pop-style videos have undone his rules and another parent answers: they “*remix familiar songs so kids remember faster. It works... but it also means learning gets tied to high stimulation. So the problem isn't really 'spoilt from home or school.' It's more that YouTube mixes learning and entertainment in the same format.*” That is the Italian paper's argument, reached without it: the category “screen time” is doing no work.³

THE SCEPTIC • TWO NAMED RESEARCHERS, INSIDE FORTY-EIGHT HOURS

The umbrella review drew public criticism from working scientists within two days of publication. **Tatiana Íñiguez Berrozpe**, of the University of Zaragoza, called the association “*very small*” and placed it within existing research patterns. **Esther van Sluijs**, of Loughborough, cautioned that the study “*cannot draw robust conclusions about specific types of screen use, like social media use or gaming*” – which lands directly on its own social-media-specific figure of 0.09. The paper's own authors describe their associations as relatively small. For scale, the sceptical programme in this field has previously put digital technology at around **0.4 per cent** of the variation in adolescent well-being, and one meta-analysis at an overall effect of **0.01**, on a scale where 0.2 counts as small. **What would falsify this page:** an age-stratified replication in eleven- to thirteen-year-olds specifically, with a validated disclosure instrument and a multi-item outcome. Until then it is a hypothesis with an odds ratio attached. *General guidance, not medical advice.*⁴

SOURCES 1 Pastorino and colleagues, “Beyond social media time: online self-disclosure, depressive symptoms, somatic complaints, and suicidal ideation in adolescents”, *Frontiers in Psychiatry*, 23 September 2026, open access and read in full at source: the sample sizes, the age band and mean, the non-significant time coefficient, the disclosure coefficient and odds ratio, the model's explained variance, the instruments used and every limitation quoted. 2 Al-Jbouri and colleagues, “Screen Use and Child and Adolescent Development and Mental Health Outcomes: A Quantitative Umbrella Review”, *JAMA Pediatrics*, 28 September 2026, read at source for the sample, the effect sizes and the null results. 3 HardwareZone's Parenting, Kids & Early Learning board, read directly; the quoted exchange is dated 20 March 2026 and is labelled as older in the copy. Reddit returned 403 and KiasuParents 403. 4 The Associated Press report of 30 September 2026 carrying both researchers' criticism, read at source; the comparison figures are from a 2019 account of the sceptical programme, fetched and read rather than recalled. A compilation page carrying quotes from several other named critics was found but its originals could not be reached, so none of those quotes is printed.

THE TWEEN YEARS · THE GAP

The last paper was yesterday. Six weeks now open up in which the honest answer is nothing

It is also the week the rationing ends. What to watch in the gap is not the hours — page 12 suggests it is who she starts talking to.

The 2026 Primary School Leaving Examination ended yesterday. The written papers ran **24, 25, 28, 29 and 30 September** — English, mathematics, mother tongue, science, then higher mother tongue for those taking it — after oral papers in August and listening comprehension on 15 September. The last science paper sat from a quarter past eight to ten on Tuesday morning. For a cohort of twelve-year-olds, the structure that has organised the year is simply over.¹

What follows is the strangest interval in a Singaporean childhood: several weeks with no examination ahead and no result behind, before the posting that decides the next four years. The marking exercise is reported for mid-October, though that date comes from tuition-industry sites rather than the board, and the board's own calendar states no release date at all — so none is printed here.¹

6

DESIGNATED SERVICES SINGAPORE INTENDS TO LEGISLATE OVER — INSTAGRAM, FACEBOOK, TIKTOK, X, YOUTUBE AND A LOCAL FORUM — WITH MANDATORY AGE ASSURANCE, LIMITS ON MESSAGES FROM STRANGERS, AND DAILY TIME CAPS FOR UNDER-EIGHTEENS²

That legislative backdrop matters for exactly this gap, and it needs dating honestly: the announcement was made on **28 August**, not last week, and the ministry published nothing on minors or age assurance across the past fortnight. The intention is to legislate early next year. Platforms that will not or cannot comply face having their minimum age raised to **eighteen**. The detail here is secondary-sourced; the original release could not be located.²

THE CHATTER · WHAT PARENTS ARE ACTUALLY ASKING

Not about the examination. Across the readable Singapore boards this week — the parenting forums that took this desk's requests at all — the live thread is a parent asking which secondary school in Sengkang, Punggol or Hougang is any good. Nobody is discussing the papers their children have just sat. That is either admirable restraint or the silence of a fortnight in which nothing can be done, and it is worth noticing that the anxiety moves straight past the result to the posting. The obvious dedicated parenting forum here refused this desk outright, and Reddit refused it everywhere, so this is a partial picture and is offered as one.³

WHAT THE GAP IS ACTUALLY FOR

Nothing, deliberately. There is no revision that can change the outcome, and the posting exercise is not influenced by what happens now.

The thing that changes is the rationing. Screen limits kept for a year of preparation tend to lapse the afternoon the last paper ends. Page 12's finding, held at the strength it deserves, suggests the variable to watch through the gap is not hours but whether she begins confiding in people she has not met.

THE SCEPTIC · DO NOT OVER-READ A SIX-WEEK WINDOW

Two cautions, in opposite directions. The first is against alarm: the whole screen-effects literature is arguing about single-digit percentages of variance, and the most-cited sceptical work puts it under one. A fortnight of unsupervised afternoons after a year of pressure is, on that evidence, unlikely to shape anybody. The second is against complacency dressed as evidence: *small on average* is not *small for her*, and an effect near zero across two point eight million children can still be large in the tail that matters. The defensible position is neither — it is that the research cannot presently tell a parent very much about one twelve-year-old, and that the people claiming otherwise in either direction are going beyond their data. *General guidance, not medical advice.*⁴

SOURCES 1 The Singapore Examinations and Assessment Board's official 2026 PSLE examination timetable, updated 13 February 2026, read as PDF at source via the board's own short link: every paper, date and sitting time printed here comes from that document. The marking-exercise date circulating on tuition sites is not in it and is flagged rather than printed as fact; the board states no results-release date, so none is given. 2 The ministerial announcement of 28 August 2026 on teen safeguards and the six designated services, reported by several Singapore outlets; the ministry's own newsroom was checked directly for the fortnight to 30 September and carries nothing on minors or age assurance, so the item is dated to August in the copy and marked secondary-sourced. 3 HardwareZone's parenting board, read directly; KiasuParents returned 403 and reddit.com 403 on every route, which is stated in the copy. 4 The comparison figures are those cited on **p12**, fetched and read at source; the reading is this desk's.

THE BIG PICTURE • THE LONG END

The Fed got a dovish surprise on Wednesday. The thirty-year sold off anyway

Core inflation undershot on both horizons and thirty points of October hike risk came out in a session. Duration fell anyway — the long end has stopped trading the policy rate.

Wednesday ran the experiment cleanly. At half past eight in Washington the August personal-income release put **core PCE at +0.2 per cent** on the month and **+3.0** on the year, against consensus of 0.3 and 3.3 — an undershoot on both horizons. Odds on an October hike fell from about seventy per cent to **thirty-nine**. Then the ten-year rose to **5.29 per cent** and the thirty-year to **5.64**, both 2026 highs.¹

THE CURVE • EIGHT SESSIONS, READ AS DATA

CLOSE	2Y	10Y	30Y	2S10S
21 Sep	4.76	4.96	5.29	+20bp
25 Sep	4.81	5.17	5.49	+36bp
28 Sep	4.92	5.24	5.56	+32bp
29 Sep	4.89	5.26	5.59	+37bp
30 Sep	4.88	5.29	5.64	+41bp

The Treasury's own par curve, parsed as data. Over the eight sessions the three-month moved **3bp** and the thirty-year **35**. Not advice.²

That shape has a name. Both ends rose and the long end rose faster: a **bear steepener**, which is what a market does when it wants paying to hold time rather than policy risk. The front end did the opposite — the three-month fell five basis points on Wednesday, pricing the easier Fed. The distance between those two behaviours is term premium, and it widened all week.²

-4bp

THE TWENTY-YEAR NOW YIELDS **MORE** THAN THE THIRTY (5.68 AGAINST 5.64), AN INVERSION HOLDING ON **61 OF 188 SESSIONS** THIS YEAR — A SUPPLY FINGERPRINT, NOT A MACRO SIGNAL²

The cause is not only American. Brent closed at **US\$97.92**, up 1.83 per cent, and four central banks tightened inside nine days — the ECB on 10 September, the Fed in mid-September to 3.75–4.00, the Bank of Japan on the 18th to **1.25** on a seven-to-

two vote, and Singapore's own authority in July. The BoJ's move came three months after its last, the shortest interval since 1990.³

5.7152

THE NAV OF THE LARGEST GLOBAL AGGREGATE BOND ETF IN DOLLAR ACCUMULATING FORM, AGAINST A 52-WEEK RANGE OF **5.71 TO 5.89**. NOT NEAR THE FLOOR OF ITS RANGE. THE FLOOR. DOWN 1.53 PER CENT THIS YEAR⁴

WHAT IT MEANS • FOUR INSTRUMENTS, ONE FORCE

World equities — VWRA at **US\$192.88**, 1.8 per cent off its record. The risk has stopped being earnings and become the discount rate: against a 5.29 per cent ten-year the equity risk premium is thinner than the index price admits. **Global bonds** — on the floor of their own year, because the long end is repricing supply, not policy. **Gold** — **US\$4,157**, down **6.39 per cent on the month**, in the middle of a Middle East war: real yields and a firm dollar are beating the haven bid. **Bitcoin** — **US\$83,719**, flat on the day. No liquidity impulse to trade, so it waits. Not advice; explanation.⁴

HIGHER-FOR-LONGER

BEAR STEEPENER

TERM PREMIUM

THE SCEPTIC • AGAINST THIS DESK'S OWN REGIME CALL

Three objections, and the third is serious. **First**, the inflation data already turned: a flat core year-on-year beside a 3.4 per cent headline is what a *level* effect looks like, not a persistence problem — in which case mid-September was the last hike. **Second**, term premium is the name strategists give the residual they cannot otherwise explain. Nothing fiscal happened in eight sessions; Brent went to ninety-eight and three central banks moved. This may be an inflation selloff in a fiscal costume, ending when oil does. **Third**: the equity market disagrees, and it holds the cash flows. In a genuinely tight regime the S&P would not be one per cent from its high and Singapore's exports would not be at their fastest since 1988. One of the two markets is wrong, and this desk has sided with the bond market.⁵

SOURCES 1 Bureau of Economic Analysis, Personal Income and Outlays for August 2026, 8.30am EDT on 30 September, read at source; the implied October odds from Investing.com's fed-funds monitor at 4.55pm EDT — the earlier ~70 per cent is a secondary report of CME FedWatch, whose page was unreachable, so only the direction is claimed. 2 The US Treasury's daily par yield curve, downloaded as CSV and parsed for every session shown. 3 The ECB, 10 September; the Bank of Japan, 18 September; Brent at the 30 September close. 4 VWRA via Investing.com; the bond NAV from iShares, dated 29 September; gold and bitcoin via TradingEconomics. 5 This desk's reading.

DOWN THE RABBIT HOLE · FREEDIVING & SPEARFISHING

The technique runs out at the same depth the physics does. That wall is where the hobby begins

Around thirty metres the lungs reach residual volume and the ordinary way of clearing your ears stops working. Everything enthusiasts spend money on is an answer to those two facts arriving together.

Two worlds share a breath. Freedivers go down a rope to a plate and come back; spearfishers hunt horizontally over reef and blue water. They share the physiology, the gear and the one absolute rule — never alone, one up and one down — and barely share a forum. The disciplines are formally named: static apnea, the only one scored in time; dynamic, scored in distance; and the depth events, with fins, without, or pulling down the line.¹

The mechanism is a reflex you already have. Submersion and breath-hold trigger bradycardia, peripheral vasoconstriction and a blood shift into the chest. In one study of eighteen subjects holding for a mean of 137 seconds, **spleen volume fell 20.8 per cent** — squeezing stored red cells into circulation — and haemoglobin rose 2.4 per cent, neither of which happened when the same subjects rebreathed for the same time.²

~30 m

WHERE A DIVER OF AVERAGE LUNG PROPORTIONS — SIX LITRES TOTAL, ONE AND A HALF RESIDUAL — REACHES RESIDUAL VOLUME. BELOW IT THE CHEST CANNOT COMPRESS FURTHER WITHOUT RISK, AND PLASMA FLOODS THE CAPILLARIES AROUND THE LUNGS TO BRACE THEM³

And at almost exactly that depth the standard ear-clearing technique fails. Valsalva — pinch and blow — uses the breathing muscles to push air up, which works with a tank in a horizontal position and stops working head-first on compressed lungs. Frenzel uses the tongue and larynx instead and reaches **thirty or forty metres**. Past that you carry a mouthful of air down, charged before the lungs compress out of reach.⁴

THE CHATTER · THE MOMENT PEOPLE DESCRIBE

The hook is not the depth. It is where buoyancy inverts and you stop swimming: “After residual volume, buoyancy flips. Positive becomes negative. Divers start to sink without effort, drawn downward like a stone in a vast blue sea.” The other version is the purchase — a competitor on buying his first carbon fins: “it really felt like the moment my commitment finally took shape.” Both come from a fin maker’s own journal, which is where a surprising amount of this sport’s writing lives. Reddit refused this desk entirely.⁵

How big is this world? Nobody publishes an honest answer. AIDA entries reached a high of **2,979 athletes in 2023**, the body has officiated **330 world records since 1992**, and its depth championship, running at Limassol as this issue goes to press, fields **178 athletes from 55 countries** — the largest depth field the sport has held. What does not exist is a certification count, or a credible market size. **No market number is printed here**, and the only survey of which card the community holds is **thirteen years old**.⁶

THE SCEPTIC · THE HARD NUMBERS ARE THE INJURY NUMBERS

There is no clean annual fatality count, and the literature says so. What is documented is worse in a different way. Adverse events in depth competition run at **three to four per hundred dives**, hypoxic complications in up to **ten per cent**. Across AIDA international team events 1998 to 2004, between **9.7 and 11.1 per cent of all performances** were disqualified for signs of severe brain hypoxia. At one world championship in 2006, fifty-seven divers produced **thirty-five adverse events**: three lost consciousness underwater, four surfaced coughing blood. The killing zone is the last few metres and the surface, because oxygen partial pressure falls on the ascent — so he feels fine until he does not.⁷

SOURCES 1 AIDA International’s own discipline definitions, read at source. 2 Splenic contraction and haemoglobin study, *Frontiers in Physiology*, read at source, with the control condition as described. 3 Residual-volume and blood-shift mechanics from a fin maker’s technical explainer and corroborated in Muth and Tetzlaff, *Deutsche Zeitschrift für Sportmedizin*, 2024; the heart-rate figures circulating with them come from the manufacturer alone and are not printed. 4 Molchanovs’ own teaching page on Frenzel against Valsalva, quoted. 5 Both quotations from alchemy.gr’s journal, 16 May 2022. 6 AIDA’s own competitive and calendar pages; the thirteen-year-old certification survey is a forum readership poll with no stated sample. 7 Muth and Tetzlaff 2024 for every incidence figure and the ascent mechanism.

THE DEEP END · WHERE THEY GATHER, AND WHAT THEY BUY

The world body has eleven videos on YouTube. One photographer has four hundred and fifty-nine thousand subscribers

The centre of gravity has moved to Southeast Asia, where the depth season runs October to May. And the busiest spearfishing board has seventeen thousand posts about guns and a safety section last used in 2022.

The community map is thinner than the sport's profile suggests. The largest spearfishing board carries **2,037 members, 6,738 threads and 88,353 posts**, of which the speargun section alone holds seventeen thousand, active this July. Its safety section's most recent post is dated **25 November 2022** and titled "I survived to a frightening and near-fatal diving accident!" Reddit was unreachable all morning, so no subreddit is characterised here.¹

Read the channel pages and the sport's public face resolves: **AIDA Freediving, 19,500 subscribers and eleven videos; Daan Verhoeven — a competition photographer, not an athlete — 459,000 subscribers;** a retired French champion who posts twice a year, 377,000. In public the sport belongs entirely to its artists.²

The gear, and the argument inside it. Long blades are the signature object, and the settled community view is the opposite of the sales pitch. From a competitor writing for a carbon house: "it was commonly said that the deeper you wanted to go, the stiffer your blades should be. That time is over... my legs were killing me on the way up. Since I switched to softer fins, these deep dives are so much easier." Softer blades allow smaller amplitude at higher frequency: less drag, less oxygen.³

The ladder. The plastic fin every dive centre owns costs about **A\$170** and is the one to skip. Fibreglass is the real entry: Molchanovs SPORT blades at **US\$375**, Australian-made DiveR composites at **A\$510**. Carbon runs Alchemy V3 at **€460**, Molchanovs PRO at **US\$800**. And both makers' guides converge on the same unglamorous conclusion: the **foot pocket**, a €190 item, does more for a diver than the blade bolted to it.³

Spearfishing is where the custom tier lives. The mechanism is band stretch: on a conventional gun a maker estimates **25 to 35 per cent** of the band's energy is lost to the oval shape of the rubber and the wishbone; on a roller, starting the band under the stock and running it through bearings cuts that to **five or ten**. A quality entry railgun runs **A\$400–490**, high-end and inverted rollers A\$770 to **A\$1,400**, the American bluewater end to **A\$3,200**.⁴

THE CHATTER · THE CUTTING EDGE, AND WHAT TO BELIEVE OF IT

Two carbon launches inside a month, contradicting each other. **Alchemy's ICON** (€800, 7 September) makes foot-pocket stiffness and internal volume *user-adjustable* on a competition fin, which addresses the most-reported complaint in the category; the fifteen-layer carbon and thirty-degree blade angle are not new. **VB Labs' Axiom**, announced 19 August with no price, leads on **40 per cent higher stiffness**. Note the collision: the industry's own best-argued position is that softer is more efficient. One of those claims is wrong, only one has a price attached, and no independent instrumented comparison of any two blades in this sport exists.⁵

THE SCEPTIC · THE FIN MAKER SAYS BUY PLASTIC

The strongest argument against spending is made by the company charging most. Alchemy's own writer, February 2025: "If you're a beginner, plastic fins are your best friend... Switching to carbon fins prematurely can be counterproductive. Carbon fins are incredibly responsive, meaning they'll amplify both good and bad technique. If your kick is messy or uneven, carbon fins will punish you." And two days before this issue — three weeks after launching the most expensive fin of the year — the same author: "A product can be beautifully engineered, technologically advanced and completely unnecessary... and still fail because it solves a problem nobody actually had." Whether that is candour or the most sophisticated marketing in the hobby is worth leaving open.⁶

SOURCES 1 Spearfishing.world's own board statistics and section indexes, read on 30 September 2026, including the dates given; DeeperBlue's forum read the same day. reddit.com returned a block on every route attempted. 2 Subscriber and video counts read directly from the YouTube channel pages, 30 September 2026. 3 Alchemy's journal, 16 May 2022, quoted at length, for the stiffness argument and the foot-pocket verdict; prices read from the makers' own stores and from Adreno, Australia's largest spearfishing retailer, on 30 September 2026. 4 Koah Spearguns' own technical description for the roller mechanism; the price ladder from Adreno and the makers' own pages. 5 Alchemy's ICON page and DeeperBlue's reports of 7 September and 19 August 2026; the absence of any named source in the Axiom coverage is this desk's observation. 6 Alchemy's journal, 10 February 2025 and 28 September 2026, both quoted verbatim.

THE DEEP END · THE CULTURE, AND THE COST

A thousand dollars buys the entry. Four thousand buys the endgame. The course is the part that keeps you alive

The Southeast Asian depth season opens this week, which is where the records now get set — and where a Singapore diver can be certified for less than the flight.

The vocabulary tells you where the difficulty is. Mouthfill and reverse pack are equalisation; packing is forcing extra air past full lungs, and carries documented costs. Samba is loss of motor control; blackout is what follows. One up, one down is the law. A lanyard ties the diver to the line so a safety team can retrieve someone beyond their reach — a rule written after the sport's first competition fatality in 2013, with mandatory sonar and compulsory incident reporting from 1 January 2015.¹

The economics are unusually honest for a gear hobby, because the courses are cheap. An online beginner module is US\$80; a two-day Level 1 in Singapore about S\$570, the same in Amed, Bali, about US\$275. Which is the shape of a sensible first ninety days: lurk, buy the course, then the pocket and the suit, then the blades.²

And the calendar has moved east. The Southeast Asian depth season runs October to May and is now where records happen. This week alone: a cup at Panglao, a depth event at Mabini, an official weekend at Cebu. Japan ran two events in September, and Singapore held its own pool and static opens at Changi Beach Club on 19–21 September.³

THE CHATTER · THE BEGINNER MISTAKES, AND THE SINGAPORE PROBLEM

The classic errors are consistent: hyperventilating before a dive, which delays the urge to breathe past the point the oxygen runs out; diving alone; over-weighting; chasing depth before equalisation; no lanyard. And for this city there is a structural one. The world's fastest-growing freediving agency is **headquartered in Singapore** — where, under the Parks and Trees Act, merely carrying a hunting device in a marine park is an offence. The capital of breath-hold education is a place where the water you train in is a swimming pool. A first kit of genuine quality runs about US\$1,000; the endgame, US\$4,300.⁴

THE SCEPTIC · THE TREADMILL, THE CARD, AND THE TURTLE

Four things the community concedes when pressed. **The card is a receipt.** Four agencies certify the same two-day skill set, one school will issue whichever you prefer off the same course, and one charges instructors US\$2,200 to cross over from another. **The suit is the real recurring cost** — open-cell neoprene is the warmest thing you can wear and the most fragile: lubricated to put on, torn by a fingernail, permanently compressed with use, so a suit that fitted in year one is loose and cold by year three. Tailor-made starts near US\$310, and nobody advertises the replacement cycle. **The depth discipline has a body count:** no-limits was abandoned in 2019 after repeated serious injury and death. **And spearfishing's conservation critique is peer-reviewed** — the method selects for long-lived, slow-growing species and removes the largest and most fecund individuals. This month it acquired a case: an instructor trainer was named in reports over a critically endangered hawksbill turtle speared in Raja Ampat. *One publication carried it, no charges had been filed, and the man denied involvement.*⁵

NEXT DESCENTS — YOU PICK TOMORROW'S HOBBY

Amateur radio (licensed operating & POTA) — the licence, the antenna you build, and activating a park.

Leathercraft — the hand saddle-stitch, veg-tan, and a tool ladder you buy once.

Scale model building (armour & aircraft) — the airbrush ladder and a weathering culture with schools.


SOURCES 1 Packing risks from Tetzlaff and colleagues, *Frontiers in Physiology*, 9 July 2021, and Muth and Tetzlaff 2024; the lanyard's purpose quoted from AIDA's medical officer writing in *Alert Diver Europe*, 26 April 2022; the 2015 rule package as reported by DeeperBlue, 6 August 2014, following the 2013 fatality. 2 Course prices read at the schools' own pages, with kit totals as itemised on p16. 3 AIDA's own calendar; the Bahamas organiser's own site, which now lists a 2027 course season and no competition. 4 Molchanovs' about page; Singapore's Parks and Trees Act 2005, section 9. 5 Cross-over pricing from the agency's instructor pages; the selectivity literature from Lloret and colleagues, 2008 and 2020. The Raja Ampat case is DeeperBlue's, 16 September 2026, alone, and is flagged unproven.

THE LONG GAME · THE DOSE

The curve bends hard in the first hundred and fifty minutes, then very nearly stops

A meta-analysis published yesterday puts the plateau near thirty MET-hours a week. Everything the optimiser buys past that, he is buying for some reason other than mortality.

Seventeen prospective cohorts, **176,862 participants, 31,783 deaths.** At roughly **150 minutes a week** of moderate activity, all-cause mortality ran about **38 per cent lower** than in the sedentary. The association is inverse and non-linear: steep at low doses, then flattening, with estimates turning imprecise beyond **30 MET-hours a week.**¹



The shape, not the number. A *Meridian* rendering of the published dose-response relationship; the curve is illustrative of the reported form, not a reproduction of the authors' figure.¹

And then the limitation that decides how to read it: **every participant already had coronary heart disease.** A screened, personally-trained man in his forties is not in this sample. What transfers is the shape, not the 38 per cent — in every population where this has been measured, the return on the first 150 minutes is enormous and on the next 150 small. The instinct that more is more is the part the curve will not support. Only **8 of the 17** studies reached the highest thresholds.¹

7.5 hours

OF BRISK WALKING A WEEK IS ROUGHLY WHERE THE PLATEAU SITS — 30 MET-HOURS AT ABOUT FOUR METS. HALF THAT IN GENUINELY MODERATE-TO-VIGOROUS WORK. PAST IT, YOU ARE BUYING SOMETHING OTHER THAN ALL-CAUSE MORTALITY¹

The second paper is a first: **23 experts across 12 countries**, three Delphi rounds, **122 statements**, of which consensus was reached on **72** — sixty per cent. A warm-up is defined as a structured preparatory intervention in three layers, general then specific

then a conditioning activity intense enough to induce post-activation enhancement, and “the comprehensiveness of a warm-up protocol is a key factor associated with its effectiveness”. It is also monitorable: readiness and fatigue can be read during the warm-up itself, which makes it a diagnostic rather than a preamble.²

THE CHATTER · THE COMMUNITY'S REAL JOB IS CHECKING THE UNITS

Reddit was unreachable, and Singapore's readable fitness board had no longevity thread at all — its live discussion is a gym chain going into voluntary liquidation, which is its own comment on the market. The argument is in the comments under a practitioner post on what counts as a good VO₂max, where Norwegian seventy-somethings are set against American percentile tables built from **22,379** laboratory tests. **Lawrence McDaniel**, seventy-four, reports 51 at altitude on ten hours a week. And **Gene Barth** caught a unit error in the published piece itself, the stated figure implying implausible energy expenditure. That is the discourse at its most useful: not debating the conclusion, but checking the sums of the person they follow.³

THE SCEPTIC · THE PANEL'S OWN NULL, AND YOUR INSTRUMENTS

The most striking line in the consensus is a negative: while some warm-up exercise may reduce **long-term** injury risk, “there is no support for acute injury protection”. Twenty-three experts, and the most widely believed thing about warming up did not survive. On the meta-analysis, the standing objection to seventeen observational cohorts is reverse causation — sicker people move less, and the sickest die soonest — and the plateau is always the least reliable part of a curve, which is exactly where an optimiser will read. **On the instruments:** Garmin appears to be the only brand with roughly seven per cent error on VO₂max across three independent studies; for several popular rings and straps **no independent peer-reviewed validation exists at all.** That reached this desk through summaries rather than the primary papers, and is flagged. Not medical advice.⁴

SOURCES 1 Tao, Hu, Chen and Strain, “Non-Occupational Physical Activity and All-Cause Mortality in Patients with Coronary Heart Disease: A Dose-Response Meta-Analysis”, *European Journal of Preventive Cardiology*, 30 September 2026, advance article read at the publisher's site: the cohort and participant counts, the death count, the 38 per cent figure at 150 minutes, the non-linear form, the imprecision beyond 30 MET-hours, and the authors' own statement that only eight studies reached the highest thresholds. The MET conversion of that plateau into walking hours is this desk's. PubMed served a cookie wall, so the abstract was read on the publisher's site instead. 2 Boulosa and colleagues, *International Journal of Sports Physiology and Performance* 21(9), 20 July 2026; abstract read through Europe PMC after the publisher returned 403. 3 A practitioner Substack post of 1 September 2026 and its comment thread, quoted under the commenters' own names. 4 The consensus paper for the quoted null; the device-validation figures reached this desk through summaries rather than the papers themselves.

THE DIARY · THE DATED AGENDA

Two Saturdays, two stadiums, and the one date to set a reminder for is in December

Every row below was re-checked against a venue, promoter or official seller. Several candidates failed that test and were dropped; they are named at the foot.

SINGAPORE

BOOK NOW · PLAN AHEAD

- 2-3 OCT

★ **The Weeknd — After Hours Til Dawn**

NATIONAL STADIUM, 8PM, TWO NIGHTS · FROM **S\$158**, ONE OUTLET ONLY; THE PROMOTER PUBLISHES NO PRICES¹
- 19 SEP-8 NOV

teamLab: Infinite Existence

SUPPER HOUSE, KEPPEL ROAD · THE STUDIO'S **FIRST SOLO SHOW IN SINGAPORE** · **FREE**, THU-SUN — THE BEST FREE HOUR IN THE CITY FOR A TWELVE-YEAR-OLD²
- 2-3 OCT

SAM30 Festival Weekend

TANJONG PAGAR DISTRI PARK · FREE GALLERIES FRIDAY EVENING; A RYOJI IKEDA LATE NIGHT **S\$29**³
- 11 OCT

Hiroshi Sugimoto: Form Is Emptiness — closes

SINGAPORE ART MUSEUM, EXTENDED ONCE · 63 WORKS AND FOURTEEN FOSSILS FROM HIS OWN COLLECTION — THE FOSSILS ARE THE HOOK · **S\$15**³
- 17-18 OCT

★ **BIGBANG — XX : COSMOS**

NATIONAL STADIUM, 7PM · THE ONLY FULLY-PRICED ROW HERE: **S\$158** AT CAT 6 TO **S\$498** VIP STANDING · A SECOND NIGHT ADDED AFTER THE FIRST SOLD⁴
- 21 OCT

Junior and the Holiday Magic

DRAMA CENTRE THEATRE, NO CLOSING DATE PUBLISHED · S\$28-52 · **AIMED AT FOUR AND UP** — NOT FOR A TWELVE-YEAR-OLD⁵
- 31 OCT

★ **Jujutsu Kaisen in Concert**

ESPLANADE THEATRE, 3PM AND 8PM — THE SINGAPORE DEBUT, TWELVE MUSICIANS AGAINST VISUALS FROM ALL THREE SEASONS · **TAKE THE MATINEE** · S\$58-199, SINGLE-SOURCED; RUNNING TIME DISPUTED⁶
- 25 NOV

Guns N' Roses

NATIONAL STADIUM, 8PM, A WEDNESDAY · FROM S\$138, SINGLE-SOURCED · LATE AND MIDWEEK⁴
- 28-29 NOV

BABYMONSTER — CHOOM

INDOOR STADIUM, INSIDE THE SCHOOL HOLIDAYS · FROM S\$168, FIRST NIGHT REPORTED GONE — ONE OUTLET ONLY¹

JAPAN & THE REGION

THE ONES WITH A TICKET CLOCK

- 5 DEC

★ **Hatsu Basho tickets open**

FOR THE 10-24 JANUARY TOURNAMENT AT THE RYOGOKU KOKUGIKAN — **THE DATE TO DIARISE**. CHAIR SEATS RUN **¥4,000-5,500**; NOVEMBER'S FUKUOKA TOURNAMENT IS ALREADY **SOLD OUT, ALL FIFTEEN DAYS**⁷
- 28 NOV

★ **Niseko United opens**

ALL FOUR RESORTS, CLOSING 5 MAY · NIGHT SKIING FROM 12 DECEMBER · PEAK DAY PASS **¥13,500** — IT OPENS AS THE SINGAPORE SCHOOL HOLIDAYS BEGIN⁸
- 29 SEP-20 DEC

Hiroshige: New Perspectives on Edo Japan

TOKYO NATIONAL MUSEUM · ALONGSIDE A KOFUN-PERIOD HORSE SHOW FOR THE YEAR OF THE HORSE — THE KINDER OF THE TWO FOR A CHILD⁹
- 11-13 DEC

★ **Tokyo Comic Con, tenth edition**

MAKUHARI · THREE-DAY PASS **¥9,900** — AND **UNDER-FIFTEENS ¥500**, THE BEST-VALUE TICKET IN THIS ISSUE. NICOLAS CAGE MAKES HIS TOKYO DEBUT¹⁰
- 8 JAN-4 APR

Neverending Journeys — Final Fantasy at 40

HYOKEIKAN, UENO · A FIRST COLLABORATION BETWEEN MUSEUM CULTURAL PROPERTIES AND THE SERIES · ON SALE, BUT **NO PRICE IS PRINTED** — THE TICKET SITE REFUSED THIS DESK⁹

PLANNING NOTES · WHAT THE CLOCK IS DOING

The Grand Prix, eight days out. Hospitality and every three-day grandstand this desk could see have gone. What is left is walkabout: **Zone 4 Friday at S\$198**, marked selling fast, Saturday S\$298, Sunday S\$368. Friday is the family answer — practice day, thinner crowds, an earlier finish. **Do not buy on resale platforms:** the organiser says it will confiscate such tickets without compensation. **Dropped for failing verification:** a November arena date absent from its venue's own listings; the Singapore marathon, whose site still serves last year; Christmas Wonderland and the Tokyo illuminations, both still on 2025 dates. **Post Malone stays off this page** — postponed in August, still no new date.¹⁰

SOURCES Every row re-verified on 1 October 2026 against a venue, promoter or official seller. 1 The artist's own tour page and Live Nation Singapore; prices are one outlet's, Ticketmaster Singapore having returned 403. 2 teamLab's own site. 3 The Singapore Art Museum's own pages. 4 The Kallang's own event pages, which carry BIGBANG's full price table; the first night's sell-out is one outlet's claim and the venue does not state it. 5 Singapore Repertory Theatre's own show page. 6 The concert's official tour site; the price and conflicting running times from local coverage, SISTIC serving nothing readable. 7 The Sumo Association's own ticketing pages. 8 Niseko United's own lift-ticket page. 9 The Tokyo National Museum's own pages; the Final Fantasy ticket site returned an access denial, and the convention's own ticket page carries its prices. 10 The Grand Prix organiser's own ticket and reseller pages, read this morning.

THE TABLE • SINGAPORE'S NEWEST ROOMS

Five openings were offered to this page. Three had already run in it, and the ledger caught them

One of the two survivors has been reviewed by somebody who paid. The other was tasted on the house. In a market that shut two thousand kitchens in seven months, that is the whole honest inventory.

The one with a verdict attached is **AO**, on Claymore Drive, opened in September under culinary director Araki Takeshi, formerly of Esora. It is Japanese craft applied to flavours from anywhere — amaebi carpaccio at **S\$26**, a sukiyaki wagyu maki roll at **S\$28**, a sea-bream ramen at the same — in a room of wood and stone, and you enter *through* the bakery next door, which shares its address and its owner. A reviewer who dined and appears to have paid scored it **four out of five**, and supplied the dissent in the same breath: the tables sit close together and there is no private room. Daily, eleven to ten; not halal-certified; go before noon or after eight if you want the room.¹

Yamamoto's Hamburg returns to this page for its verdict run, eight issues after it was first flagged. It soft-launched inside Lau Pa Sat on 22 September and opens properly on the 5th — the first Singapore outlet of a Tokyo hamburg-steak chain founded in Ebisu in 2005, with around twenty branches in Japan, occupying its own air-conditioned unit in the middle of a hawker centre. Hamburgs run **S\$15.90 to 19.90**, the signature stuffed with gorgonzola and mushroom cream under a miso-dashi demi-glace; a meal lands at **S\$20–24** before service. It has been reviewed positively — but *on a media tasting*, the reviewer's pick being the tamagoyaki version: "*if you're only ordering one dish, make it this*". No published dissent exists, which is a caveat and not an endorsement.²

THE BUZZ • THREE NAMES CUT BY THE LEDGER, NOT BY TASTE

This desk opened with five and printed two. The other three — a bakery sharing that same Claymore Drive address, a Bangkok coffee brand newly landed at Changi, and a Korean tonkatsu group arriving at Raffles City — were each offered as a new opening, and each has appeared on this page before: two in No. 113, the third in No. 116. Under the rule the editor imposed last month a venue gets two runs, at least seven issues apart. None had waited, and all three were cut at the gate.³

That rule exists because a reader counted one restaurant running four times in thirteen issues, and because this desk keeps being handed "new" lists that are nothing of the kind. Among the openings offered this week: a room that opened in **March 2025**; a second outlet of an established name; a hotel café of twenty-six years relaunched under another; two more second outlets and a beach-club residency. An announcement is not an opening, and the only filter that works is a list of what you have already printed.³

THE HONESTY LEDGER

Ate and judged: AO alone. **Tasted on the house:** Yamamoto's Hamburg — disclosed by the reviewer, with no published dissent to weigh against it. Neither venue is in any Michelin listing this desk could read, and neither has a website; both rest on food-media reporting. Reddit, Burpple and Instagram all refused this desk, so **no forum sentiment appears here and none was invented**.

2,101

FOOD BUSINESSES THAT CEASED OPERATING IN SINGAPORE BETWEEN JANUARY AND JULY, **UP 25.1 PER CENT** YEAR ON YEAR — AGAINST 3,148 ACROSS THE WHOLE OF 2025 AND 3,047 IN 2024⁴

THE SCEPTIC • CHURN, OR COLLAPSE

The obvious reading is a market breaking, and the closure column supports it: a patisserie closed to "*rising operating costs and high rents*", Fika after seventeen years, a churrascaria after twenty. **The better argument is the contrarian one**, from Tan Hsien Wei, senior lecturer at Temasek Polytechnic's business school: churn is not decline. "*Competition has intensified as diners have more choices. Social media accelerated dining trend cycles, testing business models rigorously.*" He argues consolidation can be "*deliberate and strategic*" — groups pruning to protect a brand rather than failing. The register supports him: **2,594 new food businesses registered** across the same seven months. Openings still outrun closures, and by a clear margin. What the figures cannot tell you is which column this month's imports will end up in, and anyone who claims otherwise eight days after a soft launch is guessing.⁴

IMPORTS

VERDICT RUN

NO FORUM READ

SOURCES 1 A signed blog review of 26 September 2026 in which the writer dines and records the verdict and the criticism quoted; the pedigree, prices and hours from the restaurant's food-media coverage, the venue having no site of its own. 2 A local review site's piece of 25 September, which discloses a media tasting. 3 state/table- Ledger . json, this publication's own record, read before selection. 4 Company-registry cessation and registration data as reported by two outlets; the Temasek Polytechnic lecturer and the patisserie founder are quoted from that reporting.

THE UNDERCURRENT · PAKISTAN

The regulator tried to stop the middle class free-riding. It made a battery worth four times a grid connection

Islamabad cut what it pays households for their surplus solar to about eleven rupees a unit. The unit they keep is worth forty to fifty-five. The customers did not comply. They left.

On 9 February 2026, Pakistan's electricity regulator approved the Prosumer Regulations and ended net metering. The old arrangement, in place since 2015, had let a household with panels export its daytime surplus and take back an equivalent unit at night, credited at something near retail — roughly **25 to 27 rupees** a unit. The grid was, for that household, a free battery.¹

Under net billing, imports are charged at the retail tariff of **40 to 60 rupees** and exports are bought at the national average purchase price, about **10 to 13**. The contract term fell from seven years to five. A non-refundable concurrence fee of a thousand rupees a kilowatt appeared. No new connections would be allowed where distributed generation had already reached 80 per cent of a transformer's capacity.¹

The intention was defensible, and the minister stated it plainly in the Senate the following day: "Solar is generating electricity at five to six rupees per unit. Is providing net metering at 27 rupees fair to other consumers?" Pakistan's tariff is stuffed with capacity payments — fixed, dollar-indexed, take-or-pay obligations to independent power producers, amounting to **14.21 rupees a unit, or 82 per cent of what the consumer pays**. Every household that stops buying leaves that same fixed bill to be divided among fewer units. On its own terms the reform was an attempt to stop a subsidy flowing upward.²

A unit kept in a battery displaces an import worth forty to fifty-five rupees. The same unit sold to the grid earns eleven.

What the regulation actually did was published, four months later, by the government's own economists. The Pakistan Institute of Development Economics set it out without euphemism: "When a prosumer stores surplus electricity rather than exporting it, each stored unit displaces an imported unit worth PKR 40-55 instead of earning PKR 11. The Regulations therefore unintentionally encourage behind-the-meter storage adoption."³

4.6 GWh


OF LITHIUM-ION STORAGE IMPORTED INTO PAKISTAN IN 2025 ALONE — **60 PER CENT** OF EVERYTHING IMPORTED SINCE 2018, AND A **220 PER CENT** RISE ON THE YEAR BEFORE. RESIDENTIAL IMPORTS NEARLY QUADRUPLLED, FROM **0.62 TO 2.45 GWH**⁴

The conditions were already laid. Between 2020 and 2024 the average electricity tariff rose **155 per cent**, to around 44 rupees a unit — roughly double India's or Bangladesh's, and the highest among Pakistan's neighbours — under pressure from an international lending programme to remove subsidies. Over the same period Chinese module prices collapsed to about **27 US cents a watt**, and by 2025 one Pakistani account puts them at eight. In 2024 the country imported **17 gigawatts** of panels, more than the whole of Africa.⁵

By the 2025 financial year, **7.3 million** of Pakistan's roughly 40 million households had solar on the roof — nearly one in five, up from three million two years earlier. But only about **282,000** of them had a battery: one in twenty-six. The whole of the next decade of the country's energy politics sits in that gap, and in February the state accidentally told everyone how to close it.⁴

SOURCES 1 The Prosumer Regulations 2026, approved 9 February 2026 following a hearing on the 6th, with the tariff structure, contract term, concurrence fee and transformer threshold as described; a draft amendment of 25 February, signed by a named director and backdated to the 9th, protects existing net-metering holders at their old export rate until their agreements expire. 2 The minister's remarks in the Senate as reported in Urdu by *Jang*, 10 February 2026, translated for this desk; the capacity-payment figures are from the regulator's State of Industry Report 2025 as reported by two Pakistani outlets after the report's own PDF proved too large to fetch. 3 Pakistan Institute of Development Economics, Knowledge Brief No. 2026:149, May 2026, by a research economist and a staff economist, read in full at source; the passage is quoted verbatim. 4 Renewables First, *From Solar Panels to Storage: Pakistan's Battery Boom Begins*, September 2026, PDF read in full: the import series year by year, the sectoral split, the household and storage counts. Its import figures are built from Chinese customs mirror data rather than Pakistani customs, and its sectoral splits are stated assumptions — this is an estimate of a market the Pakistani state does not measure, and is treated as one here. 5 The same report for the tariff rise and module prices; the eight-cent figure and the 2024 panel-import volume are from *Dawn*.

The mechanism is not complicated, which is what makes it hard to stop. A household with a ten-kilowatt array already spent one and a half to two and a quarter million rupees on it, and generates twelve to fifteen hundred units a month. Adding a fourteen- to sixteen-kilowatt-hour lithium pack, landed from China at somewhere between seventy and a hundred and ten US dollars a kilowatt-hour, converts the surplus it used to sell at eleven rupees into an import it no longer has to buy at fifty. Published estimates put the payback on solar with storage at **three to five years** for a home, four to six for a business — and that is *after* a measured **48 per cent** cost increase from surcharges and duties on the batteries themselves.¹



Then the household disappears. It keeps its meter and now pays a fixed charge — from January 2026 every domestic consumer bar the lifeline tier pays between **200 and 675 rupees per kilowatt of sanctioned load every month**, where previously such charges applied only above three hundred units. But it buys almost nothing. National grid demand fell **8.66 per cent** in the 2026 financial year. In Islamabad's distribution company, April sales were down 10.77 per cent overall and **19.69 per cent residential**.³

111 TWh

GRID SALES IN FY2025, DOWN FROM **125** THREE YEARS EARLIER — AN 11 PER CENT FALL WHILE THE NUMBER OF CONSUMERS GREW 5.3 PER CENT. CONSUMER-OWNED GENERATION NOW PRODUCES ABOUT **46 PER CENT** OF WHAT THE GRID SELLS³

And the state cannot see it. Official installed capacity is **41 gigawatts**; add distributed solar and the real figure is roughly **79**. How much roughly is the problem: the minister has said 22 gigawatts of total solar, a think tank says 38, another count says 31, and

the regulator's net-metering register says 6 to 8. A researcher at Renewables First put the consequence in one sentence: "*We are making billion-rupee decisions based on data that misses nearly half.*"⁴

The counter-move is already drafted, and it is an admission. Having lost the volume, the Power Division now wants the *timing*: an adviser has proposed paying households **18 to 22 rupees a unit** to discharge their own batteries into the grid between five and ten in the evening, when peak demand passes 26,000 megawatts. That is more than the state pays the same household for its daytime export, and roughly twice what it pays under the rule that caused the problem.³

When the Strait closed in April, the official grid lost up to five thousand megawatts overnight. A thirty-eight-gigawatt system the government does not count carried on.

Not a figure of speech. On **13 and 14 April 2026**, when the Hormuz closure triggered a force-majeure declaration on gas cargoes, four to five thousand megawatts of generation vanished and up to five hours of load-shedding returned. The households with panels and packs did not notice. The ones without did.⁵

Who those are is the part that decides whether this is a success story. In the old city of Lahore, a reporter for the Urdu outlet Lok Sujag found the households that cannot follow. **Muhammad Imran**, on the second floor inside Mochi Gate for twenty-five years, his parents below and two brothers above: "*One brother has built a cage for pigeons on the roof; and there are tall buildings on both sides of our house, because of which very little direct sunlight reaches here. So solar panels cannot be installed at our place, and I cannot see any alternative for electricity.*" **Hidayatullah Khan**, renting a ground floor beside Gulberg, would have installed a system long ago; his landlord will not permit it on the roof. *That reporting dates from 2024, and is used here as the older document it is.*⁶

SOURCES 1 Household system costs, generation and savings from a Pakistani business outlet of 11 August 2026, which attributes them to a study and to industry estimates without naming either and is treated here as single-sourced; the payback bands and the 48 per cent cost increase are from the Institute for Energy Economics and Financial Analysis. 2 Renewables First, *From Solar Panels to Storage*, September 2026, for the annual import series charted, with its own methodological caveat carried in the caption. 3 The fixed-charge bands as notified from January 2026; grid sales, the consumer-growth figure and the distribution-company monthly falls from a Renewables First study reported in *The News*, 21 May 2026, and a Pakistani business outlet of 17 August 2026. 4 The fixed-charge bands as notified from January 2026; grid sales and the distribution-company falls from a Renewables First study reported 21 May 2026. 5 The competing capacity estimates as published by the minister, by Renewables First and by the regulator, printed together because the disagreement is the finding. 6 *The News*, 16 April 2026. 6 Lok Sujag, first published 13 July 2024 as its own page states, translated for this desk.

The strongest argument against this story comes from the government's own institute, and it deserves to be stated at full strength. The Pakistan Institute of Development Economics examined whether net metering explains the distribution companies' finances and concluded that it does not: "net metering alone cannot explain the financial difficulties experienced by DISCOs. The conventional structural drivers of inefficiency including consumption loss, theft, poor recoveries, deferred investments and tariff distortions continue to dominate."¹

The numbers behind it are decisive. Net-metered exports in 2024-25 came to **1,997 gigawatt-hours**, putting about **53 billion rupees** of retail revenue at risk, which net billing reduces to roughly 22 billion. A saving of some **32 billion rupees**, against a circular debt of **1.8 trillion**. The institute's stress index ranks the worst-afflicted utilities as those serving Balochistan, rural Sindh and Khyber Pakhtunkhwa — places with almost no solar, distribution losses of **17.55 per cent** against an 11.43 benchmark, and recovery below forty per cent in places. The utility with by far the most solar is not among them.¹

The grid was dying before the batteries arrived. What the batteries changed is that it can no longer be saved by the only instrument the state has, which is the tariff.

So the honest sentence is not that storage is killing Pakistan's grid. It is that storage has made the grid impossible to repair by raising prices — and raising prices is the entire policy. The institute concedes the trap in its own conclusion: net billing "may alter the form of financial pressure on DISCOs rather than eliminate it, potentially deepening volumetric sales losses over time", because a household that stores rather than exports removes not only the export settlement but the evening sales, which are the valuable ones.¹

Which is why a National Battery Framework is now being written: a steering committee, safety standards, interconnection protocols and a

registration regime — because the state cannot presently see where five gigawatt-hours a year of imported storage is going.²

1 in 26

SOLARISED PAKISTANI HOUSEHOLDS THAT CURRENTLY HAVE A BATTERY. IF THAT RATIO MOVES TOWARD THE RATIO OF HOUSEHOLDS WITH PANELS — NEARLY **ONE IN FIVE** — THE GRID'S REMAINING CUSTOMERS ARE THE PEOPLE WHO COULD NEVER LEAVE³

That is the equity of it, and it is the part nobody has solved. A strategist quoted this year put it plainly: "If everyone leaves the grid, those left behind will end up paying more — and in many cases, these are the poorer families who cannot afford to buy panels and batteries." A researcher at Pakistan's Sustainable Development Policy Institute called net billing "a tax on the sun" and warned of a self-reinforcing cycle in which the grid becomes a poverty trap. Both are describing one calculation from opposite ends.⁴

WHAT THIS DESK DOES NOT KNOW

The denominator. If the true installed base is 22 gigawatts rather than 38, the defection is half as advanced as the think tanks believe. Every count disagrees with every other count, and the state publishes the smallest one.

The buyer's intent. A country with chronic load-shedding has always bought backup power. That the chemistry has shifted to lithium tells you something has changed; it does not tell you whether the household bought a battery to leave the grid or merely to survive it.

What can be said is what Pakistan is, which is the live experiment. Singapore is building its energy future on the opposite bet — imported power, a regional grid, a tariff set centrally, as **p5** records twice this morning. The curve that broke the Pakistani settlement is not a Pakistani curve. Lithium at seventy to a hundred and ten dollars a kilowatt-hour and panels at pennies a watt are arriving everywhere, and they arrive as a question no regulator has answered: what a utility is for, once its best customers can generate and store more cheaply than it can bill them.⁵

Islamabad is about to attempt the only answer available to it. Having failed to keep its customers as customers, it proposes to become their tenant — renting back, at dusk, the storage they bought to escape it. Whether that is the invention of a national virtual power plant or the last act of a monopoly is a question the next two budgets will settle.²

SOURCES 1 Pakistan Institute of Development Economics, Knowledge Brief No. 2026:149, May 2026, read in full at source: both quoted passages, the export volume and revenue-at-risk figures, the stress-index rankings and the loss and recovery rates. The circular-debt figure is the regulator's, as reported by two Pakistani outlets in January 2026 and updated to February. 2 The time-of-use proposal was made by a named adviser and reported by a Pakistani business outlet on 4 August 2026; the battery framework is the Power Division's, as reported in the same period. 3 Renewables First, September 2026, as on **p21**. 4 The strategist is quoted in a September 2025 report by Climate & Capital Media; the institute researcher in *The News*, 9 March 2026. 5 This desk's argument.

MERIDIAN

THE LAST WORD

What the rule did instead

Pakistan's regulator set out in February to stop comfortable households free-riding on a bankrupt grid, and wrote a rule that made a lithium battery worth four times a grid connection. The households did not object. They bought batteries. By the government's own reckoning a unit kept at home is worth forty to fifty-five rupees and the same unit sold back is worth eleven, and once that was true the policy's intention stopped mattering.

Singapore's fare formula permitted 14.7 per cent this week and granted 7.0, which sounds like restraint until you notice that the unpaid remainder is now larger than the formula's own annual output. A nuclear threat arrived in Brussels in writing, promising retaliation for a blockade that appears in no published plan. A patch that ran ninety-one per cent faster was vetoed in fifteen hours, and the tester who tried to reproduce it on three cards watched it collapse. Micron has sold supply for a year that has not happened yet, on deposits, against cleanrooms that will not exist until 2030.

None of those is a story about a rule being broken. Each is a story about a rule working exactly as written and producing something nobody in the room intended — which is the ordinary condition of governing anything, and the reason the interesting question is never what a policy is for. It is what the policy makes cheap.

Twelve Desks · One World MERIDIAN · No. 119 · 1 October 2026

A note on method. This edition was researched with working network access, and its load-bearing figures were read on the pages that carry them: the US Treasury's own daily par yield curve, downloaded and parsed as data; the European Central Bank's reference rates; the Public Transport Council's fare-review release with its annexes; SP Group's tariff page and its fourth-quarter release, the second of which is served under a slug naming the wrong quarter of the wrong year; Micron's SEC filing and prepared remarks; a llama.cpp pull request opened and vetoed inside a day; the Pakistan Institute of Development Economics' own knowledge brief; two peer-reviewed papers read in full; and the Singapore examinations board's published timetable. Where a figure could not be read at its source, it is attributed to whoever did publish it and flagged as single-sourced. **Refused this desk today:** reddit.com on every route, which is why several desks report a null result rather than a forum sentiment; Reuters, the Associated Press, CNN and CNBC at article level; Ticketmaster Singapore, SISTIC, Burpple, Instagram and the Michelin Guide's Singapore listing; CandlePowerForums, BladeForums and three watch forums; the Straits Times, CNA and the Business Times; and Billboard, behind a paywall returning a payment-required error. **Corrections in this issue:** No. 118 reported SP Group's tariff release as unpublished and late, and it was neither (p5); and reported a Japanese release as having opened at number one when it opened at number two (p11). A third item — the date of Hungary's prosecutor-general vacancy — was checked against its source and No. 118's reading stands.