NO. 130 · SINGAPORE

THE CURRENT BRIEFING OF

MERIDIAN

Twelve desks. One world.

MONDAY 12 OCTOBER 2026



THREE STRIKES ON ONE AIRPORT IN A WEEK, AND NO ALERT EITHER TIME (P3)

THE WORLD

No alert was issued. Thirteen died at King Khalid, and Washington now says it may join

SINGAPORE

The haze ladder finally got its numbers. The index stopped eight points below the first rung

THE RABBIT HOLE

Here the offence is purpose, not possession — and the last four charged carried screwdrivers

Twelve Desks · One World

Source: Stringer / Reuters, via France 24. File photograph of 19 September 2026 showing smoke at King Khalid International Airport, Riyadh, during an earlier strike in the same campaign — not the attack of 10 October. Cropped; no other alteration.

THIS ISSUE · NO. 130

Contents

p3 The World — no alert was issued, and thirteen died
p4 The World — a corridor that opens nothing
p5 Singapore — eight points short of the first rung
p6 The Atelier — the rim the glass will not stick to
p7 Technology — six billion announced, three itemised
p8 Technology — the speed-up that costs the prompt
p9 The Kit — the specification is behind an email form
p10 The Good Life — seventy-seven pieces, built twice
p11 Screen & Sound — eleven critics against two hundred and four
p12–13 The Family Desk — the phone is the hinge; the hormone stops explaining
p14 The Macro Desk — the Fed hinted, and Singapore's banks took it
p15–17 The Rabbit Hole — locksport, and a law about purpose
p18 Fitness — triple the protein, and p equals 0.726
p19–20 The Diary — Brazil twice in a week; a second star
p21–23 The Long Read — human feed, and the restaurant that lost the customer
THE SPINE: the null — the threshold not crossed, the effect not found, the instrument that did not fire

FROM THE DESK

"To the curious incident of the dog in the night-time." "The dog did nothing in the night-time." "That was the curious incident."

Arthur Conan Doyle, "Silver Blaze", *The Strand Magazine*, December 1892[1]

Holmes solves the case on a negative: a watchdog that should have barked and did not. This issue is full of them. A warning system that issued nothing. A ladder whose first rung was never reached. Two trials whose fashionable variable returned an interval straddling zero. The spine is **the null** — and the discipline is to notice it rather than fill it in.

THE STRIP · AT FRIDAY'S CLOSE

US 2-year	4.80%	▼ 3bp
US 10-year	5.24%	▼ 4bp
VWRA	$193.50	▲ 0.39%
Global agg	$5.731	▲ 0.10%
DXY	102.21	▲ 0.27%
USD/SGD	1.2807	▲ 0.04%
Brent	$104.72	▲ 2.42%
Gold, COMEX	$4,216.30	▲ 1.30%
Bitcoin	$82,546	▼ 2.31%

Friday to Friday. The Straits Times Index closed 5,401.89, down 4.13% on the week (p14). Bitcoin added about 1.1% over the weekend. Not advice.[2]

FROM THE ARCHIVE · NO. 128, 10 OCTOBER 2026

We have been here before, two issues ago. The Atelier of No. 128 opened on a document that no longer exists: IEEE 1789-2015, the recommended practice for modulating current in high-brightness LEDs, **"marked Inactive-Reserved as of 26 March 2026"**, with nothing put in its place.[3]

Its rule was the usable kind. Flicker not to exceed frequency times 0.08 — **10 per cent at 120 hertz**. A number a buyer could hold a lamp against. Withdrawn, it leaves the category measuring nothing.[3]

Which is this issue in miniature. A standard that lapses, a ladder never climbed, a licence that fires and a condemnation that does not (p5).

SOURCES 1 Arthur Conan Doyle, "Silver Blaze", first published in *The Strand Magazine*, December 1892 and collected in *The Memoirs of Sherlock Holmes* (1894); the exchange is between Inspector Gregory and Holmes. 2 Levels as set out on the Macro Desk, from the Treasury's own par yield series and one commercial provider's closes for 9 October 2026; that dependence is stated there. 3 This magazine's own No. 128 of 10 October 2026, read in the archive.

THE WORLD · RIYADH & DAMMAM

No alert was issued. Thirteen died at the terminal

Saturday's missile reached the passenger hall of King Khalid without any warning at all. On Sunday a second struck Terminal 4, and Dammam the same afternoon.

The count came from the airport's own regulator. Saudi Arabia's civil aviation authority put Saturday's missile strike on the passenger terminal at **King Khalid International** at **12 dead and 309 injured** — four Saudi citizens and eight residents. By Sunday evening the Indian embassy had reported the death of an injured Indian national, taking it to **13**.[1]

What makes it the kingdom's deadliest attack in more than a decade is the hour, and a silence. The strike, Reuters reported from Riyadh, **"occurred without warning, with no alert issued by Saudi Arabia's threat warning system"**.[2]

Sunday was worse in a quieter register. At about a quarter to five the coalition said it had intercepted a cruise missile near the airport — and that a projectile had struck **Terminal 4** and started a fire. The regulator confirmed further injuries, then disclosed that **Dammam** had been hit the same afternoon.[3]

Brigadier General **Yahya Sare'e** claimed four operations against the two airports and two Aramco facilities, answering what he put at **2,058** Saudi airstrikes since the escalation — one party's figures on that party's own channel. The doctrine he named has no top: **"airport for airport, blockade for blockade, and escalation for escalation"**. Asked whether he had changed his mind about joining the strikes, the American president said: **"We may."**[4]

A capital layered with interceptors did not tell the hall that anything was coming.

RIYADH — **ARAB NEWS**

A crime, and a promise. Coalition spokesman Maj. Gen. **Turki al-Malki** calls it a **"full-fledged war crime"** that **"will be met with a firm response"**.[5]

BEIJING — **GLOBAL TIMES**

A phase change with a price tag. **Liu Zhongmin** reads **"a new phase of high-intensity confrontation"**, the Red Sea, the Gulf and Hormuz capable of going at once.[7]

THE CHATTER · WHAT THE AIRCRAFT ACTUALLY DID

Thursday's strike is the one with numbers attached. Cirium counted **169 of 378** scheduled Riyadh departures cancelled or not flown — **44.71 per cent** — with **Saudia** losing **97 of 140**, and the board recovered by six that evening.[7]

THE SCEPTIC · THE AIRCRAFT WILL NOT FIX IT

Leonardo Jacopo Maria Mazzucco of the Arab Gulf States Institute, on 9 October, dismantles the purchase everyone assumes is the answer. An F-35 buy **"would offer only a partial response to the vulnerabilities exposed"** and **"would not close persistent gaps... particularly against low-cost, distributed threats, such as drones"**.[8] The gap is not in the fighter inventory; it is in the layer that issued nothing on Saturday. Against that, **Nawaf Obaid** of King's College War Studies insists the ratio favours the defender — **"one will always come through"** — and that **"militarily, the Saudi-led coalition has the upper hand"**.[6]

44.71%

RIYADH DEPARTURES CANCELLED OR NOT FLOWN ON 8 OCTOBER — AND NORMAL BY EVENING[7]

The test is quick. If the toll stops at 13 and both airports hold a normal schedule inside 48 hours, this was attrition with a bad weekend. If either moves, two words become policy.[5]

SOURCES 1 Arab News, "Riyadh hospitals treating 309 injured in deadly airport attack", updated 11 October 2026 20:33, read at source: the aviation authority's statement giving 309 injured and 12 killed, and the nationalities. 2 Reuters, Timour Azhari and Enas Alashray, dateline Riyadh, as carried by The Jakarta Post on 11 October 2026 and read there at source; the deadliest-in-a-decade characterisation and the absent alert are the wire's own standfirst, as is the Houthi media office's claim to have shut the airport. 3 Arab News, "Coalition says new attack at Riyadh airport causes fire", updated 11 October 2026 22:53, read at source: the Terminal 4 strike and fire, the intercepted cruise missile, the confirmation of further injuries, the Dammam attack, the Indian embassy's statement, and the drone intercepted late on Sunday. 4 Al Masirah, the Ansar Allah broadcaster, for Sare'e's announcement, the quoted doctrine and the 2,058 figure — one party's numbers on that party's own outlet, not independently verified here. 5 France 24 with Agence France-Presse, 11 October 2026, for Maj. Gen. Turki al-Malki's "full-fledged war crime" posted on X and for the American president's two-word answer; Arab News of the same day for the "firm response" formulation and the foreign minister's calls with Jordan and the European Union. 6 Al Jazeera's own analysis of 11 October 2026 for both quotations from Nawaf Obaid. 7 Cirium schedule data for 8 October 2026 as reported by Khaleej Times, for the 169-of-378 count, the 44.71 per cent and the Saudia figure. 8 Leonardo Jacopo Maria Mazzucco, Arab Gulf States Institute, "Would the F-35 address Saudi Arabia's air defense gaps?", 9 October 2026, read at source.

AROUND THE DESKS · HORMUZ & THE HAGUE

The corridor opens. The strait does not

Tehran agreed coordinates with Muscat on 7 October, then said reopening "depends entirely" on Washington. In The Hague, three days after the sanctions landed, nothing has broken.

Each clause of the Iranian deputy foreign minister's statement takes something away. The corridor agreement with Oman **"does not provide for an immediate opening of the Strait of Hormuz"**; reopening **"depends entirely on the United States fulfilling its obligations"**; Tehran **"is not obligated to reopen the Strait of Hormuz at this stage"**. Coordinates were agreed on **7 October** for lodging with the International Maritime Organization. Washington puts transit at roughly **80 per cent** of February levels — February being when the war began. One source carries those quotations, a Russian state wire, and this desk says so.[1]

In The Hague the delta is an absence. Three days after the United States designated the **International Criminal Court** itself, no bank, insurer or cloud provider has been reported to have dropped it. The European Union's high representative issued a statement on **10 October**; the **Blocking Statute has not been extended** to cover the Court. Two scholars writing in Doha name the gap: **"Declarations will not prevent banks from closing ICC accounts"**.[2]

And off Obock, the worst single loss recorded on the Eastern Route. A boat carrying about **305** people bound for Yemen capsized on the night of **8–9 October**: **139 dead**, **84 missing**, **15 children** among them. The migration agency's director general: **"No one gets into a boat like that unless they believe there is no other safe way."** More than **20,000** movements have crossed into Djibouti from Ethiopia since 10 September.[3]

A condemnation is not an indemnity. Three days on, no bank has been reported to have moved.

MOSCOW — **TASS**

Leverage, held on purpose. The corridor is presented as a technical arrangement for the IMO register, with reopening conditional on American performance — and the quotations that establish the whole position exist, so far as this desk can read, only here.[1]

THE RECORD · A CORRECTION DECLINED

A researcher on this issue reported that yesterday's page three was wrong: that no Russian-diesel general licence exists. The archive outranks memory, so this desk read the documents. Both are real, both signed by the same director on 9 October — **ICC-related General Licence No. 13** at **08:41:14**, and **Russia-related General Licence No. 135** at **14:43:59**, authorising the import of **diesel fuel of Russian Federation origin**. Both expire **7 April 2027**. No. 129 was right, and no correction is printed.[4]

139

DEAD OFF OBOCK, WITH 84 STILL MISSING FROM ABOUT 305 ABOARD[3]

IN BRIEF

Simon, and a bond that may have missed. The hurricane bottomed at **928 millibars** overnight on 11 October and was **947mb** by that afternoon, without landfall. Mexico's **US$175m** parametric bond attaches at **937mb or lower inside a defined box** off Jalisco. Specialist reporting says the pressure likely never fell far enough *while the storm was in the box*. A storm can be too strong in the wrong place.[5]

The licence nobody expected to need. Of the four issued with the designation, **No. 16** authorises transactions for the **"continued detention… health, safety, and welfare"** of nine named detainees — permission for the court Washington had just sanctioned to go on guarding its prisoners.[4]

One thread runs through all of it. The instrument exists, is signed, is coordinated — and does not fire.

SOURCES 1 TASS, 11 October 2026, for deputy foreign minister Kazem Gharibabadi's three quoted statements, for the intention to register the corridor with the International Maritime Organization, and for spokesman Esmail Baghaei's 7 October account of the agreed coordinates; the American estimate of transit at about 80 per cent of February levels is attributed there to the Secretary of State. The quotations could not be matched at a second outlet — several Western wires refused our requests — and are single-sourced to a Russian state agency. The Brent close and the year-on-year move are from a market data service, read separately. 2 OFAC's recent-actions page for 9 October 2026, read at source, for the designation of the International Criminal Court as an entity and the issuance of General Licences 13, 14, 15 and 16; the European External Action Service's press index for the high representative's statement of 10 October, whose full text sat behind a browser challenge and is not quoted here. The quoted sentence is from Sergey Vasiliev and Triestino Mariniello's opinion piece for Al Jazeera, 11 October 2026. The absence of any reported commercial termination is this desk's own finding as of 11 October and is a negative, not a certainty. 3 The International Organization for Migration's own release, "Death toll in Djibouti shipwreck climbs to 139", for the counts, the location, the date and the director general's quotation, and for the movement figure since 10 September; Al Jazeera of 11 October for the fifteen children and the agency's Djibouti chief of mission. 4 The two licence documents themselves, downloaded from the Treasury's Office of Foreign Assets Control and read in full: 31 CFR part 528 General Licence No. 13, signed 2026.10.09 08:41:14 -04:00; and 31 CFR parts 587 and 589 General Licence No. 135, "Authorizing Transactions Related to the Sale, Delivery, Offloading, and Importation of Diesel Fuel of Russian Federation Origin", signed 2026.10.09 14:43:59 -04:00. Both carry the same expiry; Licence No. 16 and its nine detainees were read in the same set. 5 National Hurricane Center public advisories for the eastern Pacific system, numbers 14A and 18, read at source for the 932mb and 947mb readings and the position; the bond's structure, attachment pressure and the observation that the pressure may never have fallen far enough inside the parametric box are from the specialist catastrophe-bond press, and the outcome is not yet determined.

SINGAPORE · THE HAZE LADDER

Eight points short of the first rung

The response ladder published on Friday begins at a PSI of 200. Central Singapore reached 192, and the Grand Prix ran at 154.

Twice this month this desk reported that Singapore's upper haze bands carried no stated consequences. On **9 October** that stopped being true. The environment agency convened the **Haze Task Force — 28 agencies** — and published what happens where. Above a 24-hour PSI of **200**, the Very Unhealthy band: grassroots organisations begin distributing N95 masks, **two to a pack, against an NRIC**, with vending machines at selected transport hubs and Social Service Offices; air-conditioned shelters open at community clubs; and preschools must stop **all outdoor activity and all strenuous indoor activity**.[1]

Above **300** — Hazardous — shelters stay open to **10pm**, and the education ministry **"will consider closing schools when the air quality forecast for the next day is Hazardous"**. The clause that matters in October concerns examinations: candidates sit their papers in enclosed rooms with air purifiers, and if a closure collides with a national examination the affected papers are **rescheduled and the examination period extended**. That is a contingency written for the fortnight we are in.[1]

The defence ministry published its own thresholds the same day and went further in one respect: the armed forces act on a PSI above 200 **or** a 1-hour PM2.5 above **150µg/m³**. Two ladders, two trigger definitions, one city.[2]

None of it switched on. The index stopped at 192, and the machine stayed in its box.

None of it switched on. The agency's own 7pm reading on Friday was a 24-hour PSI of **136–192** and a 1-hour PM2.5 of **69–123µg/m³**, Band 2, Elevated, with the warning that the index **"may exceed 200"**. It did not. At **5am this morning** the live feed read **148** west and central, **138** east, **117** south, **105** north, with central 1-hour PM2.5 at **95**. Unhealthy everywhere; Very Unhealthy nowhere.[3]

THE CHATTER · THE RACE THE SMOKE DID NOT STOP

The minister for sustainability and the environment was reported on Saturday as saying the country **"could expect the haze level to exceed 200, going into the Very Unhealthy range"**. On Sunday the central index stood at **154 at six in the morning** and the Marina Bay night race went ahead over **62 laps**, won in **1:42:00.785**. Reddit was unreachable all session; the only readable discussion was a thin Hacker News item on a report that Kuala Lumpur was the world's most polluted city and Singapore third — **10 points, two comments**.[4]

THE SCEPTIC · THE LADDER POINTS THE WRONG WAY

Alan Tan, professor of environmental law at the National University of Singapore, to Agence France-Presse on 1 October: Singapore will likely **"wait for Indonesia to finger certain companies first before we take action against these same companies under the act"**, and **"Regrettably, not much has changed by way of how the root causes of fires remain unaddressed in Indonesia."** Investigations into four Indonesian companies over the 2015 haze remain open. On the regional treaty, **Mohd Hazmi Mohd Rusli** of USIM: **"It is legally binding, but it does not provide a mechanism for one member state to simply sue another."** A week quantifying masks and shelters is a week not spent on the instrument that has produced nothing in eleven years.[5]

192

THE CENTRAL PSI PEAK ON FRIDAY, AGAINST A LADDER THAT BEGINS AT 200[3]

A threshold never crossed is not useless; it is untested. Singapore now knows exactly what it would do, and still does not know whether any of it works.

SOURCES 1 The National Environment Agency's own release, "Singapore Ready To Step Up Haze Response", 9 October 2026, read in full at source: the 28-agency Haze Task Force, the 7pm readings of 136–192 PSI and 69–123µg/m³ 1-hour PM2.5 in Band 2, the statement that the 24-hour index may exceed 200, the mask-distribution mechanism and the eligibility lists, the two-pieces-per-pack and NRIC conditions, the shelter arrangements, the preschool restrictions and the exemption for children with chronic heart and lung conditions, the extension of shelter hours above 300, the quoted school-closure clause, and the examination provisions including rescheduling and extension of the examination period. 2 The Ministry of Defence's release, "The SAF Adjusts Activities According to Haze Situation", 9 October 2026, read at source for both threshold pairs and the N95 issue. 3 The environment agency's own real-time feeds, queried directly by this desk at 05:00 on 12 October 2026: the 24-hour PSI by region and the 1-hour PM2.5 by region, timestamped 2026-10-12T05:00:00+08:00. The Friday peak is from the agency's release at note 1. 4 The race result is from the championship's own classification page for the Singapore Grand Prix, read at source: 62 laps, the winning time and the gaps to second and third. The minister's remark is reported by a motorsport outlet and is not a transcript; the 6am index on race day is from the same report and was not independently re-read. The Hacker News item and its score were read at source. Reddit refused this desk and its researchers throughout, so no Singapore forum sentiment is reported here rather than invented. 5 Agence France-Presse, Jan Hennop, 1 October 2026, as carried by a science news service and read at source, for both quotations from Professor Alan Tan and for the statement that investigations into four Indonesian companies over the 2015 haze remain open, and for the quotation from Professor Mohd Hazmi Mohd Rusli of Universiti Sains Islam Malaysia.

BUY IT ONCE · ENAMELWARE

The glass will not stick to the rim

Enamel is glass fired onto steel at around 850°C. It cannot cover a rolled edge — which is where, in this climate, an heirloom begins to rust.

Start with the furnace. Enamel is a **frit** — silicates, borates and oxides melted and quenched — fused to steel at **790–860°C**. Riess, in Ybbsitz since 1690, states its recipe plainly: **"Glas auf Eisen"**, iron with quartz, clay, feldspar, soda and potash, at **850°C**. The steel is sorted by carbon — plain at **0.035–0.065 per cent**, decarburised below **0.0040** — because hydrogen dissolved in firing migrates out afterwards and pops crescent flakes off the surface.[1]

Then the coats, where price hides. The **ground coat** carries **cobalt, nickel and manganese oxides**: the steel oxidises in the kiln, the molten glass dissolves that oxide and reacts with the cobalt, and that reaction — not pressure, not glue — holds enamel on. The **cover coat** is a titanium-white layer that seals it. Riess states **two to four layers**. A single-coat piece has no cobalt bond beneath and no seal above; it chips to bare steel on the first knock.[2]

Which brings us to the rim, and a sentence most makers would rather not write. Riess concedes that enamel **does not adhere well to narrow edges such as the underside of a rolled rim**, so bare iron shows there — and in this climate that is where an heirloom starts to rust. Falcon's founder gives the alternative: after the first firing **"the rim is hand-painted… and the whole thing is re-fired"**. A bare rim, a stainless band, or a second trip through the kiln. That choice is most of the price.[2]

The trade cannot agree on what to do when the glass fails. Falcon: **"If your enamelware has a chip you can still use it as normal."** Riess: an outside chip is cosmetic, but for one on the inside face it recommends replacing the piece. Two houses, opposite advice, both in writing.[2]

THE CHATTER · TWO SHELF BRANDS, ONE LINE

Falcon tells two continents different things: its American answers page says **"manufactured in China"**, the European one **"manufactured in East Asia"** — same question, softer noun, atop a 1920s Black Country heritage. **Golden Rabbit's** about page redirects to **Crow Canyon Home**, and every Crow Canyon mug is filed under the vendor **"Golden Rabbit Enamelware"**. A Hong Kong house making enamel since 1951 says it plainly: **"Many of our customers carry our products under their own private label."**[3]

THE SCEPTIC · FOUR DOLLARS AGAINST FORTY

Against the published measurables a compliant cheap mug and an expensive one are the same object: **Mohs 5–7**, impact to **ISO 4532**, lead and cadmium to **ISO 4531**. The money buys coat count, rim treatment and a named country. And the budget tier cannot write its own copy: one camp brand advertises enamel **"twice kiln-hardened at 1000°F" — 538°C**, well below the point at which enamel fuses to steel.[4]

0.5 µg/mL

LEAD RELEASE ALLOWED FROM A MUG — SIX TIMES STRICTER THAN FOR A PLATE[4]

NEXT ON THE BENCH

Espresso machines — the boiler and the group, not the badge.

Dehumidifiers — the appliance that decides a wardrobe's life here.

Sewing machines — all-metal gear trains, and who still services them.

Pay for the rim treatment and the coat count; dry the rim before it goes away; and ignore the heritage on the box, because the heritage is almost never where the kiln is.

SOURCES 1 The International Enamellers Institute's own technical pages for the definition of frit, the 790–860°C firing range, the carbon bands and the fishscaling mechanism; Riess's own site for its 850°C figure, raw materials and founding date. 2 Riess's own frequently-asked-questions page, read at source, for the two-to-four-layer statement, the admission that enamel does not adhere to the underside of a rolled rim, and the advice to replace a piece chipped on its inside surface; Falcon Enamelware's own answers page for the contrary advice; a published interview with Falcon's founder for the quoted account of the hand-painted, re-fired rim. The cobalt and nickel adhesion mechanism and the titanium cover coat are from the institute pages above. 3 Falcon's American and European answers pages for the two differing manufacturing statements; Crow Canyon Home's own site, where the Golden Rabbit about page redirects and the quoted vendor field appears on every mug; and a Hong Kong industrial-history record of an enamelware factory founded 1951 for the private-label quotation. 4 The institute's cookware page for Mohs 5–7 to EN 101 and the references to ISO 4532 and ISO 4531; the United States Food and Drug Administration's compliance policy guide on pottery and ceramics, read at source, for the 0.5 microgram per millilitre limit for cups and mugs against 3.0 for flatware; and an American camp-equipment maker's own technical page for the quoted 1000°F claim, whose conversion is ours. 5 Prices read on the makers' own shops on 11 October 2026, and Falcon's own stockist list and Riess's own retailer list for the absence of any Singapore outlet. Reddit refused this desk all session, so no r/BuyItForLife sentiment is reported rather than invented.

TECHNOLOGY · THE GENESIS LEDGER

Six billion announced, three billion itemised

Eleven companies pledged $2.4bn in compute credits on Wednesday. Add every other component the day listed and you reach about half the headline.

The pledges are real and in writing. At a Washington summit on **8 October**, **NVIDIA** committed **$1 billion over five years** to American science across three buckets, with no hardware named. The Energy Department announced **12 new Phase II awards worth $159m** alongside six Phase I, taking the first-year cohort to **297 projects**.[1]

The full ledger runs to **$2.4bn across 11 partners**: NVIDIA $1bn; **AMD $500m**; OpenAI $200m; **Anthropic $150m**; Google $150m; AMP and Emerald AI $100m each; **$50m apiece from AWS, Armada, Crusoe and Micron**. One published its own terms: Anthropic confirms **$150m over three years** in credits across **more than 15 agencies**. Micron put in the same fifty million as a hyperscaler — a twentieth of NVIDIA's.[1]

Now the subtraction. The day was billed at **$6bn**. Add the components a trade outlet itemised — **$2.4bn** in credits, **$159m** in Phase II awards, **$96m** for early-career researchers, up to **$350m** in Pentagon quantum money — and the total is roughly **$3bn**. Half the headline is unexplained.[2]

Compute credits are not money. They are a discount on a service the issuer prices itself.

Elsewhere a number landed that a calculator can check. **TSMC** published September revenue of **NT$511,857m** on 8 October, up **54.6 per cent**. Add July and August and the quarter is **NT$1,494,243m**; at the **32.0** rate TSMC's own guidance page assumes, that is **US$46.70bn — 1.95 per cent above the top** of its own range. It reports on **15 October**, having already beaten its ceiling.[3]

THE CHATTER · THE SILENCE IS THE FINDING

A $2.4bn federal compute pledge did not register. On Hacker News the main report reached **3 points and zero comments**; Anthropic's own announcement, **2 points and one comment**. Over the same three days the front page carried Yandex bot traffic at **575 points** and a computer built from lightbulbs at **570**. The public that would normally take this apart did not look up.[4]

THE SCEPTIC · ONLY ONE PARTY CONFIRMED TERMS

Of eleven pledging companies, the reporting records one spokesperson confirming terms, and one company publishing its own. Everything else is a figure in a list. A credit has no maturity, no drawdown schedule and no published unit price; it can be throttled by queue, expired by policy or revalued by a rate card, and none of that is reported as a withdrawal. Ask in April how many of the **297** projects have drawn compute.[1]

1.95%

TSMC'S JUL–SEP REVENUE ABOVE ITS OWN GUIDANCE CEILING, AT ITS OWN ASSUMED RATE[3]

IN BRIEF

Qualcomm v Arm — the jury started early and went home empty. The Delaware jury deliberated **about four hours on Friday without a verdict** and resumes **Tuesday 13 October**. So deliberations began four days before the date this desk printed; the 13th is a resumption, not a start. Still unresolved: a separate bench trial, and no ruling on whether the five-year royalty-suspension clause is enforceable. The **$1.5bn** windfall figure in circulation does not appear in the wire copy.[5]

Two announcements, one shape: a pledge counted before it is drawn, a quarter won before it is reported. Still unreleased, meanwhile: **Mistral Large 4**'s weights, promised "by the end of the month".[6]

SOURCES 1 NVIDIA's own newsroom release of 8 October 2026 for the $1bn commitment and its three buckets; the Department of Energy's release of the same day for the twelve Phase II awards, the $159m total, the six Phase I awards and the 297-project cohort; a federal-technology trade outlet's report of 8 October for the eleven-partner consortium list and the per-company amounts, and for the detail that one spokesperson confirmed terms; and Anthropic's own announcement for its $150m over three years, the "several hundred" projects and the more-than-fifteen-agencies figure. 2 The same trade report for the $6bn billing and for the $96m early-career and up-to-$350m Pentagon quantum components; the subtraction to roughly $3bn is this desk's own. 3 TSMC's own monthly revenue page and press list, read at source, for the September figure, its year-on-year change and the 8 October release date, and for July and August; TSMC's own second-quarter page for the third-quarter guidance of US$44.6–45.8bn and the stated 32.0 exchange-rate assumption, which corrects the 32.63 figure this magazine printed in Nos. 127 and 128. The quarterly sum, the dollar conversion and the 1.95 per cent are ours. The individual September release page returned a 404 at the time of reading; the figures are from the monthly table. 4 Hacker News, searched at source over 9–11 October 2026 for the item scores and comment counts cited, and for the three front-page comparisons. 5 Reuters, Tom Hals and Stephen Nellis, 9 October 2026, as carried by a radio-station wire page and read at source, for the four hours of deliberation, the Tuesday resumption, the unresolved bench trial and the absence of a ruling on the royalty-suspension clause; the damages figure circulating elsewhere does not appear in that copy. The Delaware docket itself refused this desk. 6 Mistral's own announcement of 6 October 2026 for the parameter counts, the training hardware and the quoted release promise; the maker's public model repository, read at source on 11 October, carries no Large 4 weights.

THE HOME LAB · THE EXPERT CACHE

The speed-up that costs a quarter of your prompt

Merged on 8 October, the multi-GPU expert cache raises decode to 1.85 times. In the same run prompt processing falls to 0.73. The notes do not say so.

The sequel landed four days after we printed the original. **PR 30112**, "support MoE cache over multiple GPUs", merged into llama.cpp on **8 October**, taking the expert cache we covered in No. 129 across more than one card. On the author's own bench — two RTX 4090s, a 93.7GiB quantisation with 65GiB of experts — decode goes **34.78 to 60.38 tokens a second**, and **64.39** with the aggressive setting.[1]

Read the other column of the same table. Prompt processing goes **440.7 to 402.0**, then **321.5 — 0.91×**, then **0.73×**. The prose reports the decode multiple and does not mention that its best decode run processes prompts **27 per cent slower**. For a long-context workload that trade is the wrong way round.[1]

The dissent we printed has not been answered either. **vbooka1**, on the new thread: **"PP about 10% slower and TG became 200% slower, 33 -> 16 t/s"** — no author reply on the page. **homeofe**, on a 2080 Ti over PCIe 3.0, found the cache slower at every size that fits, **3.6 against 15.1 tokens a second**. **insunaa**, on a 5090: **"some of the decode speed seems to be bought with hallucinations"**. On a single 5090 this is a real **2.2×**; on two cards or an older bus it is a coin-flip.[2]

The week's other merge is more cheerful and makes the same point about vendor numbers. A **27.36-billion-parameter** ternary model now runs on stock llama.cpp: **1.75 bits per weight**, **5.95GB** on disk against a 53.8GB half-precision baseline, **262K context**, Apache-licensed, and **129.9 tokens a second** on a 5090 by the maker's measure. The maker also claims a fourteen-benchmark average at **98.2 per cent of full precision**.[3]

THE CHATTER · THE MODEL NOBODY CAN ACTUALLY RUN

The window's biggest thread — **1,119 points, 983 comments** — asked why the industry was not reacting to a new frontier open model. The answer was procedural. **zozbot234**: **"It's still lacking llama.cpp support, and the work on that isn't moving very fast either."** This desk checked: the support request has been **open since 10 September**, context capped at **16K**. Staged quantisations run **473.1GiB** down to **246.3GiB** — and the smallest managed **2.8 tokens a second** on a 128GB box. It is still too large for a 256GB machine once a cache is added.[4]

THE SCEPTIC · TWO CLAIMS, NO MEASUREMENT

The ternary model's 98.2 per cent does not survive an independent read. **mjechow**, on RDNA4, measured perplexity at **10.29 against 7.43** for a four-bit quantisation of the same base, and on 22 tasks counted **four failures against one**: **"I am getting more loops, cap hits and fabricated facts."** Separately, Microsoft's 7 October post announcing experimental GGUF support claims it and NVIDIA delivered **"significant performance improvements"** — and prints not one number.[5]

0.73×

PROMPT THROUGHPUT IN THE SAME RUN THAT REPORTS 1.85× DECODE[1]

ALSO MERGED

Honest sizing, for once. A Vulkan change merged 11 October moves duplicate handling into the prepass; the author's own claim is **"maybe 1% faster than baseline"**, and a second maintainer confirms it independently across AMD, Intel and NVIDIA. Ten release tags were cut on 10–11 October alone. Apple's MLX published nothing in the window; the newest release is 29 September.[6]

Three releases, one habit: the number written down is the one that went up.

SOURCES 1 The llama.cpp pull request 30112, "llama: support MoE cache over multiple GPUs", merged 8 October 2026, read at source: the review change from per-device to a global budget split by tensor-split, the author's hardware and quantisation, and the full benchmark table from which both the decode and the prompt figures and their multiples are taken. The observation that the prose omits the prompt regression is this desk's. 2 The same page and pull request 29887, both read at source, for the quoted measurements and usernames, the hit rate, the PCIe 3.0 comparison and the single-card 2.2×. The absence of an author response is as the page stood on 11 October. 3 Pull request 29600, merged 11 October 2026, and the maker's own model card, read at source, for the parameter count, the bits per weight, the file sizes, the context length, the licence, the 5090 throughput and the fourteen-benchmark average. 4 The Hacker News thread of 7 October 2026 and its score and comment count, read at source, with the quoted comment and username; and the still-open llama.cpp support pull request 28696, read at source, for the opening date, the 16K context cap, the unmapped components and the staged quantisation sizes and measured throughput. 5 The in-thread independent measurement on pull request 29600, read at source, for the perplexity comparison, the task-failure count and the quoted remark; and Microsoft's developer-blog post of 7 October 2026 for the quoted performance claim and the absence of any measurement in it. 6 Pull request 29998, merged 11 October 2026, read at source for both the author's claim and the independent confirmation; the project's own releases page for the tag count; and the MLX project's releases page for the 29 September date. Reddit refused this desk and its researchers all session, so the community read here is Hacker News and the project threads themselves.

GEAR WE'RE WATCHING · TECHNICS EAH-A1000

The specification sheet is behind an email form

Technics will tell you the driver is filled with magnetic fluid. It will not tell you what the headphone weighs — and there is still no price for it here.

The peg is not an announcement; it is the shelf. Panasonic's hi-fi line put the **EAH-A1000** on sale in the United States from **1 October** at **US$449.99**, and the first full independent reviews landed that week. The maker's page gives a **40mm Magnetic Fluid Driver**, **"up to 45 hours playback with ANC on"**, LDAC plus LE Audio and USB audio, and two colours, one called Platinum Greige.[1]

What the page does not give is more interesting: **no weight**, **no Bluetooth version**, **no cancellation figure in decibels**, and the full specification sheet **gated behind an email form**. The maker's numbers also disagree by door — the American store says 45 hours with cancellation on, the press material seeded to reviewers **40**. Both are Technics.[1]

No independent measurement of this headphone exists. The single independent figure in circulation is one a reviewer produced on his own scale: **306 grams**. Every acoustic claim attached to the A1000 is the manufacturer's.[2]

And there is no Singapore price. The maker's Singapore domain and its parent's both refused this desk, and the main local specialist dealer lists no A1000. A flagship on sale a fortnight with no published price in a market Panasonic serves directly is itself the fact.[3]

THE CHATTER · THE COLOUR, NOT THE DRIVER

The owner-side conversation that exists is thin and mostly about a word. Under the fullest review, **ORT** on 29 September: **"'GREIGE'?! How quaint of Technics to create such a flaccid portmanteau… Would this… prevent me from buying these headphones? Nah." Ian White**: **"Could be worse. Could be Pepto Pink."** The headphone forums and the measurement site were unreachable this session, so there is no owner consensus to report, and none is invented.[2]

THE SCEPTIC · THE HINGE AND THE STANDBY

Chris Boylan, who wrote the one full review with numbers in it, still recommends the headphone — then lists what a buy-once reader should weigh. Cancellation is **"slightly outmatched by some competitors"**. Volume headroom runs short on some Atmos material. It **stays powered on by default and slowly drains**. And the mechanical one: **"the folding mechanism feels slightly loose"**. On a four-hundred-dollar object meant to last a decade, the hinge decides whether it does.[2]

306g

THE ONLY PUBLISHED WEIGHT — TAKEN BY A REVIEWER ON HIS OWN SCALE[2]

STILL ON THE SHELF

Ray-Ban Meta Audio is **not sold here**: Meta's own Singapore page returns "Unavailable — Not sold in your region", while the American page lists **US$349** with a notify button and no specifications at all. A regional hold-back, not an absent market: the camera versions launched here in April.[4]

No movement on the fortnight's two gaps: the **Garmin Enduro 4** still has no Singapore price, the **Sigma 50-120mm F2.8** no independent bench.[3]

Buy-once readers should wait for a bench, not a review. Until somebody measures the cancellation, the only verified number attached to this headphone is its mass.

SOURCES 1 The maker's own United States store page for the EAH-A1000, read at source: the price, the stock-keeping code, the driver description, the quoted battery claim, the codec and connection list, the two colour names, and the absence of weight, Bluetooth version and any decibel figure, together with the email gate on the full specification sheet. The differing battery figures are from the press material seeded to reviewers and quoted in the review at note 2. 2 A full independent review published 29 September 2026, read at source, for the measured weight, the named competitor comparison on cancellation, the Atmos headroom observation, the standby drain, the quoted remark about the folding mechanism and the overall recommendation; and that article's own comment thread, read at source, for the three quoted comments and their usernames and dates. A second publication's review was readable only as far as its standfirst, and a third had published an initial impression rather than a full review at the time of reading. The two independent measurement laboratories have published nothing on this model. 3 The maker's Singapore domain and its parent's Singapore page both returned HTTP 403 to this desk, and a search of the principal local specialist dealer's catalogue returned no result for the model; no Singapore-dollar price could be found at any source. The Enduro 4 and Sigma lines restate the gaps recorded in Nos. 128 and 129 and are unchanged. 4 Meta's own Singapore and United States product pages for the audio-only frames, both read at source, for the unavailability notice, the dollar price, the notify button and the absence of published specifications; the differing figure is from a consumer roundup and is not the maker's.

THE GOOD LIFE · BOND & SPECTRE

Seventy-seven pieces, built twice

Each box set is limited to 77 numbered examples. There are two box sets — so 154 platinum-gold watches exist under a 77-piece banner.

The steel watches are honest objects, priced without drama. The **James Bond Edition**, ref. **210.30.42.20.03.004**, is **US$8,600**: 42mm steel, **14.23mm** thick, **157g**, a grade-5 titanium bezel ring, the **Calibre 8806** at 55 hours and 15,000 gauss, 300 metres. The **Spectre Edition** is the same money with a black ceramic bezel ring — and is marked **"Currently unavailable online"** on its own page, launched but not orderable.[1]

Above them sit the box sets, and the counting goes strange. Each is **"limited to 77 numbered examples"**, and there are two — so the platinum-gold watch exists in **154** pieces under a number that reads as 77. Each set also contains one of the **US$8,600** steel watches, neither sold separately. Subtract, and the platinum-gold piece is effectively **US$81,200** in a **US$89,800** parcel.[2]

Even the price is not one number. The maker prints **US$8,600**; one specialist prints **US$8,700** excluding tax; the set is **US$89,800** at two outlets and **US$89,900** at a Singapore title. No Singapore-dollar price reached this desk — the maker's Singapore path served the American store — which may be our geography rather than the maker's, and is held as such.[2]

Each set is limited to seventy-seven. There are two sets. The limit describes half of what is being made.

THE CHATTER · THE COMMENT THREADS DID THE MATHS FIRST

Under the two biggest hands-on pieces, both of 5 October, the readers got there before the reviewers. **Gavin Stevens**: **"What a wasted opportunity, they have just added a different caseback to models already on sale." Fruit-Loupe**: **"90K plus VAT for two Seamasters in a box, I can buy an extremely nice used Aston Martin for that."** The dissent from the dissent, **Ed**: **"Lot of folks bashing these watches. I own two Seamaster's. I am happy with them both."**[3]

THE SCEPTIC · A LAST GASP, AND A BORROWED FACE

Two criticisms worth the ink. **dvh**: **"If anything, it's disappointing. It borrows heavily from the Yachtmaster"**, echoed by another reader next morning. And **CoachRockne**, the one that should decide a purchase: **"This has to be the last gasp for this gen SMP"**. A price rise immediately before a generational replacement is the classic trap for a buy-once buyer.[3]

US$81,200

THE PLATINUM-GOLD WATCH, ONCE YOU SUBTRACT THE STEEL ONE IN THE SAME BOX[2]

AT AUCTION, AND ON THE BENCH

Hong Kong, online. Christie's *Watches Online: The Hong Kong Edit*, 24 September to 7 October, totalled **HKD 23,125,632** on the house's own results page. The steep multiples went to things that cannot be reordered: an **Eska** pink-gold cloisonné-enamel piece with no box or papers made a reported **HKD 1,216,000** against an 80,000–160,000 estimate. That figure is single-sourced to a trade summary; the house's own lot pages were unreadable here.[4]

Held, not printed. A Japanese maker's two new Class A integrated amplifiers are quoted at **US$21,000 each** — flagship and junior sibling at the identical figure — in one roundup. The maker's own pages publish no price, no specification, no date.[4]

The watch is good; the limit is arithmetic. Buy the steel one, or wait for the generation the threads can already see coming.

SOURCES 1 The maker's own product pages for both references, read at source on 11 October 2026: the prices, the reference numbers, the case dimensions and weights, the bezel materials, the dial description, the calibre, its power reserve and magnetic rating, the water resistance, the bracelet, and the "currently unavailable online" status on the second reference. 2 The same pages for the steel price; two watch specialists' hands-on reports of 5 October 2026 and a Singapore title's report of 12 October for the box-set figures and the differing dollar amounts, and for the quoted limitation of 77 numbered examples per set. The observation that two sets means 154 platinum-gold watches, and the subtraction that yields the effective price of the platinum-gold piece, are this desk's own. The Singapore path on the maker's own site served United States pricing to us; whether a Singapore-dollar price exists is held, not asserted. 3 Reader comments under the two hands-on reports of 5 October 2026, read at source and quoted verbatim with the usernames and dates as published on those pages. 4 Christie's own results page for the Hong Kong online sale, read at a regional mirror after the main site refused us, for the sale total and dates; the two individual lot results are from a trade summary of October auction results and are single-sourced, the house's own lot pages being rendered client-side and unreadable here. The amplifier prices are from one consumer roundup; the maker's own product pages, read at source, carry no price, no specification table and no date, so the figures are held rather than printed as the maker's.

CULTURE DESK · THE WEEKEND NUMBERS

Eleven critics against two hundred and four

An 18 per cent horror film took $27.5m and beat an Aaron Sorkin picture five to one. The aggregate verdict on the winner rests on eleven reviews.

Weekend forty-one went to the film the critics liked least. **Other Mommy** opened to an estimated **$27,500,000** from **3,365** theatres; **The Social Reckoning**, Aaron Sorkin directing Mikey Madison as Frances Haugen, took **$5,500,000** from **2,606**. Second place went to a second-week title at $11.48m and third to a four-week-old horror sequel at $8.35m. The weekend as a whole came to **$84,419,072**, down from **$98,356,419**.[1]

The scores invert the result, and the review counts invert the scores. **Other Mommy**: **18 per cent on 57 critics**, average 4.5 out of 10, audience score **47**, **CinemaScore D+** — and a **Metacritic 39 built on eleven reviews**, of which three are positive. **The Social Reckoning**: **55 per cent on 204 critics**, audience **74**, Metacritic **52 on 55**, with a user score of 7.2 resting on **eighteen ratings**. The film that won the weekend was assessed by a twentieth as many critics as the film it beat.[2]

It is also already profitable. Domestic **$27.5m** plus **$17,690,000** international gives **$45,190,000** worldwide against a reported **$19m** budget, in three days, with most territories opening before the United States — Belgium on the 7th, four more on the 8th, America on the 9th. Japan does not get it until **20 November**.[1]

Singapore has a different chart again. The larger local cinema chain lists **Other Mommy** as a new release, rated **M18**, running **1h 33m**, and ranks it **second** — behind the four-week-old horror sequel that sits third in America. The Sorkin film is not in the local top three at all. No Singapore opening date is published, so none is printed here.[3]

THE CHATTER · THE TRAILER WAS THE FILM

The audience verdict is forming where films are logged rather than reviewed, and it is less kind than 47 per cent suggests. The film's page carries **34 fans** and, at the time of reading, **no displayed average rating at all**. The substantive entries skew negative on one complaint: the marketing spent the surprise. One argues the first trailer's unannounced drop was scarier than the picture; another calls it a set of music videos. The review listings refused us, so these are summarised, not quoted.[4]

THE SCEPTIC · THE CONTRARIAN IS THE NEWSPAPER

The interesting dissent runs *for* the film. The American paper of record calls **Other Mommy** a **"sometimes-creepy, sometimes-hilarious schlockfest"** — a description of a film doing deliberately what an 18 per cent score punishes. Against it, the floor: a trade weekly scored it **10 out of 100**, among the year's **"most limp and ineffectual works"**. When an aggregate rests on eleven opinions, one critic at ten moves it several points alone. That is not a measurement; it is a small sample with a decimal point.[2]

11 v 204

CRITIC COUNTS BEHIND THE TWO AGGREGATE SCORES DECIDED THIS WEEKEND[2]

HELD, NOT PRINTED

The album chart. The Taylor Swift record is reported to hold the American album chart dated 10 October at a figure that would make it a fourteenth week at number one and a hundredth career week there. **This desk could not read the chart.** Both the chart publisher's page and the trade report returned paywall codes, so the units figure, the week count and a reported one-to-157 collapse by another album are all unverified and none is printed as fact.[5]

Two aggregates, two very different amounts of evidence, one shared decimal place. Read the denominator before the number.

SOURCES 1 The box-office tracking service's own domestic weekend page for weekend 41 of 2026 (9–11 October), read at source, for every gross, theatre count and week number quoted, and its by-year page for the two weekend totals; the same service's title page for the international figure and the worldwide sum; a second independent tracking service confirms the opening gross and theatre count. The budget figure is reported by a trade outlet and a second trade source gives $18m; the discrepancy is noted. The release-date sequence is from the second tracking service and a film database. 2 The two review-aggregation sites' own pages for both films, read at source on 11–12 October 2026, for every percentage, average, audience score, Metacritic value and, in each case, the number of critics behind it, together with the positive-mixed-negative split on the horror title; the exit-poll grade is as reported on the first site. The observation about sample size is this desk's. 3 The Singapore cinema chain's own site, read at source on 12 October, for the rating, running time and its published ranking; the other major chain refused us. No local opening date is published there. 4 The film-logging site's own title page, read at source; its review listing pages returned HTTP 403, so no reviewer handle or permalink could be captured and the sentiment above is summarised in this desk's words rather than quoted. 5 The chart publisher's page and the trade report both returned payment-required responses to this desk, so the chart position, the units figure, the week counts and the reported collapse of another album are recorded here as unverified.

RAISING THEM WELL · THE FIRST PHONE

The phone is the hinge, and the hinge is small

Among 3,486 children who had none at eleven, getting one by thirteen moved internalising problems by 0.11 of a standard deviation. Externalising did not move at all.

A study published on **30 September** does what the screen-time literature mostly does not: it takes the **acquisition** of a phone as the event. From a large American cohort it selects **3,486** children who were phone-free at the second-year visit, **mean age 11.82** — the band containing a twelve-year-old — and follows them to **13**. Of those, **1,546** acquired a phone and **1,940** did not.[1]

The results are real and modest. Internalising problems rise by a standardised **0.11** (95% CI 0.05–0.18); attention **0.09**; total problems **0.10**. **Externalising does not move: 0.04, interval crossing zero.** Academic performance, absenteeism, physical activity and in-person bullying are all null. What moves hard is the surrounding behaviour: screen time **0.19**, bedtime social-media use **0.42**, school-day sleep **−0.18**, and the odds of a social-media account more than treble — **OR 3.22**.[1]

Then the part a parent can act on. Roughly **80 per cent** of the association runs through three things: screen time (**27.9 per cent**), a social-media account (**35.5**) and media at bedtime (**16.9**). The caveats are proper: one year of follow-up, mediators and outcomes at the same visit, no causal claim.[1]

A second paper, online **10 October**, pushes the same way from the other side. In **1,510** children aged exactly **11 to 13**, smartphone *ownership* went with **lower** odds of insufficient sleep — while using it **in bed** carried adjusted odds of trouble sleeping up to **13**. Its conclusion names the variable: how the phone is used, not whether one is owned.[2]

The phone is the hinge. The door is the account and the hour it is open.

THE CHATTER · THE BUS, NOT THE RESEARCH

Parents do not argue about effect sizes; they argue about logistics. On Hacker News last month, **hefner1456** said his son wants a phone **"because all the 5th graders have them on the bus"**. In an August thread two others described giving in at **eight**; **graemep** named the forcing functions — the bus-ticket app and coordinating with friends. Nobody in any readable thread cited a study. Reddit and two parenting forums refused this desk all session.[3]

THE SCEPTIC · A TENTH OF A STANDARD DEVIATION

Candice Odgers and colleagues, writing in May, put the brake on before this paper existed. They quote the American National Academies on **"small effects and weak associations"**, cite a 2025 meta-analysis of restriction studies at a mean effect of **0.17**, note a 2024 review finding average effects **indistinguishable from zero** — and report that of **40** restriction studies, **eight were null and eight showed decreases in wellbeing**. A beta of 0.11 is a real signal and a small one. Print it that way; act on the mediators, not the panic.[4]

OR 3.22

THE JUMP IN ODDS OF OWNING A SOCIAL-MEDIA ACCOUNT, ONCE THE PHONE ARRIVES[1]

THE LOCAL RULE

In force since January. Singapore's education ministry ruled on 30 November 2025 that secondary students may not use phones or smartwatches **during school hours, including recess and co-curricular activities**. Its reason — that screen use displaces **sleep, physical activity and social interaction** — is the pathway the study above puts numbers on.[5]

General guidance, not medical advice. The account and the bedtime are the levers with numbers behind them; the handset is how they arrive.

SOURCES 1 Bren and colleagues, "Smartphone acquisition as an inflection point in early adolescence digital exposome: relevance to mental health", *Frontiers in Child and Adolescent Psychiatry*, published 30 September 2026, read at source: the design as a longitudinal secondary analysis of a large American cohort release with mixed-effects regression and a multiple-mediator path model; the analytic sample of 3,486 and its mean age; the split between acquirers and non-acquirers; every standardised coefficient and confidence interval quoted, including the null externalising estimate and the null results for academic performance, absenteeism, physical activity and in-person bullying; the odds ratio for social-media account ownership; and the three mediation proportions and the authors' own limitations. 2 Bakour and colleagues, "Smartphones, digital media use, and sleep in young adolescents", *Sleep Health*, online 10 October 2026: the sample of 1,510 aged 11 to 13, the direction of the ownership association, the adjusted odds for in-bed use, the cyberbullying finding and the quoted conclusion are from the abstract read at source; the article is paywalled and the full confidence intervals and limitations could not be read. 3 Hacker News comment threads of 14 September and 25 August 2026, read at source, quoted with the usernames as published. Reddit, and two parenting forums, refused this desk and its researchers throughout; no sentiment from them is reported or invented. 4 Lind, Schueller and Odgers, "We don't know how social media bans will affect youth but we're doing it anyway!", *Frontiers in Developmental Psychology*, 29 May 2026, read at source, for the quoted National Academies characterisation, the 2025 restriction meta-analysis effect, the 2024 review's conclusion and the count of null and negative restriction studies. 5 The Ministry of Education's own press release of 30 November 2025, read at source, for the restriction, its scope and its stated rationale. Reports of a 2027 bill appear only in secondary blogs and the relevant government page returned a 404; nothing is printed on that.

THE TWEEN YEARS · TESTOSTERONE AT TWELVE

At ten the hormone predicted calm. At twelve, nothing

In a Tokyo cohort, higher testosterone at ten tracked fewer behaviour problems, not more. Two years on it no longer predicted anything at all.

Hormones are the explanation every parent of a twelve-year-old is handed, and a paper published online on **5 October** takes it away. A Tokyo cohort measured salivary testosterone, amygdala volume by scan and parent-rated behaviour in **106 boys at ten to eleven** and **98 at twelve to thirteen**, adjusting for age, pubertal stage and intracranial volume.[1]

At the first wave the association runs opposite to folklore: higher testosterone goes with **fewer** internalising problems, exp(β) **0.64**, and fewer externalising, exp(β) **0.66**, both at adjusted **P < 0.001**. Two years later, **testosterone is no longer significantly related to behaviour at all**. What predicts instead is structure: a smaller right amygdala with higher internalising problems, exp(β) **0.71**, adjusted P < 0.001.[1]

Hold it loosely: **98 boys** at the second wave, one country, no girls, and a paywall that hid the full limitations. Treat the amygdala result as single-sourced. What survives at that sample size is the negative finding, and the negative finding is the useful one.[1]

A much larger study three days earlier points the same way. Using **4,194** children aged nine to seventeen and a decision-model analysis of working memory, an Oslo group finds the gains are driven by **processing efficiency**, with only **"weak links to individual differences in pubertal timing and tempo"**. Whether a child is early or late tells you very little about how her thinking works.[2]

At twelve the hormone has stopped doing the explanatory work it was doing at ten.

THE CHATTER · THIN, AND WE SAY SO

There is no good community read on this subject this week and this desk will not manufacture one. Reddit, a British parenting forum and the main Singapore parenting board all refused us. What is readable is two Hacker News threads carrying the same argument: **elzbardico** in September, that once puberty starts parents should give children space with friends; **spiderfarmer** in August — treat them as adults and let them stumble. Two comments are not a consensus.[3]

THE SCEPTIC · EARLY OR LATE IS NOT DESTINY EITHER

The anti-alarmist case has the bigger numbers. The Oslo analysis is **4,194** children against **98** boys, and finds puberty explaining very little of cognitive development once background is controlled. A third paper, **7 October**, reports fast *physical* development predicting faster midlife ageing while fast **social** development predicts the opposite — but it is closed access and no effect size could be obtained, so its direction is printed without a magnitude.[4]

98

BOYS IN THE SECOND WAVE — THE WHOLE BASIS OF THE AMYGDALA FINDING[1]

WHAT TO DO WITH THIS

Stop attributing. A null result removes an excuse — from the child and from the parent. At twelve the behaviour in front of you is more likely about sleep, school load, friendship and the account on the phone (p12) than about a hormone that has stopped predicting anything.

And watch the band. Both papers are specifically about eleven-to-thirteens. Guidance written for eight-year-olds or sixteen-year-olds is a different literature with different signs; the age band is not a detail.

General guidance, not medical advice. If anything here changes a decision about a child, discuss it with a doctor rather than a magazine.

SOURCES 1 Kiyota and colleagues, "Developmental shift from hormonal to neural correlates of psychobehavioral problems among early adolescent boys", *Psychoneuroendocrinology* volume 194, online 5 October 2026: the cohort and its neuroscience extension, the salivary-testosterone, imaging and parent-rated behaviour measures, the Poisson models adjusted for age, pubertal stage and intracranial volume, the two wave sizes and the 69 with complete biology, and every exponentiated coefficient and adjusted p-value quoted, including the loss of significance at the second wave and the prospective result. The paper is paywalled and is not held in the open repository, so its full limitations section could not be read; the amygdala finding is treated here as single-sourced. 2 Frogner, Pedersen and Westlye, "Adolescent development of decision making processes in n-back task performance and links to pubertal timing and tempo", *Acta Psychologica* volume 270, online 2 October 2026: the sample of 4,194 aged nine to seventeen, the longitudinal cohort, the drift-diffusion modelling, the quoted conclusion and the sex-specific result after sociodemographic adjustment. 3 Two Hacker News comment threads of 29 September and 13 August 2026, read at source and quoted with the usernames as published. Reddit, a British parenting forum and the principal Singapore parenting board all refused this desk and its researchers; no sentiment from them is reported. 4 Xie, Caspi, Moffitt and colleagues, "Do Children Who Develop Faster Go on to Age Faster in Midlife?", *Psychological Science*, 7 October 2026: the direction of both findings is as summarised in the abstract record; the article is closed access, the publisher refused this desk, and no effect size could be obtained, so none is printed.

THE BIG PICTURE · HIGHER FOR LONGER

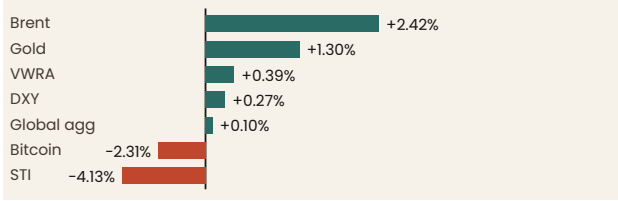
The Fed hinted again. Singapore's banks took the blow

Hawkish minutes on Wednesday pinned the long end above 5 per cent. The damage landed six thousand miles away, in three stocks that are half the index.

One document set the week. The **7 October** release of the September minutes told a market that had just seen a weak payroll print the opposite of what it wanted. Participants saw **"not sufficient progress on lowering inflation in recent months"**, against staff estimates of **3.8 per cent** headline and **3.4** core in August, and **most judged a further increase likely appropriate by year end**.[1]

The week's closes say what that did. Treasury par yields ended at **4.80 per cent** at two years and **5.24** at ten, both a few basis points *lower* on the week, leaving the two-to-ten spread at **+44bp**. Global equities rose: **VWRA +0.39 per cent**; the global aggregate bond proxy, **+0.10**. The dollar index firmed **0.27** to **102.21**, the yen through **158**, the Singapore dollar at **1.2807**. Brent added **2.42 per cent**, gold **1.30**, and bitcoin fell **2.31** to **$82,546**, recovering about **1.1 per cent** over the weekend.[2]

So the week was not risk-off; it was rate-sensitive, and the sensitivity was concentrated. **Thursday** is where it landed: the Straits Times Index fell **3.49 per cent** in a session and **4.13** on the week to **5,401.89**, the worst in Asia, because its three banks — more than half the index — were sold on a research note. A large American house argued that surging long yields mean bigger mark-to-market losses on trading books, called the quarter **"skewed negative"** and recommended trimming all three.[3]



Week to Friday 9 October 2026, per cent. A *Meridian* original, drawn from the closes at note 2.

THE CHATTER · WHAT IT MEANS FOR VWRA, GLOBAL BONDS, GOLD AND BTC

Name the regime: **higher-for-longer and dollar-bid**, not risk-off. A risk-off week does not leave a global equity fund green, and a growth scare bull-steepens a curve rather than leaving it at +44bp. For **VWRA**, a hawkish Fed is currently a sector problem, not an index problem. For **global bonds** the week is worse than it looks: on Thursday, as equities broke, the aggregate proxy moved **+0.0175 per cent**. Diversification did not hurt; it was not there. **Gold** is pricing the inflation leg. Explanation, not advice.[2]

THE SCEPTIC · ONE HOUSE SAYS HOLD, AND A DATE IS WRONG

Wednesday brings two prints at once: American September consumer prices, and the **Monetary Authority of Singapore's policy statement, confirmed on its own calendar for 14 October**. Three named houses expect a **25bp steepening**; one expects a hold. In July, four of eighteen surveyed economists called the surprise correctly, so the consensus deserves no deference. And a widely-read local retail outlet tells its readers the statement is due **"the last week of October"** — a fortnight wrong, against the authority's own date.[4]

+0.0175%

THE GLOBAL BOND PROXY'S MOVE ON THE DAY THE STRAITS TIMES INDEX FELL 3.49 PER CENT[2]

The anomaly to carry into Wednesday is gold: up **1.43 per cent** on Friday alone while the two-year yield rose five basis points and the dollar firmed — sign-inverted for any model regressing it on nominal yields. Nothing here is investment advice.[2]

SOURCES 1 The Federal Reserve's own minutes of the 15–16 September 2026 meeting, released 7 October 2026 and read at source, for the quoted characterisation of progress on inflation, the staff estimates of headline and core personal-consumption inflation for August, the unanimous support for the quarter-point increase to 3.75–4.00 per cent, the majority judgement that a further increase would likely be appropriate by year end, and the descriptions of the stance as not restrictive or only mildly restrictive; and the Committee's own 16 September statement for the target range and the absence of dissents. 2 United States Treasury daily par yield curve data for 2 and 9 October 2026, read at source, for both yields and the week-on-week changes; the two-to-ten spread and the Thursday bond-proxy move are computed from them and from consecutive closes. Every other closing level and weekly change is a per-instrument close from one commercial market feed on 9 October 2026, with the bitcoin weekend move to 11 October from the same feed; that single-provider dependence is stated here as it was in No. 129. 3 Reporting of 8 October 2026 by a financial wire, read at a syndication partner, for the research note, its quoted characterisation of the third quarter, the recommendation on all three banks, the second house's downgrade of Singapore equities and the two-day percentage falls. The index weighting of the three banks is a standing feature of the index. 4 The Monetary Authority of Singapore's own advance release calendar, read at source, for the 14 October 2026 date; the release time is not published there and is not printed here. The house calls are secondary relays of bank research and are flagged as such. The retail outlet's incorrect date was read on its own preview page on 11 October 2026.

THE INVITATION · LOCKSPORT

One pin binds first. That is the whole game

A lock opens to a pick because its chambers are imperfectly aligned. Find the order in which the pins bind and the plug turns — and then there are nine belts above you.

Inside a pin-tumbler lock a rotating **plug** sits in a fixed housing, and stacked spring-loaded pin pairs cross the boundary between them. That boundary is the **shear line**; the correct key lifts every stack until each split sits exactly on it. Picking exploits the fact that the chambers were never drilled perfectly in line: apply light rotational force with a **tension wrench** and exactly one pin binds first. Set it, the plug turns a few microns, the next binds. That sequence is the **binding order**, the lock's fingerprint.[1]

Two schools follow. **Single-pin picking** works the order deliberately; **raking** scrubs a serrated pick across all of them and hopes. Manufacturers defeat both with shaped drivers. A **spool** pin, in the hobby's own guide, **"has a smaller diameter along the center of the shaft than it does at the ends"**: it catches, the plug **counter-rotates** — **"fighting against you, as if trying to turn back"** — and you get a **false set**, which feels exactly like success. A **serrated** pin reports set when it is not. The guide's verdict: **"Picking spool pins is something that many novices find frustrating."**[1]

The moment people describe falling in is small and physical: a cheap padlock, a borrowed pick, and the first time a plug rotates under a thumb that did not use a key. The oldest English-language board, **Lock Picking 101**, carries **81,689 members**, **460,747 posts** and **60,591 topics**, peaking at **6,821** online on **8 April 2026**. The best-known figure, an anonymous former litigator who posts only his hands, has **4.68 million subscribers** and **1.25 billion views** since 2015 — and owns the shop selling the tools in the videos. At **DEF CON 34** the Lockpick Village was one of **38** villages, up from 34.[2]

The hobby is old enough to have lost institutions. The MIT guide went wide in **1991**; Germany's SSDeV was founded in **1997**. But **Keypicking.com**, for twenty years a pillar of the scene, is gone: the domain now serves an online-gambling review site.[3]

A false set feels exactly like success. That is what the spool pin is for.

THE CHATTER · THE RULE THAT DEFINES THE HOBBY

Locksport separated itself from both locksmithing and theft by writing its ethics down, unusually plainly: **"They may not open a lock that is in use." "They may open only locks that belong to them. For other locks, they need express consent of the owner."** The declared philosophy is **responsible full disclosure**, set against **security through obscurity** — and the trade body on the other side has said the rule should instead be **"security through ignorance"**.[4]

THE SCEPTIC · THE NUMBER WE COULD NOT READ

The figure everyone quotes for this hobby is its subreddit's membership, and this desk could not read it: Reddit refused us, and six mirrors too. The **285,000** members circulating for 1 October come from a third-party analytics site, not the platform, and are held as such. One counter-statistic is worth more: challenged over a video, a smart-lock maker argued picking accounts for about **4 per cent** of home break-ins.[5]

1 of 38

THE LOCKPICK VILLAGE AMONG DEF CON 34'S VILLAGES, UP FROM 1 OF 34[2]

WHAT THE NEXT TWO PAGES DESCEND INTO

The ladder — nine belts, from a transparent training lock to a Kromer (p16). **The bench** — what a first kit costs, and a custom tier where one hook outprices four of them. **The law here** — purpose, not possession (p17).

No jargon is used on the next two pages without translation, and the glossary that earns its place is printed where the vocabulary first bites.

SOURCES 1 The Open Locksport Guide's security-pins page, read at source on 11 October 2026, for the quoted definitions of the spool and serrated pins, the quoted description of counter-rotation, the false set and the quoted remark about novices; the plug, shear-line, binding-order and tension mechanism is standard and is set out on the same site. 2 The Lock Picking 101 forum's own front page, read at source on 11 October 2026, for the member, post and topic counts and the dated peak; the encyclopaedia article on the anonymous video-maker, infobox last updated 26 September 2026, for the subscriber and view counts, the 2015 start, the anonymity, the stated prior career and the ownership of the retailer; and DEF CON's own village listings for conferences 33 and 34 for the village counts and the August 2026 dates. 3 The encyclopaedia article on locksport for the 1991, 1997 and 1999 dates; the former forum domain itself, fetched on 11 October 2026, which now serves a gambling-affiliate review site, with its forum path returning a 404. The date the domain lapsed could not be established. 4 The same encyclopaedia article for the quoted ethical rules, the responsible-full-disclosure philosophy, the contrast with security through obscurity, and the quoted position attributed there to the American locksmith owners' trade body. 5 Reddit and six public mirrors all refused this desk and its researchers; the membership and growth figures are from a third-party analytics service dated 1 October 2026 and are not confirmed on the platform. The break-in percentage is a lock manufacturer's own figure, advanced in a dispute with the video-maker above, and is reported as that manufacturer's claim.

THE DEEP END · THE BELT LADDER

Nine belts, and the last one ends at a Kromer

The community ranks locks, not people. A transparent training lock and a Master Lock No. 3 share the bottom rung; eight tiers up sits a safe-lock from Freiburg.

The structure that makes this hobby legible is a ranking list of **locks**, maintained by the community and browsable as a belt ladder. Nine tiers: white, yellow, orange, green, blue, purple, brown, red and **black**, the last sub-ranked one to five. **White** holds eighteen entries, among them **any acrylic padlock**, **any cutaway lock** and the **entire Master Lock No. 1 to No. 8 range**. **Yellow** alone holds **71** locks.[1]

Then the climb, through named manufacturing. **Green** brings the **American Lock 1100**; **blue** the **Medeco Biaxial**; **purple** the **Mul-T-Lock Gamma**; **brown** the **Abloy Classic**; **red** the Abloy Disklock. Black opens with the Abloy Sentry and DOM Diamant, reaches the **Abloy Protec2** at black three, and ends at black five with one entry: the **Kromer Protector**. A beginner and a master are separated not by a grading panel but by a list of objects.[1]

Buy once, and the entry point is unusually well defined. **Sparrows' Night School, US$77**, is the best first purchase in the research: four picks, six tension wrenches, and **three cutaway practice locks** — standard, serrated and spool — so the three things you must learn to feel are in the box. The maker states its steel: **American 301 cold-max, around 280,000 PSI**. No-name kits state nothing, which is the case against them. Above it, the **Tuxedo Reloaded at $75** swaps a practice lock for a repinning kit of ten spool, ten serrated and ten mushroom pins.[2]

Thickness is where the hobby becomes a philosophy, and the market prices both sides equally. Thinner steel clears tight European keyways but flexes, muddying the feedback that *is* the skill; thicker steel transmits the binding order cleanly and will not fit. One American maker sells its flagship sets in parallel **0.025-inch and 0.020-inch** lines at **identical prices** — **$159.95, $188.50, $236.50** either way. Germany's Multipick works in metric at **0.4, 0.5 and 0.6mm**, its 39-piece Elite set **€267.75** with tax.[3]

Nobody is graded. The locks are. It is a more honest instrument than most sports manage.

THE CHATTER · THE BEGINNER'S MISTAKE, DATED TO THE DAY

The canonical error has an origin thread, opened **fifteen years ago today**. On **12 October 2011**, on a Master Lock No. 3, **eaglefrq** wrote: **"I was really putting a lot of pressure on the tension while trying to single pick." Squelchtone**, the same afternoon, gave the fix that is still the whole curriculum: **"Forcing it to turn isn't going to make it turn. insert the key and feel how much pressure it takes to turn the key."** An hour later eaglefrq had overcorrected — which is the second mistake, arriving the same day as the first.[4]

THE SCEPTIC · THE WALL IS AT WEEK FOUR

The hobby documents its own cliff. The guide names two failure modes arriving together on security pins: already-set pins dropping while you work the spool, and **oversetting**, where the force needed to pass the ledge **"push[es] the key pin past the shear line as well"**. White and yellow fall in a fortnight; green — one American padlock — is realistically a six-to-twelve-month target. Most people stop at the spool.[1]

US$77

NIGHT SCHOOL: FOUR PICKS, SIX TENSIONERS AND THE THREE CUTAWAYS THAT TEACH THE FEEL[2]

The ladder is also the shopping list. Progressive practice locks start at **two pin stacks** and climb, at **$35**; a spool cutaway is **$22.50**. Nothing on the next page costs less than that and does more.[2]

SOURCES 1 The community belt-ranking explorer, read at source on 11 October 2026, for the nine tiers and their sub-ranks, the eighteen white-belt entries including the acrylic and cutaway classes and the Master Lock range, the count of 71 yellow-belt locks, and every named lock quoted at green, blue, purple, brown, red and black, including the single black-five entry. The oversetting and un-setting mechanisms and the quoted phrase are from The Open Locksport Guide's security-pins page, also read at source. 2 The maker's own product pages, read at source on 11 October 2026, for the Night School contents and price, the stated steel and tensile figure, the Tuxedo Reloaded price and its repinning contents, the progressive-lock and cutaway prices; and the video-maker's own shop for the two lower-priced sets. 3 The American maker's own catalogue for the parallel 0.025-inch and 0.020-inch lines and their identical prices; a retailer's listings for the Pennsylvania maker's 0.015 and 0.018 variants, that maker's own site having returned HTTP 523 on every attempt across two hours; and the German maker's own Elite pages for the metric gauges and both euro figures, net and inclusive of tax. 4 The thread "How much tension" in the Beginner Questions board of Lock Picking 101, first posted 12 October 2011, read at source; all three quotations appear there with the usernames and timestamps as published.

THE DEEP END · THE LAW AND THE BENCH

Here the offence is purpose, not possession

Singapore's housebreaking-implement provision turns on what a tool is for. The last four people charged under it were carrying screwdrivers and spanners.

Read the statute, not the forum. Section **453** of the Penal Code 1871, current as at **12 October 2026**, makes it an offence to be found **"equipped with any article or instrument for use in the course of or in connection with any housebreaking"** — up to **two years**, or a fine. The purposive words are part of the offence, and the reverse burden in the next subsection is written for the weapons limb, not for implements. Picks are not a prohibited article here; that is what the statute says, and it is not advice.[1]

Enforcement is routine and instructive. In **August 2025** the police charged four men under the same section — and the implements were **screwdrivers and spanners**. This is an any-object-plus-intent law. Customs' published prohibited list mentions neither picks nor locksmith tools — an absence from a list, not a ruling.[2]

There is a scene here. **Locksport SG**, founded **2016**, has more than **50 members**; founder **Zach Tan** told a local outlet last month that the group **"only practise on locks that we own"** and that it is **"similar to solving a Rubik's Cube… which just happens to involve a lock"**. It held a beginners' workshop on **10 October**.[3]

Then the bench. A hand-made hook from a named artisan runs **$145 to $340**; one maker's collectors sets are **$650** and **$666**. The retailer's artisan collection carries **110 listings, every one sold out** — a drop culture, not a waitlist.[4]

One hand-cut hook costs between two and eight complete beginner kits. All 110 were sold out.

THE CHATTER · THE NINETY DAYS, AND WHAT THEY COST

The honest list is short: **Night School $77**, **progressive locks $35**, a **Master Lock No. 3**. Roughly **US$120–130** with freight. Days one to thirty you use far too much tension; thirty to sixty you meet the spool wall. The state of the art, for contrast: the retailer's next flagship set is **"currently in development and… not available for purchase"** — a page, a credit, no price, no date.[5]

THE SCEPTIC · THE LOCK THAT SAYS IT WILL NOT BE PICKED

The ceiling is admitted by both sides. A Canadian maker's disc-detainer cylinder puts its disks around a corner, springs the sidebar in reverse so torque binds it, and gives every disk **three or more false gates matching the true gate in every dimension except depth** — then states on its own page: **"we do not think that this will be a lock that will be picked in real life scenarios any time soon."**[6]

THE GLOSSARY THAT EARNS ITS PLACE

SPP single-pin picking. **Raking** scrubbing all pins at once. **Binding order** the sequence drilling tolerance creates. **False set** a security pin caught, not opened. **Overset** a key pin pushed past the line. **Impressioning** cutting a key from a blank by reading its marks.

NEXT DESCENTS

Bookbinding & book repair — sew, glue and rebuild a ruined book into an heirloom.

Orienteering — map, compass, forest, clock: the thinking runner's endurance sport.

Close-up card magic — years of private practice for eight seconds in someone's hands.

The genuine modern story is not a pick. In March 2025 six university researchers broke a Bluetooth deadbolt five ways — replay unlocking, expired guests still entering, forged audit logs.[7]

SOURCES 1 Singapore Statutes Online, Penal Code 1871, section 453, current version as at 12 October 2026, read in full: the quoted words of paragraph (1)(c), the penalty, and subsection (2), whose reverse burden applies to paragraph (1)(a) only. That reading is this desk's, from the text; it is description, not legal advice. 2 The Singapore Police Force's release of 10 August 2025 for the four arrests, the charge and the nature of the implements; and Singapore Customs' controlled-and-prohibited-goods page, read 11 October 2026, which lists neither picks nor locksmith tools — an absence rather than a ruling. The Miscellaneous Offences Act 1906 was downloaded in full and searched: it contains nothing on housebreaking implements and is not cited. 3 A Singapore news site's feature of 15 September 2026, read at source, for the founding year, membership, founder and all three quotations, and for the 10 October workshop, whose venue and price are unpublished; the 2022 conference village is on that conference's own archived programme. 4 The retailer's custom-pick collection, read across both pages on 11 October 2026, for the 110 listings, their sold-out status and the maker prices. 5 The makers' own product pages for the kit and practice-lock prices, and the retailer's own page for the forthcoming set, read 11 October 2026, for the quoted development status and the absence of a price or date. 6 The Canadian lock maker's own page for the false-gate count, the reverse-sprung sidebar and the quoted claim, which is the manufacturer's own and untested here; and the convention organiser's own listing for the ticket price. 7 "No Key, No Problem: Vulnerabilities in Master Lock Smart Locks", 19th USENIX WOOT Conference, 11–12 August 2025, read as published: the five attack classes, the 10 March 2025 report, the 14 March vendor response and the June 2025 firmware fix. No CVE identifiers were assigned.

THE LONG GAME · PROTEIN AND MUSCLE

Triple the protein, and the difference is nothing

In 108 women averaging 48, twelve weeks of identical training on 0.8, 1.6 or 2.2 grams per kilo produced no detectable difference in muscle mass. The p-value was 0.726.

The design can answer the question. A randomised two-by-three trial, published **14 September**, put **108 women aged 40 to 77** through twelve supervised weeks of resistance training alone or the identical programme plus cycling, at **0.8, 1.6 or 2.2 grams of protein per kilogram per day**. **83** completed, mean age about **48.6**, achieved intakes tracking the targets.[1]

On skeletal muscle mass there was no training-by-protein interaction (**p = 0.856**), **no main effect of protein (p = 0.726)** and none of training (p = 0.273). Tripling the intake changed nothing the study could measure. What moved was everything else: the concurrent group finished with body fat **2.04 points lower** (95% CI −2.94 to −1.14), the resistance group with leg-press one-rep maximum **6.68kg higher** (5.04–8.32), and the concurrent group with estimated aerobic capacity **4.53 ml·kg^{-1}·min^{-1} higher**. All three at p < 0.001.[1]

The authors print the caution themselves: **"Null protein findings do not establish equivalence among doses."** They also name the weaknesses — impedance rather than scanning, self-reported food records with **"no controlled feeding or objective biomarker of protein intake"**, **23.1 per cent attrition**. Absence of evidence from a well-run trial, not proof of equivalence.[1]

A second randomised trial the same month lands the same way on a different optimisation. **108** older adults, **93** analysed, mean age **67**, double-blind, trained twelve weeks at each person's measured **peak** performance hour or at their **trough**. Mid-thigh pull: **0.06 N·kg^{-1}** (−0.76 to 0.88). Handgrip: **−0.21**. The verdict: individualised diurnal timing is **"unlikely to provide clinically meaningful benefits"**.[2]

Two well-run trials, two fashionable levers, two intervals straddling zero.

THE CHATTER · THE INTENSITY ARGUMENT, IN PUBLIC

The liveliest readable thread of the season — **142 points, 74 comments** on an exercise-intensity paper — was mostly people correcting its framing. **jeffbee** disputed calling a thirty-second cycling effort "all-out", since any cyclist knows you pace it. **jandrewrogers** reported the opposite from experience: two years of moderate endurance plateaued, and a short high-intensity block produced the gains. **mimo84** closed it — **"the study isn't new"**.[3]

THE SCEPTIC · LEAN MASS IS NOT MUSCLE

Stuart Phillips of McMaster posted a preprint on **7 October** going further than the trial above. Head-to-head comparisons of **1.6 against up to 3.2** grams per kilogram **"found no clear advantage for lean mass"**; over three months, fat-free mass rises about **1.1kg**, supplemental protein adding **0.3kg**. And the measurement is suspect: **"nearly every trial reports fat-free or lean mass, which is not muscle"** — in one comparison scanner-derived lean mass exceeded a tracer measurement by **about 20kg**. His conclusion: **"An intake of 1.6 g/kg/day appears sufficient."**[4]

p = 0.726

THE PROTEIN MAIN EFFECT ON MUSCLE MASS, ACROSS A NEAR-THREEFOLD DOSE RANGE[1]

Pick the adaptation, not the macro: the same twelve weeks buys either two points of body fat or 6.7kg of leg press depending on whether you add the bike — a better-measured effect than anything the protein arm produced. General guidance, not medical advice.

SOURCES 1 Bagheri and colleagues, "Resistance versus concurrent training with three assigned protein targets in middle-aged and older women: a randomized 2 × 3 factorial trial", *Journal of the International Society of Sports Nutrition* volume 23, epub 14 September 2026, full text read at source: the design and duration, the three protein targets, the 108 randomised and 83 analysed, the achieved intakes, all three p-values on skeletal muscle mass, the three between-group differences with their confidence intervals, and the authors' own limitations including the quoted equivalence caveat and the quoted statement about feeding control and biomarkers, the impedance measurement, the unmatched exposure, the estimated aerobic capacity and the 23.1 per cent attrition. The mean age of completers is computed by this desk from the published group means; the paper reports no overall mean. 2 Bruggisser and colleagues, "Effects of Daily Rhythm-Aligned Training on Skeletal Muscle Adaptation in Older Adults: A Randomized Controlled Trial", *Journal of Cachexia, Sarcopenia and Muscle* volume 17, read at source: the randomised and analysed counts, the mean age, the double-blind registered design, the three effect estimates with their confidence intervals and the authors' quoted conclusion. 3 A Hacker News thread of 11 September 2026 on an exercise-intensity paper, read at source, for the score, the comment count and the four quoted or summarised comments with their usernames. 4 Stuart M. Phillips, "Show me the muscle: how much can protein really add to exercise gains?", preprint posted to a sport-science server on 7 October 2026, read at source, for all four quotations and both mass figures and the tracer comparison. A preprint is not peer-reviewed and is flagged as such. The intake range being qualified is the one popularised in longevity media. General guidance, not medical advice.

THE DIARY · THE DATED AGENDA

Brazil plays twice in one Singapore week

A flash sale on the cheapest of those nights closes on the 19th. Fukuoka's sumo sold out all fifteen days; Tokyo's opens 5 December.

SINGAPORE

BOOK NOW · PLAN AHEAD

13–17 NOV ★ **Mizuho Blue Challenge — Kallang Football Series**
NATIONAL STADIUM · 13 NOV SINGAPORE V PARAGUAY; **14 NOV JAPAN V BRAZIL**; 17 NOV SINGAPORE V BRAZIL AND JAPAN V PARAGUAY · VIP $888 TO CAT4 $98, **$28** FOR THE OPENER · **FLASH SALE, SECOND TICKET HALF PRICE, ENDS 19 OCT**[1]

17–18 OCT **BIGBANG 2026–2027 World Tour**
NATIONAL STADIUM, 7PM, THIS WEEKEND · VIP ULTIMATE STANDING $498 TO CAT6 $158[1]

10–11 NOV **My Chemical Romance**
INDOOR STADIUM, 8PM · THE APRIL DATE WAS RESCHEDULED HERE AND A SECOND NIGHT ADDED · $408 TO $128[1]

25 NOV **Guns N' Roses — World Tour 2026**
NATIONAL STADIUM, DOORS 6PM · BASE PRICES NOT PUBLISHED; THREE VIP PACKAGES ONLY[1]

4–6 DEC **Singapore International Marathon — first three-day edition**
PIT BUILDING TO NATIONAL STADIUM · **REGISTRATION CLOSED 30 SEP** · NEW TITLE SPONSOR[1]

10–13 DEC **PGL Major Singapore 2026 — finals**
INDOOR STADIUM · $144 TO **$369** FOR THE GRAND FINAL · LATE FINISHES[1]

17–22 DEC ★ **BTS — *Arirang***
NATIONAL STADIUM, FOUR NIGHTS: 17, 19, 20, 22 DEC · VIP $388 TO CAT6 $148, BOOKING FEE $20 ABOVE $300 · THE VENUE PAGE SHOWS NO SOLD-OUT STATUS[1]

8–10 JAN **Jay Chou — *Carnival II***
NATIONAL STADIUM, 2027 · CAT1 $398 TO CAT5 $238 · CARD PRESALES RAN IN MAY; **NO GENERAL-SALE DATE PUBLISHED**[1]

30–31 JAN ★ **HSBC SVNS Singapore**
NATIONAL STADIUM, 2027 · DATES CONFIRMED ON THE VENUE'S OWN PAGE, WHICH CARRIES ONLY A REGISTER-INTEREST FORM — **NO PRICES, NO ON-SALE**[1]

20 FEB **YOASOBI — Asia dome & stadium tour**
NATIONAL STADIUM, 2027 · PLATINUM VIP $498 TO CAT7 $108 · WENT ON GENERAL SALE 19 AUG WITH NO PRESALE AT ALL[1]

JAPAN & THE REGION

THE ONES WITH A TICKET CLOCK

10–24 JAN ★ **Grand Sumo, Hatsu Basho**
RYŌGOKU KOKUGIKAN, TOKYO, 2027 · **ADVANCE SALE OPENS SAT 5 DECEMBER** — NOVEMBER'S FUKUOKA BASHO IS **SOLD OUT ON ALL FIFTEEN DAYS**, NO DOOR TICKETS[2]

31 OCT–28 MAR **Mariko Mori — *All That Shines***
MORI ART MUSEUM, ROPPONGI · THE FIRST RETROSPECTIVE SINCE 2002 · ADULTS ¥2,800 WEEKDAY; **JUNIOR HIGH AND UNDER FREE**; ADVANCE BOOKING NOT YET DATED[3]

30 OCT–11 JAN **HAIKYU!! Challengers**
MORI ARTS CENTER GALLERY, ROPPONGI · DATES ON THE MUSEUM'S OWN LISTING; **PRICES NOT PUBLISHED** — THE STRONGEST TOKYO PICK FOR A TWELVE-YEAR-OLD[3]

14 NOV–28 MAR **Masterworks from the Musée d'Orsay**
TOKYO METROPOLITAN ART MUSEUM, UENO · CENTENARY SHOW, ABOUT 110 WORKS · ¥2,400, ADVANCE ¥2,200 · **CLOSES 28 MARCH, NOT 29** AS THE ROUND-UPS HAVE IT[3]

7 MAR **Tokyo Marathon 2027 — and you have missed it**
THE TWENTIETH EDITION · GENERAL ENTRY CLOSED 28 AUG, BALLOT DRAWN · **OVERSEAS VOLUNTEER REGISTRATION 4–15 NOV** IS THE WAY IN[4]

14–28 MAR **Grand Sumo, Haru Basho**
OSAKA, 2027 · ADVANCE SALES OPEN **6 FEBRUARY** — THE FALLBACK IF THE TOKYO ON-SALE GOES THE WAY FUKUOKA'S DID[2]

27–28 MAR **AnimeJapan 2027 — and it has moved**
INTEX OSAKA, NOT TOKYO BIG SIGHT · CONFIRMED ON THE EVENT'S OWN ENGLISH PAGE[5]

11–31 JAN **Australian Open 2027**
MELBOURNE · DATES ON THE TOURNAMENT'S OWN SITE; **NO ON-SALE PUBLISHED**[5]

25–27 MAR **Art Basel Hong Kong 2027**
HONG KONG · DATES ONLY ON THE FAIR'S OWN PAGE — NO VENUE, NO PRICES, NO ON-SALE[5]

9–11 APR **Hong Kong Sevens — moved out of the window**
KAI TAK STADIUM, 2027 · THE SITE SAYS ONLY "TICKETING DETAILS TO FOLLOW" · THE SINGAPORE TENNIS OPEN HAS SLIPPED TO **20–26 SEP 2027**[5]

SOURCES **Every Singapore entry re-verified on the venue operator's own event pages on 12 October 2026**, the national ticketing agent and the stadium's own site having refused this desk. 1 The operator's own pages for each event, read at source: the football series' four fixtures, its four-tier price table, its 4 September general sale and the flash-sale window and terms; the two BIGBANG nights and their price range; the rescheduling of the April rock date to November, the automatic transfer of existing tickets and the added second night; the stadium rock date, its doors time, the three VIP packages and the absence of base prices; the marathon's first three-day edition, its closed registration and its changed title sponsor; the esports finals and their four day prices; the four December stadium nights, their full price table and booking fees, and the absence of any sold-out marking; the January concert prices and the unpublished general-sale date; the rugby dates with their register-interest form and absent pricing; and the February stadium date, its prices and its no-presale general sale. 2 The Japan Sumo Association's English ticket service, read at source, for the Tokyo advance-sale date of 5 December, the Osaka date of 6 February, and the quoted sold-out status of all fifteen Fukuoka days with no day tickets, against its 19 September on-sale. 3 The Mori Art Museum's own exhibition page for the retrospective's dates, its admission table, the free entry for junior-high and younger and the unannounced advance-booking date, and its own listing for the gallery show next door; the Tokyo Metropolitan Art Museum's own page for the Orsay show's dates, work count, prices and closure days, which also corrects the 29 March closing date circulating in aggregators. 4 The marathon's own English site for the 20th edition date, the closed general entry, the drawn ballot and the overseas volunteer window. 5 Each event's own official site, read at source, for the dates quoted and for the absence of pricing or on-sale information where stated. An event recorded as cancelled in this magazine's events ledger is excluded from this agenda.

THE TABLE · OCTOBER'S OPENINGS

A second star here, and a queue that has not formed

The guide promoted a Kampong Glam dining room after three years at one star. Meanwhile the most anticipated opening in town still has no date and no confirmed address.

The week's real news is a promotion. **Seroja**, which this strip first covered at No. 63, has gone from one star to **two** in the 2026 Singapore guide after three consecutive years at one. In the same list **1887 by André** debuted **straight at two stars**, and six restaurants took first stars. One caveat: the guide's own selection page returned no readable text to this desk, so the list comes from a local title carrying it in full.[1]

The most anticipated arrival has nothing you can book. The Tokyo **nama donut** brand **I'm Donut?** teased Singapore on its global account on **28 September** — a hoarding, the question-mark logo, **"Singapore? 2026 — Coming soon"**. A food title lists an Orchard mall address with every price to be confirmed; two other outlets state plainly that **no opening date has been announced and the location is not known**. This desk prints neither. The precedent is confirmed: its Kuala Lumpur debut on **12 September** drew more than **200 people queueing before doors**.[2]

The one that is actually open is at the airport. **VE/LA**, the Bangkok coffee and matcha house founded in 2021 by **Pete Kasidit Prasitrattanaporn** when he was sixteen, opened on **Friday 25 September** at **Terminal 1, Arrival Meeting Hall West — landside**, so no boarding pass, about twenty seats, and **open 24 hours**. It is the chain's third country, its first airport and its first round-the-clock site anywhere. A Singapore-exclusive **Kaya Toast Latte is $10**; the Airecano is **$8.50** and the pecan cinnamon bun **$6.80**.[3]

Two verdict runs close the page. **GAMJA** on Boon Tat Street, from No. 113, has settled into a CBD lunch room: pork-bone stew from **$15.80**, sharing pots from **$53**, with a white, chilli-free version. **Kiap Kiap Chicken Club & Bar** on Liang Seah Street, from No. 114, runs to **1am** and **2am** at weekends, fried chicken from **$12** — and the dish a round-up singles out is not the chicken but the udon.[4]

THE CHATTER · THE AIRPORT COFFEE, TASTED

The one tasting panel to have filed on **VE/LA** did not choose the signature. Its unanimous favourite was the **Airecano**, **"aromatic, without any bitterness"**, while the Singapore-exclusive kaya latte was filed as **"rich and sweet, but with a savoury edge from the salted egg"** — which is a dessert, honestly described. The bun it called **"addictive, sticky-sweet cinnamon sugar layers"**. The founder's own line to a local site: **"Singapore feels like a place that perfects what it does."**[3]

THE SCEPTIC · THE HYPE THREAD IS MAJORITY AGAINST

The largest readable Singapore conversation about **I'm Donut?** ran on a local hardware forum on 29 September, and it is overwhelmingly unimpressed. A moderator, **zensushi**: **"With a more health-conscious population, I Donut think it will do too well."** **madcampus** produced the precedent that matters: **"Even Pablo the famous cheese tart from Osaka opened in Wisma also closed shop barely 2 years."** **Hello_Kitty**: **"after one year, no queue"**. **addict951**: **"Give it 1 year. Thx"**. The counter-view came from the same poster who supplied the precedent — **"The most hyped donut from Japan is coming to Singapore. Expect long queues."** Both can be true, eleven months apart.[5]

TWO STARS · TOO NEW TO CALL · NO DATE, NO ADDRESS · LANDSIDE, OPEN ALL NIGHT

If you want one booking from this page: the airport, at four in the morning, for a coffee you do not need a boarding pass to drink. Everything else here is either already decided by a guide or not yet decided by anyone.

SOURCES 1 A Singapore lifestyle title's coverage of the 2026 Michelin Guide Singapore selection, read at source, for the promotion to two stars after three consecutive years at one, the two-star debut, and the six new one-star entries. The guide's own selection page returned an empty body to this desk, so the ceremony date and venue are not printed and the list is attributed to the secondary carrier. 2 The brand's own global account post of 28 September 2026 for the teaser and its quoted caption; two independent titles of 29 and 30 September, both read at source, for the explicit statements that no opening date has been announced and that the location is not known; a food title's October listing for the mall address and the prices marked to be confirmed, which this desk declines to print as confirmed; and a Malaysian food site for the 12 September Kuala Lumpur debut and the reported queue. 3 The brand's own Singapore page, read at source, for the 25 September opening date, the unit number and terminal location, the landside position, the twenty-four-hour operation and the first-airport and first-round-the-clock claims; a Singapore food title's review of 25 September for the tasting-panel verdicts quoted and the prices; and a local site for the founder's quoted remark and his age at founding. 4 A Singapore lifestyle round-up, read at source, for the Korean restaurant's dishes, lunch and sharing-pot prices, address and hours, and for the chicken bar's hours, prices and the dish it singles out. Both venues were introduced in this strip at Nos. 113 and 114 respectively and are given their verdict run here. 5 A Singapore hardware forum thread of 29 September 2026, read at source; all five quotations appear there with the usernames as published. Reddit and the main local review platform both refused this desk, so no sentiment from either is reported.

THE LONG READ · HUMAN FEED

The word arrived three years after it was predicted

In 2023 a Chinese user wrote that the day of human feed might not be far off. By this summer office workers were freezing a week of dinners in trays and calling it exactly that.

The joke names itself. 人粮 — *rénliáng*, "human grain" — is built on 猫粮 and 狗粮, cat food and dog food, and it means what it sounds like. On a Sunday a Chinese office worker blanches a week of vegetables, shrimp and chicken, portions them into trays, freezes them at −18°C and microwaves one for five minutes a day. The variants in circulation are 人饲料, human fodder, and 预制人粮, pre-made human feed. Nobody is being ironic.[1]

The arithmetic is the argument. A first long feature on the practice, published on **1 July**, reports a supermarket shop of **under ¥100** feeding one person for a week — **under ¥10 a meal**, against **¥30 to 40** for a light-meal delivery and the **¥1,500 a month** one interviewee had previously spent having dinner brought to her. A thirty-year-old in Beijing is quoted saying she can never go back, and that the only skill required is the ability to chop. By **20 September** the topic had passed **ten million views** on the lifestyle platform where it lives.[1]

It did not come from nowhere. The precursor, 冻门 — "frozen door" — was already a story in the state English-language press in **February 2025**, with roughly **3.7 million posts** on the same platform and a ¥10 homemade lunch set against a ¥20–25 takeaway. That version was about time. A group on another platform called 吃饭糊弄学 — roughly, the study of fobbing yourself off at mealtimes — has more than **30,000** members. What changed between then and now is not the freezer. It is why people are using it.[2]

> ***"The day of human feed may not be far off." Posted in September 2023, in the middle of an earlier row.***

The prophecy is on the record. In **September 2023**, during an earlier public argument about pre-made restaurant food, a user on the same platform wrote that **"the day of human feed may not be far off"**. Three years later the word exists, the practice has a name and a nutritionist has had to go into a provincial daily to tell the country not to live on it.[3]

Two numbers explain the demand better than any manifesto. Chinese enterprise employees averaged **48.6 hours a week** in 2025, against a statutory cap of 44. And only **12.7 per cent** of Chinese young people cook daily, with **5.4 per cent** never cooking at all. One netizen's accounting, as reported: two hours commuting, eleven at work, eight asleep. Twenty-four hours does not stretch to a kitchen. Another line summarising the trend has travelled further than any of the statistics — **five minutes to eat, two hours to cook**.[4]

The language people use about it is not the language of a life hack. The sharpest line in circulation, as reported, divides the two things: **"有的叫生活，有的叫活着"** — some people call this living, others call it merely staying alive. Another commenter says eating **should be something that makes you happy, and in this society it has been alienated into a task**. A convert is candid about the cost: 箍秋 says the flaw is that human feed is essentially boiled vegetables, so when she wants something fried it becomes cloying. And a food-engineering writer gives the diagnosis from inside: the practice comes above all from **a loss of trust in the outside world**.[5]

What follows is not a diet story. It is a price-discovery event, and the state is on the wrong side of it.

SOURCES 1 A long feature published by a Chinese business-culture site on 1 July 2026, read at source in the original, for the practice as described, the weekly shop under ¥100, the per-meal figure, the delivery comparisons, the monthly delivery spend of the interviewee, the quoted conviction that she cannot go back, and the construction of the term by analogy with the words for cat and dog food; and a news-portal feature of 20 September 2026, read at source, for the ten-million-view figure on the lifestyle platform and for the circulating variants of the term. 2 A state English-language outlet's report of February 2025, read at source, for the precursor movement, its approximate post count and its lunch-cost comparison; and the 1 July feature above for the membership of the mealtime group. 3 A Chinese-culture magazine's feature on the earlier prepared-dish backlash, read at source, for the quoted September 2023 post. The nutritionist's intervention is cited and sourced on page twenty-three. 4 China's National Bureau of Statistics annual release for 2025, read at source, for the average weekly hours worked by enterprise employees; a youth cooking-behaviour survey reported by a Chinese data title for the daily-cooking and never-cooking shares; and the 20 September feature at note 1 for both the netizen's reported accounting of the day and the quoted summary line, each attributed there as reported rather than interviewed by this desk. 5 A commentary site's repost of 29 September 2026 of a magazine feature, and a news-portal feature of 1 August 2026, both read at source, for the quoted Chinese formulation and its gloss, the remark about alienation, the named convert's complaint and the food-engineering writer's diagnosis. All are summarised here in this desk's words except where quoted, and all are attributed as reported by those outlets rather than gathered by us.

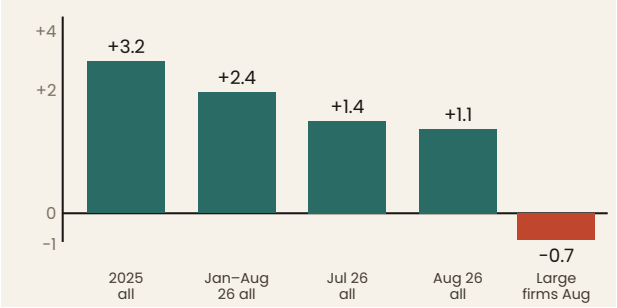
To see what people are refusing, go back to the row that started it. On **10 September 2025** a well-known entrepreneur attacked a large northwestern restaurant chain on the microblog platform for serving expensive pre-made dishes. Its founder denied it on the 11th, threw open the kitchens of **more than 370 stores** from the 12th and threatened to sue. By **16 January 2026** the chain was closing **102 outlets** — about a third of its estate — after footfall fell **45 per cent** at the worst point, and its listing plans were damaged.[1]

Here is the thing almost everyone got wrong, including the chain. The complaint was never pre-made food as such. In a 2024 survey **more than 90 per cent** of Chinese respondents said they accept prepared dishes in principle, and **more than 60 per cent** pay **¥21 to 40** for one. What people objected to was paying a restaurant's assembly premium for a factory's product — a pricing argument wearing a food-safety costume. Once a consumer has worked out the input cost of a meal, that premium is not recoverable.[2]

The state's answer, when it came, confirmed the suspicion rather than settling it. On **22–23 January 2026** the food-safety office announced it had drafted a national standard for prepared dishes and a terminology-and-classification standard. On **6 February** both went out for comment — and with them a notice on disclosure of preparation methods in restaurants that is explicitly **encouraging rather than mandatory**. The standard's definition, meanwhile, **excludes dishes made in a chain's own central kitchen**, which is most of what people were angry about. Asked for transparency, the state offered to suggest it.[3]

Asked to make restaurants say what was reheated, the state drafted a notice encouraging them to.

That is the gap this essay is about. A trust failure was answered with a voluntary instrument, so households supplied the transparency themselves — by moving the production step inside the flat. The industrial kitchen was not rejected. It was **insourced**. The freezer tray is a factory the consumer owns, and its appeal is precisely that nothing about it is hidden.[3]



Year-on-year growth in Chinese catering revenue, per cent, nominal. A *Meridian* original, drawn from the statistics bureau's own monthly and annual releases.[4]

There is a further reason the insourcing sticks, and it is the one the policy cannot touch. A prepared dish eaten at home is cheap and nobody minds; the same dish served at a table carries a charge for an act — cooking — that did not take place there. Once a household has priced the inputs, that charge stops looking like a service and starts looking like a fee. A consumer who has worked out a cost does not unwork it. Which is why the behaviour is reported as irreversible by the people doing it rather than as a phase.[2]

And the demand was not destroyed, only relocated. The delivery platforms had their best quarter in years over the same period: one reported second-quarter revenue of **¥104.6 billion**, up **14.4 per cent**, with core local commerce operating profit up **52.3 per cent** — achieved, its own commentary notes, with **markedly reduced** delivery subsidies. People are still spending on food. They have simply stopped paying for the room, the service and the reheating as separate line items.[5]

Which brings us to the figure that turns a meme into an economic event. In **August 2026**, catering revenue at China's above-designated-size units — the formal, countable restaurant sector — was **¥138.7 billion**, **down 0.7 per cent** on the year. Not slowing. Contracting.[4]

SOURCES 1 A Beijing newspaper's report of 11 September 2025, timed 19:53 and read at source, for the attack of 10 September, the founder's denial, the opening of more than 370 store kitchens from 12 September and the threat of litigation; and a business wire's report of 16 January 2026, read at source, for the closure of 102 outlets, the proportion of the estate, the 45 per cent fall in customer traffic at its worst and the damage to the listing plan. 2 A news-portal feature of 20 September 2026, read at source, for both survey figures — the share of respondents accepting prepared dishes in principle, attributed there to a 2024 research institute study, and the share paying ¥21 to 40 for a prepared dish, attributed to a market-research firm. Neither underlying study was read by this desk and both are reported as relayed. 3 The state news agency's report of 23 January 2026 for the announcement of the drafted national standard and the terminology-and-classification standard, and for the accompanying disclosure notice; and reports of 6 February 2026, read at source, for the opening of both consultations, the comment deadlines, the explicitly encouraging rather than mandatory character of the disclosure notice, and the exclusion of chain central-kitchen dishes from the standard's definition. 4 China's National Bureau of Statistics, monthly release of 15 September 2026 and annual release of 19 January 2026, both read at source: the August figure for above-designated-size catering and its year-on-year change, the January-to-August cumulative figures for all catering and for above-designated-size catering, the July figure, and the 2025 full-year growth rate. The chart is drawn from those releases alone. 5 A Chinese financial daily's report of 28 August 2026 and a state financial wire's of 29 August, both read at source, for the platform's second-quarter revenue and growth rate, the core local-commerce operating-profit growth, and its own characterisation of reduced delivery subsidies. The inference drawn from those figures is this desk's.

The restaurant count is the other half, and it is brutal. In the first half of 2026 China closed **2,055,200** outlets and opened **1,650,000** — a net loss of **410,000** — taking the national total from about **8.40 million to 7.99 million in six months**. New registrations fell **17.3 per cent**. Among listed operators the pattern is margin: one chain grew sales **7 per cent** while net profit fell **24.3**, because average spend per diner dropped from **¥65 to ¥50.50**.[1]

Now put the household's capital spending beside it. In the first seven months of 2026, Chinese **chest-freezer** sales rose **13.2 per cent** by volume and **9.9** by value to **¥7.28 billion**, while **refrigerator** volumes fell **4.9** and value **7.1**. The whole home-appliance market shrank **9.9 per cent** in the first half, to **¥425 billion**, with only chest freezers and dishwashers growing. In a market where households bought almost nothing, they bought the one machine that lets them stop eating out.[2]

Households bought almost no appliances this year. The exception was the machine that replaces a restaurant.

The sceptical reading is serious. First, the freezer line may be an artefact: the research house itself calls it a seesaw with refrigerators, driven by a lower average price as consumers buy a cheap supplementary unit instead of a new fridge — and forecasts the cooling category to decline mildly this year. Second, catering is still outgrowing goods, **+2.4 per cent** against **+1.0**. Third, the closures are a supply story, the unwinding of a 2023–24 opening wave concentrated in heavy-asset formats. Fourth, the precursor movement was documented in **February 2025**, before the squeeze, so the austerity reading may be retrofitted.[3]

The fifth objection has a hospital behind it. The head of nutrition at a Nanjing maternity hospital told a provincial paper on **23 September** that freezing and reheating strips the water-soluble vitamins, that human feed should be one component of a meal rather than a replacement, that leafy greens should never be pre-frozen, and that she does not recommend relying on it long-term. The party paper ran the same warning on **30 September**. A diet a hospital tells you not to live on has an expiry date built in.[4]

The complaint was never the factory. It was being charged a restaurant's price for a factory's work.

So what does it mean? That China's consumption problem is being described wrongly. The usual account is of households without money. What 人粮 shows is households *with* money declining a specific transaction — the one in which you pay a service premium for an assembly step you have now priced. Once a consumer has done that sum, no amount of stimulus unwinds it, because the thing being refused is not a price but a markup. Eating has been reclassified from pleasure to maintenance, and a maintenance task is bought on cost.

Three dated tests will settle it. The statistics bureau publishes again around **16 November** and the annual figures around **17 January**: this essay is right if above-designated-size catering prints negative year-on-year in at least two of September to December, and full-year catering growth lands at or below **2.4 per cent** against 2025's **3.2**; it is wrong if full-year growth reaches 3.2 again. The appliance research house publishes full-year data in **January**: right if chest-freezer value growth stays positive while refrigerators stay negative, wrong if freezers turn negative and the seesaw explanation wins. And by **31 March 2027**, if Beijing makes restaurant disclosure mandatory rather than encouraged, the central claim here — that consumers had to supply their own transparency — weakens considerably.[5]

A government that answered a trust crisis with a suggestion got the outcome a suggestion earns. Eight point four million restaurants became seven point nine nine, and the transparency arrived anyway, one freezer tray at a time.

SOURCES 1 A Chinese restaurant-industry publication's analysis of 22 August 2026, read at source on a news portal, for the first-half opening and closing counts, the net change, the fall in the national outlet total, the decline in new company registrations attributed there to a corporate-data provider, and the listed chain's revenue, profit and average-spend figures. The outlet counts are attributed in that piece to a brand-monitoring firm. 2 An appliance trade magazine's report of 7 September 2026, read at source, for the January-to-July chest-freezer volume and value growth, the value figure in yuan, and the refrigerator volume and value declines, all attributed there to a market-research house; and an economics daily's report of 3 August 2026 for the first-half home-appliance retail total and its decline, and for the statement that only chest freezers and dishwashers grew. 3 The same appliance trade report for the research house's own seesaw explanation, its attribution of freezer growth to a lower average price and supplementary-unit buying, and its mild-decline forecast for the cooling category; the statistics bureau's release of 15 September 2026 for the January-to-August catering and goods-retail growth rates; the industry analysis at note 1 for the composition of the closures; and the February 2025 state English-language report cited on page twenty-one for the precursor movement's framing. 4 A Jiangsu evening paper's health report of 23 September 2026, read at source, for the named hospital nutrition head's warnings on water-soluble vitamin loss, meal replacement, leafy greens, long-term reliance and the weight-loss and anti-inflammatory claims; and the provincial party daily's version of 30 September 2026. 5 The statistics bureau's and the research house's own publication calendars for the dates given; the comparison figures are those at notes 2 and 4 of page twenty-two. The three tests are this desk's, and are logged in this magazine's predictions ledger for grading.

MERIDIAN

THE LAST WORD

What did not happen

A threat-warning system covering a capital issued nothing at all on Saturday, and thirteen people died in an airport hall (p3). A response ladder published the same week in Singapore begins at an index of two hundred, and the index stopped at a hundred and ninety-two (p5). Tripling a protein dose across twelve weeks of identical training moved muscle mass by an amount with a p-value of 0.726 (p18).

None of those is a story about failure. They are stories about instruments that exist, were built carefully, and did not fire — a more awkward thing to report, because there is no event to photograph. A pull request that made decoding 1.85 times faster made prompts 27 per cent slower, and printed only the first number (p8).

Holmes, on page two, is the patron saint of this edition: the dog did nothing in the night-time, and that was the curious incident. Beijing answered a crisis of trust with a notice *encouraging* restaurants to say what they had reheated, and households bought freezers instead (p21). An absence is evidence. It is simply harder to put on a cover.

TWELVE DESKS · ONE WORLD

MERIDIAN · No. 130 · 12 October 2026

A note on method. This edition's load-bearing figures were read where they are published: both Treasury general licences of 9 October in full; the daily par yield curve; the Federal Reserve's September minutes; the Monetary Authority of Singapore's release calendar; the environment agency's 9 October release and its live pollutant feeds, queried directly at five this morning; Singapore Statutes Online for section 453 as it stands today, and the police release on the last prosecution under it; TSMC's own revenue table and guidance page; the merged pull requests and model cards; four peer-reviewed papers and one preprint, read for sample size and effect size; China's statistics bureau releases; a venue operator's own pages for every Singapore date printed; and the makers' own pages for every price. Where a source refused us — Reddit to all our researchers, several Western wires, a chart publisher — it is said so on the page. Figures we could not verify are held. One correction was proposed and declined, because the archive was right (p4).