

A SPECIAL EDITION OF

MERIDIAN

MiniMax H3 — an ultra deep dive

 An Asus Strix RTX 4090 graphics card, powered on

Source: Wikimedia Commons — an Asus Strix RTX 4090, the class of consumer card this issue's buyer's ladder is built around. Representative image, honestly labelled.

THE BAKE-OFF

Ranked #1 in video editing — by a configuration the open weights don't include

THE COMFYUI STORY

Native support landed day zero. Local inference stops at 768 pixels

THE LONG READ

A Hong Kong stock, two chipmakers, and China's press reaching for its favourite headline

THIS SPECIAL

MiniMax H3 — Contents

p3–4 The Invitation — an open-weight model that beat the closed ones on their own leaderboard four days after release, and the license clause that locks three of the world’s largest markets out of running it locally

p5–6 MiniMax — Company & Lineage — a Shanghai startup that can’t agree with itself on its own birth year, an \$850 million cap table, and the direct line of video models running from Video-01 to H3

p7–8 The Machine — inside a 33-billion-parameter transformer that generates picture and sound in the same pass, and the two pipeline components MiniMax kept for itself

p9–10 The Bake-Off — the real Artificial Analysis scorecard, the competitive set the community actually uses, and why the #1 result can’t be reproduced locally

p11–12 The ComfyUI Story — four new nodes, six workflow templates split three ways local and three ways API, and the joint audio-video mechanic that makes it new

p13–14 The Gear — a buy-once-cry-once ladder from a \$480 entry card to a \$13,000 workstation GPU, priced into a genuine

2026 memory shortage

p15–16 The Community Map — verified where verifiable: real Civitai workflows and GitHub issues, an unconfirmed Reddit, and no official Discord found at all

p17–18 Prompt Craft & Failure Modes — the official shot-list formula, the documented bugs, and what silently drops from a prompt instead of being refused

p19–20 Licensing & The Economics — the excluded countries, the Hollywood lawsuit behind them, and an honest local-vs-API cost worked out in real numbers

p21–22 Your First Ninety Days — a staged programme built entirely from workflows, checkpoints and prices that already exist, ending in one real financial decision

p23–24 The Sceptic — the configuration nobody can download, the open bug tracker, and how much of the “community is thrilled” story survives its own sourcing

p25–26 The Long Read — “The Kill Line”

FROM THE DESK

"Given enough eyeballs, all bugs are shallow."

— Eric S. Raymond, *The Cathedral and the Bazaar* (1999)¹

Coined for open-source software, it describes this launch almost too well: a device-mismatch bug in H3’s own VAE decoder was filed, diagnosed and patched by a stranger on GitHub within days of release ([p12](#)), and a community quantization workaround for a tensor-width quirk nobody at MiniMax had to explain shipped before this issue closed ([p12](#)). Whether “open weights” earns the fuller promise of that sentence — training data, source code, an unrestricted license — is a harder question this issue keeps returning to ([p19](#), [p23](#)).

MINIMAX H3 — AT A GLANCE

METRIC	FIGURE
Parameters, H3-Omni-Transformer	33.1B
API announced	31 Jul 2026
Open weights published	3 Aug 2026
ComfyUI native support	Day zero, PR #15224
Local (open-weight) output cap	768px short edge
Full 2K output	API only
Artificial Analysis rank, Video Editing	#1
License excludes	US, EU, UK, South Korea

Figures as researched for this edition, Aug 2026; see the sections named for full sourcing.²

WHY THIS ISSUE EXISTS

A reader flagged MiniMax H3 in chat four days after its open-weight release — already running it inside ComfyUI, already asking for the full treatment. Unlike the daily book, this special runs no news desks, no Diary, no Rabbit Hole: all twenty-seven pages are eleven sections on one model, built around what could actually be verified in the first week of its life.³

SOURCES ¹ Eric S. Raymond, *The Cathedral and the Bazaar*, 1999 (“Linus’s Law”). ² Figures per MiniMax H3 GitHub README, Reuters/SCMP launch coverage, Comfy-Org/ComfyUI PR #15224, Artificial Analysis Video Arena, and the MiniMax H3 Community License — full sourcing in each named section. ³ Reader note logged via chat and the manual special-edition trigger, 7 Aug 2026.

MINIMAX H3 — WEEK ONE

Already No. 1, Four Days In

An open-weight model just beat the closed ones on their own leaderboard — and its license locks the US, the EU, the UK and South Korea out of running it at home.

MiniMax put MiniMax H3 into the world twice in four days. On July 31, 2026, the model went live closed and API-only, reachable through MiniMax's Open Platform and the consumer Hailuo app, downloadable nowhere¹. Three days later, on August 3, MiniMax uploaded the weights to Hugging Face, governed by a new "MiniMax H3 Community License Agreement" dated effective August 2². By the time this issue went to press, that second date was four days old, and the first was seven.

What shipped is not, strictly, a video model. MiniMax built H3 around a single, dense, 33-billion-parameter "H3-Omni-Transformer" that takes text, images, video and audio in and hands the same four modalities back, in one forward pass rather than four bolted-together tools³. Feed it up to nine reference images, three video clips and three audio clips at once — twelve files, maximum, per prompt³ — and it can return as much as 15 seconds of motion, at resolutions up to 2K, with stereo audio generated alongside the picture instead of dubbed on after³. MiniMax's earlier video line — Video-01, S2V-01, Hailuo 02, Hailuo 2.3 — shipped silent or near-silent, sound handled elsewhere. H3 folds it in.

The open-weighting mattered immediately to a specific audience. The same day the weights landed, Comfy-Org shipped a repackaged mirror and native ComfyUI support, merged as pull request #15224⁴. Anyone running the open-source node graph could load H3 locally on release day, not months later once someone got around to it. That is the reader this issue is written for: someone who dropped an H3

node into a graph last week and wants to know, precisely, what they are now running.

The scale of the moment is easy to understate. At the closed launch, MiniMax's own benchmark citations placed H3 behind Google's Gemini "Omni Flash" in text-to-video, and behind both Gemini and ByteDance's Seedance 2.0 in image-to-video — but first, ahead of every rival tested, in video editing¹. When the weights went public three days later, MiniMax carried that No. 1 ranking into open release with it. So far as MERIDIAN can determine, that makes H3 the first openly-weighted video model to top a category on a major generation leaderboard, rather than trailing the closed frontier by a release cycle or two. MERIDIAN walks the full scorecard, category by category, on [p9](#).

The license carries a catch attached to that same speed. MiniMax's Community License Agreement excludes the United States, the European Union, the United Kingdom and the Republic of Korea from running, modifying or distributing the outputs of the locally-run open weights²; the hosted API stays available everywhere regardless. The South China Morning Post ties the exclusion to active US litigation — Disney, Universal and Warner Bros. Discovery are suing MiniMax and Hailuo AI, alleging its models can reproduce copyrighted characters without permission, and a US court reportedly denied MiniMax's motion to dismiss in May 2026⁵. The full license text, the pricing tiers and what the exclusion means for a working ComfyUI user sit on [p19](#).

33.1B

PARAMETERS IN THE H3-OMNI-TRANSFORMER

SOURCES 1 Reuters (Eduardo Baptista), MiniMax H3 launch coverage, 31 Jul 2026 — capabilities and launch benchmark rankings. 2 MiniMax H3 Community License Agreement, effective 2 Aug 2026, published via Hugging Face, MiniMaxAI/MiniMax-H3, 3 Aug 2026. 3 GitHub, MiniMax-AI/MiniMax-H3 README — architecture, modality limits, output specs. 4 GitHub, Comfy-Org/MiniMax-H3, pull request #15224, 3 Aug 2026. 5 South China Morning Post, scmp.com/tech/tech-trends/article/3362951 — license exclusion and litigation context.

READING THIS ISSUE

Twenty-Four Pages For A Week-Old Model

Every section that follows is built to survive being read a year from now, even though half its numbers were only confirmed this week.

A four-day-old open-weight release rarely earns a magazine's full attention. MERIDIAN makes an exception because three things are true at once: H3 is a newly-public company's flagship product¹, its license carries an exclusion tied to active US copyright litigation covering three of the world's largest markets², and it already has native, day-zero support inside ComfyUI³. Every fact in this issue is between four and seven days old, counted from the July 31 API announcement to the August 3 weight release⁴. Some of what follows will be out of date before this issue is off the press. MERIDIAN says so wherever true, rather than writing around it.

The pages that follow are ordered around the questions a working artist actually asks. Who made this, and what they shipped before it — MiniMax's funding, its January 2026 IPO, and the line from Video-01 to H3 — sits on [p5](#). What is inside the model, its architecture and limits, hedged against

MiniMax's own numbers, runs on [p7](#). How H3 compares against the field on real benchmarks, not launch-day claims, is the scorecard on [p9](#). The ComfyUI integration and workflow are on [p11](#); the hardware ladder beneath — from a rented GPU to a card that runs the full weights at home — follows on [p13](#).

The back half turns skeptical on purpose. [p15](#) and [p17](#) map what the community is doing with H3, and where prompts break. [p19](#) works through the license and pricing in full. [p21](#) is a plan for the next ninety days, not the next ninety minutes. [p23](#) is reserved for the counter-argument: reasons this release may matter less than it looks. [p25](#) runs long — an Undercurrent-style read of how MiniMax's domestic audience, in Chinese-language forums, is receiving H3, a conversation most English-language coverage hasn't touched.

Every fact in this issue is between four and seven days old.

FROM THE DESK

This issue exists because a reader wrote in four days ago: they'd just put an H3 node into a ComfyUI graph for the first time and wanted the full treatment before the story cooled. MERIDIAN does not usually turn one request into twenty-four pages. This one earned it, not because the model is finished being understood, but because so little of it is finished yet — and a reader building with it now deserves to know exactly what's confirmed, what's single-source, and what nobody has verified at all.

THE SCEPTIC

Treat this issue's own confidence with the same suspicion it asks of MiniMax's launch claims. A book assembled four to seven days after release is, by construction, working from early accounts, launch-window benchmarks and search-engine synthesis rather than months of settled use — the same caution this issue applies to H3's "#1" headline on [p9](#) applies to MERIDIAN's own reporting here. Where a claim rests on one source, this issue says so; where the community's reaction couldn't be verified at all, it says that too, rather than filling the gap with confident-sounding prose.

IN BRIEF — THE WEEK IN ONE TABLE

31 Jul H3 announced, API-only, via MiniMax's Open Platform and the Hailuo app.

3 Aug Open weights published to Hugging Face under the new Community License.

3 Aug ComfyUI native support merged same-day as PR #15224.

4–7 Aug Community quants, a Turbo LoRA, and the first documented bugs all land.

SOURCES 1 MiniMax IPO, HKEX 0100.HK, debut 9 Jan 2026 — company funding history, search-synthesis. 2 South China Morning Post, scmp.com/tech/tech-trends/article/3362951 — license exclusion and litigation context. 3 GitHub, Comfy-Org/MiniMax-H3, pull request #15224, 3 Aug 2026. 4 Reuters (Eduardo Baptista), 31 Jul 2026; MiniMax H3 Community License Agreement / Hugging Face, MiniMaxAI/MiniMax-H3, 3 Aug 2026.

SHANGHAI, EST. 2021 — OR 2022

The Company Behind the Conch

MiniMax raised roughly \$850 million before it ever rang a bell in Hong Kong — and even its own paperwork can't settle on a birth year.

MiniMax's own record does not agree with itself. Corporate registries show Shanghai Jizhi Technology and a Beijing sibling incorporated in November 2021; a separate filing, Shanghai Xiyu Technology, dates to January 2023; and the company's own messaging claims "founded in early 2022."¹ MERIDIAN prints the conflict rather than pick whichever date sounds tidiest.

What isn't in dispute is who built it and where. The company is headquartered in Shanghai. CEO Yan Junjie spent six years at SenseTime, one of China's earliest and largest computer-vision AI firms, rising to deputy head of its research institute before leaving to co-found MiniMax with Yang Bin and Zhou Yucong.²

The money arrived in stages, totalling roughly \$850 million before the company touched a public market. A Series A in June 2023, led by Tencent, brought in \$250 million. A larger round in March 2024, led by Alibaba, added \$600 million at a \$2.5 billion valuation — at the time, the largest pre-IPO raise in Chinese generative AI. An extended Series B in July 2025 added \$300 million more, pushing the valuation to \$4 billion, with Shanghai State-owned Venture Capital Group among the participants — state capital sitting alongside the private money.³

MiniMax listed on the Hong Kong Stock Exchange on 9 January 2026 under ticker 0100.HK, priced at HK\$165 a share. It jumped roughly 109% on debut, closing near HK\$345 and pushing market capitalisation to about HK\$106.7 billion — on the order of \$13.7 billion.⁴

Inside China's AI industry, MiniMax gets filed under the so-called "AI Tigers" — Zhipu (Z.ai), Moonshot AI, Baichuan, StepFun and 01.AI — distinct from ByteDance, Alibaba or DeepSeek, where frontier AI is one division inside a larger balance sheet.⁵ Outside analysts cast Zhipu as the state-aligned "national champion" and MiniMax as closer to "the Character AI of China," consumer-companionship-built, reportedly first of the group to post self-sustaining unit economics.⁵

The video brand and the company are not the same thing. Hailuo — 海螺, literally "conch" or "sea snail" — is MiniMax's consumer video app and marketing name; "MiniMax H3" is the model's formal designation, and "Hailuo 3.0" is the consumer-facing alias the app displays for it.⁶ Easy to conflate with MiniMax's other consumer products: Talkie (international, June 2023) and Xingye/ 星野 (domestic, September 2023) are AI-companion chat apps, not video tools — and by at least one account, historically staffed more heavily than the Hailuo video team.⁶

+109%

MINIMAX'S SHARE-PRICE MOVE ON ITS HKEX DEBUT, 9 JANUARY 2026 — FROM AN HK\$165 PRICING TO A CLOSE NEAR HK\$345

ZHIPU / Z.AI **MOONSHOT AI** **BAICHUAN** **STEPFUN** **01.AI** — the cohort MiniMax is grouped with; full model breakdown continues on [p6](#), spec detail on H3 itself on [p7](#).

IN BRIEF — THE REST OF THE "AI TIGERS"

Zhipu / Z.ai the state-aligned "national champion," enterprise and foundation-model contracts.

Moonshot AI maker of Kimi; reportedly powers third-party coding tools including Cursor's Composer 2.

Baichuan, StepFun, 01.AI the remaining names in the cohort local coverage nicknamed the "Six Tigers."

SOURCES ¹ Corporate registry filings for Shanghai/Beijing Jizhi Technology (Nov 2021) and Shanghai Xiyu Technology (Jan 2023), cross-checked against MiniMax's own "founded early 2022" public messaging — exact date unresolved. ² Yan Junjie biography and SenseTime tenure, company and press profiles. ³ MiniMax funding-round reporting: Series A (Tencent, Jun 2023), Series B (Alibaba, Mar 2024), extended Series B (Shanghai STVC Group et al., Jul 2025). ⁴ HKEX listing filings and market reporting on the 0100.HK debut, 9 Jan 2026. ⁵ Analyst and trade-press characterisation of China's "AI Tiger" cohort and its positioning relative to ByteDance/Alibaba/DeepSeek. ⁶ MiniMax brand materials on the Hailuo app naming; kr-asia.com reporting on Talkie/Xingye staffing relative to the Hailuo team.

THE LINEAGE

Four Names Stand Between Video-01 and H3

A separate text-model line — abab, then M1, M2, M3 — ran alongside it inside the same company, evidence of the same underlying research bet, not the model that ships this month.

H3's ancestry runs through video, not text. Video-01, MiniMax's first AI-native video model, shipped August 2024: 720p, 25fps, six-second clips, cinematic camera control.¹ That's the root of everything that follows.

S2V-01 followed in January 2025 — "Subject-to-Video" — keeping a character's face consistent from one reference image across camera moves.² Hailuo 02 (June 2025) reached 1080p and ten-second clips on a new "Noise-aware Compute Redistribution" architecture, ~3x the parameters and 4x the training

data of its predecessor.³ Hailuo 2.3 (October 2025) sharpened expression, motion and micro-expressions; its release title introduced "Media Agent," a hint at where the line was headed.⁴

H3 arrived in two steps: an API-only release 31 July 2026 (clips to 15s at up to 2K, native stereo audio, unified multimodal input, footage editing and motion transfer)⁵, then open weights on Hugging Face three days later under a Community License, with same-day native ComfyUI support.⁵ Full specs on [p7](#); licence on [p19](#).

MODEL	DATE	WHAT IT ADDED
Video-01	Aug 2024	First AI-native video model: 720p, 25fps, clips to 6s, cinematic camera control
S2V-01	Jan 2025	"Subject-to-Video" — locked character identity from one reference image across camera moves
Hailuo 02	Jun 2025	1080p, clips to 10s; new NCR architecture, ~3x params / 4x training data of prior gen
Hailuo 2.3	Oct 2025	Sharper expression, physical motion, micro-expressions; introduces "Media Agent" framing
MiniMax H3	Jul-Aug 2026	Clips to 15s, up to 2K, native stereo audio, unified multimodal input, footage editing & motion transfer

A parallel text-model line existed inside the same company — context, not H3's ancestry. abab gave way to MiniMax-Text-01 (Jan 2025), then M1 (Jun 2025, the first open-weight large-scale hybrid-attention reasoning model), then M2 (Oct 2025) and M3 (mid-2026).⁶ A different lineage, same open-weighting habit — video produced H3.

That's this magazine's read: most labs open-weight a model on the way to going public, or instead of it.

MiniMax reversed the order — raise from public investors first, on 9 January 2026, then hand a serious piece of its core technology to anyone with a GPU seven months later. Its stated rationale is community support and hardware compatibility, framing closed video-gen as slower to iterate than closed text models.⁸ Confidence, or a hardware play — too new to call. Chinese-platform dynamics on [p25](#).

UNVERIFIED

One unconfirmed report claims CEO Yan Junjie pledged to forgo salary until AGI arrives and allocate 5% of his personal equity — split between an employee pool and an open-source fund — timed to H3's open-weighting.⁷ Not traceable to a top-tier outlet; flagged as exactly that. A separate, also single-sourced account has 0100.HK shares rising over 10% intraday around the 3 August release; treat the same way.

A publicly traded company gave away its flagship model's weights seven months after its own IPO — that ordering runs backwards from how this industry usually does it.

SOURCES 1 MiniMax Video-01 launch materials and press coverage, Aug 2024. 2 MiniMax S2V-01 release notes, Jan 2025. 3 MiniMax Hailuo 02 technical announcement on the NCR architecture, Jun 2025. 4 MiniMax Hailuo 2.3 release page, Oct 2025. 5 Reuters (Eduardo Baptista byline) and SCMP reporting on the 31 Jul 2026 H3 API announcement; Hugging Face MiniMaxAI/MiniMax-H3 model card and Comfy-Org PR #15224 on the 3 Aug 2026 open-weight release. 6 MiniMax model-release history: abab 6.5/6.5s (2024), MiniMax-Text-01 (Jan 2025), M1 (Jun 2025), M2 (Oct 2025), M3 (2026) — corroboration incomplete for M3, treated as provisional. 7 Single, unconfirmed report on CEO pledge and market reaction — could not be traced to a top-tier outlet. 8 MiniMax's own stated rationale for open-weighting, company communications, search-synthesis.

ARCHITECTURE

The Machine

A single 33-billion-parameter transformer generates picture and sound together — but the layer counts most widely quoted about it turn out to describe a different, borrowed component entirely.

H3 does not stitch together separate picture and voice generators. MiniMax built it as one transformer — the H3-Omni-Transformer, a 33.1-billion-parameter dense, single-stream model¹ — where video, audio and text tokens move through the same stack of layers together, not routed to separate subnetworks or reconciled afterward. About 13 billion of those parameters sit in the AdaLN branches that modulate each diffusion timestep¹; the headline figure MiniMax publishes for the whole model is 33.1B.

MiniMax's GitHub documentation does not publish an exact layer count, hidden dimension or head total for the Omni-Transformer itself¹. A widely-circulated secondary figure — 50 layers, hidden size 5376, 56 heads — describes something else: MERIDIAN traced it to the H3-Encoder, built on Alibaba's Qwen3-VL-32B, whose hidden states the pipeline draws from its 50th layer². That tap point appears conflated with a spec for the main transformer³.


The mechanism that lets one stream carry all three modalities is a position-tracking scheme called 3D

multimodal rotary position embedding, or MM-RoPE². Every transformer needs to know where each piece of information sits relative to the rest; text models do it along a single line. H3 tracks position along three axes simultaneously: time, height and width. In practice, that lets the model treat a video clip as one continuous object rather than a filmstrip reassembled frame by frame — motion and camera movement are built into the representation, not inferred afterward.

Two more components handle compression. A VAE — variational autoencoder — compresses raw data into a compact shorthand cheap enough to reason over, then expands it back to full quality on the way out. H3 runs two, independently. H3-VisualVAE, labelled f16t4d24, compresses video 16x spatially and 4x along time into 24 latent channels — after the transformer's own patchification, the effective spatial compression works out to roughly 32x². H3-AudioVAE does the same for sound: 32kHz audio compressed to latent tokens at a 40Hz rate, channels processed independently².

33.1B

PARAMETERS IN THE H3-OMNI-TRANSFORMER



THE SCEPTIC

MiniMax markets H3 as one unified model, and structurally that framing survives scrutiny for the core transformer. But "dense, single-stream" doesn't mean self-contained end to end: the encoder is a repurposed 32-billion-parameter model built by a different company, and both VAEs are separately trained networks bolted to the transformer's input and output ends. Call it one model at the core with three attached, independently-built components at the edges — a meaningfully different claim than "one model, full stop."

SOURCES ¹ MiniMax H3 GitHub README (github.com/MiniMax-AI/MiniMax-H3), architecture section. ² Same repository, H3-Encoder / MM-RoPE / VAE technical specifications. ³ Secondary aggregation (kingy.ai-derived reporting) citing 50-layer/5376-hidden/56-head figures for H3-Omni-Transformer — uncorroborated against MiniMax's primary documentation.

OUTPUT & LIMITS

What H3 Can Actually Make

Two pipeline components that produced MiniMax's flagship demos and topped a video-editing leaderboard stay behind the API — the open weights cap out lower.

A single H3 prompt has hard ceilings. The Ref2VA variant — MiniMax's reference-conditioned mode for character and product consistency — accepts up to nine reference images, up to three video clips of two to fifteen seconds each (fifteen seconds combined), and up to three audio clips, each needing an accompanying image or video to anchor it. Across all of that, a prompt tops out at twelve files total¹. FL2VA, the text-and-first-frame variant, works to a narrower brief: a single starting frame rather than multi-reference input.

What comes out is bounded just as precisely. Clips run four to fifteen seconds at a fixed 24 frames per second, at resolutions up to 2K with a default shorter edge of 768 pixels — the same tier MiniMax's API prices separately — across six aspect ratios². One hands-on account sharpens the duration figure: a "15-second" clip is reportedly 362 frames, about 15.08 seconds, snapping to a 17-frame grid rather than round numbers³ — a single source, but concrete enough to repeat.

SPEC	VALUE
Duration	4–15 seconds, fixed 24fps
Resolution	up to 2K · default 768px short edge
Aspect ratios	21:9, 16:9, 4:3, 1:1, 3:4, 9:16
Audio	32kHz stereo, generated in the same forward pass

That last line is the genuinely new part. Nearly every other video generator MERIDIAN has covered treats sound as an afterthought: render the picture, then bolt on text-to-speech or sound effects and hope the timing lines up. H3's video and audio latents come out of the same forward pass through the same network² — the model is deciding what a door sounds like at the same moment it decides how the door swings open, not stitching the two together afterward. That is a structural difference from bolted-on TTS, not a marketing distinction — the single biggest technical claim in this release.

None of that is available at full strength in the open weights. MiniMax's own release materials disclose that two pipeline components, H3-Context-IR and

H3-Regenerate-2K, are held back from the open-source release and remain API-only⁴. H3-Regenerate-2K's job is documented well enough to describe: it upsamples the 768p output the local model actually produces into the full 2K resolution advertised for the hosted service. What H3-Context-IR does, MERIDIAN could not confirm beyond that it, too, stays server-side. The practical result is that anyone running H3 locally through ComfyUI (see [p11](#)) is capped at 768px short edge — and the benchmark run that put H3 at the top of a video-editing leaderboard used the full hosted pipeline, both excluded components included, not the configuration anyone can download⁵ (see [p9](#)). The flagship demos and the download are not, quite, the same product.

The model decides what a door sounds like at the same moment it decides how the door swings open.

IN BRIEF — WHAT STAYS BEHIND THE API

H3-Context-IR not released as open weights; function not fully documented, stays server-side.

H3-Regenerate-2K upsamples a finished 768p local render to full 2K — API-only.

SOURCES 1 MiniMax H3 GitHub README, H3-Base-Ref2VA modality limits. 2 Same repository, output specifications (duration, fps, resolution, aspect ratios, audio). 3 Atlas Cloud (atlascloud.ai) hands-on writeup, single-source, on frame-count/17-frame-grid claim. 4 Hugging Face MiniMax H3 model card / GitHub README disclosure that H3-Context-IR and H3-Regenerate-2K are excluded from the open-source release. 5 XenoSpectrum (xenospectrum.com/en/minimax-h3-open-weights-video-editing), on the benchmarked configuration's use of the API-only 2K pipeline versus the 768p-capped open weights.

BENCHMARKS

H3 Takes Three Podium Finishes — One of Them for a Model You Can't Run


Text-to-video and image-to-video results are close and clean; the video-editing crown was won by a hosted pipeline the public weights don't include.

Artificial Analysis published the numbers itself: the benchmarking outfit behind the Video Arena leaderboard stated on X that MiniMax H3 "is #1 in Video Editing and ranks top 3 in both Text to Video and Image to Video."¹ That is an unusually clean sweep for a model that shipped roughly a week before this issue closed.

The Video Arena works by blind pairwise human voting, scored with a Bradley-Terry Elo system across six pools split by task and by whether the clip carries audio.¹ In Video Editing, H3 sits at the top with an Elo of roughly 1,130. In Text-to-Video it places second at 1,239, five points behind Gemini's "Omni Flash" variant at 1,244. In Image-to-Video it places third at 1,187, behind ByteDance's Seedance 2.0 at 1,196 and Gemini Omni Flash at 1,194.¹ None of those margins are wide — this is a model finishing at or near the front of every category it entered.

The Video Editing sample size is one figure MERIDIAN could not pin down: secondary sources put the evaluation pool at either 8,211 or 5,043 samples.² Nobody disputes the ranking — accounts differ only on how many votes it rests on.

CATEGORY	RANK	ELO	WHO'S AHEAD
Video Editing	#1	~1,130	— (leads; sample size disputed)
Text-to-Video	#2	1,239	Gemini "Omni Flash" (1,244)
Image-to-Video	#3	1,187	Seedance 2.0 (1,196), Gemini "Omni Flash" (1,194)



5
POINTS SEPARATING H3
FROM GEMINI "OMNI
FLASH" AT THE TOP OF
TEXT-TO-VIDEO — A NEAR
TIE, NOT A LOSS

Put together, the three results describe a model that is competitive rather than dominant: a genuine top-three finisher in every pool Artificial Analysis measured, and the outright leader in exactly one of them. What that leader result is actually made of is the harder question.

Here is the qualification that matters more than the sample-size dispute: per XenoSpectrum's reporting, the configuration Artificial Analysis actually scored for the #1 Video Editing result ran on MiniMax's hosted pipeline — a 2K-output setup built on two undisclosed components, H3-Context-IR and H3-Regenerate-2K (see [p7](#)).³ The weights available for download are the H3-Base checkpoint, capped at 768p.³ A reader running H3 locally is running a lower-resolution sibling of the model that won.

THE SCEPTIC

The gap between "#1 on the leaderboard" and "what loads in your ComfyUI graph" is the single most important thing to understand before repeating the headline. MiniMax benchmarked its server-side pipeline, not the checkpoint it shipped publicly. That isn't fraud — vendors routinely submit their best hosted configuration to arenas like this — but it means local users should expect competent 768p editing, not the 2K result the ranking implies.

One further caveat, sourced to a single outlet and not independently verified by MERIDIAN: kingy.ai reports that H3's Text-to-Video and Image-to-Video placements lean on Artificial Analysis's internal evaluation methodology rather than the full public vote history, since H3 is too new for stable head-to-head data in those pools.⁴ If accurate, the #2 and #3 finishes are more provisional than the leaderboard's clean numbers suggest.

SOURCES 1 Artificial Analysis (@ArtificialAnlys) on X, and the Artificial Analysis Video Arena leaderboard, Aug 2026 — ranking and Elo figures for Video Editing, Text-to-Video and Image-to-Video, and arena methodology. 2 Sample-size figures for the Video Editing category are reported inconsistently across secondary sources (8,211 vs. 5,043 evaluation samples); MERIDIAN could not reconcile the discrepancy. 3 XenoSpectrum, "MiniMax H3: Open weights, video editing king," [xenospectrum.com/en/minimax-h3-open-weights-video-editing](#), Aug 2026. 4 kingy.ai analysis, Aug 2026 — single-source claim, not independently verified by MERIDIAN.

THE COMPETITIVE SET

The Models the Community Actually Puts Next to It

One rival is scheduled to be retired by its own maker before the year is out; another is the closest thing to genuinely open source H3 has to answer to.

Ask anyone running ComfyUI which models they load next to H3 and the same names come up: Seedance 2.0/2.5, Kling 2.6/3.0, Runway's Gen-4.5, Wan 2.5/2.7, and — mostly as a footnote about its own scheduled shutdown — Sora 2.¹ That list leans on review-blog coverage rather than measured head-to-head testing — a limitation MERIDIAN flags, not settled consensus.¹

Wan gets singled out for a reason unrelated to quality: it ships under Apache 2.0, open source in the sense that governs code, not just weights. H3 carries no such license, and the 2K pipeline behind the Video Editing win never leaves MiniMax's servers.

Where H3 wins, hands-on accounts converge on three things: audio as a first-class modality rather than bolted on — a real differentiator against silent Kling; native 2K/24fps, the highest default resolution among public models at launch; and reference-driven character consistency.² Where it loses: large-scale action motion, where Seedance is reported steadier, with H3 showing texture instability and visual noise under complex movement; lip-sync is mixed, near-perfect on some clips, drifting on others.³

Native audio is the real differentiator against Kling. Large-scale action motion is where it comes apart.

FRAMING MINIMAX'S OWN LAUNCH MATERIALS

H3 is presented as a category leader — "#1 in Video Editing" — with the local, open-weight release framed as the same achievement made available to everyone.⁵

FRAMING THIS ISSUE'S OWN READING

#1 in one category, competitive but not dominant in the other two, and won by a hosted configuration the open weights don't reproduce — real, but narrower than the headline.

IN BRIEF — WHERE EACH RIVAL BEATS IT

Seedance 2.0/2.5 steadier under large-scale action and complex motion.

Gemini "Omni Flash" narrowly ahead on Text-to-Video and Image-to-Video Elo.

Wan 2.5/2.7 genuinely open source (Apache 2.0) where H3 is open weights only.

WHAT MERIDIAN COULD VERIFY

Direct organic Reddit discussion of H3's benchmark standing could not be independently confirmed for this piece — no thread content this research could read and cite. Two Hacker News threads do exist and confirm developer attention: "MiniMax H3: Open weights Omni-modal model" and "MiniMax H3 Day-0 Support in ComfyUI." Their existence is confirmed; their comment content was not accessible to this research.⁴

The-decoder.com ran with "first open model to top a ranking" as its headline — confirmed real, body unreadable.⁵ Two qualifications matter: the #1 finish is Video Editing specific, not an overall crown, and it was won by a configuration the public can't download. kingy.ai's caution: "a single third-party arena run, not a peer-reviewed suite... treat rankings as a prior, not a verdict." [p23](#) takes the sceptic's case further.

SOURCES ¹ Comparison set (Seedance 2.0/2.5, ByteDance; Kling 2.6/3.0; Runway Gen-4.5; Wan 2.5/2.7, Alibaba; Sora 2) aggregated from ComfyUI/video-gen review-blog coverage; corroboration limited, not measured head-to-head testing. ² Hands-on/review coverage on H3's strengths: native audio as first-class modality, native 2K/24fps output, reference-driven character consistency. ³ Hands-on/review coverage on H3's weaknesses: physics and large-scale action motion versus Seedance, texture instability and facial distortion under complex motion, mixed lip-sync reports. ⁴ Hacker News, news.ycombinator.com/item?id=49130750 and news.ycombinator.com/item?id=49155629 — existence confirmed, comment content unreadable; organic Reddit discussion could not be independently verified. ⁵ the-decoder.com headline (body unreadable to MERIDIAN); kingy.ai, Aug 2026 — "a single third-party arena run, not a peer-reviewed suite... treat rankings as a prior, not a verdict."

NODE BY NODE

Four New Nodes, No Fanfare

ComfyUI's day-zero support for H3 landed as PR #15224 — a routine core merge by kijai, requiring nothing more than an update to ComfyUI 0.30 or later.

On 3 August 2026, three days after MiniMax pushed H3 into the wild, a single pull request appeared in Comfy-Org/ComfyUI. Its title had no ceremony in it: "feat: Support MiniMax-H3 (CORE-375)." Author: kijai, the contributor whose name has become shorthand for "this new model works in ComfyUI now."¹ A model with a joint audio channel and a 33.1-billion-parameter transformer landed the way a routine bugfix does — merged, versioned, done.

The mechanics matter more than the ceremony. H3 support ships as core code, not a separate node pack: a new `comfy_extras/nodes_minimax_h3.py` file and

a `comfy/ldm/minimax/` module, added directly to the engine.¹ Anyone running ComfyUI gets it on the next update. The one condition is a version floor: ComfyUI 0.30.0 or later — a figure that turns up consistently across independent update threads but that MERIDIAN could not confirm word-for-word against Comfy-Org's own release notes at press time.²

kijai's diff adds exactly four new node classes — worth learning by their real names, not marketing shorthand, since that's what appears in the node browser:³

NODE	WHAT IT DOES
EmptyMiniMaxH3LatentAV	Initializes the empty joint audio+video latent every generation starts from — the audio channel is stereo, shaped [batch, 32, 2, time].
MiniMaxH3ImageToVideo	Image-conditioned entry point: feeds a first frame, last frame, or first-and-last pair into the packed sequence alongside plain text-to-video.
MiniMaxH3ReferenceToVideo	Reference-conditioned entry point: locks subject, motion, or voice from a reference image, video clip, or audio clip.
MiniMaxH3SigmaShift	Sets the noise (sigma) schedule for the sampling pass — needed because video and audio denoise on different curves inside one sampler.

In practice, the four nodes recombine into three graph shapes sharing one sampler and decode chain: `EmptyMiniMaxH3LatentAV` for plain text-to-video; `MiniMaxH3ImageToVideo` for first/last-frame conditioning; `MiniMaxH3ReferenceToVideo` to lock character, style, camera or voice from a reference.³

The naming shows what's structurally new. Older ComfyUI video workflows hand a finished picture sequence to a separate TTS or lip-sync node; H3 doesn't. Per the PR description, it's a single-stream packed-token diffusion transformer denoising video

and audio latents together, one sampling loop, conditioned on Qwen3-VL-32B hidden states.⁴ Lip sync and sound-effect timing fall out of the two streams sharing every step, not a post-process — which is why `MiniMaxH3SigmaShift` exists as its own node: video and audio don't share a noise schedule.

Decode is a two-VAE handoff: a 3D causal CNN video VAE reconstructs frames, a DAC/BigVGAN-lineage audio VAE reconstructs 32kHz stereo, and ComfyUI muxes both into one file.⁴ No new installer, backend, or sampler type — four node classes and a merge.

H3 doesn't generate video and then bolt on speech — it denoises picture and sound in the same pass, step for step.

SOURCES 1 PR Comfy-Org/ComfyUI #15224, "feat: Support MiniMax-H3 (CORE-375)," author kijai, merged 3 Aug 2026, GitHub. 2 ComfyUI update and release threads, Aug 2026 — version-floor figure corroborated across multiple independent sources, not confirmed against official release notes directly. 3 PR Comfy-Org/ComfyUI #15224 diff, `comfy_extras/nodes_minimax_h3.py`, GitHub. 4 PR Comfy-Org/ComfyUI #15224 description, GitHub.

THE CATCH

768 Pixels, Then You Pay

Six official templates split cleanly down that line — three run entirely on your own GPU, three call MiniMax's hosted API, and only the API path reaches full 2K.

Six workflow templates shipped with the four nodes, in Comfy-Org/workflow_templates — three local, three API-key.¹ Local is open-weight, running on your own card from checkpoints at Comfy-Org/MiniMax-H3; API-key sits under a partner/video/MiniMax menu, calling MiniMax's hosted service on an account key instead. Both cover

text-to-video, image-to-video, reference-to-video. The API-key nodes are reported, unconfirmed by MERIDIAN, as MinimaxHailuo03TextToVideoNode, MinimaxHailuo03FirstLastFrameNode, and MinimaxHailuo03ReferenceNode — Hailuo branding carried over rather than rebadged.¹

768px

MAXIMUM SHORT-EDGE OUTPUT ON LOCAL, OPEN-WEIGHT NODES — FULL 2K RUNS THROUGH THE API ONLY

Where matters: local inference tops out at 768 pixels short edge, 768x1344 max — FL2VA and Ref2VA both distilled from the same H3-Base transformer on p7.² Full 2K, and what the bake-off (p9) judged, comes only from H3-Regenerate-2K, an API-only operation upsampling a finished 768p render server-side — no local route exists.²

GitHub Issue #15251 documents a RuntimeError — tensors on both cuda:0 and cpu — thrown by ViT3DDecoder.forward when ComfyUI partially offloads to system RAM on low-VRAM setups.³ The reporter's patched vae.py has not, as of this writing, been confirmed merged.³

THE ROUGH EDGE

Issue #15251 isn't a stray forum complaint — it's filed against ComfyUI's own repository, with a stack trace and a proposed fix attached. What it doesn't have yet is a merge.³ Readers on the lower end of the buyer's ladder, p13, should expect to meet it.

Quantized weights appeared within a day, from HF repos including molbal, joeygambino, Abiray, leejet, vantagewithai, and lilcheaty — GGUF Q2-Q5, pruned INT4/INT8, NVFP4.⁴ None use K-quants: H3's 2688-wide transformer fails the divisible-by-256 test, pushing the community to classic Q4_0/Q5_1. On a 32GB card, Q5_1 holds ~22GB resident, ~ten minutes for a 243-frame clip at 20 steps.⁴

Speed arrived two days later: larryvrh's MiniMax-H3-Turbo-Lora (5 Aug 2026) cuts sampling from ~20 steps to 4 — a claimed 5x speedup — packaged for ComfyUI by drbaph.⁵ Stock sampling clips at 4 steps since video and audio need different shift values (12 vs 3); a community dual-clock Euler sampler fixes most of it. Comfort zone: 6-8 steps, audio the weakest part.⁵

THE SPEED PATH

Four steps at five-times speed sounds like the number worth chasing. Six to eight steps, with the dual-clock sampler, is the number that currently looks like a finished frame.

IN BRIEF — THE FOUR NODES

EmptyMiniMaxH3LatentAV the empty joint latent.

MiniMaxH3ImageToVideo first/last-frame conditioning.

MiniMaxH3ReferenceToVideo locks subject/motion/voice.

MiniMaxH3SigmaShift sets the noise schedule.

SOURCES 1 Comfy-Org/workflow_templates repository listing, GitHub — six templates, three local plus three API-key, node names not primary-confirmed. 2 Comfy-Org/MiniMax-H3, Hugging Face — checkpoint listing, local 768px short-edge cap. 3 GitHub Issue #15251, Comfy-Org/ComfyUI — device-mismatch RuntimeError report and proposed patch. 4 Community GGUF/NVFP4 quantization repositories, Hugging Face (molbal, joeygambino, Abiray, leejet, vantagewithai, lilcheaty), Aug 2026. 5 larryvrh/MiniMax-H3-Turbo-Lora and drbaph/MiniMax-H3-Turbo-Lora-ComfyUI, Hugging Face, 5 Aug 2026.

THE GEAR — BUYER'S LADDER

Buy Once, Cry Twice

GDDR7 spot prices have jumped as much as 130 percent quarter over quarter through 2026 — industry estimates put relief no earlier than 2028.

GDDR7 has spiked as much as 130 percent quarter over quarter through 2026, DRAM close behind¹. Intel and Silicon Motion don't expect relief before ~2028². Every card on this ladder sits above its original MSRP — the "cry" half of buy-once-cry-once is doing more work than it has in years.

The ladder follows H3's quantisation tiers more than raw horsepower — [p11](#) covers the quant setup in full. VRAM headroom decides which build loads: ENTRY lands on a 16GB-class GGUF quant, roughly Q4_0, moving up through ComfyUI's pruned INT8 path to full NVFP4.

An RTX 4060 Ti 16GB runs ~\$480 street³ — inflated, still the cheapest door in. The RTX 4070 Ti Super

16GB launched at \$799, now \$1,179³.

High-end buyers hit a fork. NVIDIA discontinued the RTX 4090 24GB; used units sell for \$2,150–\$2,350³. The RTX 5080 16GB launched at \$999, now ~\$1,600³ — less VRAM, still a worse deal than a year ago.

The RTX 5090 32GB is the practical ceiling: FE MSRP \$1,999.99, street ~\$4,300+, AIBs \$4,600–\$5,000+³ — enough for Q5_1 GGUF (~22GB) or pruned INT8/NVFP4. Past that it's not consumer gear: the RTX PRO 6000 Blackwell 96GB lists at \$13,250, retail seen \$8,000–\$9,400⁴. MiniMax's SGLang team offers a cheaper path: two RTX 5090s, layerwise offload, 26.3GiB/GPU peak⁵.

130%

REPORTED QUARTER-OVER-QUARTER GDDR7 PRICE SPIKES IN 2026 — DIRECTLY BEHIND EVERY PRICE ON THIS PAGE

TIER	CARD	STREET PRICE	VRAM	WHAT IT RUNS
Entry	RTX 4060 Ti 16GB	~\$480 ³	16GB	~Q4_0 GGUF quant
Mid	RTX 4070 Ti Super 16GB	MSRP \$799 → ~\$1,179 ³	16GB	16GB-class quant, more headroom
High-end	RTX 4090 24GB (discontinued)	used \$2,150–\$2,350 ³	24GB	official pruned INT8 path
High-end	RTX 5080 16GB	MSRP \$999 → ~\$1,600 ³	16GB	16GB-class GGUF quant
Over the top	RTX 5090 32GB	FE MSRP \$1,999.99 → street ~\$4,300+ (AIBs \$4,600–\$5,000+) ³	32GB	Q5_1 GGUF (~22GB) or pruned INT8/NVFP4
Workstation	RTX PRO 6000 Blackwell 96GB	NVIDIA list \$13,250 (retail listings seen \$8,000–\$9,400) ⁴	96GB	full pruned pipeline
Dual-GPU	2× RTX 5090 (SGLang cookbook)	~2× RTX 5090 street (~\$8,600+)	2×32GB	layerwise offload, 26.3GiB/GPU peak ⁵

THE SCEPTIC

minimaxh3.run and a community account (kombitz) warn MiniMax hasn't released an official GGUF build, and to watch for fakes online⁶ — against the common assumption that repos like joeygambino, molbal, Abiray and DmitryDB are semi-official. They're community conversions, unverified against any published checksum.

IN BRIEF — READING THE LADDER

VRAM, not clock speed decides which quant a card can load — see [p11](#).

Every price here is a moving target — the 2026 memory shortage, not H3, set them.

SOURCES 1 DRAM/GDDR7 shortage, 60–130%+ QoQ memory cost spikes, per TechRadar and industry trackers shattered.io and tech-insider.org, 2026. 2 No relief expected before ~2028, per Intel and Silicon Motion industry estimates, 2026. 3 Street/retail pricing for RTX 4060 Ti, 4070 Ti Super, 4090 (used) and 5090, per retail listings including Amazon and WCCFTech reporting, Aug 2026 — moving targets. 4 RTX PRO 6000 Blackwell 96GB pricing: NVIDIA's list price vs conflicting retail listings, Aug 2026. 5 Dual RTX 5090 layerwise-offload figures, MiniMax H3 SGLang cookbook documentation. 6 GGUF authenticity warning, minimaxh3.run and a community account (kombitz), Aug 2026, set against community GGUF repositories on Hugging Face (joeygambino, molbal, Abiray, DmitryDB).

THE GEAR — SPEED & RENTAL

How Long, and Whether to Rent Instead

The same class of clip is reported at 175 seconds on one RTX 5090 and five and a half minutes on another — the settings, not the silicon, explain the gap.

Every generation-speed number in circulation for H3 was measured under different conditions, and MERIDIAN won't pretend otherwise by building a single comparison table out of them. Resolution, step count and precision all move independently across the reports below — treat what follows as data points, not a race.

An RTX 4090 running the non-pruned INT8 path completes a default text-to-video generation in around 100 seconds with a warm cache, though users report slowdowns when the job spills into system RAM¹. One account puts an RTX 5090 running the NVFP4/pruned path, a roughly 31.7GB footprint, at 175 seconds for a 10-second clip at 864×480² — a single source, and MERIDIAN could not corroborate it independently. A separate report has an RTX 5090 taking closer to five and a half minutes for a 5-second clip at a comparable resolution¹, which is the gap the dek above is pointing at. Two RTX 5090s in tandem, per MiniMax's own SGLang cookbook, need 559.67 seconds for a lossless 50-step, 5-second clip at 1344×768³. At the bottom of the ladder, RTX 3060 12GB owners report anywhere from 9 to 25 minutes for a short 480p clip, and the two sources disagree by nearly a factor of three⁴.

Renting is the honest alternative to any of this. RunPod's Secure Cloud lists H100 PCIe instances at

\$2.89 an hour and H200 at \$4.39⁵, on-demand, no bidding. Vast.ai's marketplace runs wider: verified datacenter H100s go for \$1.50 to \$1.87 an hour, general on-demand closer to \$2.21, and spot instances have been seen as low as \$0.34 an hour — though a spot instance can be pre-empted mid-render⁵. Against either, MiniMax's own API prices 2K output at \$0.13 per second of generated video⁶; full pricing detail, including the tiers below 2K, runs in [p19](#). MERIDIAN could not find a verified like-for-like comparison of rental cost versus API cost per finished clip — anyone quoting one online is doing their own math, not citing a benchmark.

None of this hardware is actually new. The RTX 5090 launched in January 2025, over a year before this issue went to press, and its price move from \$1,999.99 to \$4,300-plus is a function of the memory shortage, not fresh demand for the card itself. The genuine advance behind H3 running on consumer hardware at all shipped in software, not silicon: a pruning and NVFP4 quantisation technique that strips roughly 40 percent of the model's modulation-weight parameters and replaces them with a lookup-table substitution, released alongside H3's day-zero ComfyUI support⁷. That's the cutting edge of the last twelve months in this section — not a new card, a new way of not needing one.

40%

SHARE OF MODULATION-WEIGHT PARAMETERS PRUNED FROM H3'S NVFP4 PATH — THE REAL BREAKTHROUGH BEHIND CONSUMER-CARD SUPPORT

The real cutting edge of the last twelve months isn't a new card — it's a new way of not needing one.

FROM THE DESK

Renting by the hour only wins if you actually stop the clock. A generation that runs long on a rented H100 costs real money every extra minute — the flat per-second API price doesn't have that failure mode.

SOURCES 1 RTX 4090 ~100s default T2V and a separate RTX 5090 ~5.5min report, community ComfyUI benchmark accounts, Aug 2026, figures differ. 2 RTX 5090 175s NVFP4/pruned figure, per ai-muninn, single account, unverified independently. 3 Dual RTX 5090 559.67s (~9.3min) figure, MiniMax H3 SGLang cookbook documentation. 4 RTX 3060 12GB 9–25 minute range, per kingy.ai and a separate community account; figures conflict. 5 RunPod Secure Cloud and Vast.ai marketplace pricing, Aug 2026. 6 MiniMax API pricing, \$0.13/s for 2K output, official API pricing page, see also [p19](#). 7 Pruning/NVFP4 quantisation technique (~40% modulation-weight parameters), per H3 technical release notes and ComfyUI day-zero support documentation, Aug 2026.

FIELD REPORT

Verified Where Verifiable

Civitai lists six named community releases within four days of launch; what happened on Reddit could not be confirmed at all.

MiniMax announced H3's API on 31 July 2026 and open-weighted the model three days later under a new MiniMax H3 Community License, with the fuller lineage covered at [p5](#). ComfyUI shipped day-zero support that same afternoon: pull request #15224, "feat: Support MiniMax-H3 (CORE-375)," merged into what is now Comfy-Org/ComfyUI — the repository moved from comfyanonymous/ComfyUI during this window, and old links now redirect¹. Four new nodes arrived with it — EmptyMiniMaxH3LatentAV, MiniMaxH3ImageToVideo, MiniMaxH3ReferenceToVideo, MiniMaxH3SigmaShift — wired to the single-stream, joint video-audio architecture detailed at [p7](#). A same-day merge like that isn't automatic; it means ComfyUI's own maintainers treated H3 as priority work rather than something the community would patch in later, and it set the clock running on everything below. This page is written not quite a week after that first date, and it says so plainly: some of what follows is confirmed with names and dates, one claim rests on a single unverified account, and one entire platform's reaction could not be checked at all.

Civitai is the clearest signal here. Two articles carrying the platform's own byline went up within days of launch: "MiniMax H3: Open-Weights Video, Live to Generate Today" and "MiniMax H3: 2K, Stereo Sound, 15s — Full ComfyUI API Test"². Workflow builders moved just as fast behind them — "MiniMax-H3 Multishot," "MiniMax H3 All-in-One Reference Automatic Prompt Workflow," "Minimax H3 T2V/I2V/REF2V Advanced Filmmaking Workflow," "MiniMax-H3 Pseudo Image Generation Workflow," and a hybrid build, "PlagueKind — Minimax H3 — LTX2.3 Workflow — I2V," that treats H3 as one stage inside a multi-model pipeline rather than a self-contained tool³. A speed LoRA followed on 5 August:

creator larryvrh's Turbo LoRA, which claims to cut generation from roughly twenty sampling steps to four⁴. Five named workflows plus one LoRA, dated, inside four days, is the strongest evidence in this whole map that people are opening ComfyUI and building — not just reading about the release.

6

NAMED CIVITAI WORKFLOWS & TOOLS PUBLISHED WITHIN 4 DAYS OF LAUNCH

GitHub adds texture Civitai can't: disagreement and bug reports, not only enthusiasm. PR #15224 itself drew 28 comments, mostly automated CodeRabbit review plus human sign-off from alexisrolland and comfyanonymous¹ — ordinary traffic for a same-day model merge. The issues filed after are more telling. #15251 reports device-mismatch RuntimeErrors in the H3 VAE during partial CPU offload — exactly the failure mode of the lower-VRAM setups discussed at [p13](#), with the ComfyUI-side technical detail behind it at [p11](#)⁵. #15318, filed roughly a day before this was written, asks for multiple reference-frame support, since only first- and last-frame anchoring works today⁵. A feature request a full day after launch, not a bug report on day one, is a small but real sign that people are past installation and into wanting more from the tool. Two third-party repositories round out the picture: HM-RunningHub/ComfyUI_RH_MinMaxH3 and ethanfel/ComfyUI-MiniMax-H3-Guide, a helper node for prompt preparation⁵.

Five named workflows and one speed LoRA, dated, inside four days — the strongest evidence in this map that people are building, not just reading about the release.

SOURCES 1 Comfy-Org/ComfyUI (formerly comfyanonymous/ComfyUI), pull request #15224, "feat: Support MiniMax-H3 (CORE-375)," merged 3 Aug 2026, GitHub. 2 civitai.com/articles/33334, "MiniMax H3: Open-Weights Video, Live to Generate Today"; civitai.com/articles/33333, "MiniMax H3: 2K, Stereo Sound, 15s — Full ComfyUI API Test," Civitai. 3 Civitai model listings civitai.com/models/2833322, /2663838, /2834586, /2832542, /2834514, workflow titles as cited, Civitai. 4 comfyui-wiki.com/en/news/2026-08-06-minimax-h3-turbo-lora, on creator larryvrh's Turbo LoRA, released 5 Aug 2026. 5 Comfy-Org/ComfyUI issues #15251 and #15318, GitHub; HM-RunningHub/ComfyUI_RH_MinMaxH3 and ethanfel/ComfyUI-MiniMax-H3-Guide repositories, GitHub.

THE GAPS

No Confirmed Thread. No Official Server.

Six tutorials reached YouTube before the model turned a week old — proof the fastest layer of this community is the one furthest from the source.

Reddit is the largest gap in this map, and Meridian is printing that plainly instead of filling it in. Direct verification of any H3 thread, comment, or score was not possible with the research tools available this week¹. The only available signal is a single secondary source: a Patreon post from creator "LumosAI" claiming "Reddit is already flooded" with H3 discussion, paired against LTX 2 as an upscaling step¹. That is one unverified claim, not a finding — no dedicated r/MiniMax subreddit turned up either. Reddit conversation may well exist at real volume, this fast, for a model this capable — but this issue cannot vouch for a single thread of it.

Discord is thinner still. No dedicated official MiniMax or Hailuo server appeared anywhere in this research². What exists sits one step removed: Banodoco, a genuine, long-running open-source AI-video Discord that has covered models like Hunyuan, with no H3-specific activity there cited; and a third-party "Comfy Deploy" server carrying more than 5,800 members, a deployment tool's community rather than core ComfyUI's own². WanGP's Discord lists H3 as one of several supported models, not a dedicated home for it³. None of it is confirmed H3-specific traffic, and this page will not round that up to more than it is.

YouTube is the opposite case: fast, public, and trivial to confirm. Within days of the open-weight release, tutorials were already stacked up — "MiniMax H3 — Full ComfyUI Workflow Tutorial (Local+Free Cloud)," "MiniMax H3 is Insane! 2K AI Video Model — Full Tutorial & Workflow," and four more covering music-video builds, speed tricks, and low-VRAM setups⁴. At least one creator had already packaged a tutorial-and-workflow bundle for sale on Gumroad⁴ — the

monetization layer moves at the same speed as the tutorial layer, faster than either Reddit or Discord could be confirmed to move.

THE SCEPTIC

Five workflows and a LoRA are evidence of builders, not evidence of a movement. Every platform confirmed active on this spread — Civitai, GitHub, YouTube — is one where posting costs nothing and shows up instantly. The harder tests of a community, a subreddit vote count or a Discord member graph measured over weeks rather than days, are exactly the ones four days can't yet answer.

Two conversations sit outside this page on purpose — the Chinese-platform scene at [p25](#), and the license exclusion debate at [p19](#)/[p23](#). The plain read: Civitai and GitHub show real activity with names and dates; YouTube shows a tutorial economy already trading; Reddit's reaction is one unverified claim; an official Discord doesn't appear to exist. Not hype, not dismissal — early, and where checkable, real.

PLATFORM	STATUS	EVIDENCE
Civitai	Confirmed, active	2 articles; 6 named workflows & LoRA in 4 days
GitHub	Confirmed, ongoing	PR #15224 (28 comments); issues #15251, #15318
YouTube	Confirmed, active	6+ tutorials within days; Gumroad bundle live
Reddit	Unconfirmed	One secondary claim only; no thread verified
Discord — official	Not found	No dedicated MiniMax/Hailuo server located
Discord — adjacent	Confirmed, not H3-specific	Banodoco; Comfy Deploy; WanGP

IN BRIEF — TWO CONVERSATIONS THIS PAGE SKIPS

The Chinese-platform scene — Zhihu, Bilibili, V2EX — runs on its own terms at [p25](#).

The license exclusion debate gets full treatment at [p19](#) and [p23](#).

SOURCES 1 Patreon, "LumosAI," post claiming active Reddit discussion pairing MiniMax H3 with LTX 2 (unverified, single-source); no dedicated r/MiniMax subreddit located. 2 disboard.org, Comfy Deploy Discord server listing (5,800+ members, comfyui tag); Banodoco, github.com/banodoco and x.com/banodoco. 3 WanGP project Discord server listing, third-party ComfyUI-adjacent video tool. 4 YouTube video titles as cited, watch IDs lcmzsy5CdYc, OcsrFO_RHxs, a2IJ-sfi5Sg, 3LrjnNgU9uc, RR3rtc9jdcl, VNI9QTSHxEU; Gumroad tutorial-and-workflow bundle listing.

READING THE MANUAL

H3 Wants a Screenplay, Not a Caption

Nine reference images, three reference videos, three audio clips — the ceiling is generous, but character continuity past frame one still needs a written backup, not just a picture.

MiniMax and Comfy-Org built native H3 support into ComfyUI at launch — four new nodes, six workflow templates, minimum version ComfyUI 0.30.0¹ — and documentation specifies a shot list, not one descriptive sentence: on-screen action, ambient sound, and music written out separately in the same prompt block.¹

Reference media gets the same discipline: images, videos, and audio are tagged in the order cited — <Picture 1>, <Video 1>, <Audio 1> — each tag assigned one job: identity, style, motion, camera, or voice.¹ A prompt can carry up to nine reference images, three reference videos (each with an optional soundtrack), and three standalone audio clips, across three modes

— Text-to-Video, First & Last Frame, and Reference-to-Video.¹

A different shorthand circulates outside Comfy-Org's own pages: "Subject + Action + Environment + Camera + Lighting + Style + Sound + Ending."² A reasonable checklist, but not MiniMax's language — and treating it as gospel risks missing the rule that matters: action separated from ambient sound from music.

One gap matters for anyone working locally: H3-Context-IR and H3-Regenerate-2K, two hosted-service pipeline stages, aren't part of the open-weight release.⁶ A workflow built from prompt structure alone won't reproduce whatever those stages add upstream.

9 · 3 · 3

MAX REFERENCE IMAGES · VIDEOS · AUDIO CLIPS PER PROMPT

THE FORMULA

Action — what happens on screen, shot by shot.

Ambient sound — the space's noise floor, described apart from the action.

Music — its own line, never folded into the action beat.

Reference tags — cite in the order used, <Picture 1> / <Video 1> / <Audio 1>, one explicit job per tag: identity, style, motion, camera, or voice.

Caps — 9 images · 3 videos (soundtrack optional, per video) · 3 standalone audio clips.

Structure gets a prompt read correctly; it doesn't guarantee the render holds together. H3's motion is steadier than MiniMax's earlier generations, but reviewers pushing complex action still report texture instability and mid-shot facial distortion — physics multiple outlets describe as trailing the model's strongest results.³ One reviewer's rotating anatomical shot was discarded for unusable anatomy — a single account, unconfirmed at scale.⁴

The steeper problem shows up over time: across a 15-second clip, testers report clothing changing between shots and backgrounds shifting character.⁵ MiniMax's fix, Omni Reference, asks for more than a reference image — a text description of traits that need to hold, spelled out rather than left for the image to imply.⁵

Omni Reference holds a face for fifteen seconds only if the prompt says, in words, what to hold onto.

SOURCES 1 docs.comfy.org / blog.comfy.org, official H3 integration notes: ComfyUI ≥0.30.0, four new nodes, six workflow templates, timed-shot prompt structure, reference-tag jobs, generation caps and modes. 2 Third-party prompt-formula guides (fal.ai, PixelDojo, SuperMaker, atlascloud) — not affiliated with Comfy-Org or MiniMax. 3 seart.ai and curiousrefuge.com reviews of H3 motion and temporal artifacts under complex action. 4 Single reviewer account of an anatomy failure in a heart/blood-flow test clip — SINGLE-SOURCE, unconfirmed at scale. 5 deevidei.ai and Runware documentation on 15-second identity drift and the Omni Reference mode's text-description requirement. 6 Hugging Face MiniMax H3 model card — H3-Context-IR and H3-Regenerate-2K disclosed as excluded from the open-weight release, API-only.

WHEN IT BREAKS

The Model Doesn't Refuse — It Substitutes

Three open GitHub issues — a device mismatch, a corrupted latent, an unresolved VAE decode failure — mark where the local ComfyUI install still catches on its own edges.

Ask H3 for something its moderation layer won't allow, and it rarely says no outright. One account of testing gore prompts found the outcome quieter: the generation completes, reads as a success, and the disallowed element simply isn't there.¹ A different failure mode from refusal — a "successful" render still needs a second look.

Named copyrighted characters get different treatment: a prompt invoking Darth Vader hits a

hard, input-level block, rejected in ~7 seconds¹ — fast enough testers tell it apart from generation-stage substitution by latency alone. License terms behind that split: [p19](#).

Audio is native, not bolted on, but native doesn't mean guaranteed: documentation promises no accurate lip-sync, pronunciation, or speaker identity⁵ — a mismatch that reads fine in a thumbnail and shows wrong at full playback. Every dialogue clip needs checking individually.

THE SCEPTIC

One account claims H3's local, open-weight build filters less aggressively than the hosted API.³ Plausible, given how moderation is typically implemented — but MERIDIAN found no second source confirming it. Treat it as a lead, not a fact to build a workflow around.

The moderation layer isn't where most early adopters get stuck — the install is. Three open GitHub issues, filed against the nodes on [p11](#), describe distinct, unresolved failures (detailed below).² #15251 is a device-mismatch RuntimeError in the VAE decoder on 8GB cards during partial offload.² #15254 throws an AttributeError saving latents, and #15246 is a VAE decoding error that persists even with custom nodes disabled.² #15318 is a feature request: the node only anchors first/last-frame keyframes; multi-frame reference raises a ValueError.²

The most useful troubleshooting habit: run the smallest official workflow first, then change one memory flag at a time — `--lowvram`, `--reserve-vram`, `--disable-smart-memory` — rather than stacking all three.³ Out-of-VRAM doesn't always error cleanly — one documented case was a silent SIGKILL, no traceback³ — which looks like a script bug and is neither.

AMD support is a single, unresolved report: one Hugging Face thread describes a CUDA invalid-argument error loading H3's video VAE on an RX 9070 XT, no confirmed working configuration.⁴ Unproven, not broken.

SYMPTOM	LIKELY CAUSE	WHAT TO TRY
RuntimeError: device mismatch in VAE decoder	Partial CPU offload on ≤8GB cards (#15251)	Avoid partial offload; unmerged community patch to vae.py exists
AttributeError saving latents ('NestedTensor'...)	Latent-save step, root cause unresolved (#15254)	No confirmed fix — avoid saving latents mid-pipeline
VAE decoding error after sampling	Unresolved, persists with custom nodes disabled (#15246)	Isolate with the smallest official workflow; still open upstream
Job vanishes from /queue, no error shown	Silent OOM SIGKILL	Toggle one flag at a time: <code>--lowvram</code> , <code>--reserve-vram</code> , <code>--disable-smart-memory</code>

SOURCES 1 atlascloud.ai, "The Kiss Clears, Blood Doesn't" — documented moderation behaviour: silent substitution of disallowed elements and ~7-second hard input-level block on named copyrighted characters. 2 GitHub, comfyanonymous/ComfyUI issues #15251, #15254, #15246, #15318. 3 kingy.ai — memory-flag troubleshooting order, silent-SIGKILL OOM behaviour, and single-source claim re: local build filtering less than the hosted API (uncorroborated). 4 Hugging Face discussion #17 on the H3 repository — RX 9070 XT / ROCm load failure, single source, no confirmed resolution. 5 docs.comfy.org / Hugging Face MiniMax H3 model card — native audio generation, no guarantee of lip-sync, pronunciation, or speaker-identity accuracy.

THE LICENSE

Open Weights, Four Closed Countries

The ban traces to a Hollywood copyright suit still moving through a US courtroom — not to anything printed in the license itself.

Download H3's open weights and the Community License Agreement lets you run them locally free, provided your product earns under \$20 million a year.¹ Run those same weights from inside the US, EU, UK, or South Korea, and the license forbids it outright.¹ SCMP confirmed the exclusion directly against the license text — one of the better-corroborated facts in this issue.²

The hosted API stays open everywhere, excluded countries included.¹ The restriction attaches to local execution, not MiniMax's own compute — a studio in Los Angeles can still call MiniMax's servers and pay per second.

MiniMax's license text doesn't explain why these four territories. The explanation comes from reporting: Disney, Universal, and Warner Bros. Discovery are suing MiniMax's Hailuo AI product over reproducing copyrighted characters, with a motion to dismiss

reportedly denied in May 2026.³ SCMP's framing, corroborated by TechTimes, Albase, and KuCoin,⁴ is inference layered on a license clause, not MiniMax's stated rationale.

The same defensiveness shows downstream: H3's tendency to silently substitute or block IP-sensitive content rather than refuse it, covered on [p17](#). The exclusion arrived with the license on 3 August 2026,¹ four days before press — "too new to call," not settled.

The rest is conventional. Above \$20M revenue needs written authorization (api@minimax.io).¹ Commercial UIs must display "MiniMax H3." A no-distillation clause bars training a competing model on H3's outputs.¹ Governing law: Hong Kong SAR. Excluded-territory developers can apply directly for licensed access.

TERM		DETAIL
Free-use revenue ceiling	Under \$20M USD annual revenue — no fee, no authorization needed	
Above the ceiling	Requires prior written MiniMax authorization (api@minimax.io)	
Attribution	"MiniMax H3" must be displayed prominently on commercial product UI	
Excluded territories	US, EU, UK, South Korea — barred from locally running open weights; hosted API unaffected	
No-distillation	H3 outputs cannot be used to train a competing model, globally	
Governing law	Hong Kong SAR	

ON HUGGING FACE

MiniMax's own model card carries a discussion thread titled "Any License Question Ask here!" — existence and title confirmable; MERIDIAN could not retrieve the replies inside it.⁵ Separate reporting suggests the territory exclusion, not H3's capabilities, draws the most attention in the days after launch — live and ongoing, not a settled footnote.⁶ Reddit reportedly picked up the story too, but MERIDIAN could not verify specific threads or quotes there, and found no dedicated MiniMax subreddit.

SOURCES 1 MiniMax H3 Community License Agreement, terms per GitHub README, github.com/MiniMax-AI/MiniMax-H3, corroborated by multiple secondary outlets, Aug 2026. 2 South China Morning Post, scmp.com/tech/tech-trends/article/3362951, confirming territory exclusion directly against license text. 3 SCMP reporting on Disney/Universal/Warner Bros. Discovery litigation against MiniMax/Hailuo AI and the May 2026 motion-to-dismiss denial. 4 TechTimes, Albase.com, and KuCoin News Flash, corroborating the litigation-to-exclusion link. 5 Hugging Face, MiniMaxAI/MiniMax-H3, Discussions #12, "Any License Question Ask here!" — thread title confirmed, contents not independently accessed. 6 TechTimes and Albase.com framing of the exclusion as the dominant early community topic.

THE ECONOMICS

What A Second Of Video Actually Costs

Renting a cloud GPU can undercut MiniMax's own API by roughly five to one — but only if the card never sits idle, which nobody's does.

MiniMax's official API bills 2K output at \$0.13 per second — a fifteen-second clip runs \$1.95.¹ The 768p tier is cheaper but contested: some sources put it at \$0.08 per second, others at \$0.09, and MERIDIAN could not resolve which is current.¹ Some platforms also describe 768p as closed-beta, requiring a sales conversation rather than self-serve access.¹

Resellers quote their own rates, and they scatter widely — none are MiniMax's own pricing. Segmind lists roughly \$0.1625/s at 2K; Atlas Cloud charges \$0.14/s at 2K and \$0.10/s at 768p; AIHubMix runs cheaper at \$0.124/s and \$0.077/s; EmpirioLabs runs far more expensive, at \$0.26/s and \$0.18/s.² Reference media carries its own line item, per a single unverified source: the first five images are free, each additional one costs roughly \$0.04, and audio references reportedly carry no fee.³

\$0.13/s

MINIMAX'S OWN 2K API RATE — 768P RUNS \$0.08–\$0.09/s, SOURCES DISAGREE ON WHICH

SOURCE	2K RATE	768P RATE
MiniMax (official)	\$0.13/s	\$0.08–\$0.09/s (disputed)
Segmind	~\$0.1625/s	—
Atlas Cloud	\$0.14/s	\$0.10/s
AIHubMix	~\$0.124/s	~\$0.077/s
EmpirioLabs	\$0.26/s	\$0.18/s

Put those numbers against hardware — nobody has published this comparison, so it's MERIDIAN's own construction from the rates above and [p13](#)'s buyer's-ladder prices, illustrative not authoritative. Fifty ten-second, 768p-class clips cost ~\$42.50 at the API's disputed mid-rate. The only benchmarked local speed found: an RTX 5090 on H3's pruned NVFP4 build, 175 seconds per clip — other reports give different, non-comparable times.⁴

Renting an H100 to match that pace is cheap: RunPod's on-demand H100 runs \$2.89–\$4.55/hr.⁵

Fifty clips, ~2.4 rental hours, is \$7–\$11 against the API's \$42.50 — roughly a fifth the cost, assuming the card stays busy the whole billed hour, optimistic without a continuous queue.

Buying is the slow-money option. A street RTX 5090 runs ~\$4,300⁶ — itself inflated by the 2026 memory shortage, no relief expected before 2028.⁶ Against the API, that hardware pays back after ~5,000 clips; against renting, 19,000–31,000, since rental is already cheap per clip. Neither figure counts electricity or depreciation.

Owning the hardware only wins the economics at a volume most individual creators will never hit.

SOURCES 1 MiniMax API pricing, search-aggregated (not independently fetched from minimax.io's pricing page); 768p rate reported as either \$0.08 or \$0.09 per second, unresolved. 2 Segmind, Atlas Cloud, AIHubMix, and EmpirioLabs published rate cards, search-aggregated, Aug 2026 — reseller rates, not MiniMax's own. 3 Reference-image/audio add-on pricing, single-source (atlascloud.ai), unverified. 4 Single-source RTX 5090 NVFP4/pruned-build benchmark (ai-muninn.com), 175 seconds per ten-second 864x480 clip; other benchmarks report different, non-comparable generation times at different settings. 5 RunPod Secure Cloud on-demand and serverless GPU rental pricing, runpod.io. 6 RTX 5090 street pricing (~\$4,300) per Amazon/WCCFTech listings cited in the Gear section, [p13](#), 2026 DRAM/GDDR7 shortage context per shattered.io, tech-insider.org, and TechRadar.

THE PROGRAMME

Your First Ninety Days

The workflow templates that shipped with the model already split into two systems living under one name — a free local path and a metered one.

LOCAL · 768P

API · 2K

MiniMax H3 landed inside ComfyUI as core code, not a plugin. Comfy-Org/ComfyUI #15224 — kijai, merged 3 August 2026 — added four nodes to comfy_extras: EmptyMiniMaxH3LatentAV, MiniMaxH3ImageToVideo, MiniMaxH3ReferenceToVideo, MiniMaxH3SigmaShift.¹ Any install method picks them up on the next update; the real requirement is ComfyUI 0.30.0 or newer.¹ Full node detail on [p11](#).

WEEK ONE — GET IT RUNNING

DAYS 1–7

1–2

Update ComfyUI, confirm the four H3 nodes loaded

SHIPS IN CORE ≥0.30.0 — NO SEPARATE INSTALLER

3–4

Load the official day-zero workflow templates

SIX TEMPLATES: THREE LOCAL, THREE API-KEY²

5–7

Learn local-768p vs API-2K, then run the smallest workflow first

LOCAL CAPS AT 768×1344, FULL 2K IS API-ONLY³; OOM CAN SIGKILL SILENTLY, NO TRACEBACK⁴

Local H3 — FL2VA and Ref2VA at Comfy-Org/MiniMax-H3, built on the 33.1B-param H3-Base — tops out at 768×1344. Full 2K runs through partner API nodes, specifically H3-Regenerate-2K, an upsampler, not a native 2K generator; neither it nor H3-Context-IR was released as open weights.³

Native ComfyUI support means two systems under one name — a free local path capped at 768p, and a metered one that isn't.

FROM THE FORUMS

Kingy.ai: test the smallest workflow first, then change one memory flag at a time — **--lowvram**, **--reserve-vram**, **--disable-smart-memory**. A device-mismatch VAE bug, filed as Issue #15251, was still unresolved four days after release.⁴

WEEKS TWO TO FOUR — TUNE THE RIG

DAYS 8–28

8–14

Move to a GGUF or quantized checkpoint sized to your card

Q5_1 FOR 24–32GB, Q4_0 FOR 16GB⁵

15–21

Try the Turbo LoRA once the base workflow feels familiar

TWENTY STEPS DOWN TO FOUR — BUT NEEDS A DUAL-CLOCK SAMPLER⁶

22–28

Start using Reference-to-Video for character consistency

MINIMAXH3REFERENCETOVIDEO — UP TO NINE REFERENCE IMAGES²

H3's transformer runs 2688-wide tensors not divisible by 256, ruling out K-quant formats — the community settled on classic Q4_0/Q5_1.⁵ On an RTX 5090, Q5_1 held ~22GB resident, ~ten minutes for a 243-frame clip at 20 steps; 16GB cards on Q4_0 ran slower.⁵ Match the quant to the card on [p13](#).

42.5GB

DISK SPACE FOR THE SMALLEST OFFICIAL LOCAL CHECKPOINT, THE PRUNED INT8 BUILD

larryvrh's MiniMax-H3-Turbo-Lora (5 Aug 2026, packaged by drbaph) cuts sampling from ~20 steps to 4 — near a five-times speedup.⁶ Loading it is trivial; running it well is not. Stock sampling produces severe noise at four steps — the dual-clock Euler sampler (video shift 12, audio shift 3) is required. Four-step results still read softer than six-to-eight.⁶

SOURCES 1 Comfy-Org/ComfyUI PR #15224, "feat: Support MiniMax-H3 (CORE-375)," merged 3 Aug 2026, author kijai; ComfyUI ≥0.30.0 version requirement, cross-checked across setup guides. 2 Comfy-Org/workflow_templates repository, six MiniMax H3 templates (three local, three API-key), Aug 2026; docs.comfy.org/blog.comfy.org official prompt guidance, reference-media caps and generation modes. 3 ComfyUI partner API nodes and H3-Regenerate-2K node behaviour, search-synthesis cross-checked; HuggingFace MiniMax H3 model card, confirming H3-Context-IR and H3-Regenerate-2K withheld from the open-source release. 4 Kingy.ai, ComfyUI MiniMax H3 troubleshooting notes, Aug 2026; Comfy-Org/ComfyUI Issue #15251, device-mismatch VAE decoder bug. 5 Community GGUF/NVFP4 repositories (molbal, joeygambino, Abiray, realrebelai, leejet, vantagewithai, lilcheaty) and RTX 5090 Q5_1 benchmark, Aug 2026. 6 larryvrh/MiniMax-H3-Turbo-Lora, 5 Aug 2026 preview; drbaph/MiniMax-H3-Turbo-Lora-ComfyUI packaging notes.

MONTHS TWO AND THREE

Sound, Frames, and the Bill

By day sixty the decision stops being which node to click and becomes whether to keep paying Nvidia, start paying MiniMax, or start paying a cloud host by the hour.

H3 does not bolt audio onto video after the fact. One packed-token diffusion transformer denoises video latents (24 channels, 16x spatial compression) and stereo audio latents (32 channels, 40Hz) in one sequence, conditioned on Qwen3-VL-32B.¹ The streams run different noise schedules, reconciled by MiniMaxH3SigmaShift — sync falls out of generating both from the same steps.¹ Native sound doesn't guarantee lip sync or stable speaker identity — audio remains the weaker stream.²

MONTH TWO — BEYOND SILENT T2V DAYS 29–59

29–38 **Generate audio and video jointly, on purpose**
ONE PACKED STREAM, NOT A BOLTED-ON VOICE TRACK¹

39–59 **Run First/Last-Frame workflows, then learn the prompt-structure conventions**
MINIMAXH3IMAGETOVIDEO ANCHORS START/END FRAMES; TIMED SHOTS TAGGED <PICTURE 1>/<VIDEO 1>/<AUDIO 1>³

Official guidance structures a prompt as a timed shot: action, sound, and music as separate clauses, reference media tagged in citation order — Picture 1, Video 1, Audio 1 — each assigned a job.³ A third-party formula (fal.ai, PixelDojo, SuperMaker, Atlas Cloud) isn't official but has spread widely.³ Failure-mode rundown on [p17](#).

THE SCEPTIC

Atlas Cloud describes H3's moderation as substitution, not refusal — gore generates with blood quietly dropped, while copyrighted characters trigger a hard block in ~7 seconds.⁴ A separate, uncorroborated claim holds the local path is more permissive than the API⁴ — don't build a workflow around it.

MONTH THREE — DECIDE WHO PAYS DAYS 60–90

60–85 **Price your own card, then price cloud rental by the hour**
ENTRY ~\$480 TO FLAGSHIP ~\$4,329 STREET; RUNPOD A100 \$1.39/HR TO H100 \$2.89/HR, SPOT FROM \$0.34/HR⁵

86–90 **Weigh API 2K against local 768p plus regeneration**
2K \$0.13/S VS 768P \$0.08–0.09/S PLUS \$0.05/S REGEN⁶

The DRAM/GDDR7 shortage has pushed GPU prices 60-130%+ above cost through 2026, relief not expected before ~2028.⁶ A flagship RTX 5090, discontinued at MSRP, resells ~\$4,329 — nine times an entry \$480 RTX 4060 Ti.⁵

\$0.05/s

COST TO UPSCALE A LOCAL 768P CLIP TO FULL 2K VIA THE API REGENERATION ENDPOINT

MiniMax charges \$0.13/s for 2K and \$0.08-0.09 at 768p.⁶ A locally-generated 768p clip upgrades to full 2K via Regeneration for \$0.05/s, without the \$0.13/s full-2K rate.⁶ That bridges local output ([p13](#)) and the 2K MiniMax advertises; break-even math on [p19](#).

By day ninety the question isn't which node to click. It's whether you're still paying Nvidia, or paying MiniMax by the second.

SOURCES 1 Comfy-Org/ComfyUI PR #15224 description, joint audio-video mechanic, MiniMaxH3SigmaShift, Qwen3-VL-32B conditioning, merged 3 Aug 2026. 2 drbaph/MiniMax-H3-Turbo-Lora-ComfyUI release notes on audio vs video quality; deevid.ai and Runware documentation on lip-sync/identity limitations. 3 docs.comfy.org/blog.comfy.org official prompt-structure guidance; MiniMaxH3ImageToVideo first/last-frame conditioning; fal.ai/PixelDojo/SuperMaker/Atlas Cloud third-party prompt formula. 4 Atlas Cloud, "The Kiss Clears, Blood Doesn't," content-moderation report, Aug 2026; kingy.ai claim on local-vs-API permissiveness, single-source, uncorroborated. 5 RunPod Secure Cloud/Serverless pricing and Vast.ai marketplace listings, Aug 2026; GPU street prices (Amazon, WCCFTech and others), Aug 2026. 6 MiniMax official API pricing and Atlas Cloud reseller pricing, Aug 2026; shattered.io/tech-insider.org/TechRadar reporting on the DRAM/GDDR7 shortage, with Intel/Silicon Motion relief estimates near 2028.

CROSS-EXAMINATION

The #1 Ranking Runs on Parts You Can't Download

MiniMax's open licence also carves out the U.S., the E.U., the U.K. and South Korea — and the timing tracks an active Hollywood lawsuit almost too neatly.

Artificial Analysis's Video Arena — blind, pairwise human votes scored with a Bradley-Terry Elo — placed MiniMax H3 first in Video Editing, Elo roughly 1,130, ahead of every other model in the pool. It placed second in Text-to-Video (1,239, just behind Gemini "Omni Flash" at 1,244) and third in Image-to-Video (1,187, behind ByteDance's Seedance 2.0 and Gemini again).¹ Those are real, hard-won numbers, not a cherry-picked demo reel.

Here is the part the launch headlines skipped. The configuration that earned the Video Editing crown ran MiniMax's hosted pipeline at full 2K, built on two components never released as open weights: H3-Context-IR and H3-Regenerate-2K, both API-only.² The checkpoints anyone can download — H3-Base — cap out at 768p and ship without either piece.³ Install H3 locally tonight and you are running something adjacent to the model that won, missing the exact two parts that pushed it to #1.

Even the win has loose threads. The Video Editing sample size is unsettled — 8,211 or 5,043 evaluation pairs, depending on the source, a gap MERIDIAN could not close. One account goes further,

describing the Text-to-Video and Image-to-Video placements as resting on "internal evaluation methodology" rather than full public vote history⁴ — plausible, single-source.

Then there is what "open" is actually promising. MiniMax released model weights, not training data or training code, and the 2K production pipeline stays entirely server-side.³ Community comparisons already draw that line explicitly, framing H3 against Alibaba's Wan 2.5/2.7 as "open weights" versus "fully open source" — Wan ships under Apache 2.0, H3 does not.⁵ The licence narrows "open" further still: the U.S., the E.U., the U.K. and South Korea are excluded from using, modifying or deploying anything generated locally by the open weights.⁶ The reason, per the South China Morning Post, ties directly to active U.S. litigation: Disney, Universal and Warner Bros. Discovery are suing MiniMax's Hailuo AI product, alleging its models reproduce copyrighted characters without permission, with a motion to dismiss reportedly denied in May 2026.⁶ The excluded territories line up, almost exactly, with where the plaintiffs filed.

THE SCEPTIC

One assessment of the arena result puts it plainly: "a single third-party arena run, not a peer-reviewed suite... treat rankings as a prior, not a verdict."⁷ That's fair — blind pairwise voting still beats most alternative video-model benchmarks on offer. But "a prior, not a verdict" is exactly the right frame, and it is not the frame most launch coverage used. The #1 badge is doing more work in the press than the underlying comparison supports — for a configuration nobody outside MiniMax's own servers can run.

2 of 2

API-ONLY COMPONENTS (H3-CONTEXT-IR, H3-REGENERATE-2K) BEHIND THE #1 RESULT — NEITHER SHIPPED AS OPEN WEIGHTS

The reader installing H3 in ComfyUI tonight is not running the model that won.

SOURCES 1 Artificial Analysis (@ArtificialAnlys), X post, x.com/ArtificialAnlys/status/2083042088338538594. 2 XenoSpectrum, "MiniMax H3: open weights, video editing," xenospectrum.com/en/minimax-h3-open-weights-video-editing. 3 MiniMax-AI/MiniMax-H3, GitHub README, github.com/MiniMax-AI/MiniMax-H3. 4 kingy.ai, analysis of Artificial Analysis arena methodology, single-source. 5 Community comparison threads framing H3 against Wan 2.5/2.7 as "open weights vs open source," search-aggregated. 6 South China Morning Post, scmp.com/tech/tech-trends/article/3362951. 7 kingy.ai, commentary on benchmark verdict framing, single-source.

THE RECEIPTS

The Bug Tracker Tells a Different Story

And most of the "community is thrilled" narrative behind H3 travelled through search-engine synthesis, not readable threads — a caution this section owes itself too.

Issue #15251, filed within days of launch, is a plain `RuntimeError`: a device mismatch in the H3 VAE decoder, triggered by partial CPU offload on an 8GB card.¹ A patched `vae.py` has been posted but isn't confirmed merged. Three more issues followed the same week: an `AttributeError` saving latents, a VAE decoding error persisting with custom nodes disabled, and a feature request — the model anchors only first/last-frame references, anything else raises a `ValueError`.¹

Out-of-memory failures are worse than an error message. By one single-sourced account, they sometimes present as a silent SIGKILL: no traceback, no entry in ComfyUI's history, the job simply vanishes from the queue.² That's one report MERIDIAN can't confirm at scale, but the workaround — test the smallest workflow first, change one memory flag at a time — reads like advice for a model still finding its footing on consumer hardware.

On motion, reviewers are more consistent: under large-scale action H3 reportedly shows texture instability, facial distortion and visual noise that Seedance 2.0 avoids more reliably, and its physics performance trails the field even where H3 otherwise competes.³ Quantization tells a similar story: GGUF, pruned INT4/INT8 and NVFP4 builds for Blackwell cards all appeared the day after release, from independent Hugging Face contributors, not MiniMax.⁴ One RTX 5090 benchmark found real differences: NVFP4's rel2 error ran about 0.094 against W4A8's 0.073, and INT8 builds showed less mid-motion artifacting and better shape retention through fast pans than NVFP4.⁵ No source compares any local quantized build against the hosted full-precision API — a gap that matters given H3-Context-IR and H3-Regenerate-2K never left the API tier.

THE SCEPTIC

Strip away the audio-video framing and ask what's actually unprecedented. Denoising video and audio together in one packed-token stream is real and, per the engineers who merged H3 into ComfyUI, new to this class of open release.⁶ But the motion and physics complaints above are described as reduced relative to the prior MiniMax generation, not eliminated.³ Reduced-not-eliminated is improvement, not reinvention. [p5](#)'s account of the company's lineage is the place to weigh whether H3 is a genuine leap or Hailuo's next number wearing an open-weight coat. The honest answer is some of both — and the launch moment blurs which half is which.

THE SCEPTIC

That caution applies to this cross-examination's own supply chain, and to most of this issue's community reporting. Several outlets behind the "day-zero love" narrative could not be fetched directly: two Hacker News threads confirmed to exist but unreadable, organic Reddit discussion described only through a secondary Patreon reference, no official MiniMax or Hailuo Discord found.⁷ What's solid: the GitHub merge itself, the Civitai workflow uploads within a day, the independent quantization builds — real tools shipped fast. What's thinner is any claim about what a broad community "feels"; most of it travels through search engines summarizing other blog posts. Apply the same scrutiny to the hype that this page applies to the benchmark.

None of this makes H3 a bad model. The joint audio-video pipeline functions as advertised, ComfyUI support landed day-zero [p11](#), and a heavily quantized build runs on cards costing a fraction of full precision [p13](#). It makes H3 a model whose loudest claim needs

an asterisk, and its licence [p19](#) a document written as much for a courtroom as a changelog. Judge it on what MERIDIAN could verify [p9](#), not the banner headline.

SOURCES 1 GitHub, Comfy-Org/ComfyUI issues #15251, #15254, #15246, #15318. 2 [kingy.ai](#), single-source account of OOM/silent-kill behaviour. 3 [seart.ai](#) and [curiousrefuge.com](#) hands-on comparison notes, cross-referenced motion/physics findings. 4 Hugging Face community quantization releases ([joeygambino](#), [molbal](#), [leejet](#), [Abiray](#), [coolthor](#), [ilcheaty](#)), MiniMaxAI/MiniMax-H3 repository. 5 [ai-muninn.com](#), RTX 5090 quantization benchmark, single-source. 6 GitHub, Comfy-Org/ComfyUI PR #15224 merge description. 7 Method note: Hacker News [news.ycombinator.com/item?id=49130750](#) and [id=49155629](#) (existence confirmed, content unreadable); Patreon "LumosAI" (secondary, unconfirmed); Hugging Face discussions tab (fetch blocked).

UNDERCURRENT

The Kill Line

A Shanghai video model open-sourced its weights on 3 August; within days a Hong Kong-listed stock had moved, two non-Nvidia chipmakers had shipped day-one support, and China's financial press had reached for its favorite headline again.

On 3 August, MiniMax's Hong Kong-listed shares (0100.HK) moved hard enough to become a headline of their own — estimates range from just over 10 percent intraday to more than 20 percent, with one report putting the close near HK\$249.40. Outlets converge on the cause: the stock's Stock Connect inclusion landed the same day as H3's open-weight release, and financial media treated the two as one story.¹ One widely used headline needed no translation: H3接棒DeepSeek — H3 takes DeepSeek's baton.¹

Sina Finance ran the fuller version: MiniMax"入通"即大涨，H3接棒DeepSeek，中国开源AI再划一道"斩杀线" — MiniMax surges on Stock Connect entry, H3 takes DeepSeek's baton, China's open-source AI draws another "kill line."¹ Kill line — 斩杀线 — is trading-desk shorthand from the DeepSeek moment eighteen months earlier: the threshold past which a closed competitor's advantage stops being worth paying for. Applying it here to a video model rather than a chatbot is most of the argument. 36Kr made the longer case, headlined 天下苦闭源模型久矣 — the world is weary of closed models — arguing MiniMax wants the role DeepSeek claimed in text.²

Worth being precise about what repeated. DeepSeek undercut Western pricing on language and reasoning; H3 applies the identical mechanics — open weights, an aggressive price cut, same-week developer access — to video, where MiniMax was already racing ByteDance's Seedance on capability

and cost. Chinese coverage made the comparison blunt: one Tencent News/Leiphone headline called H3 "比Seedance便宜七成的模型" — a model 70 percent cheaper than Seedance — and asked whether the entire industry was about to change shape.³

The subplot underneath the stock price is chips. ITHome's day-zero coverage tracked which hardware could run H3 the moment weights landed: Moore Threads' MTT S5000 and Intel shipped adaptation support immediately, with Huawei's Ascend line implicated in the broader domestic-silicon framing even where not named.⁴ The subtext: proof that China's frontier AI does not depend on chips it cannot buy.

Zhihu answered fastest, and from the most angles. Within days the question thread on H3 had drawn seven-plus long-form columns — a technical teardown, a deployment walkthrough, and pricing pieces arguing the multimodal landscape now has to shift.⁵ The column that actually spread was narrower: a report that an INT4-quantized build of H3 would run on a consumer RTX 3060 with 12GB of memory, hardware a hobbyist already owns.⁵

Bilibili moved on the same clock, demonstration over analysis: eight-plus videos within one to four days — deployment tests, low-VRAM guides, ComfyUI builds, a direct comparison against Seedance across six scenarios.⁶ Not speculative enthusiasm — people running the actual weights, on consumer cards, the same week the weights existed.

12GB

VRAM ZHIHU REPORTED WAS ENOUGH TO RUN AN INT4-QUANTIZED H3 ON A CONSUMER RTX 3060⁵

An open-weight flagship video model that runs on non-Nvidia silicon from day one is a small, specific proof against the idea that China's frontier AI still depends on chips it cannot buy.

SOURCES 1 Sina Finance, "MiniMax"入通"即大涨，H3接棒DeepSeek，中国开源AI再划一道"斩杀线"; market figures also per Eastmoney/caifuhao and Investing.com China, Aug 2026. 2 36Kr, "天下苦闭源模型久矣，MiniMax要做多模态领域的Deepseek," 36kr.com/p/3923196846894725. 3 Tencent News/Leiphone, "比Seedance便宜七成的模型，宣布开源：视频生成行业变天了？" 4 ITHome, day-zero hardware-adaptation coverage: Moore Threads MTT S5000, Intel; Huawei Ascend/AMD implied, Aug 2026. 5 Zhihu, multiple columns incl. "【AIGC】MiniMax H3 解析（初版），" "MiniMax H3 开源了：全模态模型架构、下载、部署与文生视频实战，" and "MiniMax H3 开源即顶流：INT4 跑，3060 12G显存可玩." 6 Bilibili, tutorial/demo cluster incl. "MiniMax H3六大场景实测对比Seedance," BV1ToMD6IE3z.

UNDERCURRENT

The Corner That Stayed Quiet

Zhihu produced seven-plus columns within days and Bilibili eight tutorials; Xiaohongshu, by MERIDIAN's search, produced none — a gap that says as much about H3's audience as the noise does.

V2EX is smaller and less inclined to declare victories — a handful of threads, specific rather than promotional, including one post whose title undersells its own point.¹ A second thread ran an organic three-way benchmark of Seedance 2.5, H3, and Kling 3.0 — the kind of comparison that happens when users treat a release as one option, not an event.²

Provisioning tells its own sovereignty story: alongside the Hugging Face release, MiniMax mirrored H3's weights on ModelScope, the Alibaba-backed hub for China's domestic network³ — a small choice, consistent with a launch built to need no routing around.

MERIDIAN's search turned up no distinct Xiaohongshu presence — no lifestyle try-ons, no consumer-creator wave, none of the amplification that turns a release into a cultural moment.⁶ Likeliest

explanation: H3 launched API-first, for developers and studios, Hailuo a secondary front.

One thread runs against the triumphant framing: woshipm.com, a product-manager community, used H3 to turn a poster into a 15-second ad at a third of standard price, while flagging the open-source promise hadn't been delivered at time of writing.⁴ Accurate: MiniMax announced API-only 31 July, didn't publish weights until 3 August — a three-day gap between pledge and artifact.⁵

Put together: a financial press with a name for this pattern, a hardware ecosystem proving same-day compatibility, a technical community moving teardown-to-deployment-to-democratization in a week. The DeepSeek playbook, replayed in video, with a Hong Kong ticker instead of a repository — but a story days old, too young to say whether the enthusiasm converts or the stock move holds. MiniMax's IPO history: [p5](#).

V2EX · DEVELOPER FORUM

"折腾了一晚 MiniMax H3，感觉真正有意思的并不是 2K" — "Spent all night messing with MiniMax H3, the really interesting part isn't the 2K." MERIDIAN could not retrieve the thread body — V2EX blocked automated access — so what survives is the claim itself: the headline resolution spec wasn't what a working developer ended up caring about.

THE SCEPTIC

Two figures are single-source: the exact size of the stock move, and H3's price relative to Seedance beyond what Chinese headlines asserted. Hobbyist enthusiasm isn't, by itself, evidence of production adoption. And a launch this recent hasn't had a bad week yet.

None of that undercuts what happened this week: a company that priced its Hong Kong IPO seven months ago gave away its newest flagship free, and within days traders, chipmakers, and hobbyists running INT4 checkpoints all treated it as

confirmation — the kill line, once drawn, gets redrawn. The English-language builder read runs at [p15](#); this one belongs to Zhihu, Bilibili, V2EX, and a ticker in Hong Kong.

SOURCES 1 V2EX, "折腾了一晚 MiniMax H3，感觉真正有意思的并不是 2K," v2ex.com/t/1231362. 2 V2EX, comparison thread (Seedance 2.5 / H3 / Kling 3.0), v2ex.com/t/1232329. 3 ModelScope, MiniMax/MiniMax-H3 mirror, modelscope.cn/models/MiniMax/MiniMax-H3. 4 woshipm.com, "MiniMax H3上手：用1/3价格把海报变成15秒动态广告，但开源还没兑现," Aug 2026. 5 Reuters (Eduardo Baptista byline) and South China Morning Post, H3 announcement timeline, 31 Jul–3 Aug 2026. 6 MERIDIAN platform search, Aug 2026 — methodology note on Xiaohongshu absence.

THE LAST WORD

A reader started using MiniMax H3 inside ComfyUI four days after it shipped and asked for the full treatment. Eleven sections and one Hong Kong stock chart later, the honest summary is this: the model is genuinely capable, and the loudest claim made about it — that it is the first open model to top a video-generation leaderboard — was won by a configuration nobody outside MiniMax's own servers can run. Both things are true at once. So is the license clause locking three of the world's largest markets out of the open weights, tied to a live Hollywood copyright suit most launch coverage left out entirely. This issue has tried to hold all of it in view rather than print the flattering half.

SOURCES, IN BRIEF — github.com/MiniMax-AI/MiniMax-H3 and github.com/Comfy-Org/ComfyUI (model card, PR #15224, issues #15251/#15254/#15246/#15318); huggingface.co/MiniMaxAI/MiniMax-H3 and huggingface.co/Comfy-Org/MiniMax-H3 (weights, license); docs.comfy.org and blog.comfy.org (ComfyUI integration); Reuters (Eduardo Baptista) and scmp.com (launch timeline, license litigation context); artificialanalysis.ai (Video Arena rankings); xenospectrum.com and kingy.ai (benchmark-configuration reporting); civitai.com (community workflows); runpod.io and vast.ai (cloud pricing); zhihu.com, bilibili.com, v2ex.com, 36kr.com, sina.com.cn and modelscope.cn (the Chinese-platform Undercurrent). Full per-page footnotes throughout. Figures and prices as researched and dated Aug 2026; single-source claims flagged in place; sources that disagreed are shown disagreeing, not silently resolved; nothing in this issue was invented.